

Input output

Polling vs interrupt, Interrupt controllers, interrupt descriptor table, interrupt handlers, Direct memory access, hard disks

Agenda

Agenda

- Overview of IO devices (OSTEP Ch. 36): Polling, Interrupts, Direct memory access

Agenda

- Overview of IO devices (OSTEP Ch. 36): Polling, Interrupts, Direct memory access
- Interrupt handling (xv6 Ch. 3): interrupt controllers, interrupt descriptor table

Agenda

- Overview of IO devices (OSTEP Ch. 36): Polling, Interrupts, Direct memory access
- Interrupt handling (xv6 Ch. 3): interrupt controllers, interrupt descriptor table
- Hard disk drives (OSTEP Ch. 37): disk geometry, disk scheduling

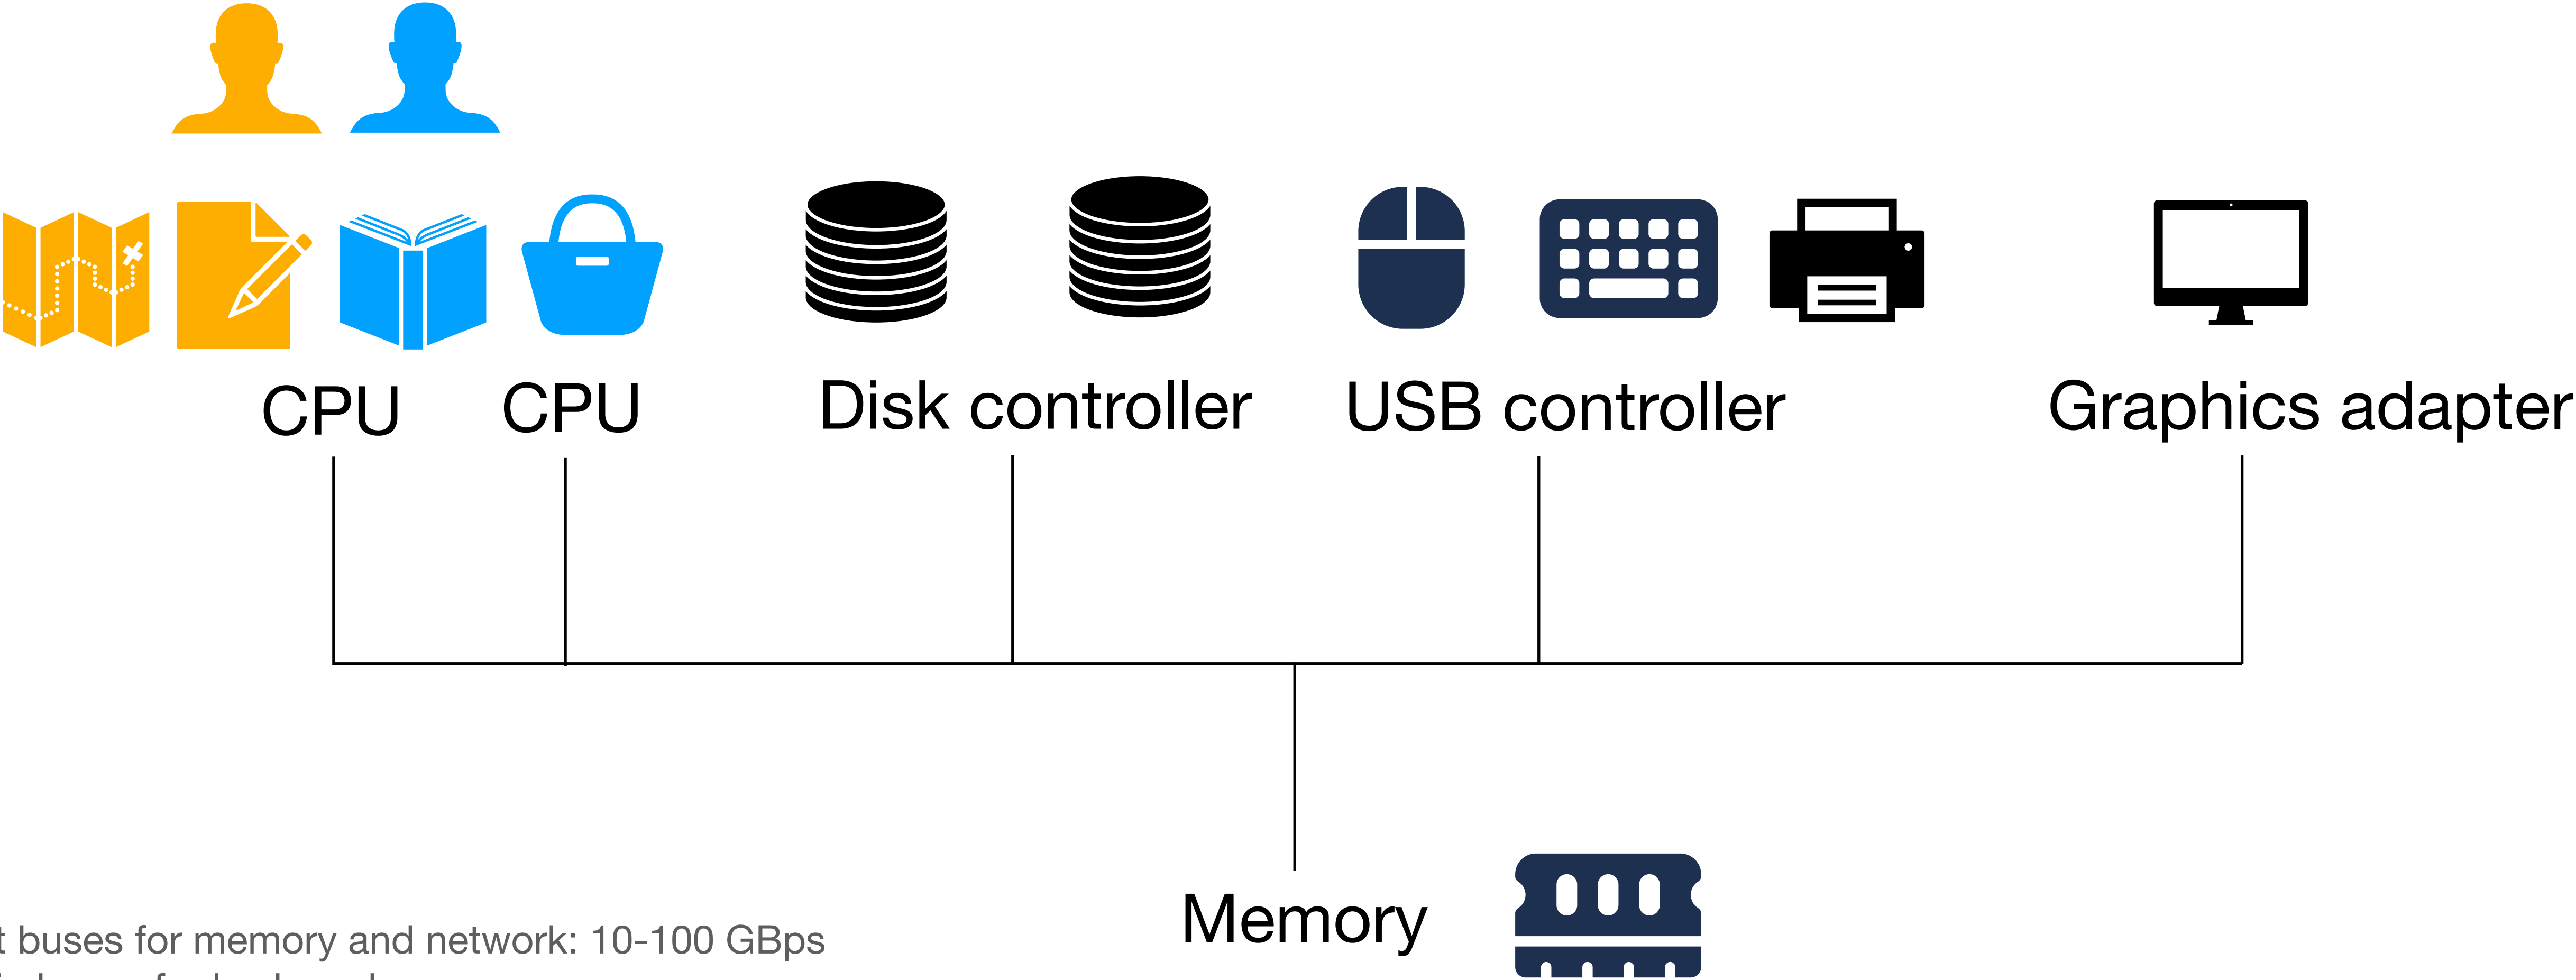
Agenda

- Overview of IO devices (OSTEP Ch. 36): Polling, Interrupts, Direct memory access
- Interrupt handling (xv6 Ch. 3): interrupt controllers, interrupt descriptor table
- Hard disk drives (OSTEP Ch. 37): disk geometry, disk scheduling
- Redundant Array of Inexpensive Disks (OSTEP Ch. 38): improve capacity, throughput, fault tolerance

Overview of IO devices

OSTEP Ch. 36

Computer organization



Fat buses for memory and network: 10-100 GBps
Thin buses for keyboard, mouse

Fitting into the OS

Hide device specific details in device driver

Fitting into the OS

Hide device specific details in device driver

- Abstraction allows OS and applications to stay device-neutral

Fitting into the OS

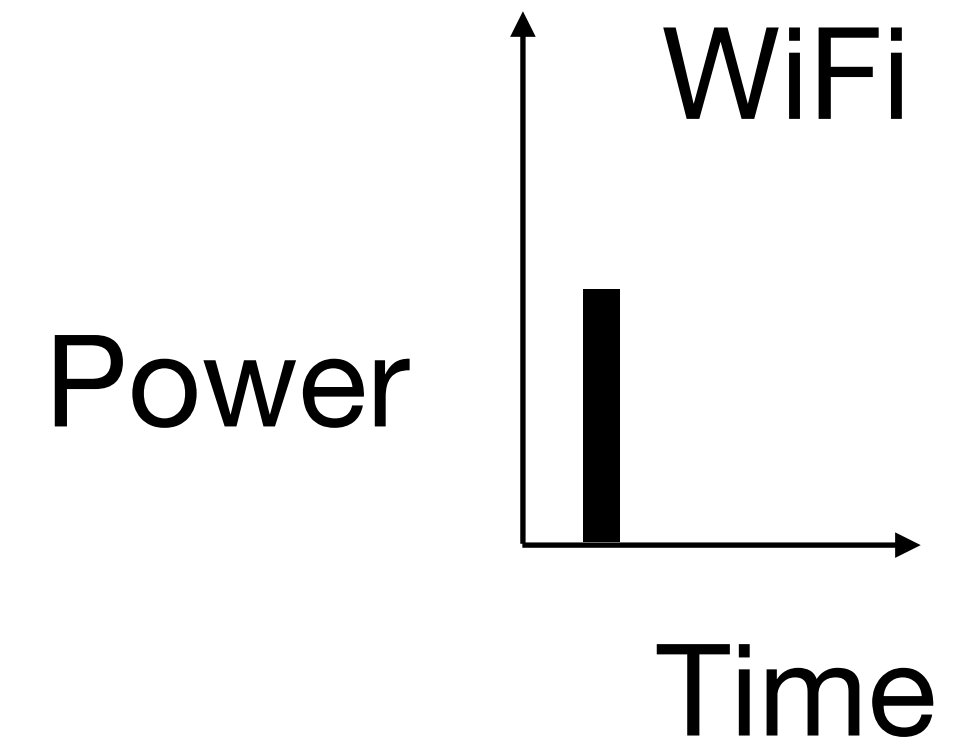
Hide device specific details in device driver

- Abstraction allows OS and applications to stay device-neutral
- Abstraction can hurt.

Fitting into the OS

Hide device specific details in device driver

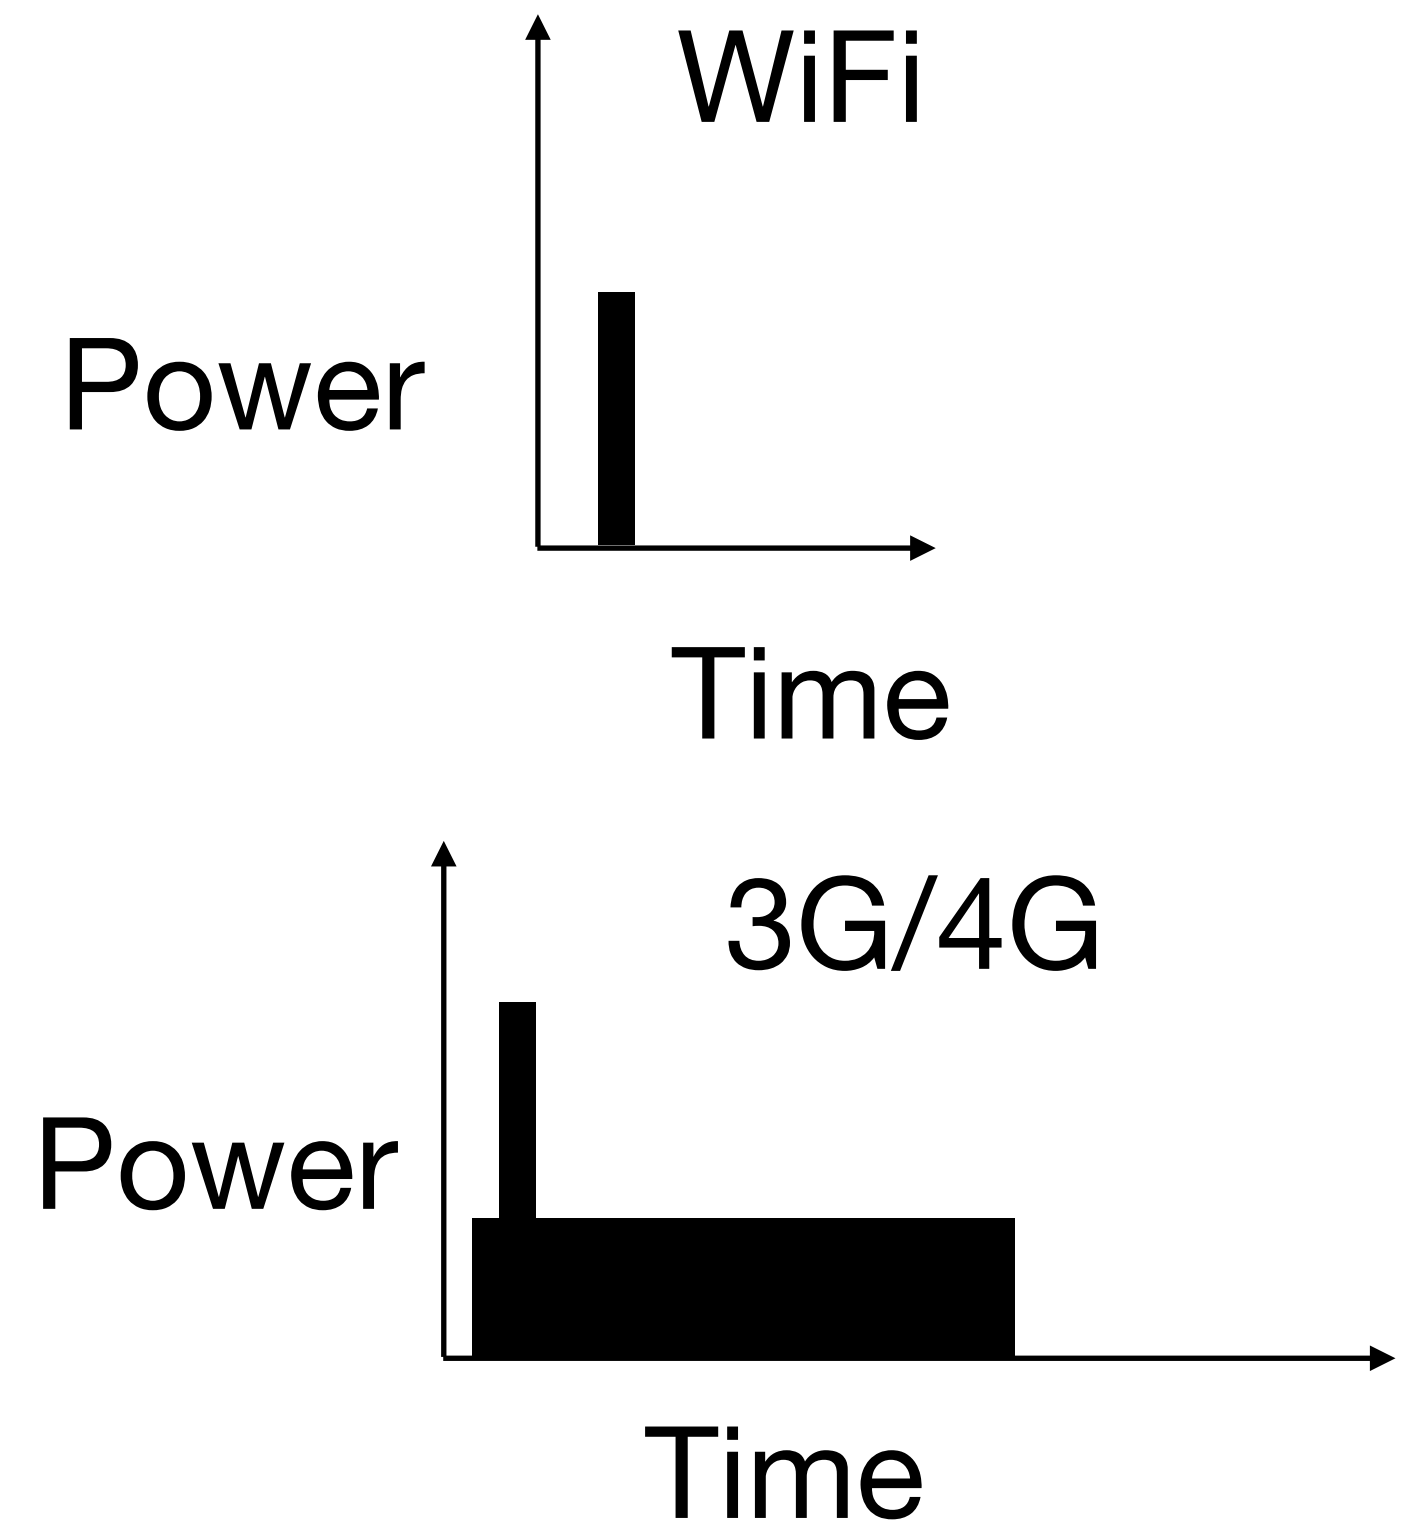
- Abstraction allows OS and applications to stay device-neutral
- Abstraction can hurt.



Fitting into the OS

Hide device specific details in device driver

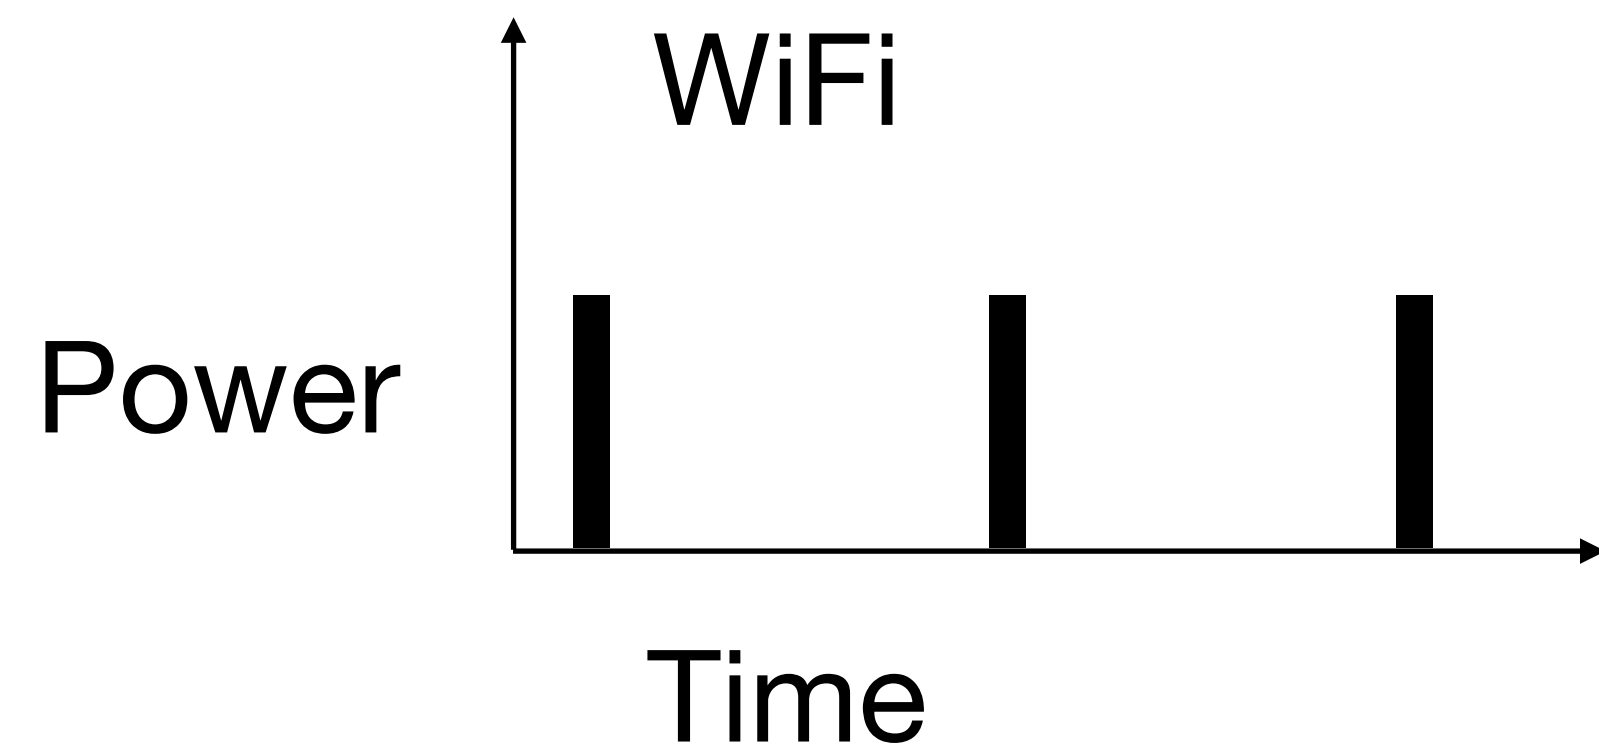
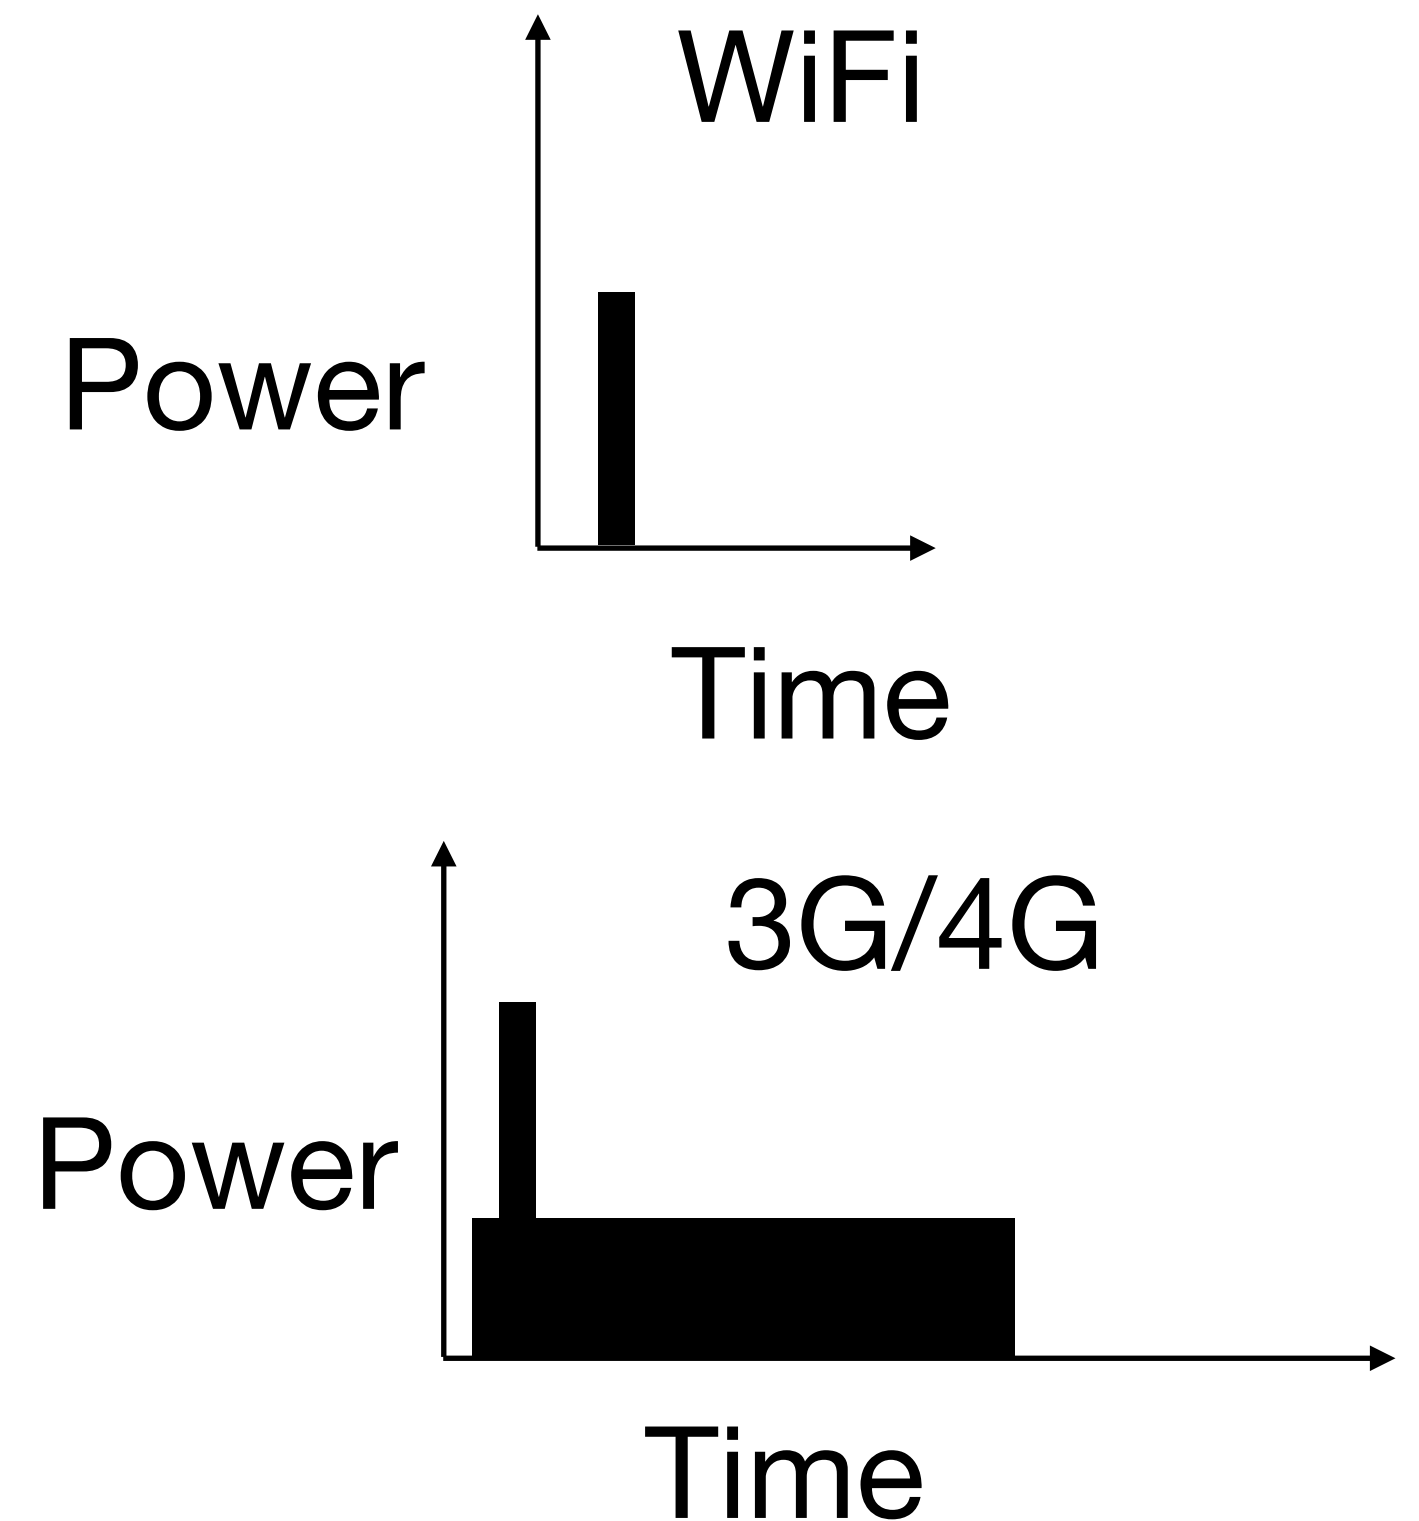
- Abstraction allows OS and applications to stay device-neutral
- Abstraction can hurt.



Fitting into the OS

Hide device specific details in device driver

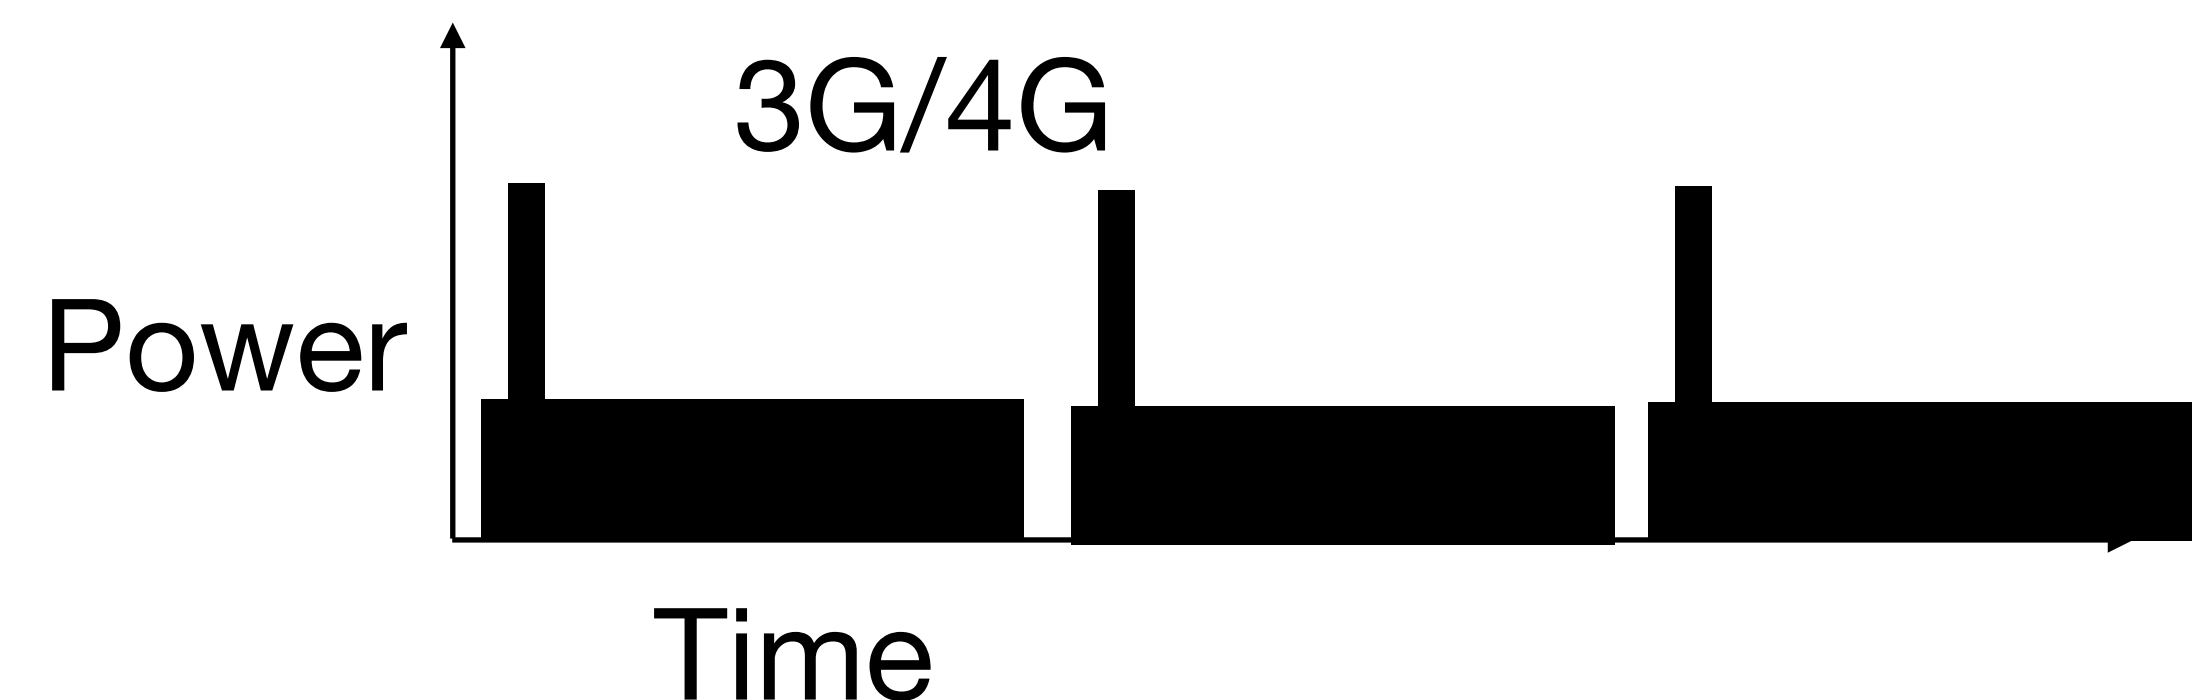
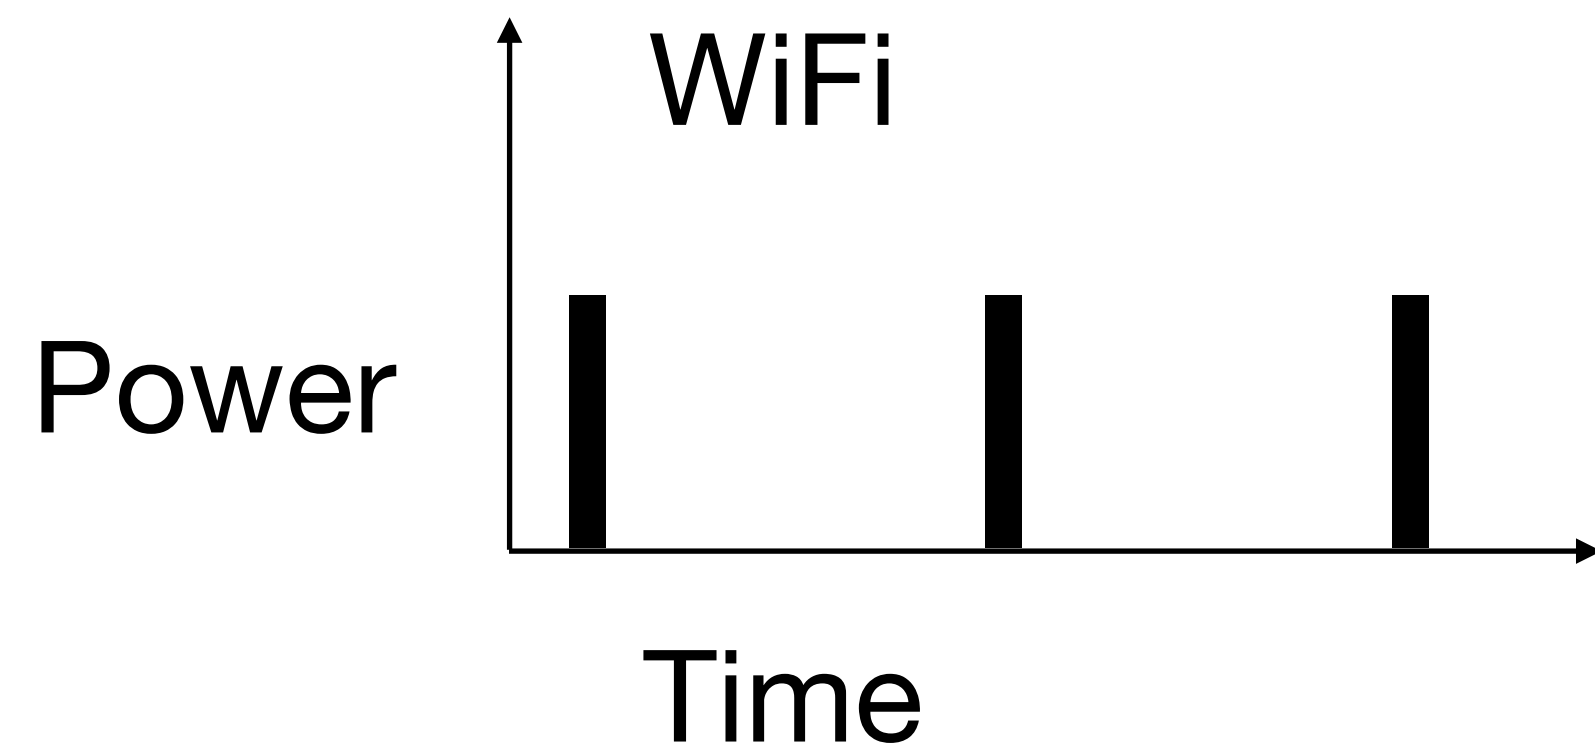
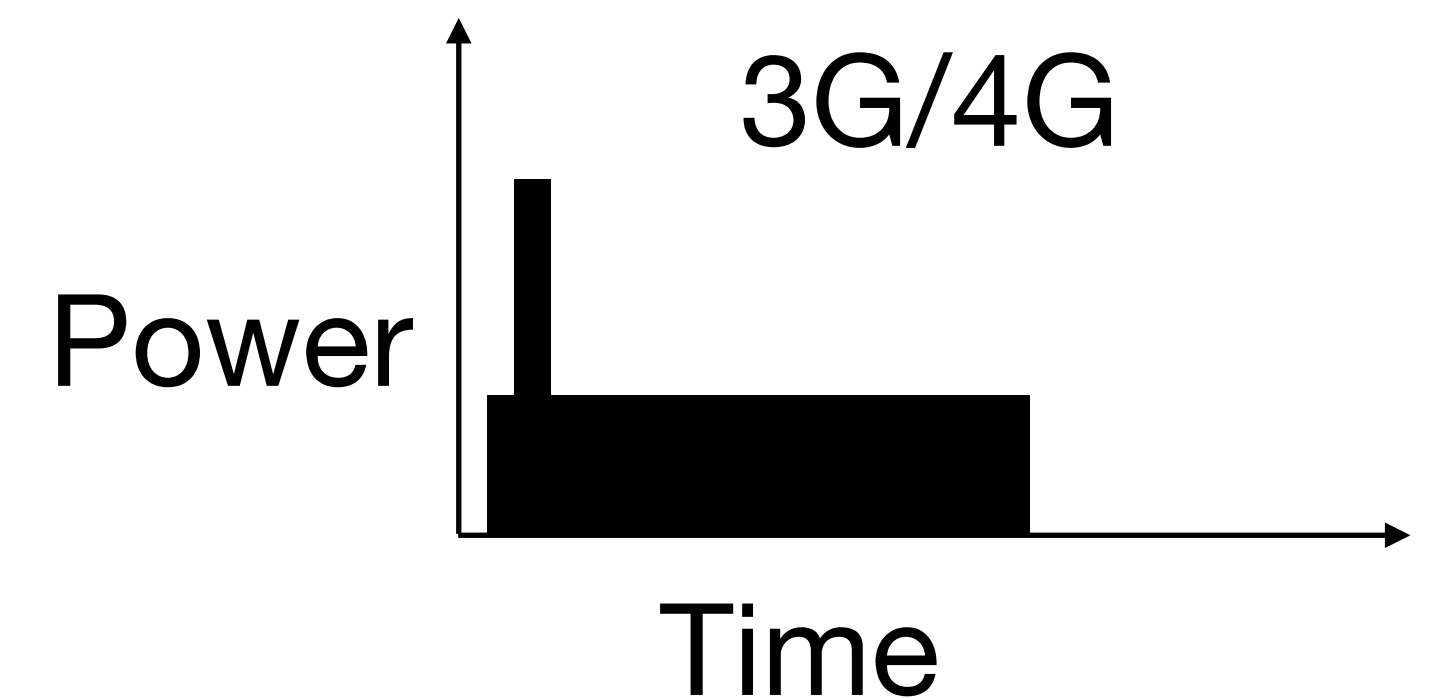
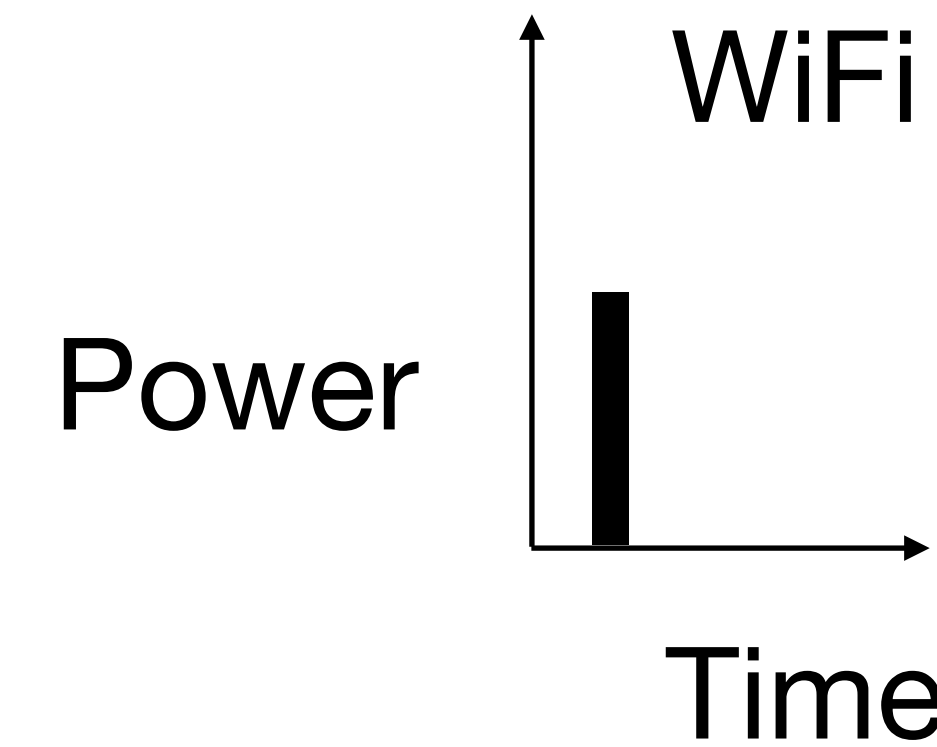
- Abstraction allows OS and applications to stay device-neutral
- Abstraction can hurt.



Fitting into the OS

Hide device specific details in device driver

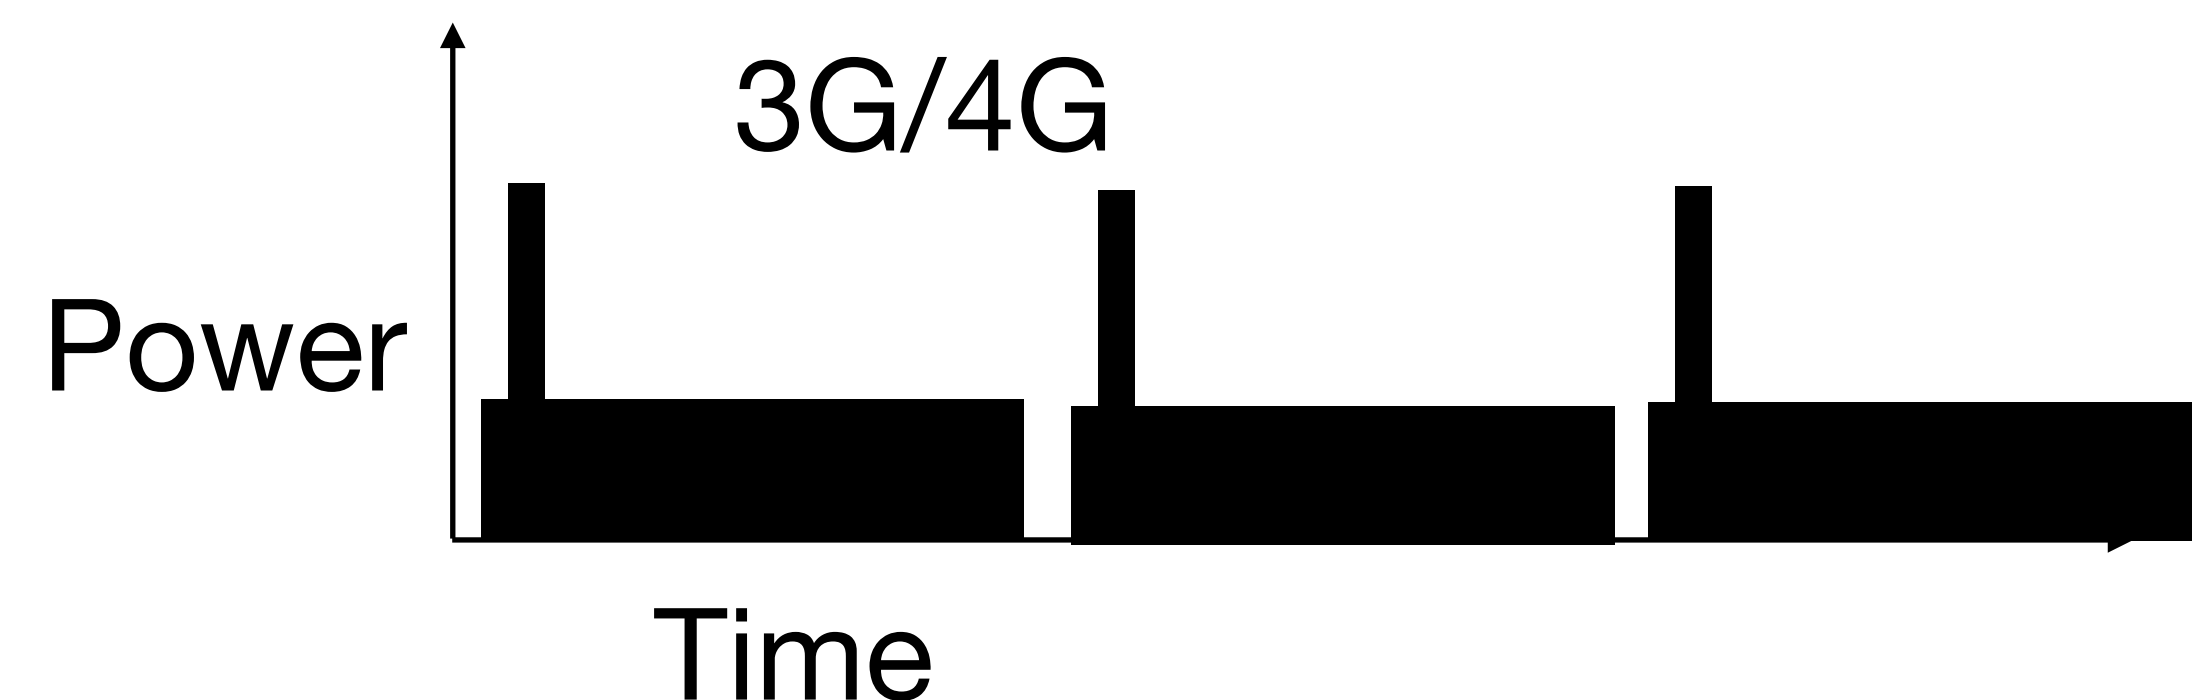
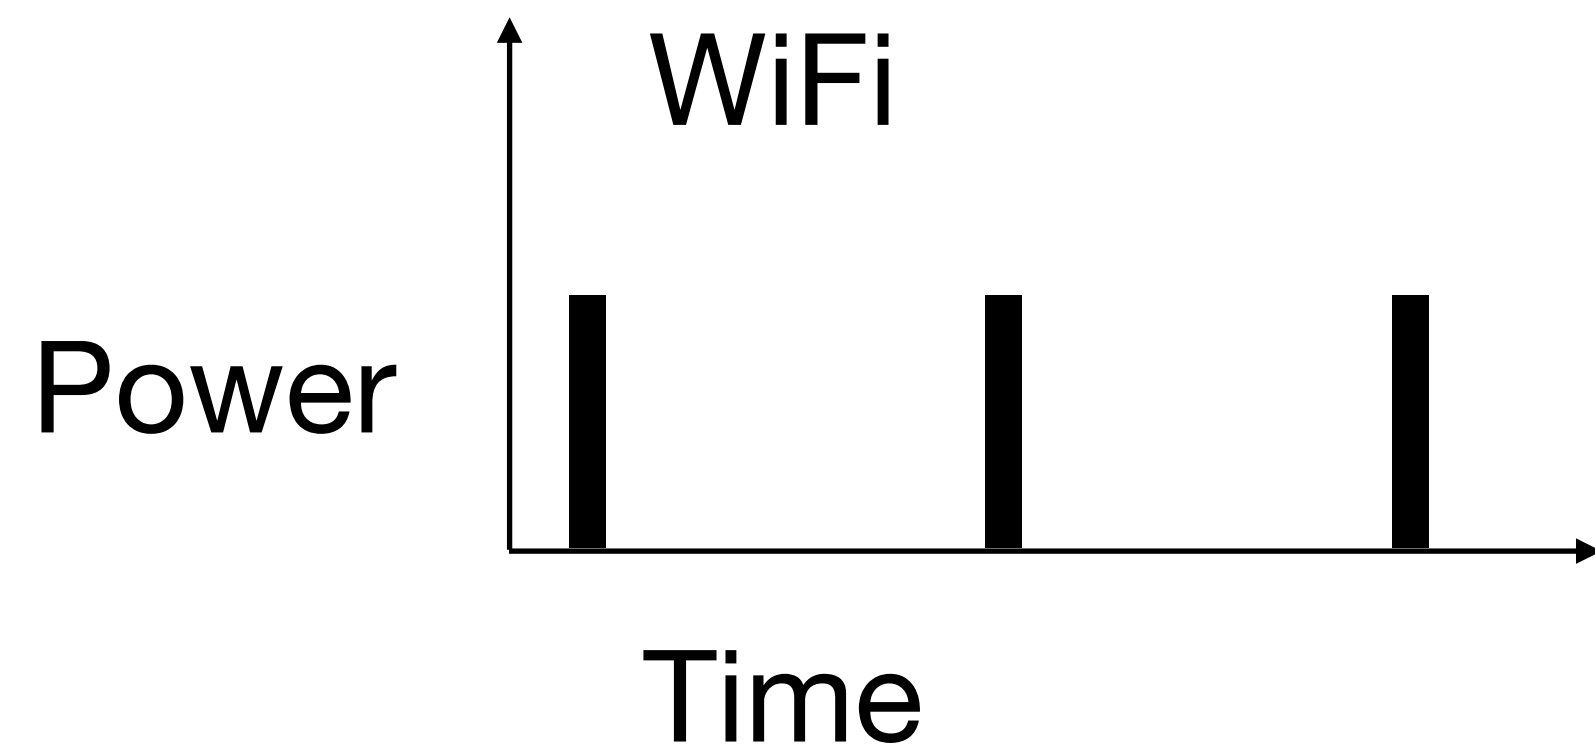
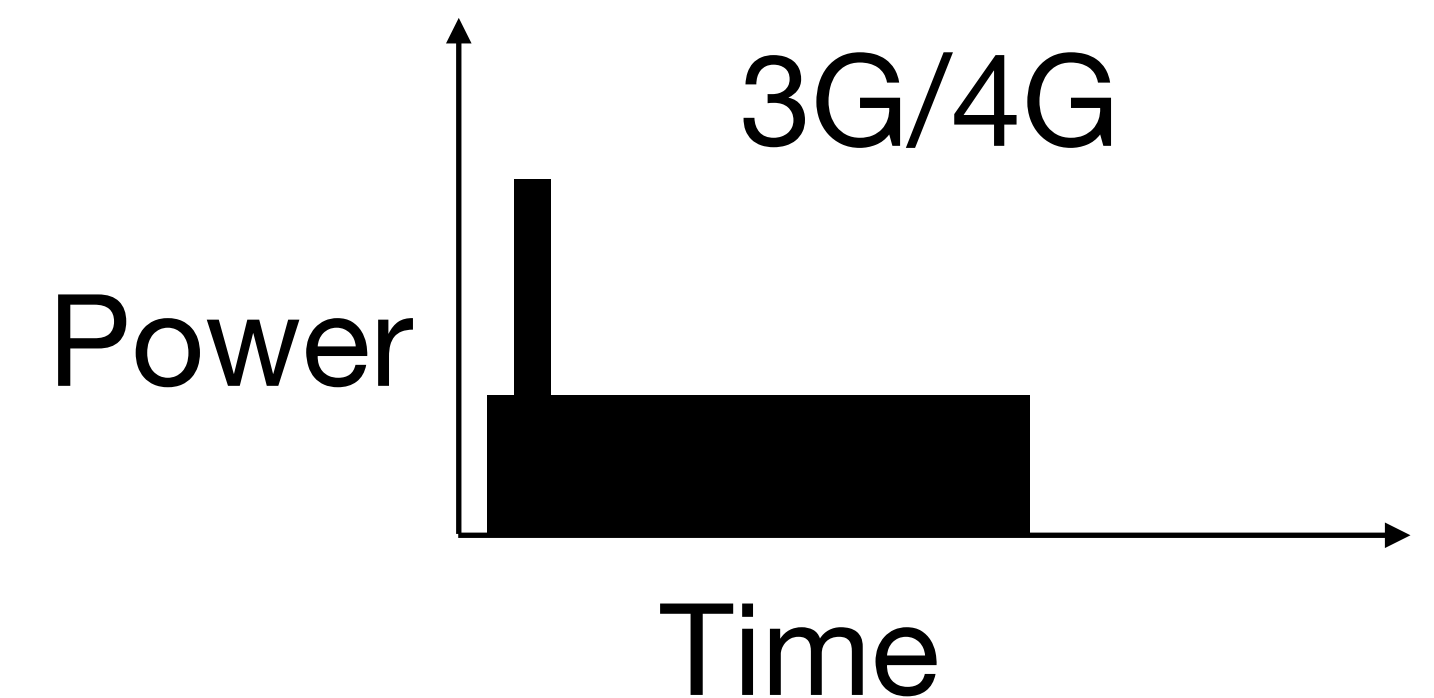
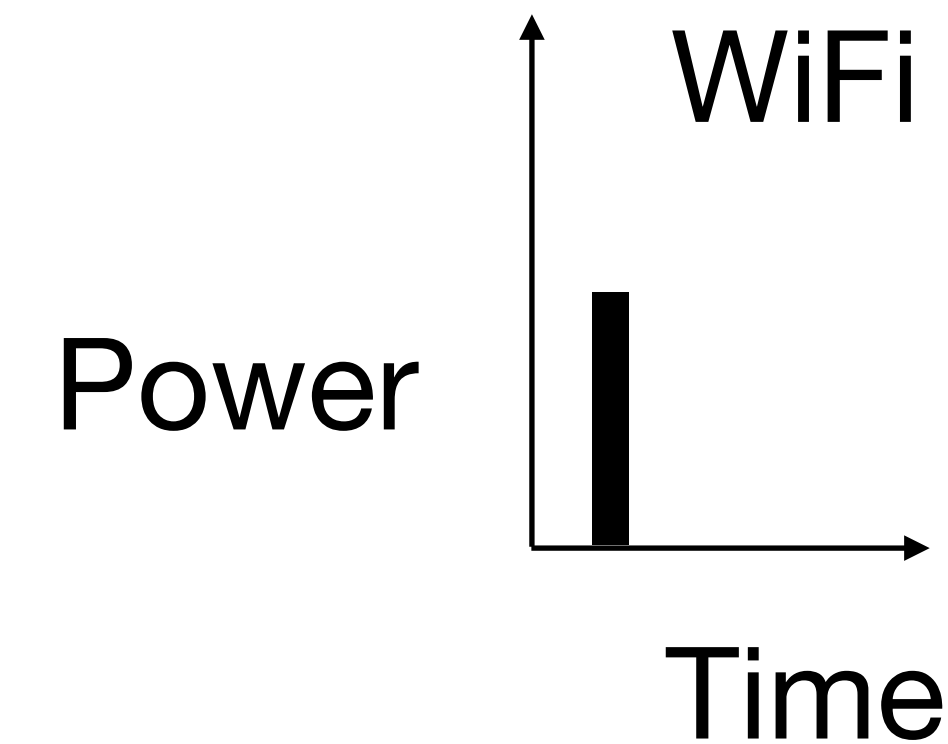
- Abstraction allows OS and applications to stay device-neutral
- Abstraction can hurt.



Fitting into the OS

Hide device specific details in device driver

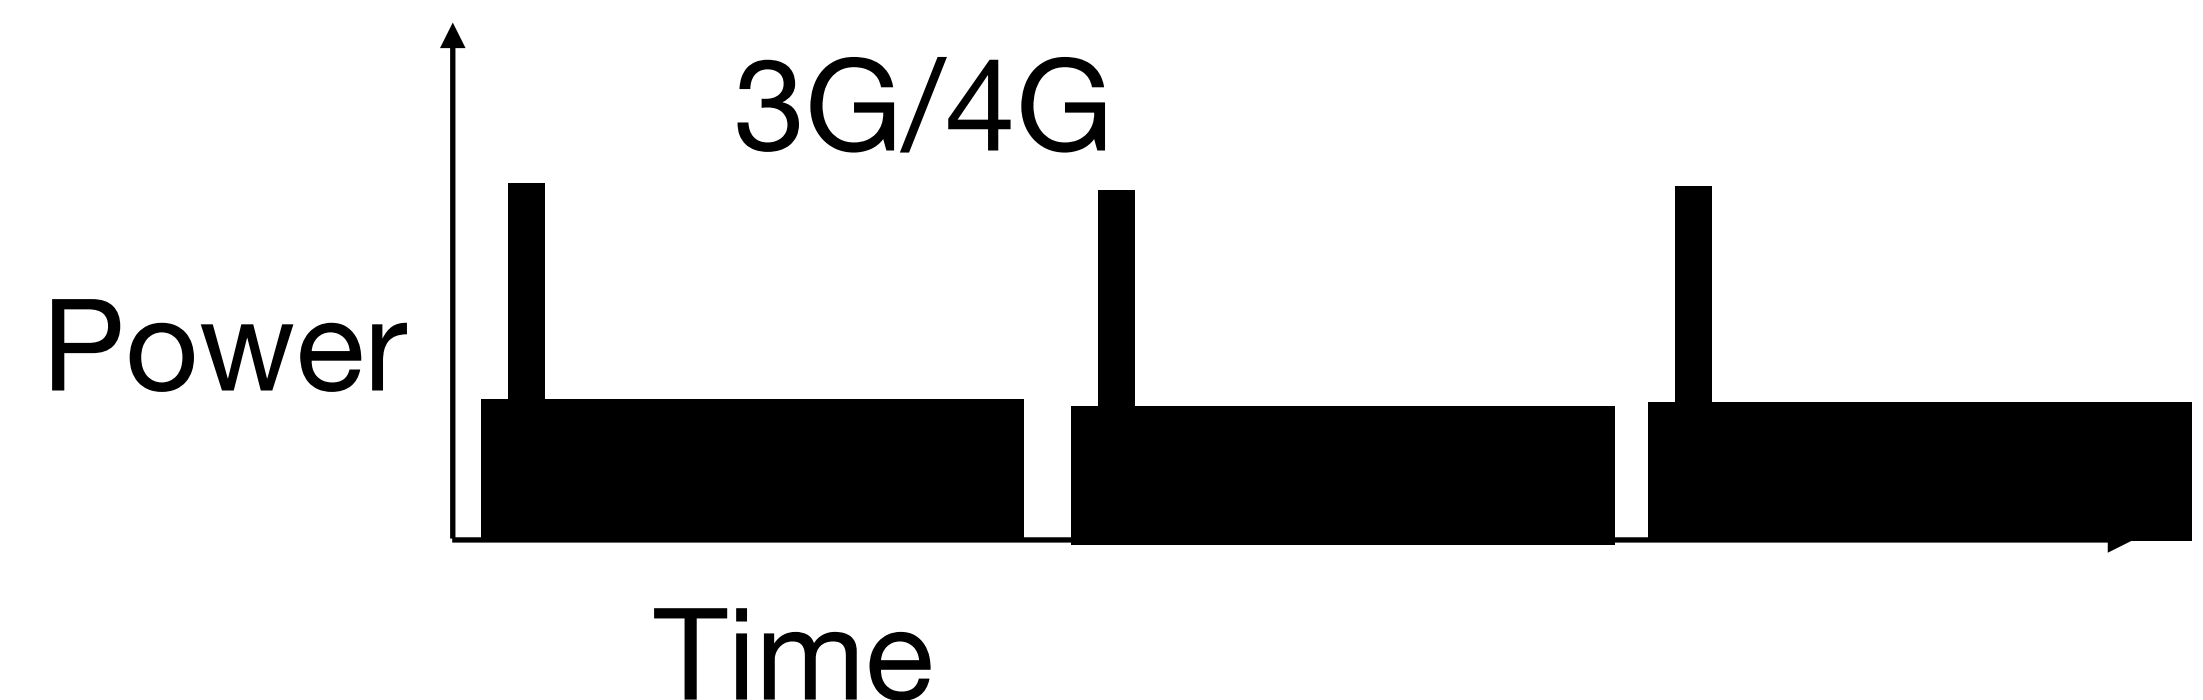
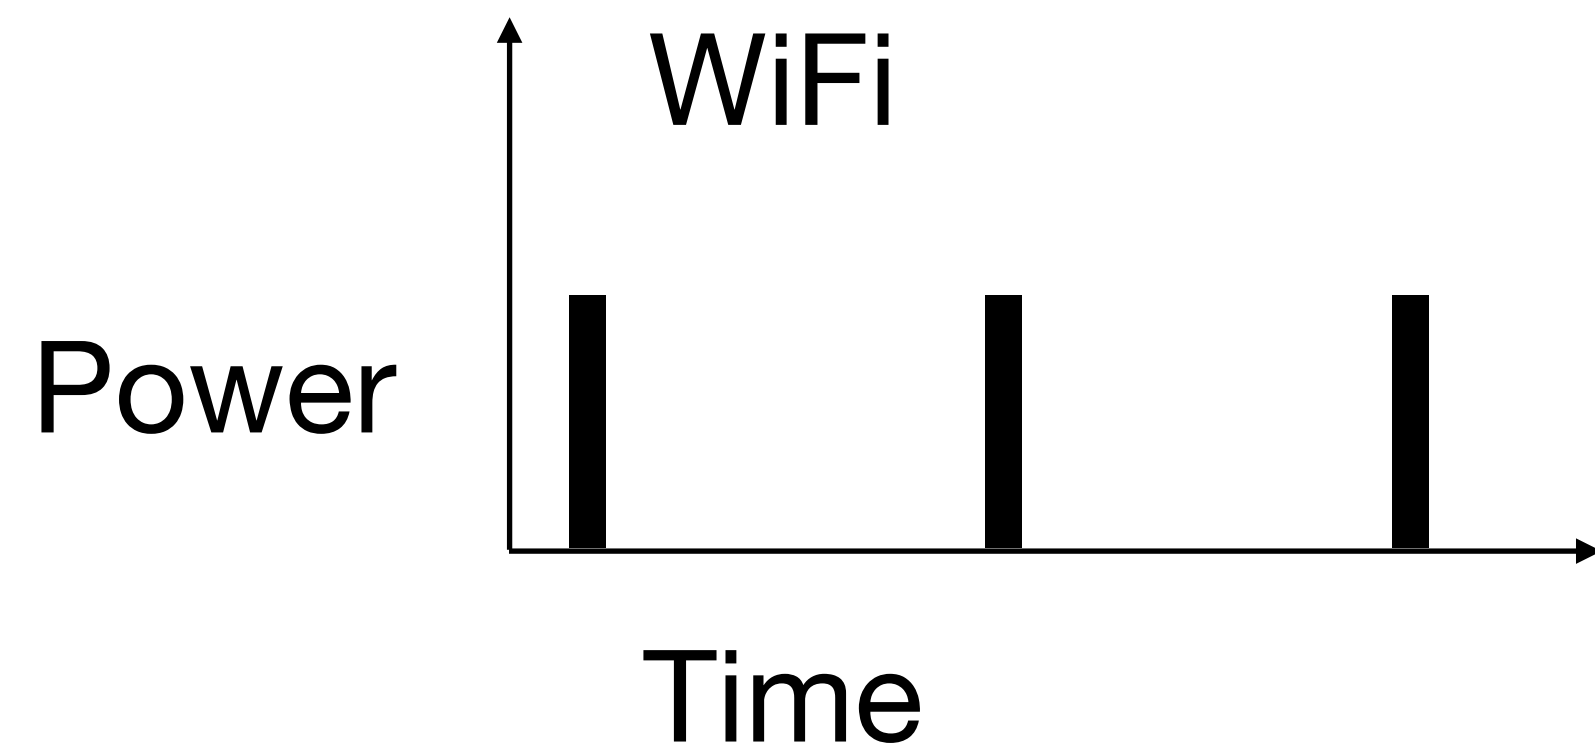
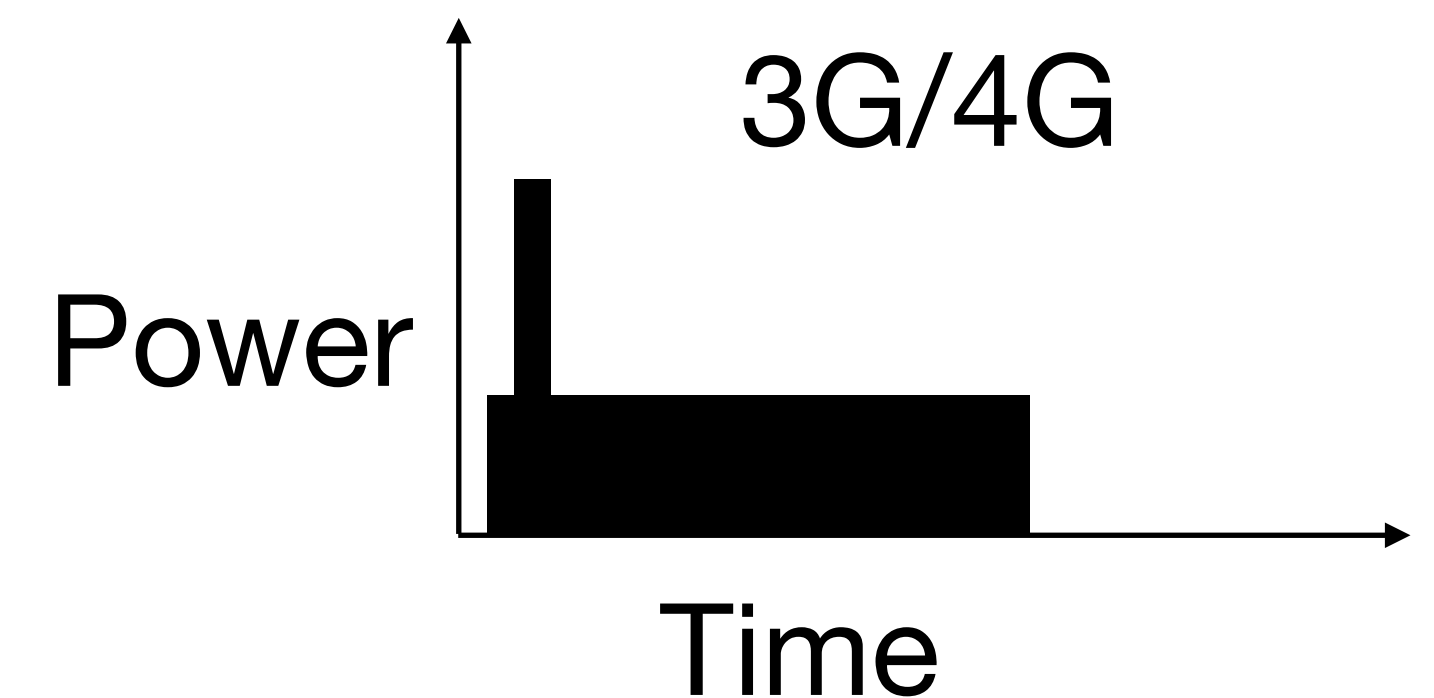
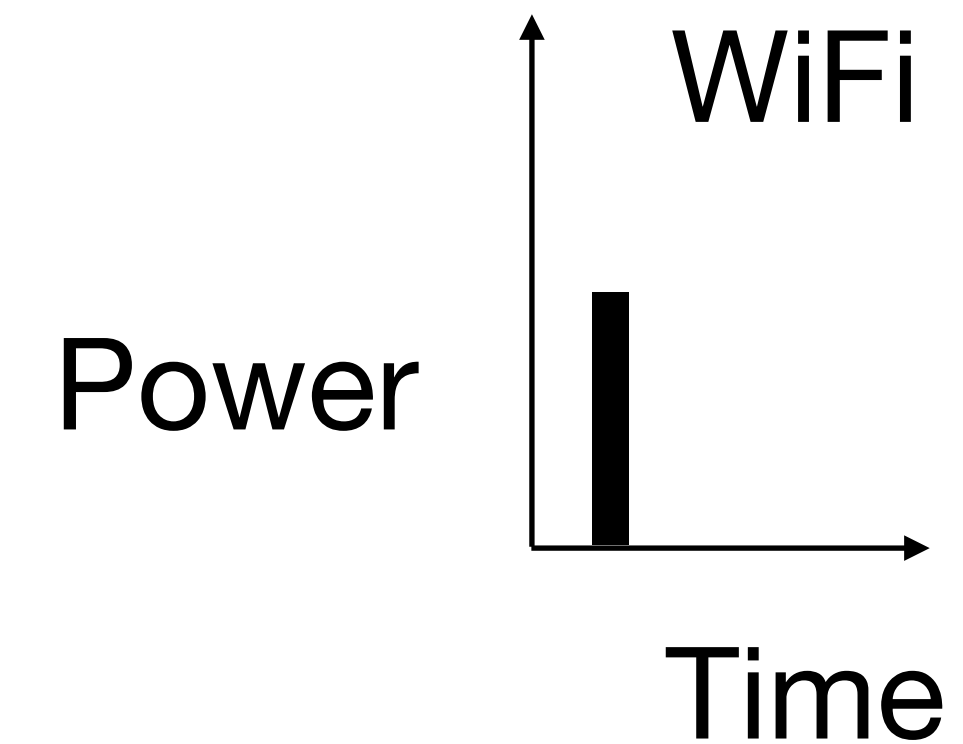
- Abstraction allows OS and applications to stay device-neutral
- Abstraction can hurt.
 - Example: 3G/4G are inefficient for small periodic pings



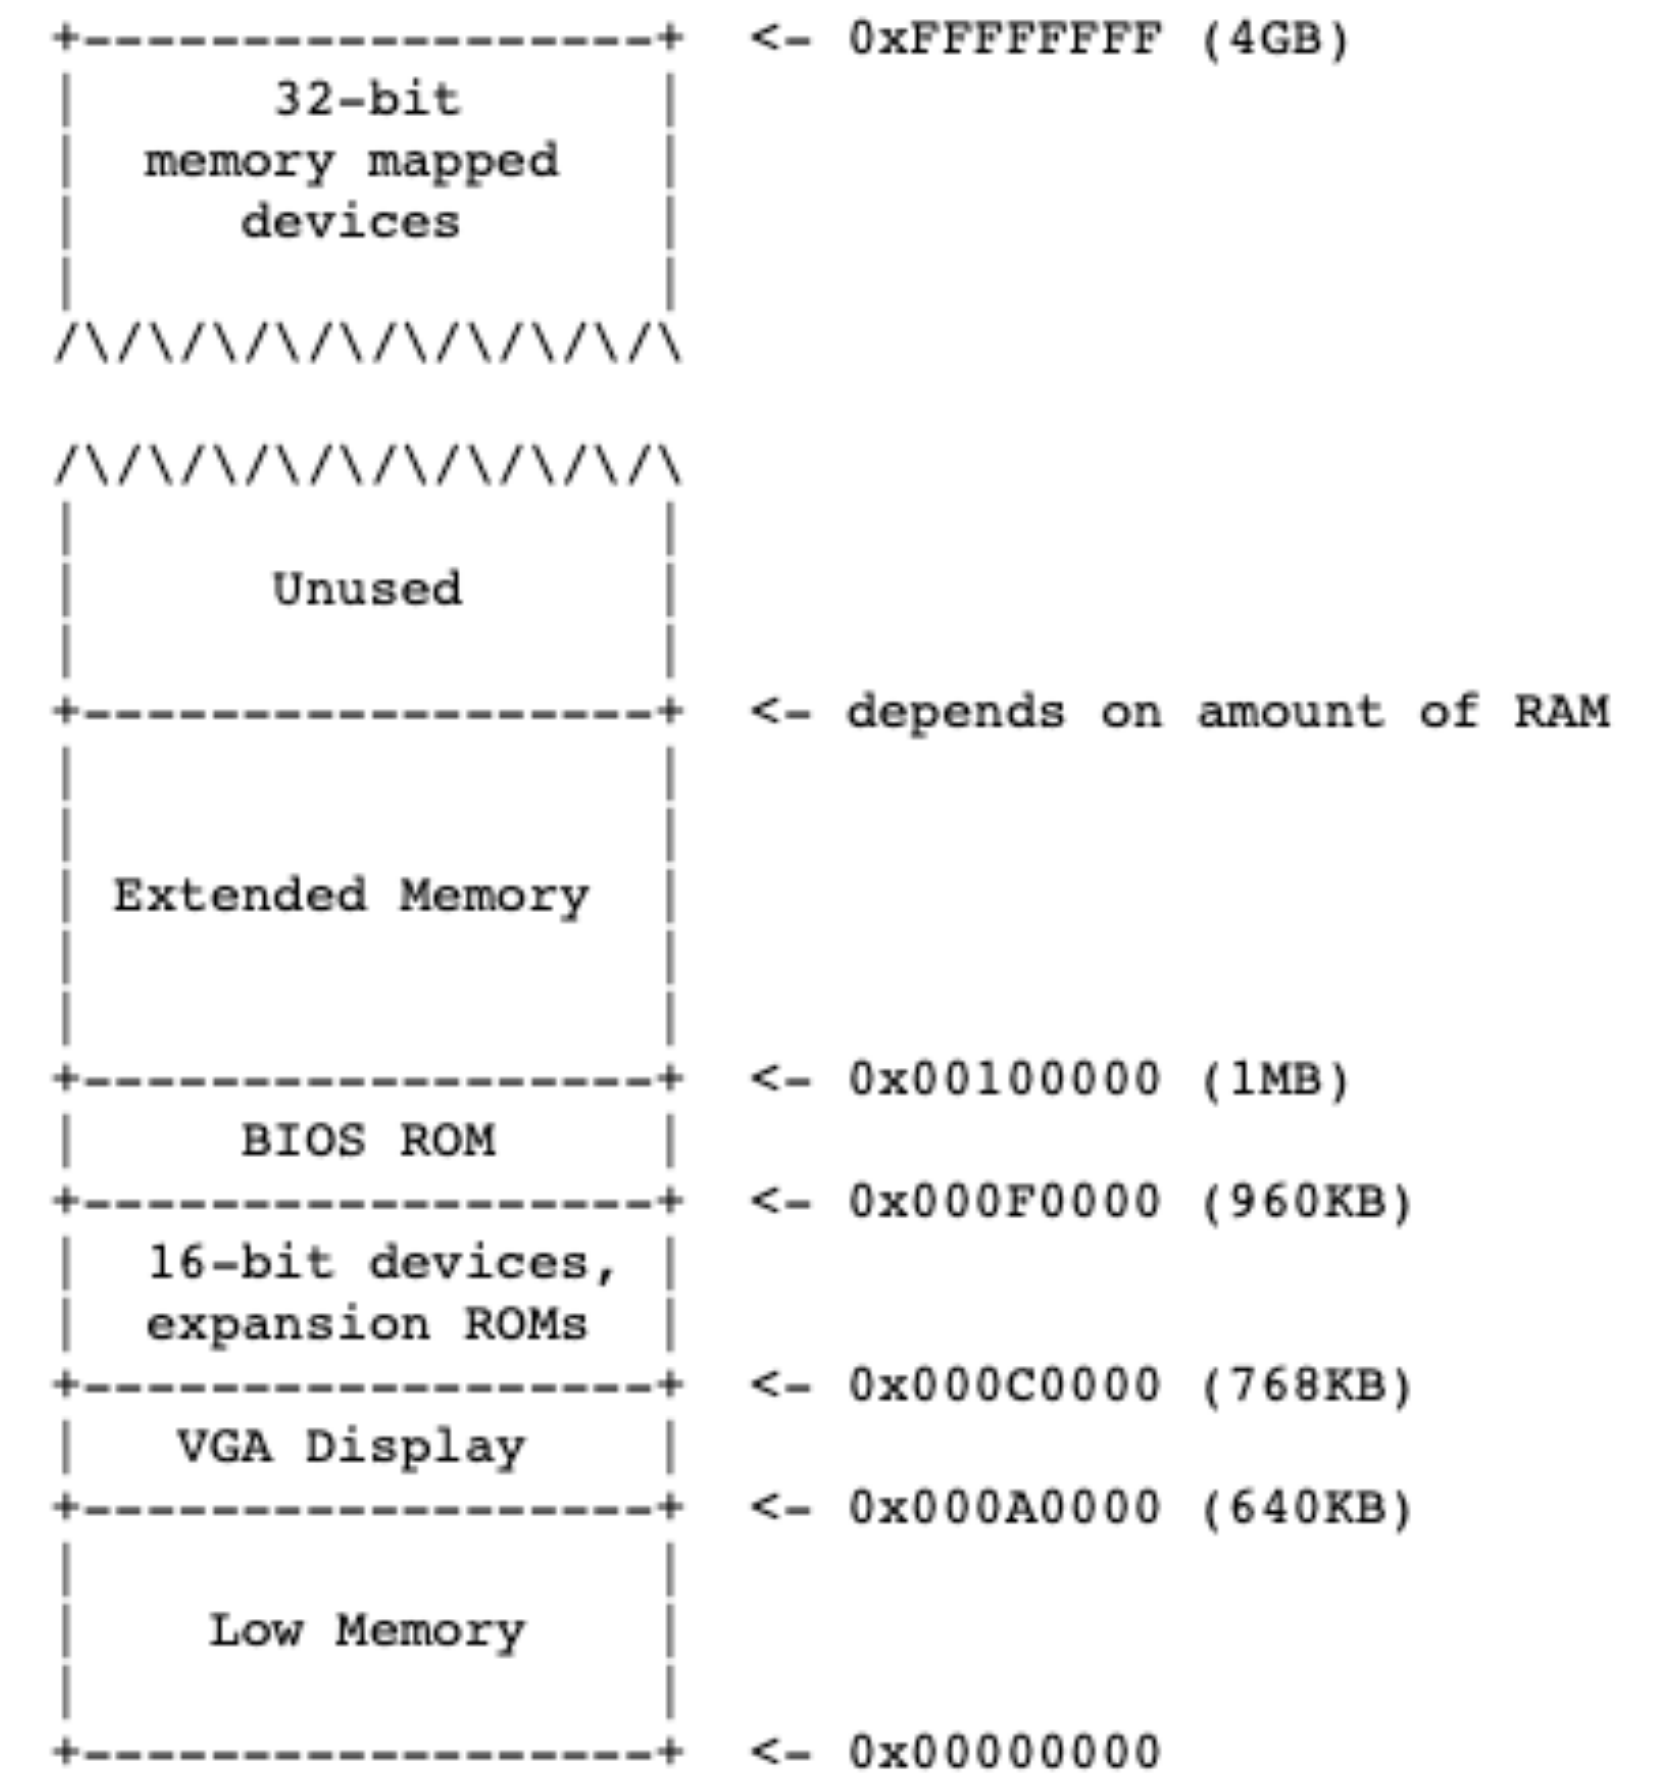
Fitting into the OS

Hide device specific details in device driver

- Abstraction allows OS and applications to stay device-neutral
- Abstraction can hurt.
 - Example: 3G/4G are inefficient for small periodic pings
- > 70% of OS code is device drivers. Tend to have most number of bugs

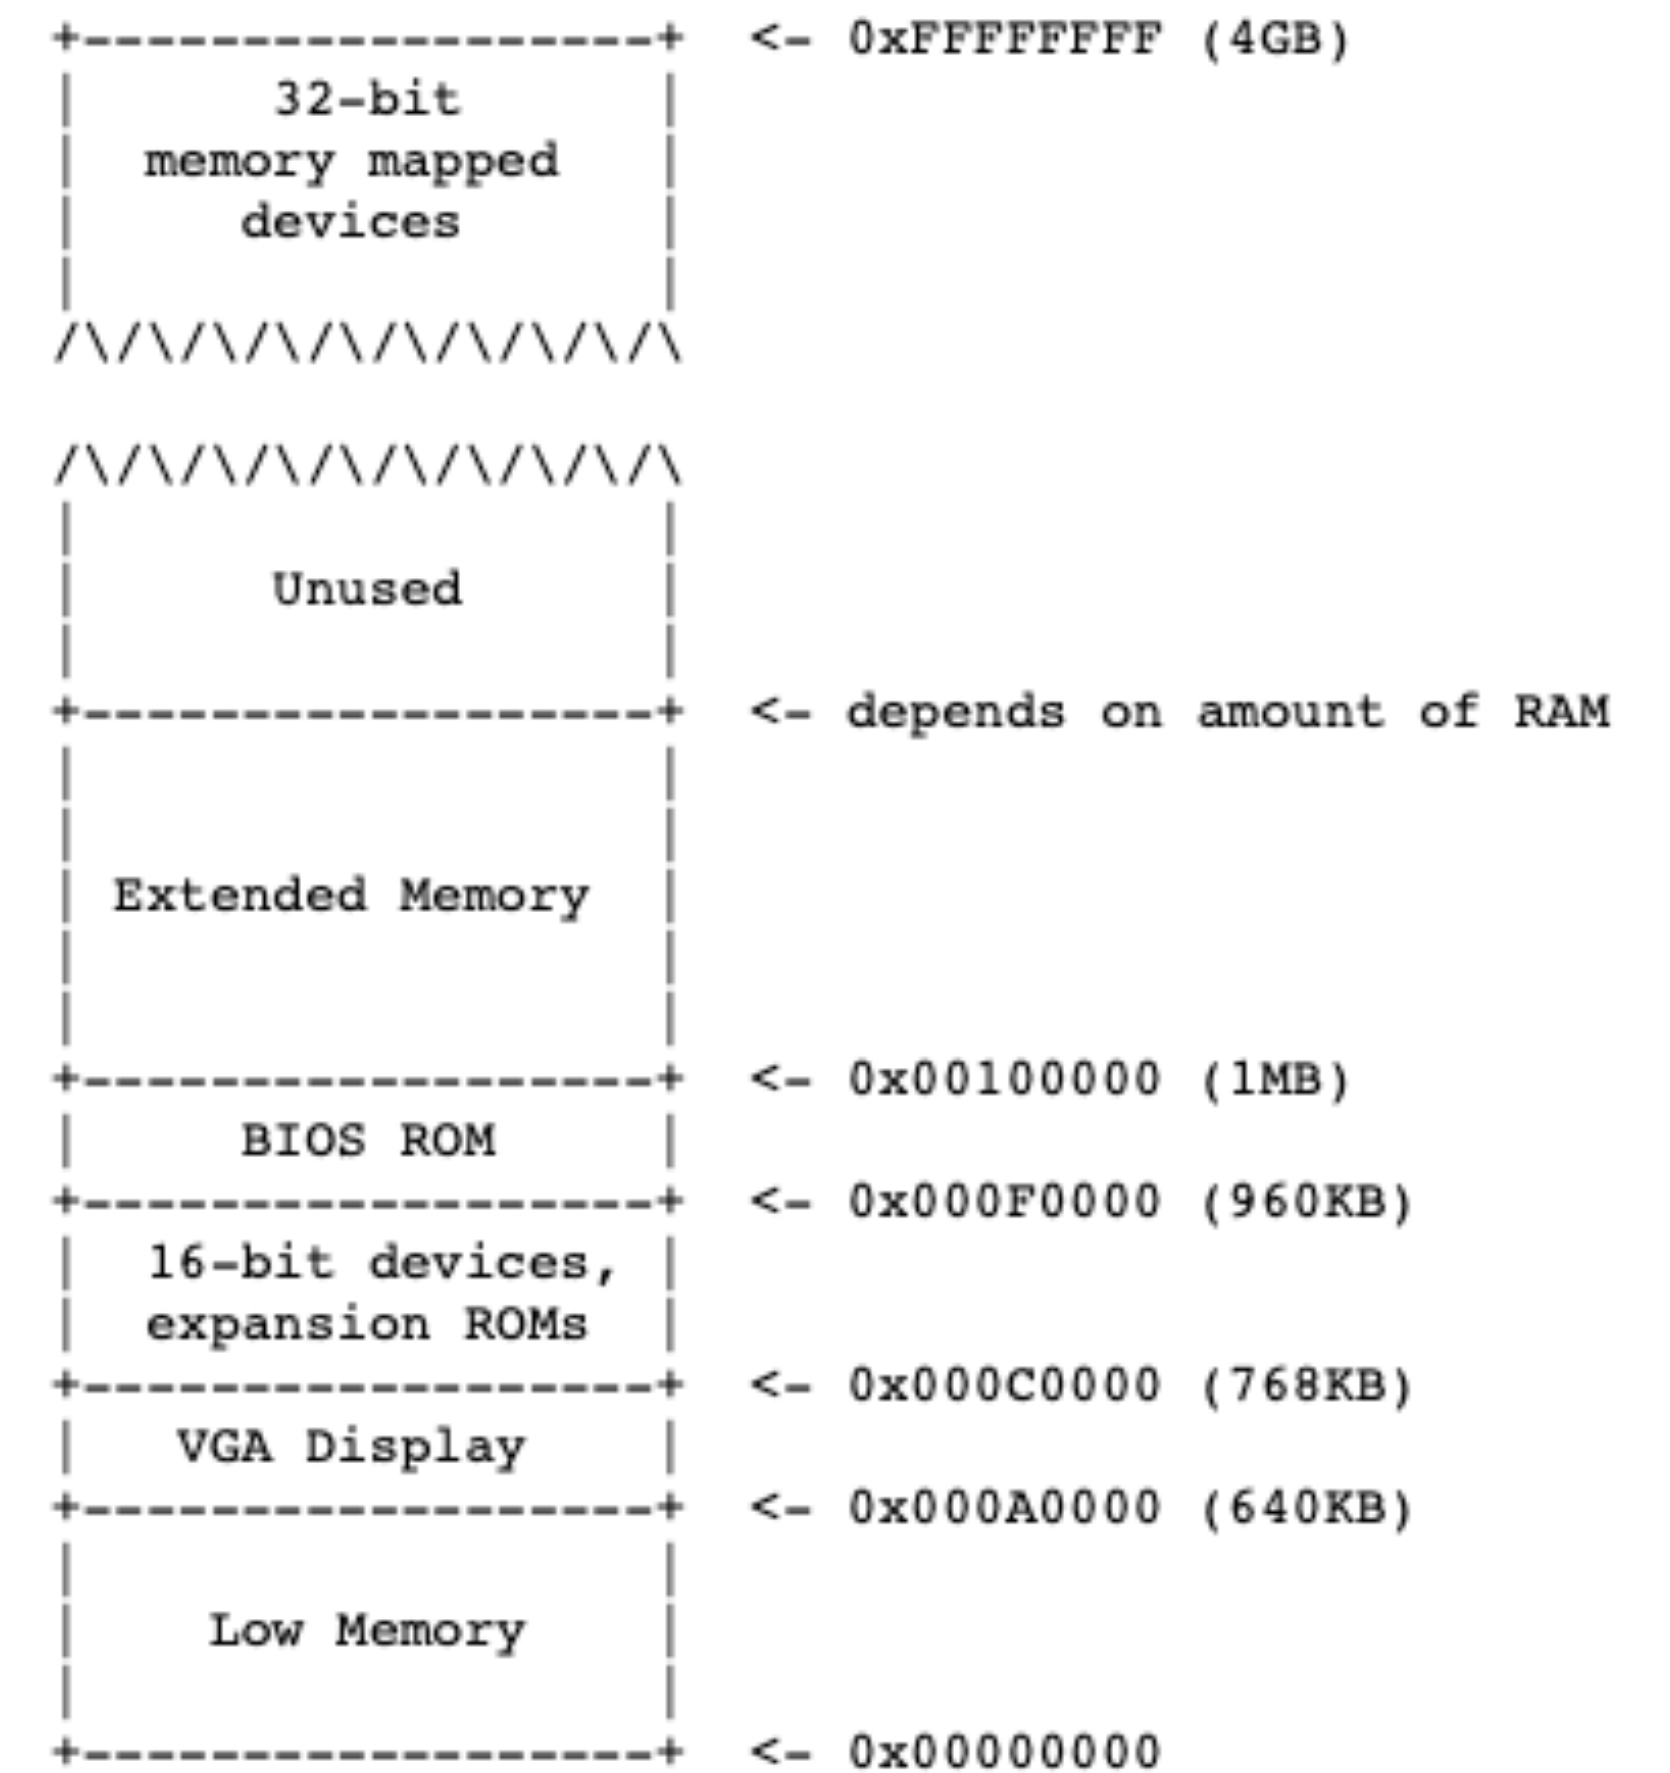


Memory-mapped IO and Port-mapped IO



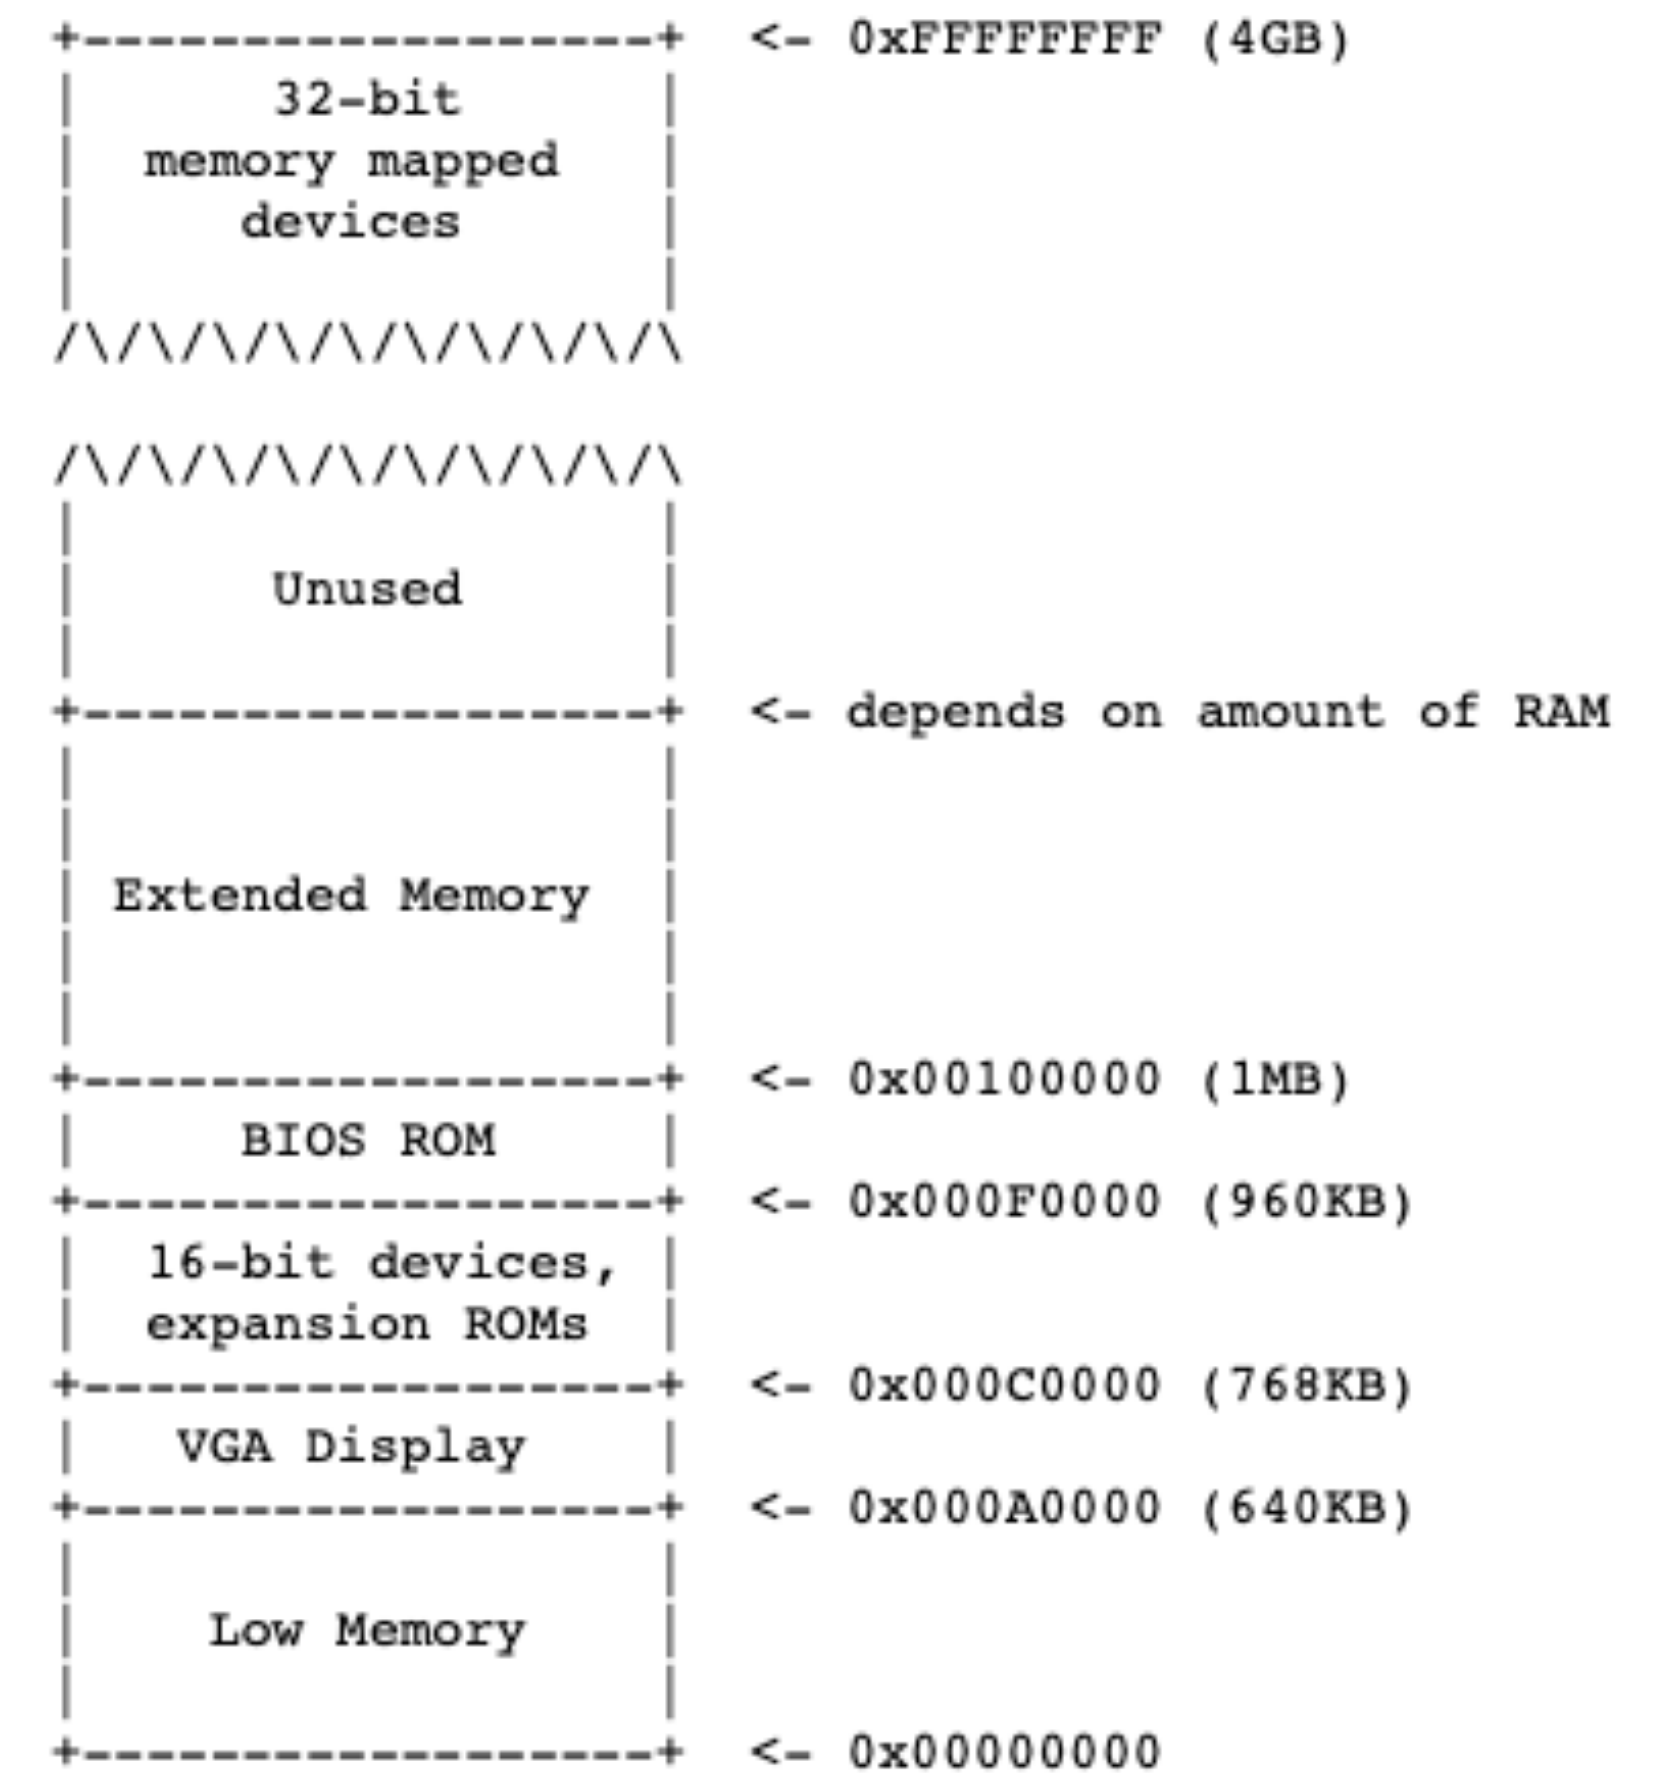
Memory-mapped IO and Port-mapped IO

- Memory mapped:



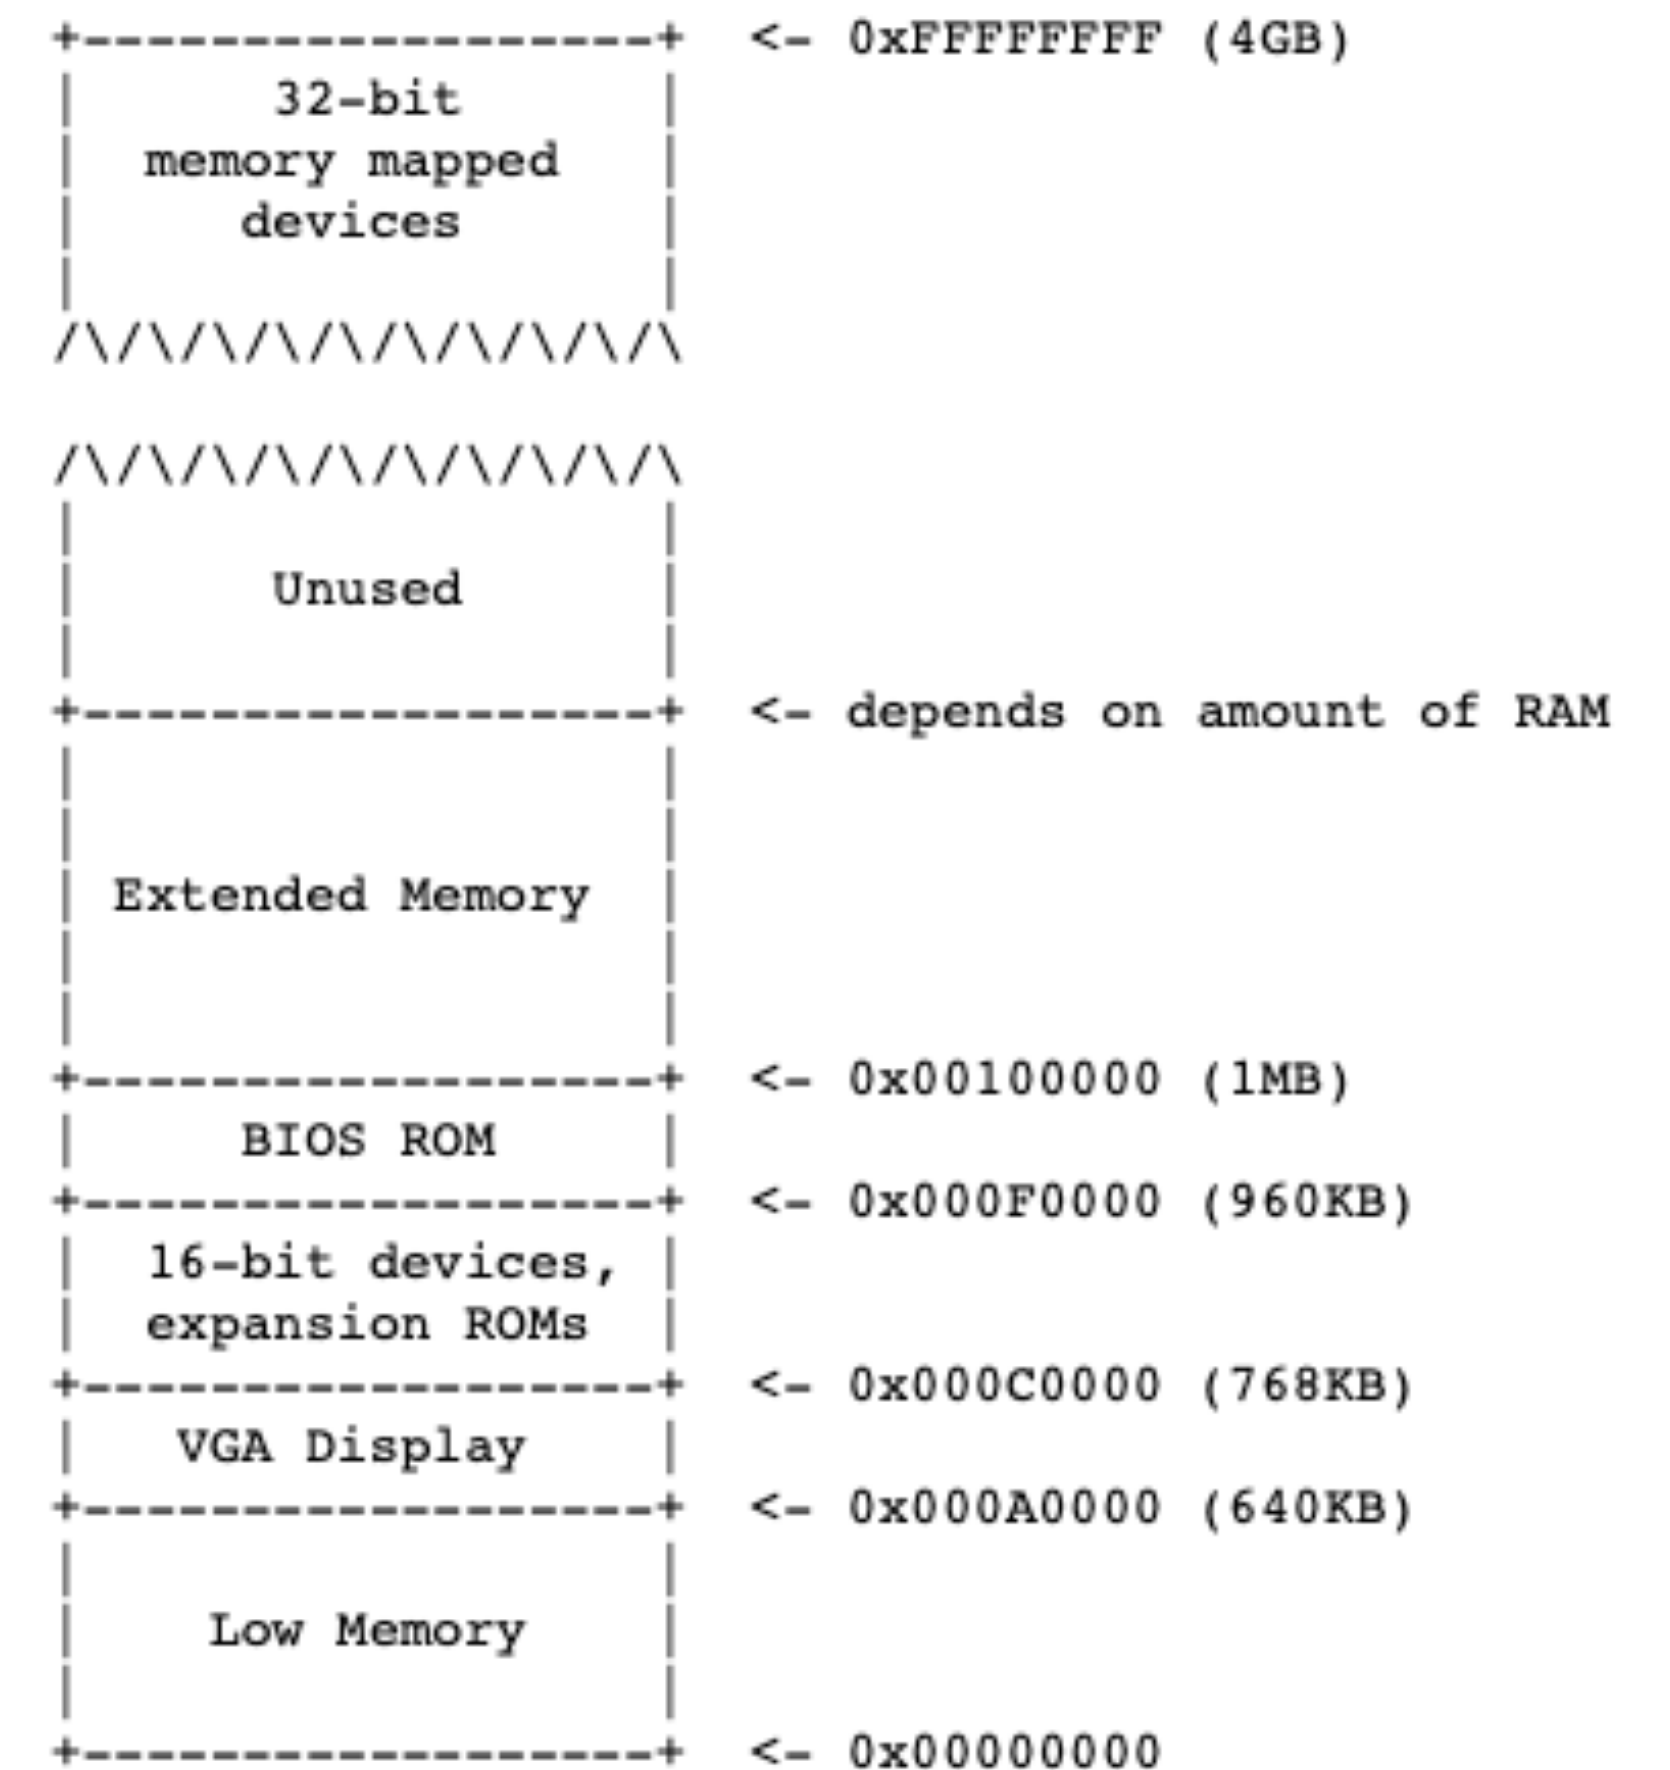
Memory-mapped IO and Port-mapped IO

- Memory mapped:
 - Regular memory access instructions



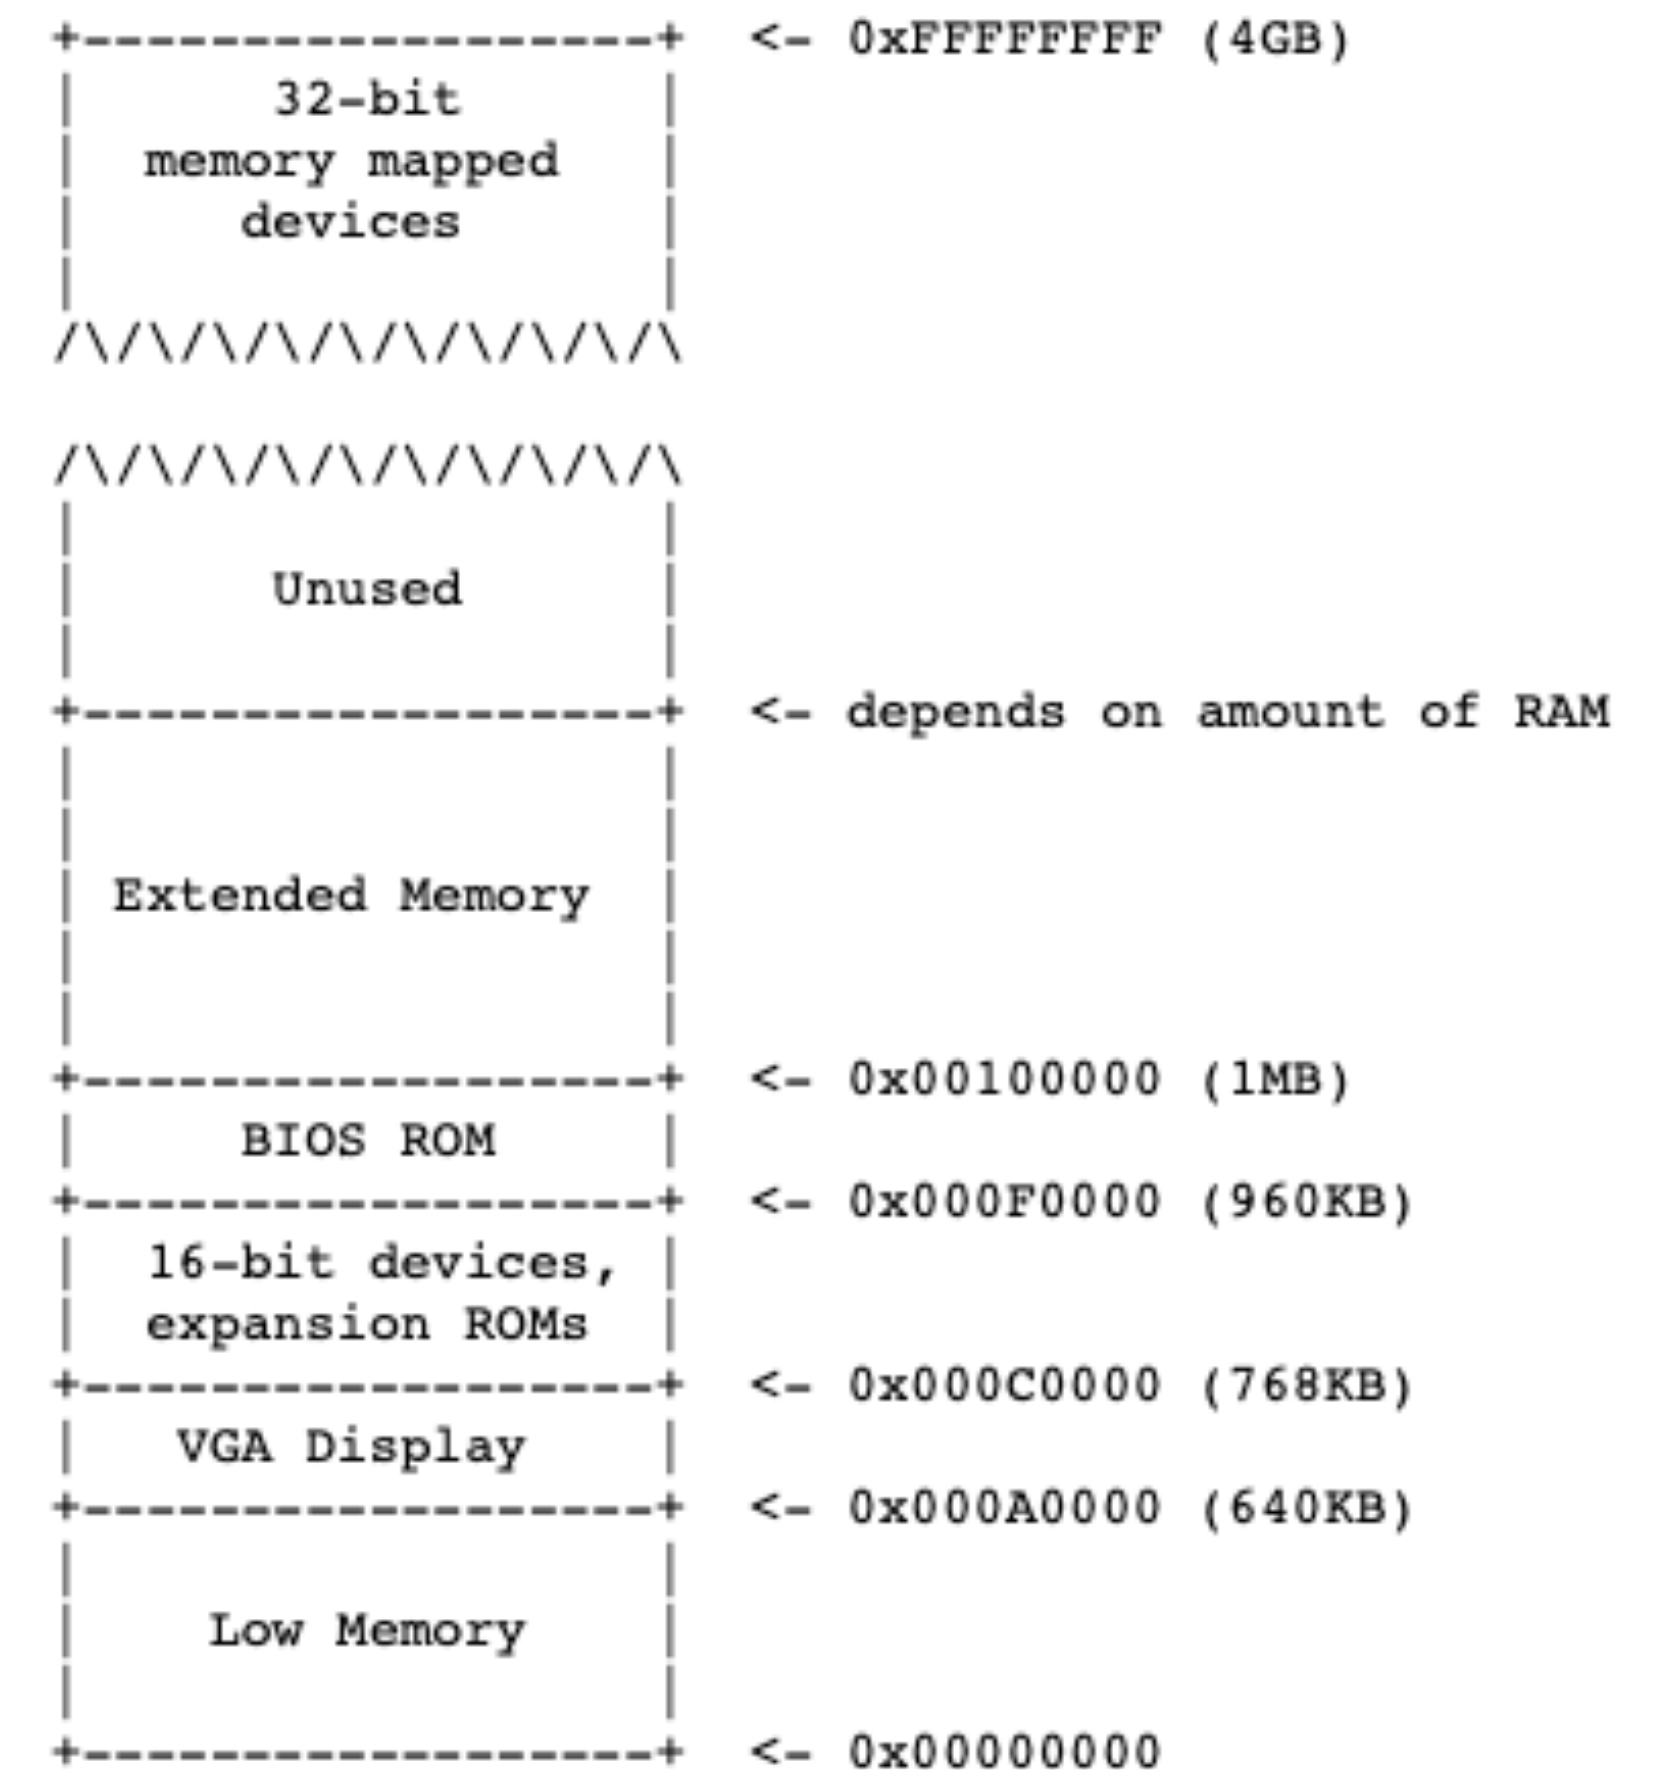
Memory-mapped IO and Port-mapped IO

- Memory mapped:
 - Regular memory access instructions
 - Reads and writes are routed to appropriate device



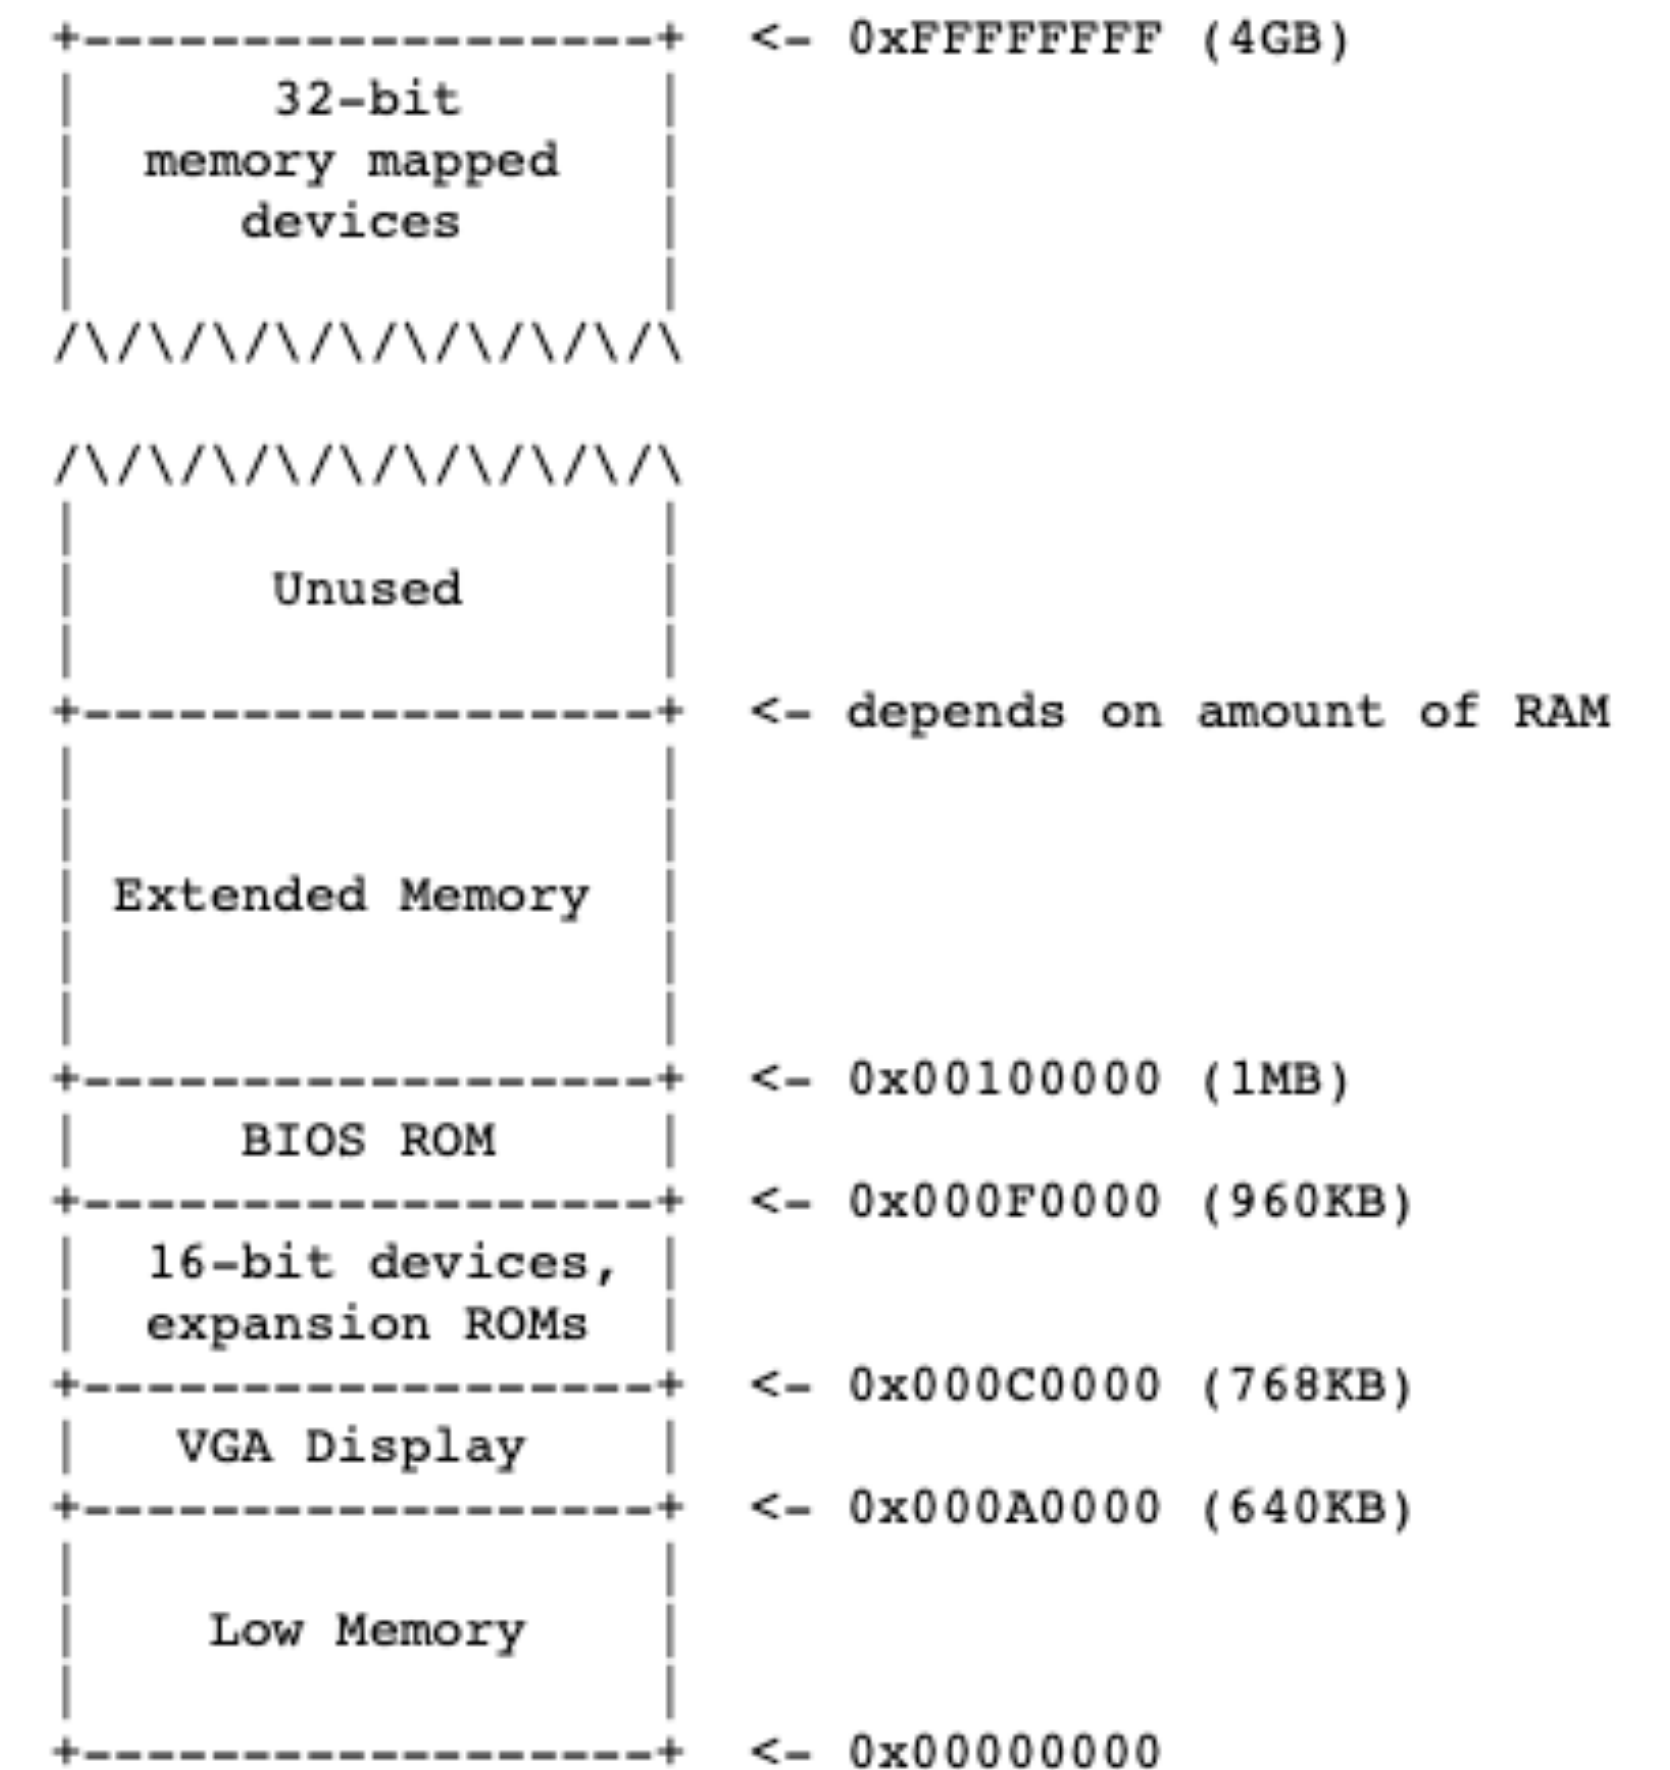
Memory-mapped IO and Port-mapped IO

- Memory mapped:
 - Regular memory access instructions
 - Reads and writes are routed to appropriate device
 - Does not behave like memory! Reading same location twice can change due to external events



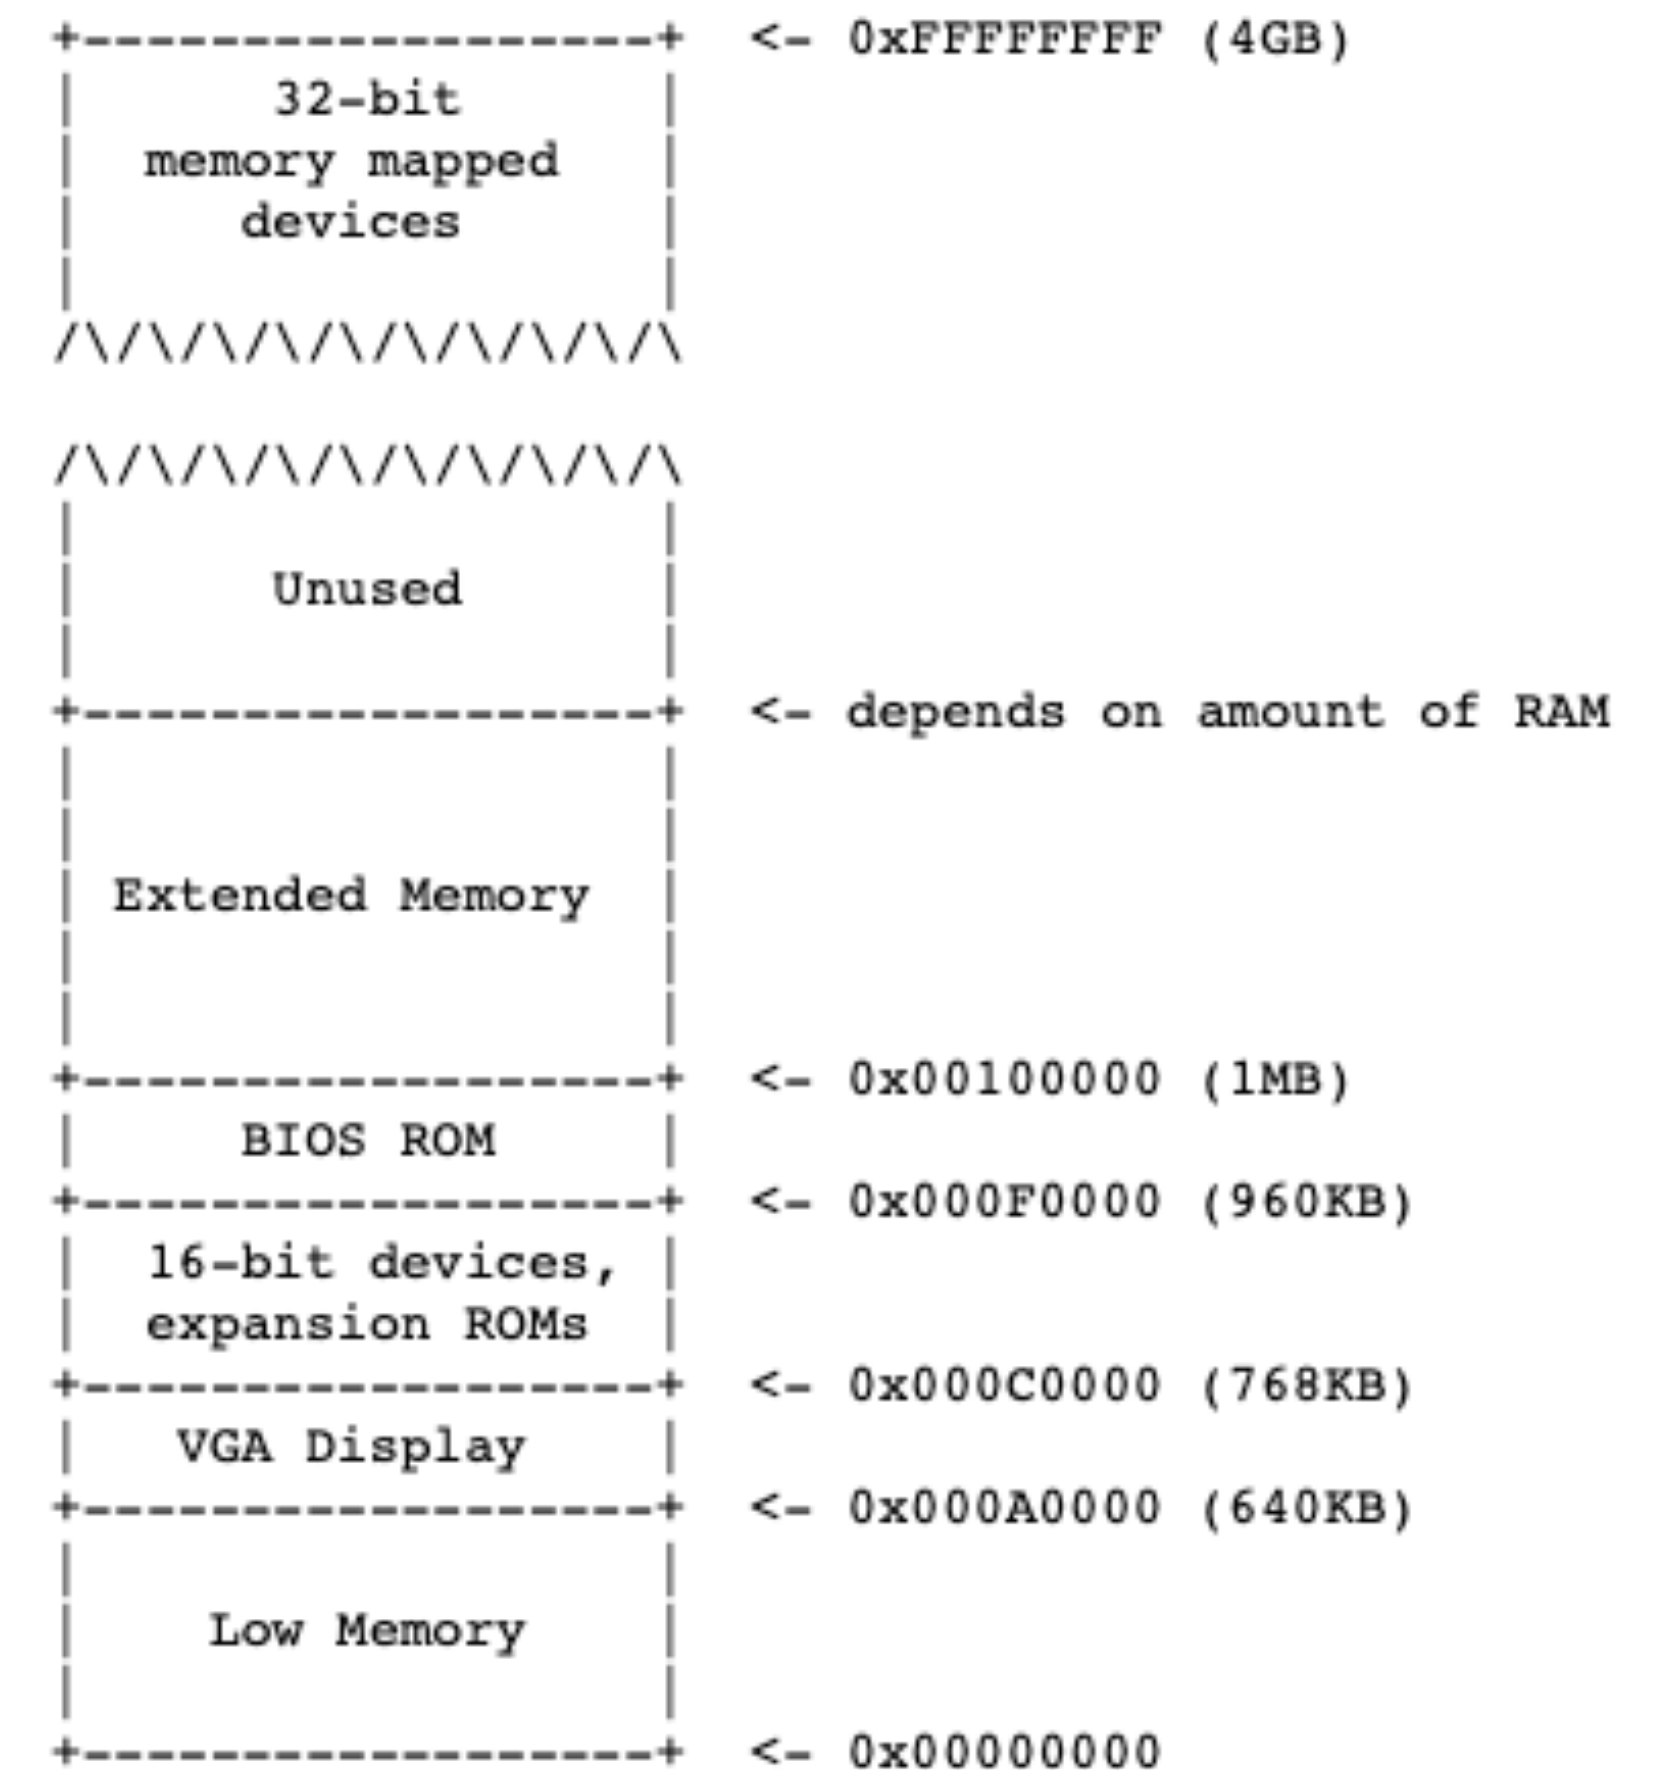
Memory-mapped IO and Port-mapped IO

- Memory mapped:
 - Regular memory access instructions
 - Reads and writes are routed to appropriate device
 - Does not behave like memory! Reading same location twice can change due to external events
- Port mapped:



Memory-mapped IO and Port-mapped IO

- Memory mapped:
 - Regular memory access instructions
 - Reads and writes are routed to appropriate device
 - Does not behave like memory! Reading same location twice can change due to external events
- Port mapped:
 - Special IN and OUT instructions



Canonical protocol

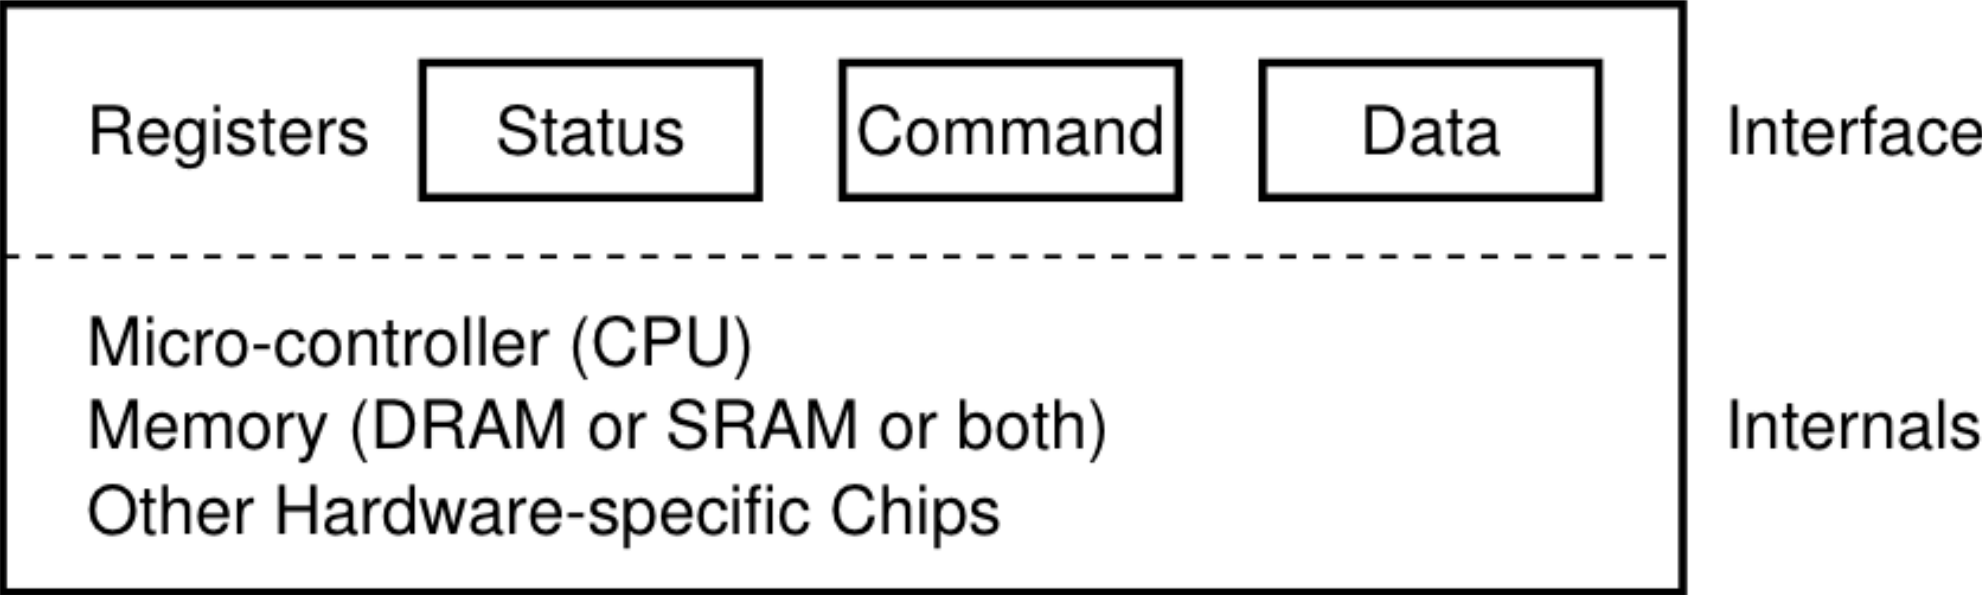


Figure 36.3: A Canonical Device

Canonical protocol

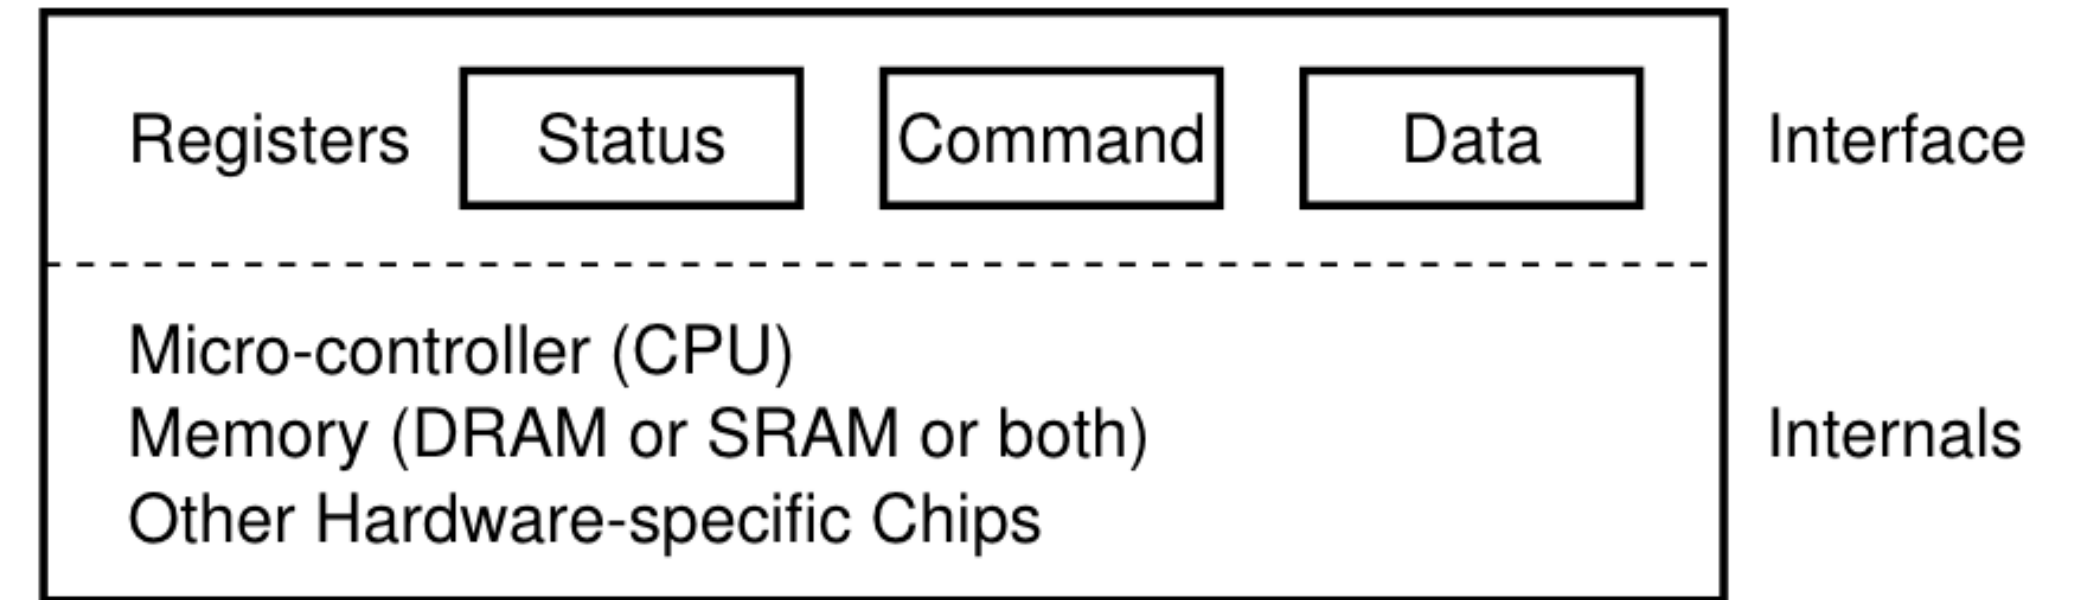


Figure 36.3: A Canonical Device

bootmain.c

```
void waitdisk(void){
    // Wait for disk ready.
    while((inb(0x1F7) & 0xC0) != 0x40);
}

// Read a single sector at offset into dst.
void readsect(void *dst, uint offset) {
    // Issue command.
    waitdisk();
    outb(0x1F2, 1);    // count = 1
    ..
    outb(0x1F7, 0x20); // cmd 0x20 - read sectors

    // Read data.
    waitdisk();
    insl(0x1F0, dst, SECTSIZE/4);
}
```

Canonical protocol

- Poll device until it is ready

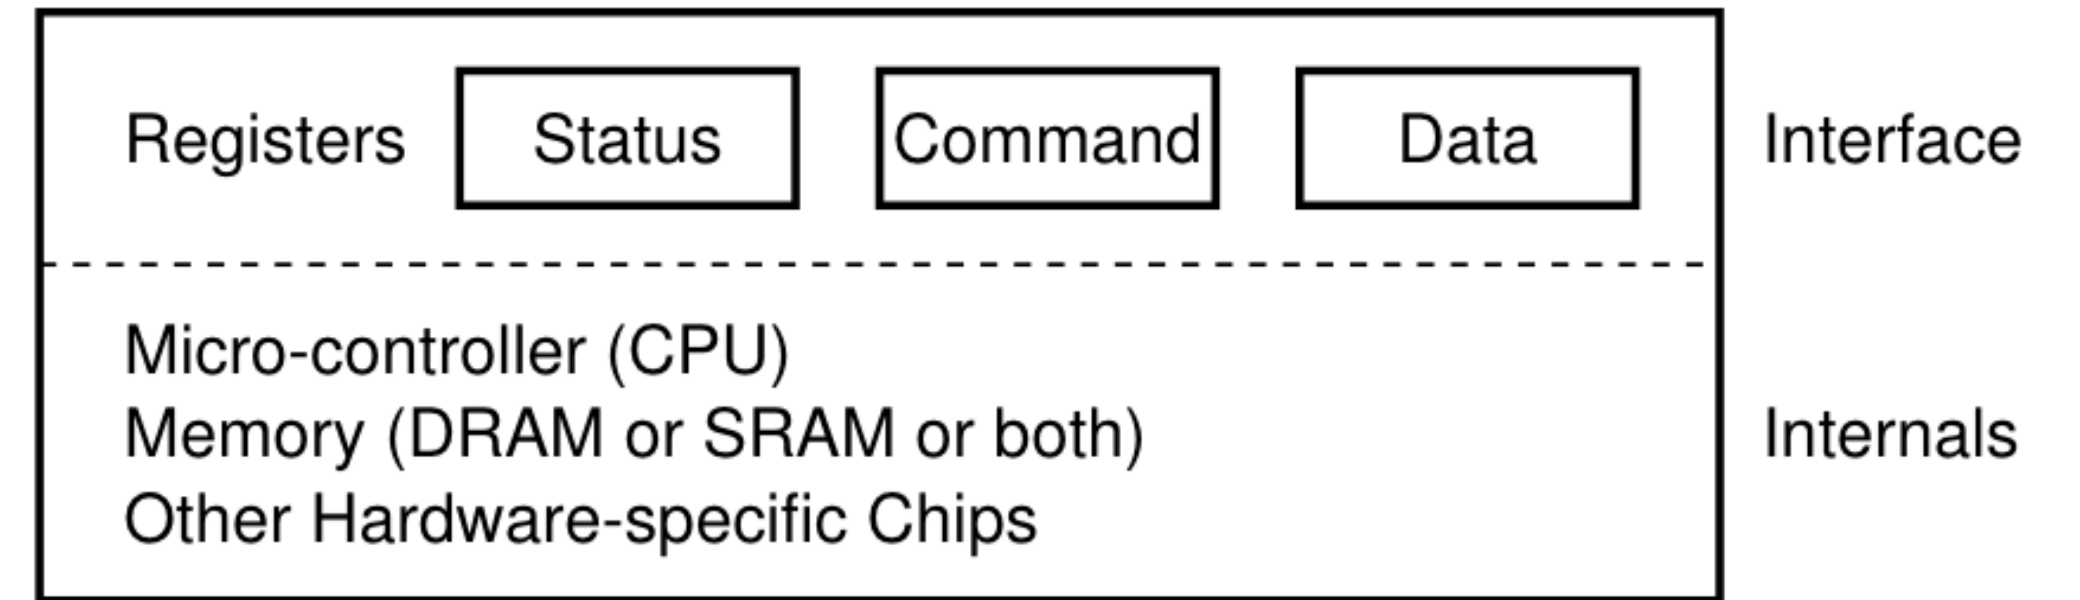


Figure 36.3: A Canonical Device

bootmain.c

```
void waitdisk(void){
    // Wait for disk ready.
    while((inb(0x1F7) & 0xC0) != 0x40);
}

// Read a single sector at offset into dst.
void readsect(void *dst, uint offset) {
    // Issue command.
    waitdisk();
    outb(0x1F2, 1);    // count = 1
    ..
    outb(0x1F7, 0x20); // cmd 0x20 - read sectors

    // Read data.
    waitdisk();
    insl(0x1F0, dst, SECTSIZE/4);
}
```

Canonical protocol

- Poll device until it is ready
- CPU cannot do anything else.

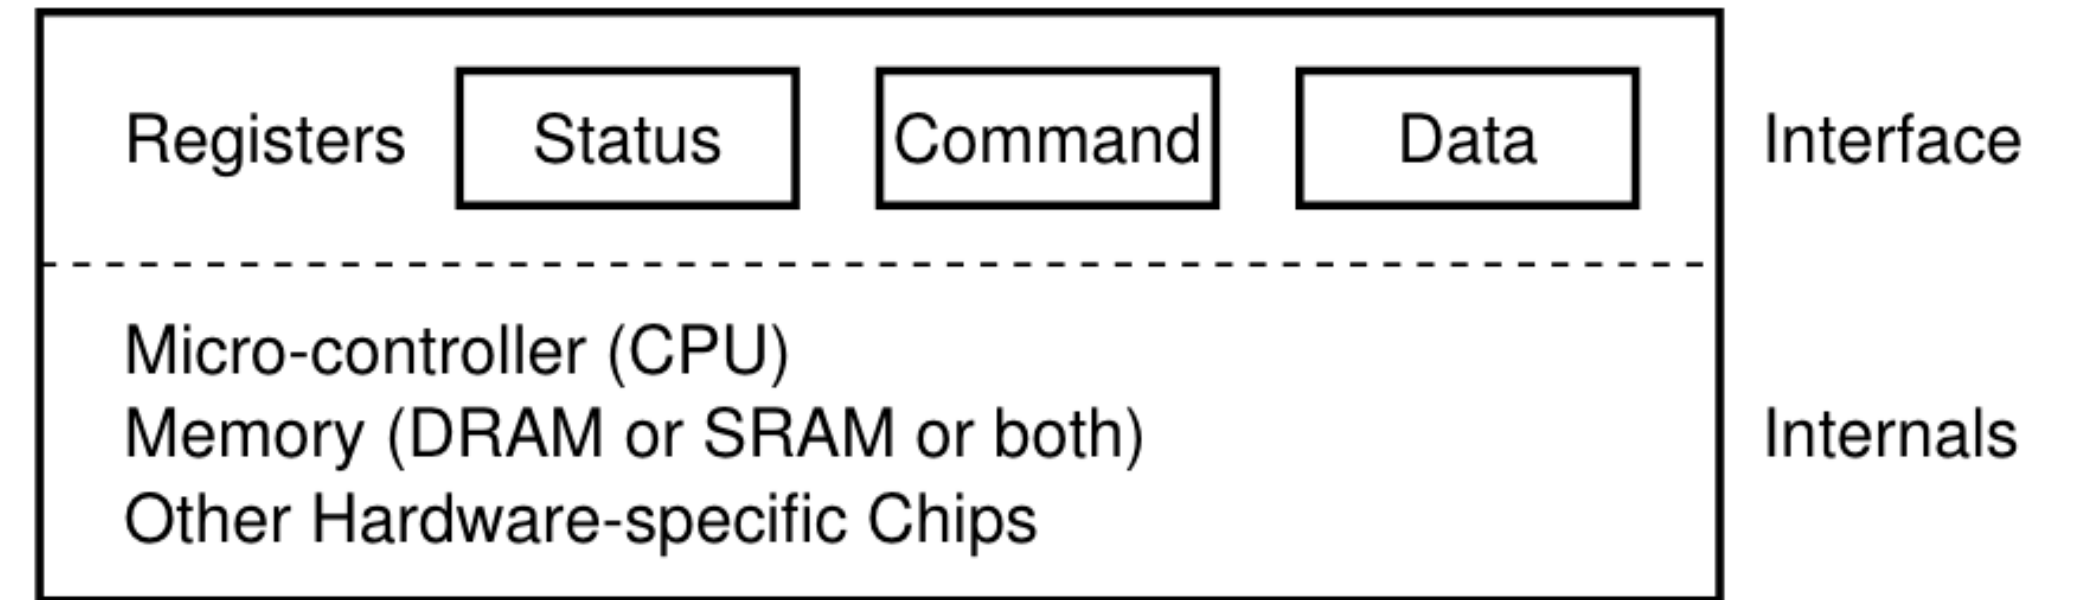


Figure 36.3: A Canonical Device

bootmain.c

```
void waitdisk(void){
    // Wait for disk ready.
    while((inb(0x1F7) & 0xC0) != 0x40);
}

// Read a single sector at offset into dst.
void readsect(void *dst, uint offset) {
    // Issue command.
    waitdisk();
    outb(0x1F2, 1);    // count = 1
    ..
    outb(0x1F7, 0x20); // cmd 0x20 - read sectors

    // Read data.
    waitdisk();
    insl(0x1F0, dst, SECTSIZE/4);
}
```

Canonical protocol

- Poll device until it is ready
- CPU cannot do anything else.
- Example: CPU needs to spend ~1 million instructions waiting for disk

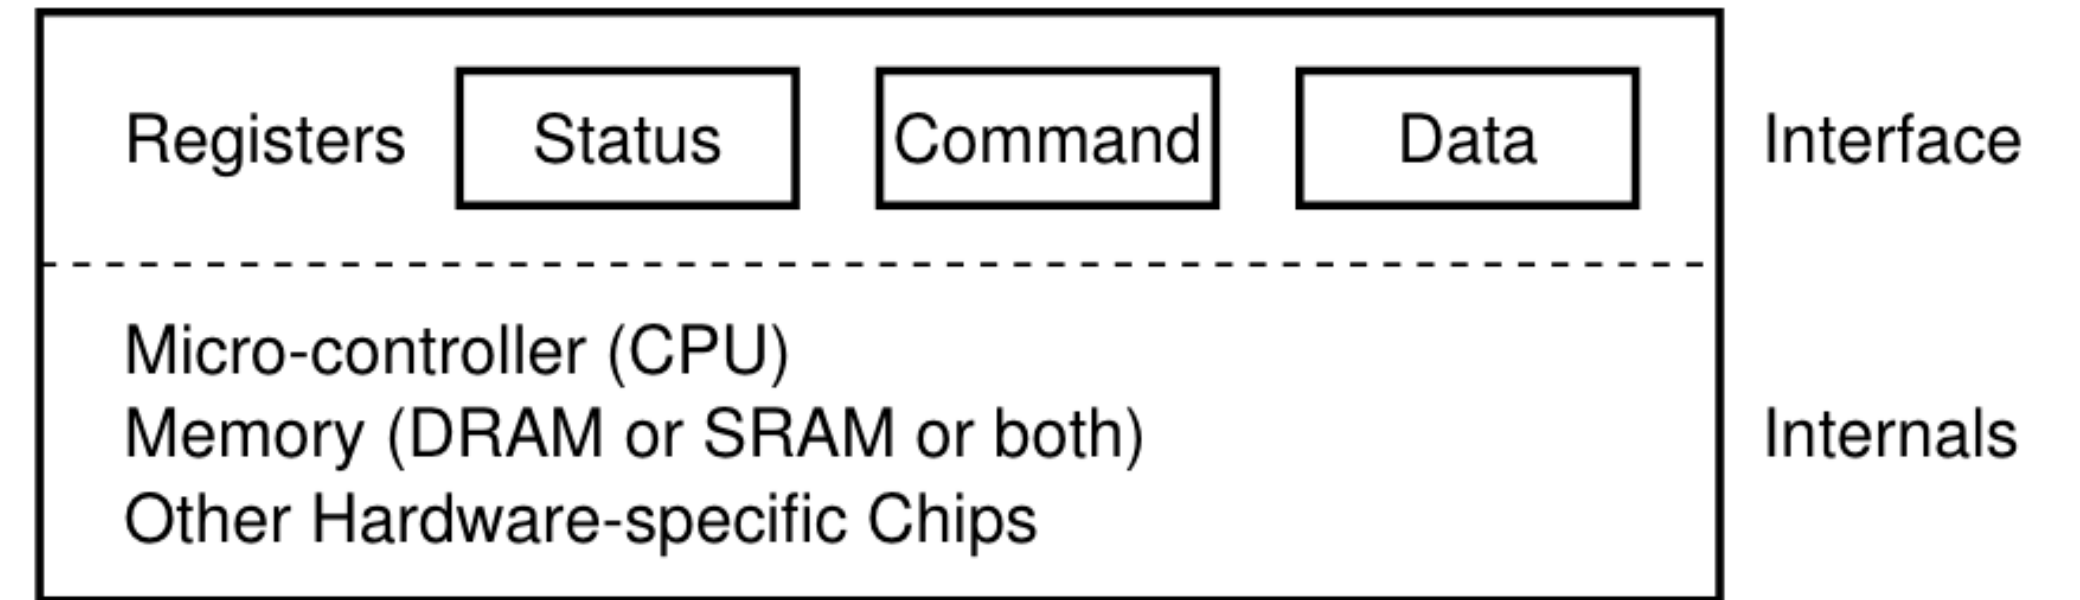


Figure 36.3: A Canonical Device

bootmain.c

```
void waitdisk(void){
    // Wait for disk ready.
    while((inb(0x1F7) & 0xC0) != 0x40);
}

// Read a single sector at offset into dst.
void readsect(void *dst, uint offset) {
    // Issue command.
    waitdisk();
    outb(0x1F2, 1);    // count = 1
    ..
    outb(0x1F7, 0x20); // cmd 0x20 - read sectors

    // Read data.
    waitdisk();
    insl(0x1F0, dst, SECTSIZE/4);
}
```

Canonical protocol

- Poll device until it is ready
- CPU cannot do anything else.
- Example: CPU needs to spend ~1 million instructions waiting for disk
- Ok for bootloader. It does not have anything else to do.

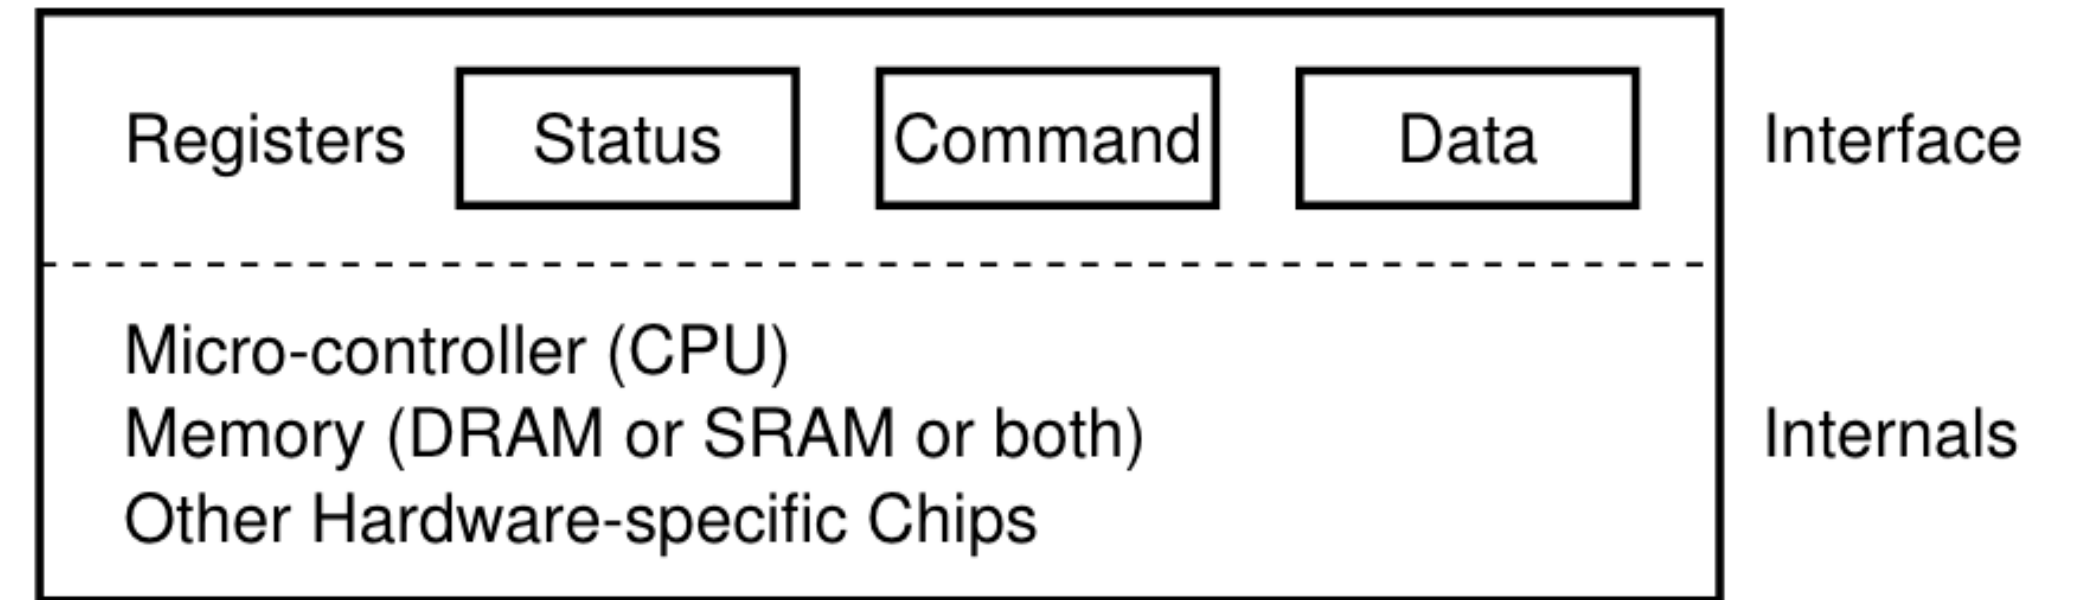


Figure 36.3: A Canonical Device

bootmain.c

```
void waitdisk(void){
    // Wait for disk ready.
    while((inb(0x1F7) & 0xC0) != 0x40);
}

// Read a single sector at offset into dst.
void readsect(void *dst, uint offset) {
    // Issue command.
    waitdisk();
    outb(0x1F2, 1);    // count = 1
    ..
    outb(0x1F7, 0x20); // cmd 0x20 - read sectors

    // Read data.
    waitdisk();
    insl(0x1F0, dst, SECTSIZE/4);
}
```


Canonical protocol

- Poll device until it is ready
- CPU cannot do anything else.
- Example: CPU needs to spend ~1 million instructions waiting for disk
- Ok for bootloader. It does not have anything else to do.
- Not ok for OS. It can run other processes.

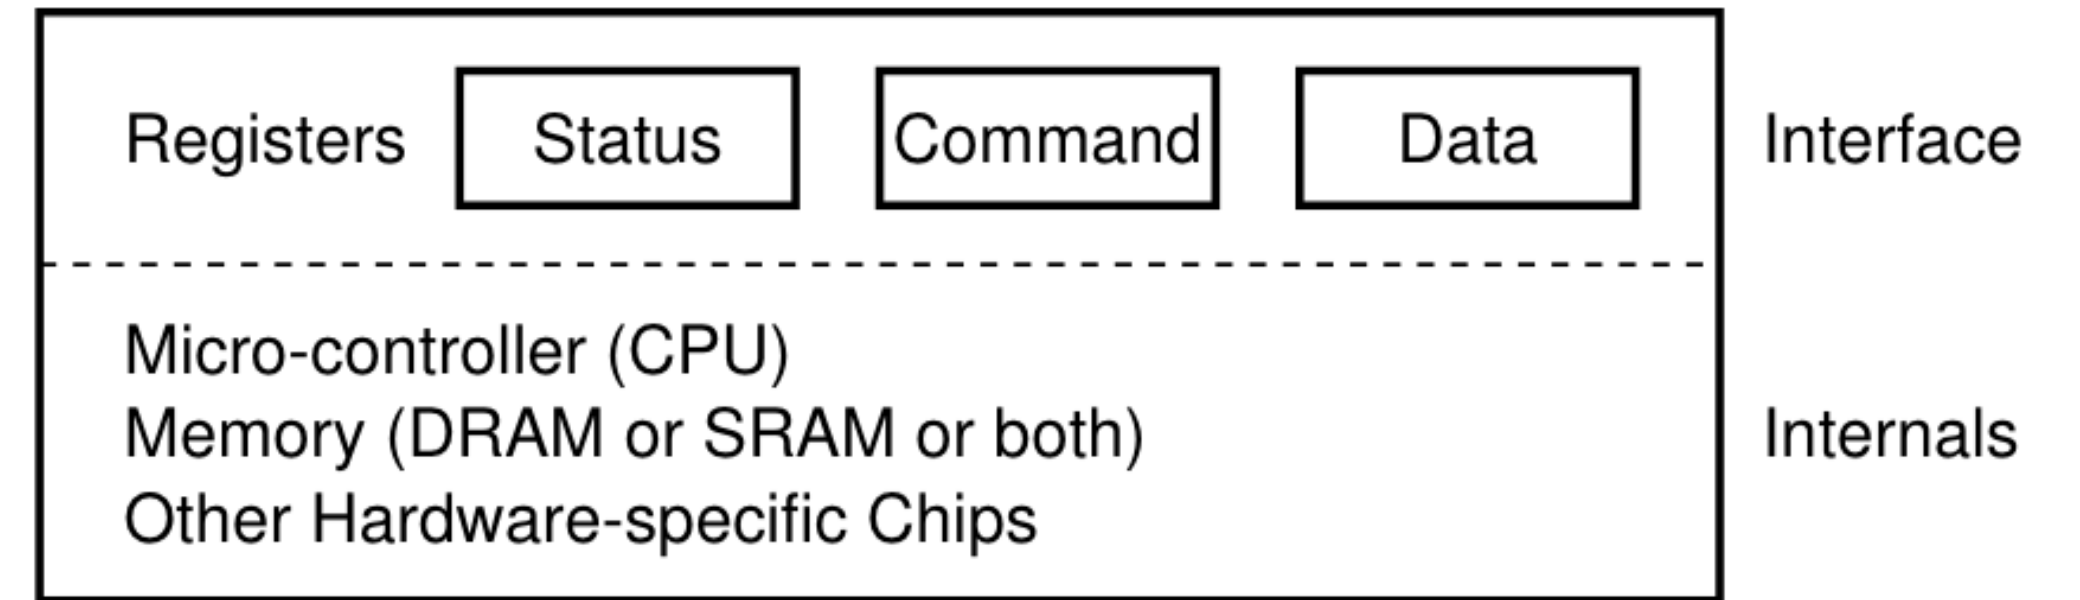


Figure 36.3: A Canonical Device

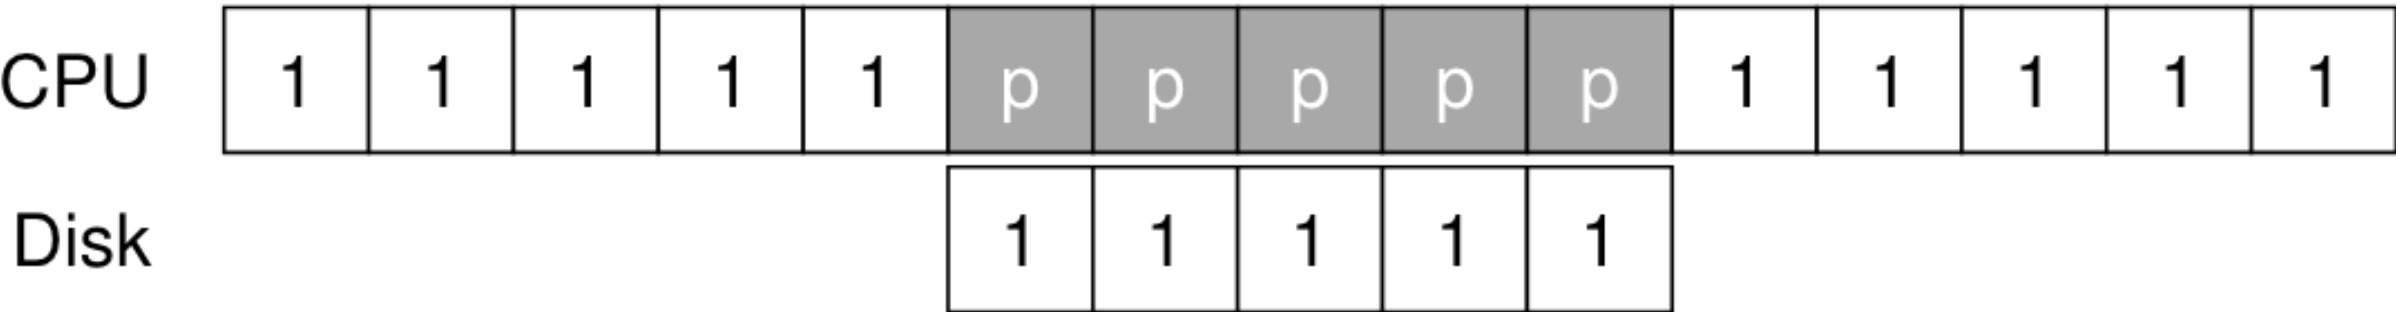
bootmain.c

```
void waitdisk(void){
    // Wait for disk ready.
    while((inb(0x1F7) & 0xC0) != 0x40);
}

// Read a single sector at offset into dst.
void readsect(void *dst, uint offset) {
    // Issue command.
    waitdisk();
    outb(0x1F2, 1);    // count = 1
    ..
    outb(0x1F7, 0x20); // cmd 0x20 - read sectors

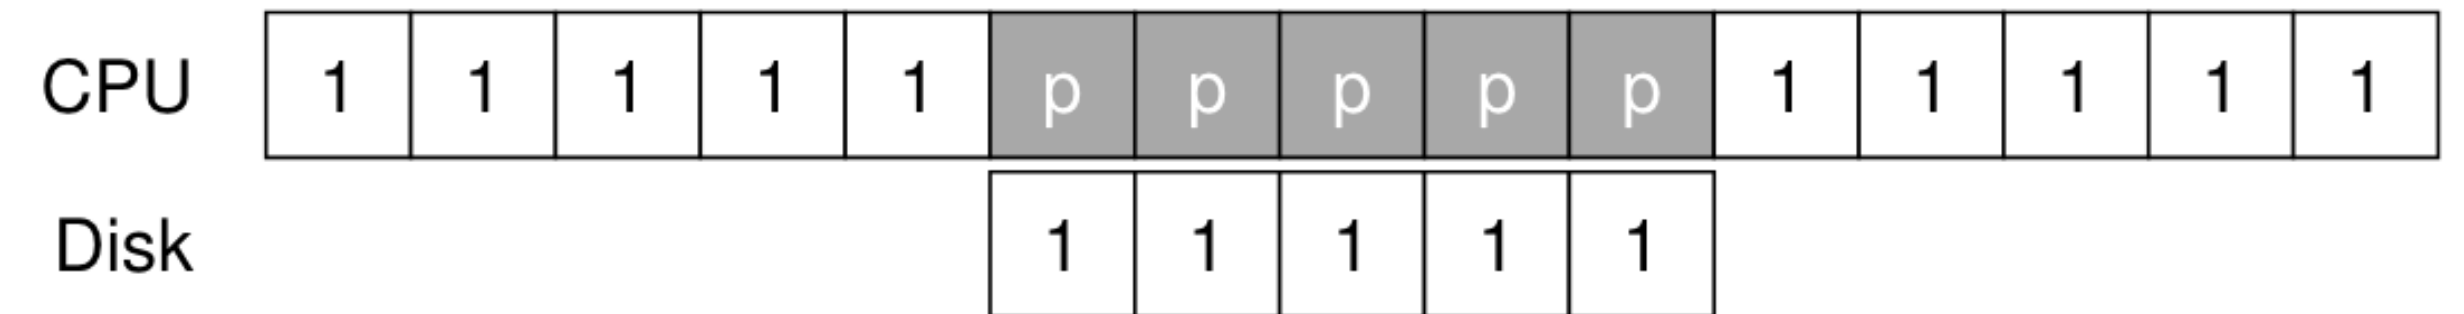
    // Read data.
    waitdisk();
    insl(0x1F0, dst, SECTSIZE/4);
}
```

Lowering CPU overheads with interrupts



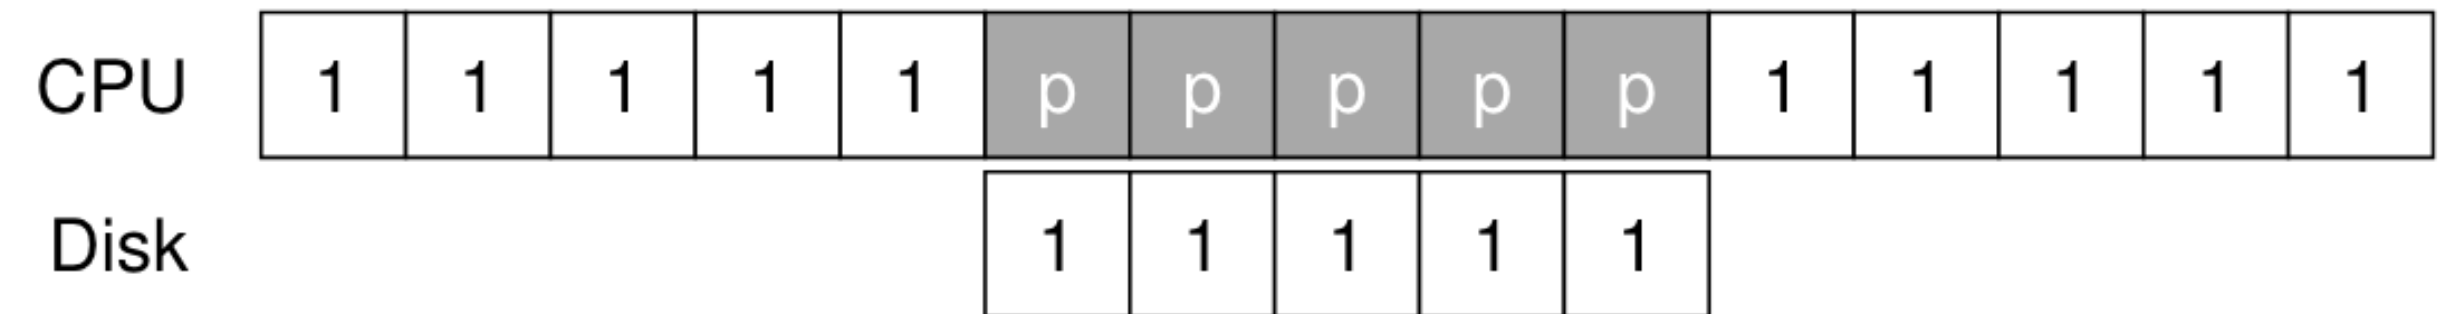
Lowering CPU overheads with interrupts

- Device sends an interrupt that it is ready



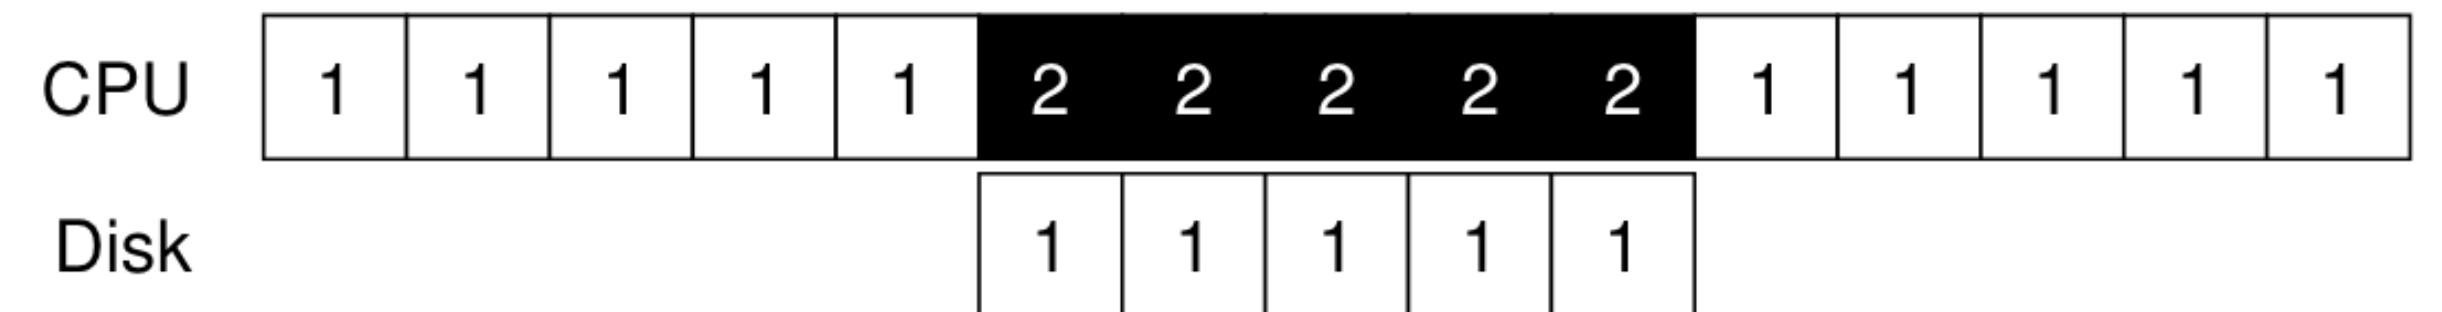
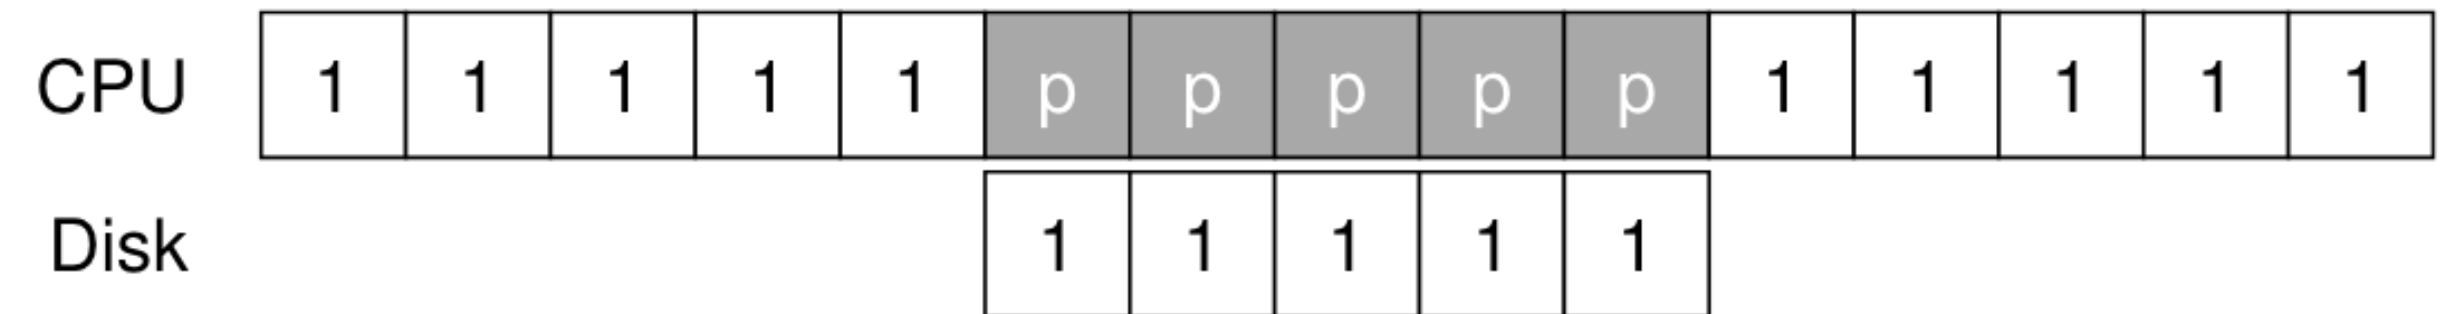
Lowering CPU overheads with interrupts

- Device sends an interrupt that it is ready
- CPU runs another process in the meantime



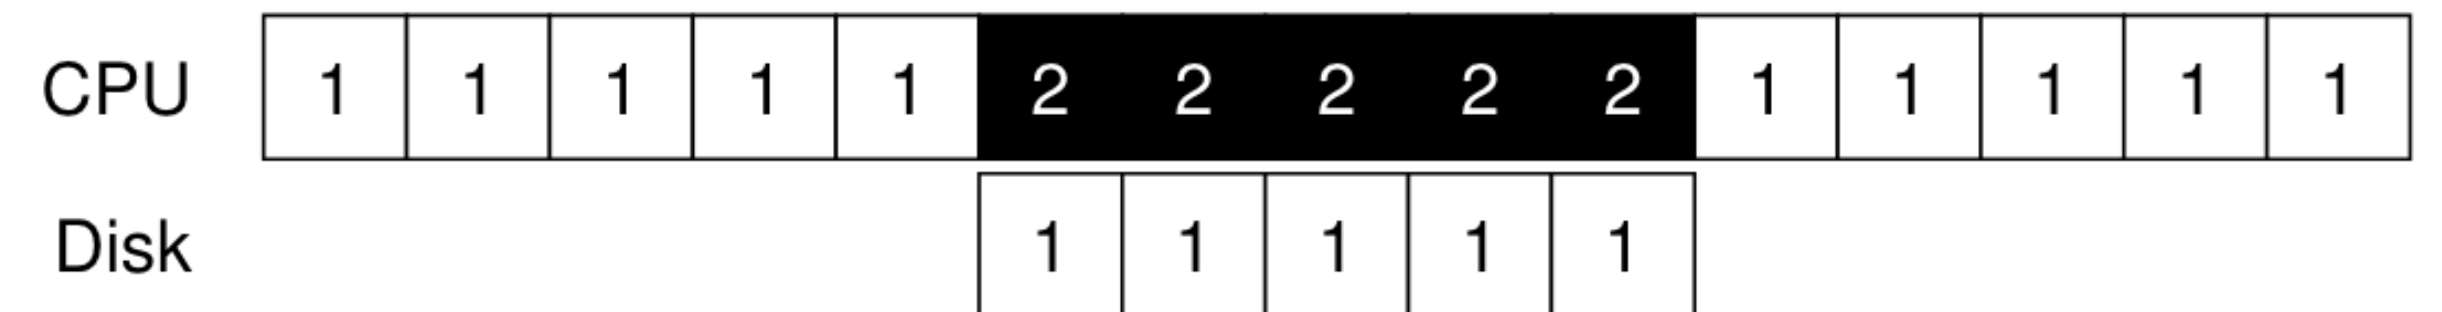
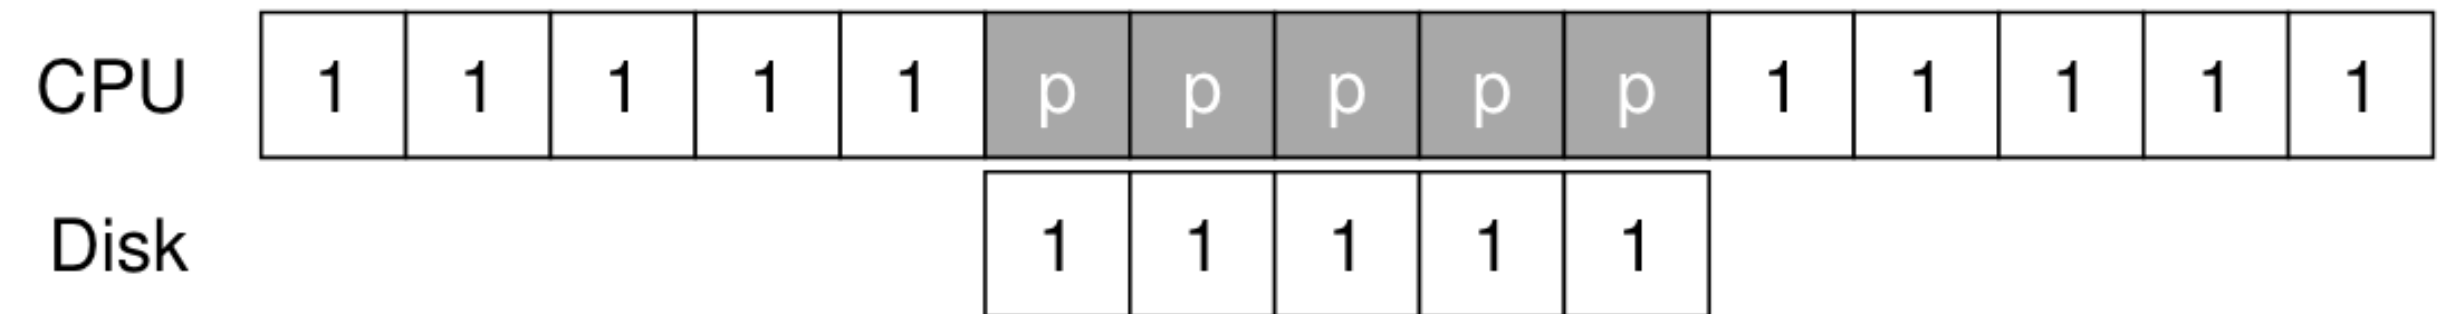
Lowering CPU overheads with interrupts

- Device sends an interrupt that it is ready
- CPU runs another process in the meantime



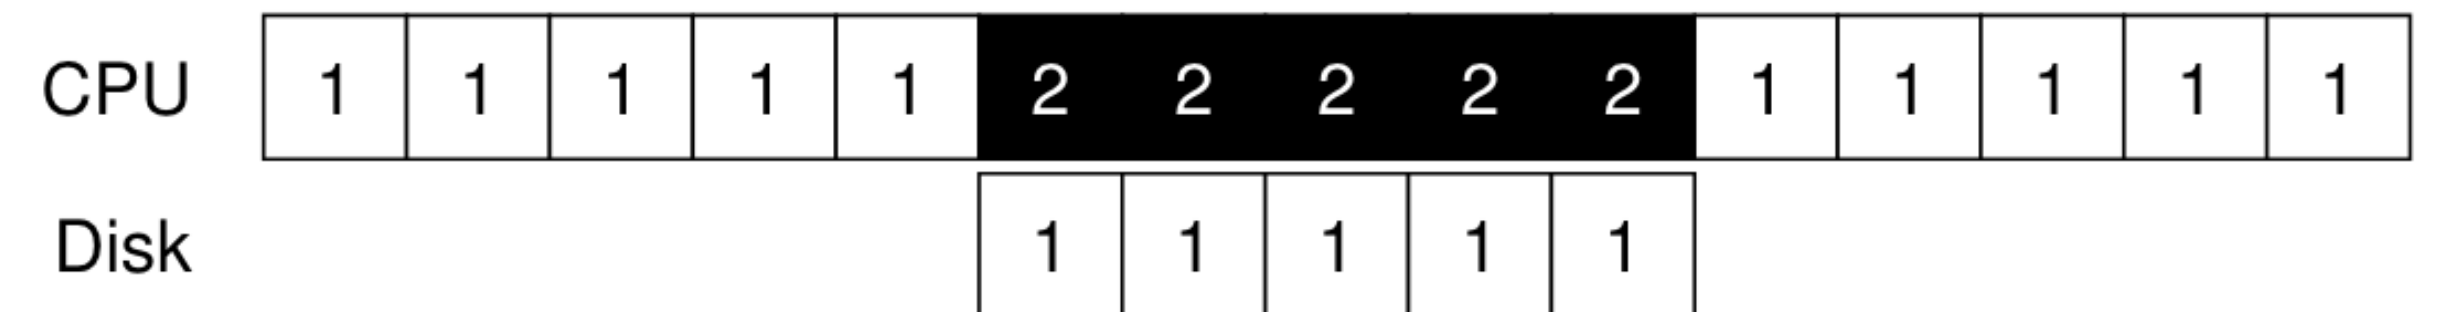
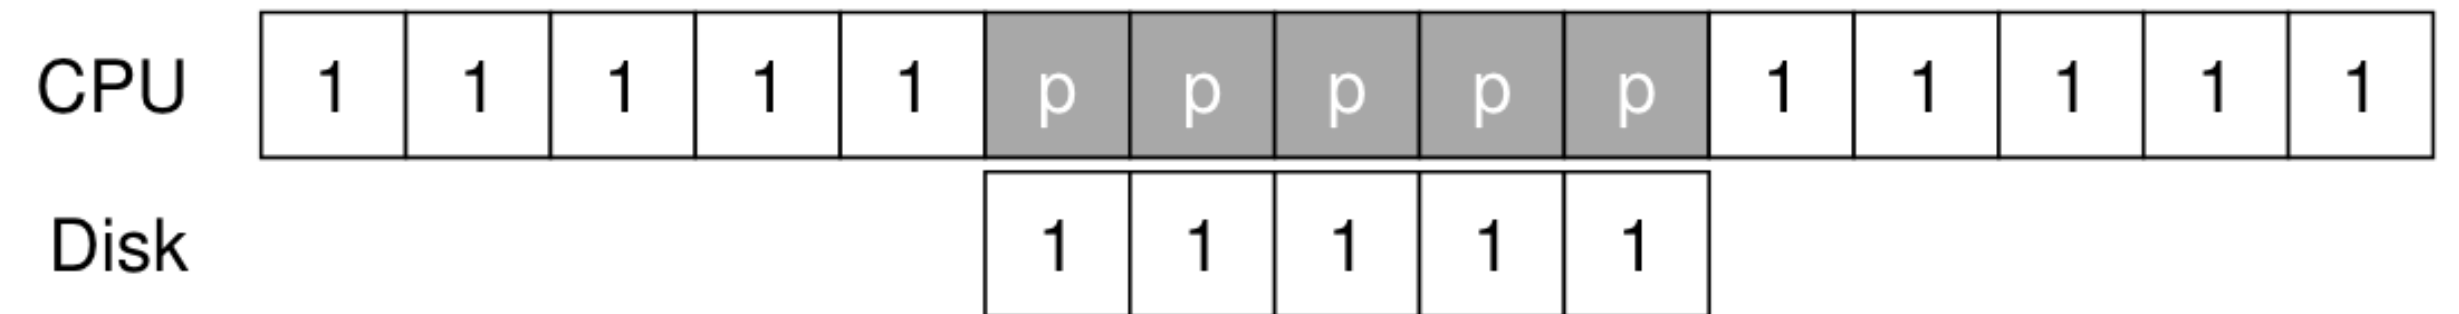
Lowering CPU overheads with interrupts

- Device sends an interrupt that it is ready
- CPU runs another process in the meantime
- Better CPU utilisation



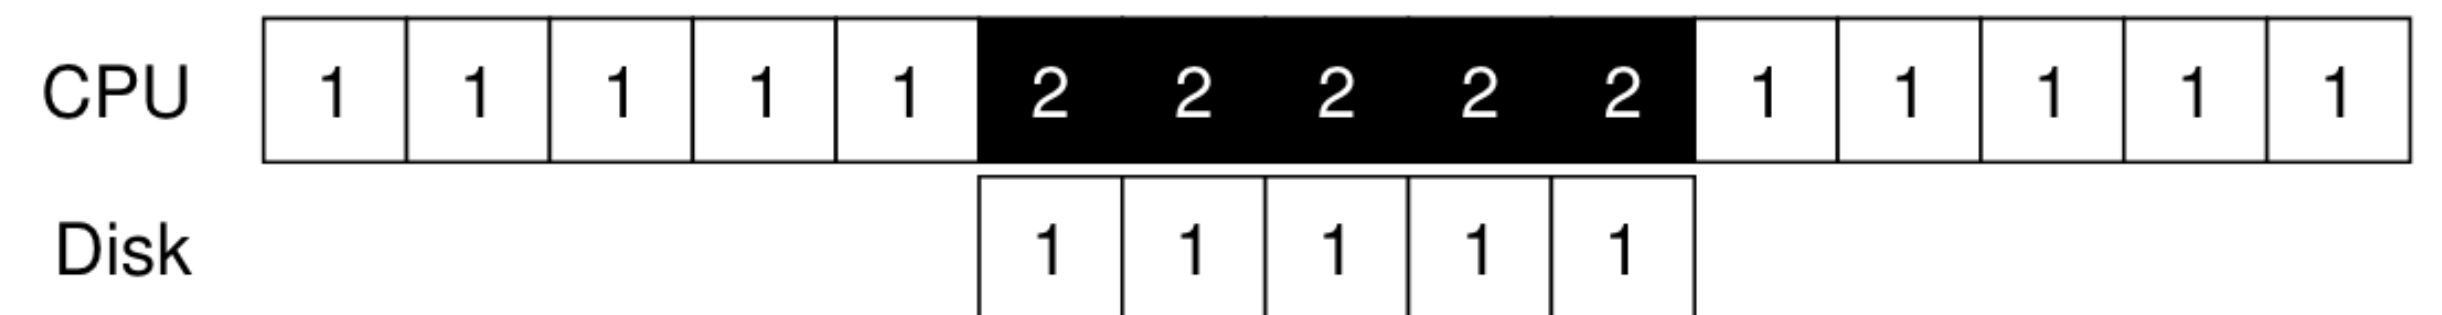
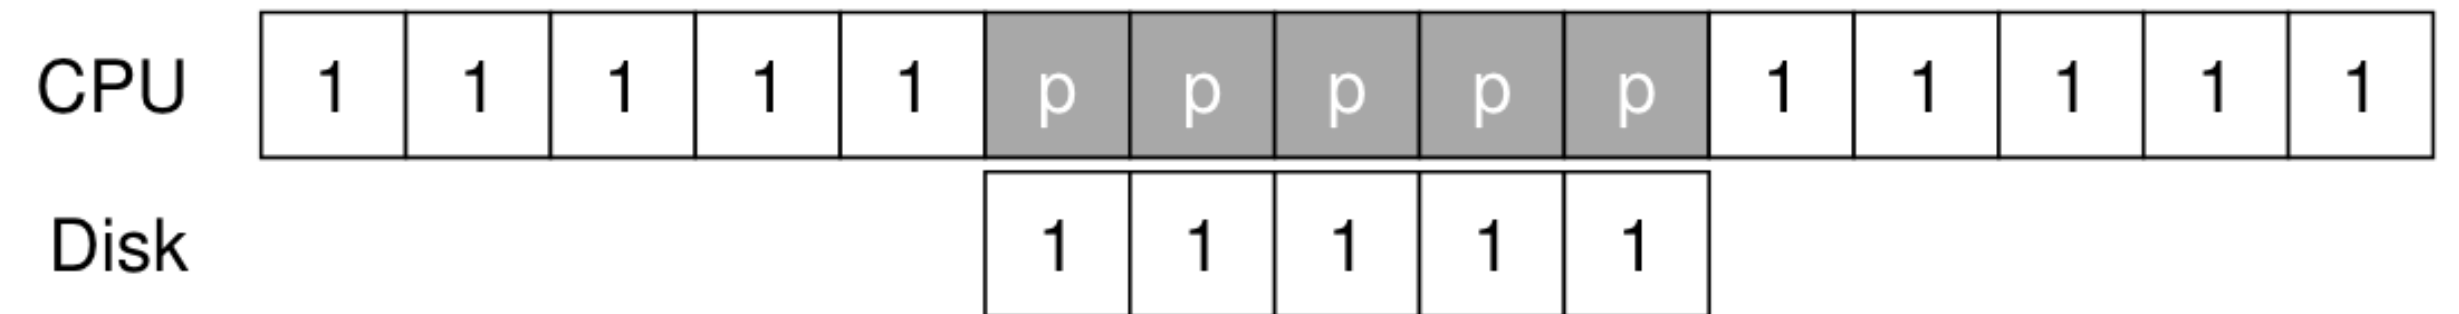
Lowering CPU overheads with interrupts

- Device sends an interrupt that it is ready
- CPU runs another process in the meantime
- Better CPU utilisation
- Not a good idea if device is fast.



Lowering CPU overheads with interrupts

- Device sends an interrupt that it is ready
- CPU runs another process in the meantime
- Better CPU utilisation
- Not a good idea if device is fast.
 - If first poll finds that the device is ready, unnecessary overhead of switching processes



More efficient data movement

Direct Memory Access (DMA)

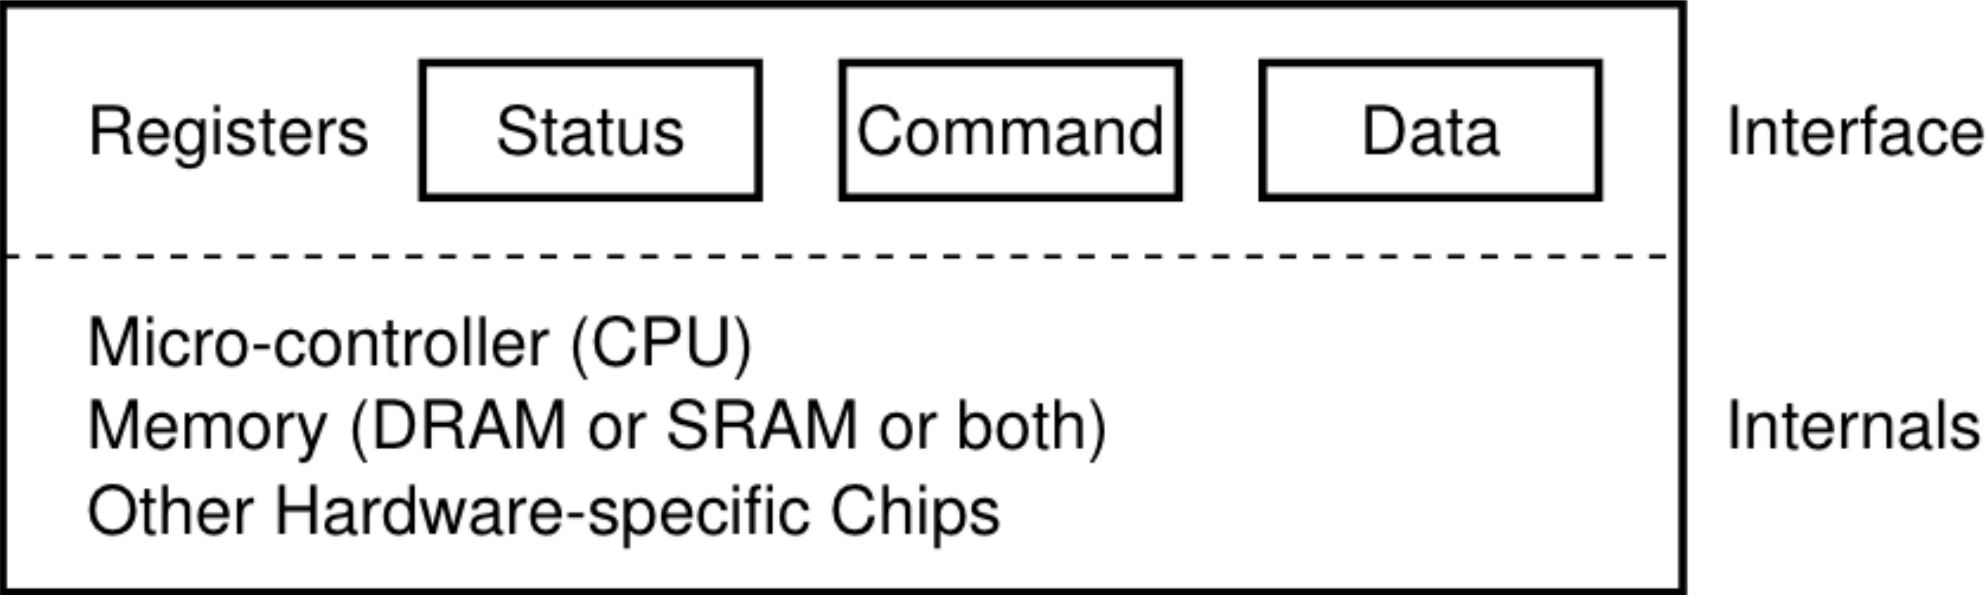


Figure 36.3: A Canonical Device

More efficient data movement

Direct Memory Access (DMA)

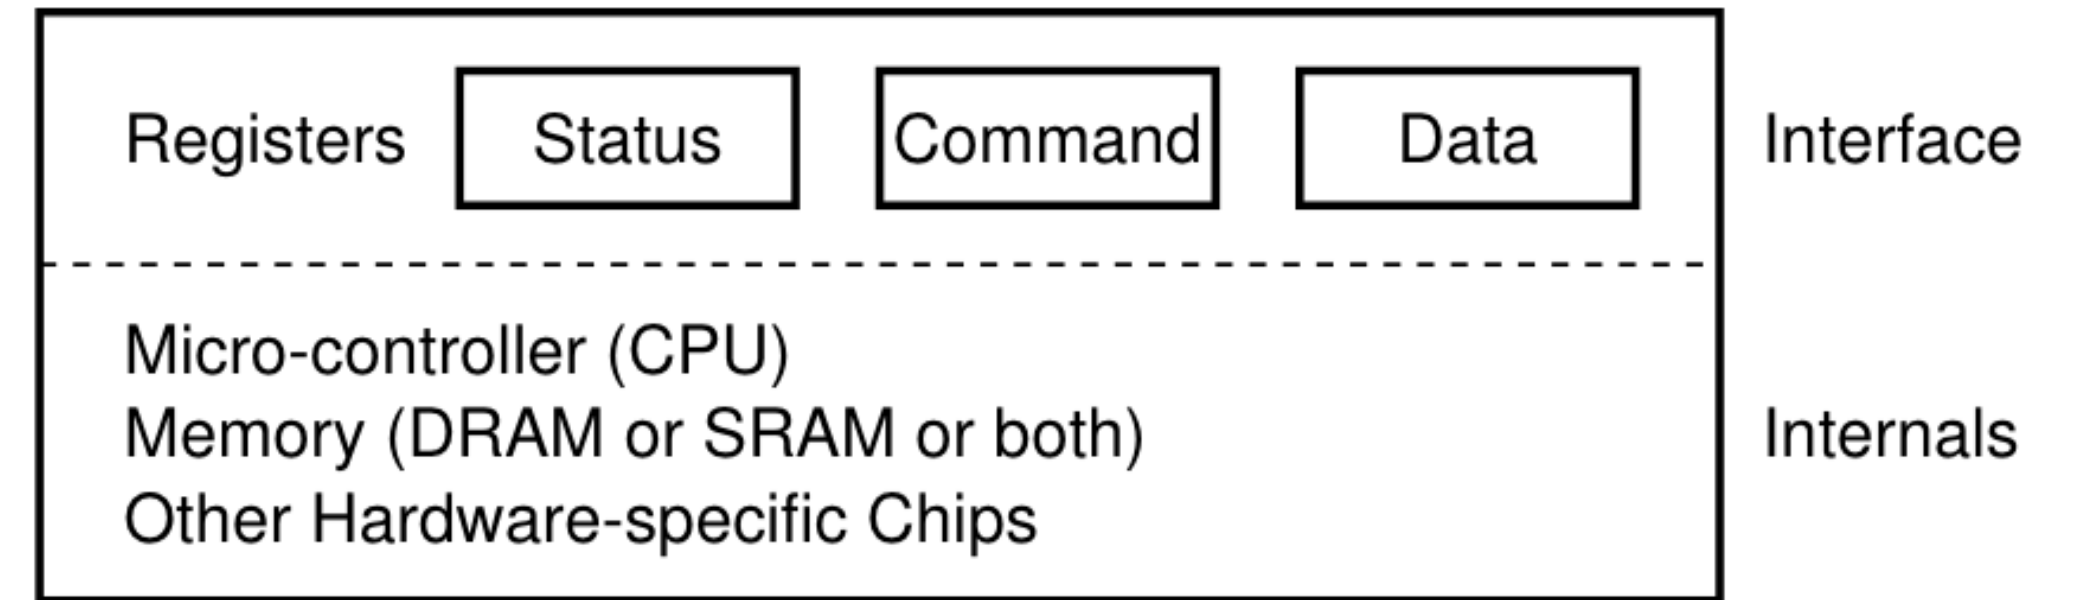


Figure 36.3: A Canonical Device

bootmain.c

```
void waitdisk(void){
    // Wait for disk ready.
    while((inb(0x1F7) & 0xC0) != 0x40);
}

// Read a single sector at offset into dst.
void readsect(void *dst, uint offset) {
    // Issue command.
    waitdisk();
    outb(0x1F2, 1);    // count = 1
    ..
    outb(0x1F7, 0x20); // cmd 0x20 - read sectors

    // Read data.
    waitdisk();
    insl(0x1F0, dst, SECTSIZE/4);
}
```


More efficient data movement

Direct Memory Access (DMA)

CPU	1	1	1	2	2	c	c	c	1
Disk				1	1				

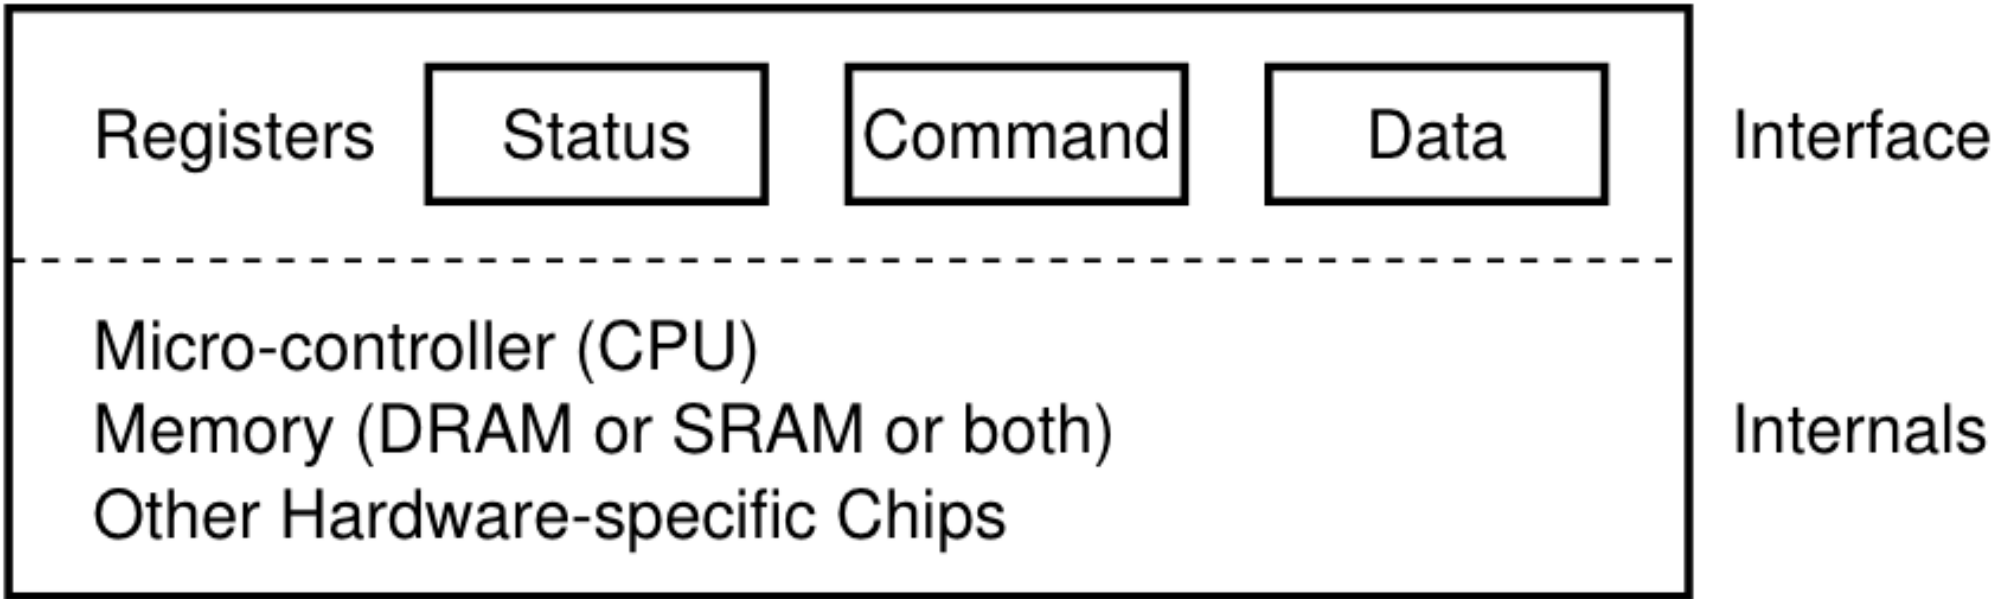


Figure 36.3: A Canonical Device

bootmain.c

```
void waitdisk(void){
    // Wait for disk ready.
    while((inb(0x1F7) & 0xC0) != 0x40);
}

// Read a single sector at offset into dst.
void readsect(void *dst, uint offset) {
    // Issue command.
    waitdisk();
    outb(0x1F2, 1);    // count = 1
    ..
    outb(0x1F7, 0x20); // cmd 0x20 - read sectors

    // Read data.
    waitdisk();
    insl(0x1F0, dst, SECTSIZE/4);
}
```

More efficient data movement

Direct Memory Access (DMA)

CPU	1	1	1	2	2	c	c	c	1
Disk				1	1				

CPU	1	1	1	2	2	2	2	2	1
DMA						c	c	c	
Disk				1	1				

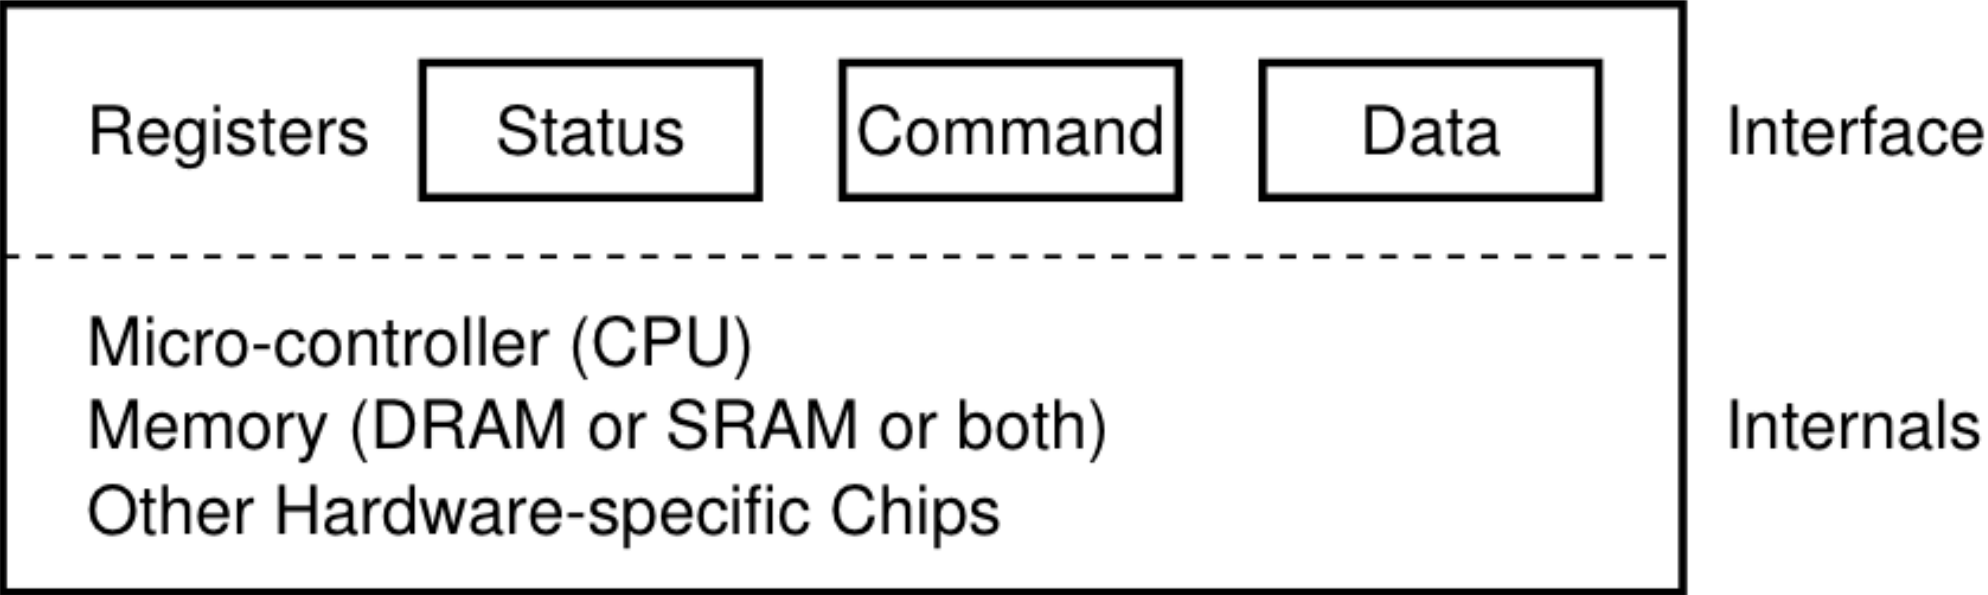


Figure 36.3: A Canonical Device

bootmain.c

```
void waitdisk(void){
    // Wait for disk ready.
    while((inb(0x1F7) & 0xC0) != 0x40);
}

// Read a single sector at offset into dst.
void readsect(void *dst, uint offset) {
    // Issue command.
    waitdisk();
    outb(0x1F2, 1);    // count = 1
    ..
    outb(0x1F7, 0x20); // cmd 0x20 - read sectors

    // Read data.
    waitdisk();
    insl(0x1F0, dst, SECTSIZE/4);
}
```

Interrupt controllers, interrupt handling

xv6 Ch. 3 “Code: interrupts”

Calculator analogy



20
10
30
50
30
10
20
10



Calculator analogy



20
10
30
50
30
10
20
10

- 2 0 = (move pointer to 10)



20
10
30
50
30
10
20
10

Calculator analogy



20
10
30
50
30
10
20
10

- 2 0 = (move pointer to 10)
- + 1 0 = (move pointer to 30)



20
10
30
50
30
10
20
10

Calculator analogy



20
10
30
50
30
10
20
10

- 2 0 = (move pointer to 10)
- + 1 0 = (move pointer to 30)
- + 3 0 = (move pointer to 50)



Calculator analogy



20
10
30
50
30
10
20
10

- 2 0 = (move pointer to 10)
- + 1 0 = (move pointer to 30)
- + 3 0 = (move pointer to 50)

Interrupt

← Give me the calculator!



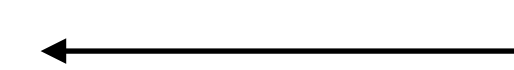
Calculator analogy



20
10
30
50
30
10
20
10

- 2 0 = (move pointer to 10)
- + 1 0 = (move pointer to 30)
- + 3 0 = (move pointer to 50)

Interrupt



Give me the calculator!

- $3 * 2 = 6$



Calculator analogy



20
10
30
50
30
10
20
10

- 2 0 = (move pointer to 10)
- + 1 0 = (move pointer to 30)
- + 3 0 = (move pointer to 50)



Interrupt
← Give me the calculator!

- $3 * 2 = 6$

End of Interrupt
← Ok, you can have it back

Calculator analogy



20
10
30
50
30
10
20
10

- 2 0 = (move pointer to 10)
- + 1 0 = (move pointer to 30)
- + 3 0 = (move pointer to 50)
- + 5 0 = (move pointer to 30)



Interrupt
← Give me the calculator!

• 3*2 = 6

End of Interrupt
← Ok, you can have it back

Calculator analogy



20
10
30
50
30
10
20
10

- 2 0 = (move pointer to 10)
- + 1 0 = (move pointer to 30)
- + 3 0 = (move pointer to 50)
- + 5 0 = (move pointer to 30)
- + 3 0 = (move pointer to 10)



Interrupt
←

Give me the calculator!

- $3 * 2 = 6$

End of Interrupt
←

Ok, you can have it back

Calculator analogy



20
10
30
50
30
10
20
10

- 2 0 = (move pointer to 10)
- + 1 0 = (move pointer to 30)
- + 3 0 = (move pointer to 50)
- + 5 0 = (move pointer to 30)
- + 3 0 = (move pointer to 10)
- + 1 0 = (move pointer to 20)



Interrupt
←

Give me the calculator!

- $3 * 2 = 6$

End of Interrupt
←

Ok, you can have it back

Calculator analogy



20
10
30
50
30
10
20
10

- 2 0 = (move pointer to 10)
- + 1 0 = (move pointer to 30)
- + 3 0 = (move pointer to 50)

Interrupt

← Give me the calculator!

- $3 * 2 = 6$

End of Interrupt

← Ok, you can have it back

- + 5 0 = (move pointer to 30)
- + 3 0 = (move pointer to 10)
- + 1 0 = (move pointer to 20)
- + 2 0 = (move pointer to 10)



Programmable interrupt controllers (PIC)

Example: Intel 8259A

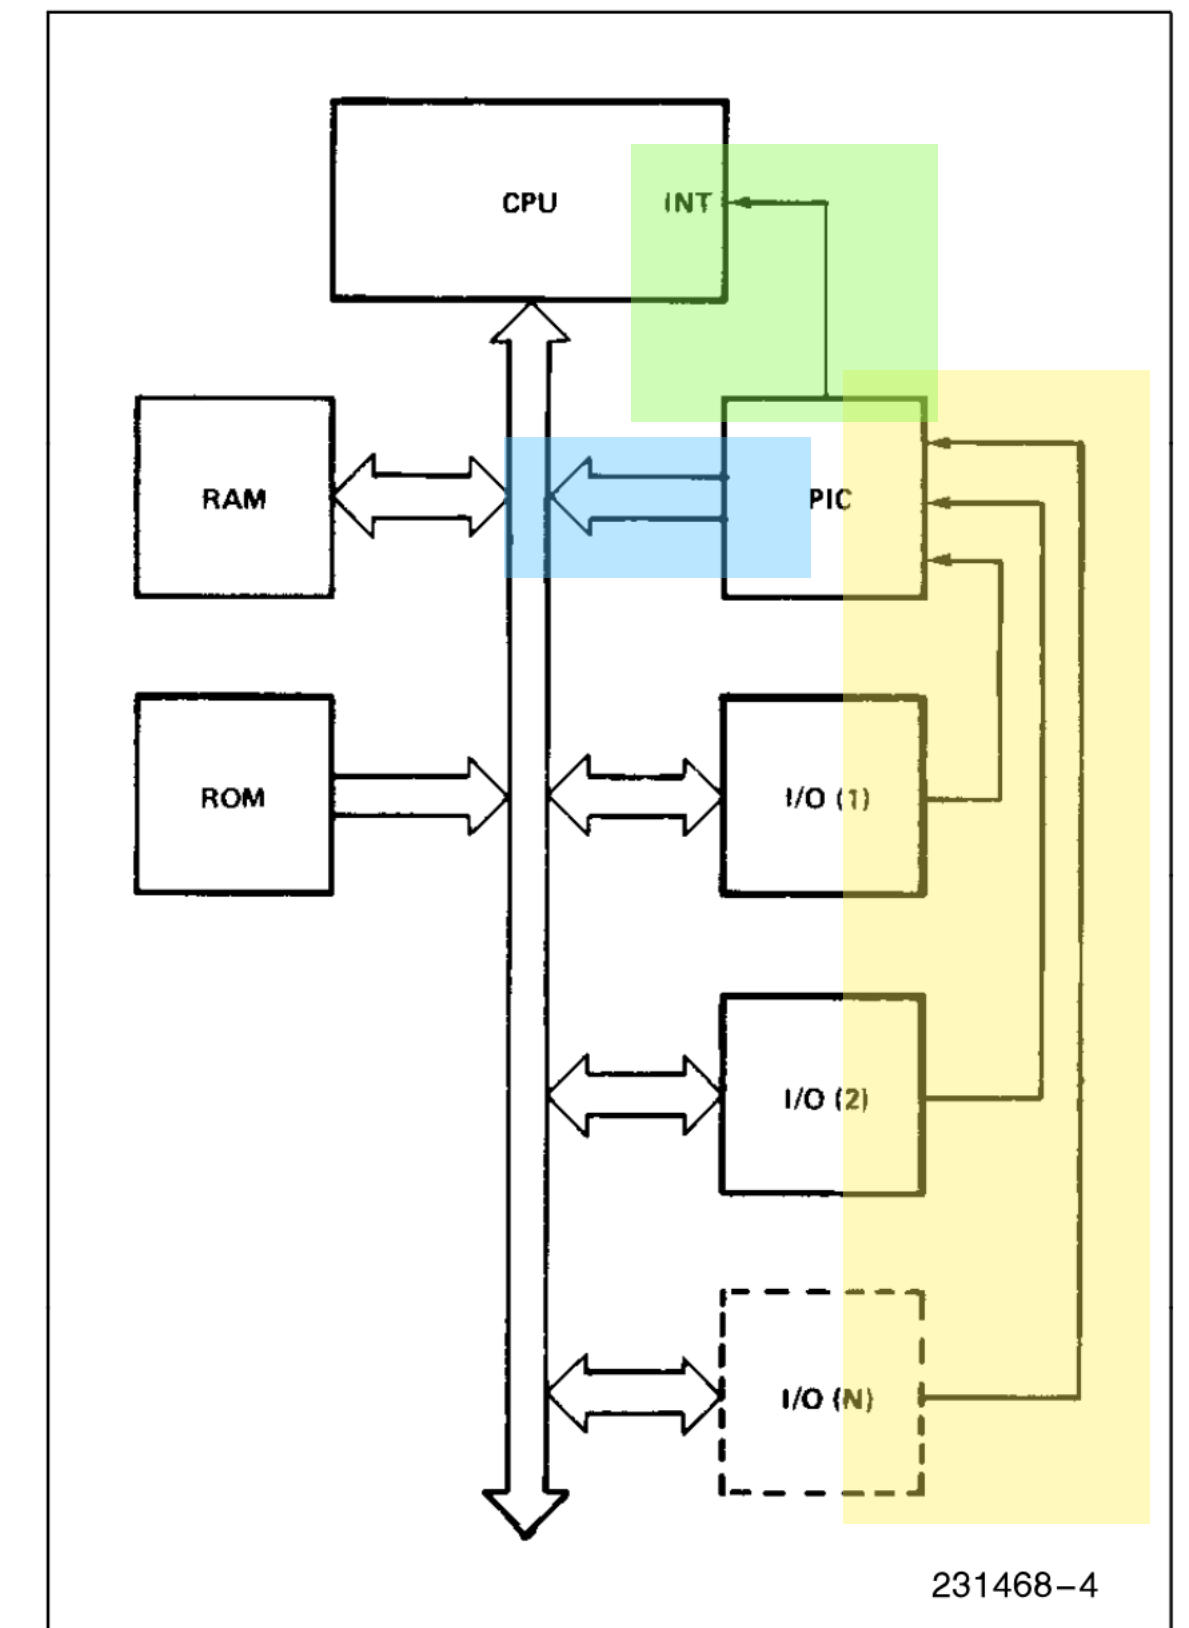
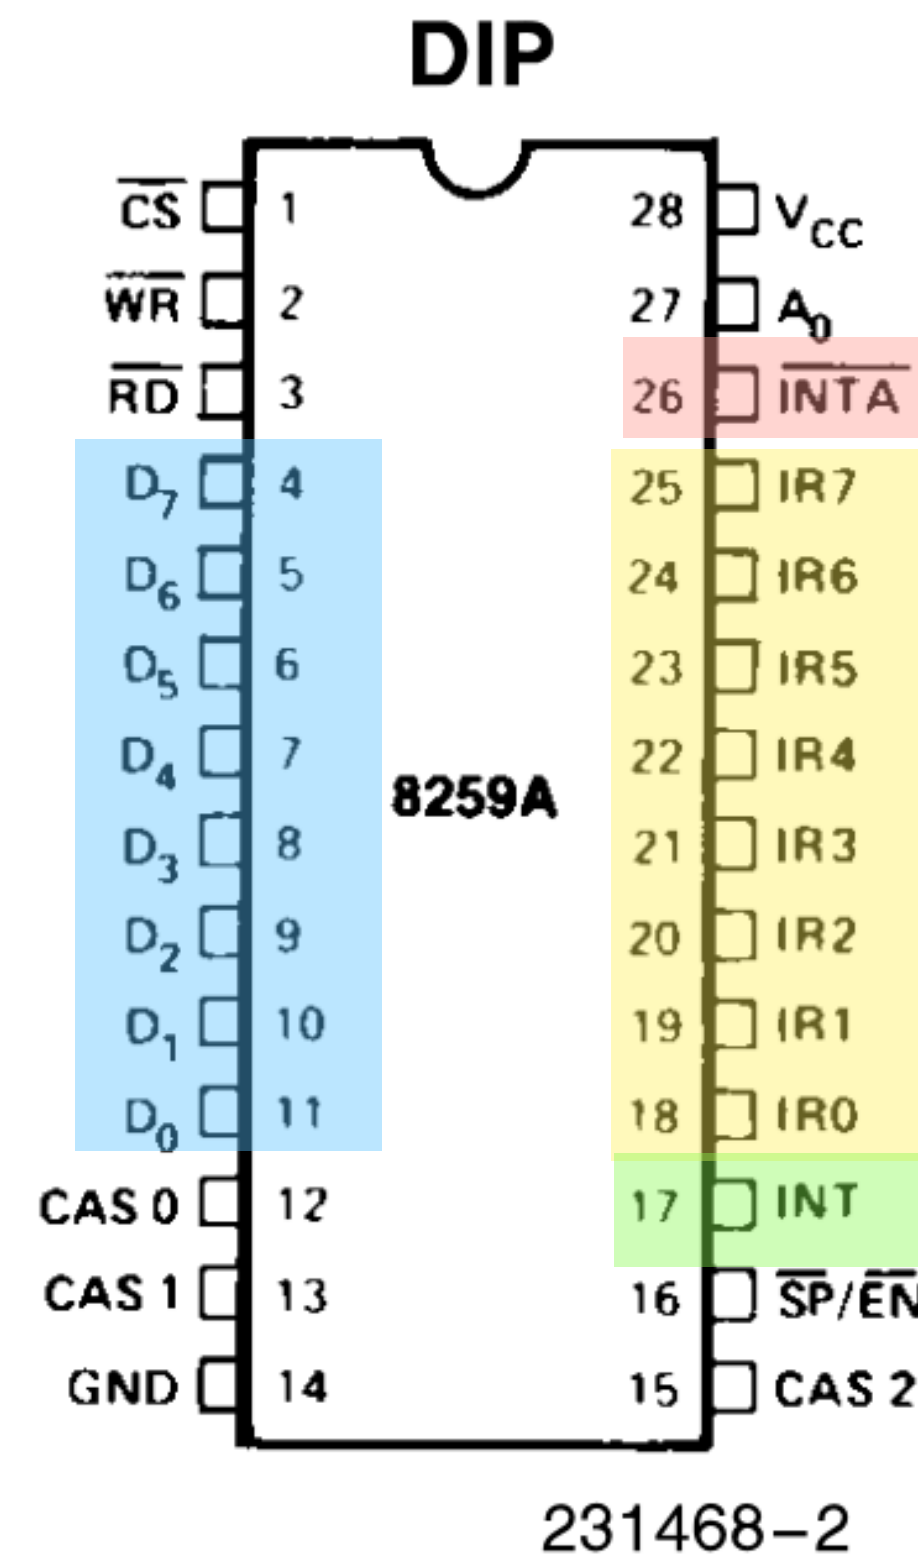


Figure 3b. Interrupt Method

Figure 3b. Interrupt Method

Figure 3b. Interrupt Method

Figure 3b. Interrupt Method

Programmable interrupt controllers (PIC)

Example: Intel 8259A

- Devices connect to **IR0-IR7 pins**.
Device enables its pin to raise interrupt
- **INT pin** connects to CPU.
- PIC sends an 8-bit “interrupt vector” to CPU via **D0-D7 pins**
- CPU acknowledges that it is now working on interrupt on **INTA pin**

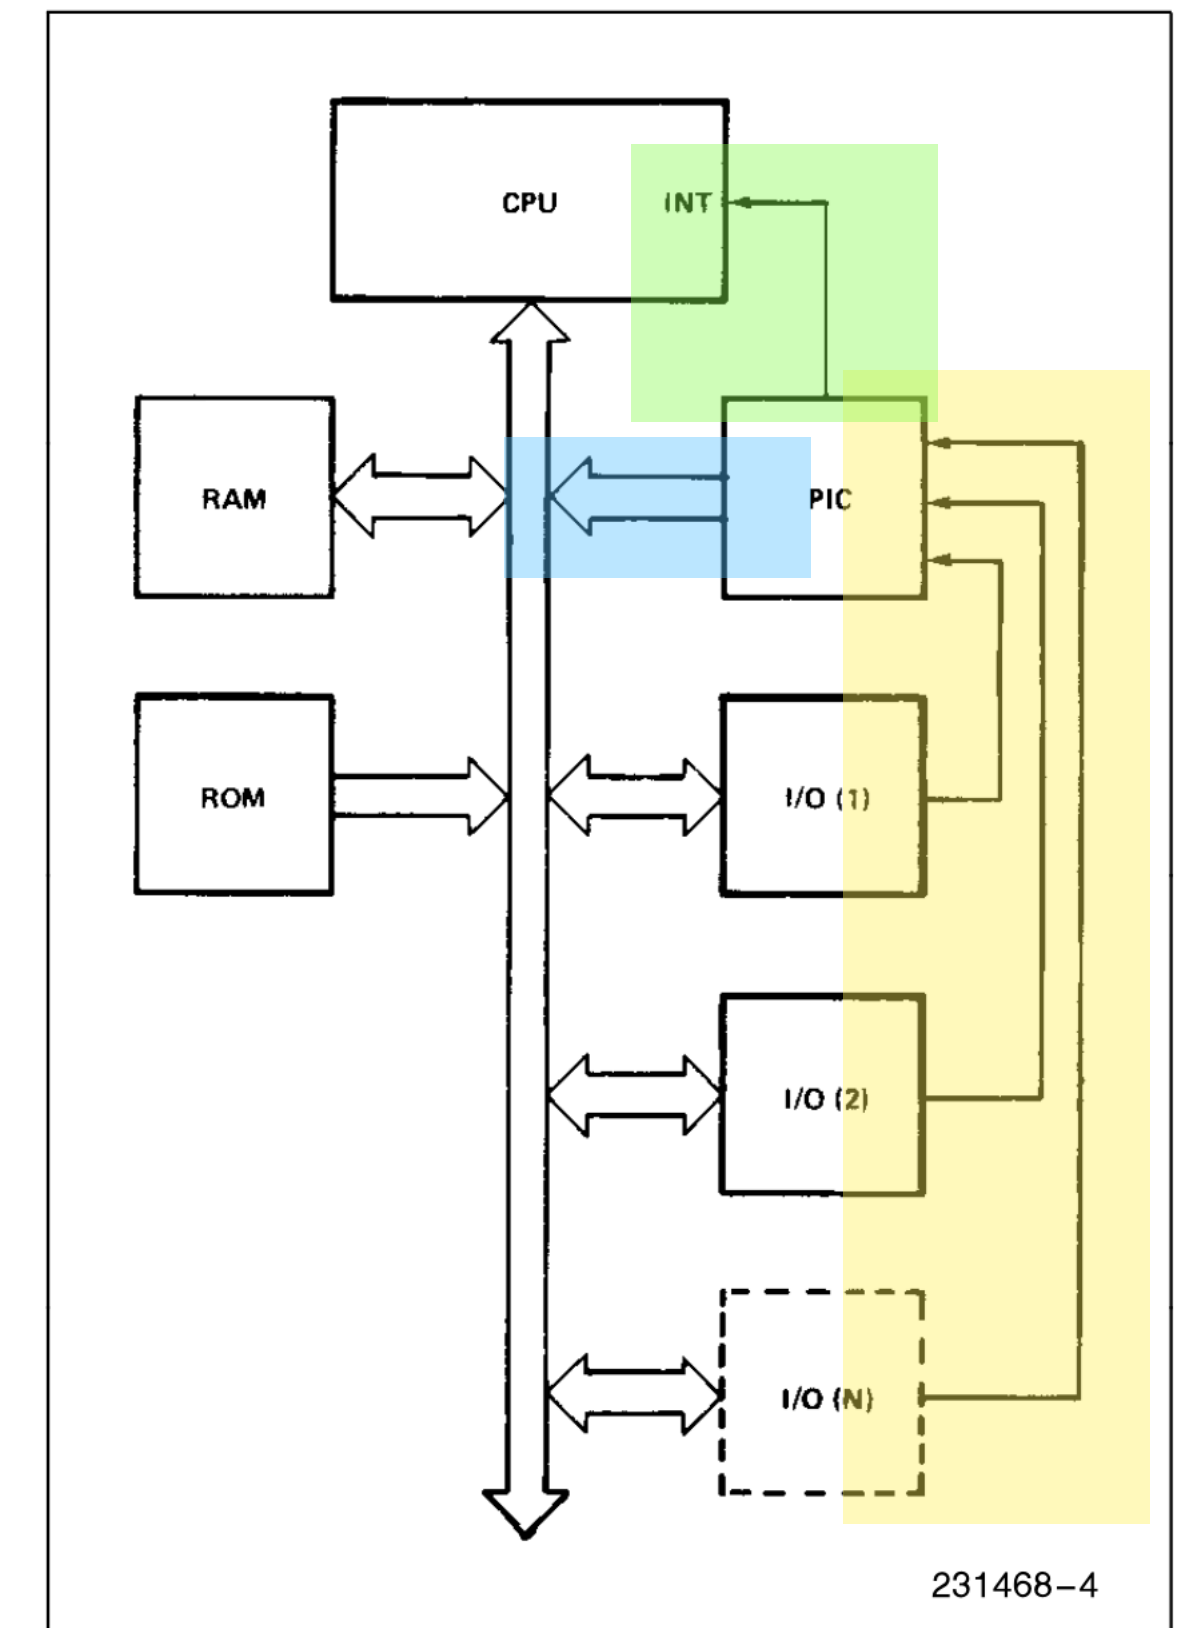
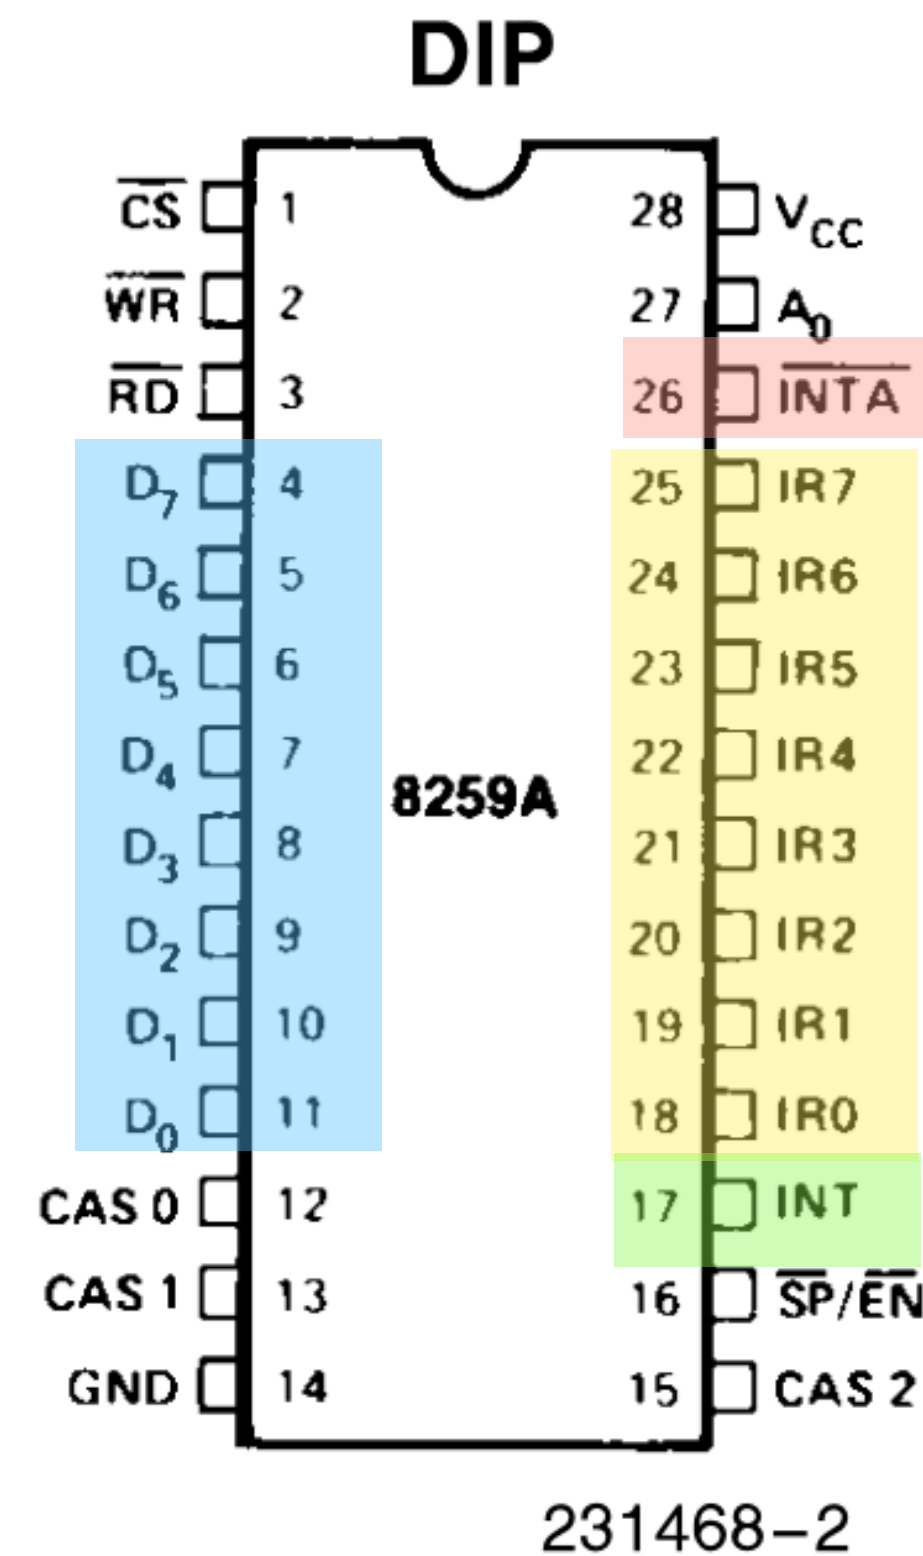
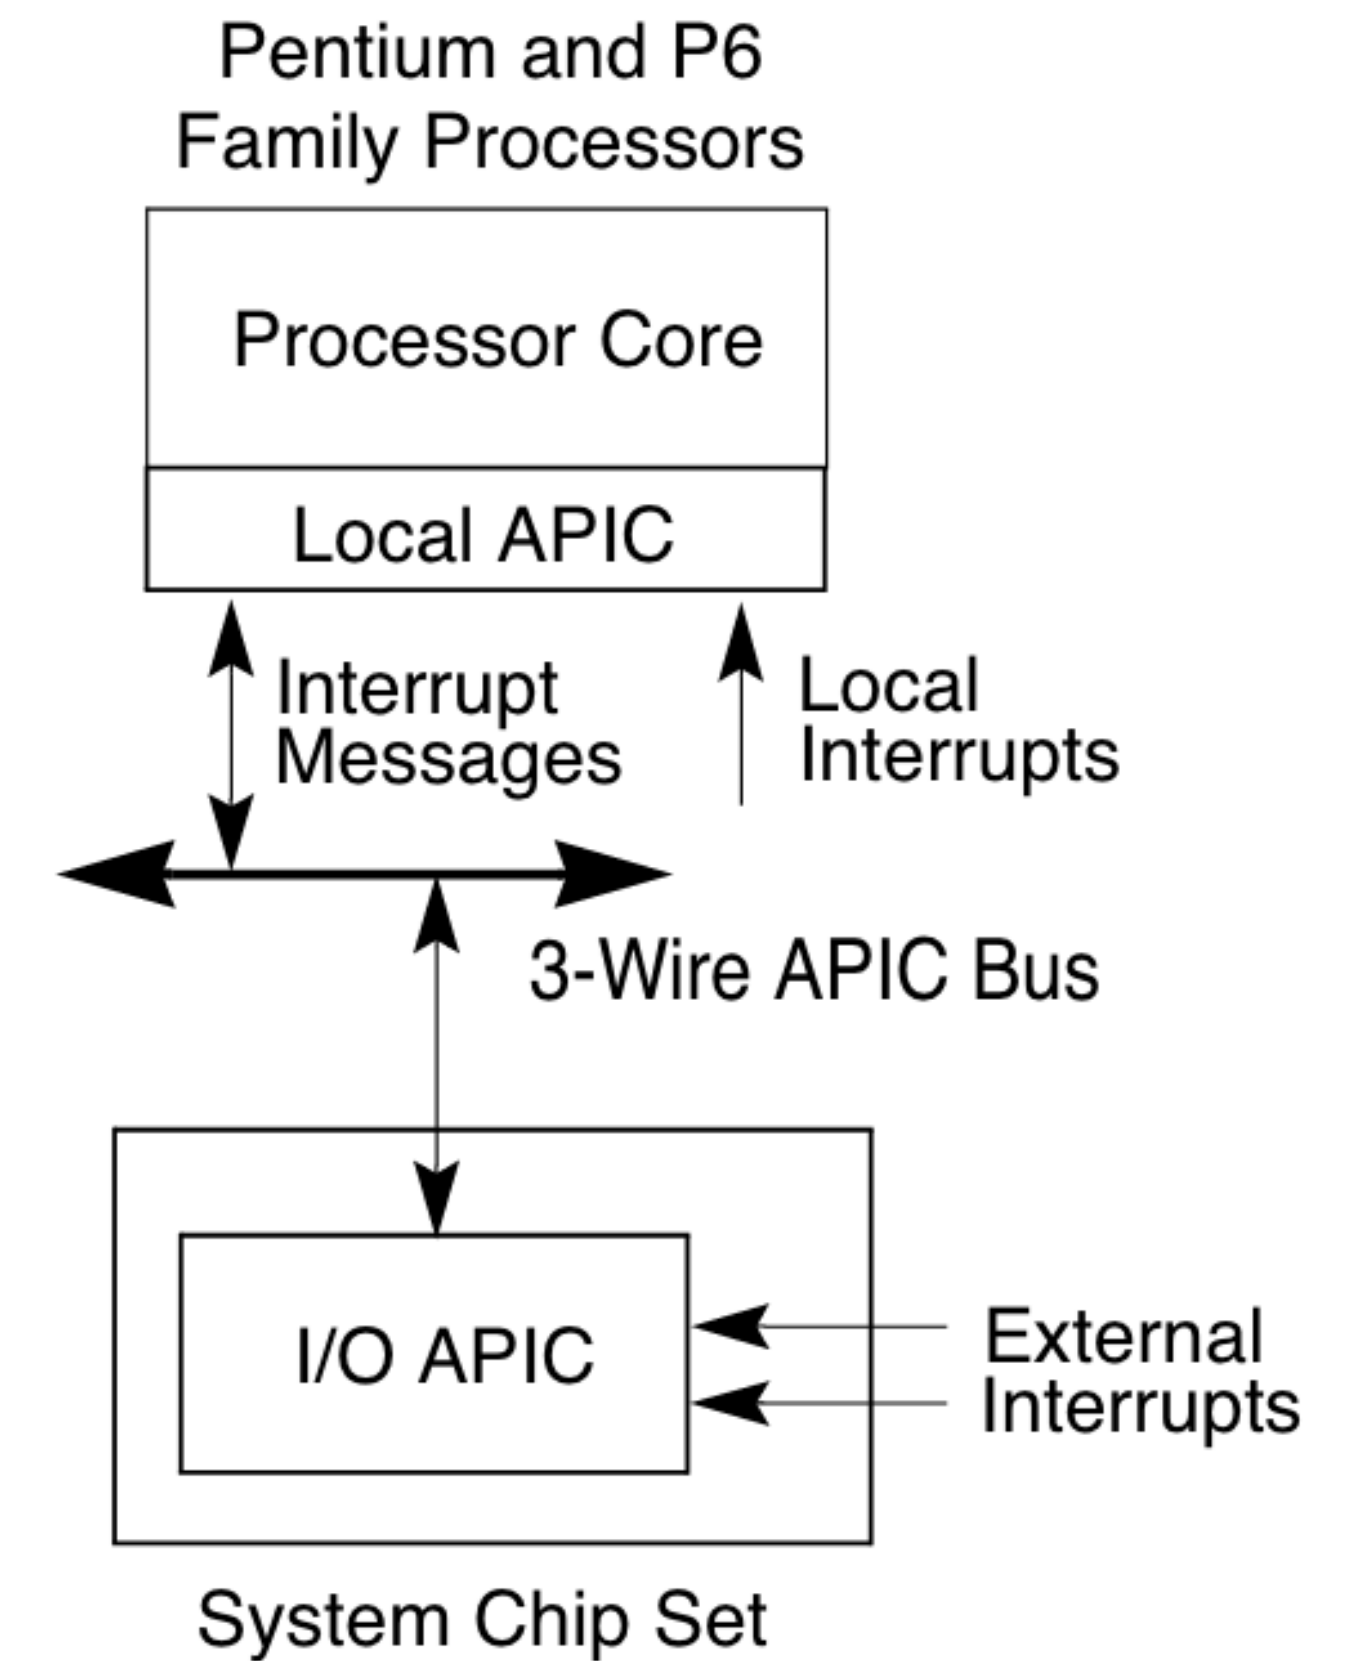


Figure 3b. Interrupt Method

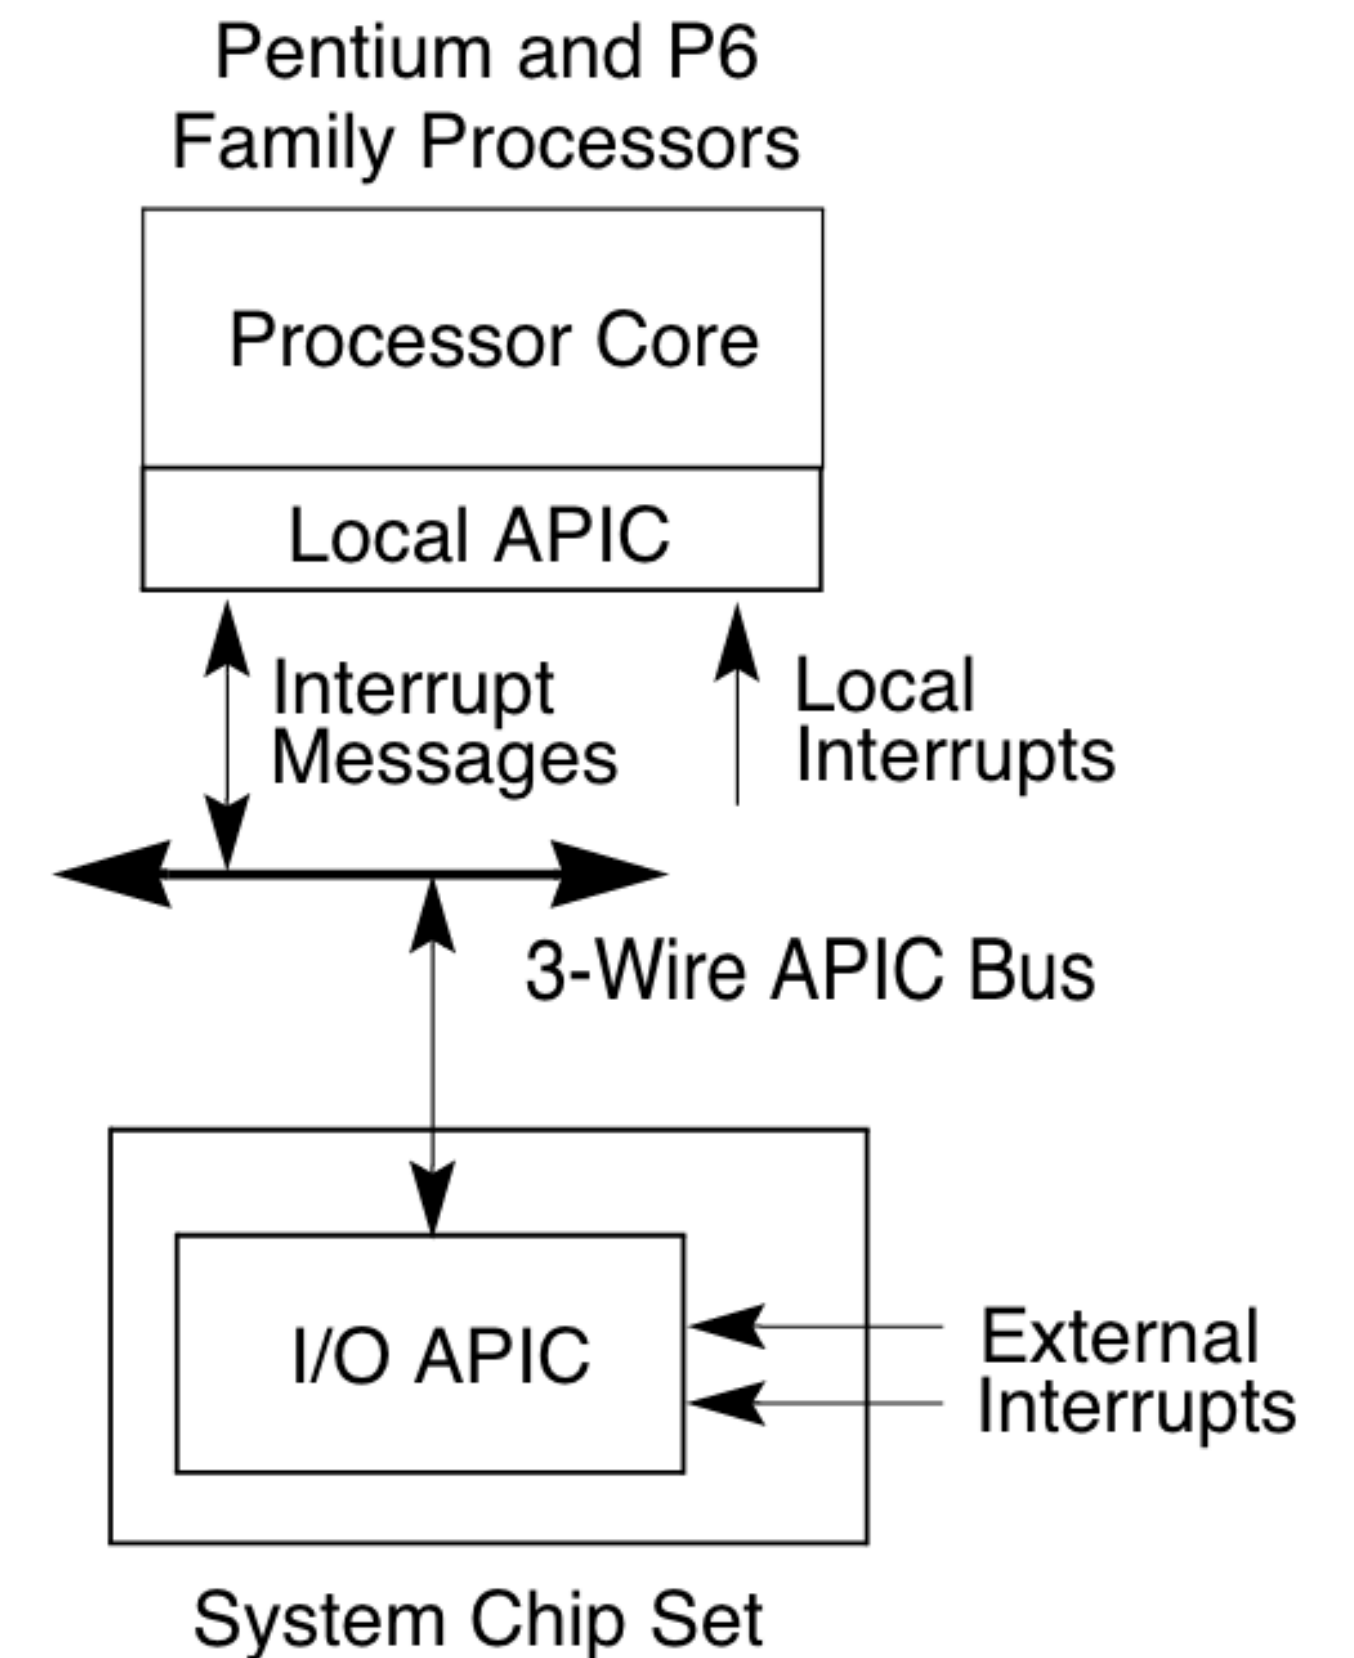
Figure 3b. Interrupt Method

Advanced programmable interrupt controllers (APIC)



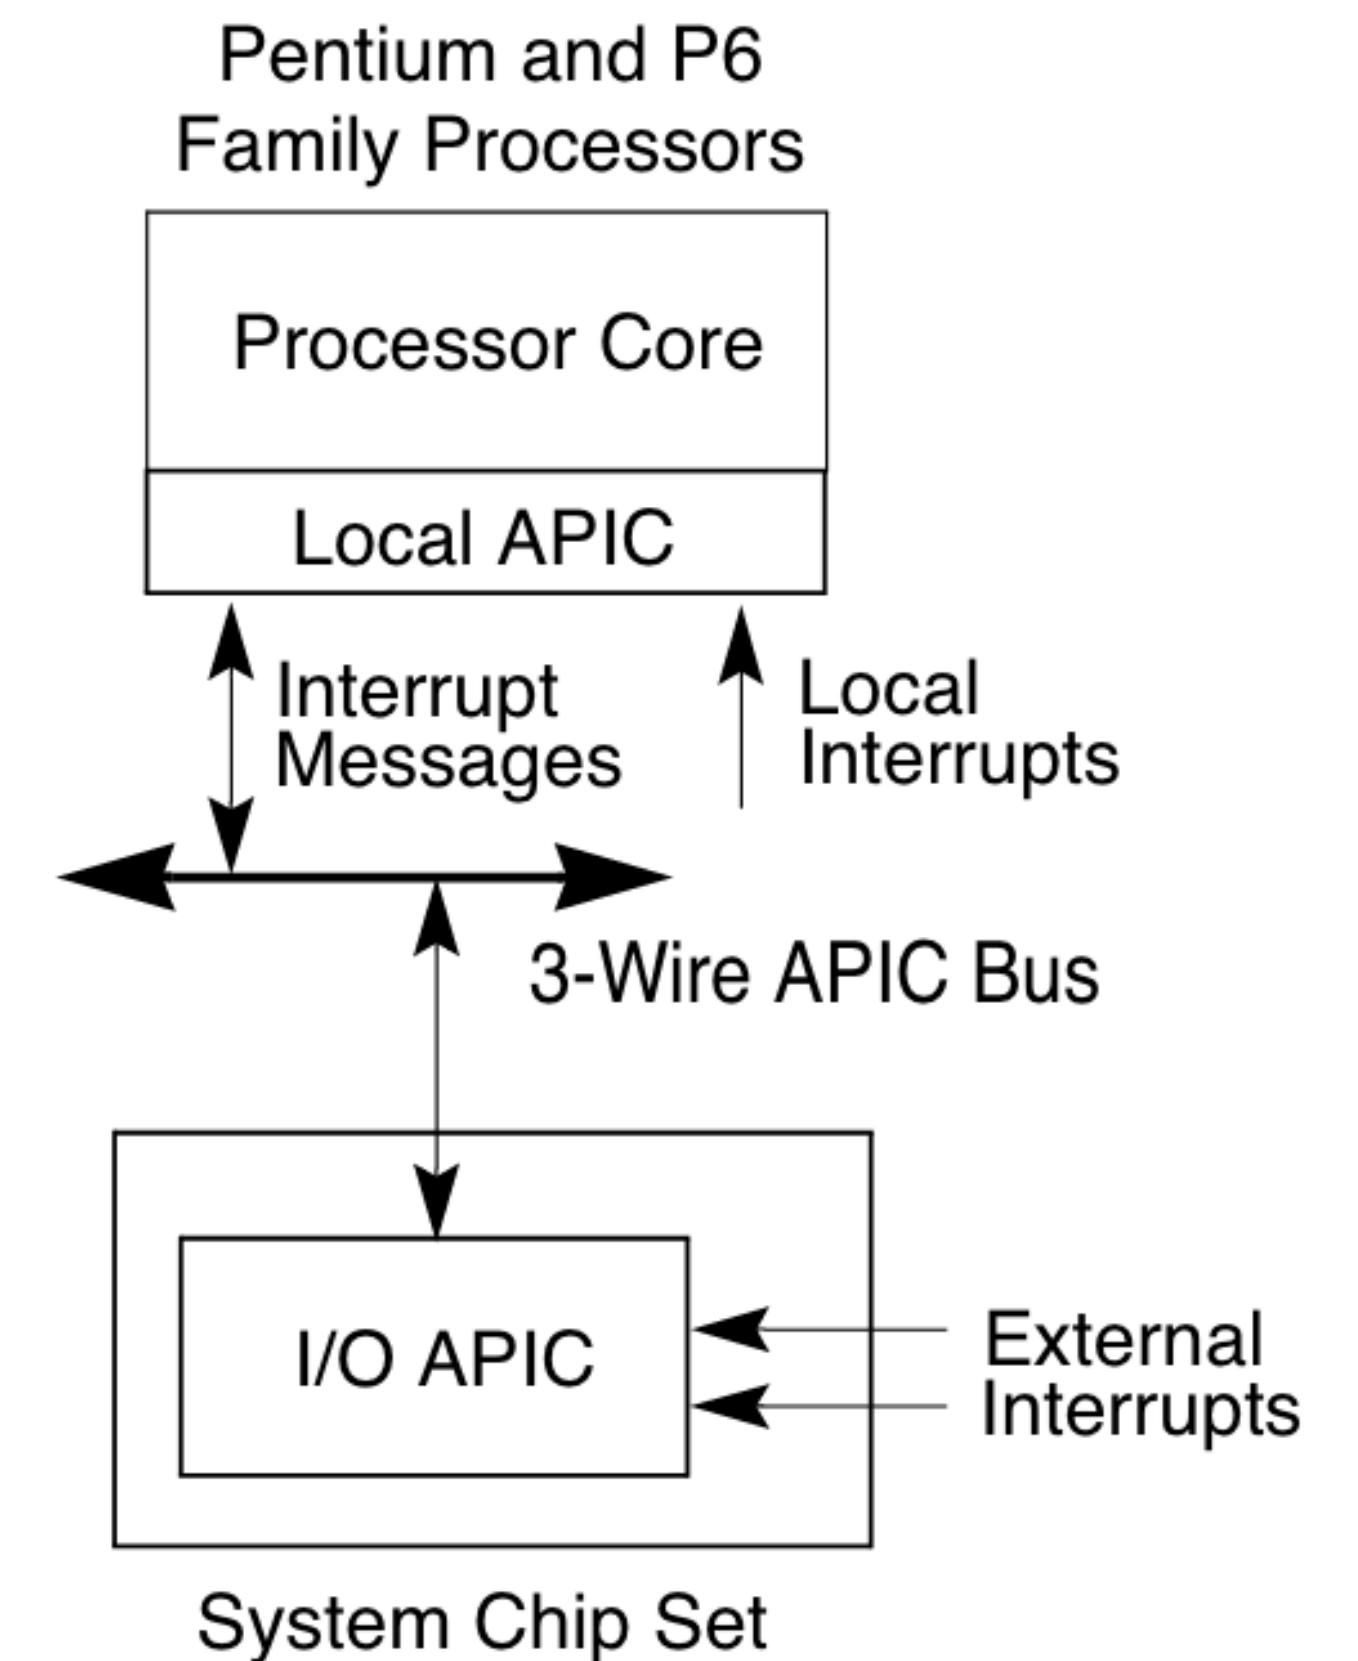
Advanced programmable interrupt controllers (APIC)

- Each CPU can have local APICs for handling *local interrupts* like timer, thermal sensor, etc.



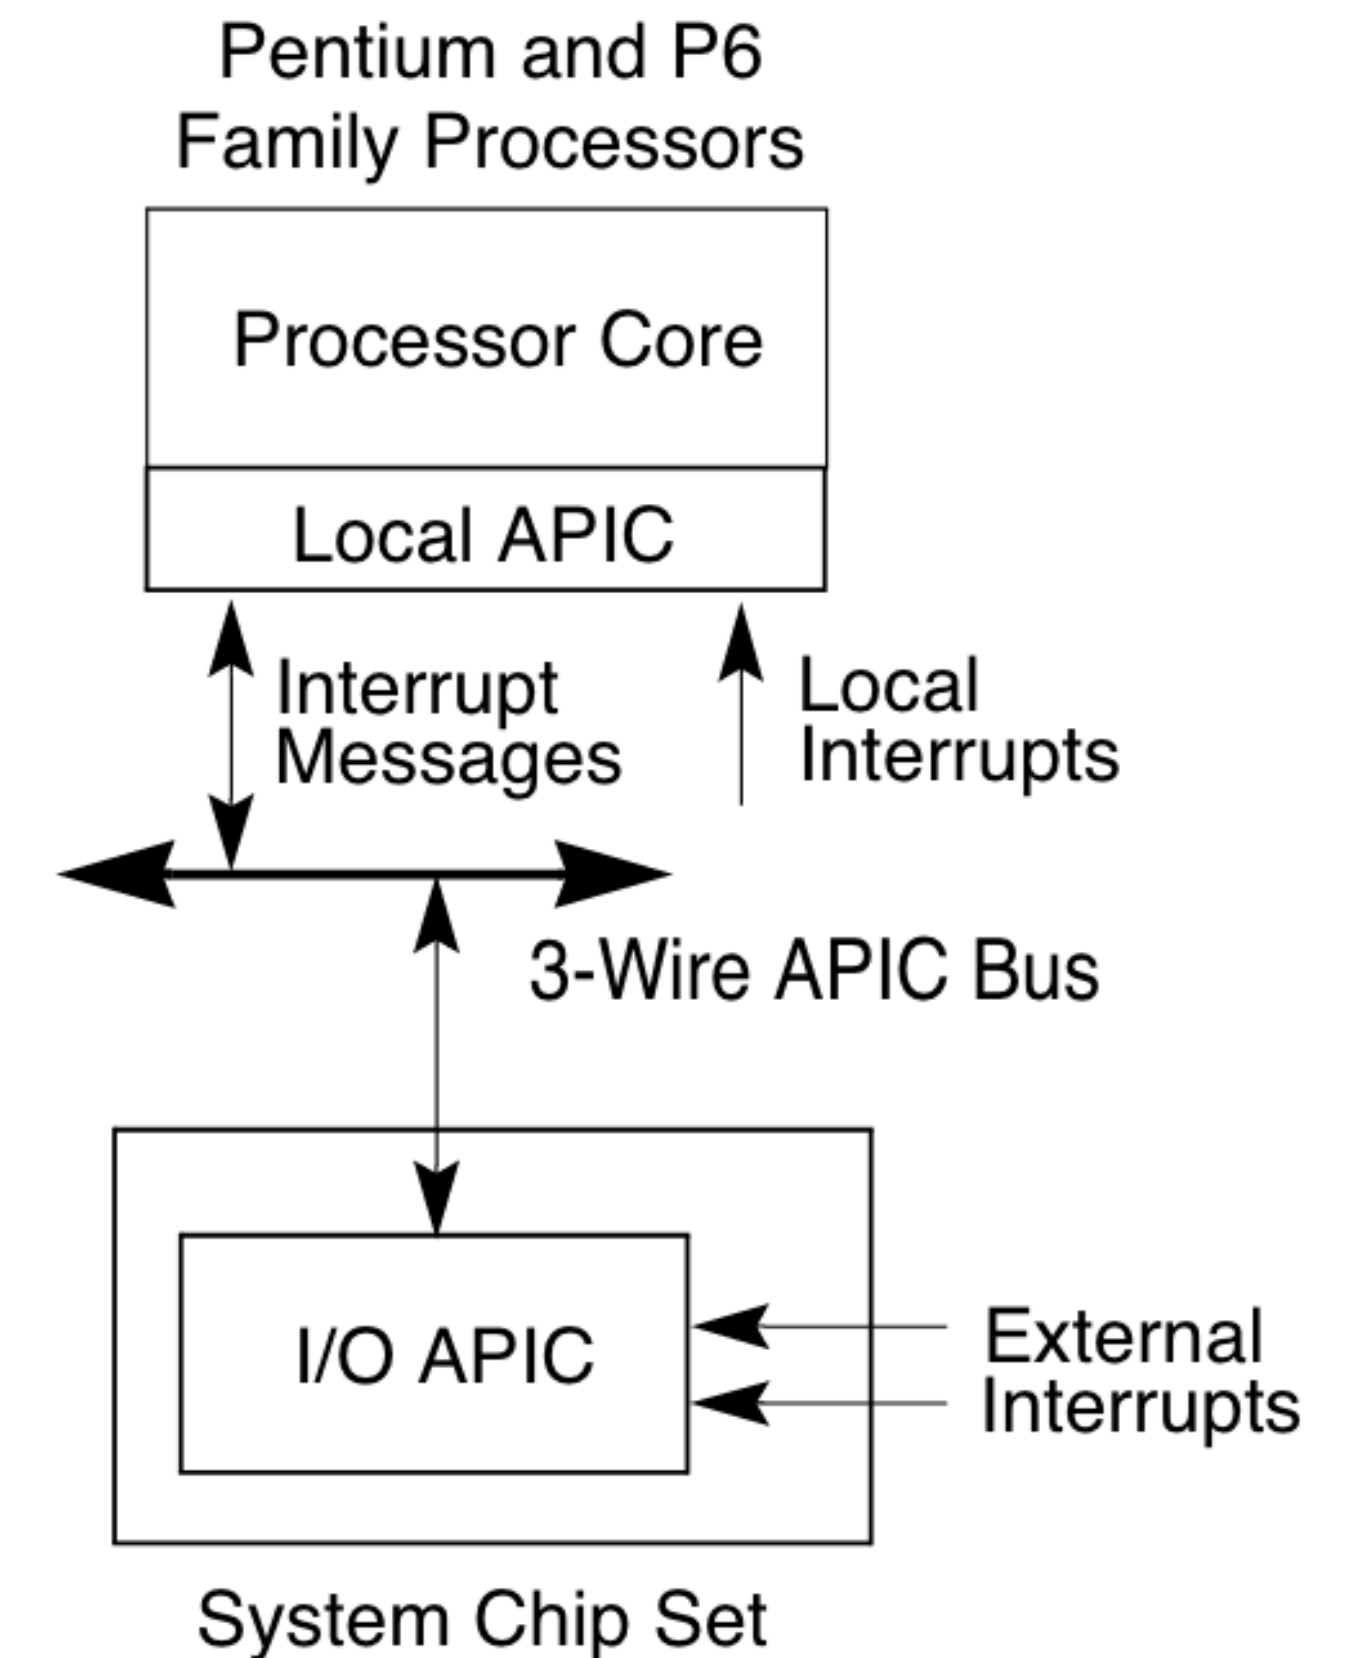
Advanced programmable interrupt controllers (APIC)

- Each CPU can have local APICs for handling *local interrupts* like timer, thermal sensor, etc.
- A separate IO APIC receives external interrupts like keyboard, mouse, disk, etc and forwards it to a particular CPU



Advanced programmable interrupt controllers (APIC)

- Each CPU can have local APICs for handling *local interrupts* like timer, thermal sensor, etc.
- A separate IO APIC receives external interrupts like keyboard, mouse, disk, etc and forwards it to a particular CPU
- Example: Route keyboard interrupts to CPU-0, disk interrupts to CPU-1



Code walkthrough

- main.c calls lapicinit, picinit, ioapicinit
- lapicinit enables timer interrupt at every 10ms. lapicw is just writing to memory location (MMIO)
- picinit just disables PIC using outb instructions (PMIO)
- ioapicinit initialises IO APIC with MMIO
- Bootloader had disabled interrupt with cli. We will not receive interrupts yet.

Interrupt enable flag

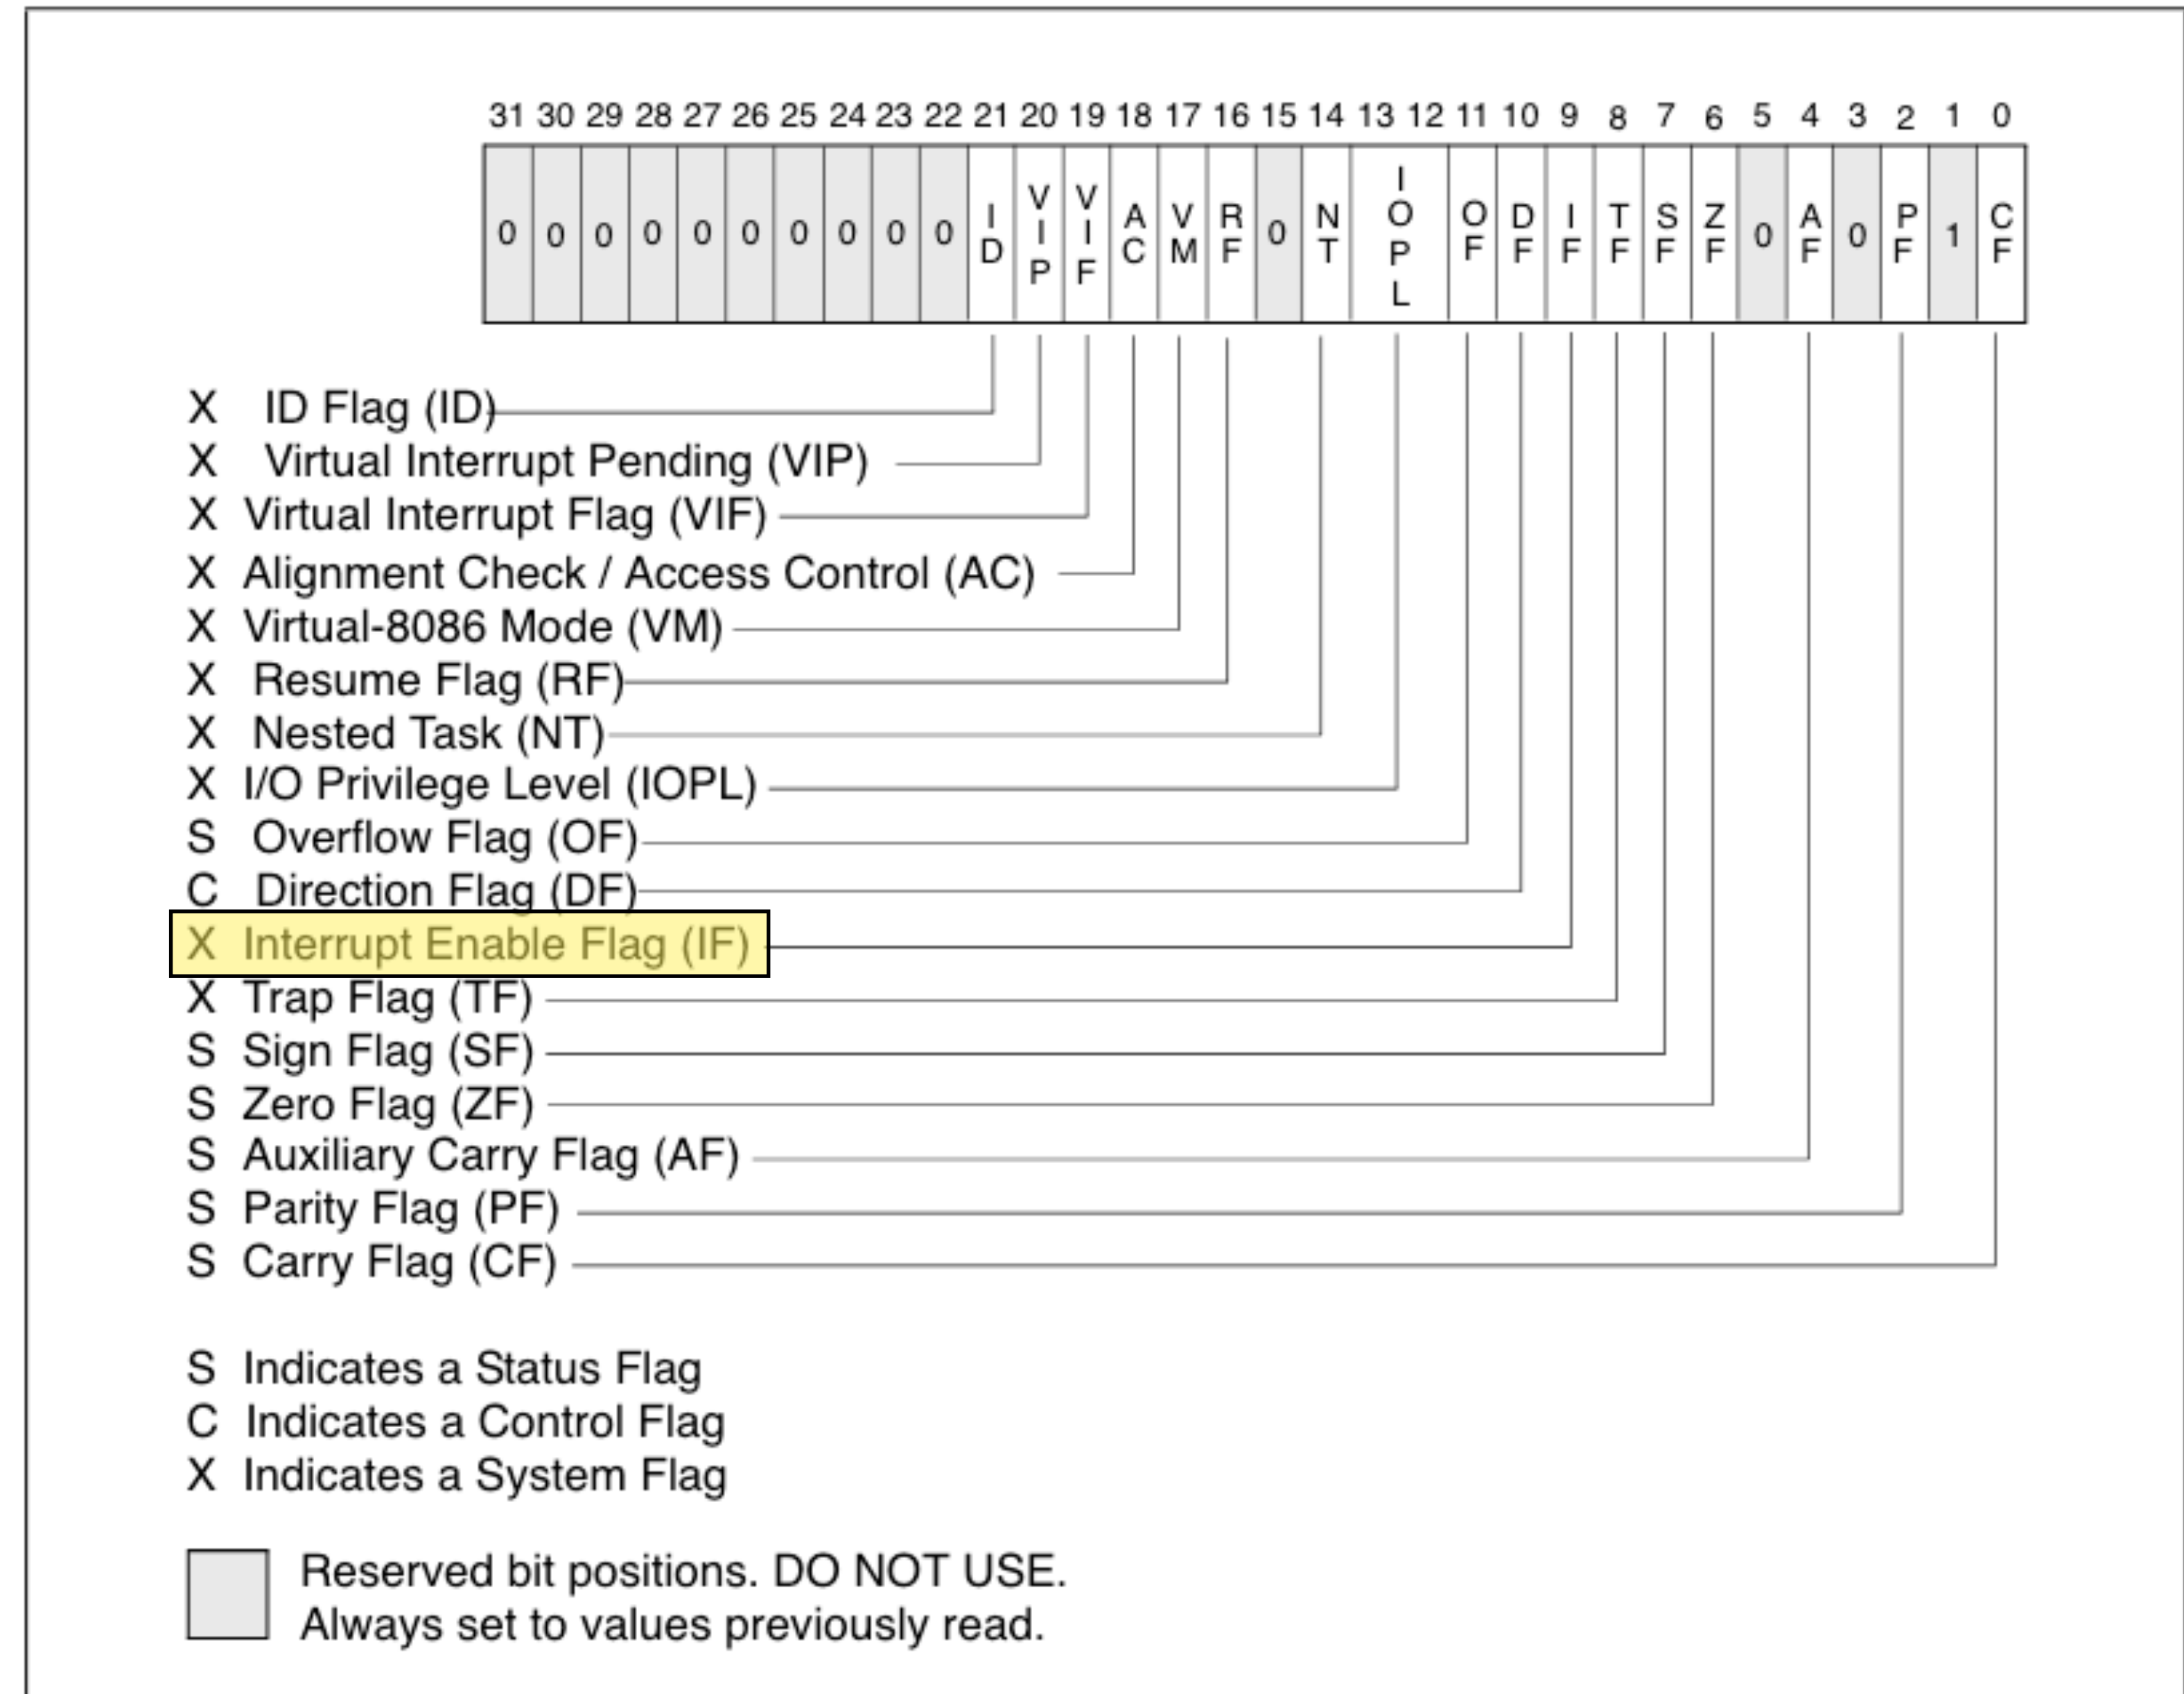


Figure 3-8. EFLAGS Register

Interrupt enable flag

- cli: Clear interrupt flag

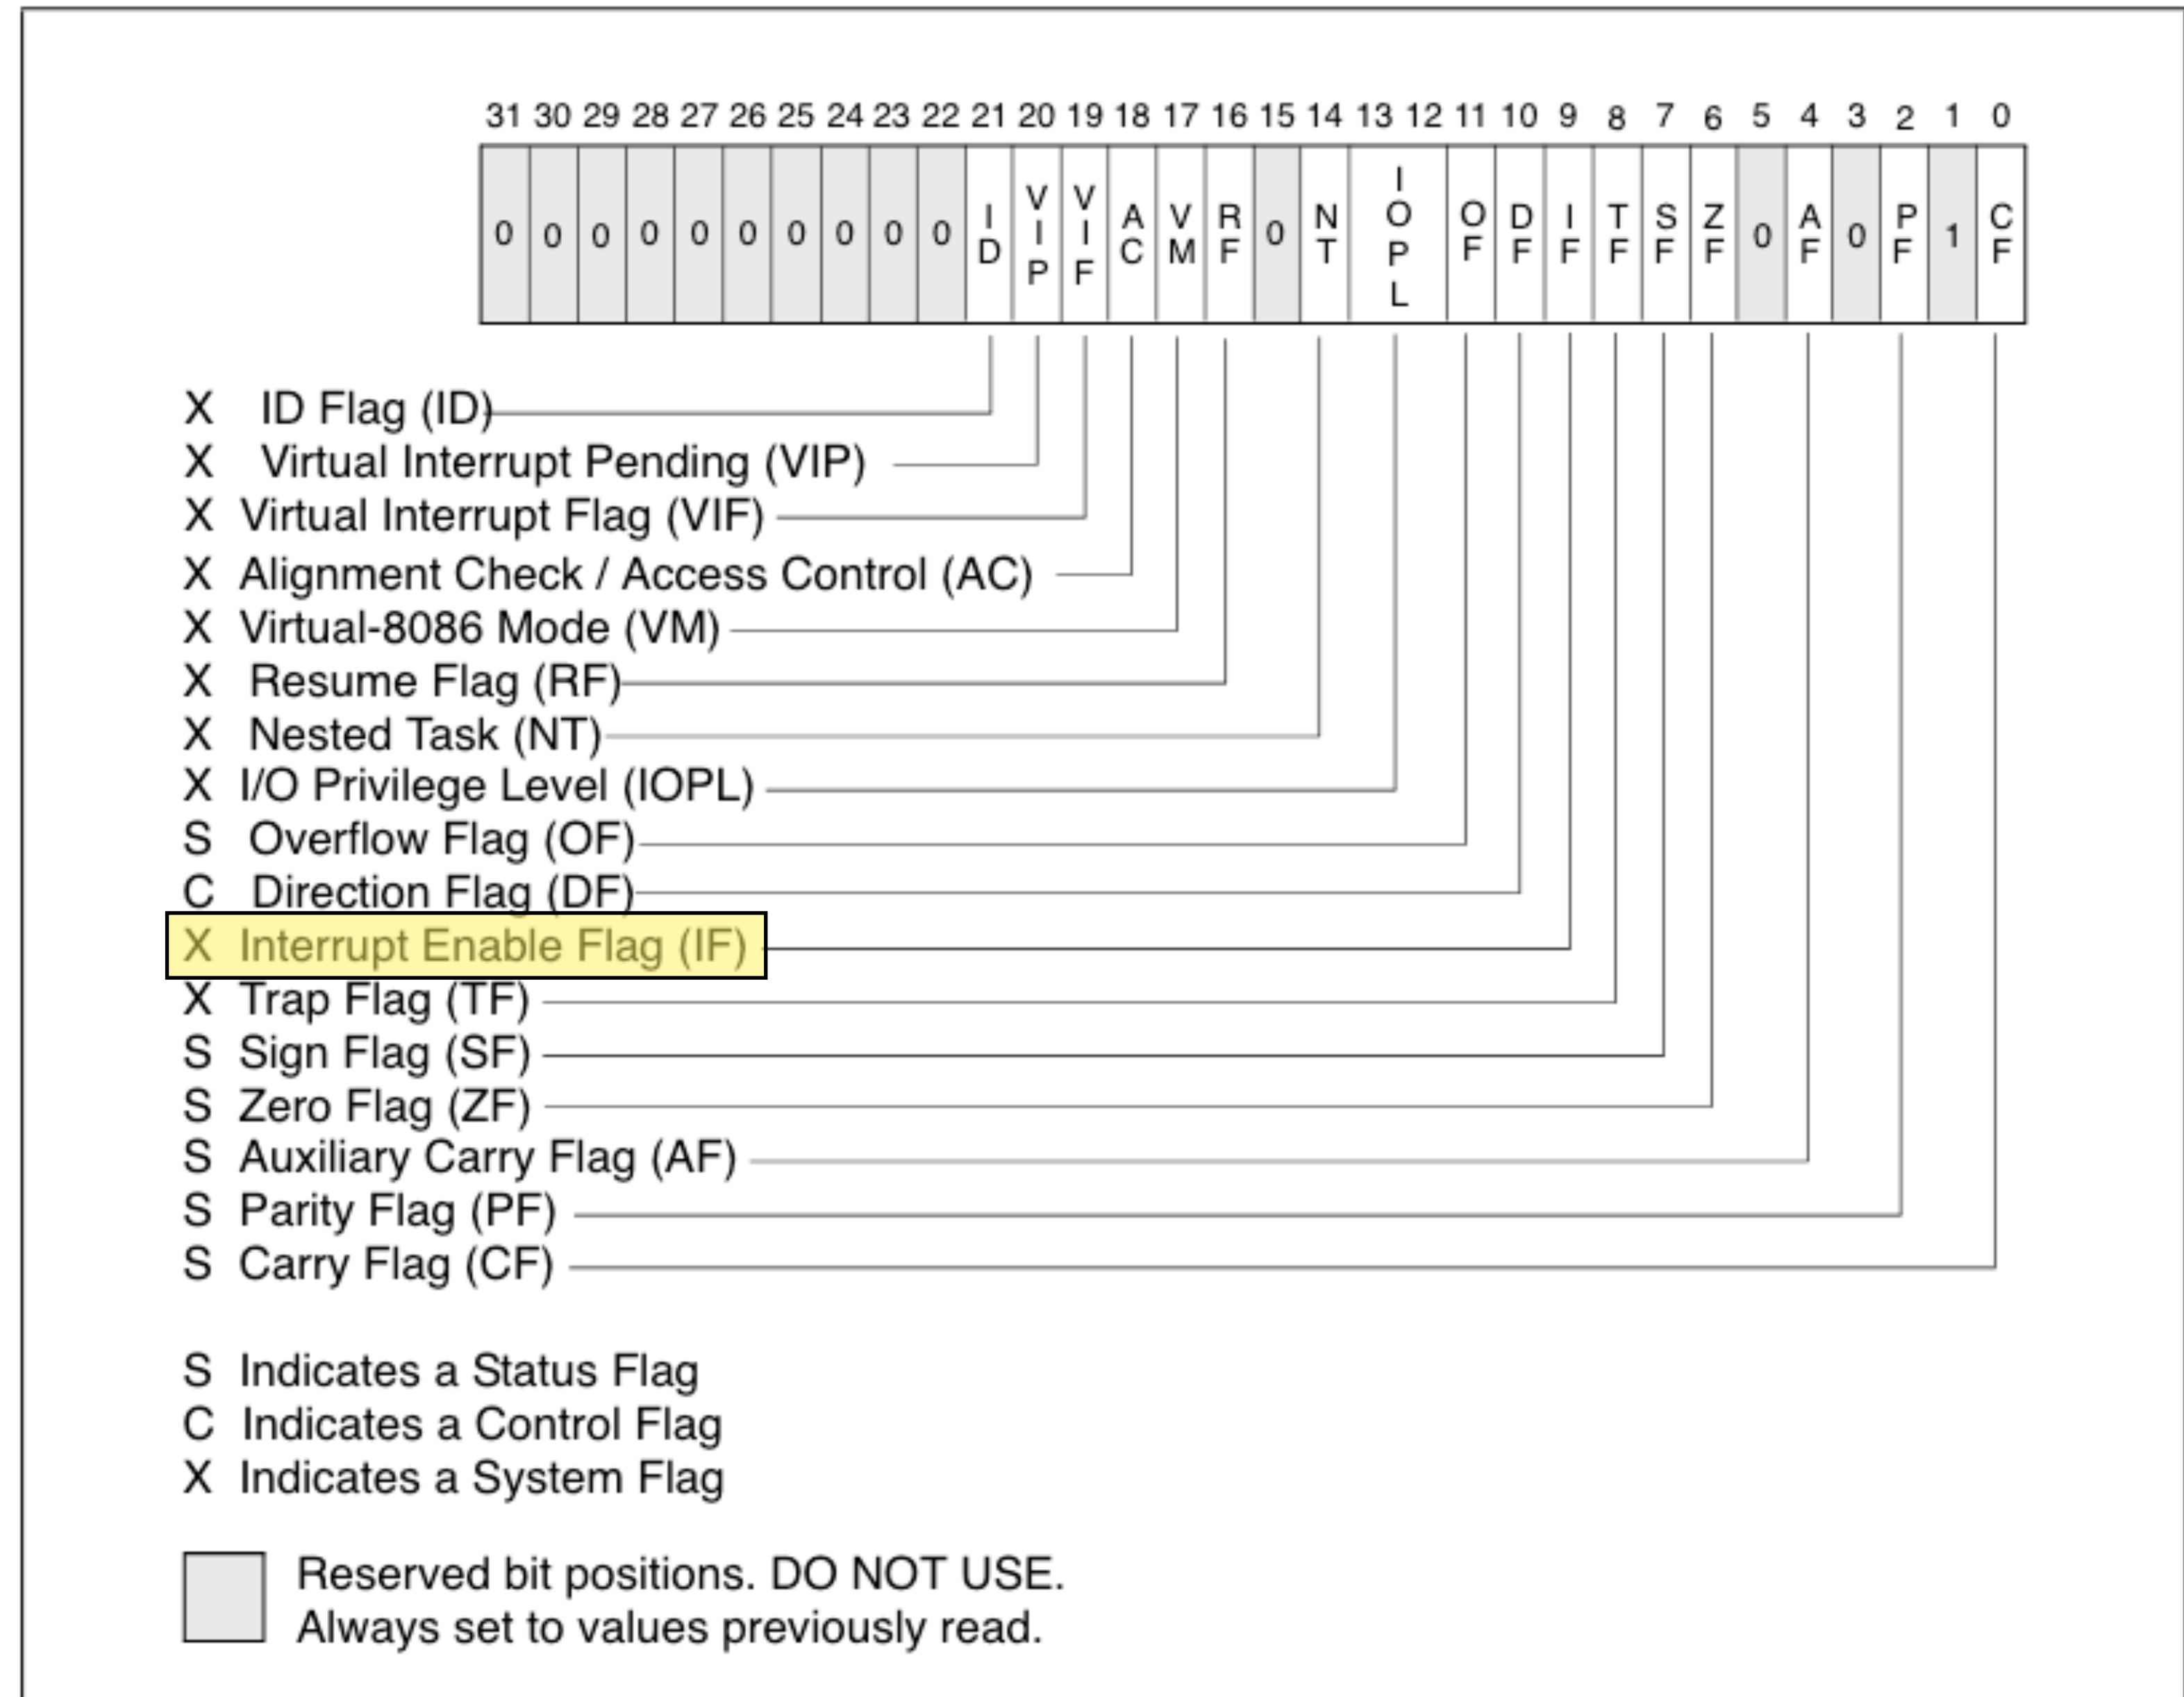


Figure 3-8. EFLAGS Register

Interrupt enable flag

- cli: Clear interrupt flag
 - PICs are not allowed to interrupt

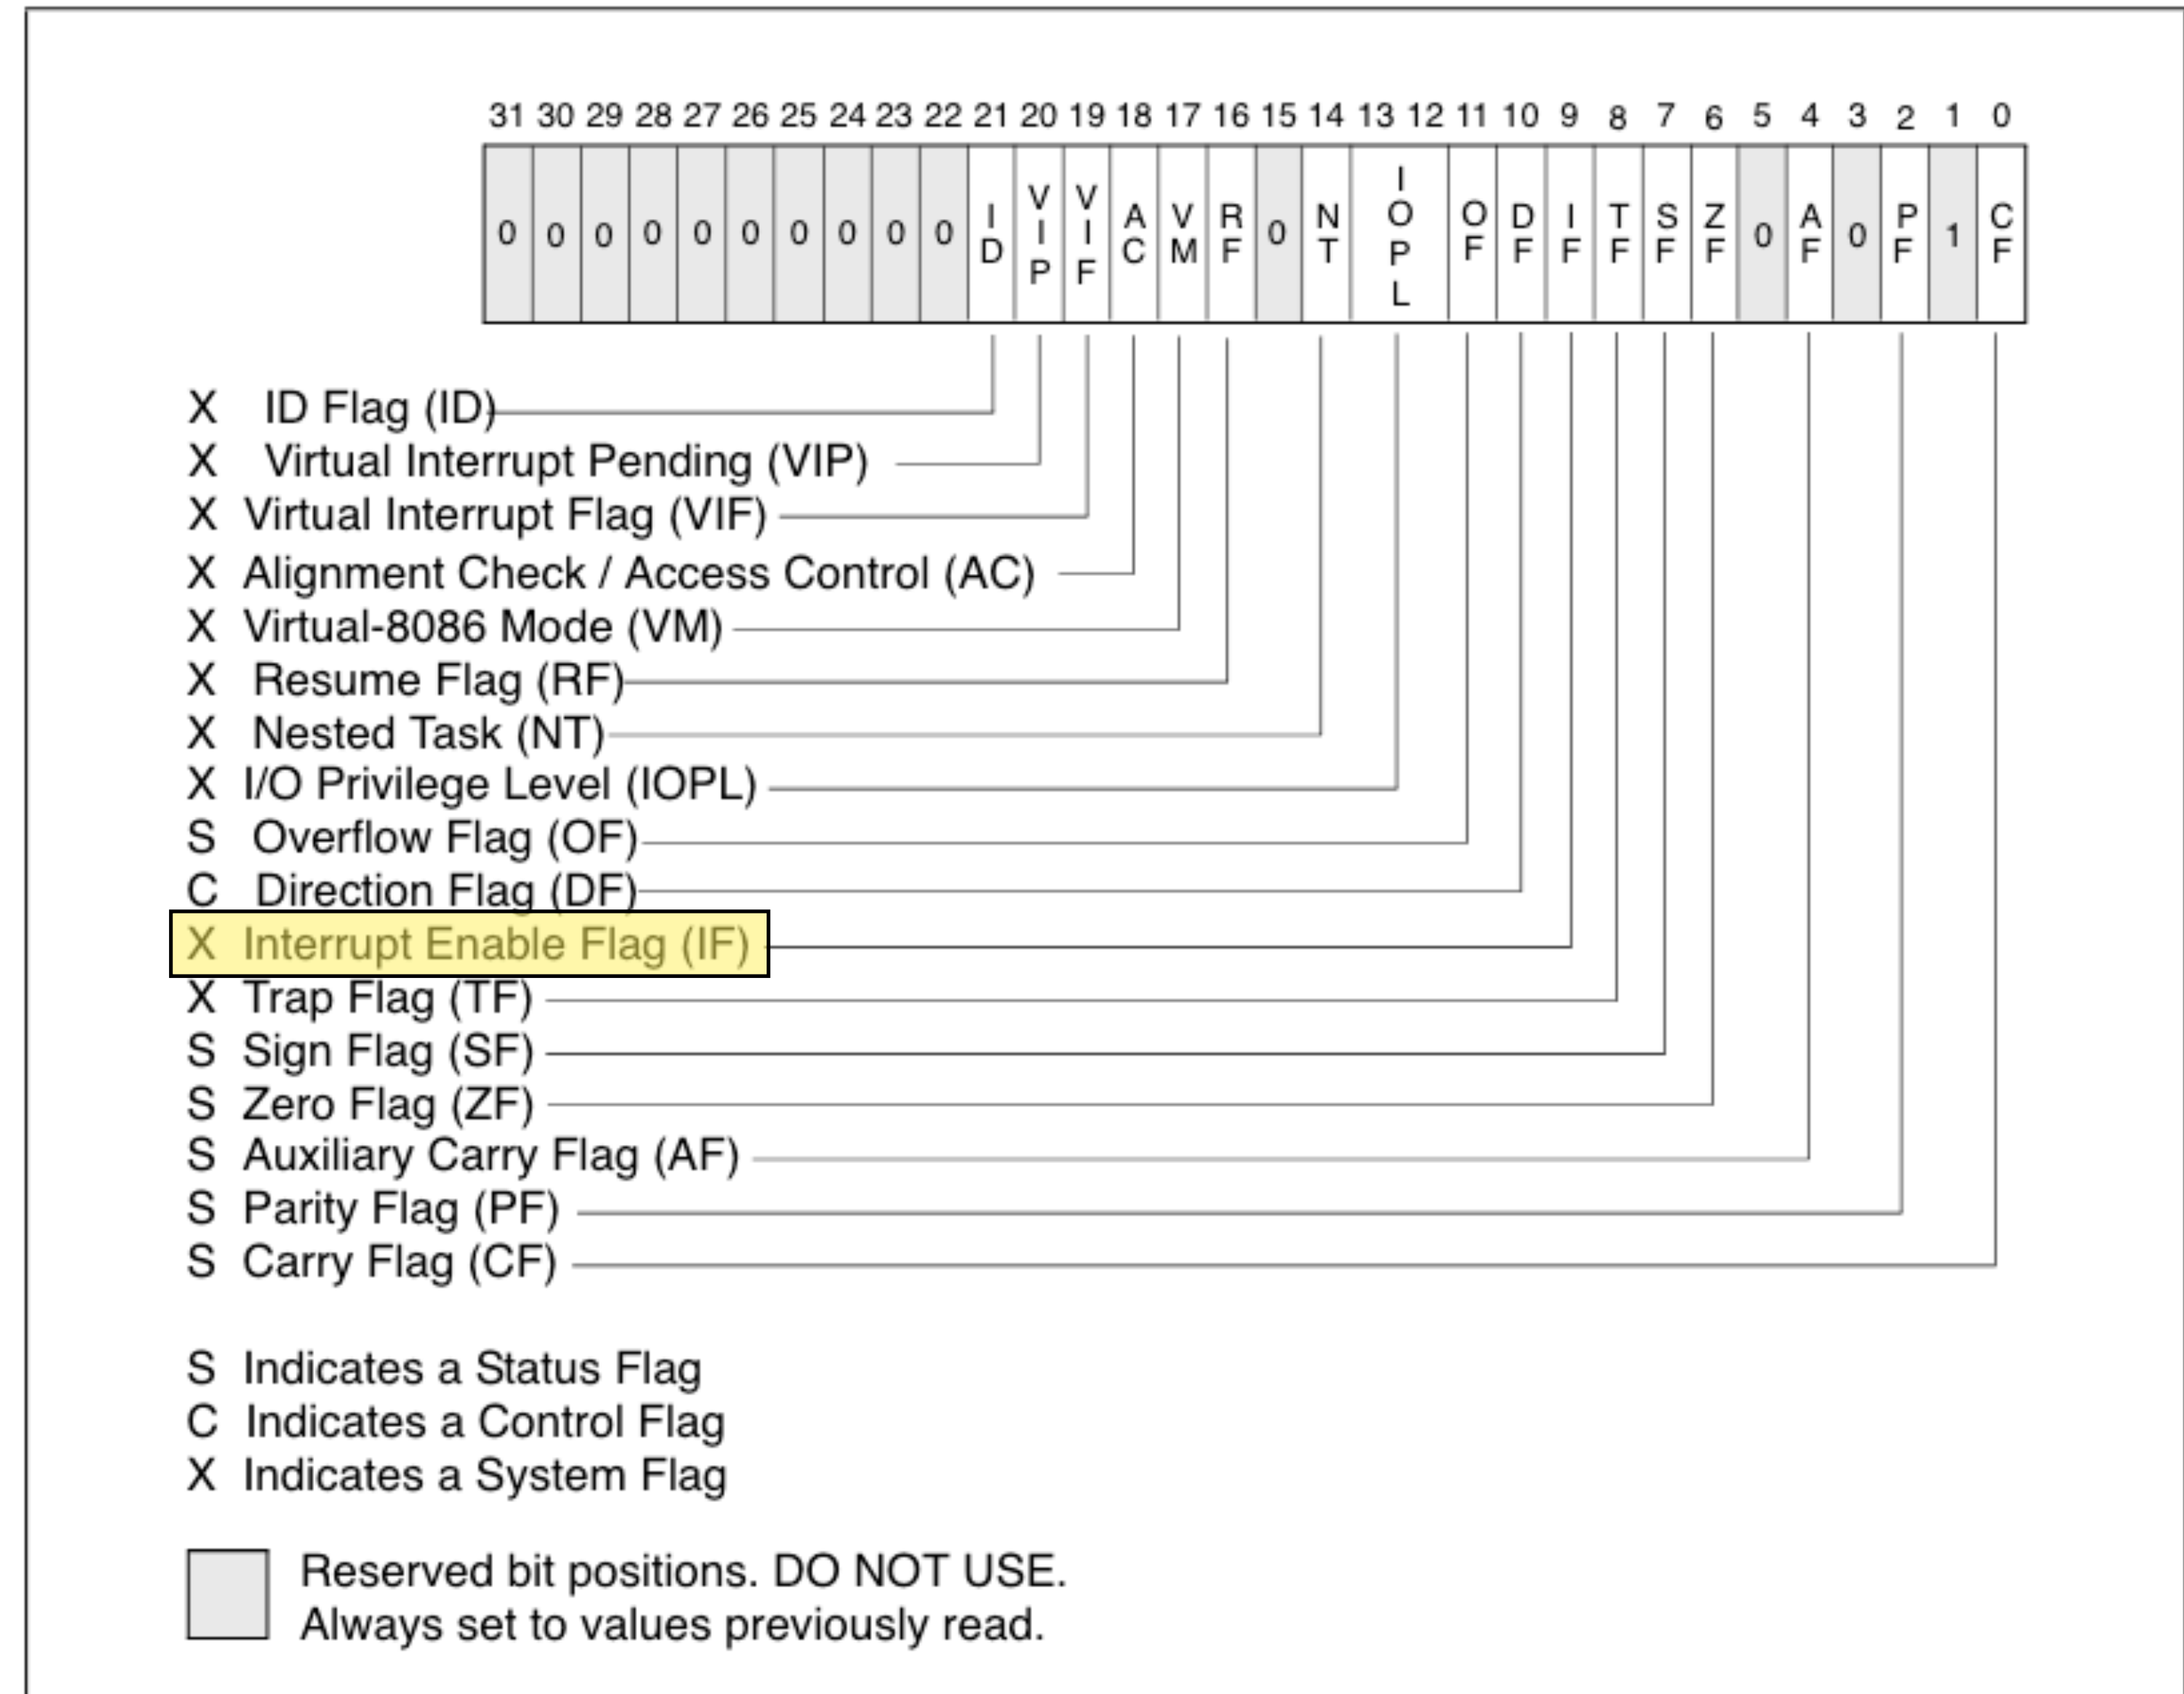


Figure 3-8. EFLAGS Register

Interrupt enable flag

- cli: Clear interrupt flag
 - PICs are not allowed to interrupt
- sti: Set interrupt flag

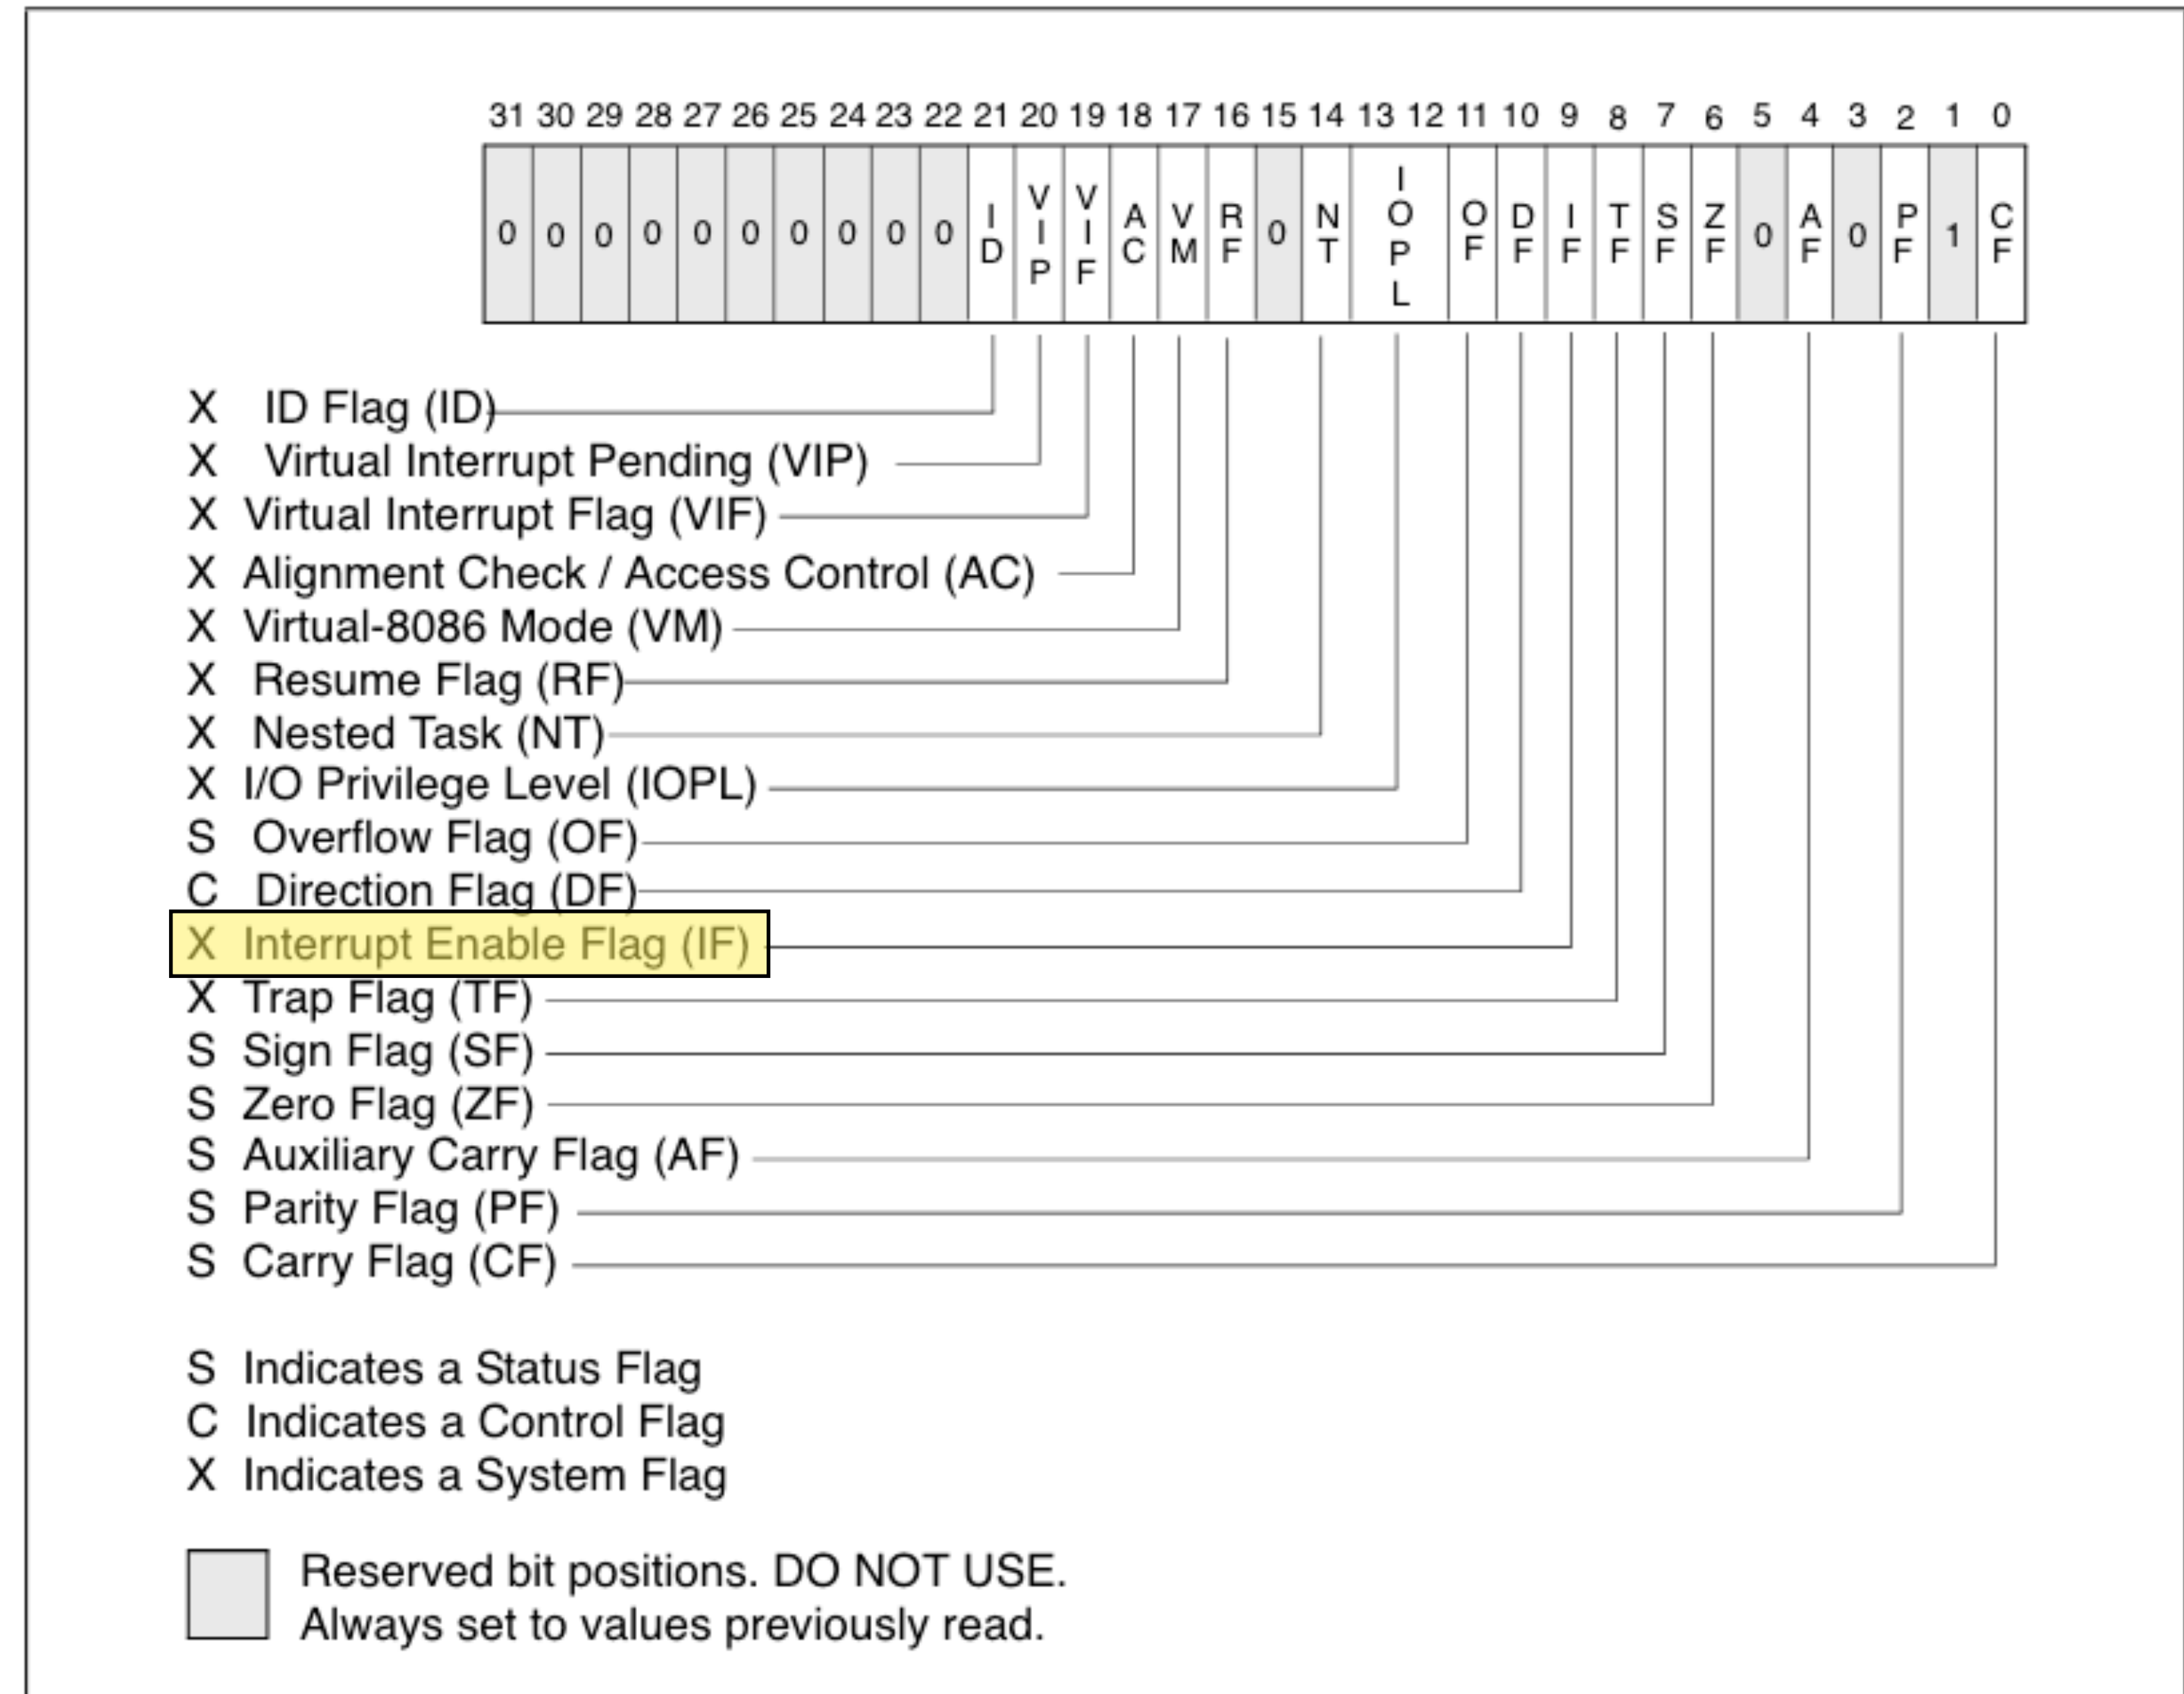


Figure 3-8. EFLAGS Register

Interrupt handling in a nutshell

Interrupt handling in a nutshell

- OS sets up “interrupt descriptor table” (IDT)

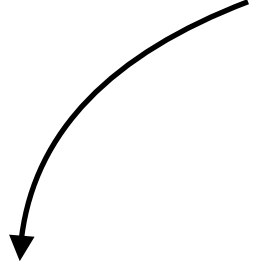
Interrupt handling in a nutshell

- OS sets up “interrupt descriptor table” (IDT)
- Points IDTR to IDT using LIDT instruction

Interrupt handling in a nutshell

- OS sets up “interrupt descriptor table” (IDT)
- Points IDTR to IDT using LIDT instruction

IDTR: Interrupt descriptor table register

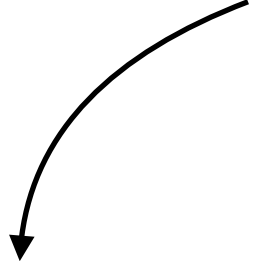


S.No.*	cs	eip
0x01
0x02	0x8	0xFF
...

Interrupt handling in a nutshell

- OS sets up “interrupt descriptor table” (IDT)
- Points IDTR to IDT using LIDT instruction
- When interrupt occurs, jump %eip to interrupt handler, handle interrupt, tell LAPIC about end of interrupt, resume what we were doing

IDTR: Interrupt descriptor table register



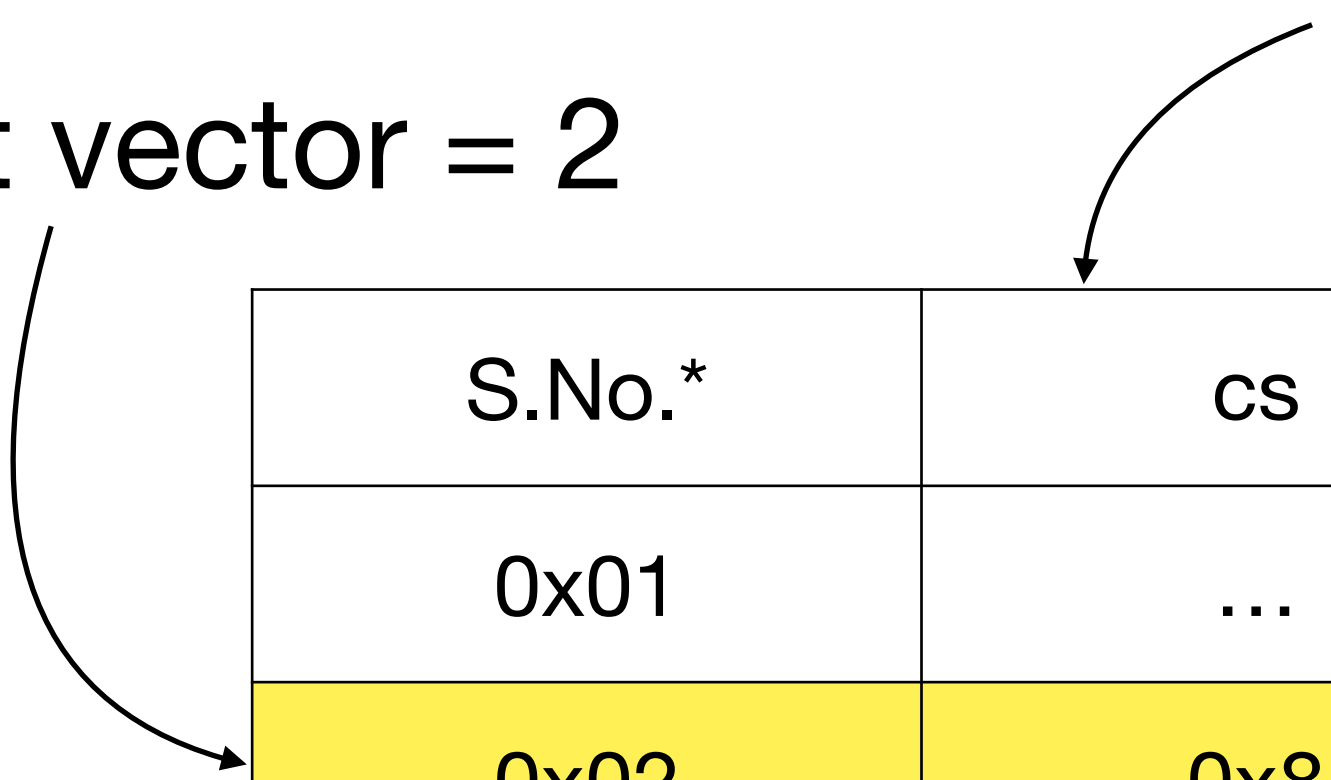
S.No.*	cs	eip
0x01
0x02	0x8	0xFF
...

Interrupt handling in a nutshell

- OS sets up “interrupt descriptor table” (IDT)
- Points IDTR to IDT using LIDT instruction
- When interrupt occurs, jump %eip to interrupt handler, handle interrupt, tell LAPIC about end of interrupt, resume what we were doing

IDTR: Interrupt descriptor table register

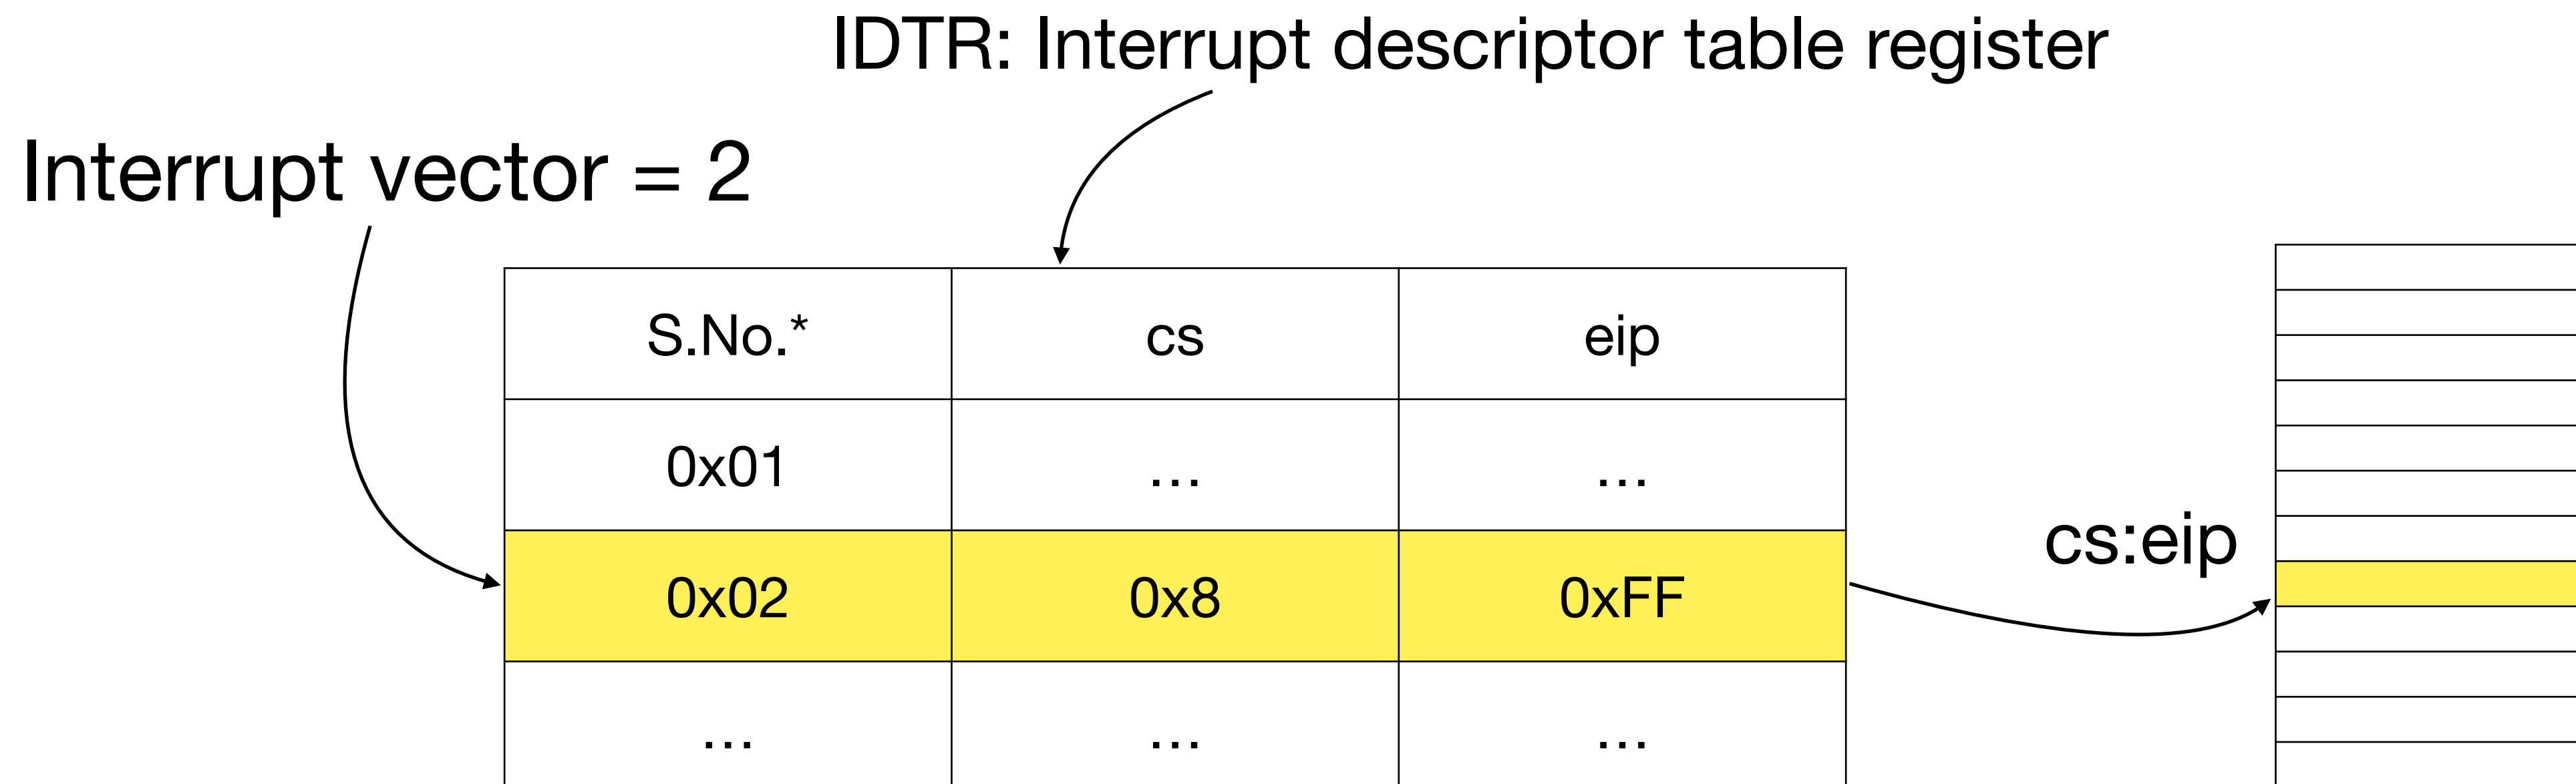
Interrupt vector = 2



S.No.*	cs	eip
0x01
0x02	0x8	0xFF
...

Interrupt handling in a nutshell

- OS sets up “interrupt descriptor table” (IDT)
- Points IDTR to IDT using LIDT instruction
- When interrupt occurs, jump %eip to interrupt handler, handle interrupt, tell LAPIC about end of interrupt, resume what we were doing



Interrupt descriptor table

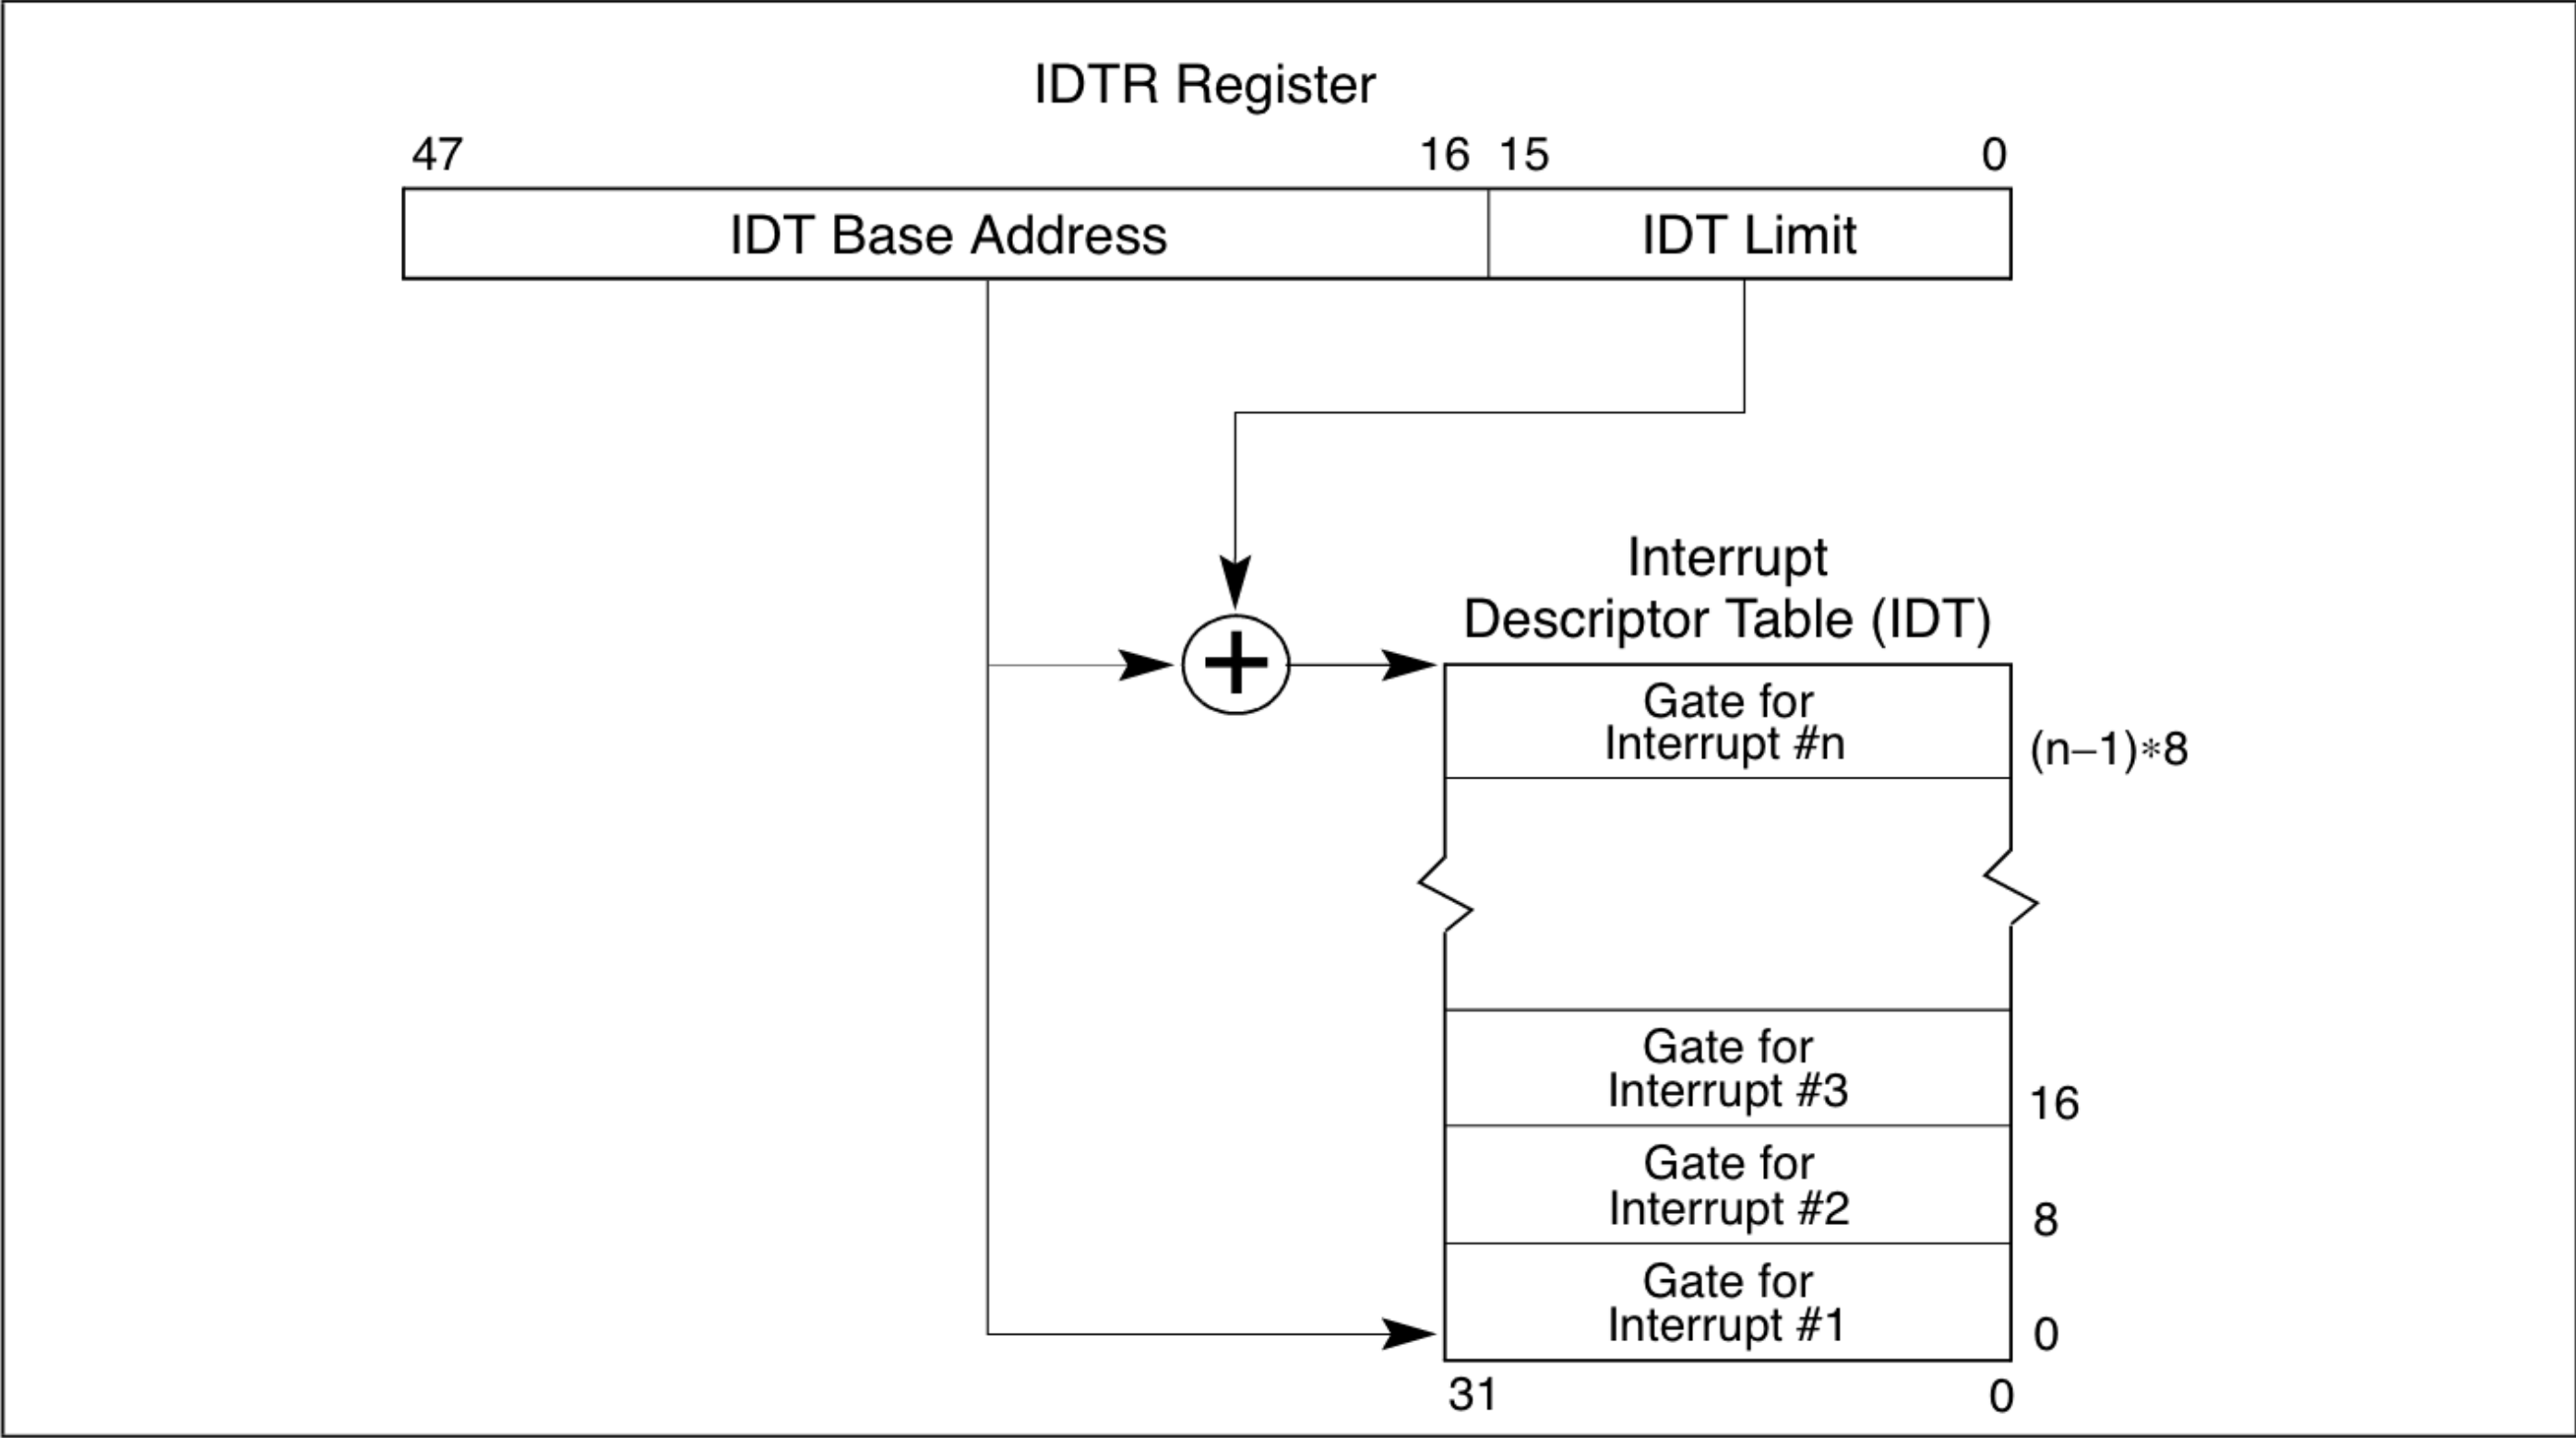


Figure 6-1. Relationship of the IDTR and IDT

Interrupt descriptor table

- Interrupt descriptor table register (IDTR) points to interrupt descriptor table in memory

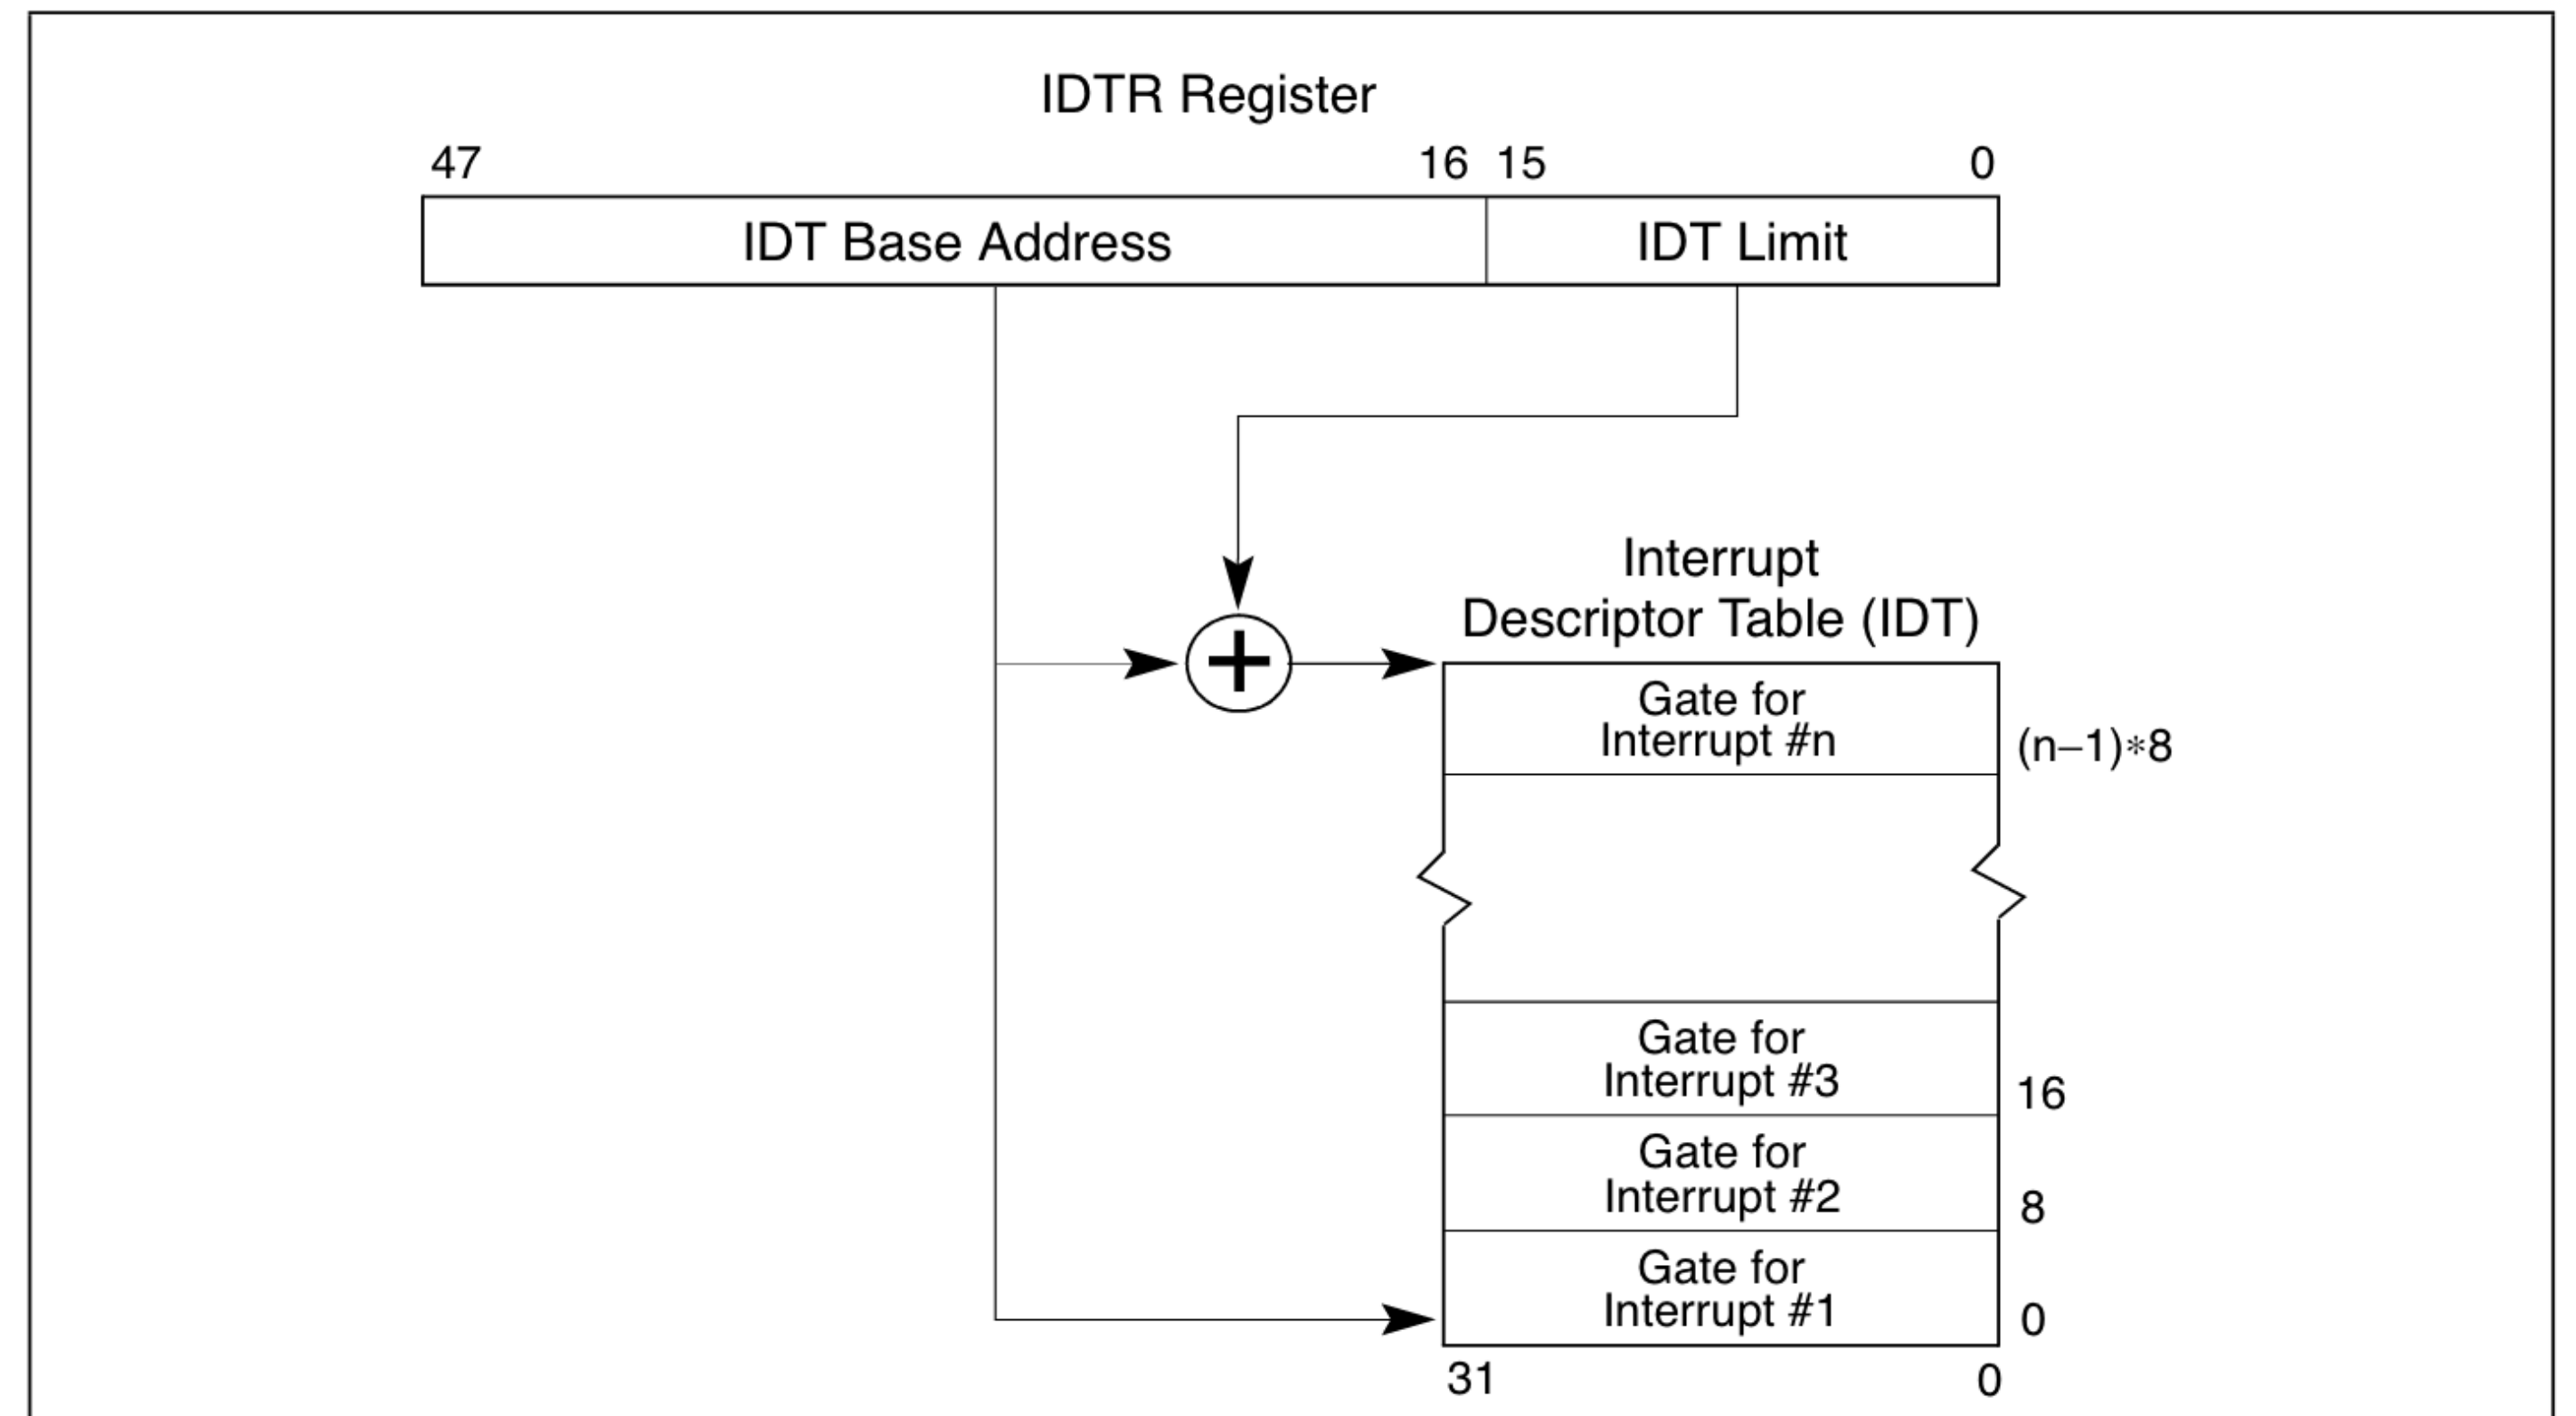


Figure 6-1. Relationship of the IDTR and IDT

Interrupt descriptor table

- Interrupt descriptor table register (IDTR) points to interrupt descriptor table in memory
- OS sets up IDT and initialises IDTR using LIDT instruction

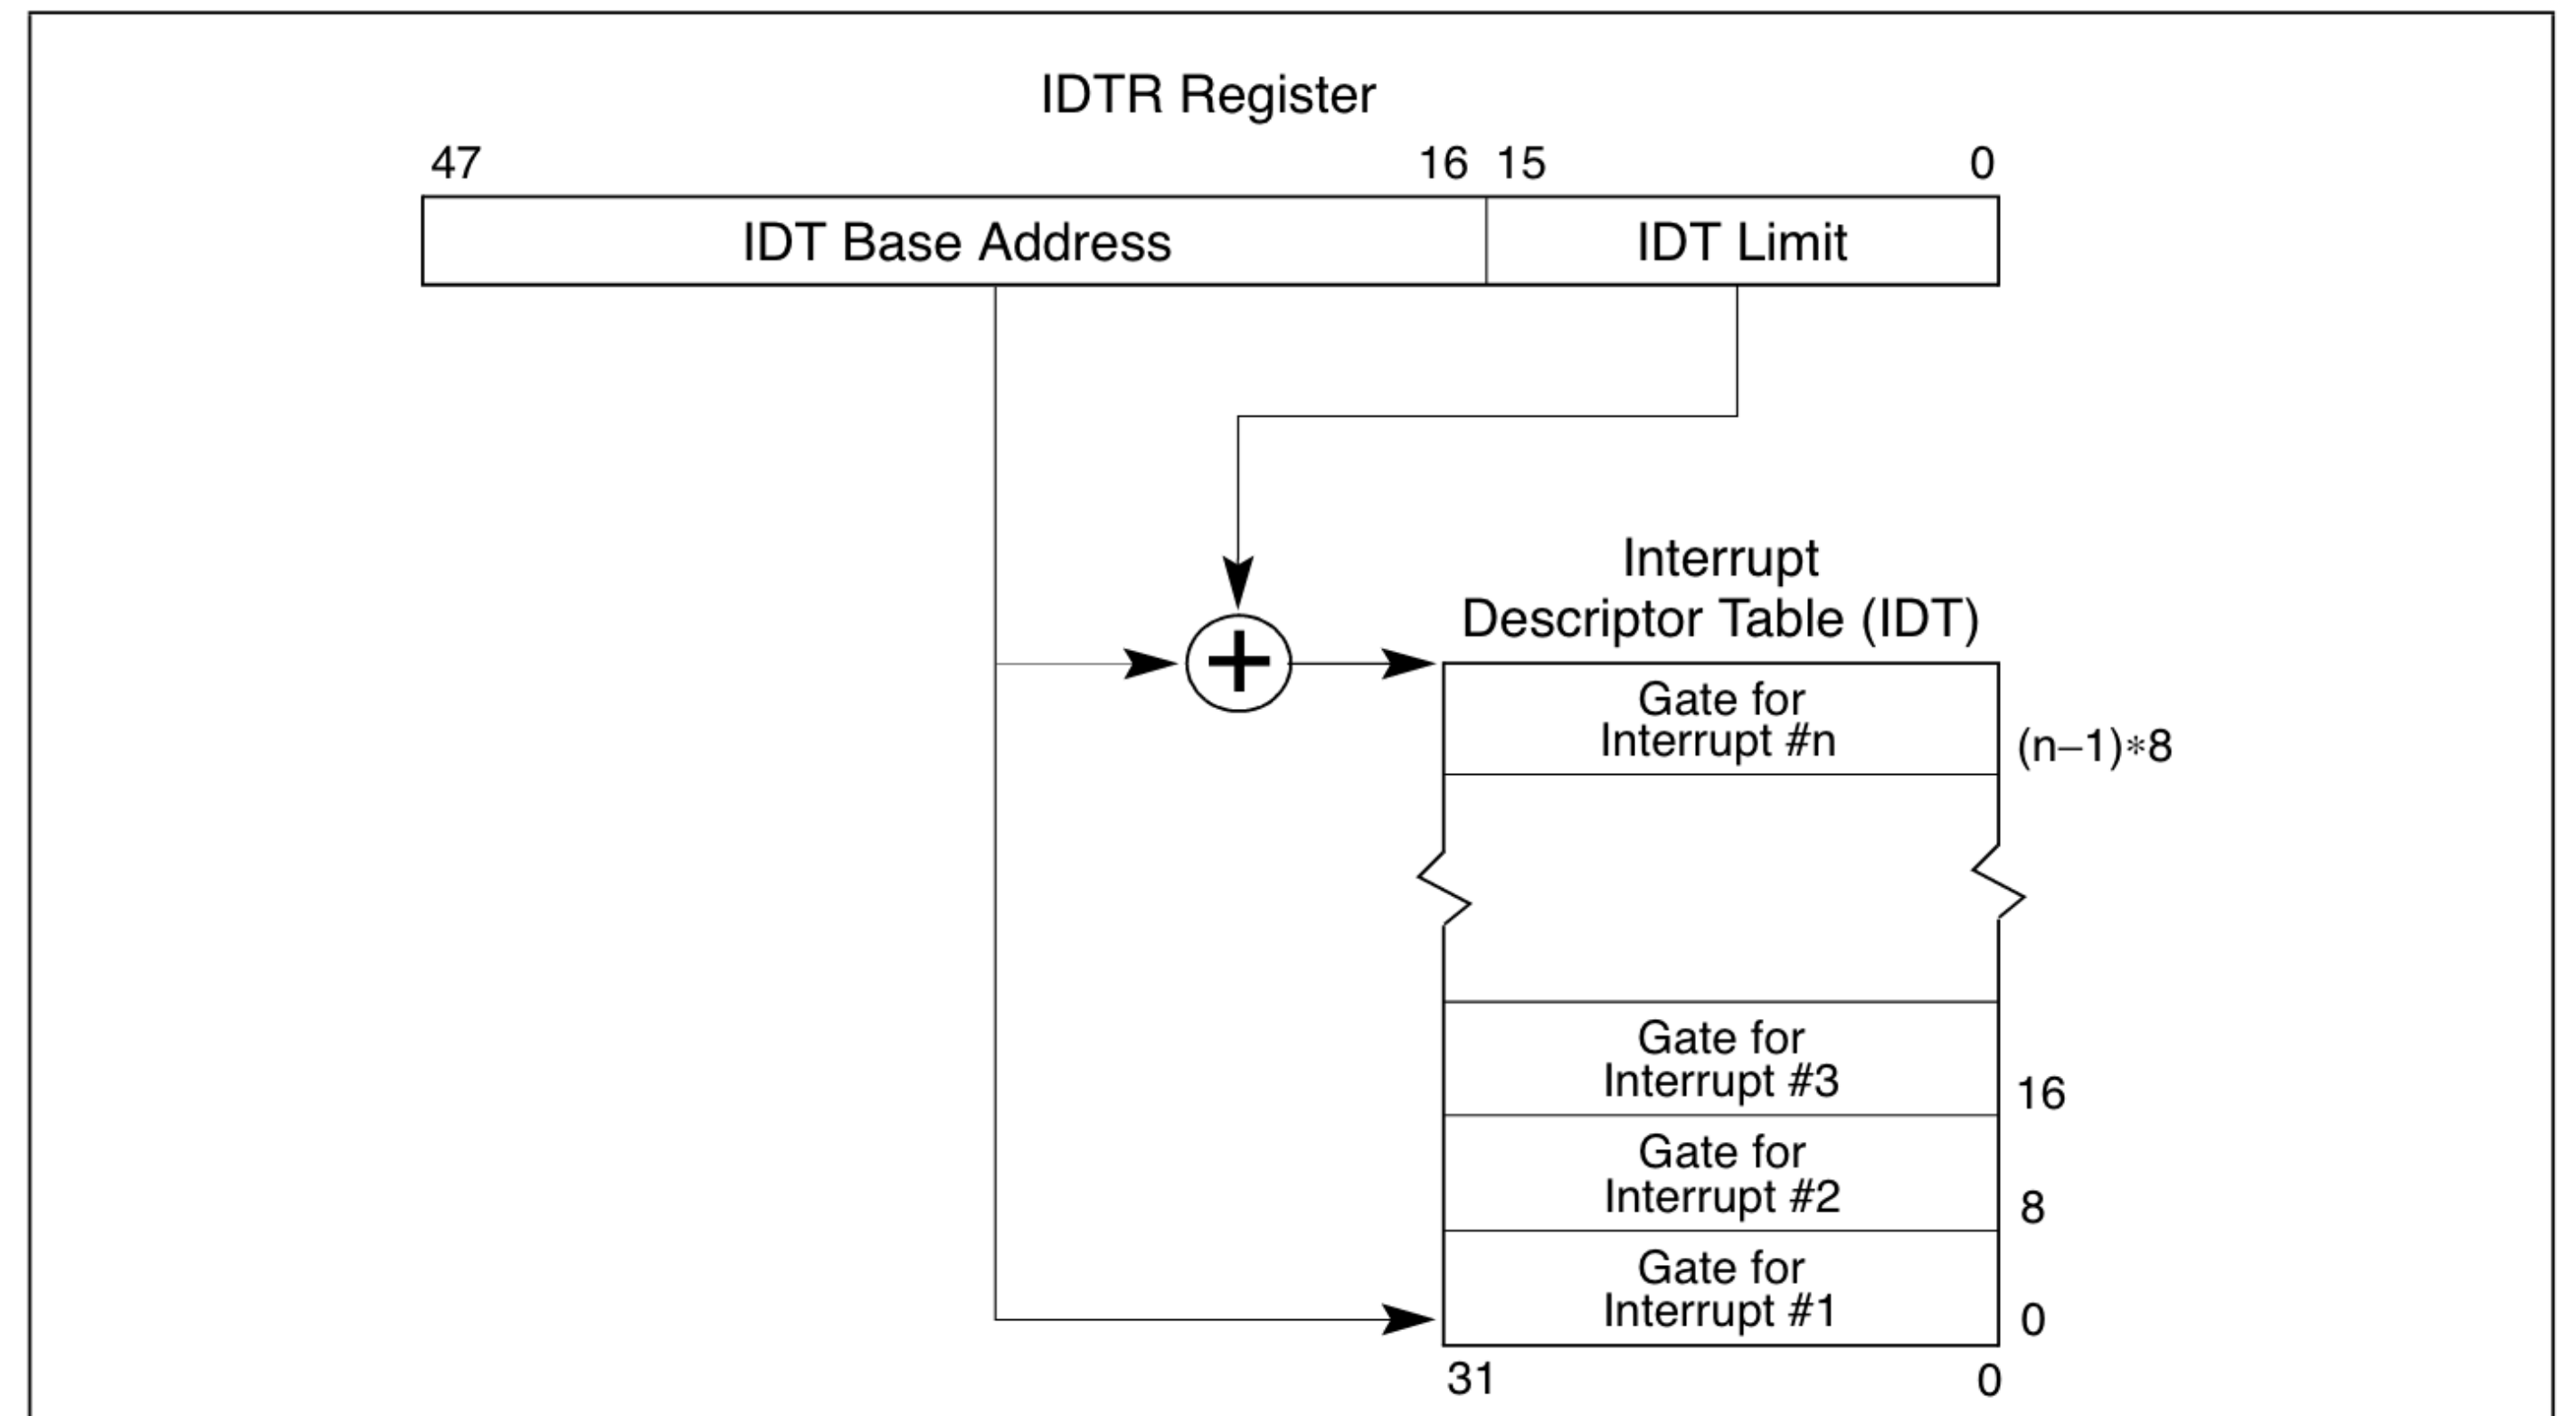


Figure 6-1. Relationship of the IDTR and IDT

Interrupt descriptor table

- Interrupt descriptor table register (IDTR) points to interrupt descriptor table in memory
- OS sets up IDT and initialises IDTR using LIDT instruction
- Interrupt descriptor table has one entry for each interrupt vector (upto $2^8=256$)

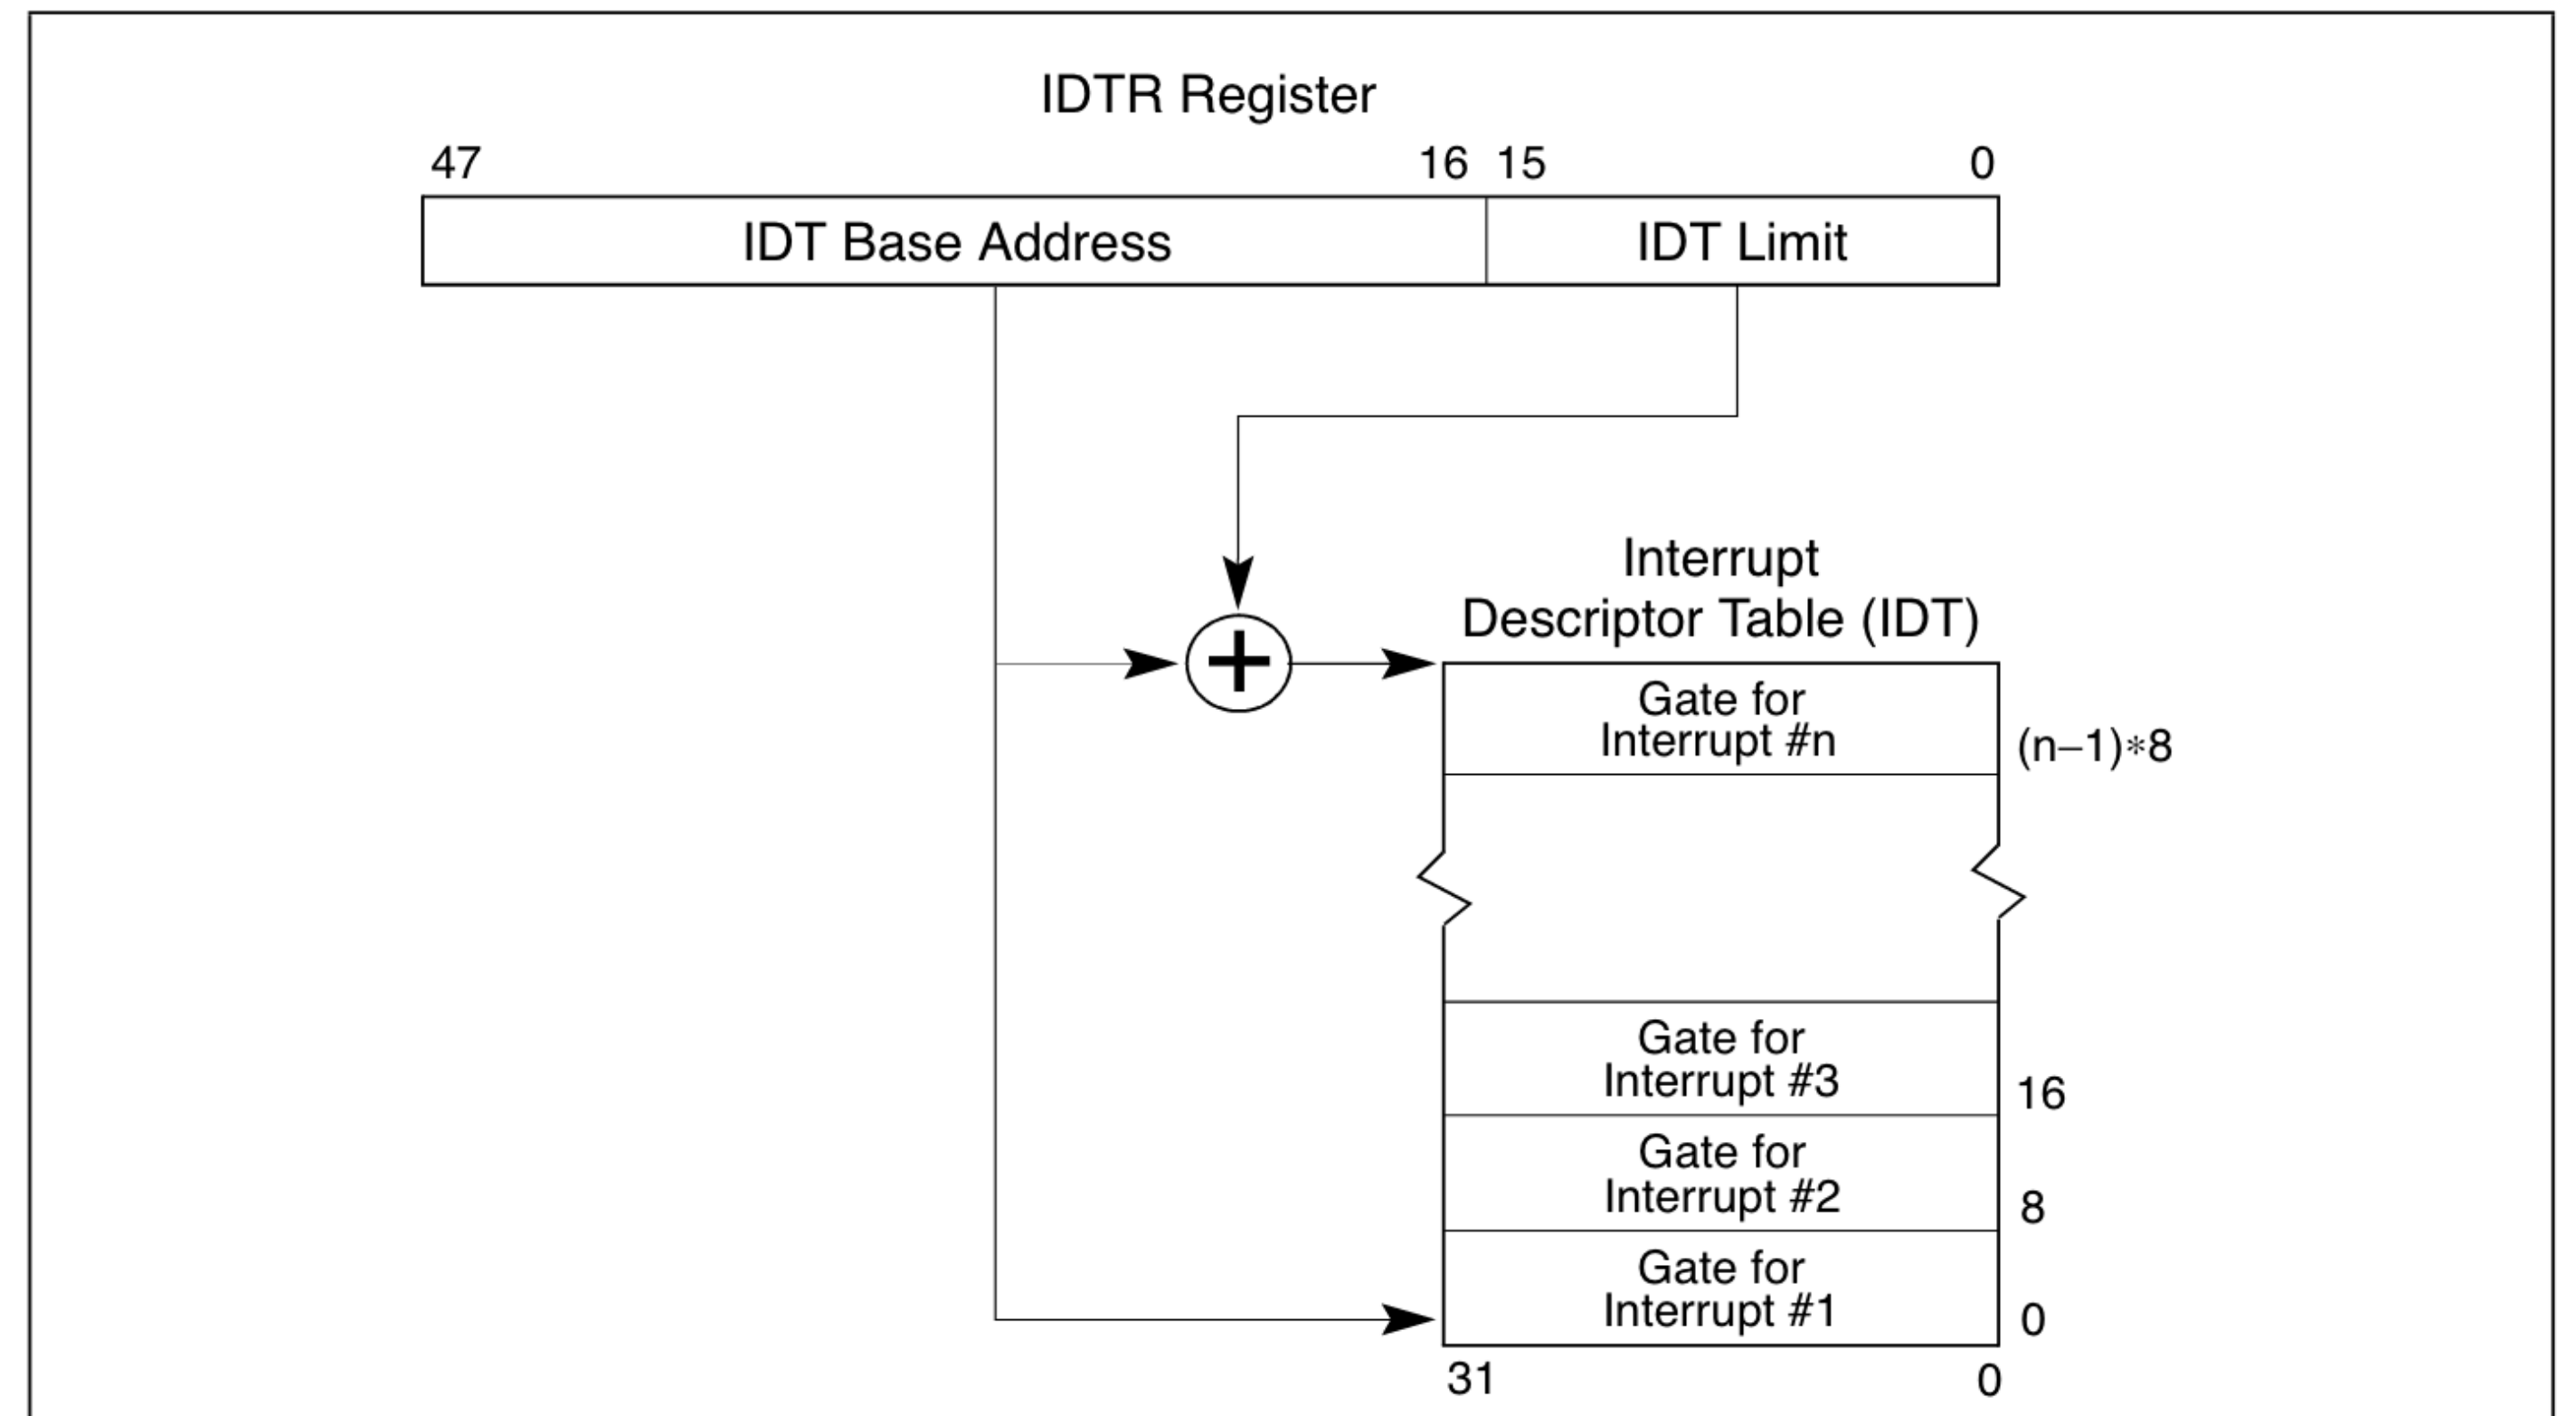
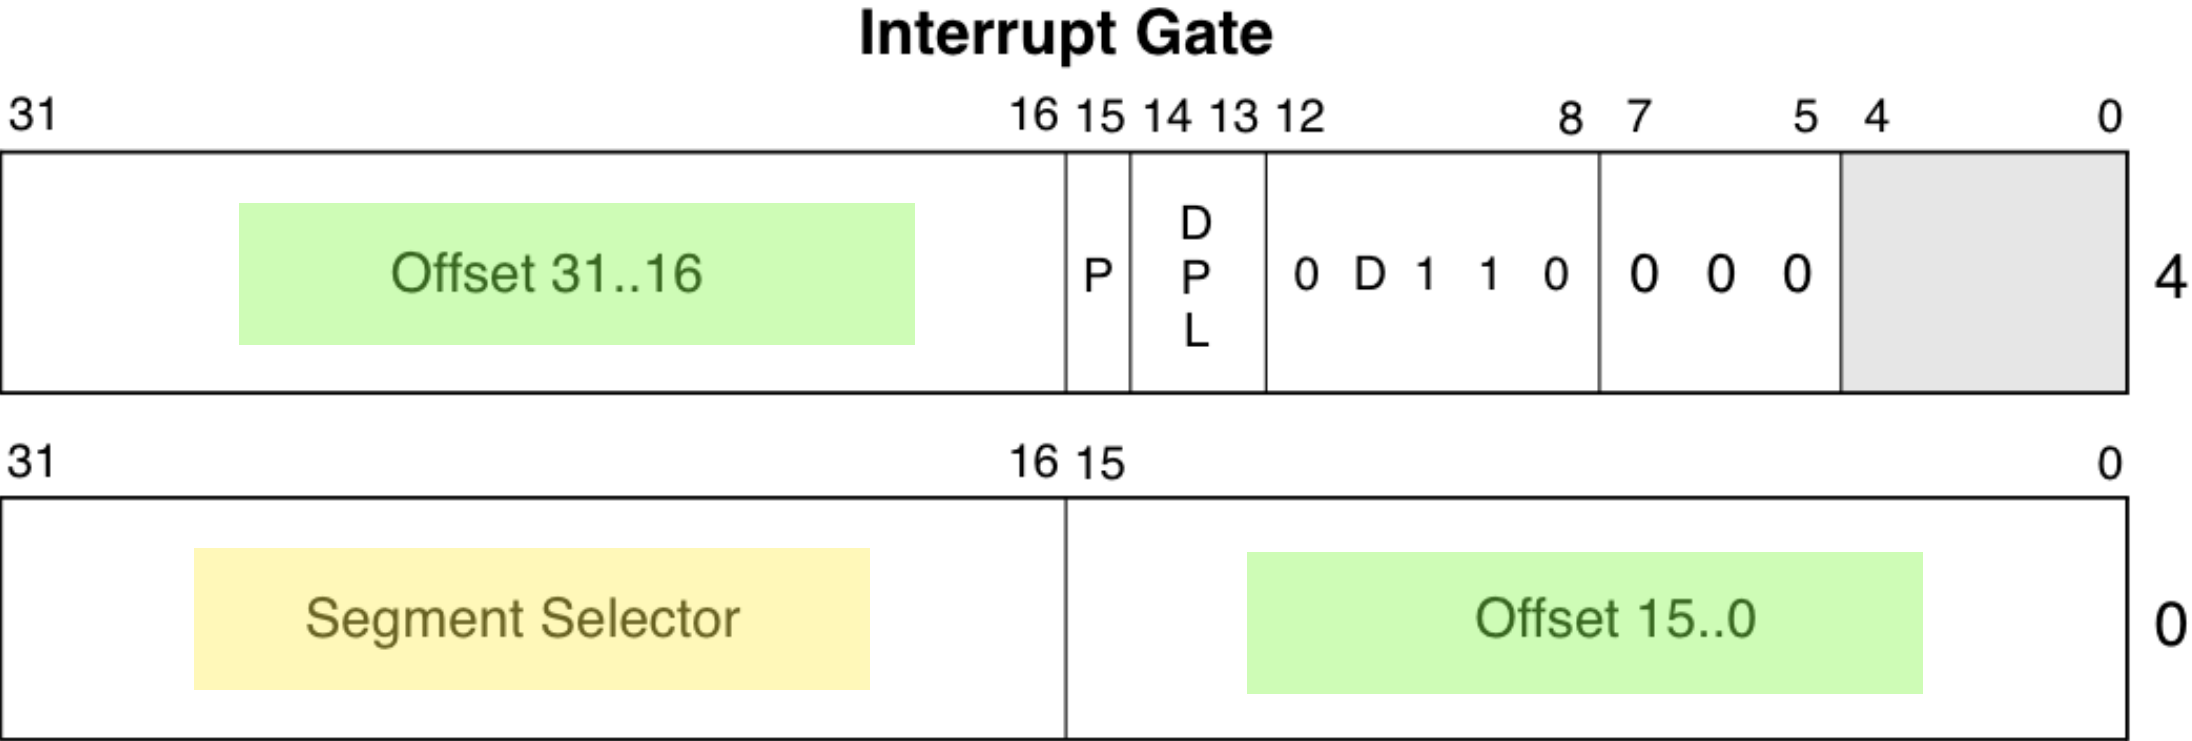
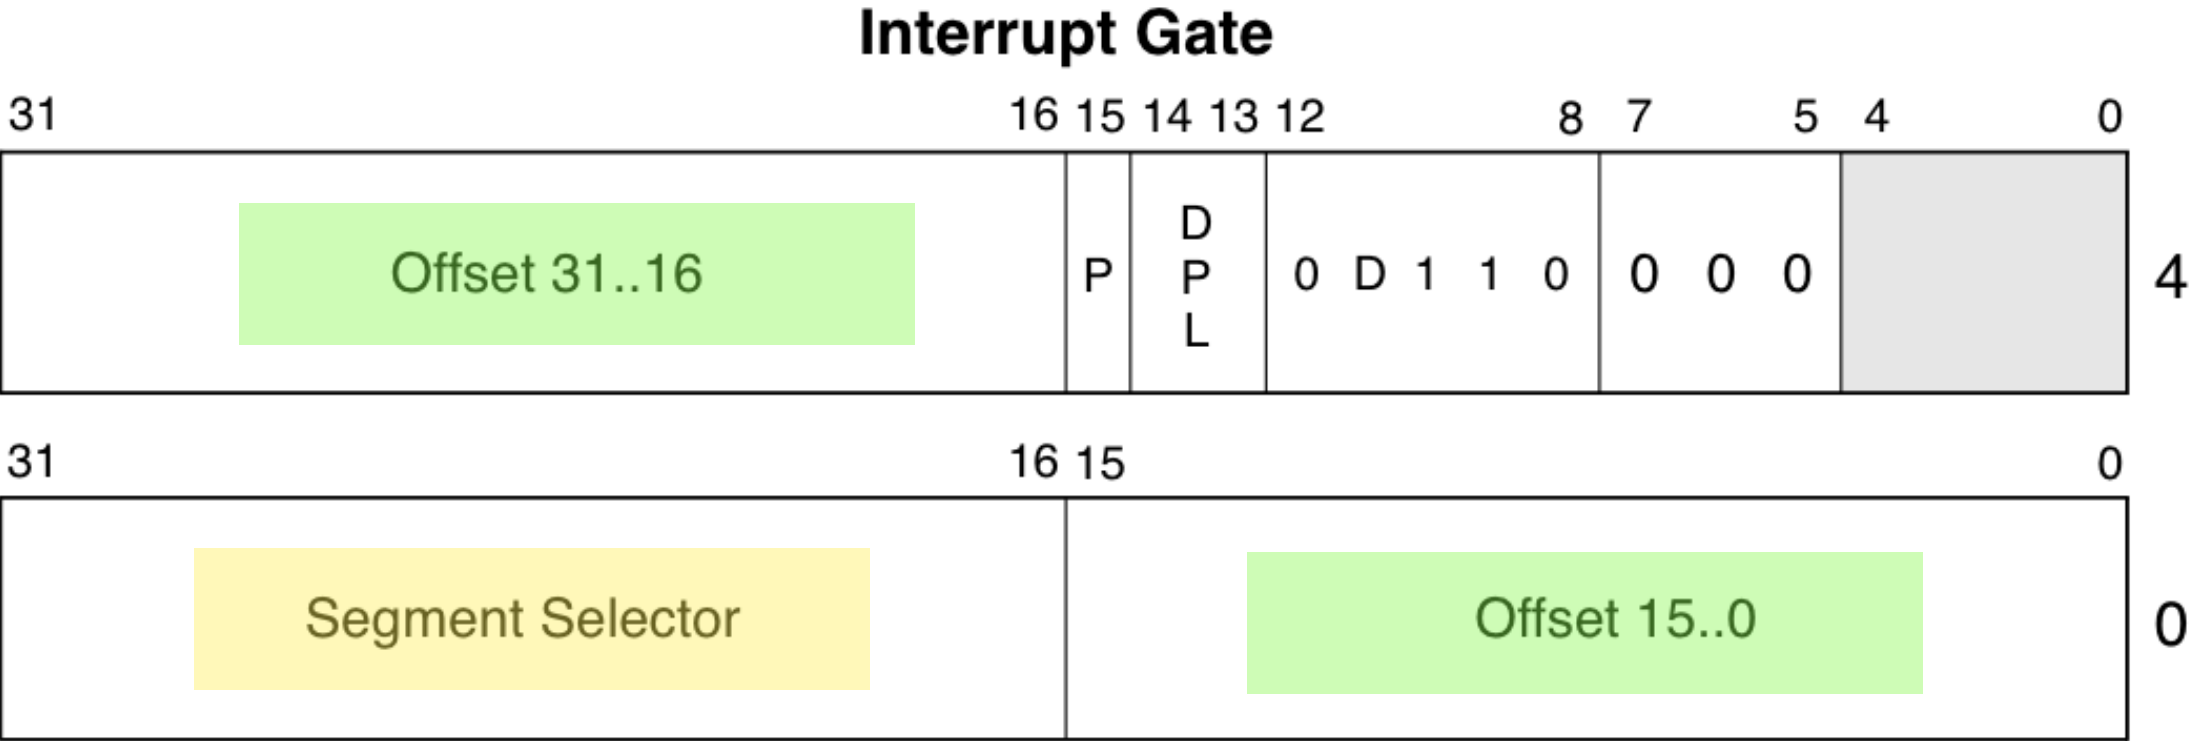


Figure 6-1. Relationship of the IDTR and IDT

Interrupt descriptor table (2)

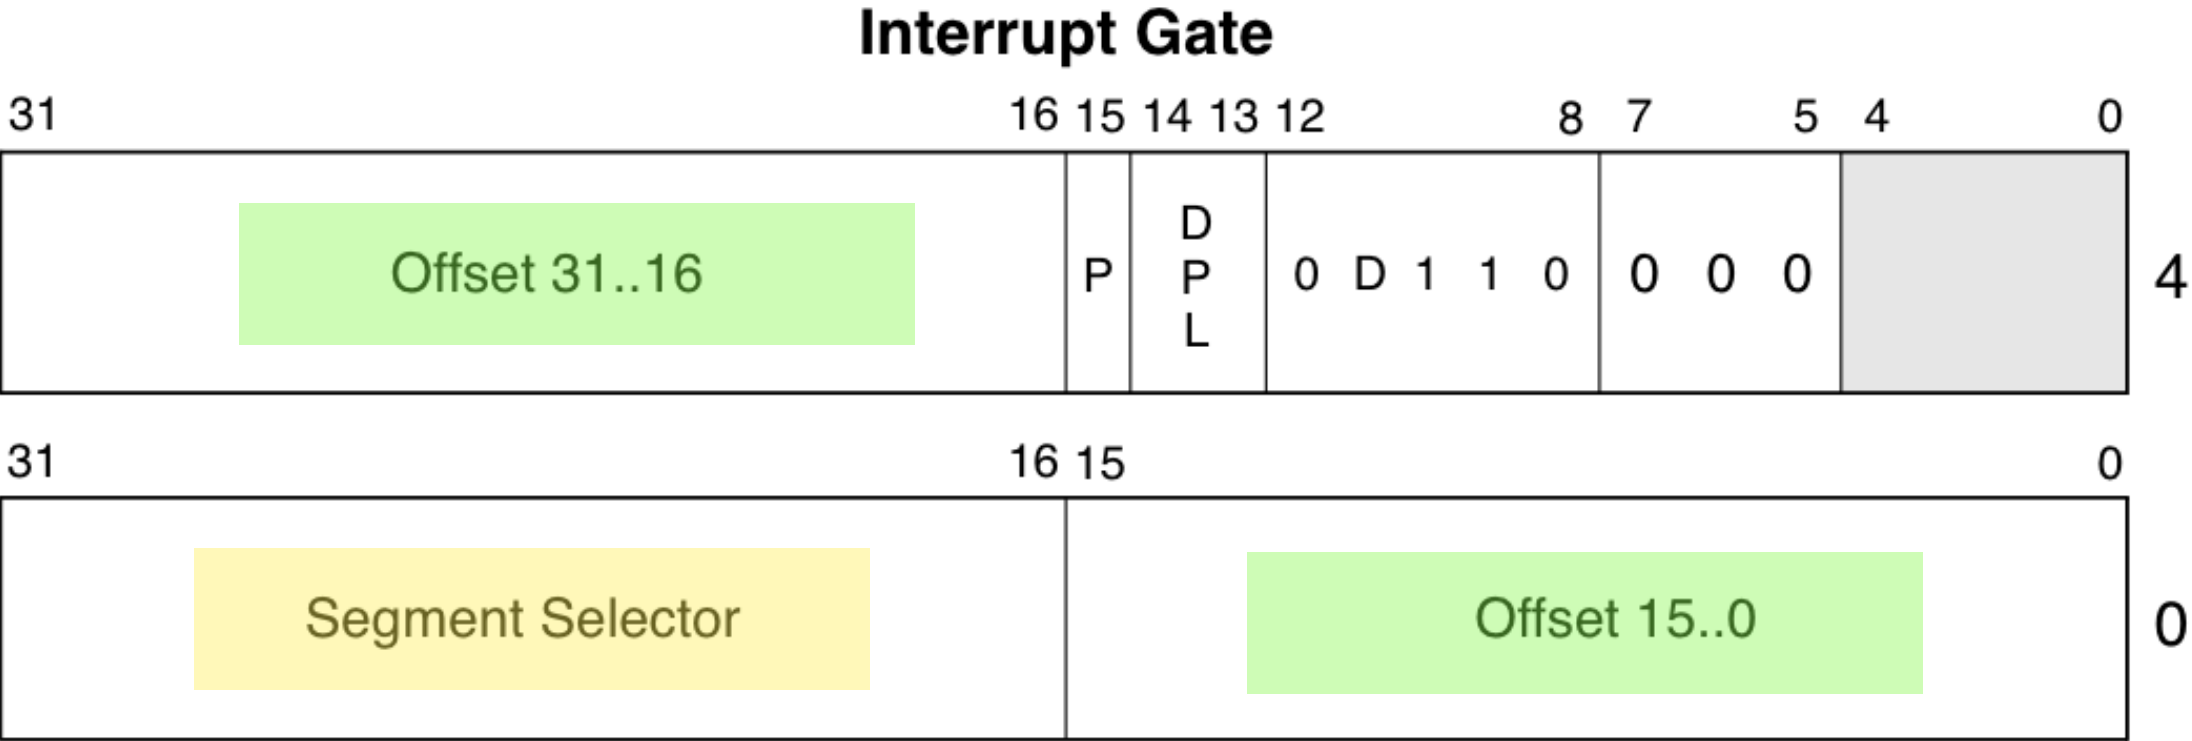


Interrupt descriptor table (2)



- Each IDT entry is 64-bits. Contains code segment and eip

Interrupt descriptor table (2)



- Each IDT entry is 64-bits. Contains **code segment** and **eip**
- When interrupt appears, hardware changes CS and EIP to the one pointed by IDT entry

Interrupt descriptor table (2)

- Each IDT entry is 64-bits. Contains **code segment** and **eip**
- When interrupt appears, hardware changes CS and EIP to the one pointed by IDT entry

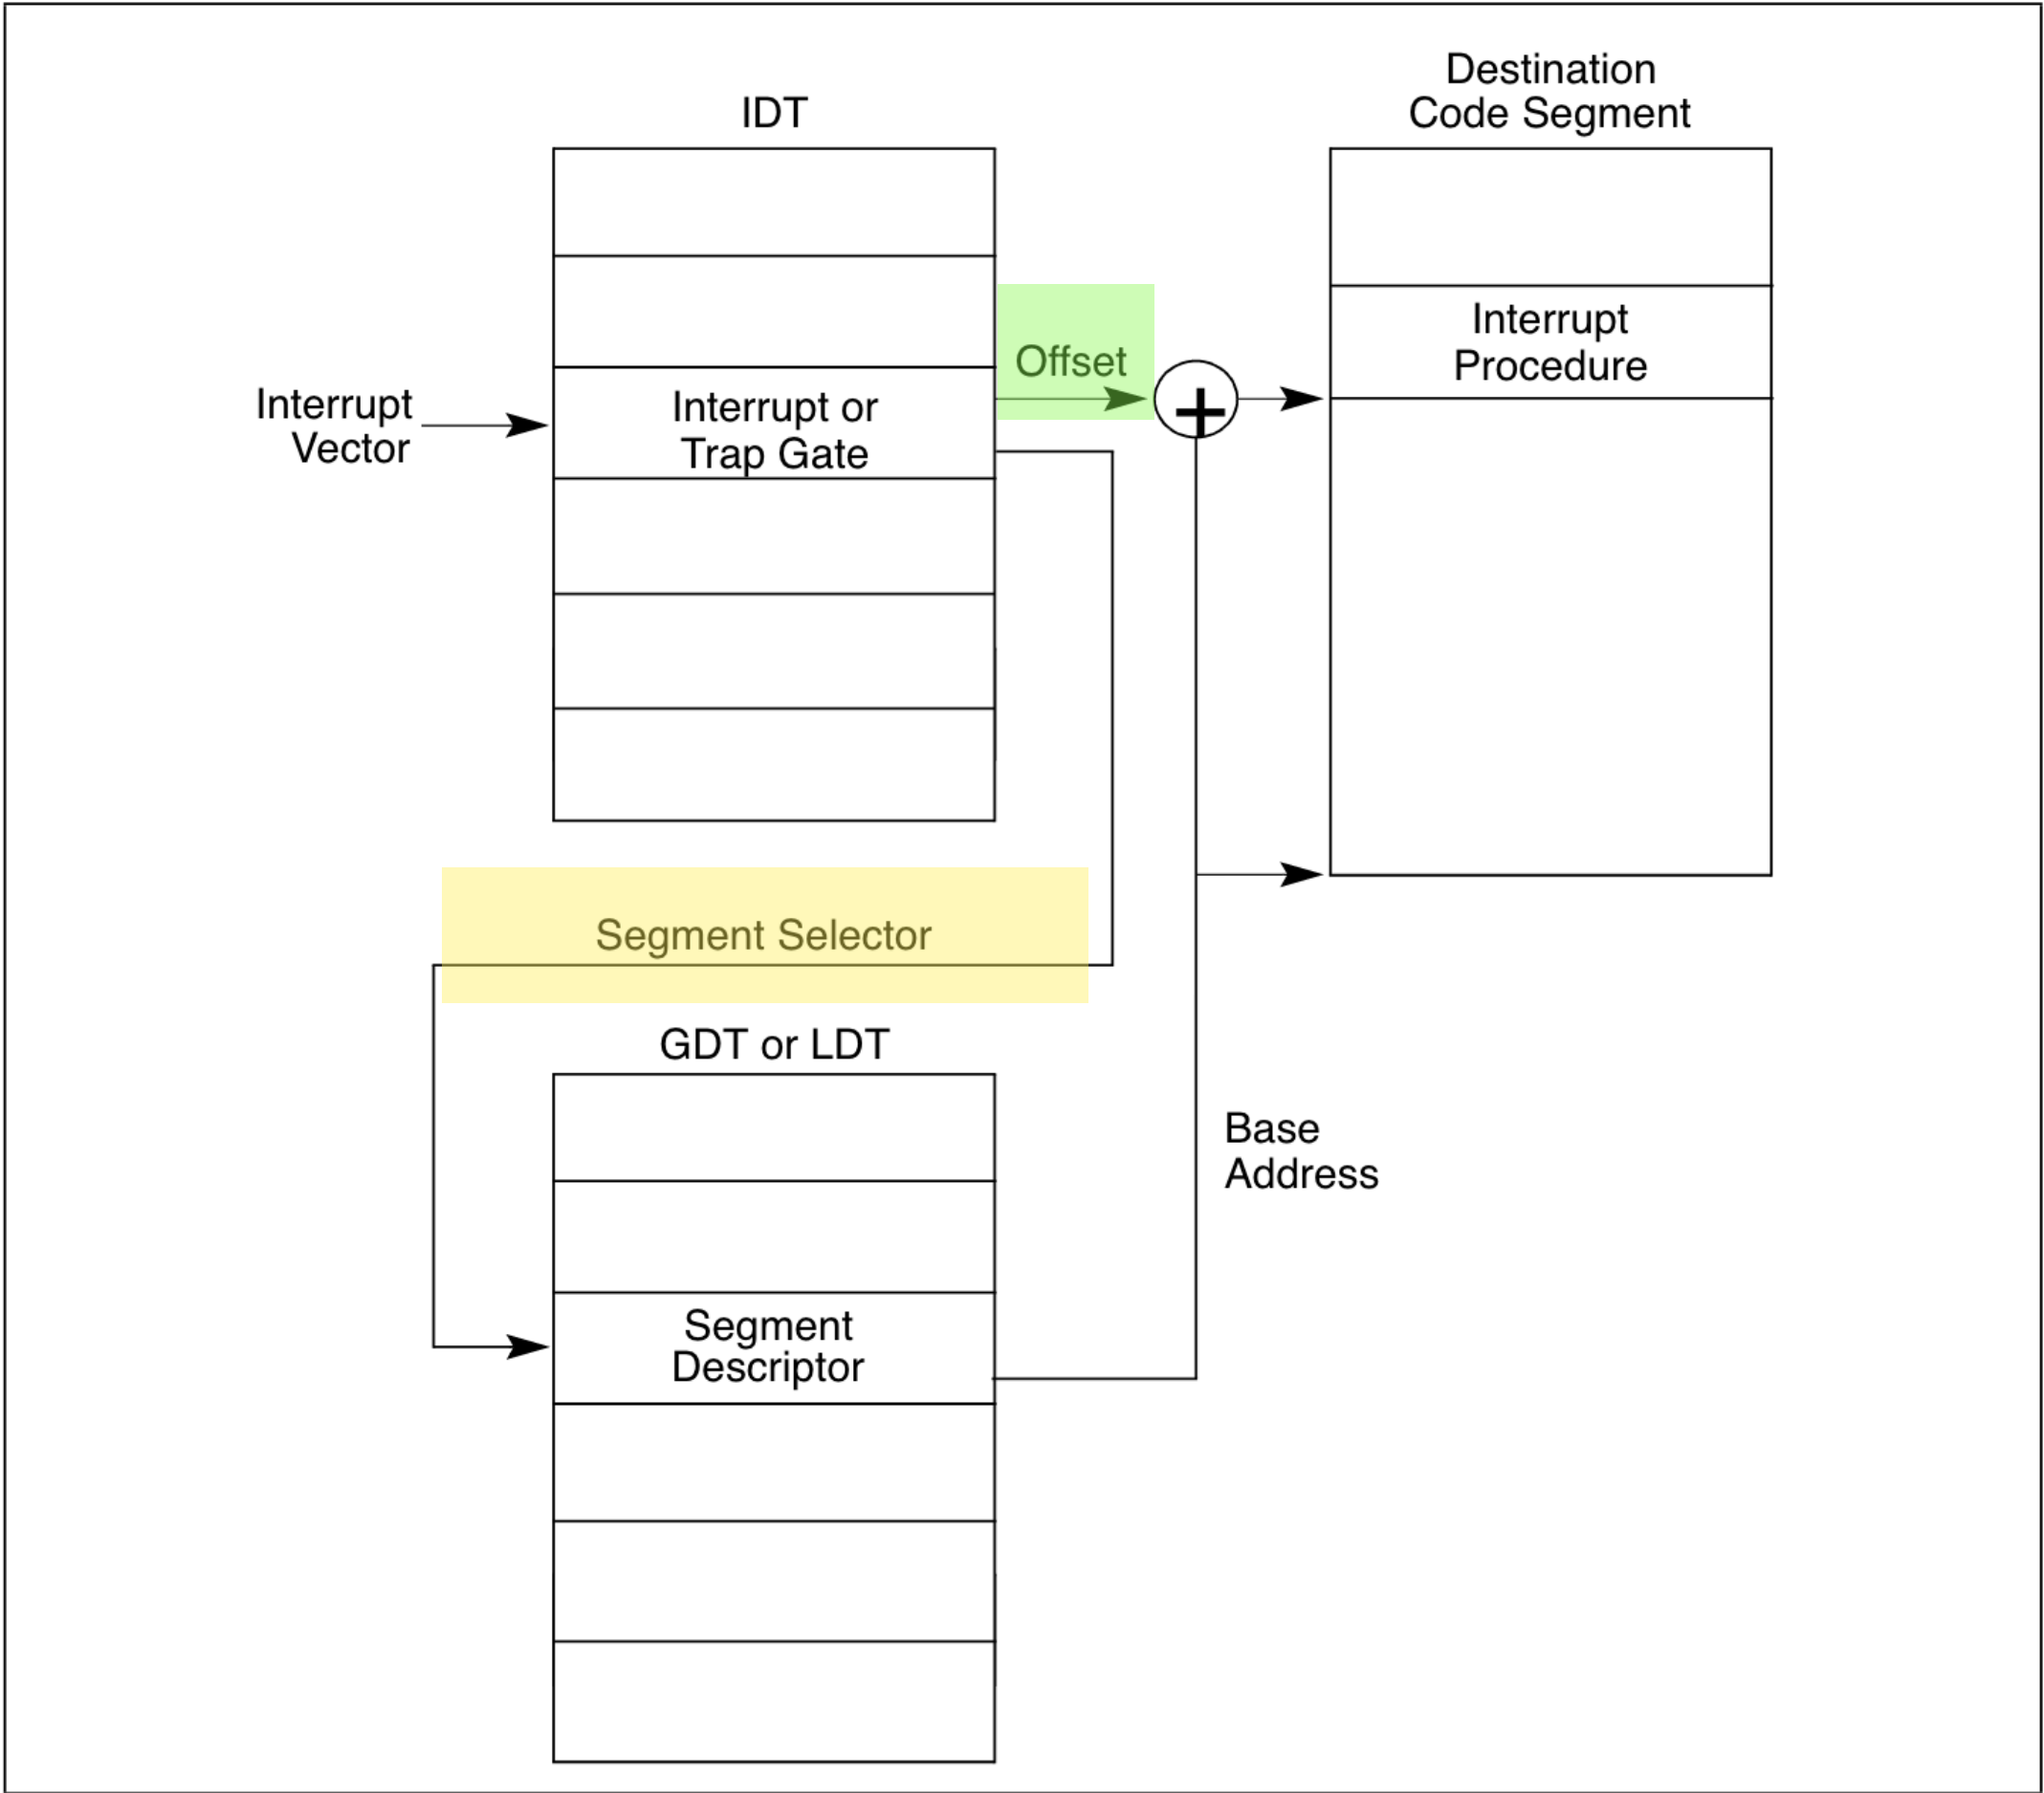
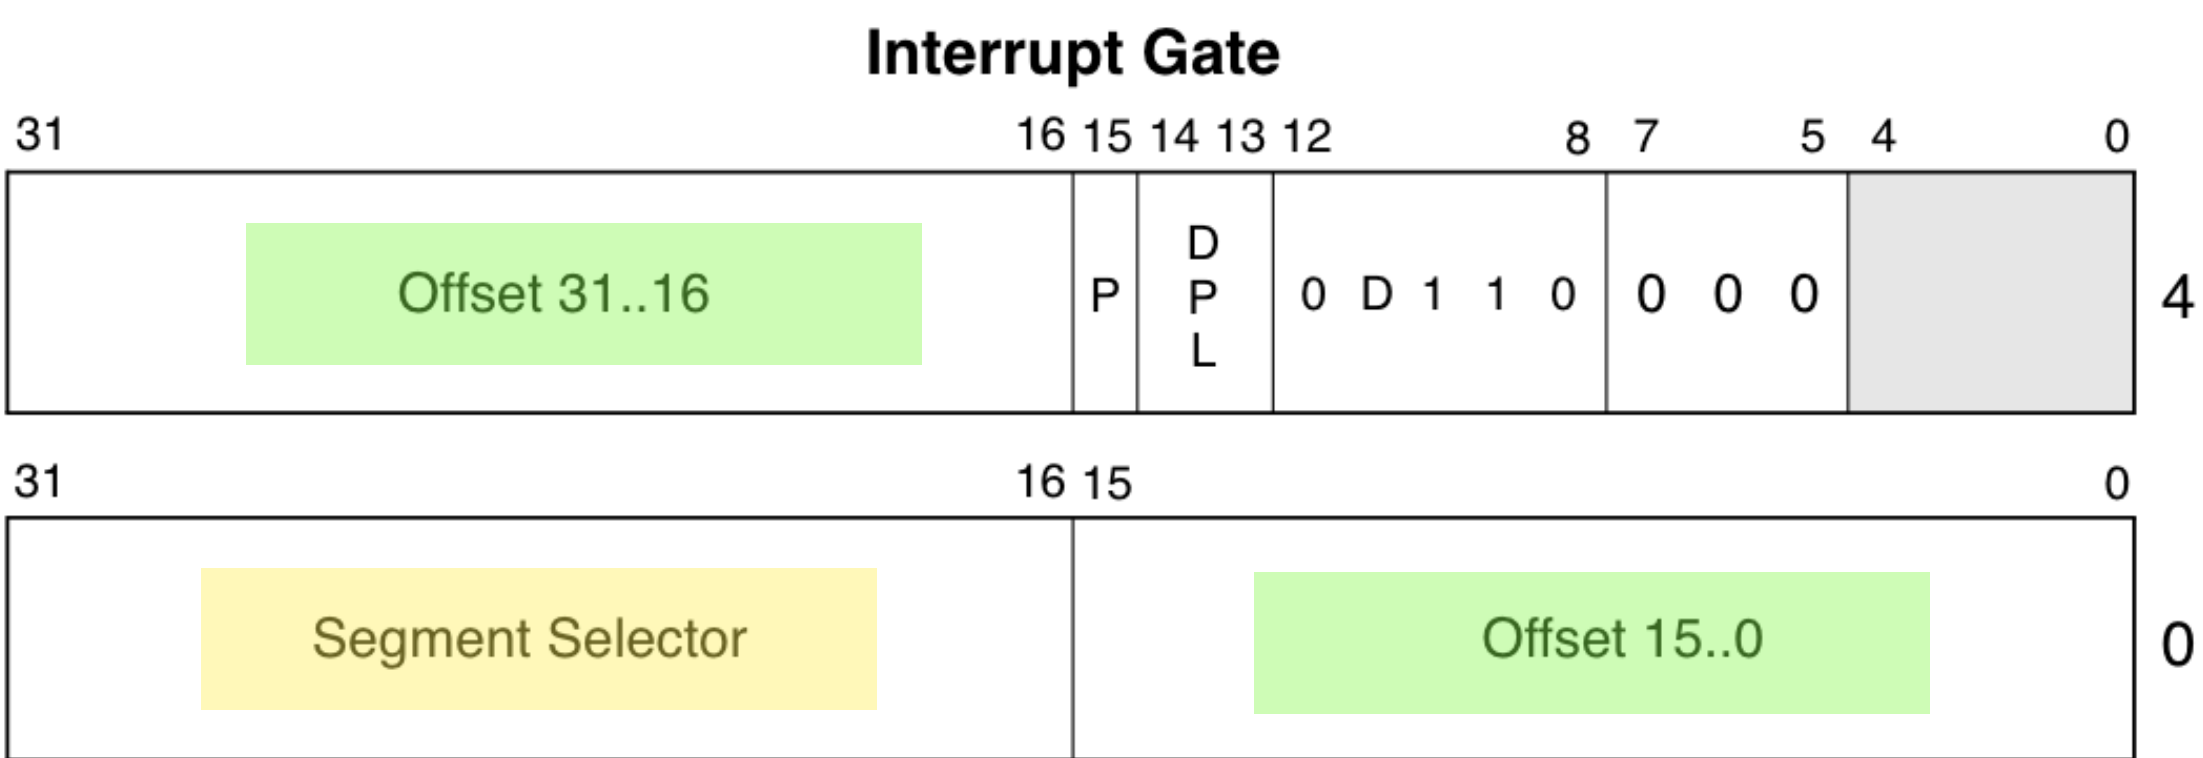
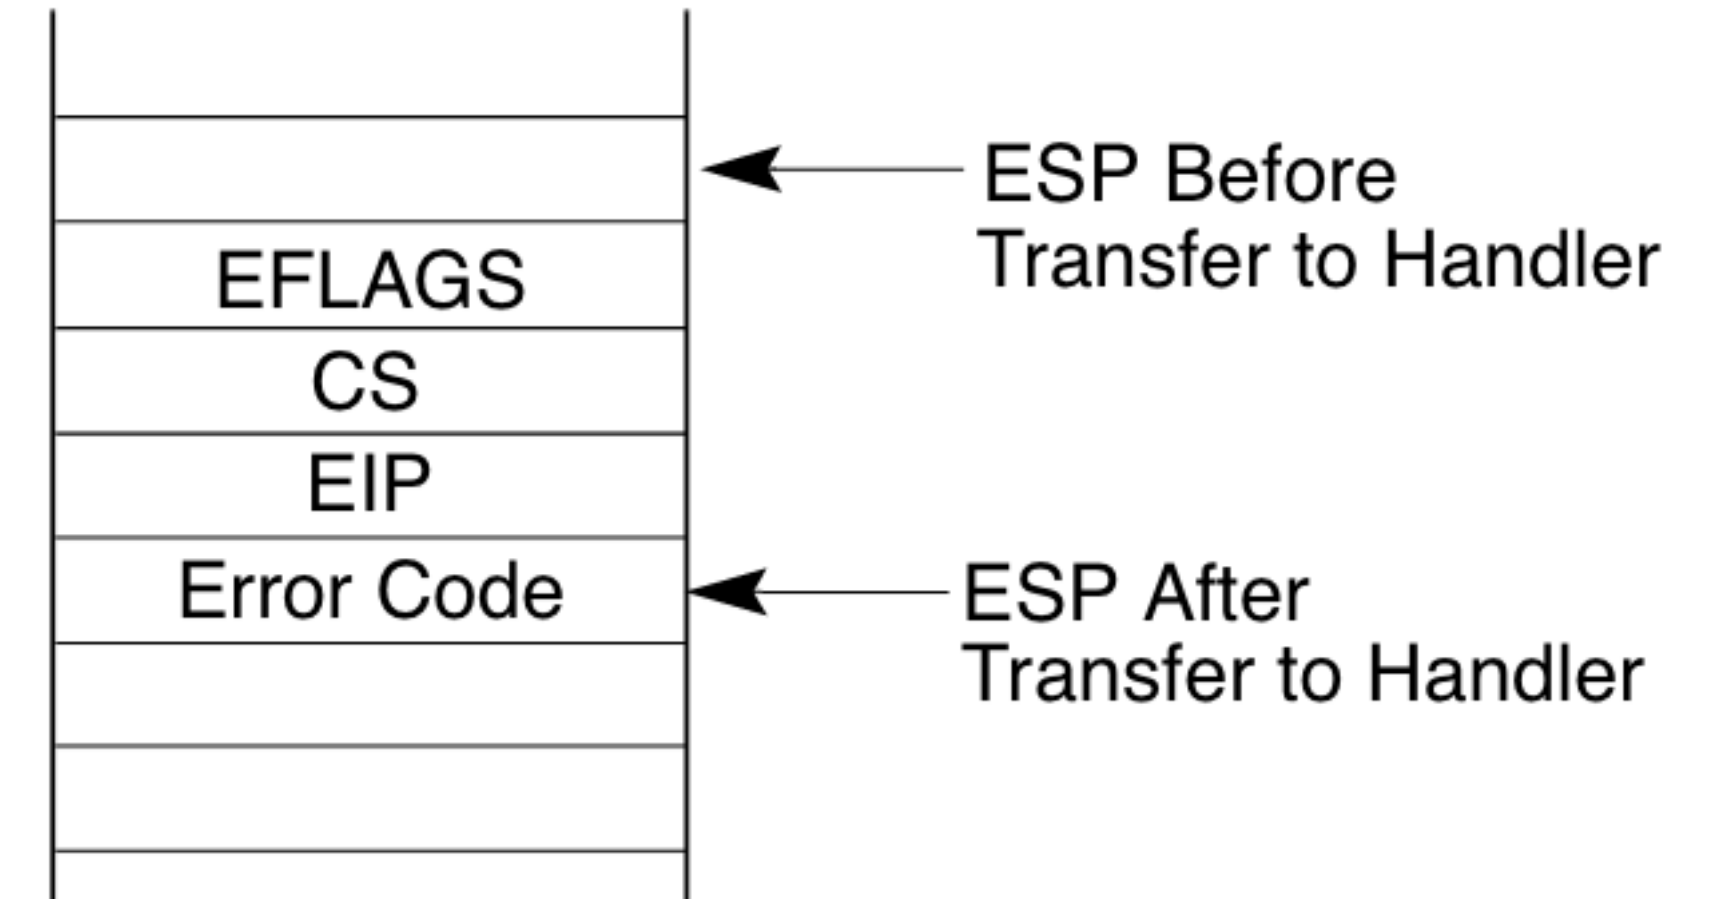


Figure 6-3. Interrupt Procedure Call

Interrupt handling

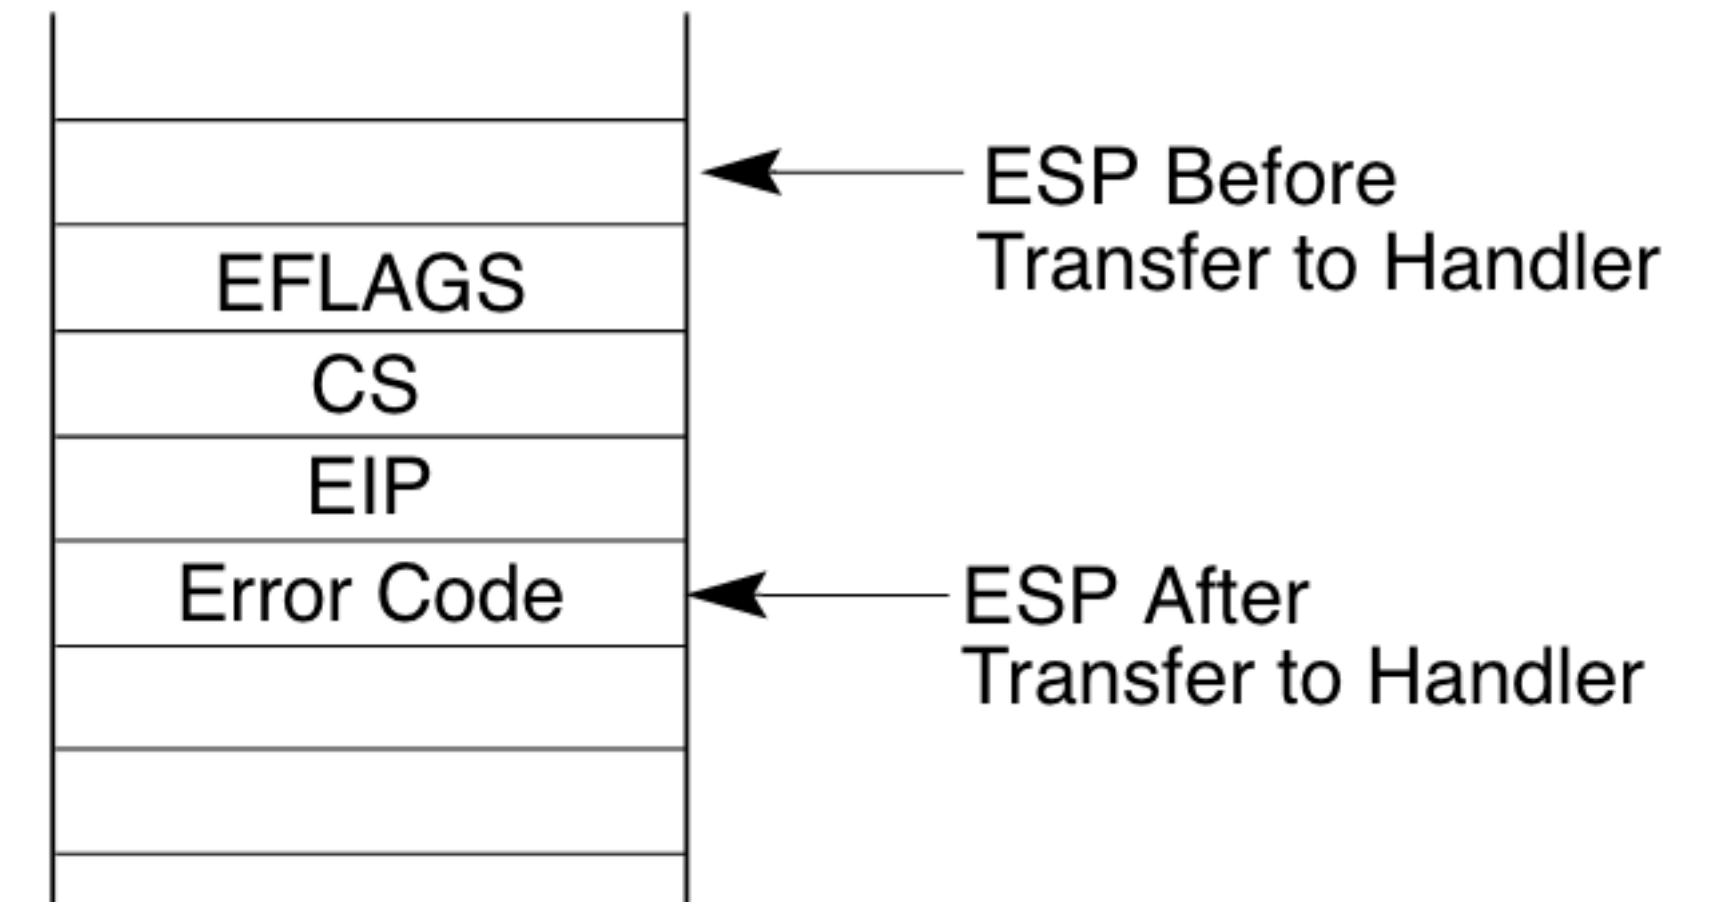
Interrupted Procedure's
and Handler's Stack



Interrupt handling

- On an interrupt, hardware pushes old EFLAGS, CS and EIP on the stack

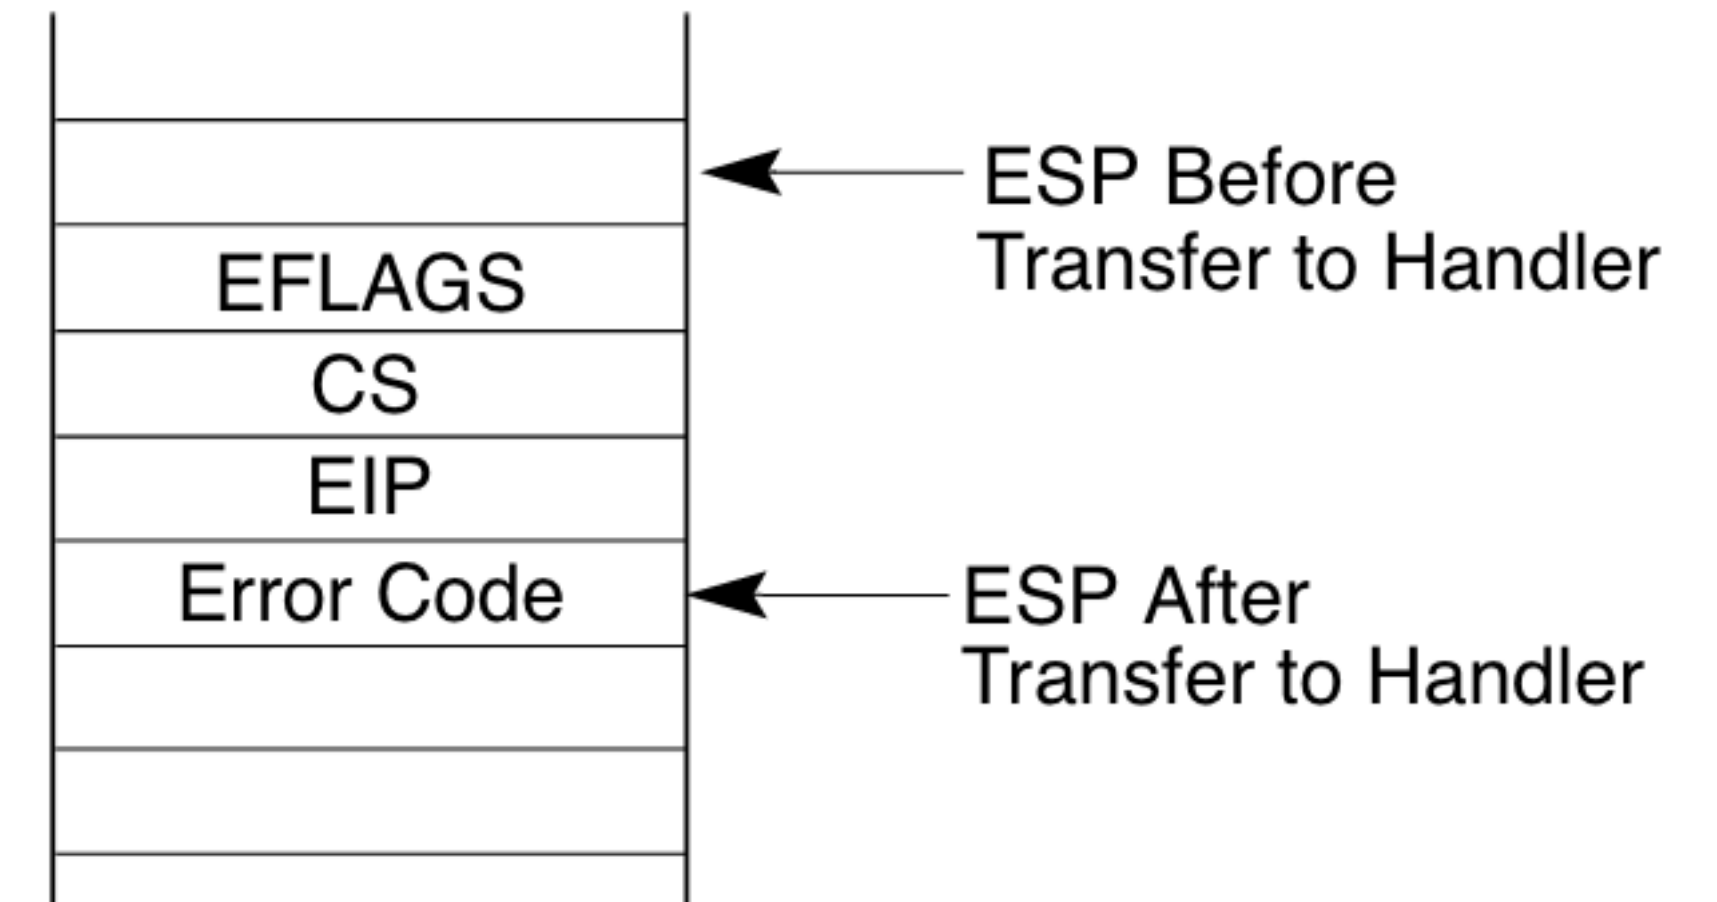
Interrupted Procedure's
and Handler's Stack



Interrupt handling

- On an interrupt, hardware pushes old EFLAGS, CS and EIP on the stack
- Jumps CS and EIP according to IDT

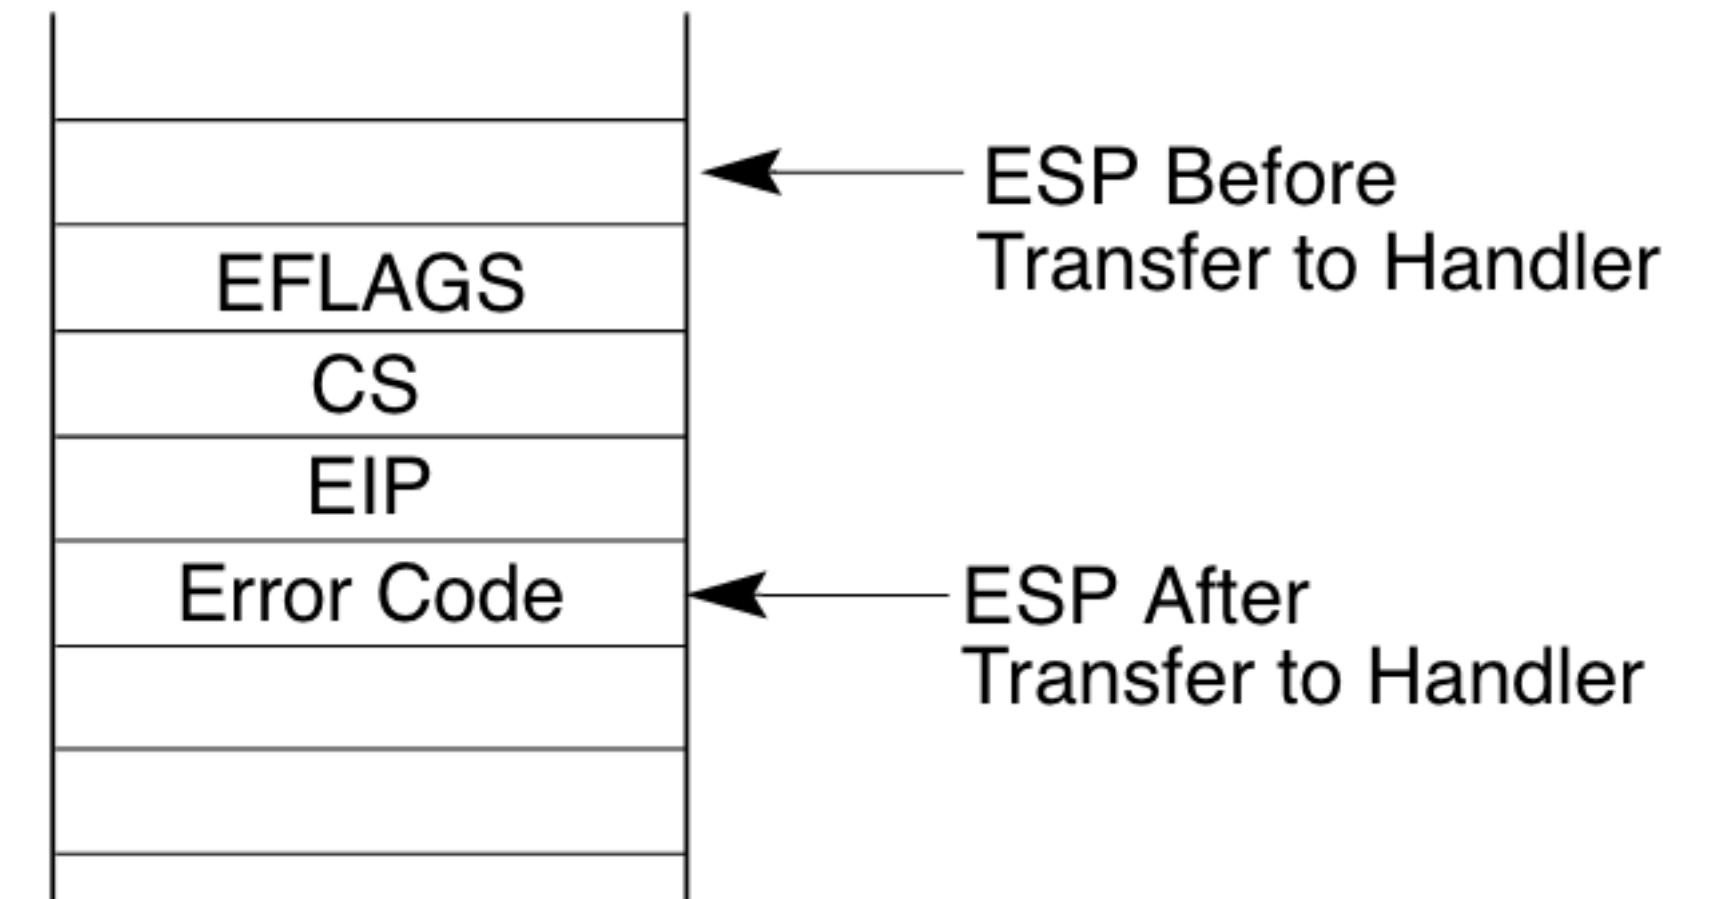
Interrupted Procedure's
and Handler's Stack



Interrupt handling

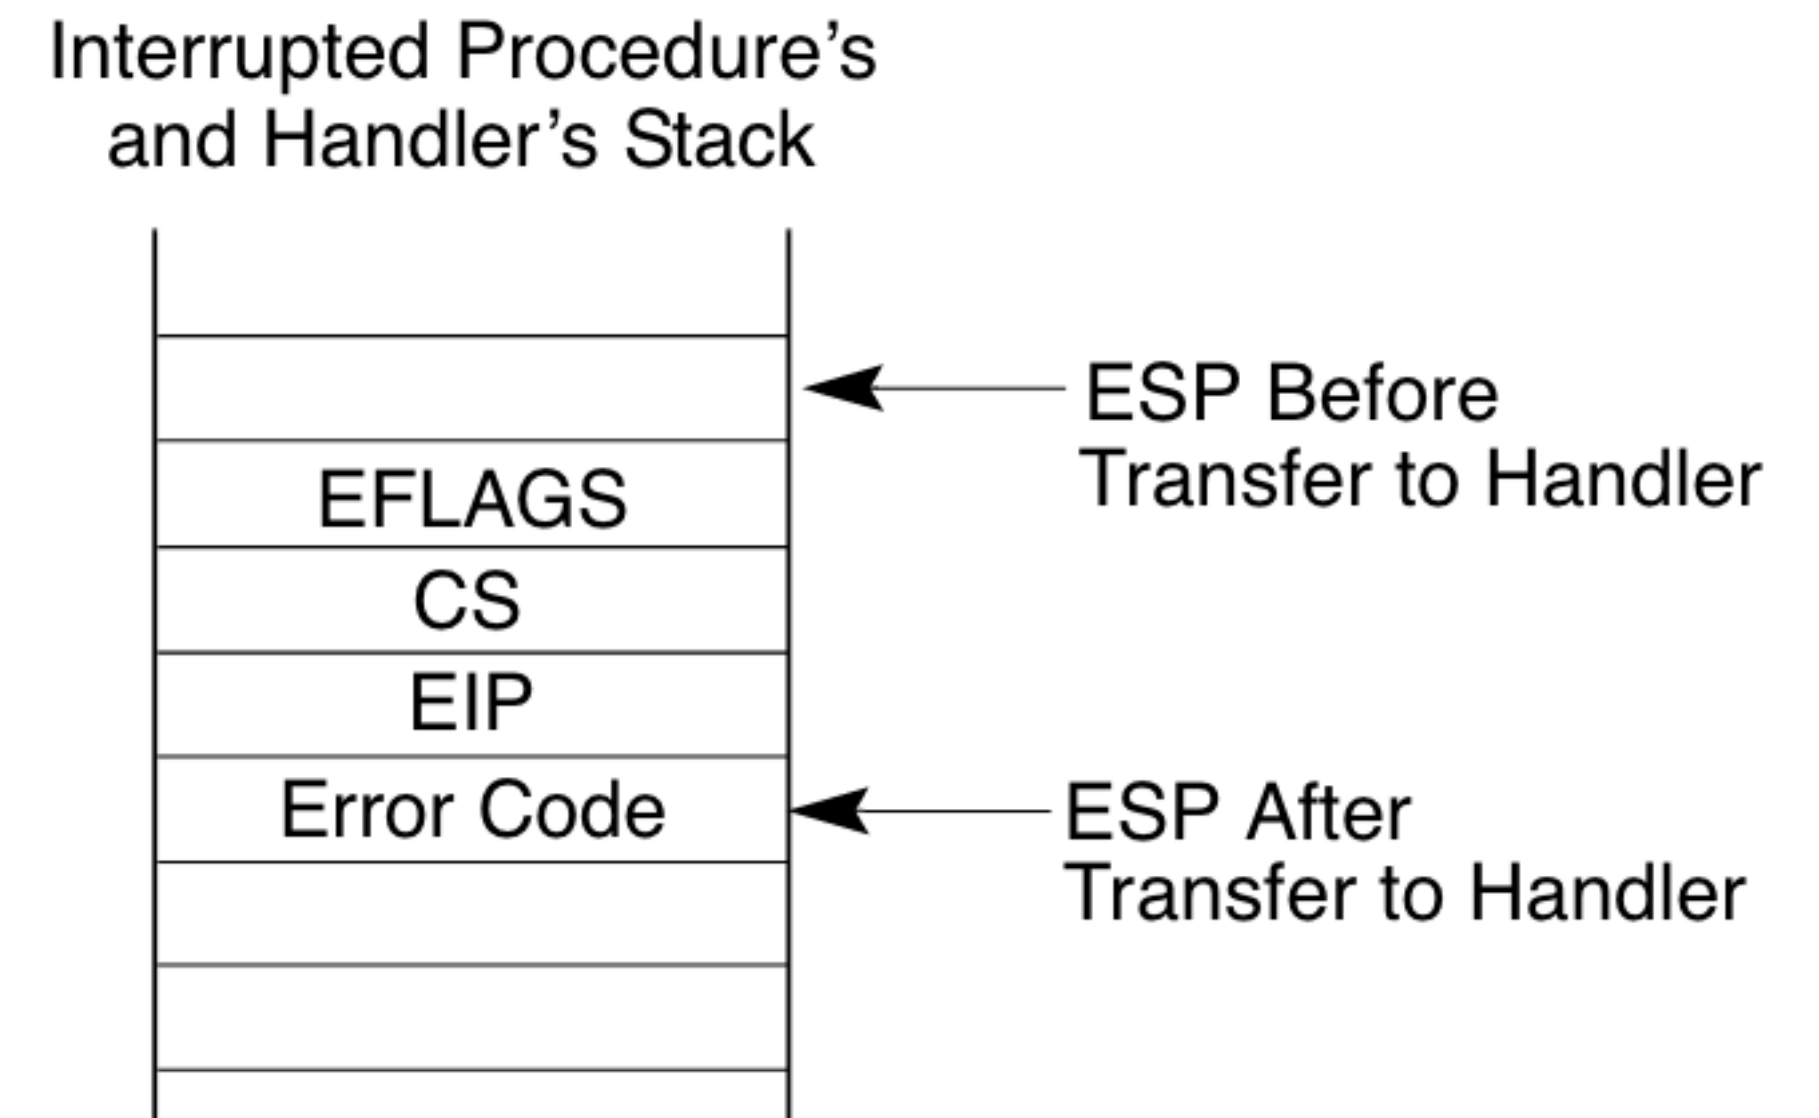
- On an interrupt, hardware pushes old EFLAGS, CS and EIP on the stack
- Jumps CS and EIP according to IDT
- IRET instruction (similar to RET instruction) restores CS, EIP, EFLAGS, ESP

Interrupted Procedure's
and Handler's Stack



Interrupt handling

- On an interrupt, hardware pushes old EFLAGS, CS and EIP on the stack
- Jumps CS and EIP according to IDT
- IRET instruction (similar to RET instruction) restores CS, EIP, EFLAGS, ESP
- Interrupt handler may push more registers, like eax etc. on the stack.



Code walkthrough

p4-traps

- vectors.pl creates 256 IDT entries. 'i'th entry write 'i' on top of the stack and jumps to 'alltraps'
- main.c calls tvinit and idtinit to setup interrupt descriptor table to populate the 256 entries and point IDTR to IDT. It calls sti to receive interrupts.
- 'alltraps' in trapasm.S runs 'pushal' to save general purpose registers. Then it calls 'trap' with the trapframe.
- 'trap' in 'trapasm.S' reads trapno saved by vectors.S to find out which interrupt occurred. It handles timer and spurious interrupts. It signals EOI to LAPIC when it is done with interrupt.
- trapasm recovers registers with popal, backs up esp above err code and trap number, executes IRET to jump back to whatever OS was doing earlier

Visualizing interrupt handling

eip → for(;;)
;

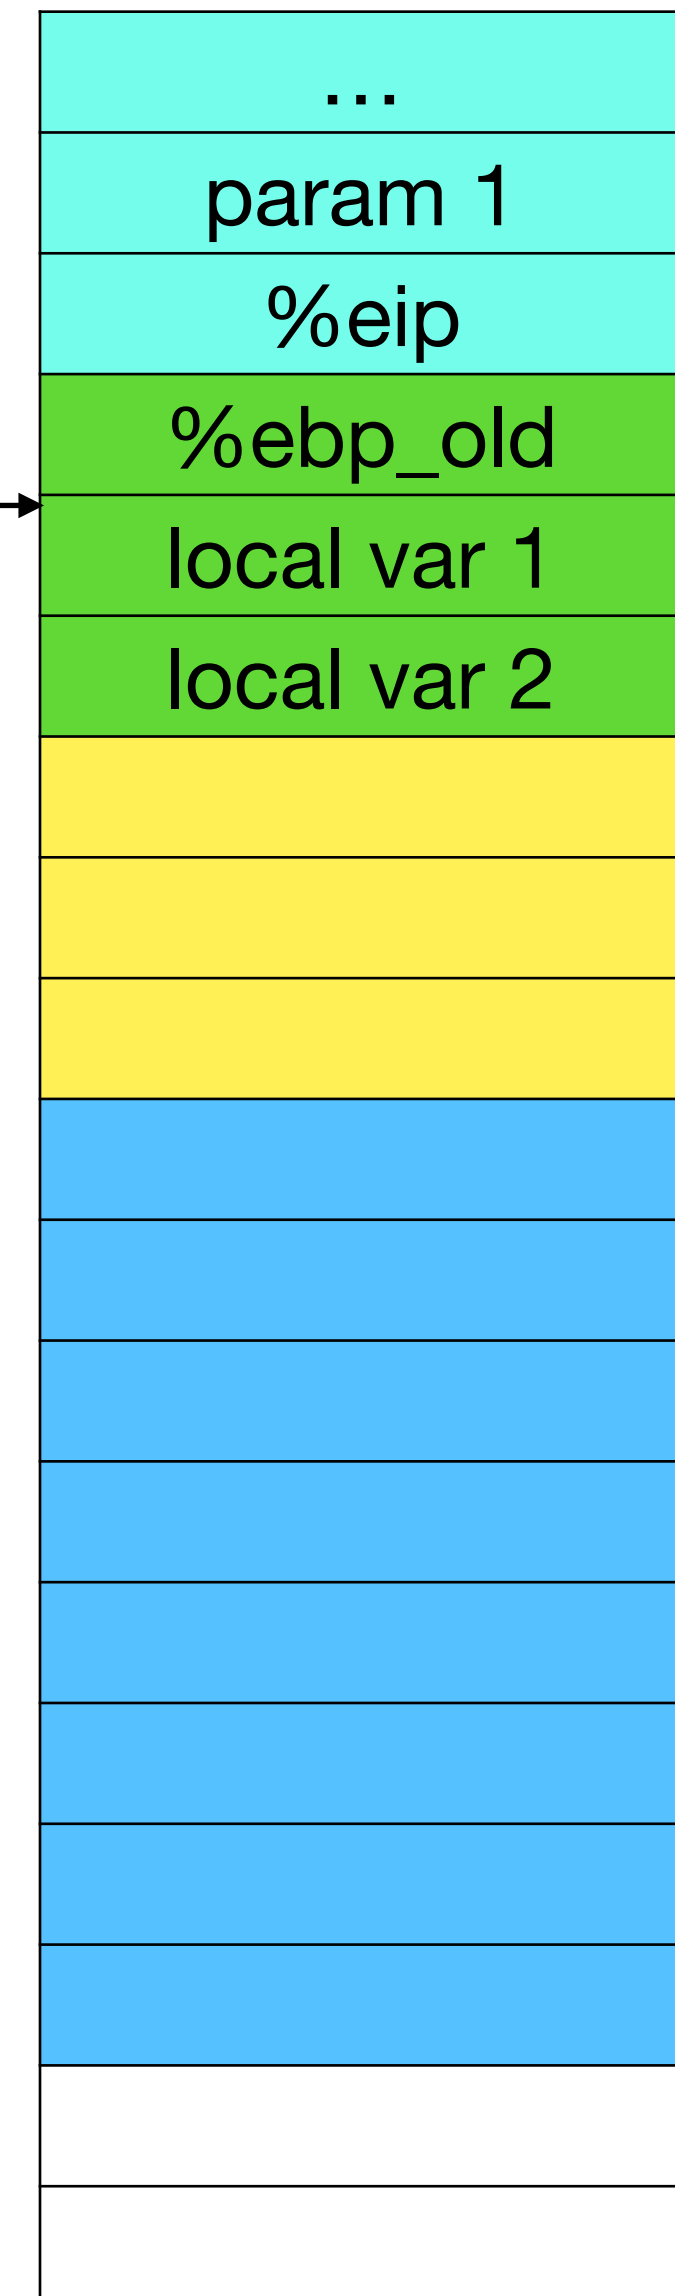
```
trap.c
void
trap(struct trapframe *tf)
{
    switch(tf->trapno){
    case T_IRQ0 + IRQ_TIMER:
        ticks++;
        cprintf("Tick! %d\n", ticks);
        lapiceoi();
        ..
    }
    return
}
```

```

vectors.S
.globl vector0
vector0:
    pushl $0
    pushl $0
    jmp alltraps

```

```
trapasm.S
alltraps:
    pushal
    pushl %esp
    call trap
    addl $4, %esp
    popal
    addl $0x8, %esp
    iret
```



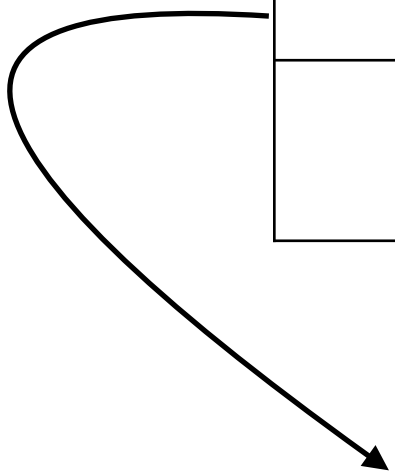
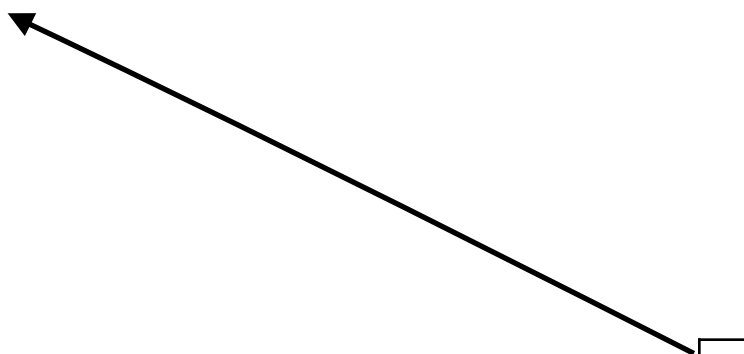
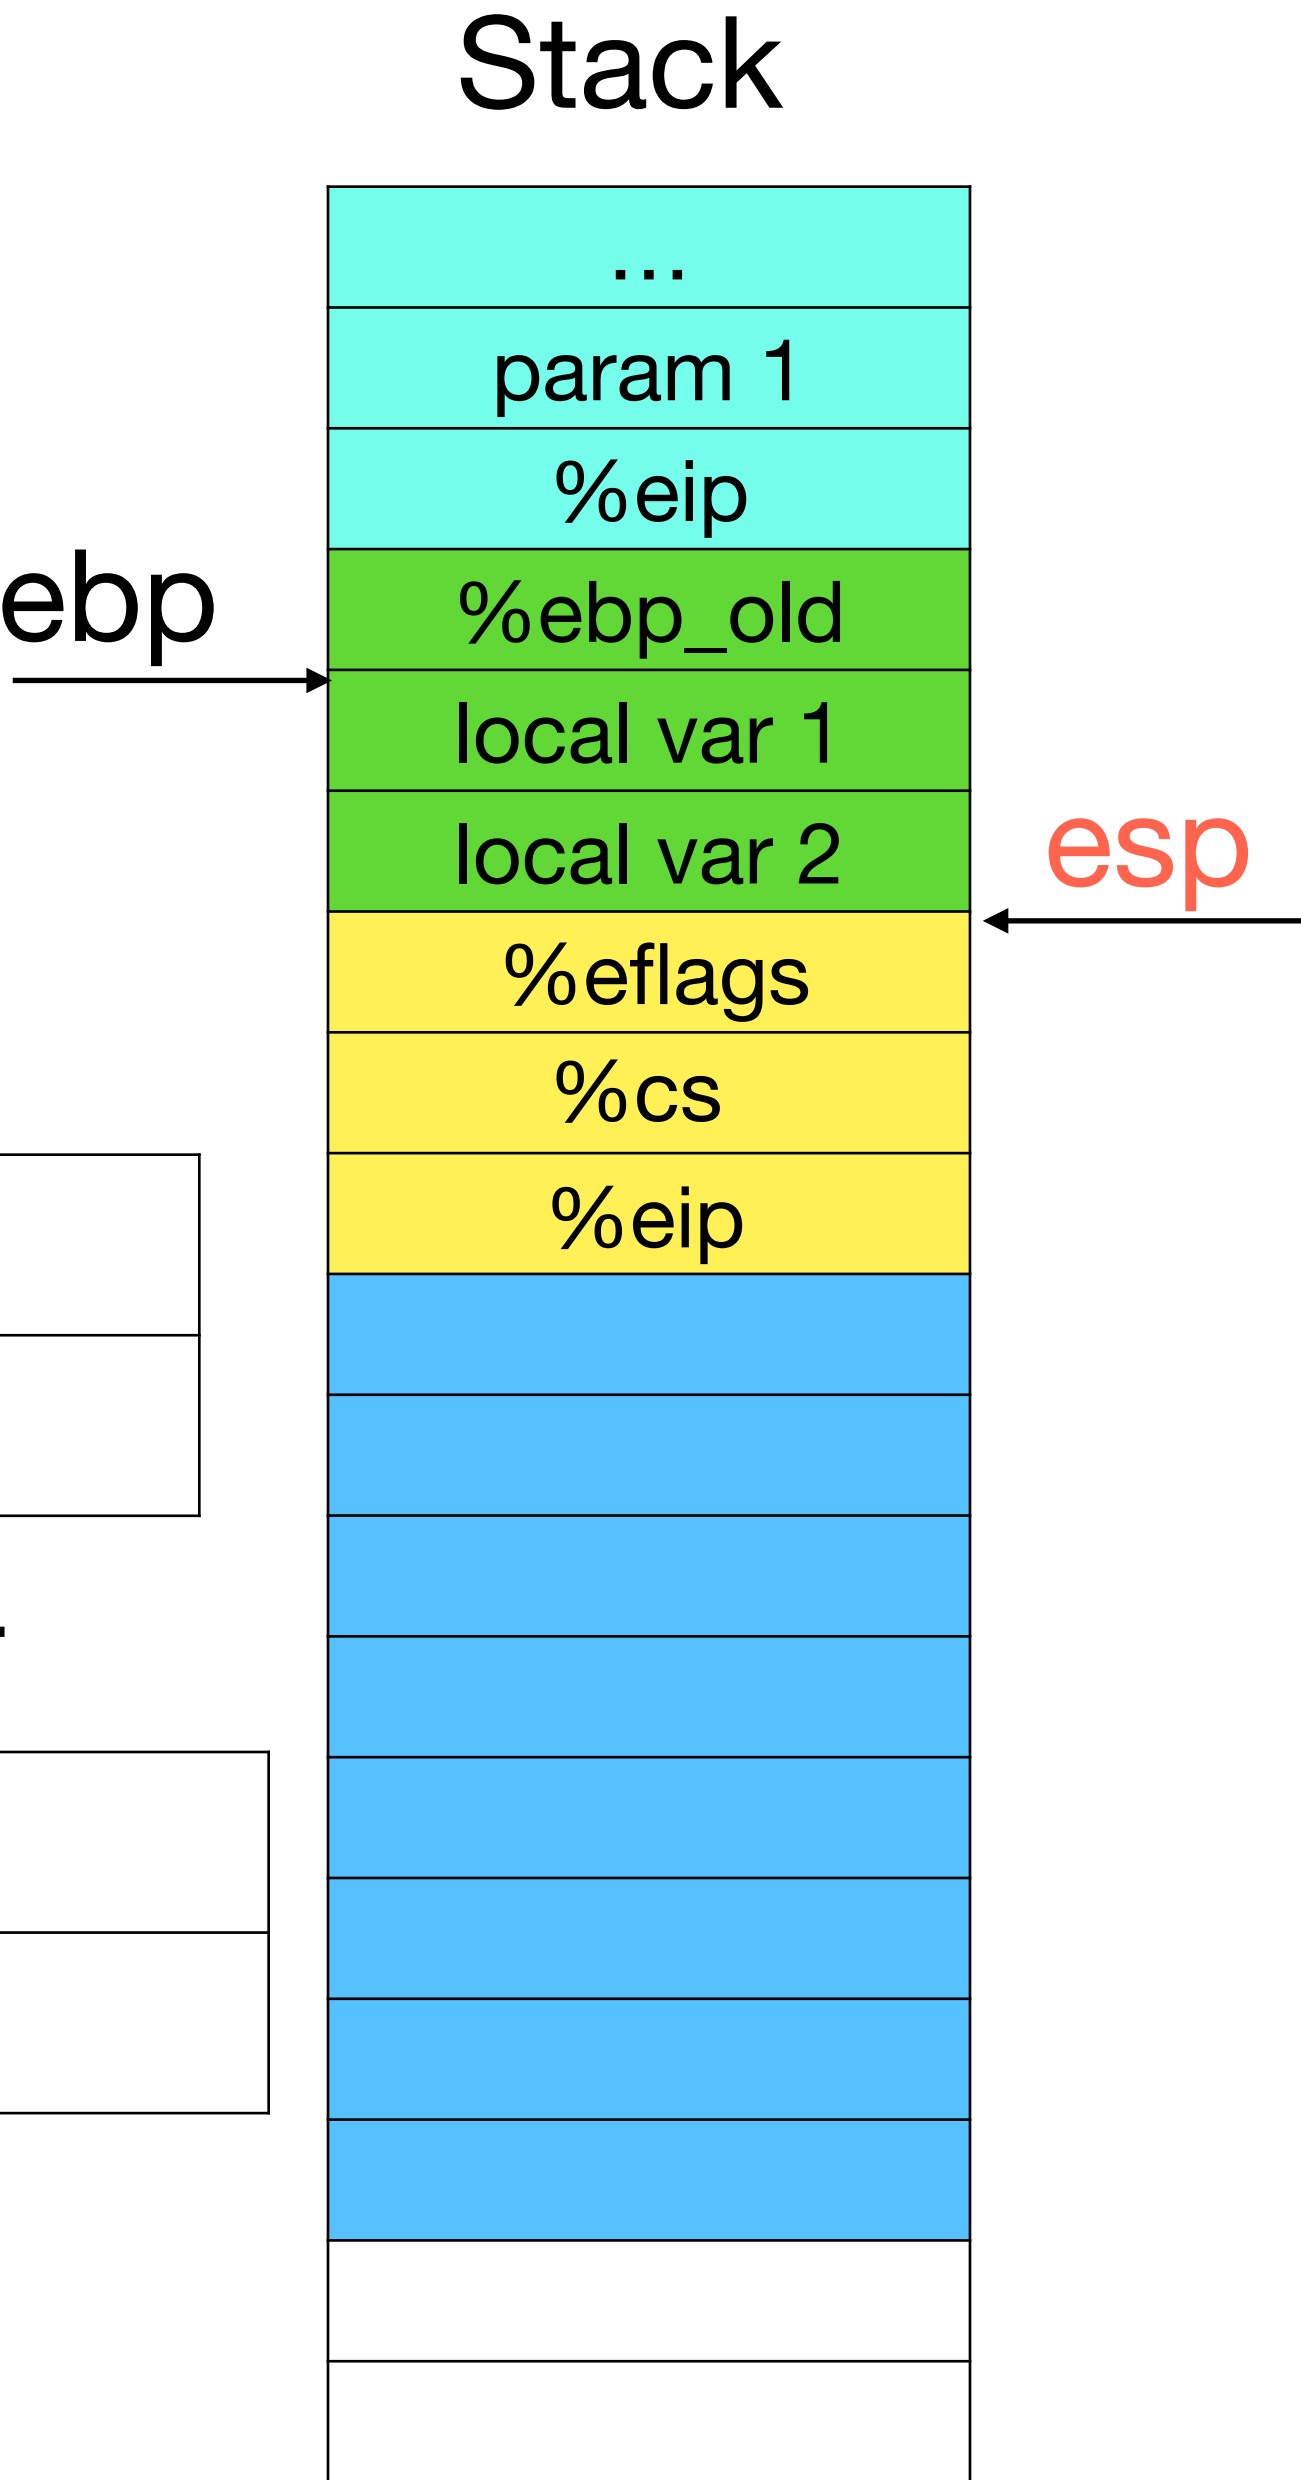
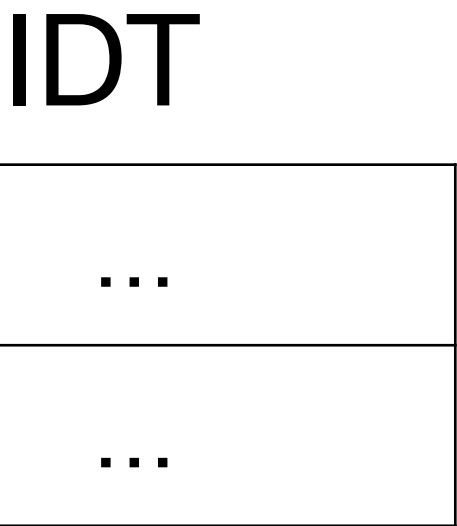
Visualizing interrupt handling

```
eip → for(;;)
;

trap.c
void
trap(struct trapframe *tf)
{
    switch(tf->trapno){
    case T_IRQ0 + IRQ_TIMER:
        ticks++;
        cprintf("Tick! %d\n\0", ticks);
        lapiceoi();
    ..
    return
```

```
vectors.S
.globl vector0
vector0:
    pushl $0
    pushl $0
    jmp alltraps
```

```
trapasm.S
alltraps:
    pushal
    pushl %esp
    call trap
    addl $4, %esp
    popal
    addl $0x8, %esp
    iret
```



ebp →

esp ←

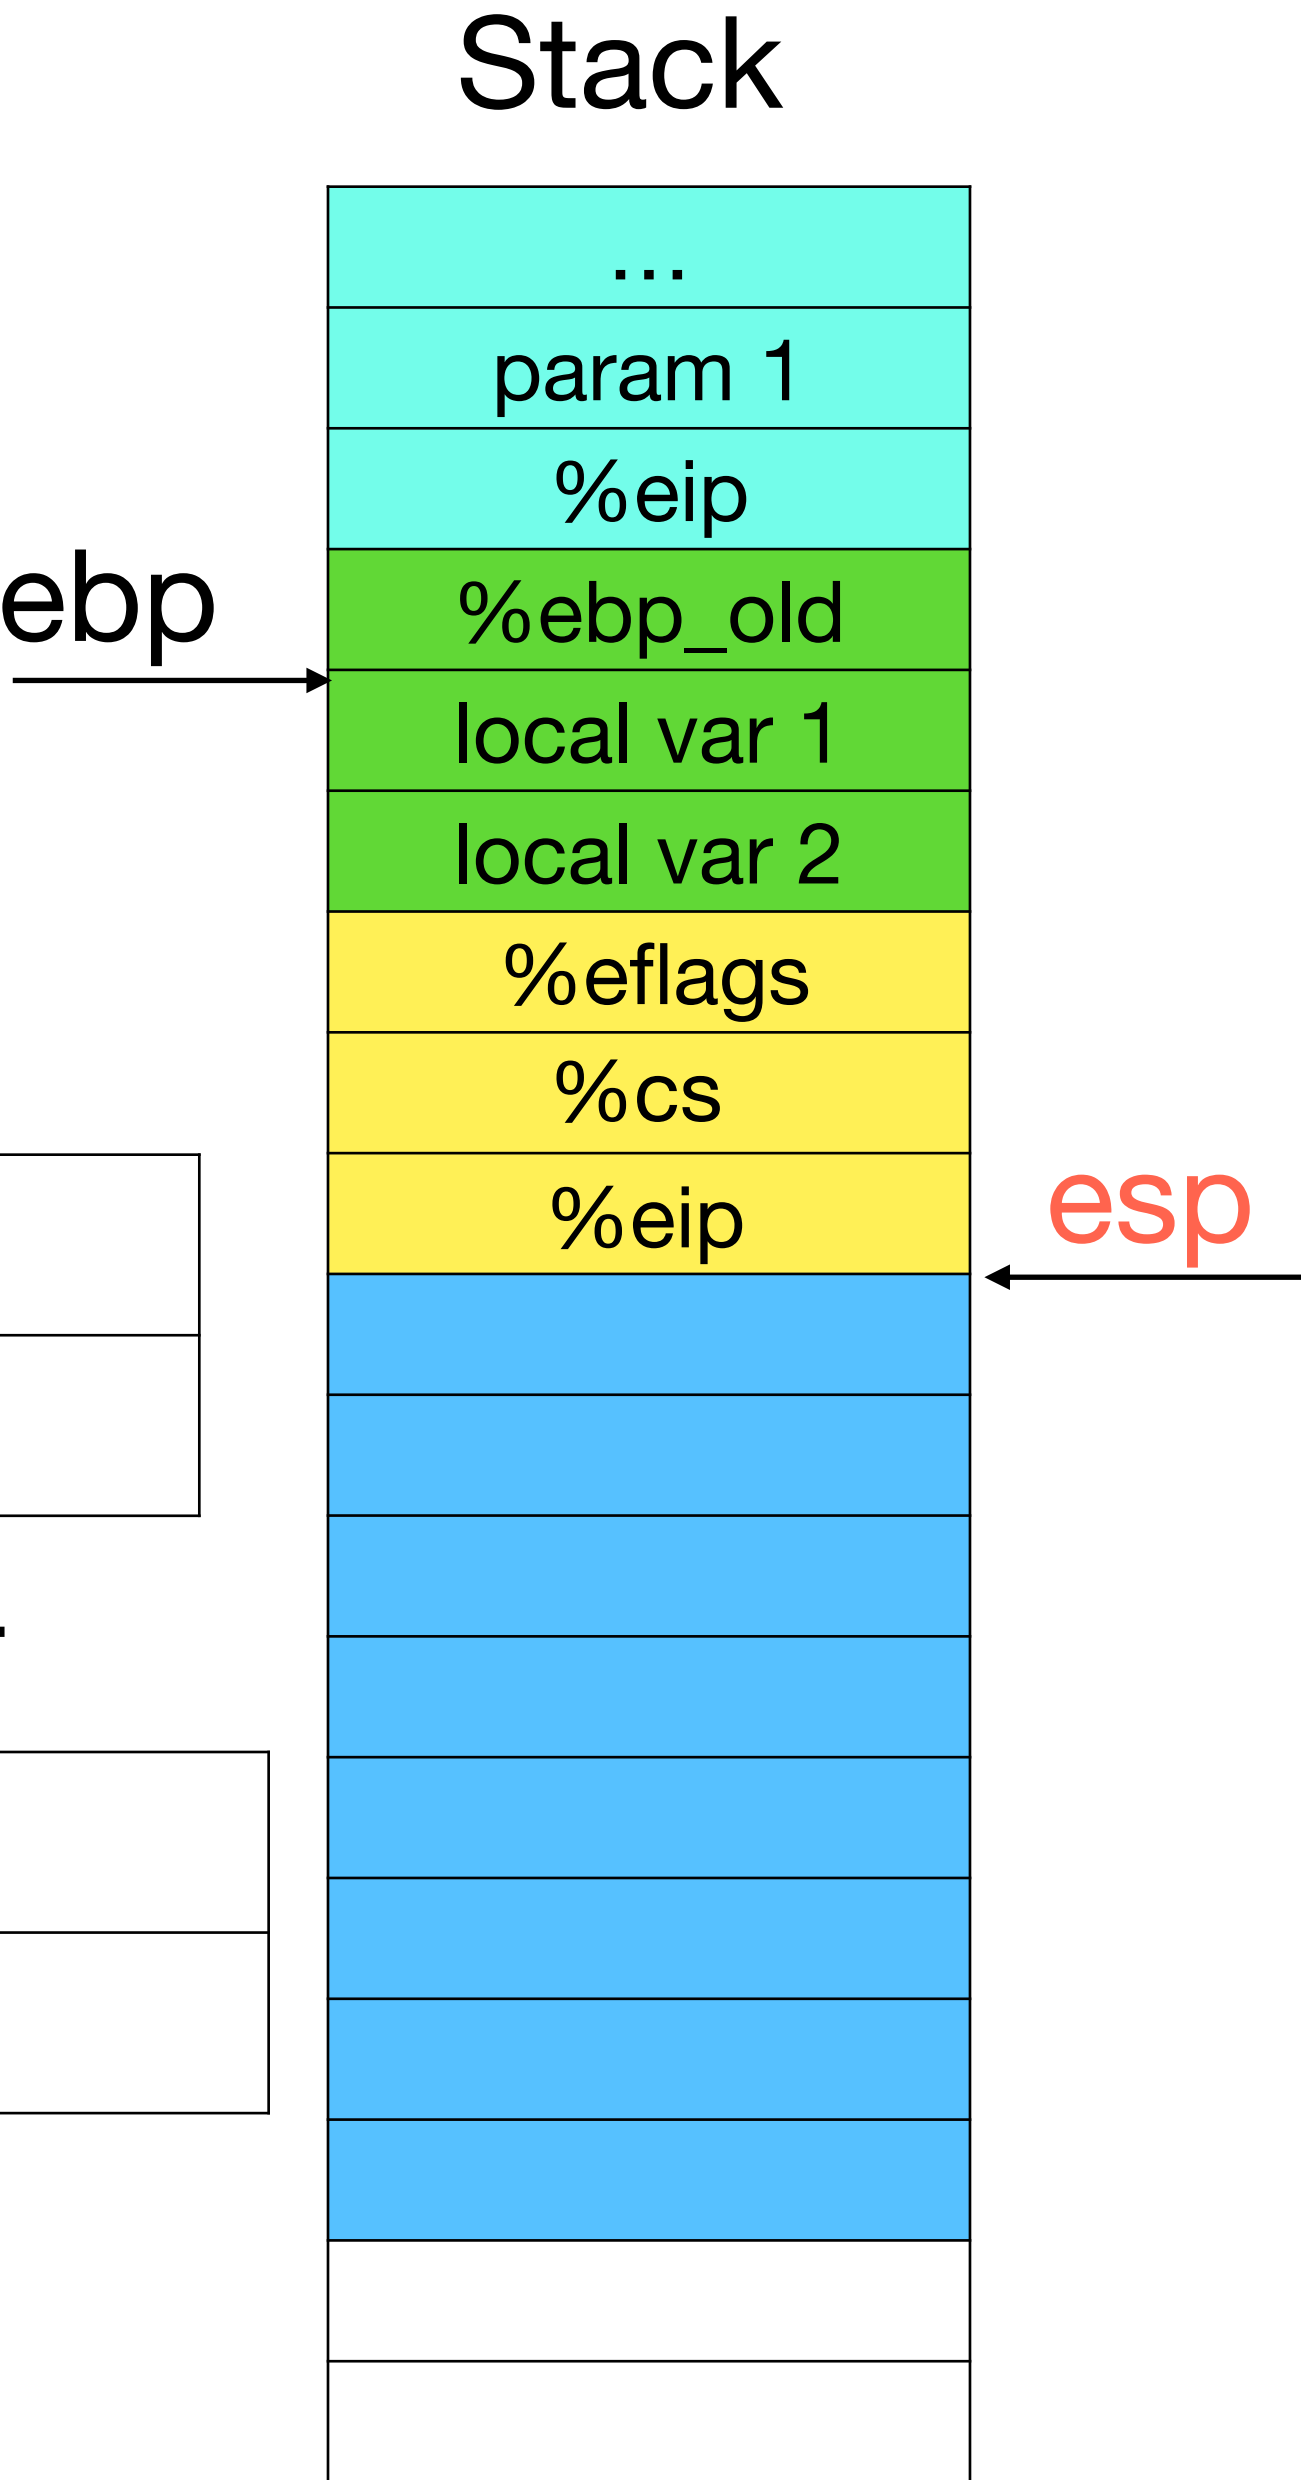
Visualizing interrupt handling

```
eip → for(;;)
;

trap.c
void
trap(struct trapframe *tf)
{
    switch(tf->trapno){
    case T_IRQ0 + IRQ_TIMER:
        ticks++;
        cprintf("Tick! %d\n\0", ticks);
        lapiceoi();
    ..
    return
```

```
vectors.S
.globl vector0
vector0:
    pushl $0
    pushl $0
    jmp alltraps
```

```
trapasm.S
alltraps:
    pushal
    pushl %esp
    call trap
    addl $4, %esp
    popal
    addl $0x8, %esp
    iret
```



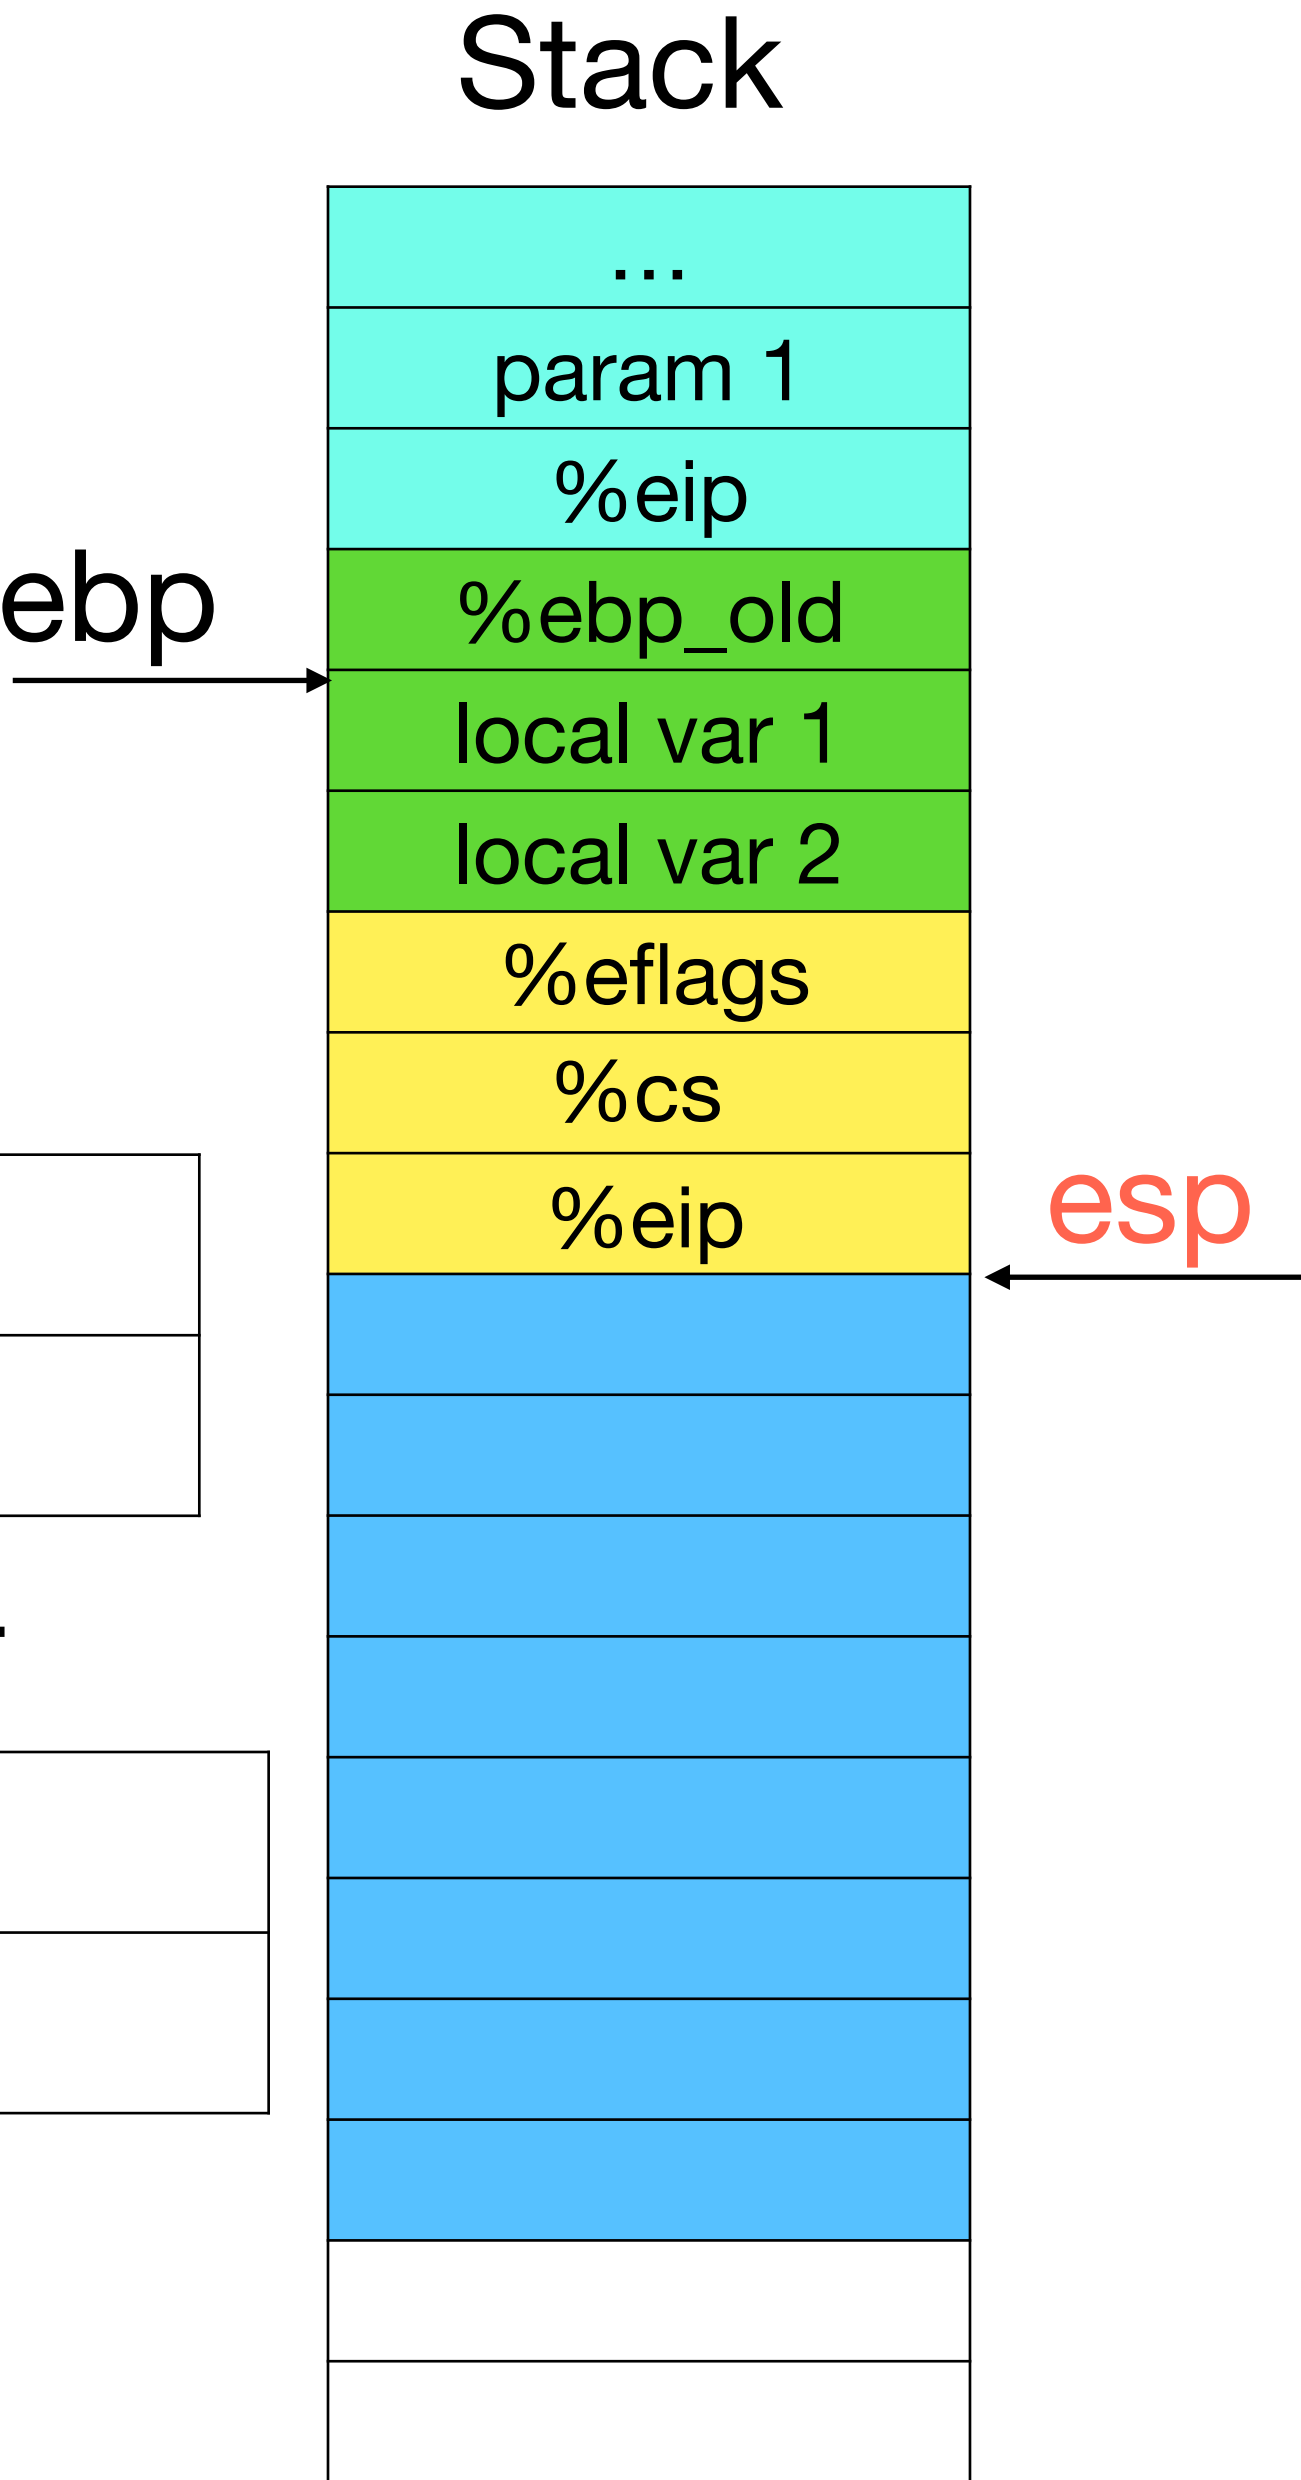
Visualizing interrupt handling

```
eip → for(;;)
;

trap.c
void
trap(struct trapframe *tf)
{
    switch(tf->trapno){
    case T_IRQ0 + IRQ_TIMER:
        ticks++;
        cprintf("Tick! %d\n\0", ticks);
        lapiceoi();
    ..
    return
```

```
vectors.S
.globl vector0
vector0:
    pushl $0
    pushl $0
    jmp alltraps
```

```
trapasm.S
alltraps:
    pushal
    pushl %esp
    call trap
    addl $4, %esp
    popal
    addl $0x8, %esp
    iret
```



Visualizing interrupt handling

eip → for(;;)
;

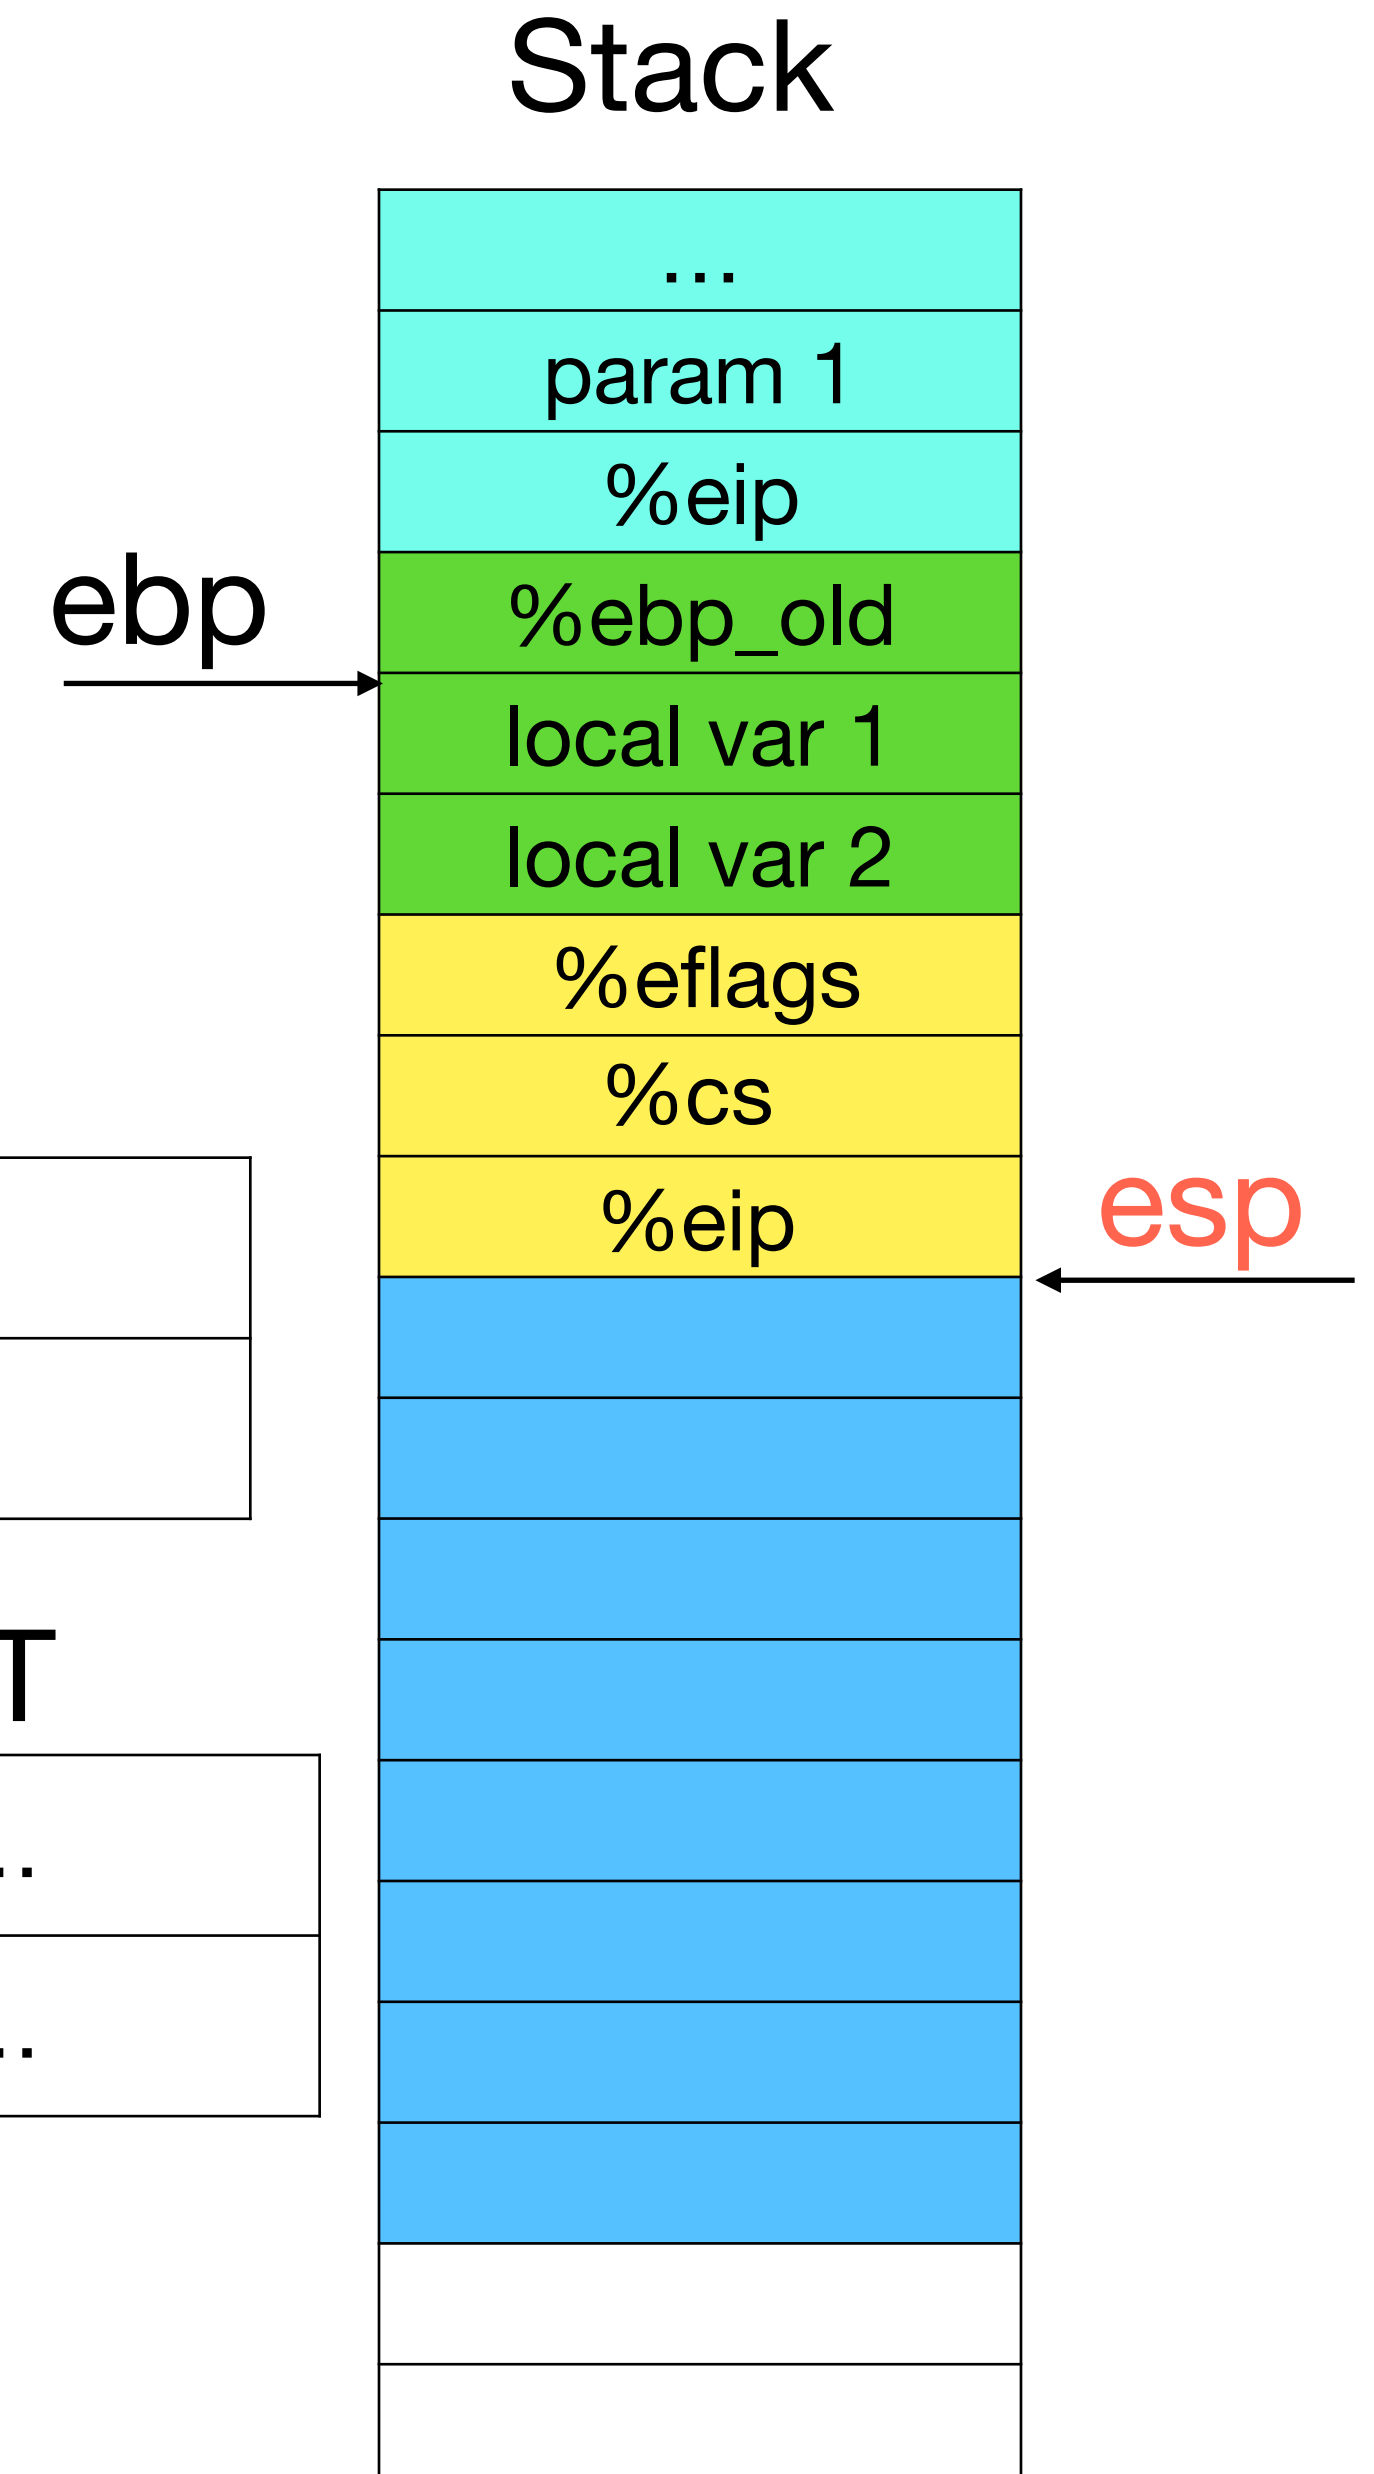
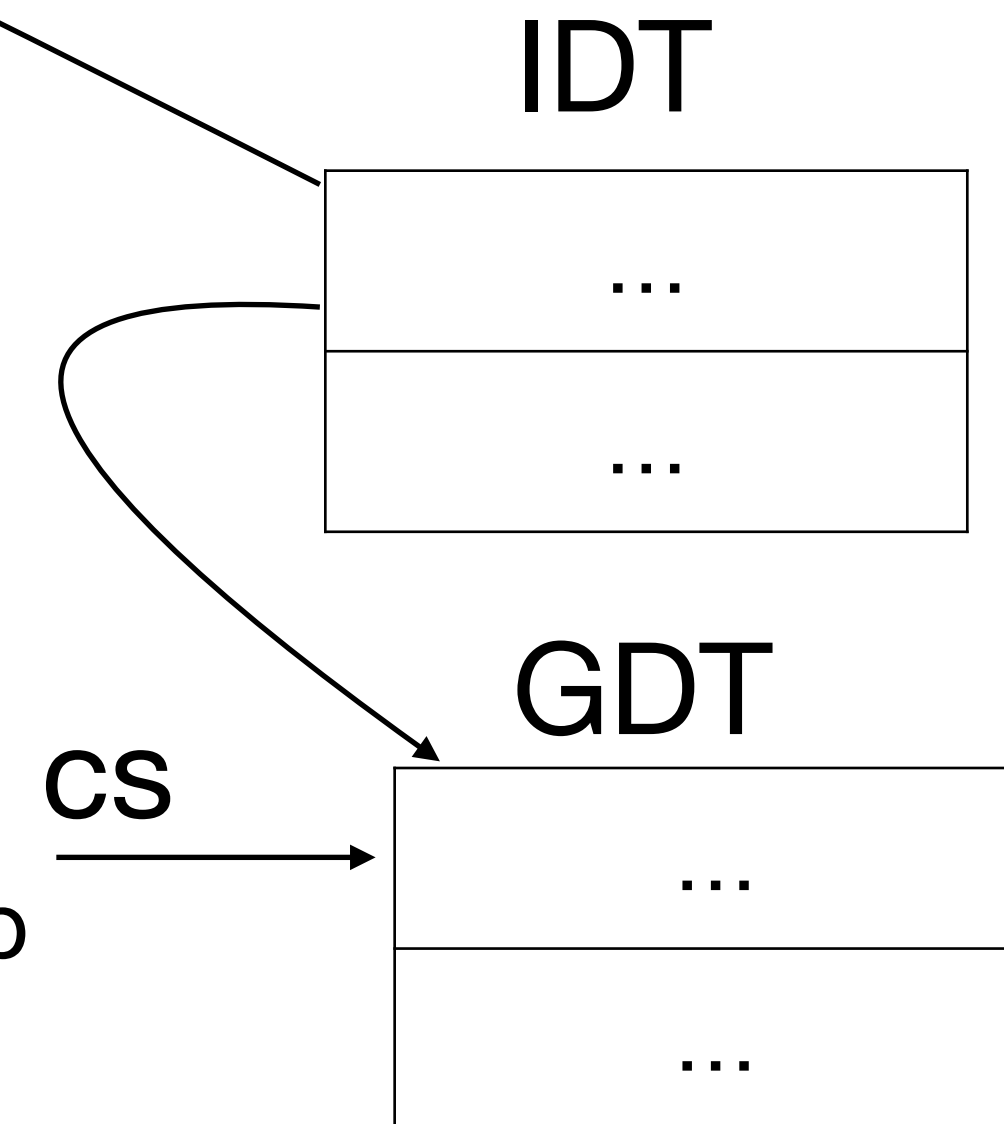
```
trap.c
void
trap(struct trapframe *tf)
{
    switch(tf->trapno){
    case T_IRQ0 + IRQ_TIMER:
        ticks++;
        cprintf("Tick! %d\n", ticks);
        lapiceoi();
        ..
    }
    return
}
```

```

vectors.S
.globl vector0
vector0:
    pushl $0
    pushl $0
    jmp alltraps

```

```
trapasm.S
alltraps:
    pushal
    pushl %esp
    call trap
    addl $4, %esp
    popal
    addl $0x8, %esp
    iret
```



Visualizing interrupt handling

```
for(;;)  
;
```

trap.c

```
void  
trap(struct trapframe *tf)  
{  
    switch(tf->trapno){  
    case T_IRQ0 + IRQ_TIMER:  
        ticks++;  
        cprintf("Tick! %d\n\0", ticks);  
        lapiceoi();  
    ..  
    return
```

vectors.S

```
.globl vector0  
vector0:  
    pushl $0  
    pushl $0  
    jmp alltraps
```

eip

trapasm.S

```
alltraps:  
    pushal  
    pushl %esp  
    call trap  
    addl $4, %esp  
    popal  
    addl $0x8, %esp  
    iret
```

IDT

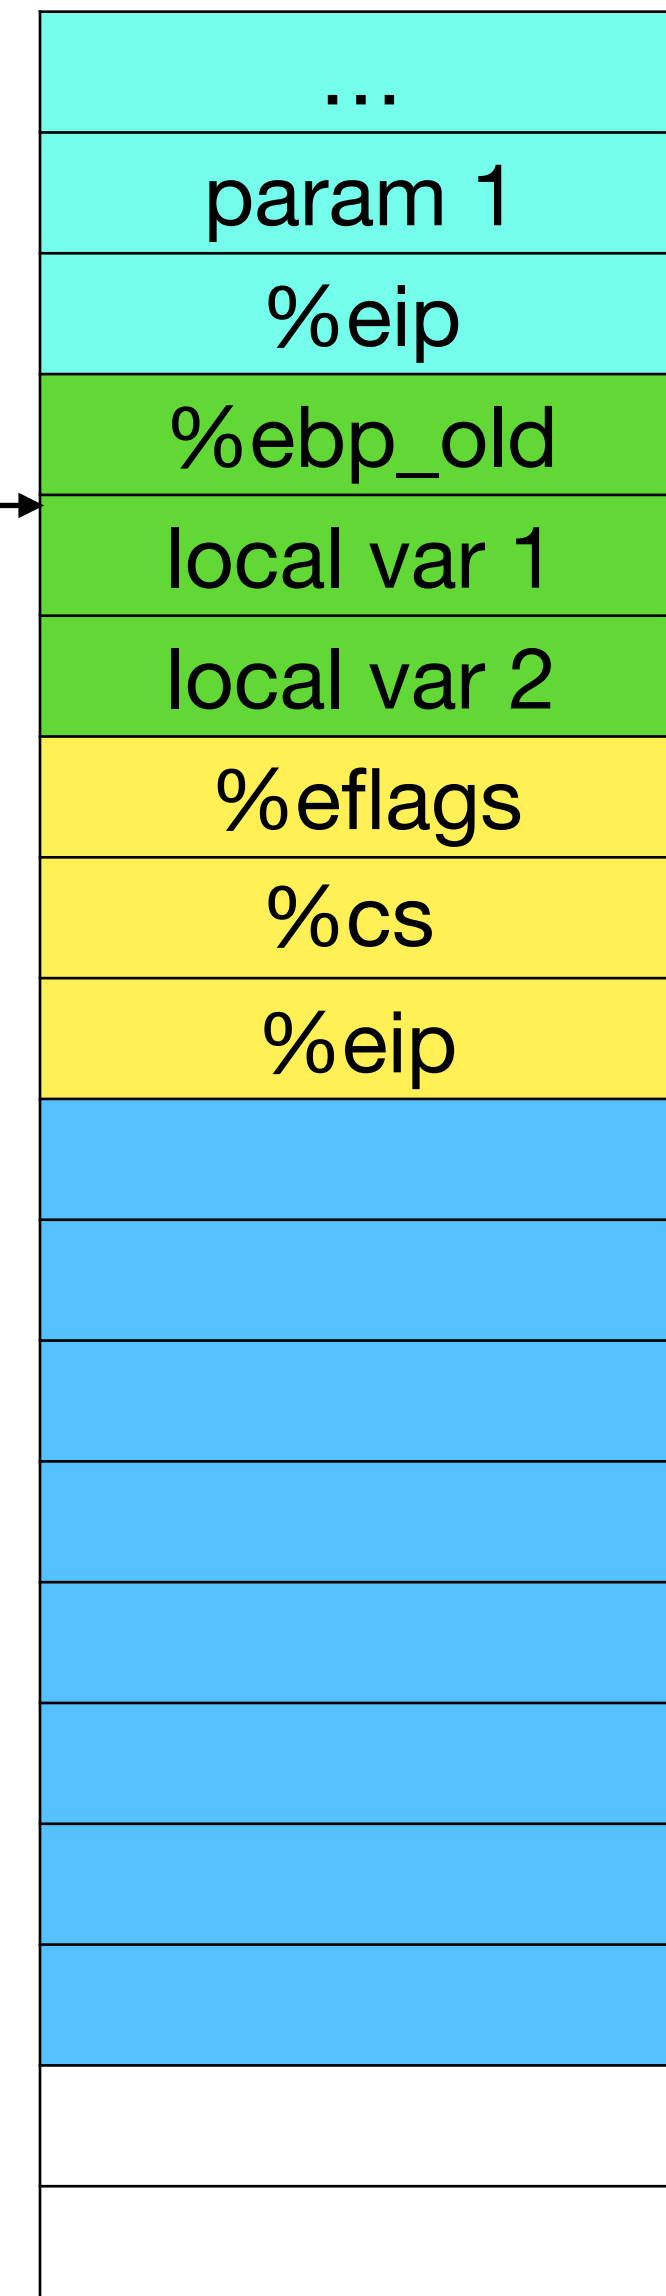


GDT



CS

Stack



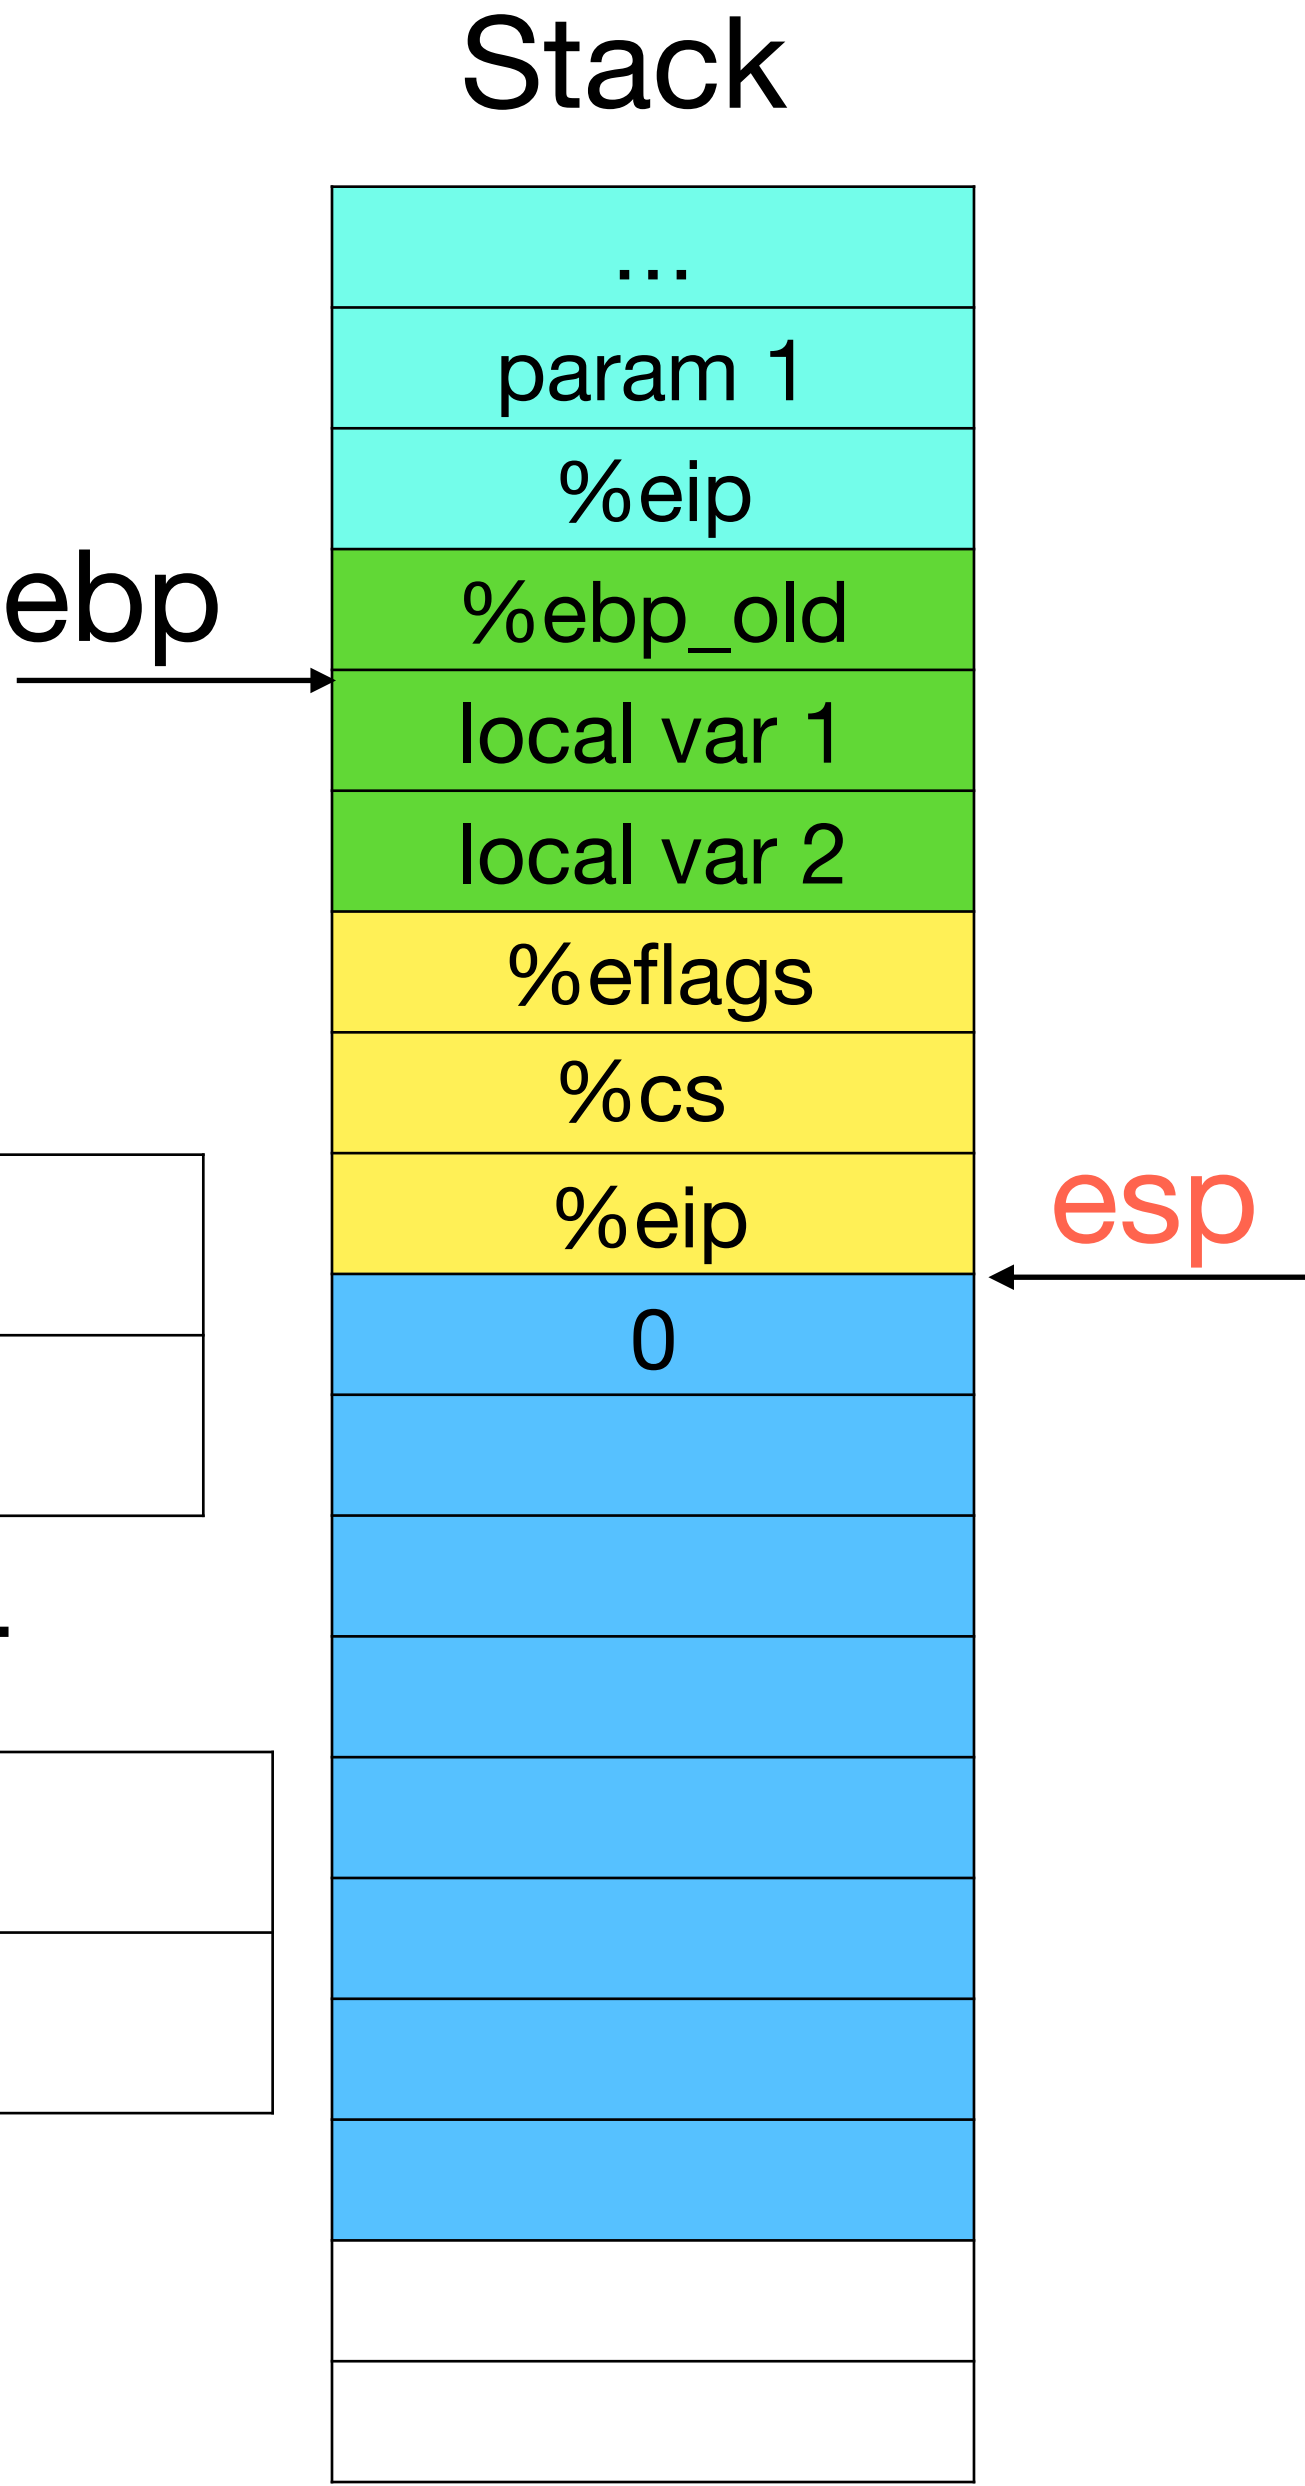
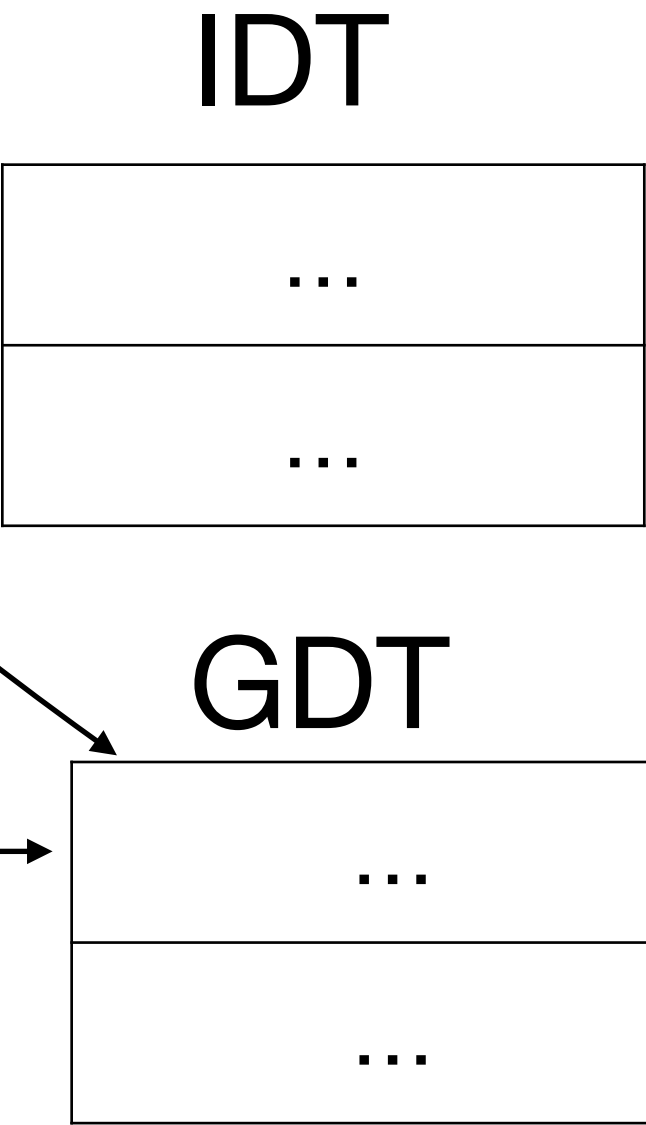
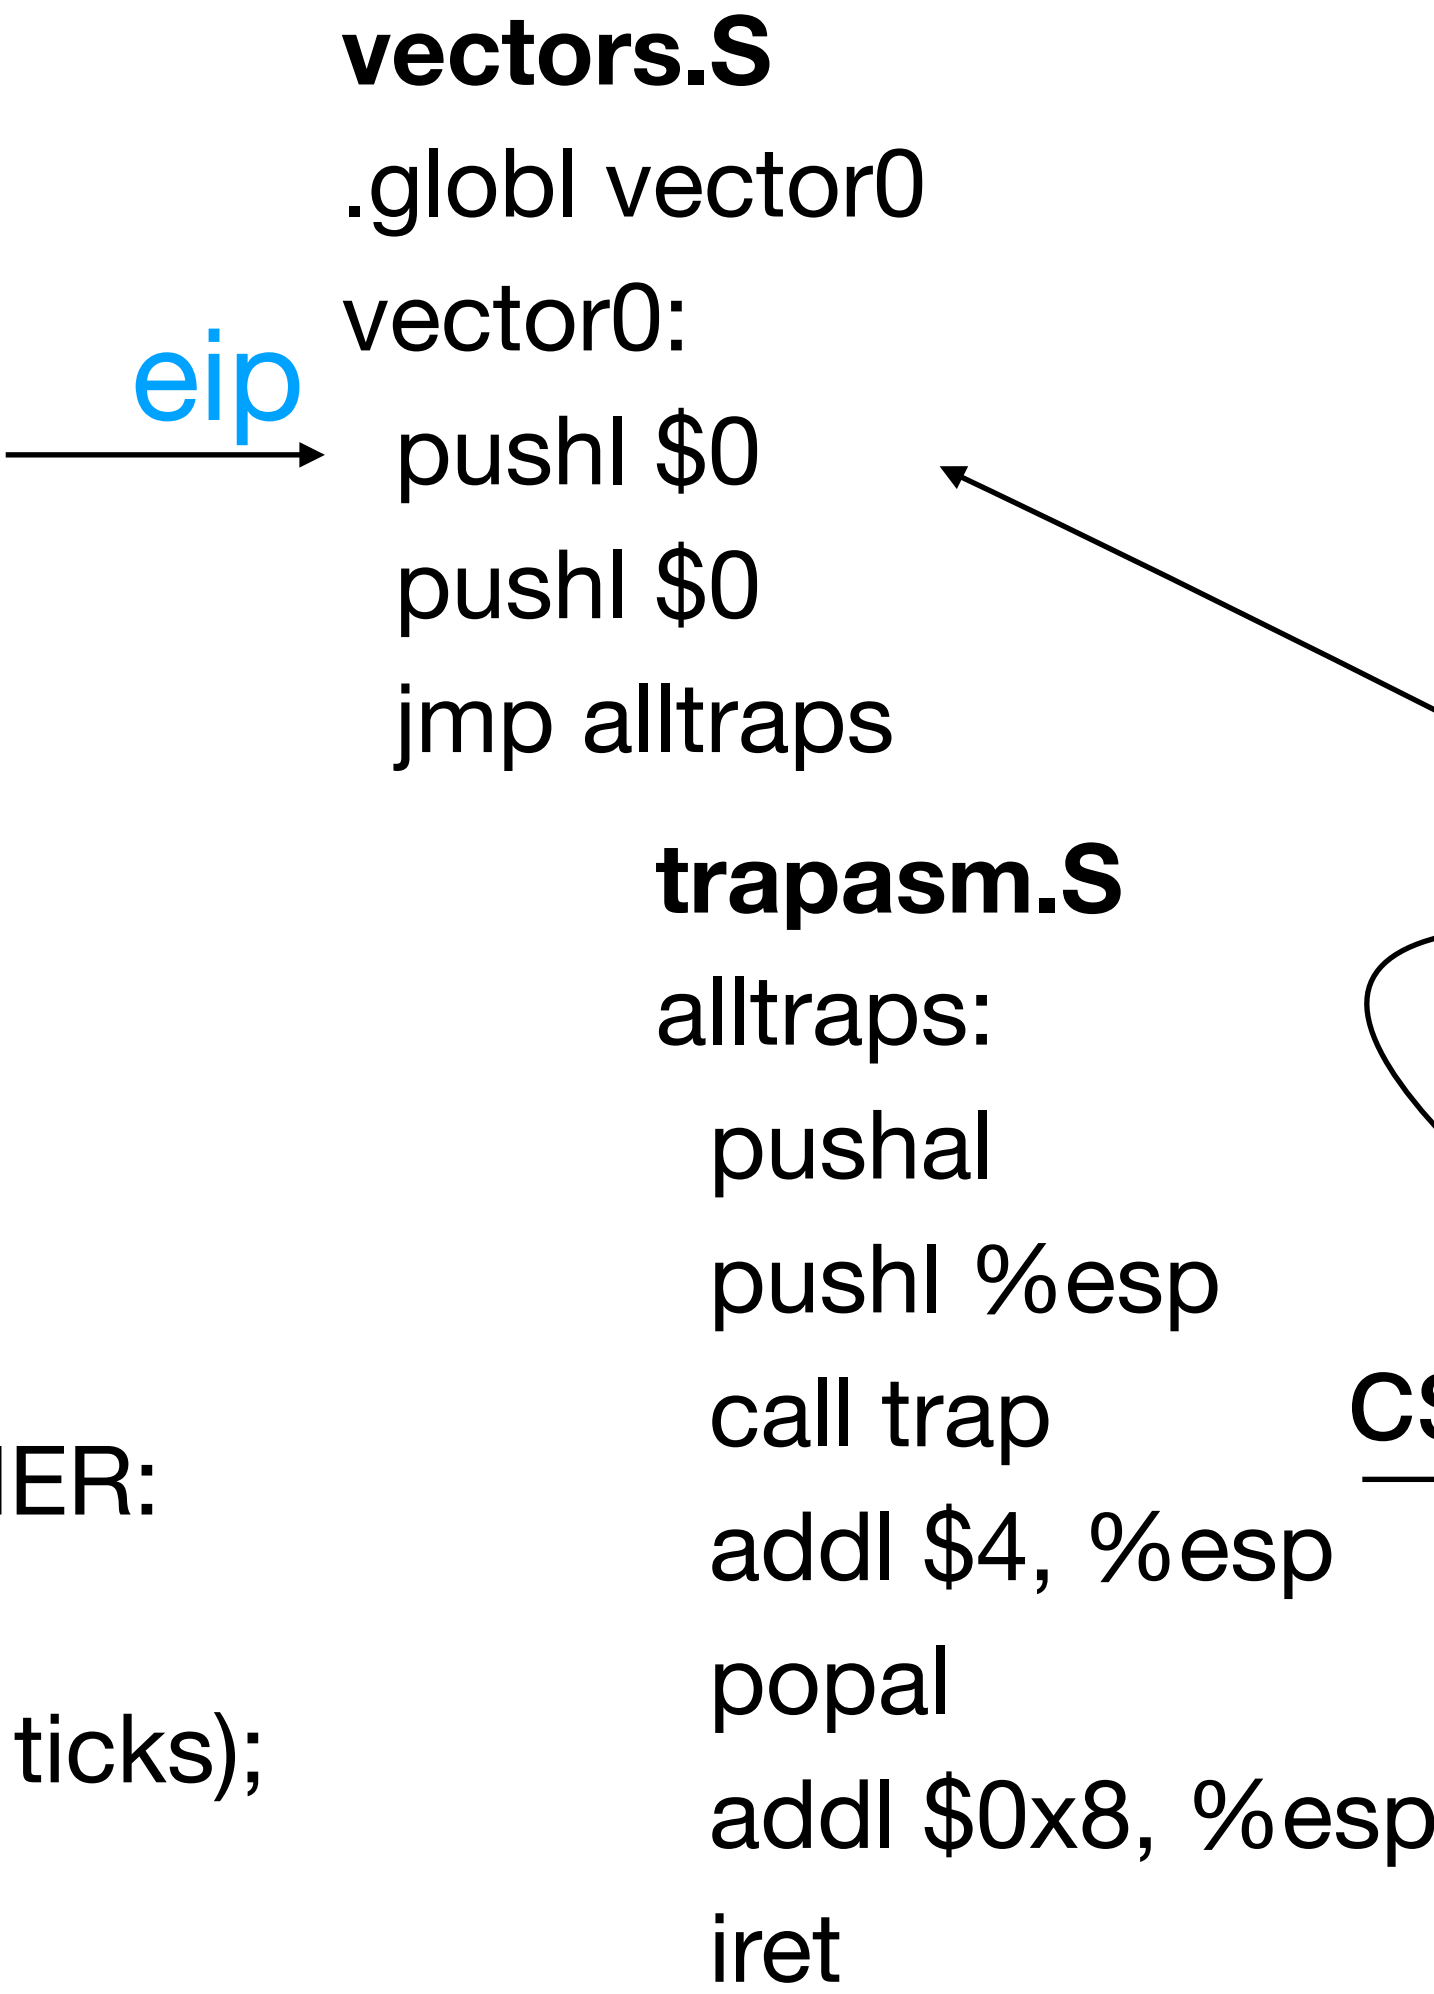
ebp

esp

Visualizing interrupt handling

```
for(;;)  
;
```

```
trap.c  
void  
trap(struct trapframe *tf)  
{  
    switch(tf->trapno){  
    case T_IRQ0 + IRQ_TIMER:  
        ticks++;  
        cprintf("Tick! %d\n\0", ticks);  
        lapiceoi();  
    ..  
    return
```



Visualizing interrupt handling

```
for(;;)  
;
```

trap.c

```
void  
trap(struct trapframe *tf)  
{  
    switch(tf->trapno){  
    case T_IRQ0 + IRQ_TIMER:  
        ticks++;  
        cprintf("Tick! %d\n\0", ticks);  
        lapiceoi();  
    ..  
    return
```

vectors.S

```
.globl vector0  
vector0:  
    pushl $0  
    pushl $0  
    jmp alltraps
```

eip

trapasm.S

```
alltraps:  
    pushal  
    pushl %esp  
    call trap  
    addl $4, %esp  
    popal  
    addl $0x8, %esp  
    iret
```

IDT

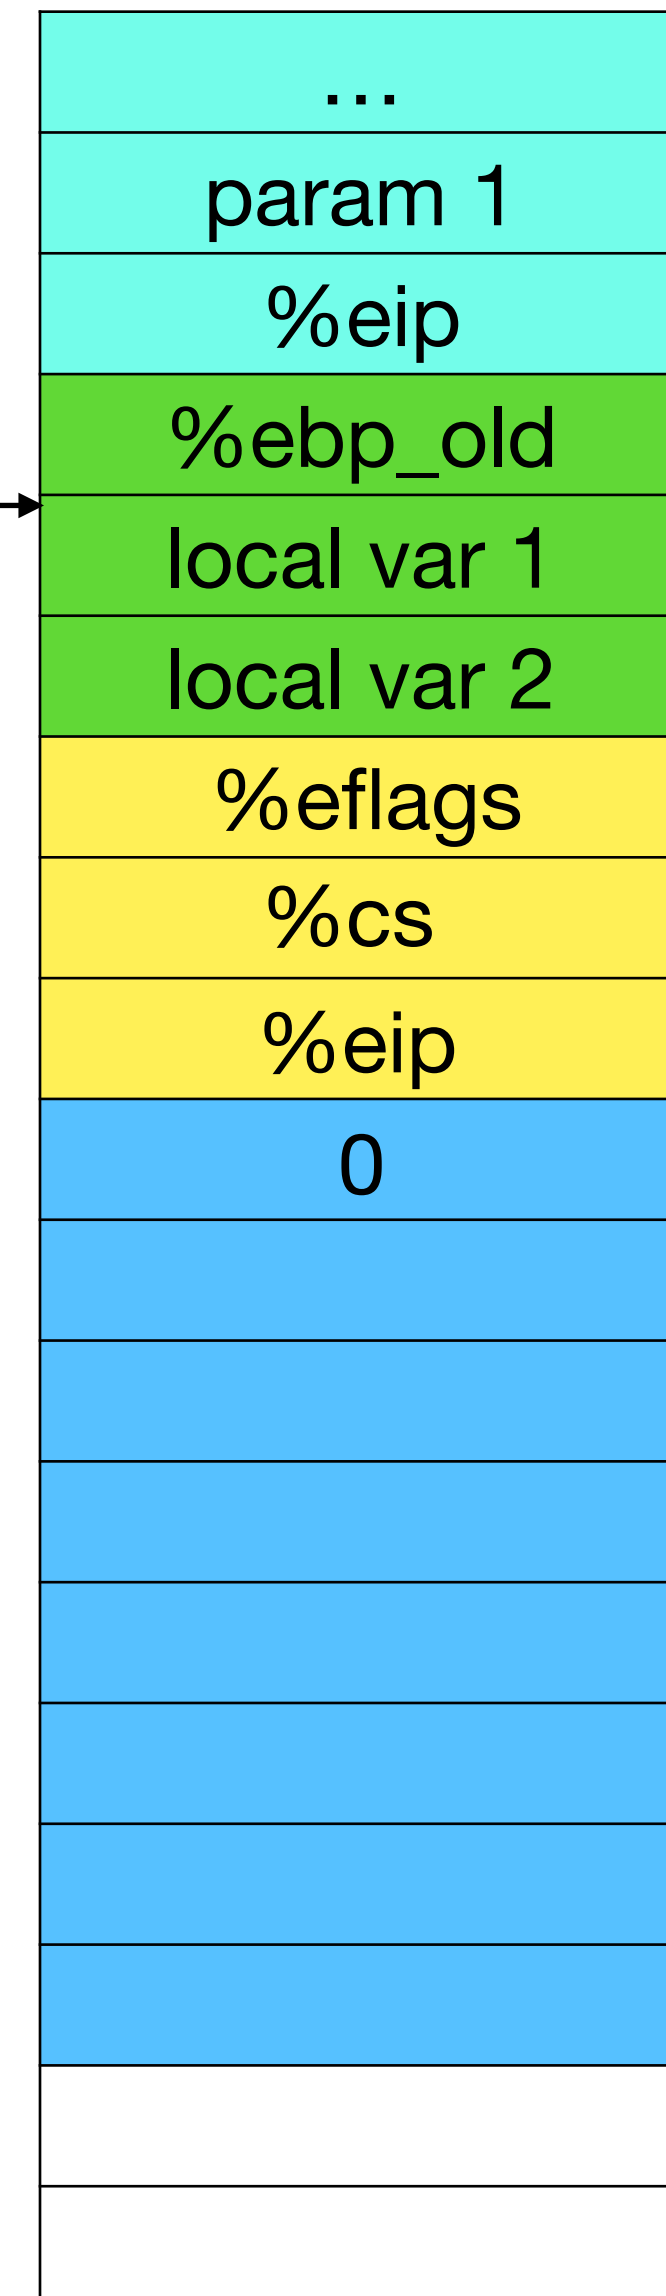


GDT



CS

Stack



ebp

esp

Visualizing interrupt handling

```
for(;;)  
;
```

trap.c

```
void  
trap(struct trapframe *tf)  
{  
    switch(tf->trapno){  
    case T_IRQ0 + IRQ_TIMER:  
        ticks++;  
        cprintf("Tick! %d\n\0", ticks);  
        lapiceoi();  
    ..  
    return
```

vectors.S

```
.globl vector0  
vector0:  
    pushl $0  
    pushl $0  
    jmp alltraps
```

eip

trapasm.S

```
alltraps:  
    pushal  
    pushl %esp  
    call trap  
    addl $4, %esp  
    popal  
    addl $0x8, %esp  
    iret
```

IDT

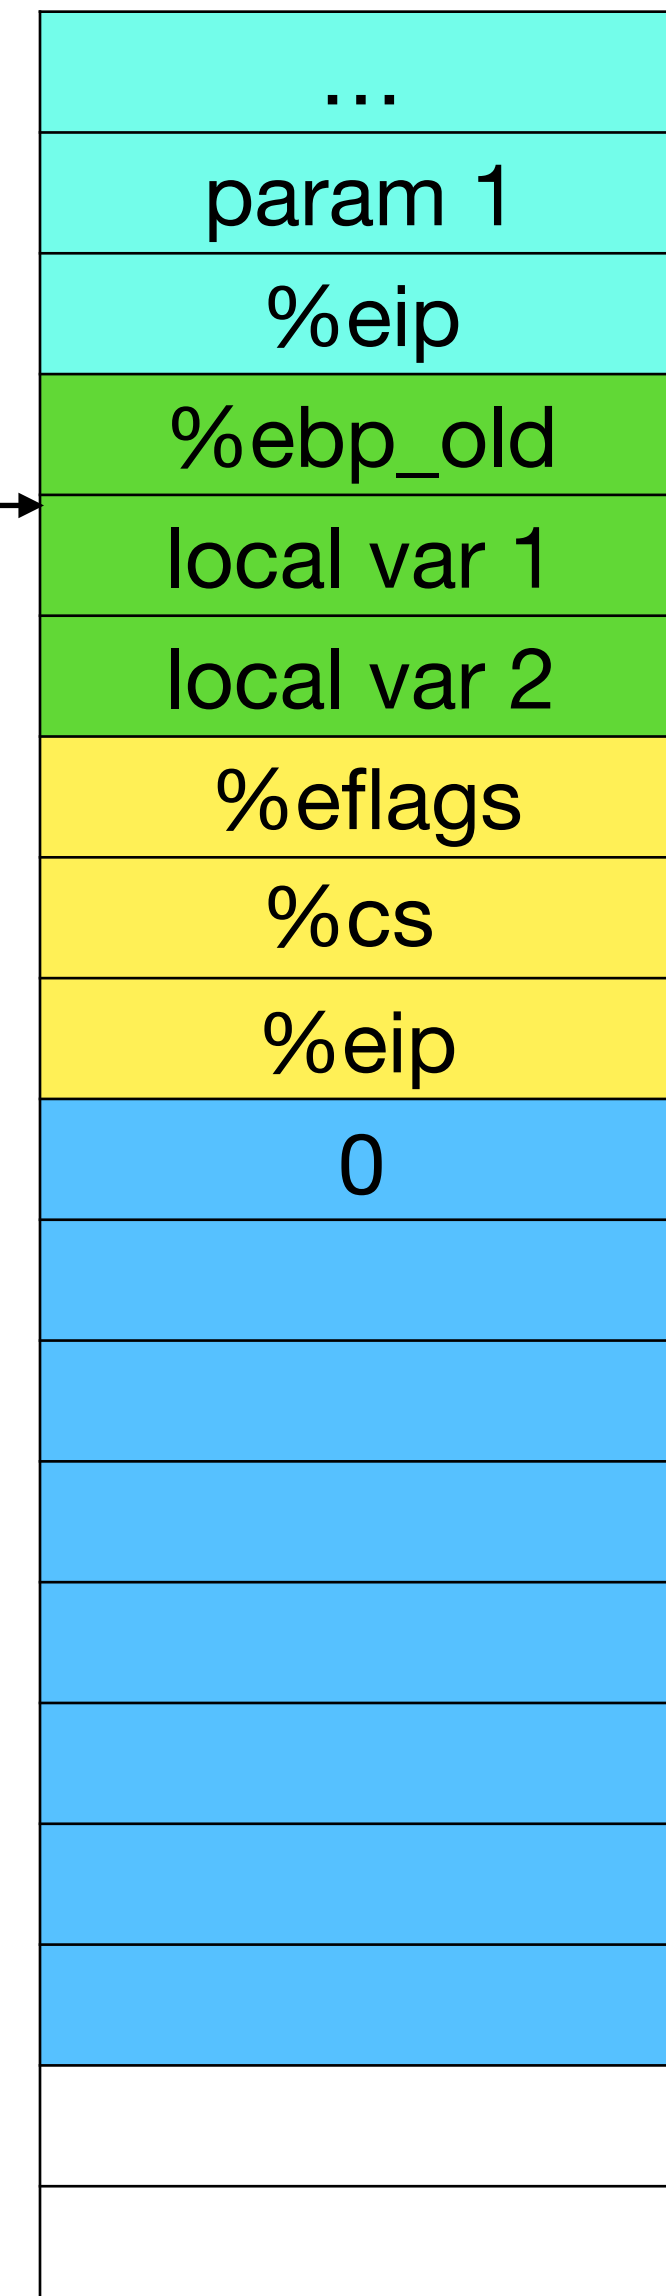


GDT



CS

Stack



ebp

esp

Visualizing interrupt handling

```
for(;;)
;
```

```
trap.c
void
trap(struct trapframe *tf)
{
    switch(tf->trapno){
    case T_IRQ0 + IRQ_TIMER:
        ticks++;
        cprintf("Tick! %d\n\0", ticks);
        lapiceoi();
    ..
    return
```

```
vectors.S
.globl vector0
vector0:
    pushl $0
    pushl $0
    jmp alltraps
```

```
trapasm.S
alltraps:
    pushal
    pushl %esp
    call trap
    addl $4, %esp
    popal
    addl $0x8, %esp
    iret
```

IDT

GDT

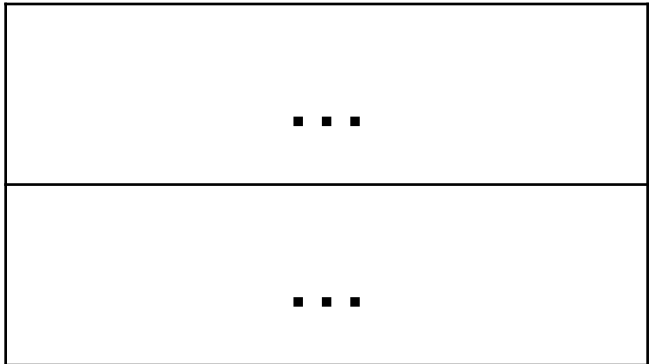
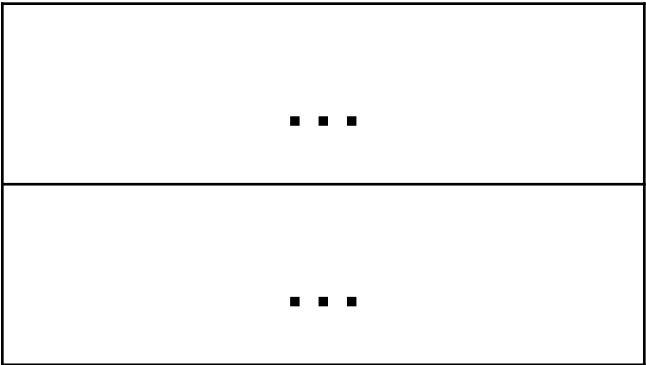
Stack

ebp

esp

eip

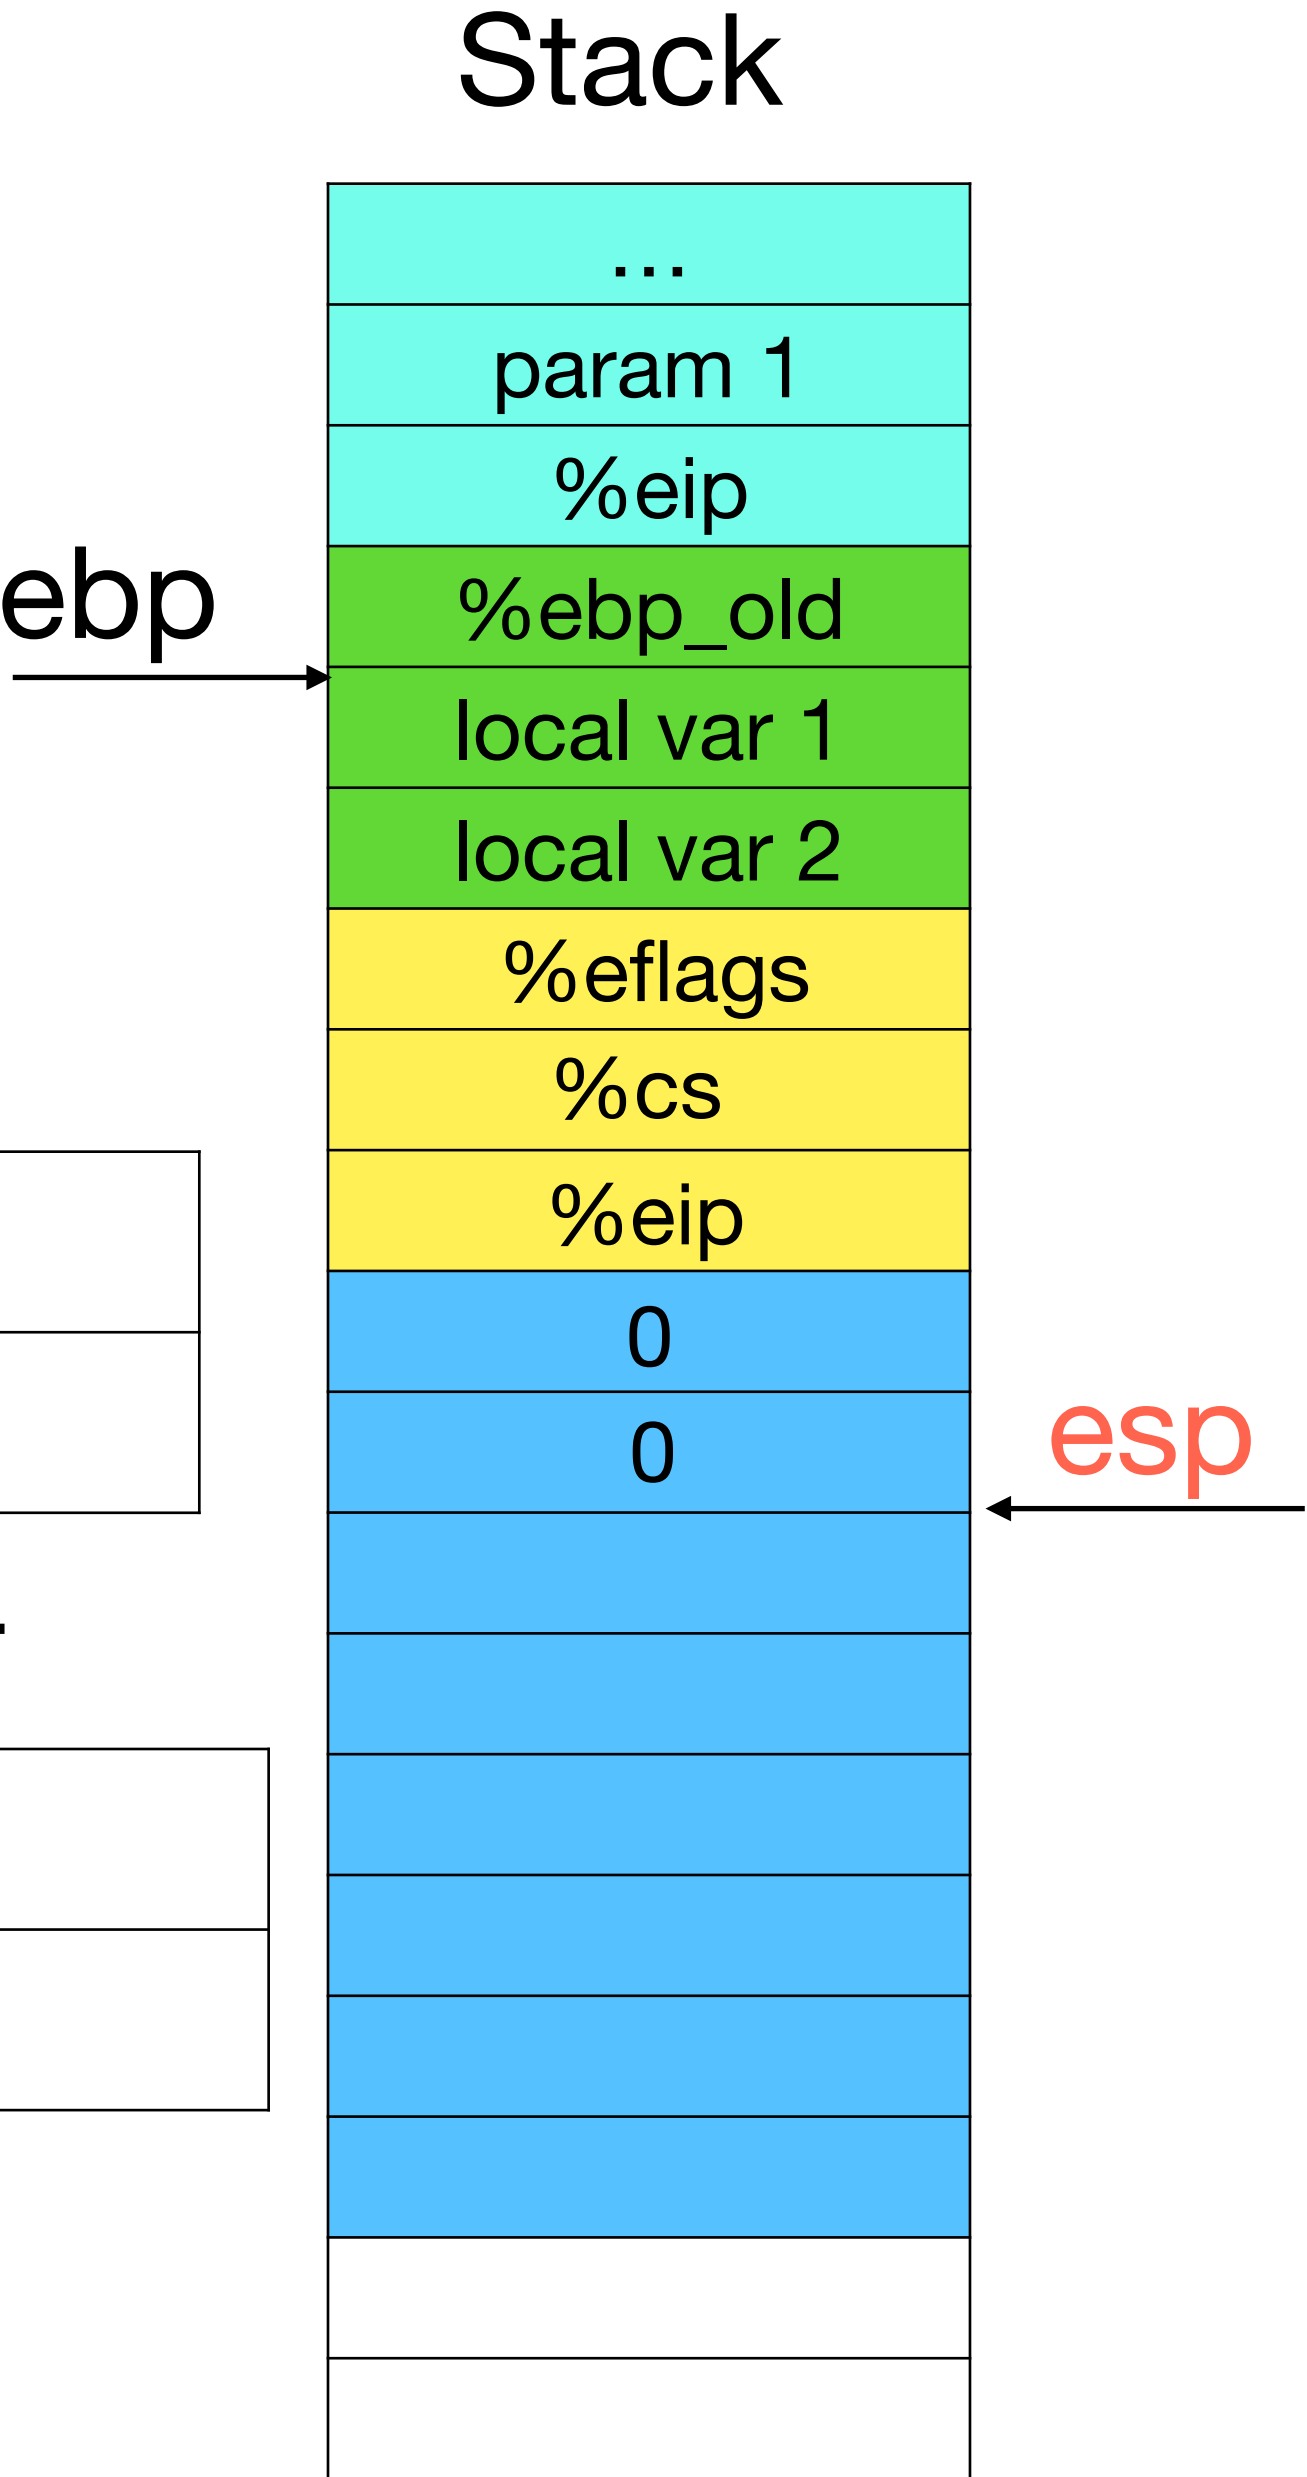
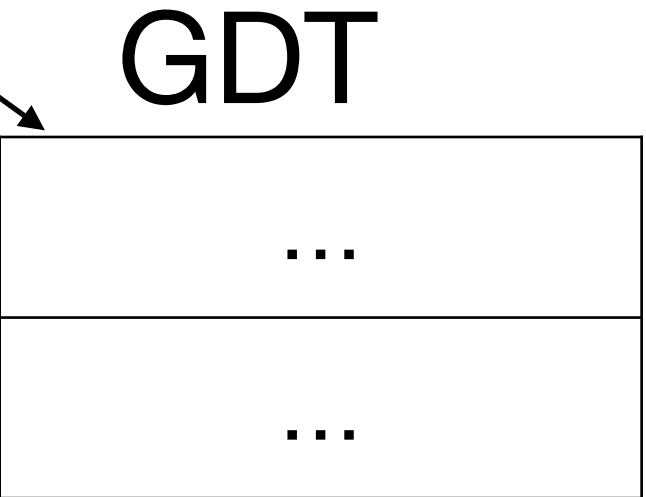
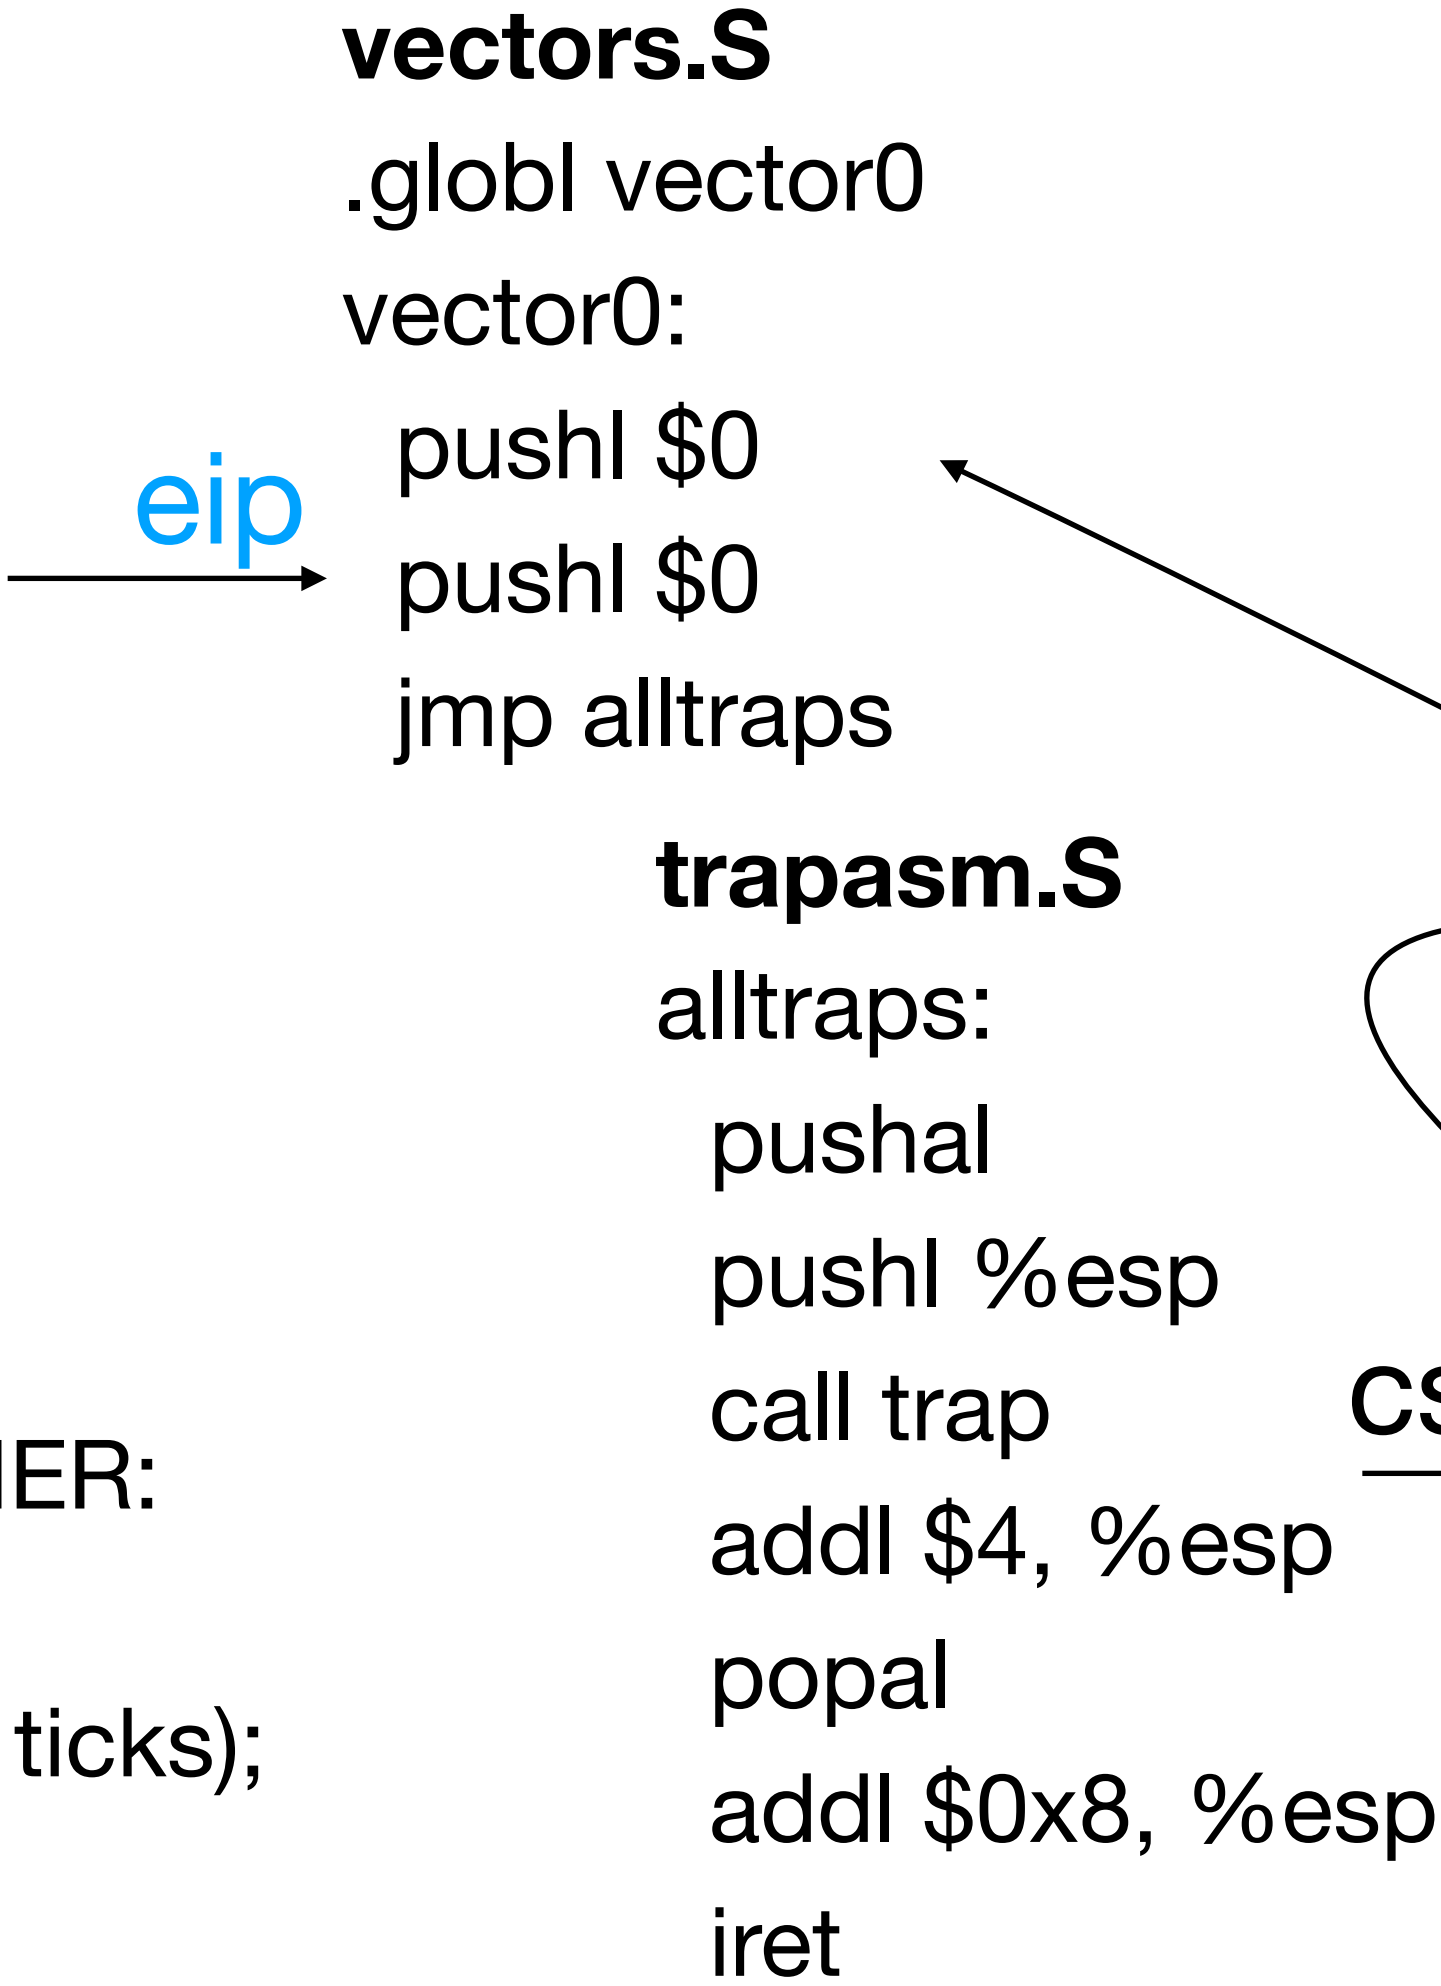
CS



Visualizing interrupt handling

```
for(;;)
;
```

```
trap.c
void
trap(struct trapframe *tf)
{
    switch(tf->trapno){
    case T_IRQ0 + IRQ_TIMER:
        ticks++;
        cprintf("Tick! %d\n\0", ticks);
        lapiceoi();
    ..
    return
```



Visualizing interrupt handling

```
for(;;)
;
```

```
trap.c
void
trap(struct trapframe *tf)
{
    switch(tf->trapno){
    case T_IRQ0 + IRQ_TIMER:
        ticks++;
        cprintf("Tick! %d\n\0", ticks);
        lapiceoi();
    ..
    return
```

```
vectors.S
.globl vector0
vector0:
    pushl $0
    pushl $0
    jmp alltraps
```

```
trapasm.S
alltraps:
    pushal
    pushl %esp
    call trap
    addl $4, %esp
    popal
    addl $0x8, %esp
    iret
```

eip

ebp

CS

esp

IDT

GDT

Stack



Visualizing interrupt handling

```
for(;;)  
;
```

trap.c

```
void  
trap(struct trapframe *tf)  
{  
    switch(tf->trapno){  
    case T_IRQ0 + IRQ_TIMER:  
        ticks++;  
        cprintf("Tick! %d\n\0", ticks);  
        lapiceoi();  
    ..  
    return
```

vectors.S

```
.globl vector0  
vector0:  
    pushl $0  
    pushl $0  
    jmp alltraps
```

trapasm.S

```
alltraps:  
    eip → pushal  
        pushl %esp  
        call trap  
        addl $4, %esp  
        popal  
        addl $0x8, %esp  
        iret
```

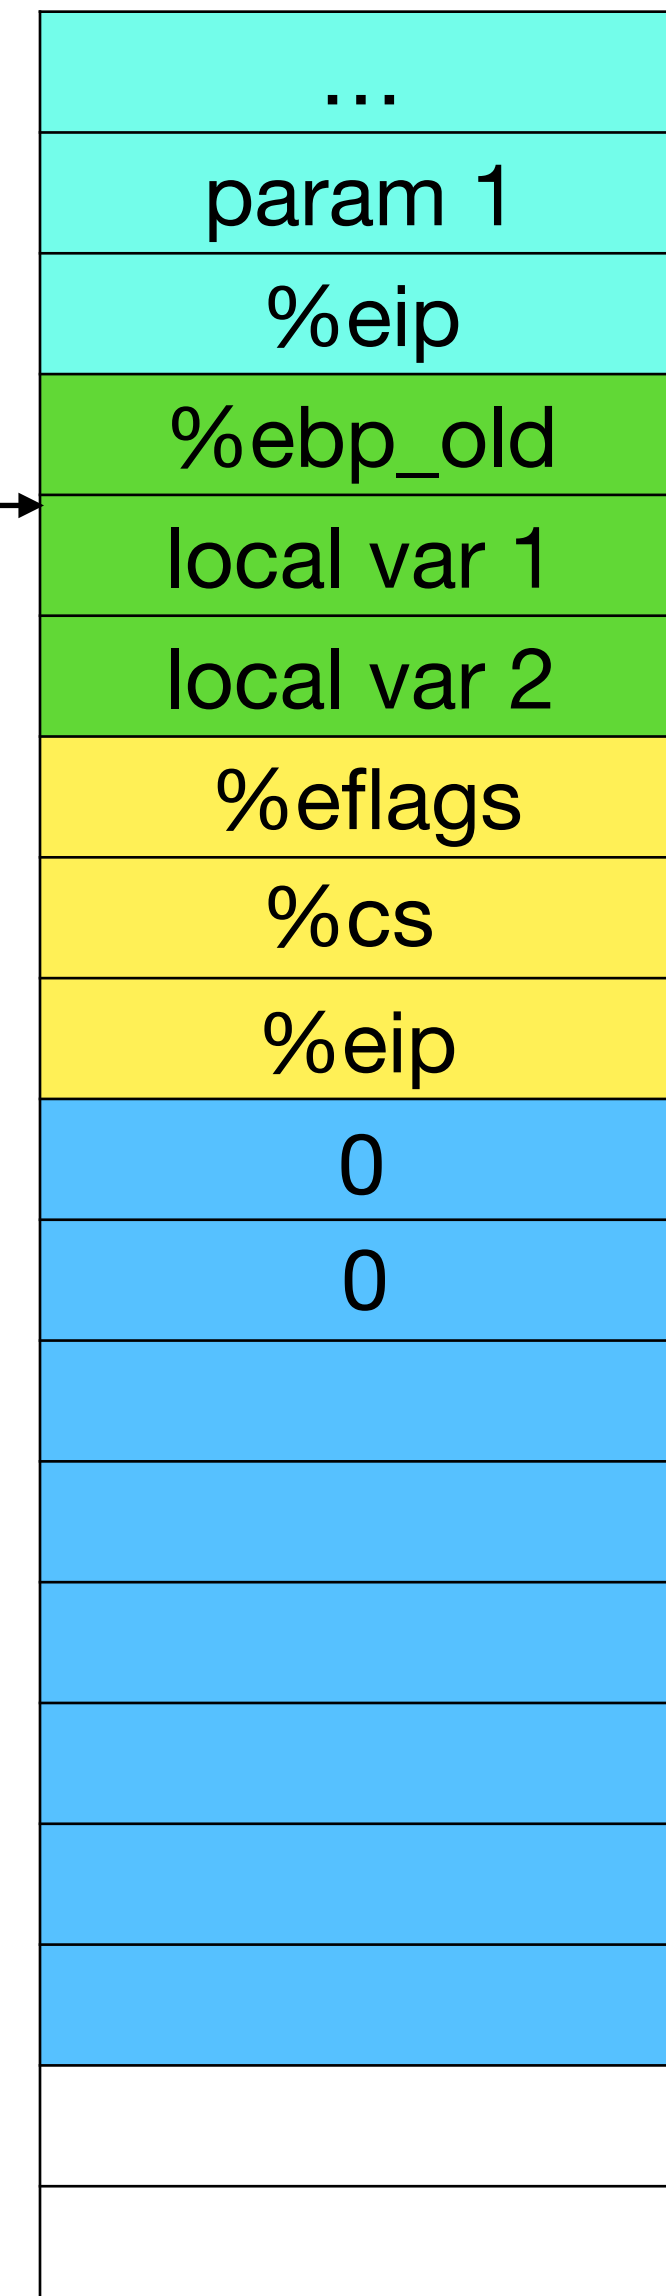
IDT



GDT



Stack



ebp →

esp ←

Visualizing interrupt handling

```
for(;;)  
;
```

trap.c

```
void  
trap(struct trapframe *tf)  
{  
    switch(tf->trapno){  
    case T_IRQ0 + IRQ_TIMER:  
        ticks++;  
        cprintf("Tick! %d\n\0", ticks);  
        lapiceoi();  
    ..  
    return
```

vectors.S

```
.globl vector0  
vector0:  
    pushl $0  
    pushl $0  
    jmp alltraps
```

trapasm.S

```
alltraps:  
    eip → pushal  
        pushl %esp  
        call trap  
        addl $4, %esp  
        popal  
        addl $0x8, %esp  
        iret
```

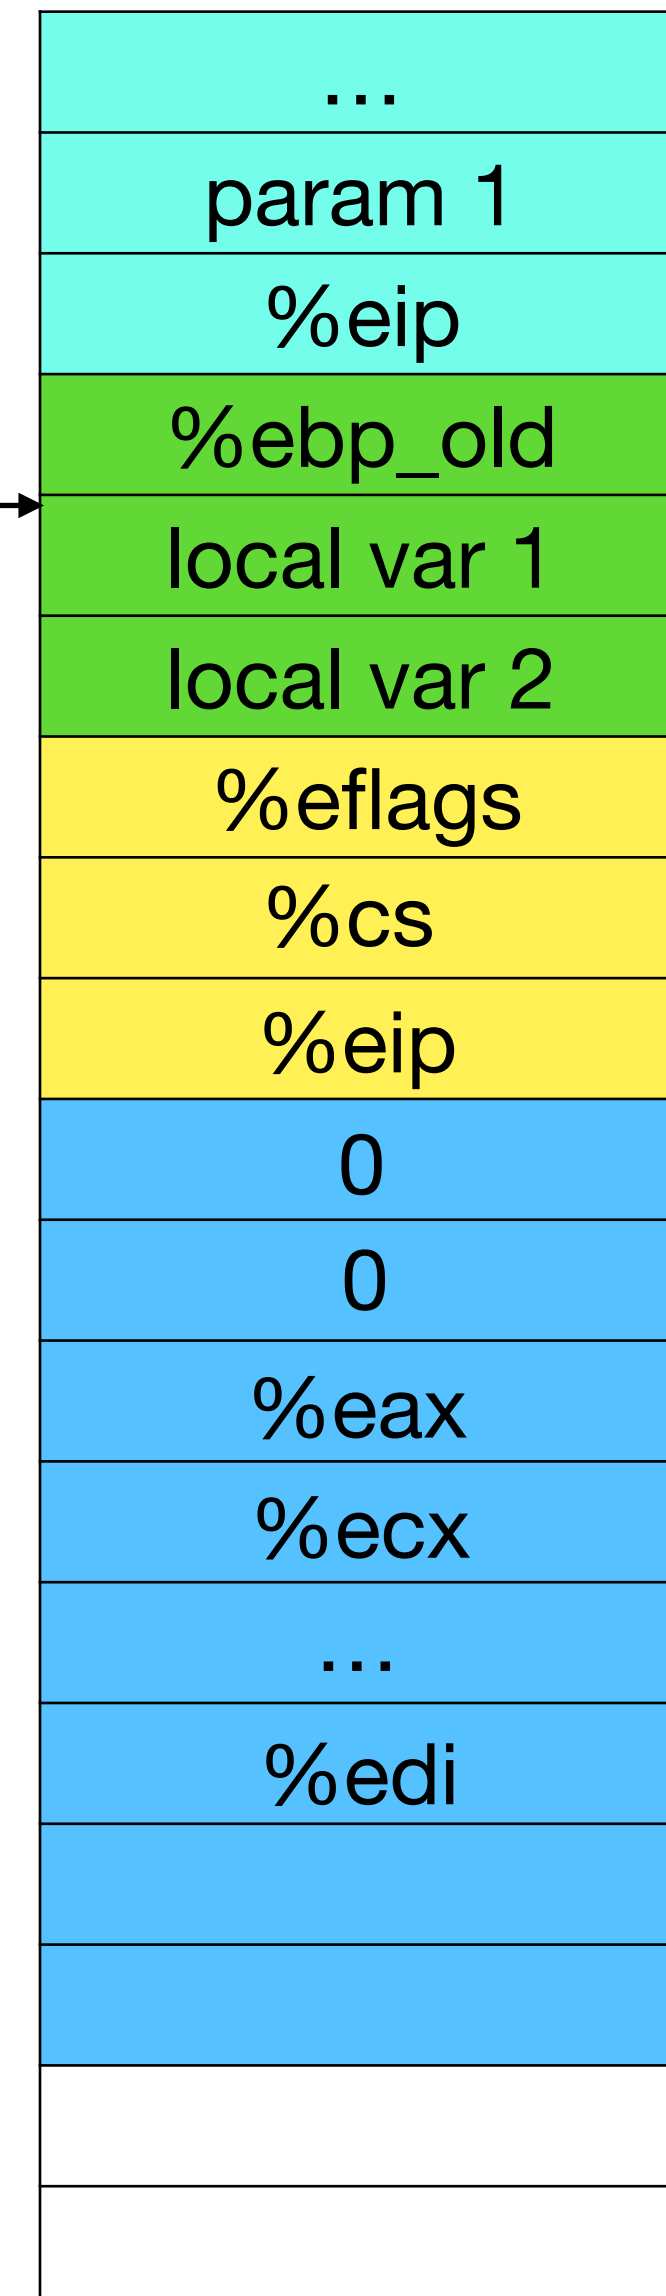
IDT



GDT



Stack



ebp →

esp ←

Visualizing interrupt handling

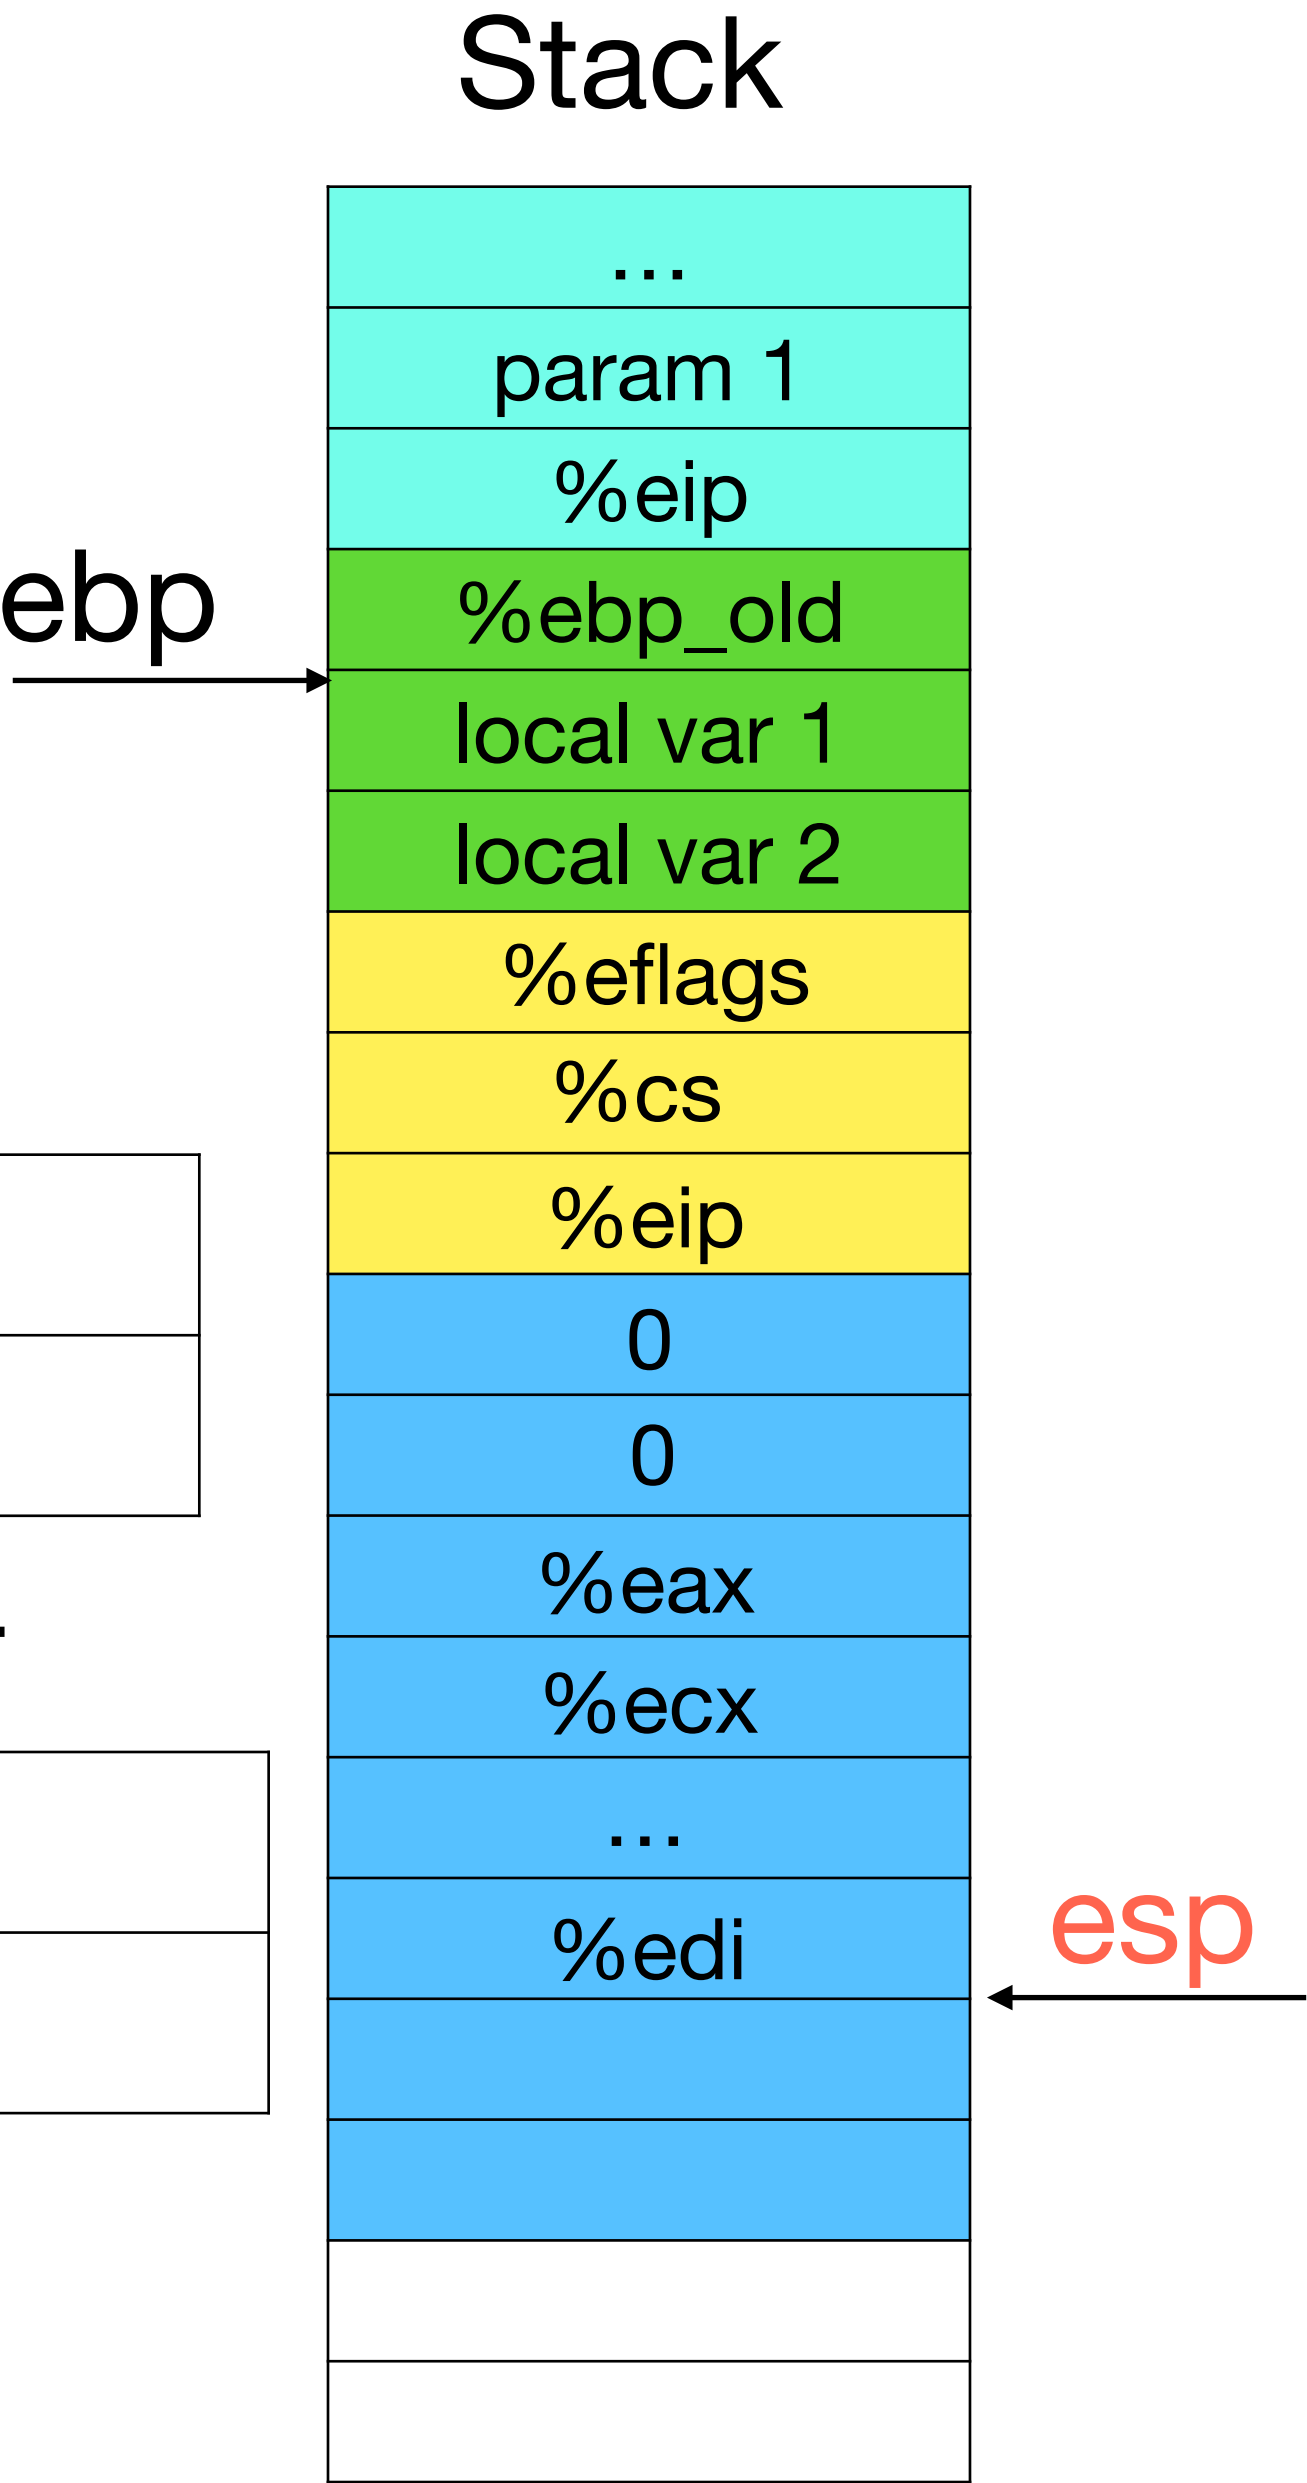
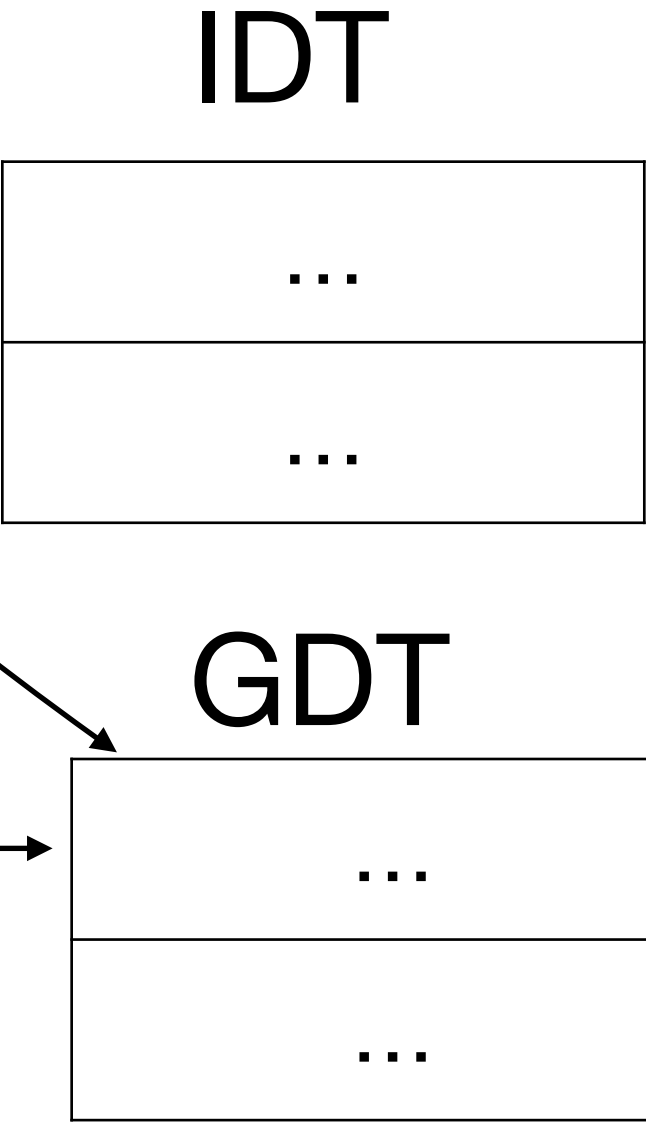
```
for(;;)
;

trap.c
void
trap(struct trapframe *tf)
{
    switch(tf->trapno){
    case T_IRQ0 + IRQ_TIMER:
        ticks++;
        cprintf("Tick! %d\n\0", ticks);
        lapiceoi();
    ..
    return
```

```

vectors.S
.globl vector0
vector0:
    pushl $0
    pushl $0
    jmp alltraps

trapasm.S
alltraps:
    eip → pushal
    pushl %esp
    call trap
    addl $4, %esp
    popal
    addl $0x8, %esp
    iret
```



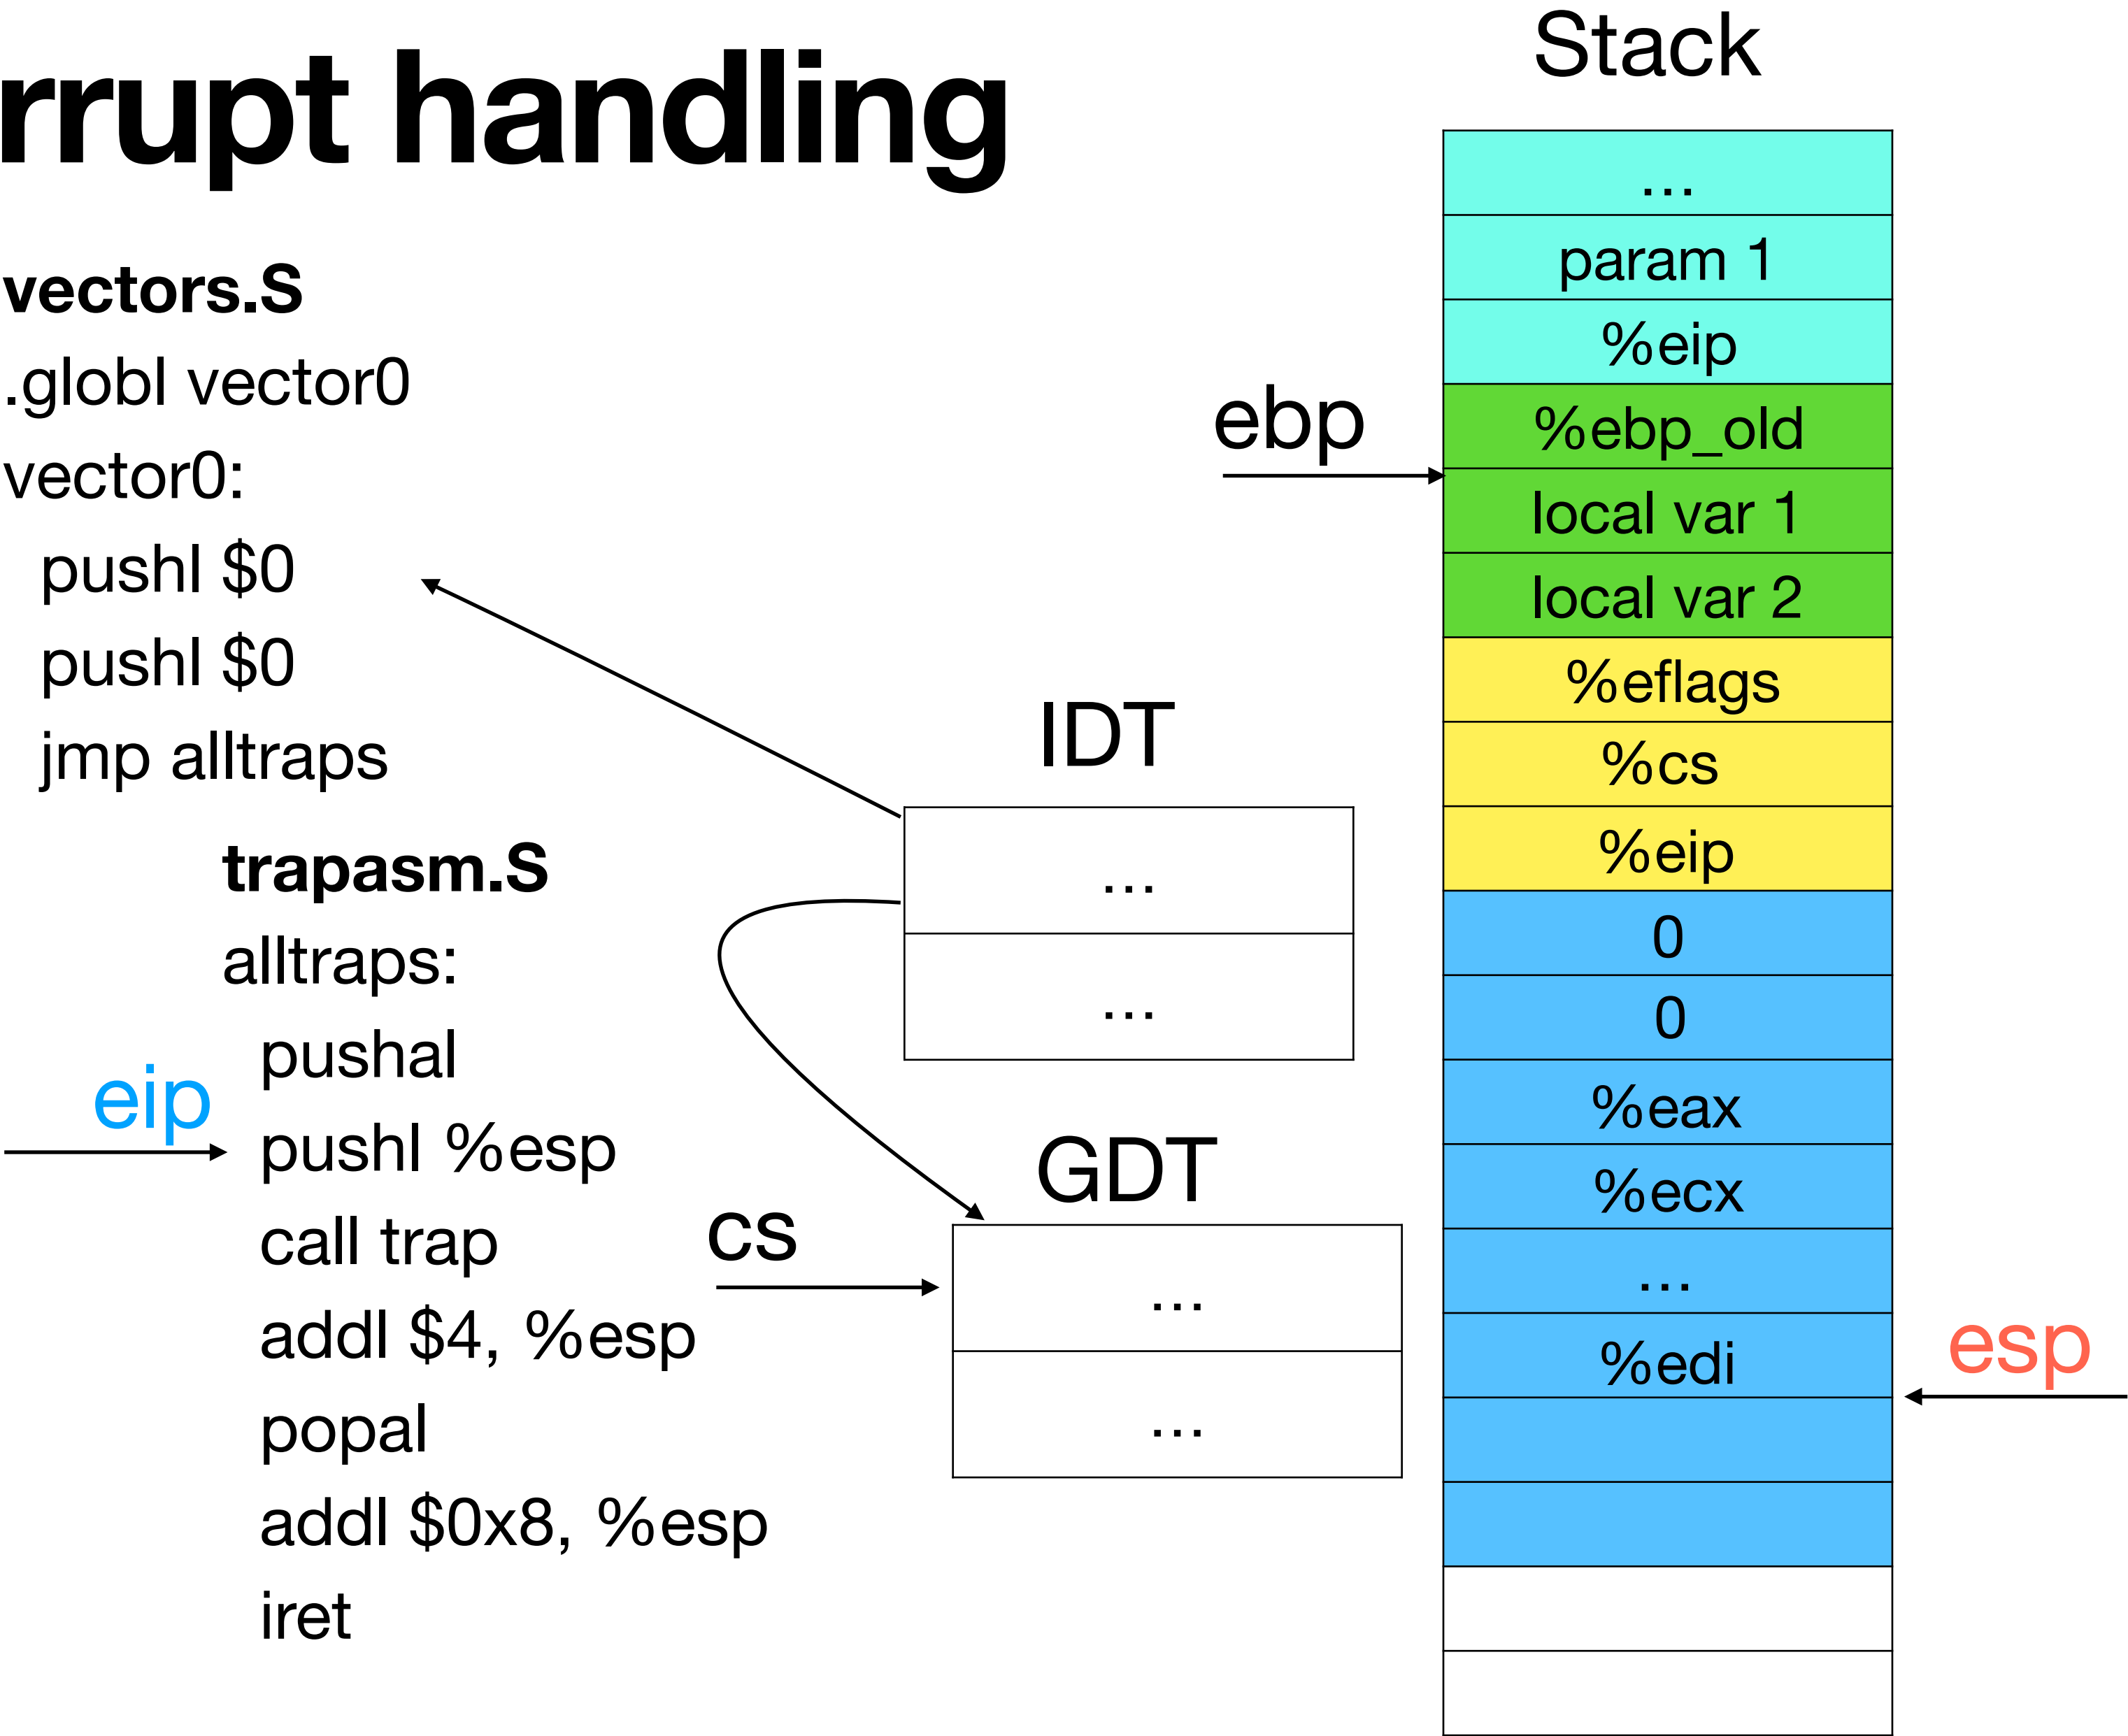
Visualizing interrupt handling

```
for(;;)
;
```

```
trap.c
void
trap(struct trapframe *tf)
{
    switch(tf->trapno){
    case T_IRQ0 + IRQ_TIMER:
        ticks++;
        cprintf("Tick! %d\n\0", ticks);
        lapiceoi();
    ..
    return
```

```
vectors.S
.globl vector0
vector0:
    pushl $0
    pushl $0
    jmp alltraps

trapasm.S
alltraps:
    pushal
    pushl %esp
    call trap
    addl $4, %esp
    popal
    addl $0x8, %esp
    iret
```



Visualizing interrupt handling

```
for(;;)  
;
```

trap.c

```
void  
trap(struct trapframe *tf)  
{  
    switch(tf->trapno){  
    case T_IRQ0 + IRQ_TIMER:  
        ticks++;  
        cprintf("Tick! %d\n\0", ticks);  
        lapiceoi();  
    ..  
    return
```

vectors.S

```
.globl vector0  
vector0:  
    pushl $0  
    pushl $0  
    jmp alltraps
```

trapasm.S

```
alltraps:  
    pushal  
    pushl %esp  
    call trap  
    addl $4, %esp  
    popal  
    addl $0x8, %esp  
    iret
```

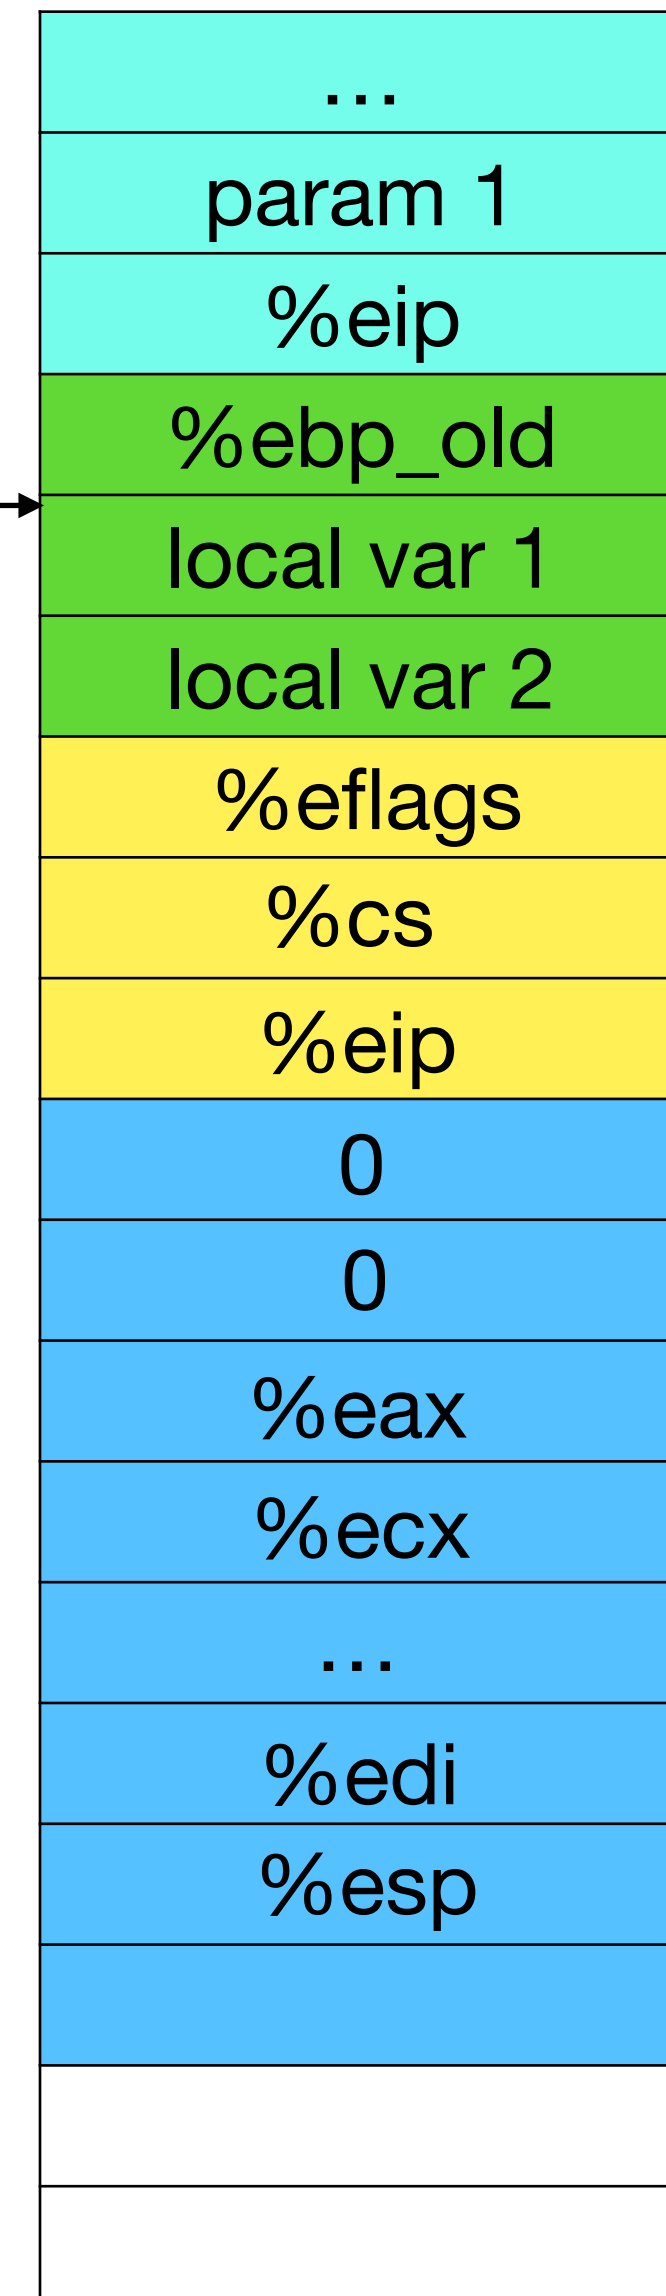
IDT



GDT



Stack



ebp

eip

CS

esp

Visualizing interrupt handling

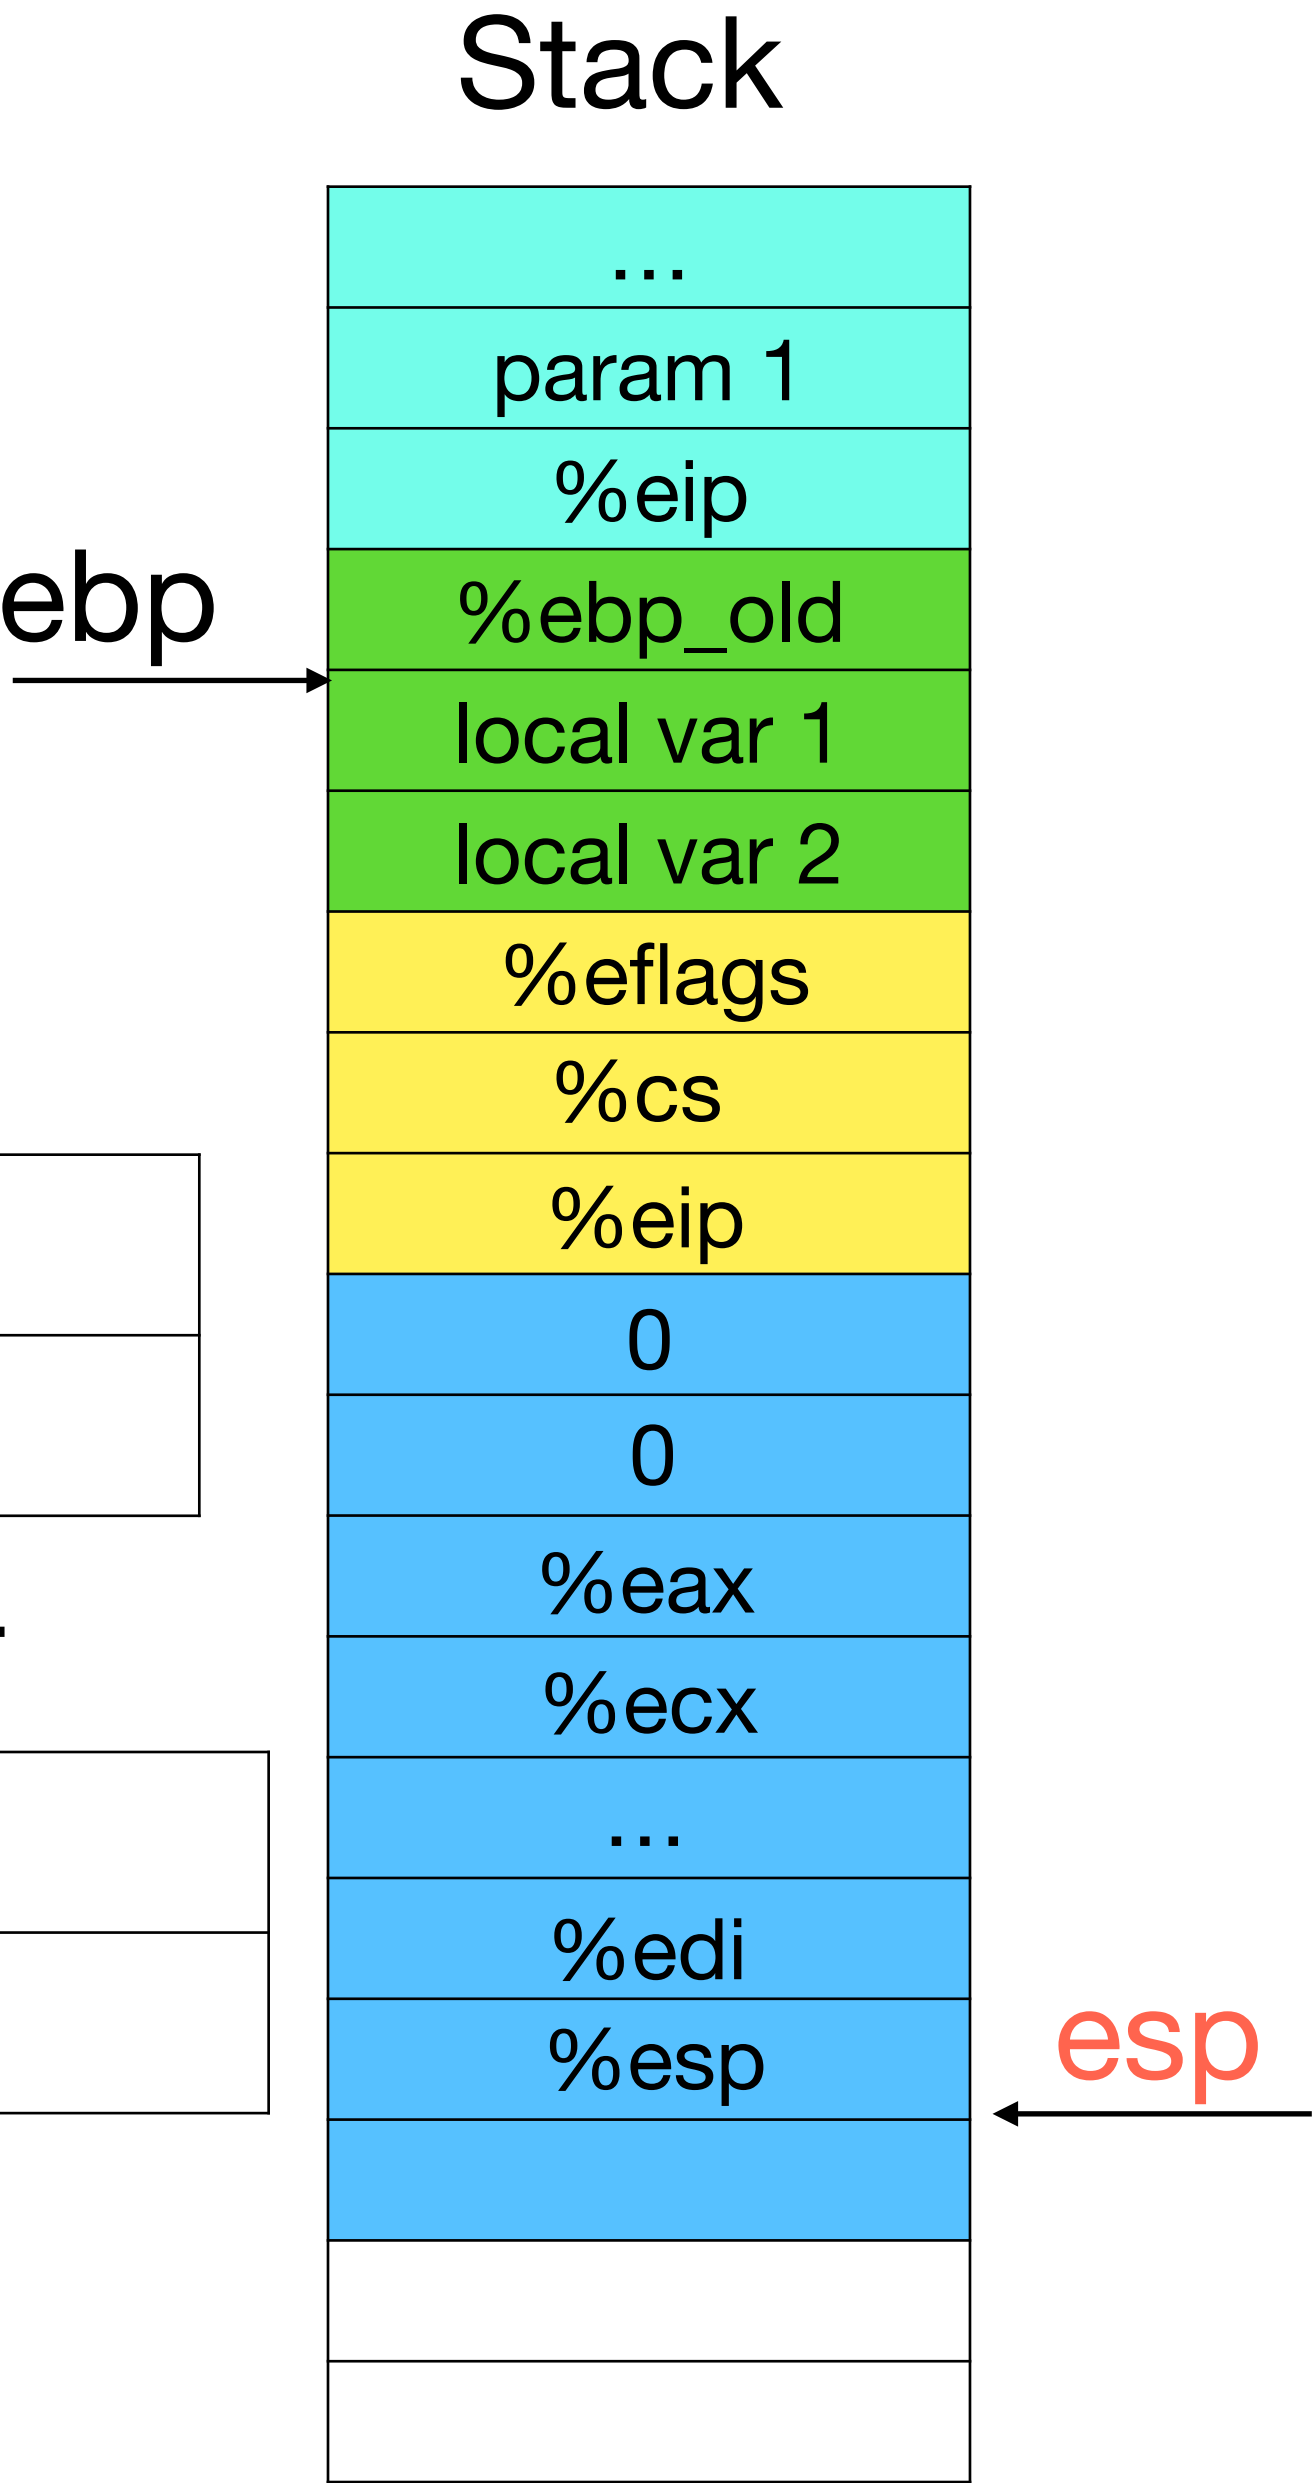
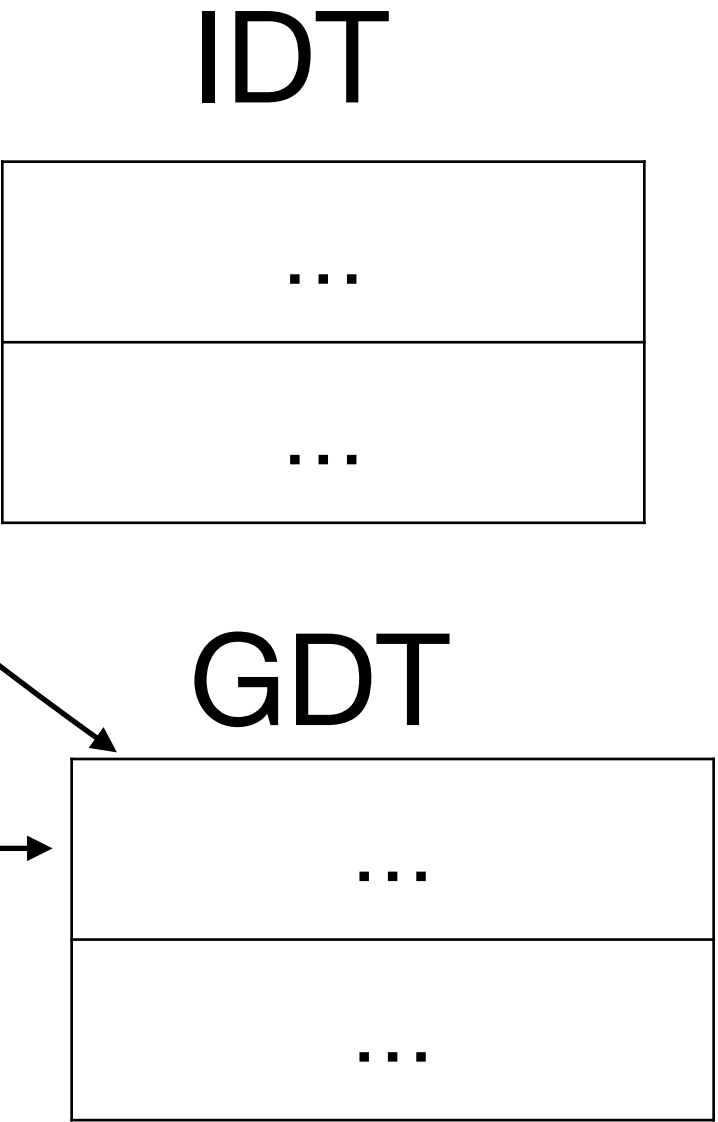
```
for(;;)
;

trap.c
void
trap(struct trapframe *tf)
{
    switch(tf->trapno){
    case T_IRQ0 + IRQ_TIMER:
        ticks++;
        cprintf("Tick! %d\n\0", ticks);
        lapiceoi();
    ..
    return
```

```

vectors.S
.globl vector0
vector0:
    pushl $0
    pushl $0
    jmp alltraps

trapasm.S
alltraps:
    eip → pushal
    pushl %esp
    call trap
    addl $4, %esp
    popal
    addl $0x8, %esp
    iret
    CS →
```



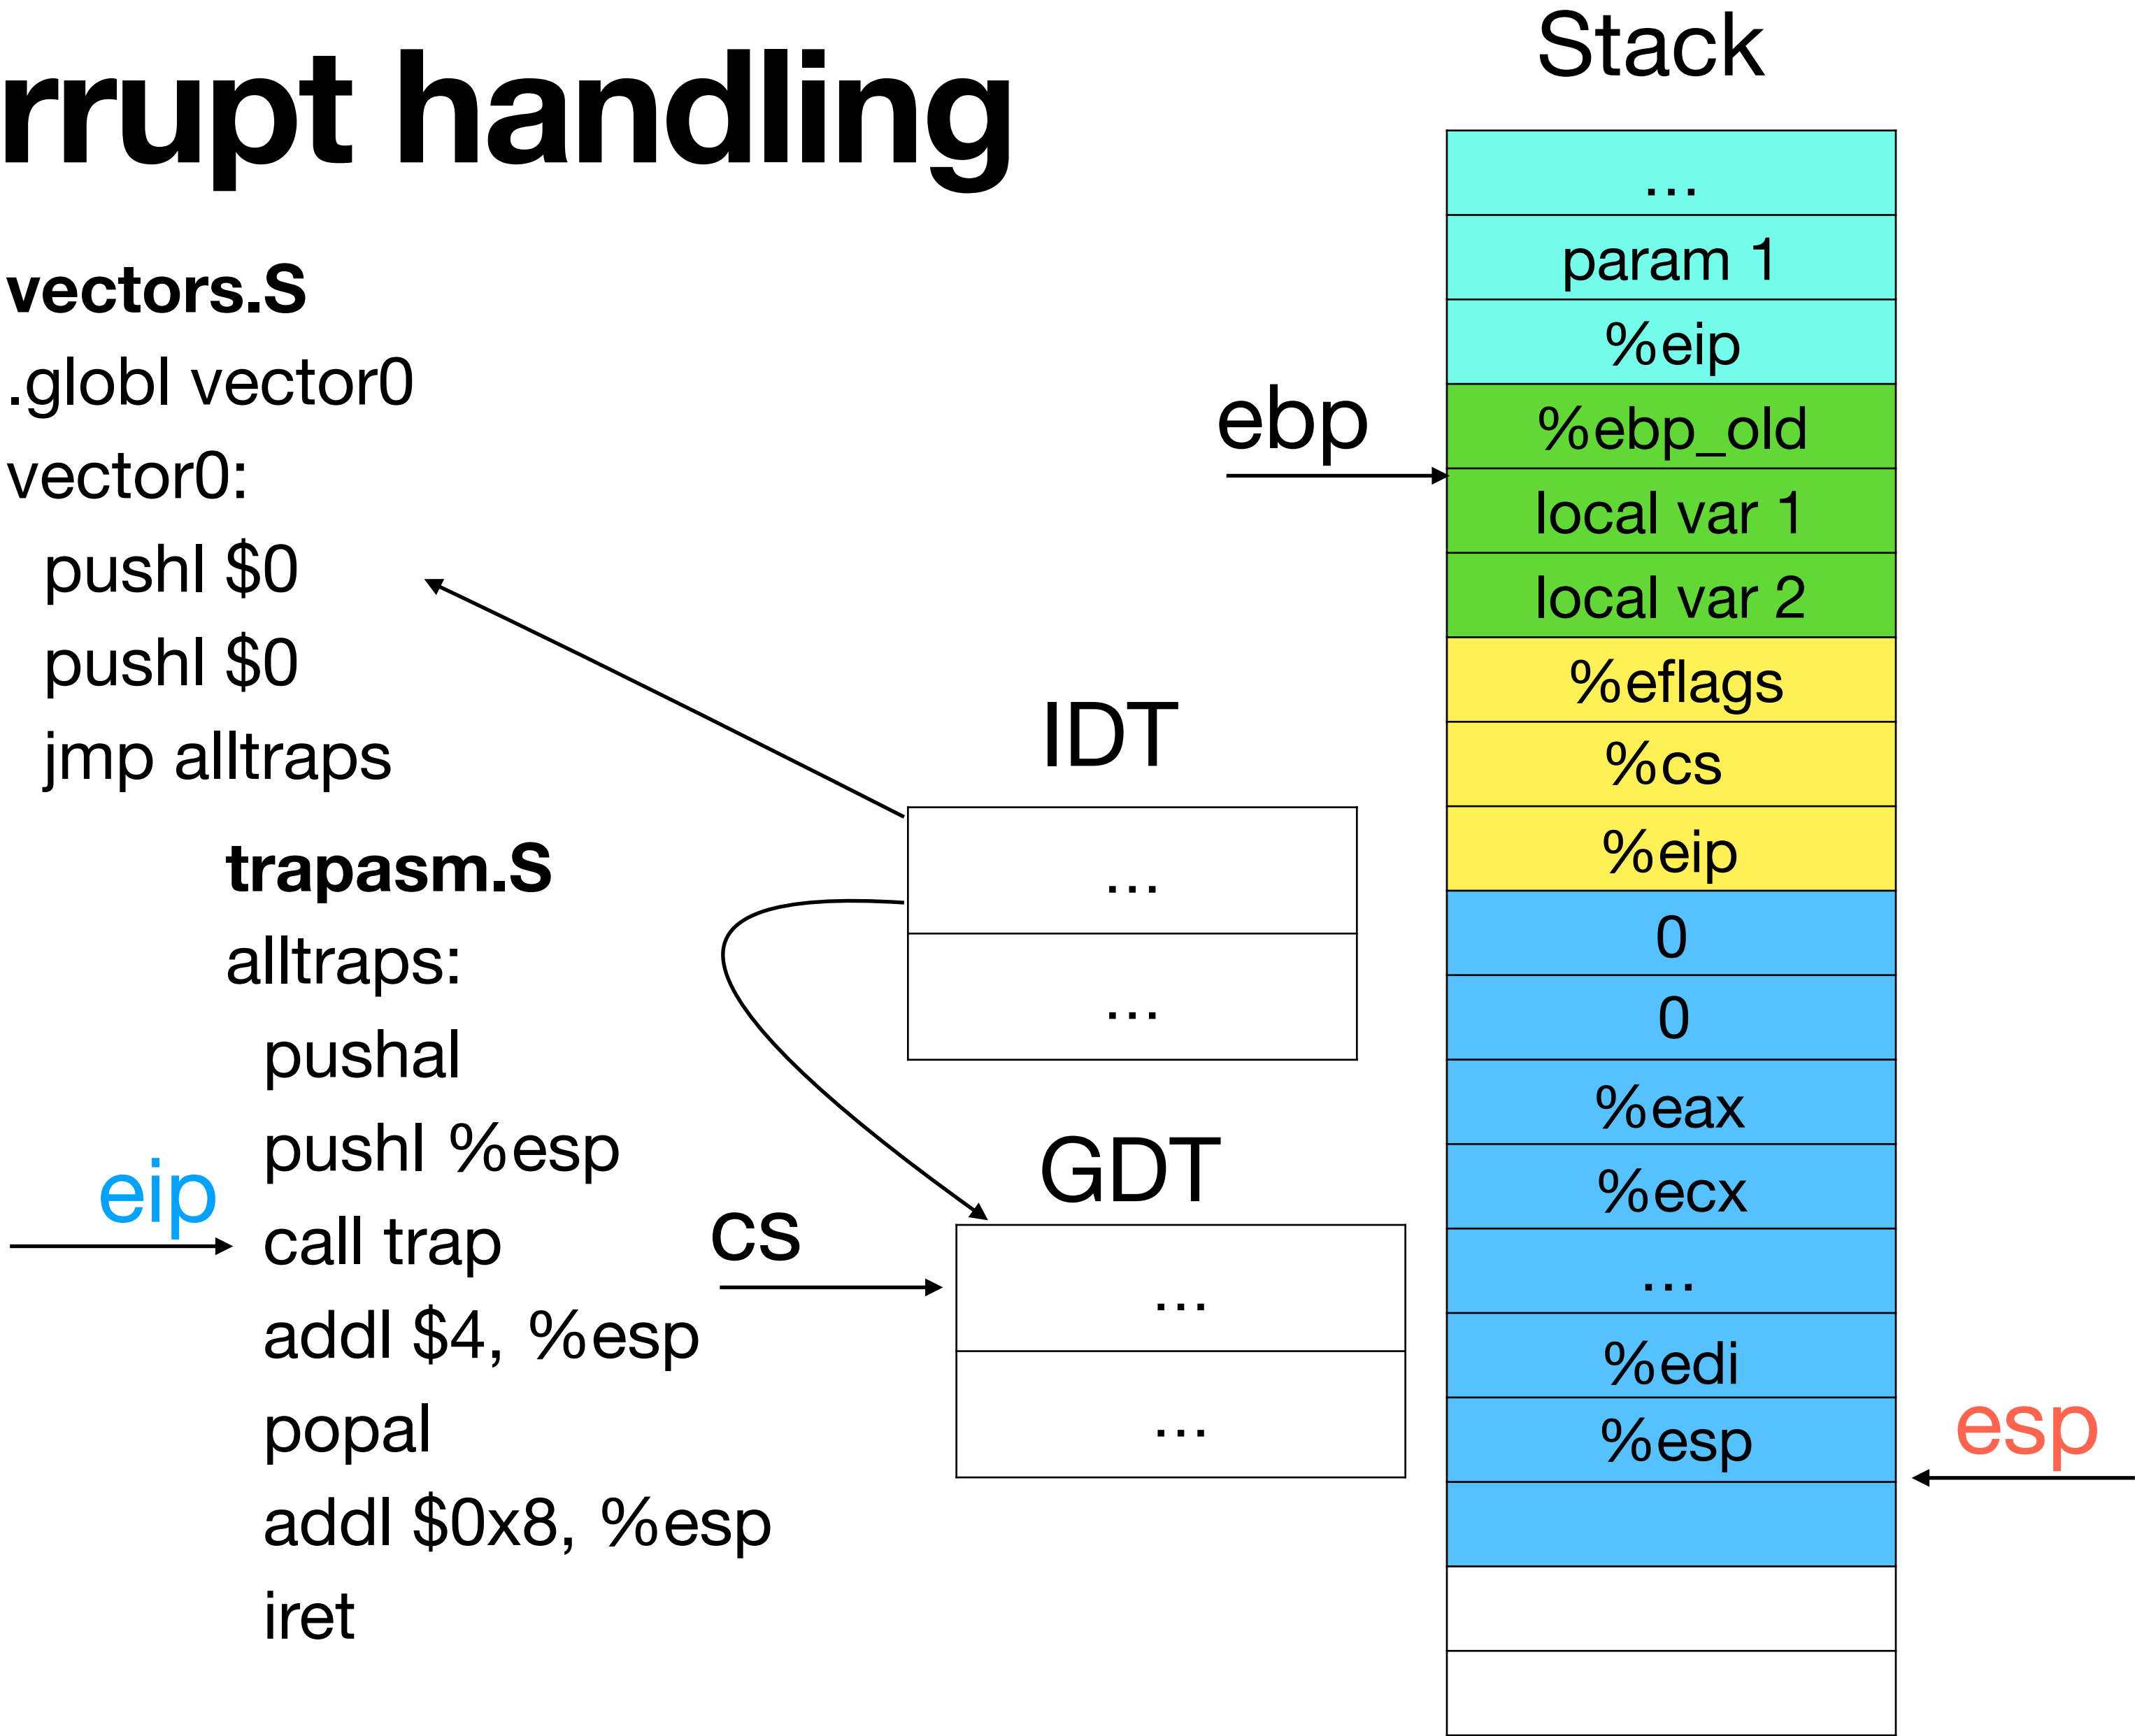
Visualizing interrupt handling

```
for(;;)
;

trap.c
void
trap(struct trapframe *tf)
{
    switch(tf->trapno){
    case T_IRQ0 + IRQ_TIMER:
        ticks++;
        cprintf("Tick! %d\n\0", ticks);
        lapiceoi();
    ..
    return
```

```
vectors.S
.globl vector0
vector0:
    pushl $0
    pushl $0
    jmp alltraps

trapasm.S
alltraps:
    pushal
    pushl %esp
    call trap
    addl $4, %esp
    popal
    addl $0x8, %esp
    iret
```



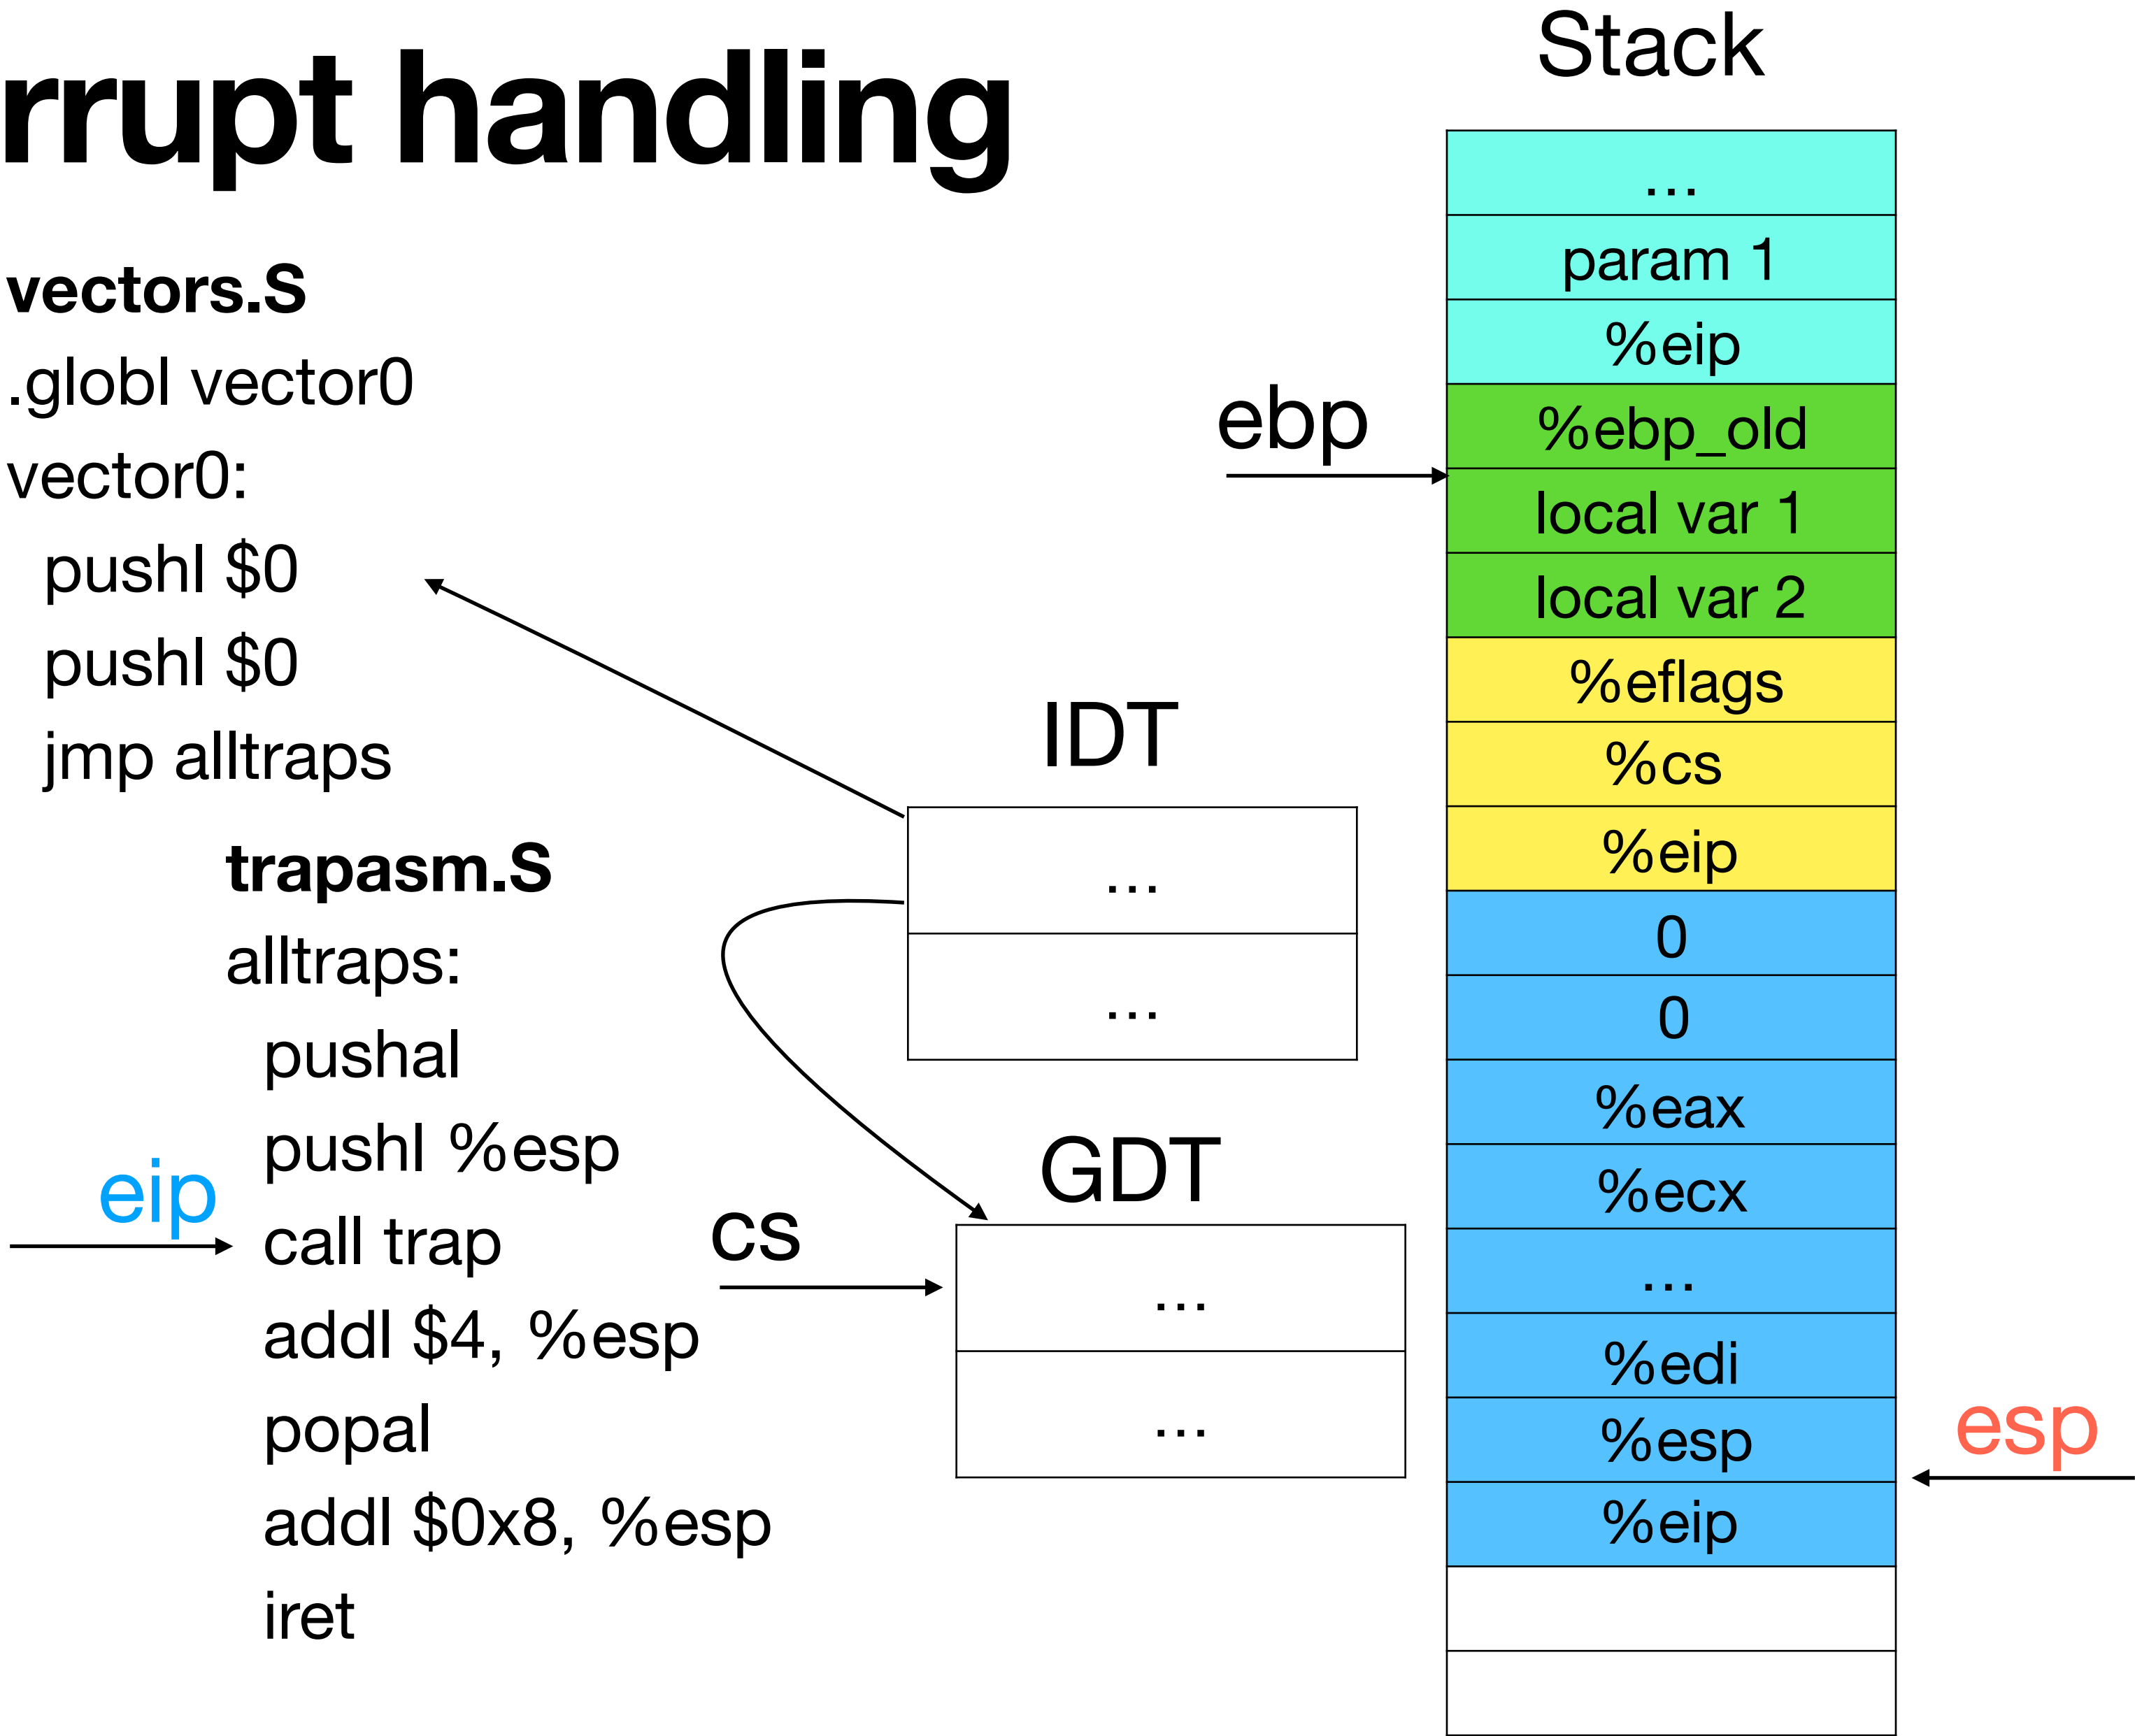
Visualizing interrupt handling

```
for(;;)
;

trap.c
void
trap(struct trapframe *tf)
{
    switch(tf->trapno){
    case T_IRQ0 + IRQ_TIMER:
        ticks++;
        cprintf("Tick! %d\n\0", ticks);
        lapiceoi();
    ..
    return
```

```
vectors.S
.globl vector0
vector0:
    pushl $0
    pushl $0
    jmp alltraps

trapasm.S
alltraps:
    pushal
    pushl %esp
    call trap
    addl $4, %esp
    popal
    addl $0x8, %esp
    iret
```



Visualizing interrupt handling

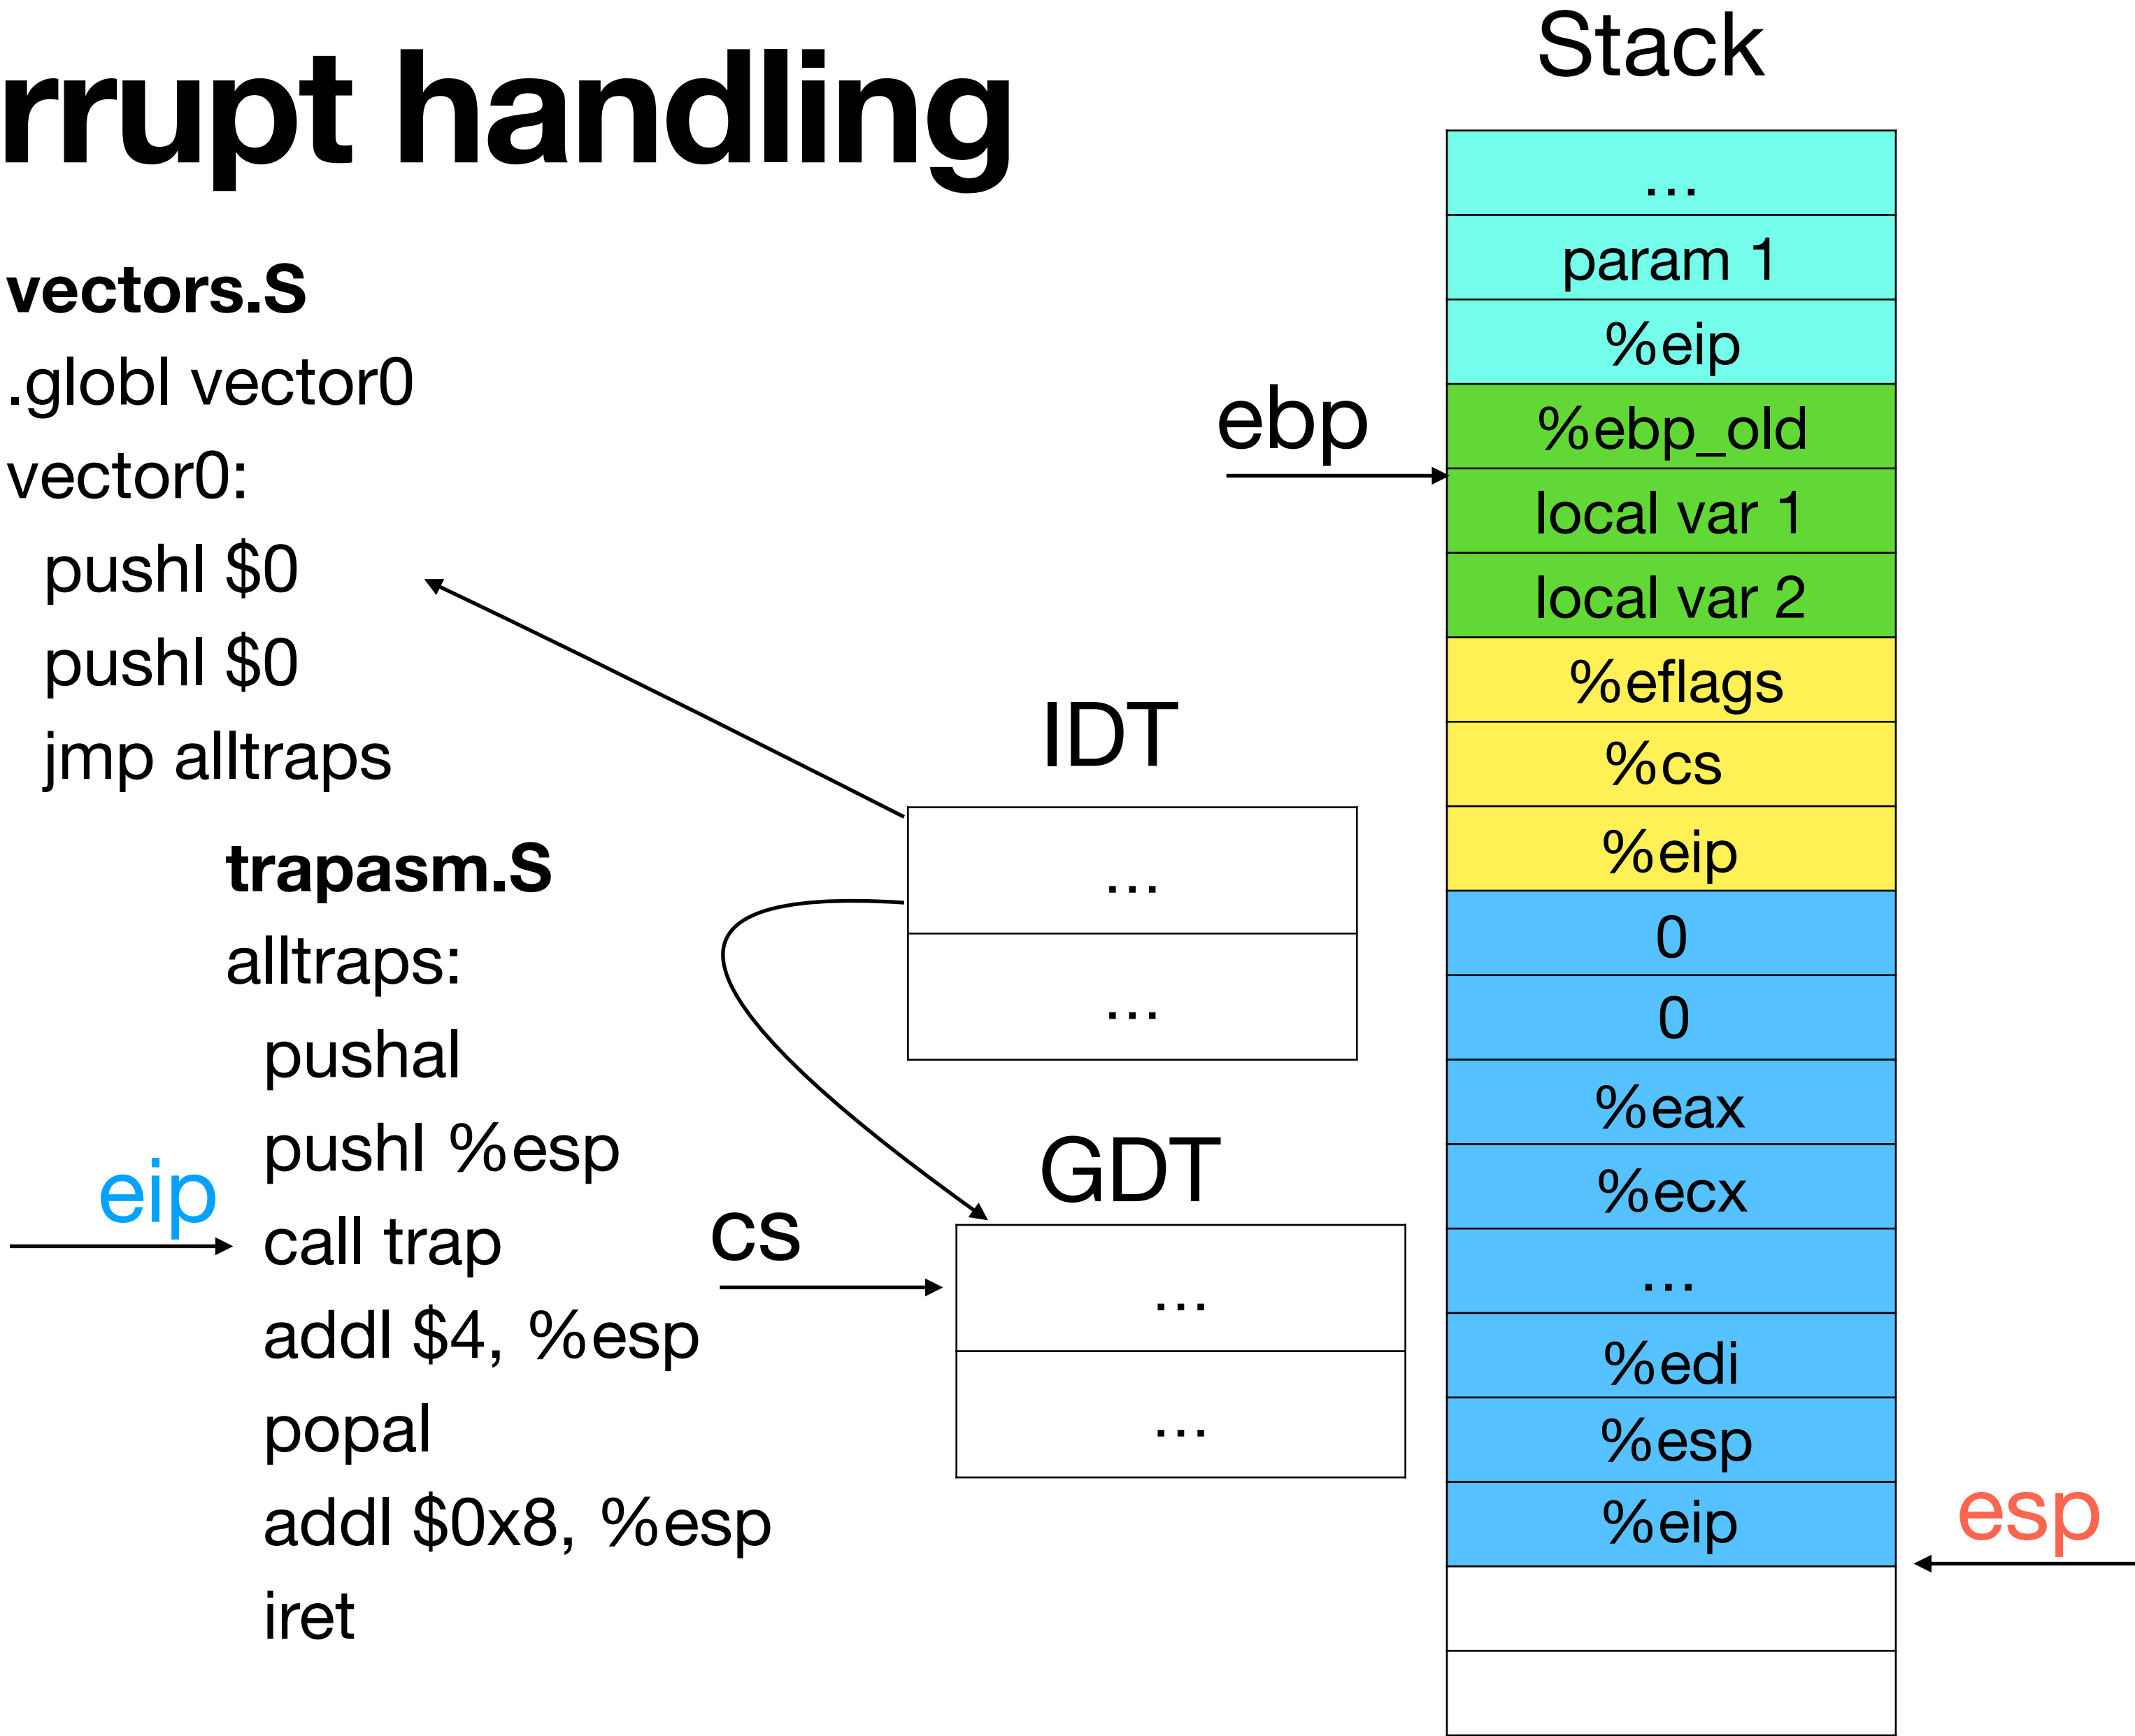
```
for(;;)
;

trap.c
void
trap(struct trapframe *tf)
{
    switch(tf->trapno){
    case T_IRQ0 + IRQ_TIMER:
        ticks++;
        cprintf("Tick! %d\n\0", ticks);
        lapiceoi();
    ..
    return
```

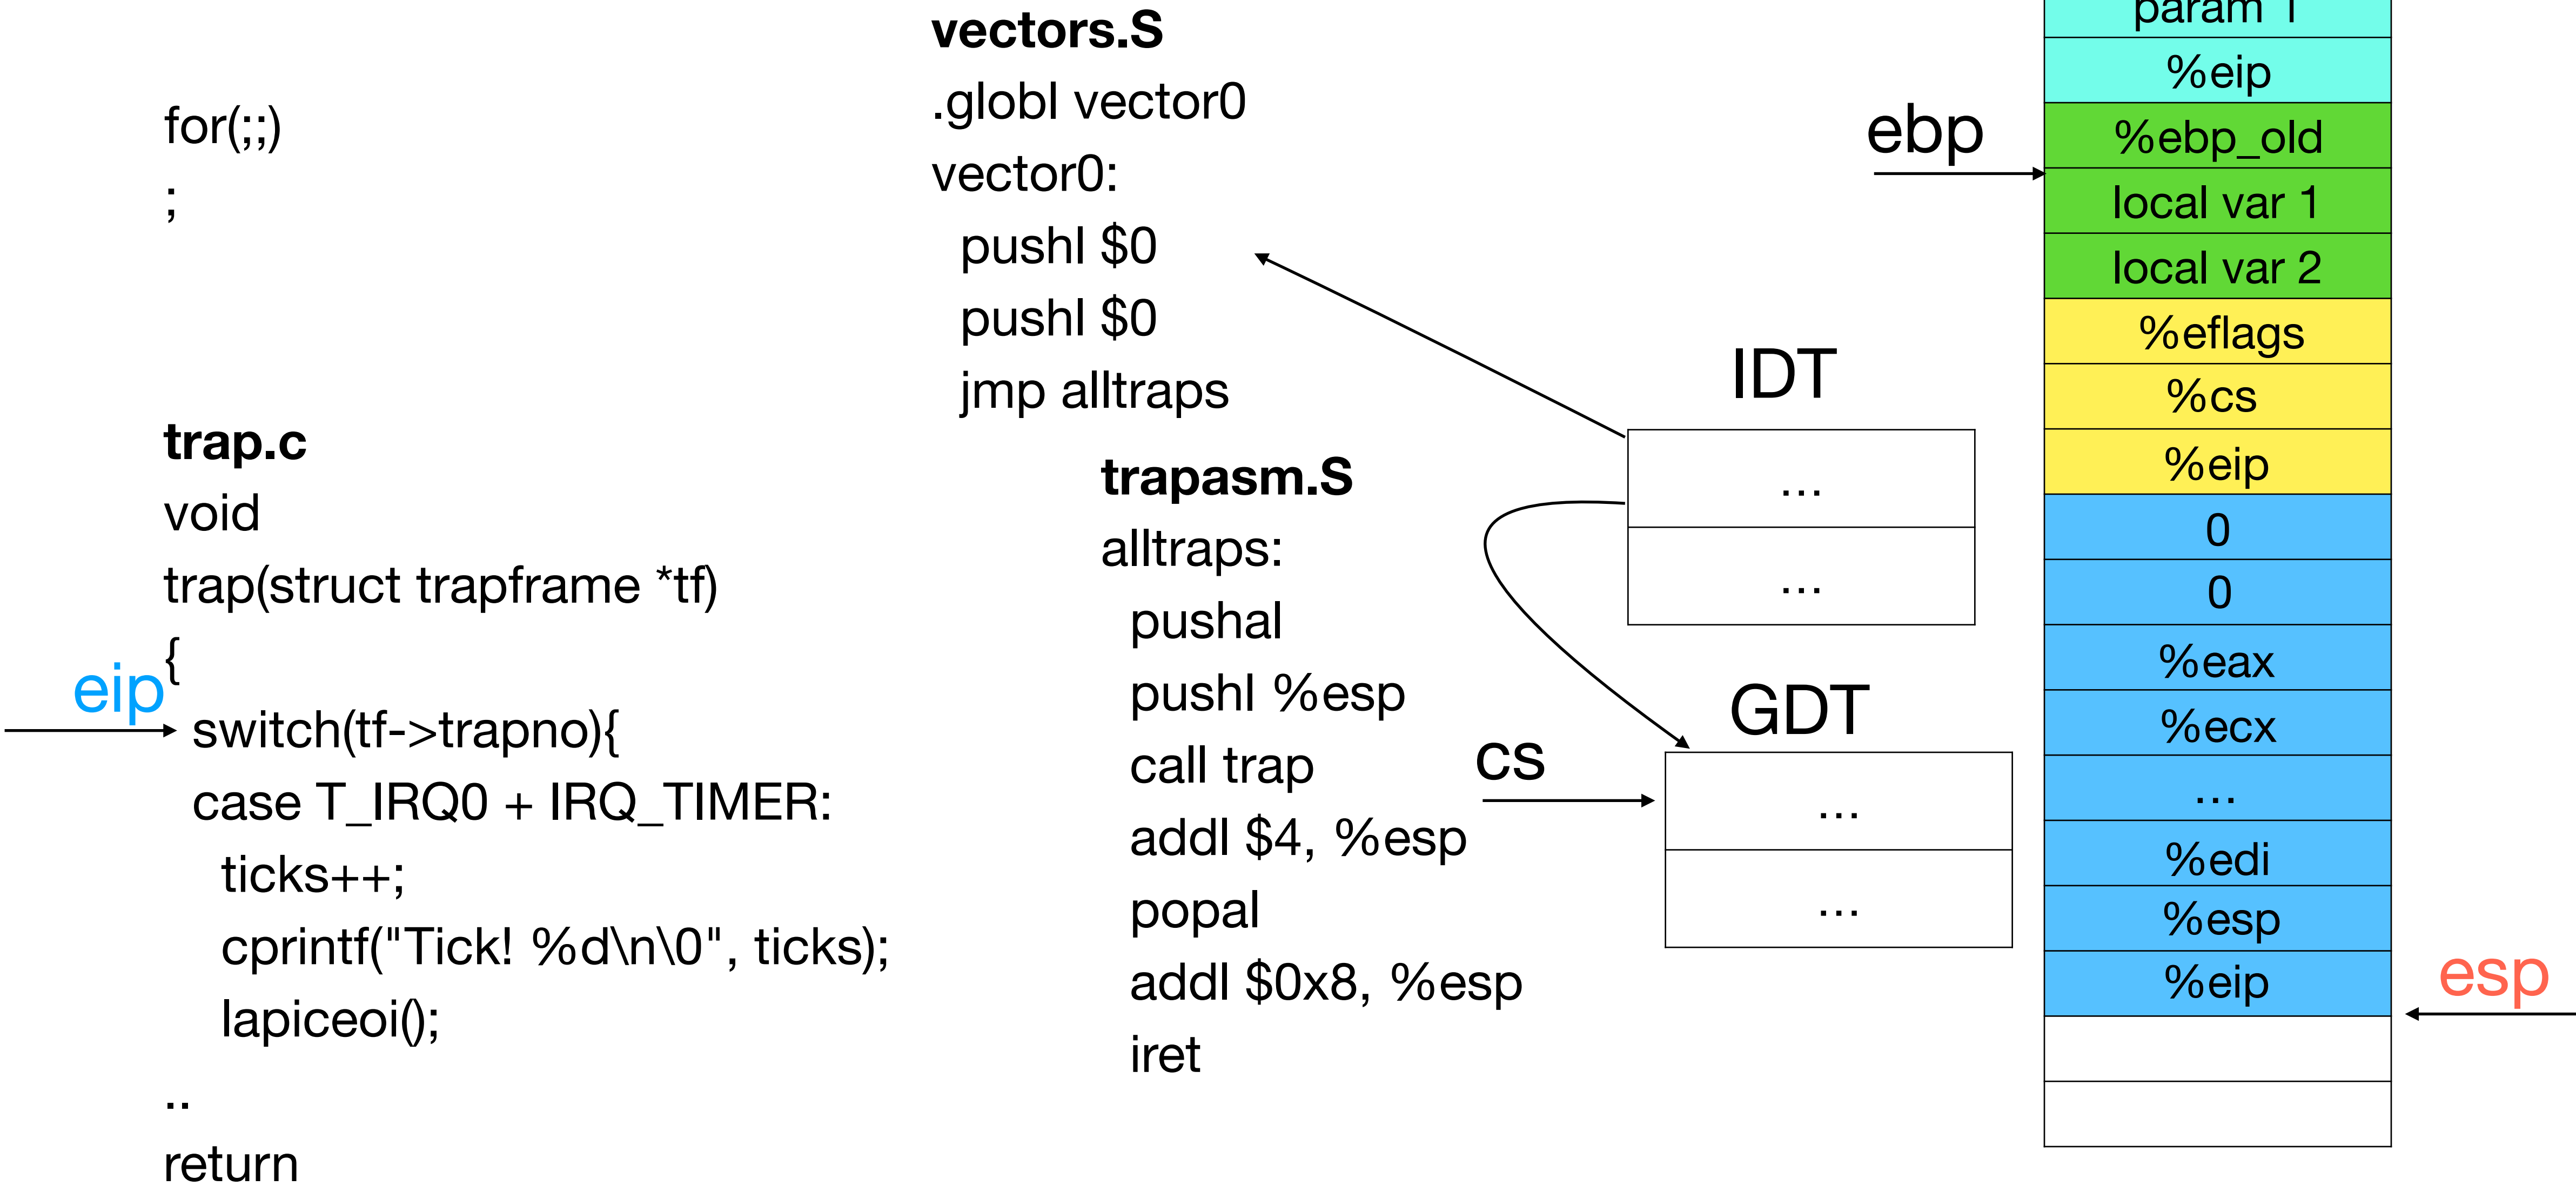
```

vectors.S
.globl vector0
vector0:
    pushl $0
    pushl $0
    jmp alltraps

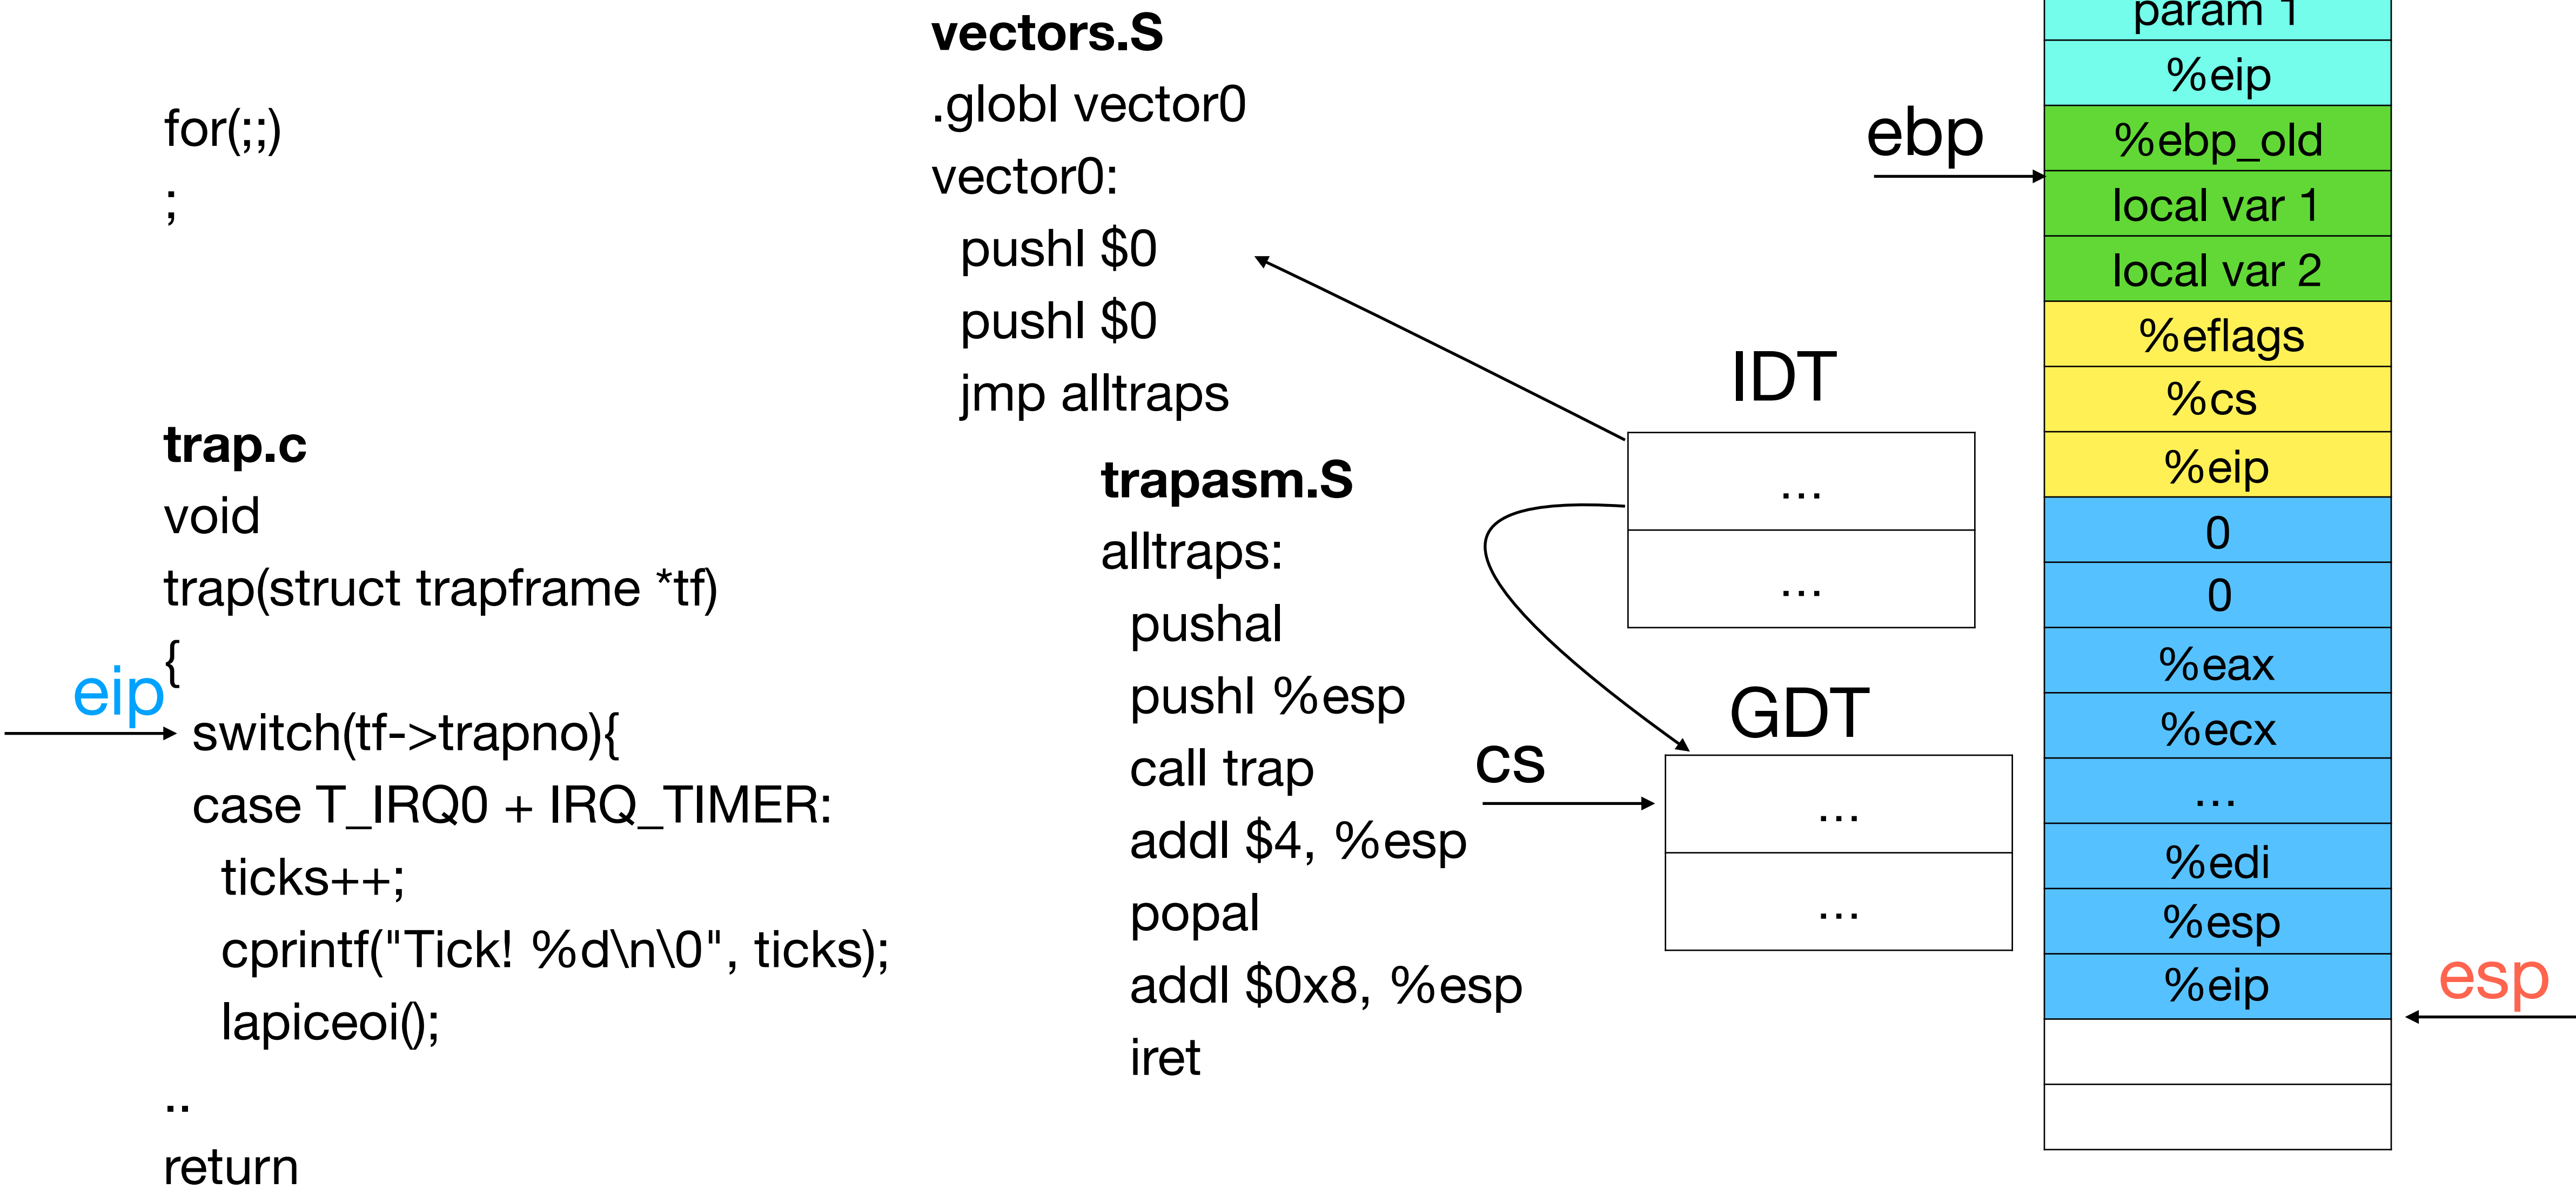
trapasm.S
alltraps:
    pushal
    pushl %esp
    call trap
    addl $4, %esp
    popal
    addl $0x8, %esp
    iret
```



Visualizing interrupt handling



Visualizing interrupt handling



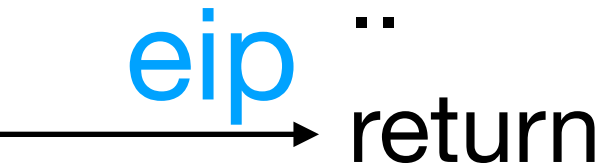
Visualizing interrupt handling

```
for(;;)
;
```

```
trap.c
void
trap(struct trapframe *tf)
{
    switch(tf->trapno){
    case T_IRQ0 + IRQ_TIMER:
        ticks++;
        cprintf("Tick! %d\n\0", ticks);
        lapiceoi();
    }
```

```
vectors.S
.globl vector0
vector0:
    pushl $0
    pushl $0
    jmp alltraps
```

```
trapasm.S
alltraps:
    pushal
    pushl %esp
    call trap
    addl $4, %esp
    popal
    addl $0x8, %esp
    iret
```



Visualizing interrupt handling

```
for(;;)
;
```

```
trap.c
void
trap(struct trapframe *tf)
{
    switch(tf->trapno){
    case T_IRQ0 + IRQ_TIMER:
        ticks++;
        cprintf("Tick! %d\n\0", ticks);
        lapiceoi();
    }
```

```
vectors.S
.globl vector0
vector0:
    pushl $0
    pushl $0
    jmp alltraps
```

```
trapasm.S
alltraps:
    pushal
    pushl %esp
    call trap
    addl $4, %esp
    popal
    addl $0x8, %esp
    iret
```

IDT

GDT

ebp

CS

Stack

...
param 1
%eip
%ebp_old
local var 1
local var 2
%eflags
%cs
%eip
0
0
%eax
%ecx
...
%edi
%esp
%eip

esp

eip
return

Visualizing interrupt handling

```
for(;;)
;
```

```
trap.c
void
trap(struct trapframe *tf)
{
    switch(tf->trapno){
    case T_IRQ0 + IRQ_TIMER:
        ticks++;
        cprintf("Tick! %d\n\0", ticks);
        lapiceoi();
    ..
    return
```

```
vectors.S
.globl vector0
vector0:
    pushl $0
    pushl $0
    jmp alltraps
```

```
trapasm.S
alltraps:
    pushal
    pushl %esp
    call trap
    addl $4, %esp
    popal
    addl $0x8, %esp
    iret
```



Visualizing interrupt handling

```
for(;;)
;
```

```
trap.c
void
trap(struct trapframe *tf)
{
    switch(tf->trapno){
    case T_IRQ0 + IRQ_TIMER:
        ticks++;
        cprintf("Tick! %d\n\0", ticks);
        lapiceoi();
    ..
    return
```

```
vectors.S
.globl vector0
vector0:
    pushl $0
    pushl $0
    jmp alltraps
```

```
trapasm.S
alltraps:
    pushal
    pushl %esp
    call trap
    addl $4, %esp
    popal
    addl $0x8, %esp
    iret
```



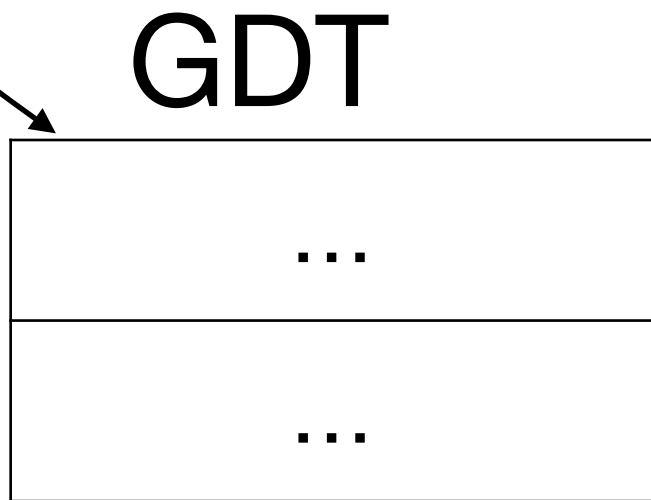
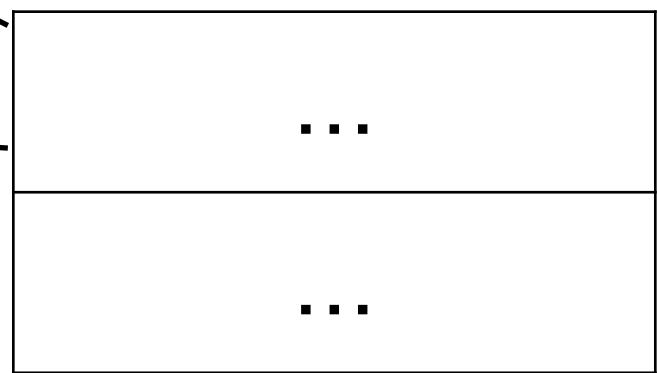
Visualizing interrupt handling

```
for(;;)
;

trap.c
void
trap(struct trapframe *tf)
{
    switch(tf->trapno){
    case T_IRQ0 + IRQ_TIMER:
        ticks++;
        cprintf("Tick! %d\n\0", ticks);
        lapiceoi();
    ..
    return
```

```
vectors.S
.globl vector0
vector0:
    pushl $0
    pushl $0
    jmp alltraps

trapasm.S
alltraps:
    pushal
    pushl %esp
    call trap
    addl $4, %esp
    popal
    addl $0x8, %esp
    iret
```



CS

Stack

ebp



esp



Visualizing interrupt handling

```
for(;;)
;

trap.c
void
trap(struct trapframe *tf)
{
    switch(tf->trapno){
    case T_IRQ0 + IRQ_TIMER:
        ticks++;
        cprintf("Tick! %d\n\0", ticks);
        lapiceoi();
    ..
    return
```

```
vectors.S
.globl vector0
vector0:
    pushl $0
    pushl $0
    jmp alltraps

trapasm.S
alltraps:
    pushal
    pushl %esp
    call trap
    addl $4, %esp
    popal
    addl $0x8, %esp
    iret
```

eip

CS

IDT

GDT

ebp

esp

Stack

...
param 1
%eip
%ebp_old
local var 1
local var 2
%eflags
%cs
%eip
0
0
%eax
%ecx
...
%edi
%esp
%eip

Visualizing interrupt handling

Stack

...
param 1
%eip
%ebp_old
local var 1
local var 2
%eflags
%cs
%eip
0
0
%eax
%ecx
...
%edi
%esp
%eip

ebp

esp

IDT

...
...

GDT

...
...

vectors.S

```
.globl vector0
vector0:
    pushl $0
    pushl $0
    jmp alltraps
```

trapasm.S

```
alltraps:
    pushal
    pushl %esp
    call trap
    addl $4, %esp
    popal
    addl $0x8, %esp
    iret
```

eip

CS

```
for(;;)
```

```
;
```

trap.c

```
void
trap(struct trapframe *tf)
{
    switch(tf->trapno){
    case T_IRQ0 + IRQ_TIMER:
        ticks++;
        cprintf("Tick! %d\n\0", ticks);
        lapiceoi();
    ..
    return
```


Visualizing interrupt handling

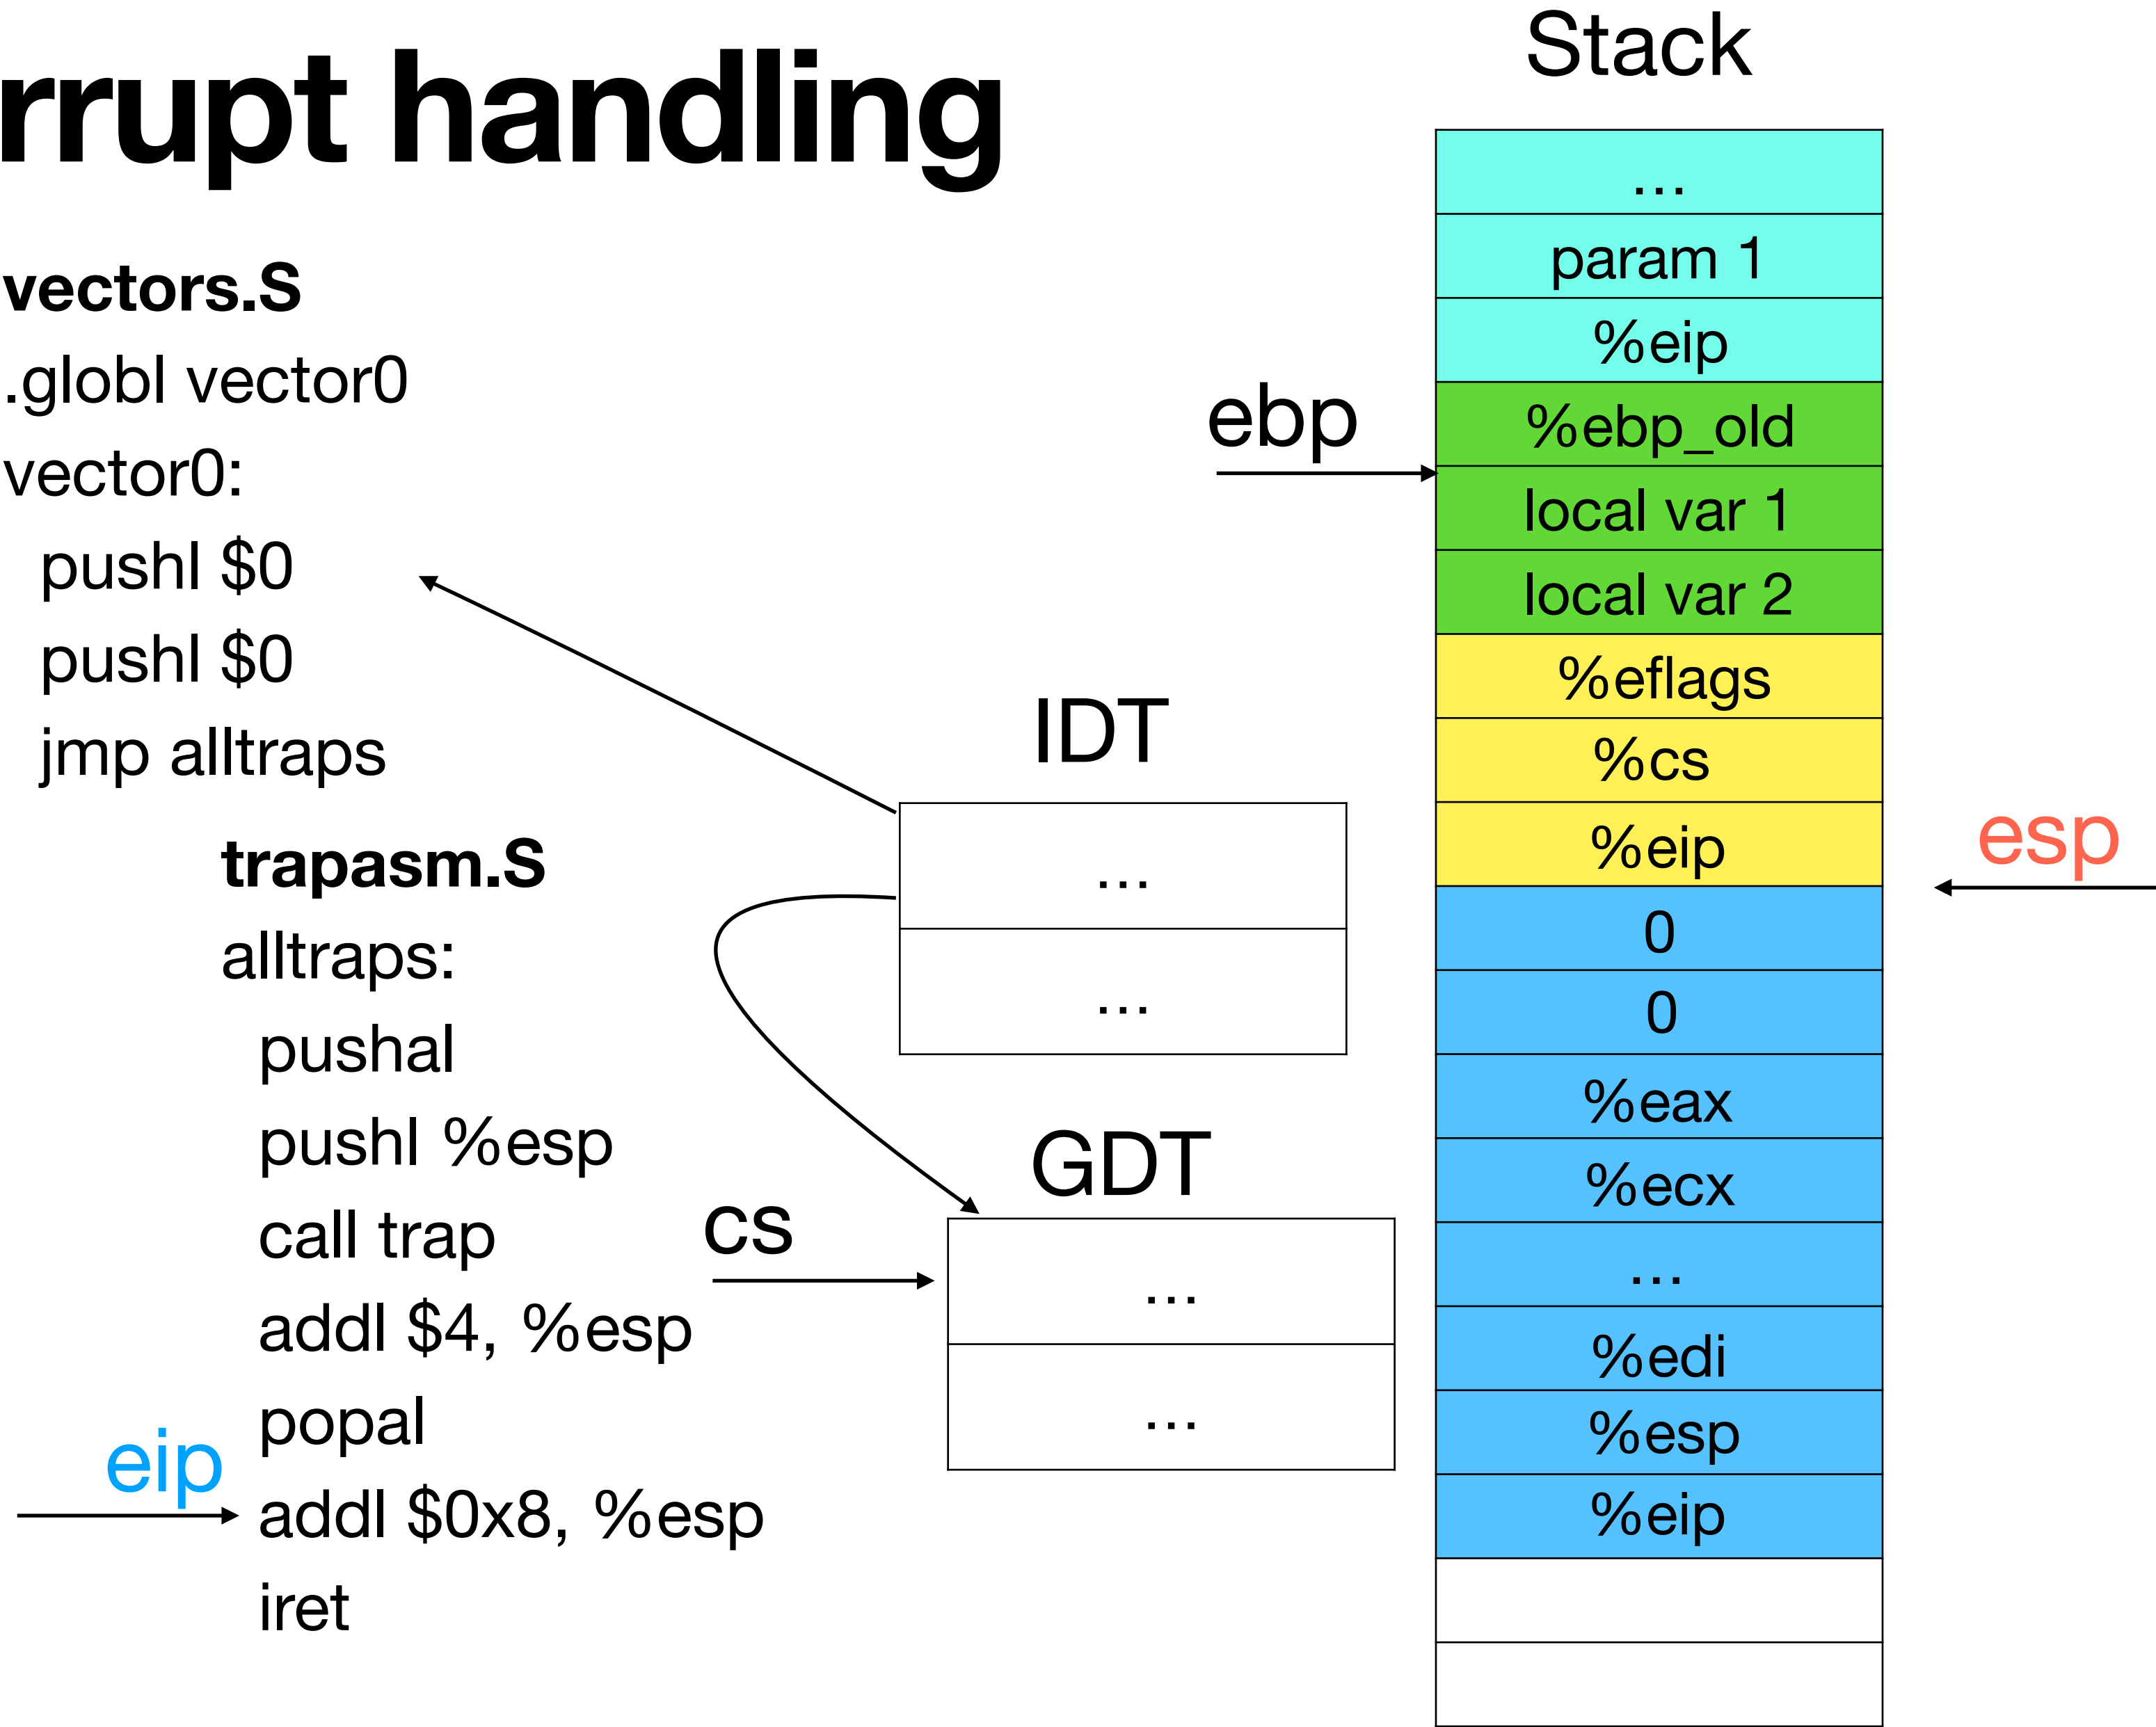
```
for(;;)
;

trap.c
void
trap(struct trapframe *tf)
{
    switch(tf->trapno){
    case T_IRQ0 + IRQ_TIMER:
        ticks++;
        cprintf("Tick! %d\n\0", ticks);
        lapiceoi();
    ..
    return
```

```

vectors.S
.globl vector0
vector0:
    pushl $0
    pushl $0
    jmp alltraps

trapasm.S
alltraps:
    pushal
    pushl %esp
    call trap
    addl $4, %esp
    popal
    addl $0x8, %esp
    iret
```



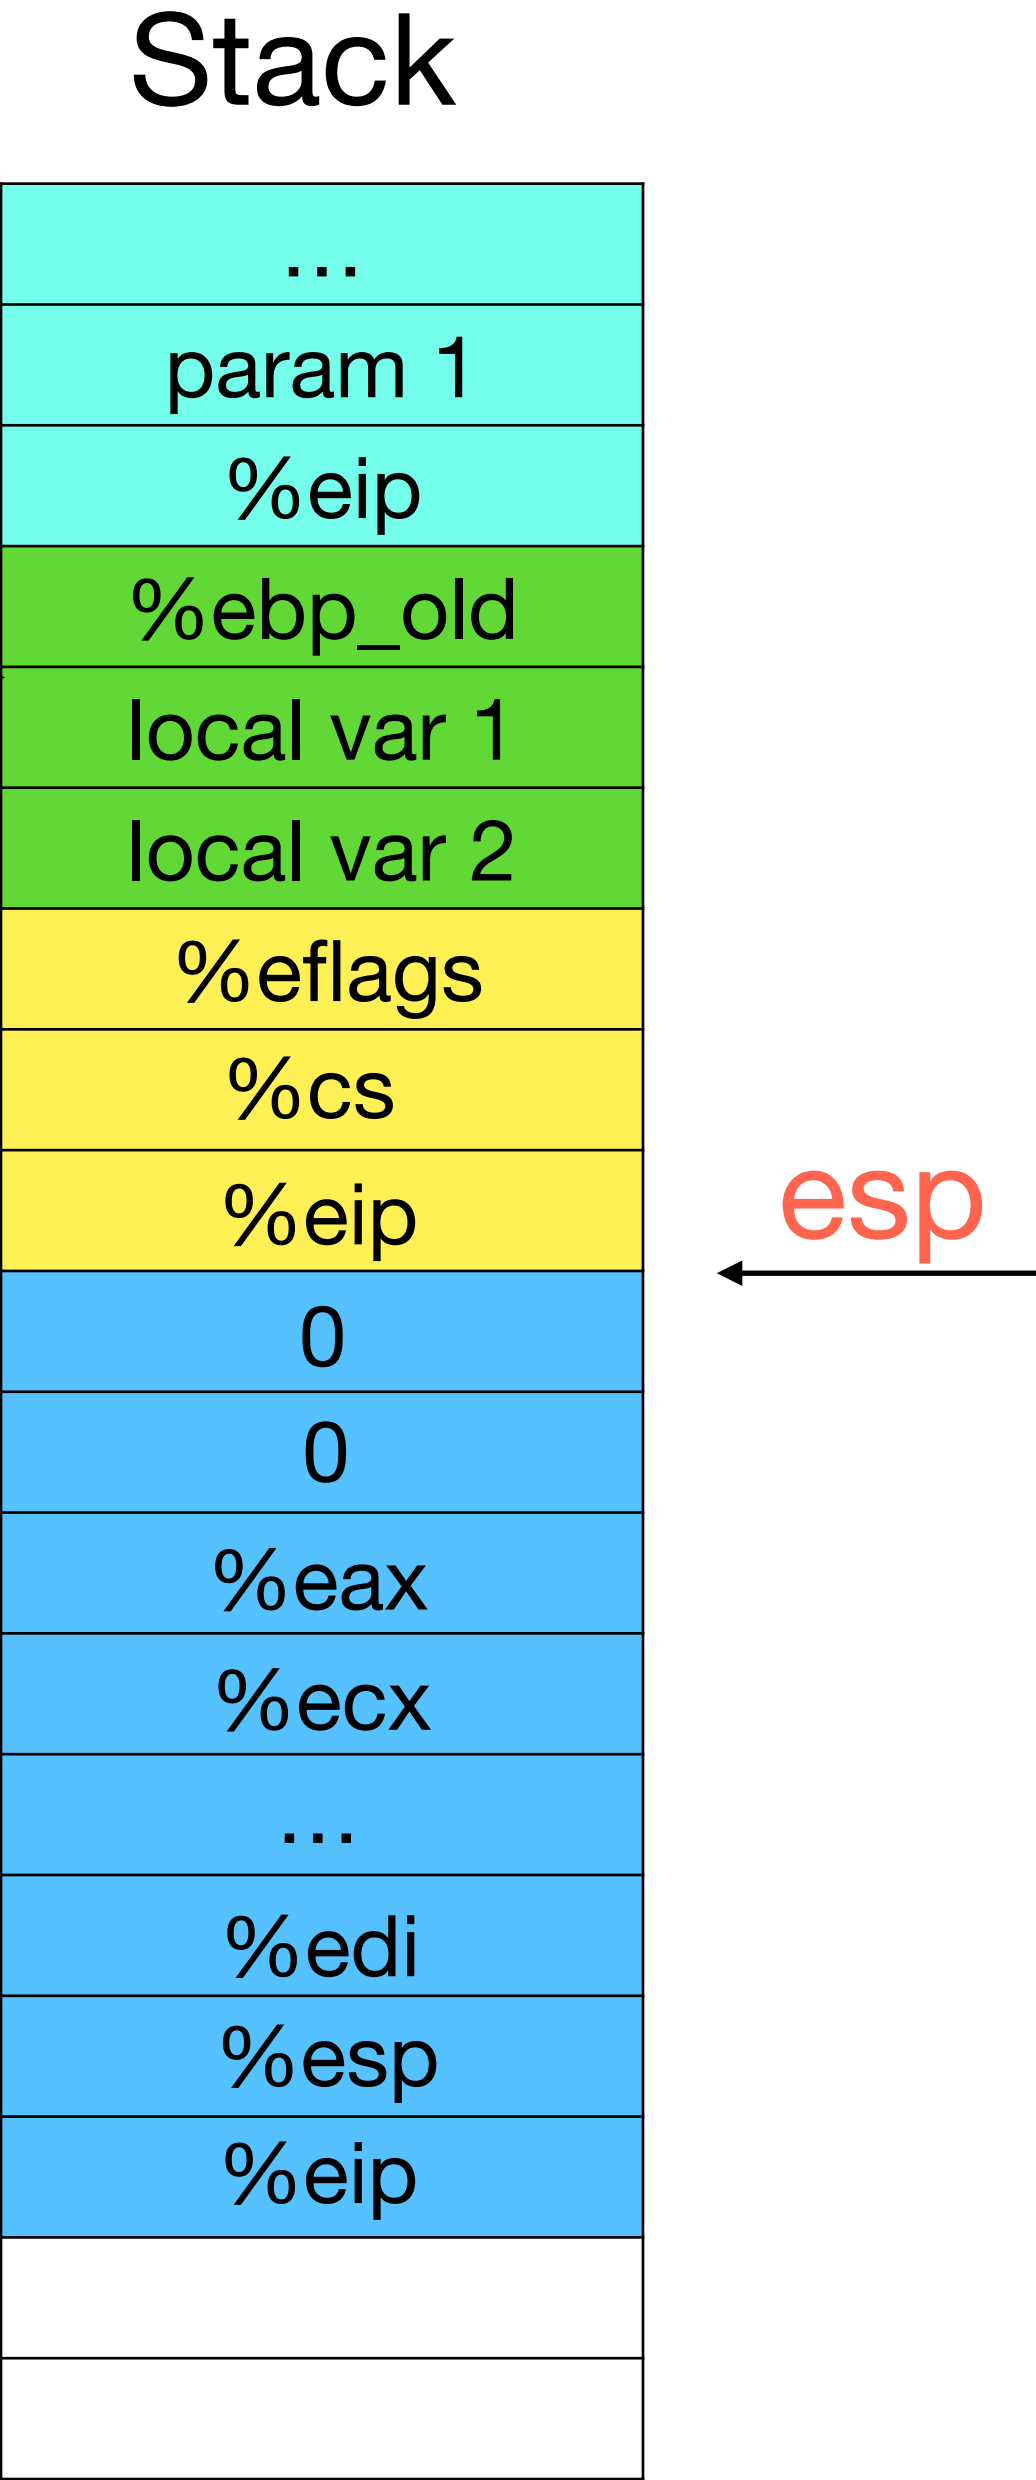
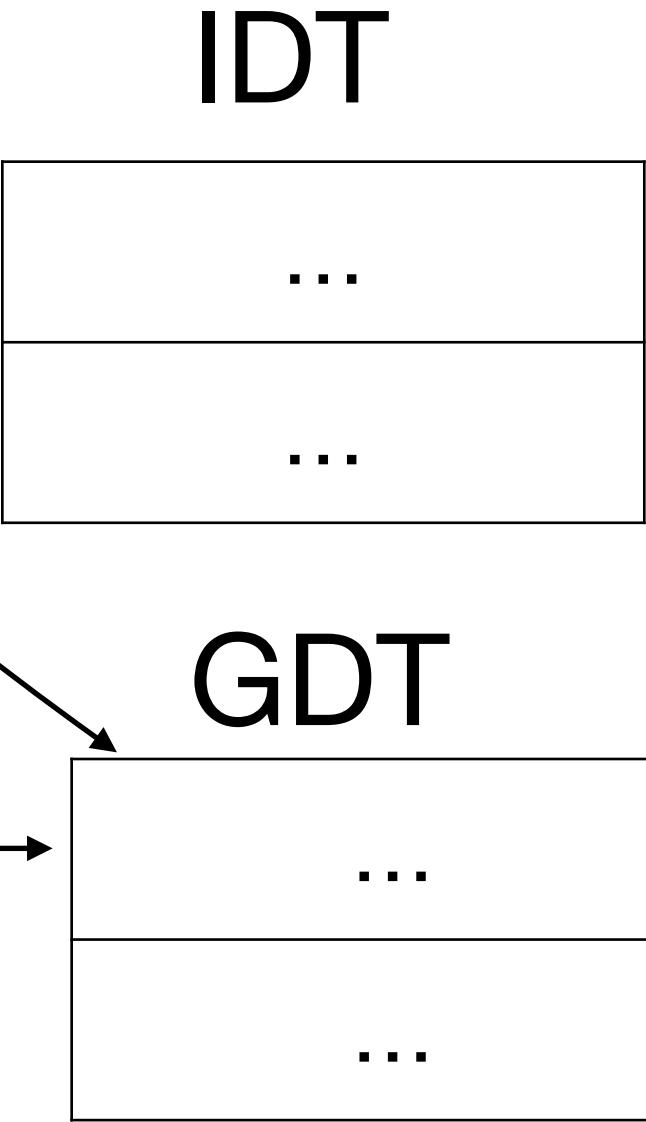
Visualizing interrupt handling

```
for(;;)
;
```

```
trap.c
void
trap(struct trapframe *tf)
{
    switch(tf->trapno){
    case T_IRQ0 + IRQ_TIMER:
        ticks++;
        cprintf("Tick! %d\n\0", ticks);
        lapiceoi();
    ..
    return
```

```
vectors.S
.globl vector0
vector0:
    pushl $0
    pushl $0
    jmp alltraps
```

```
trapasm.S
alltraps:
    pushal
    pushl %esp
    call trap
    addl $4, %esp
    popal
    addl $0x8, %esp
    iret
```



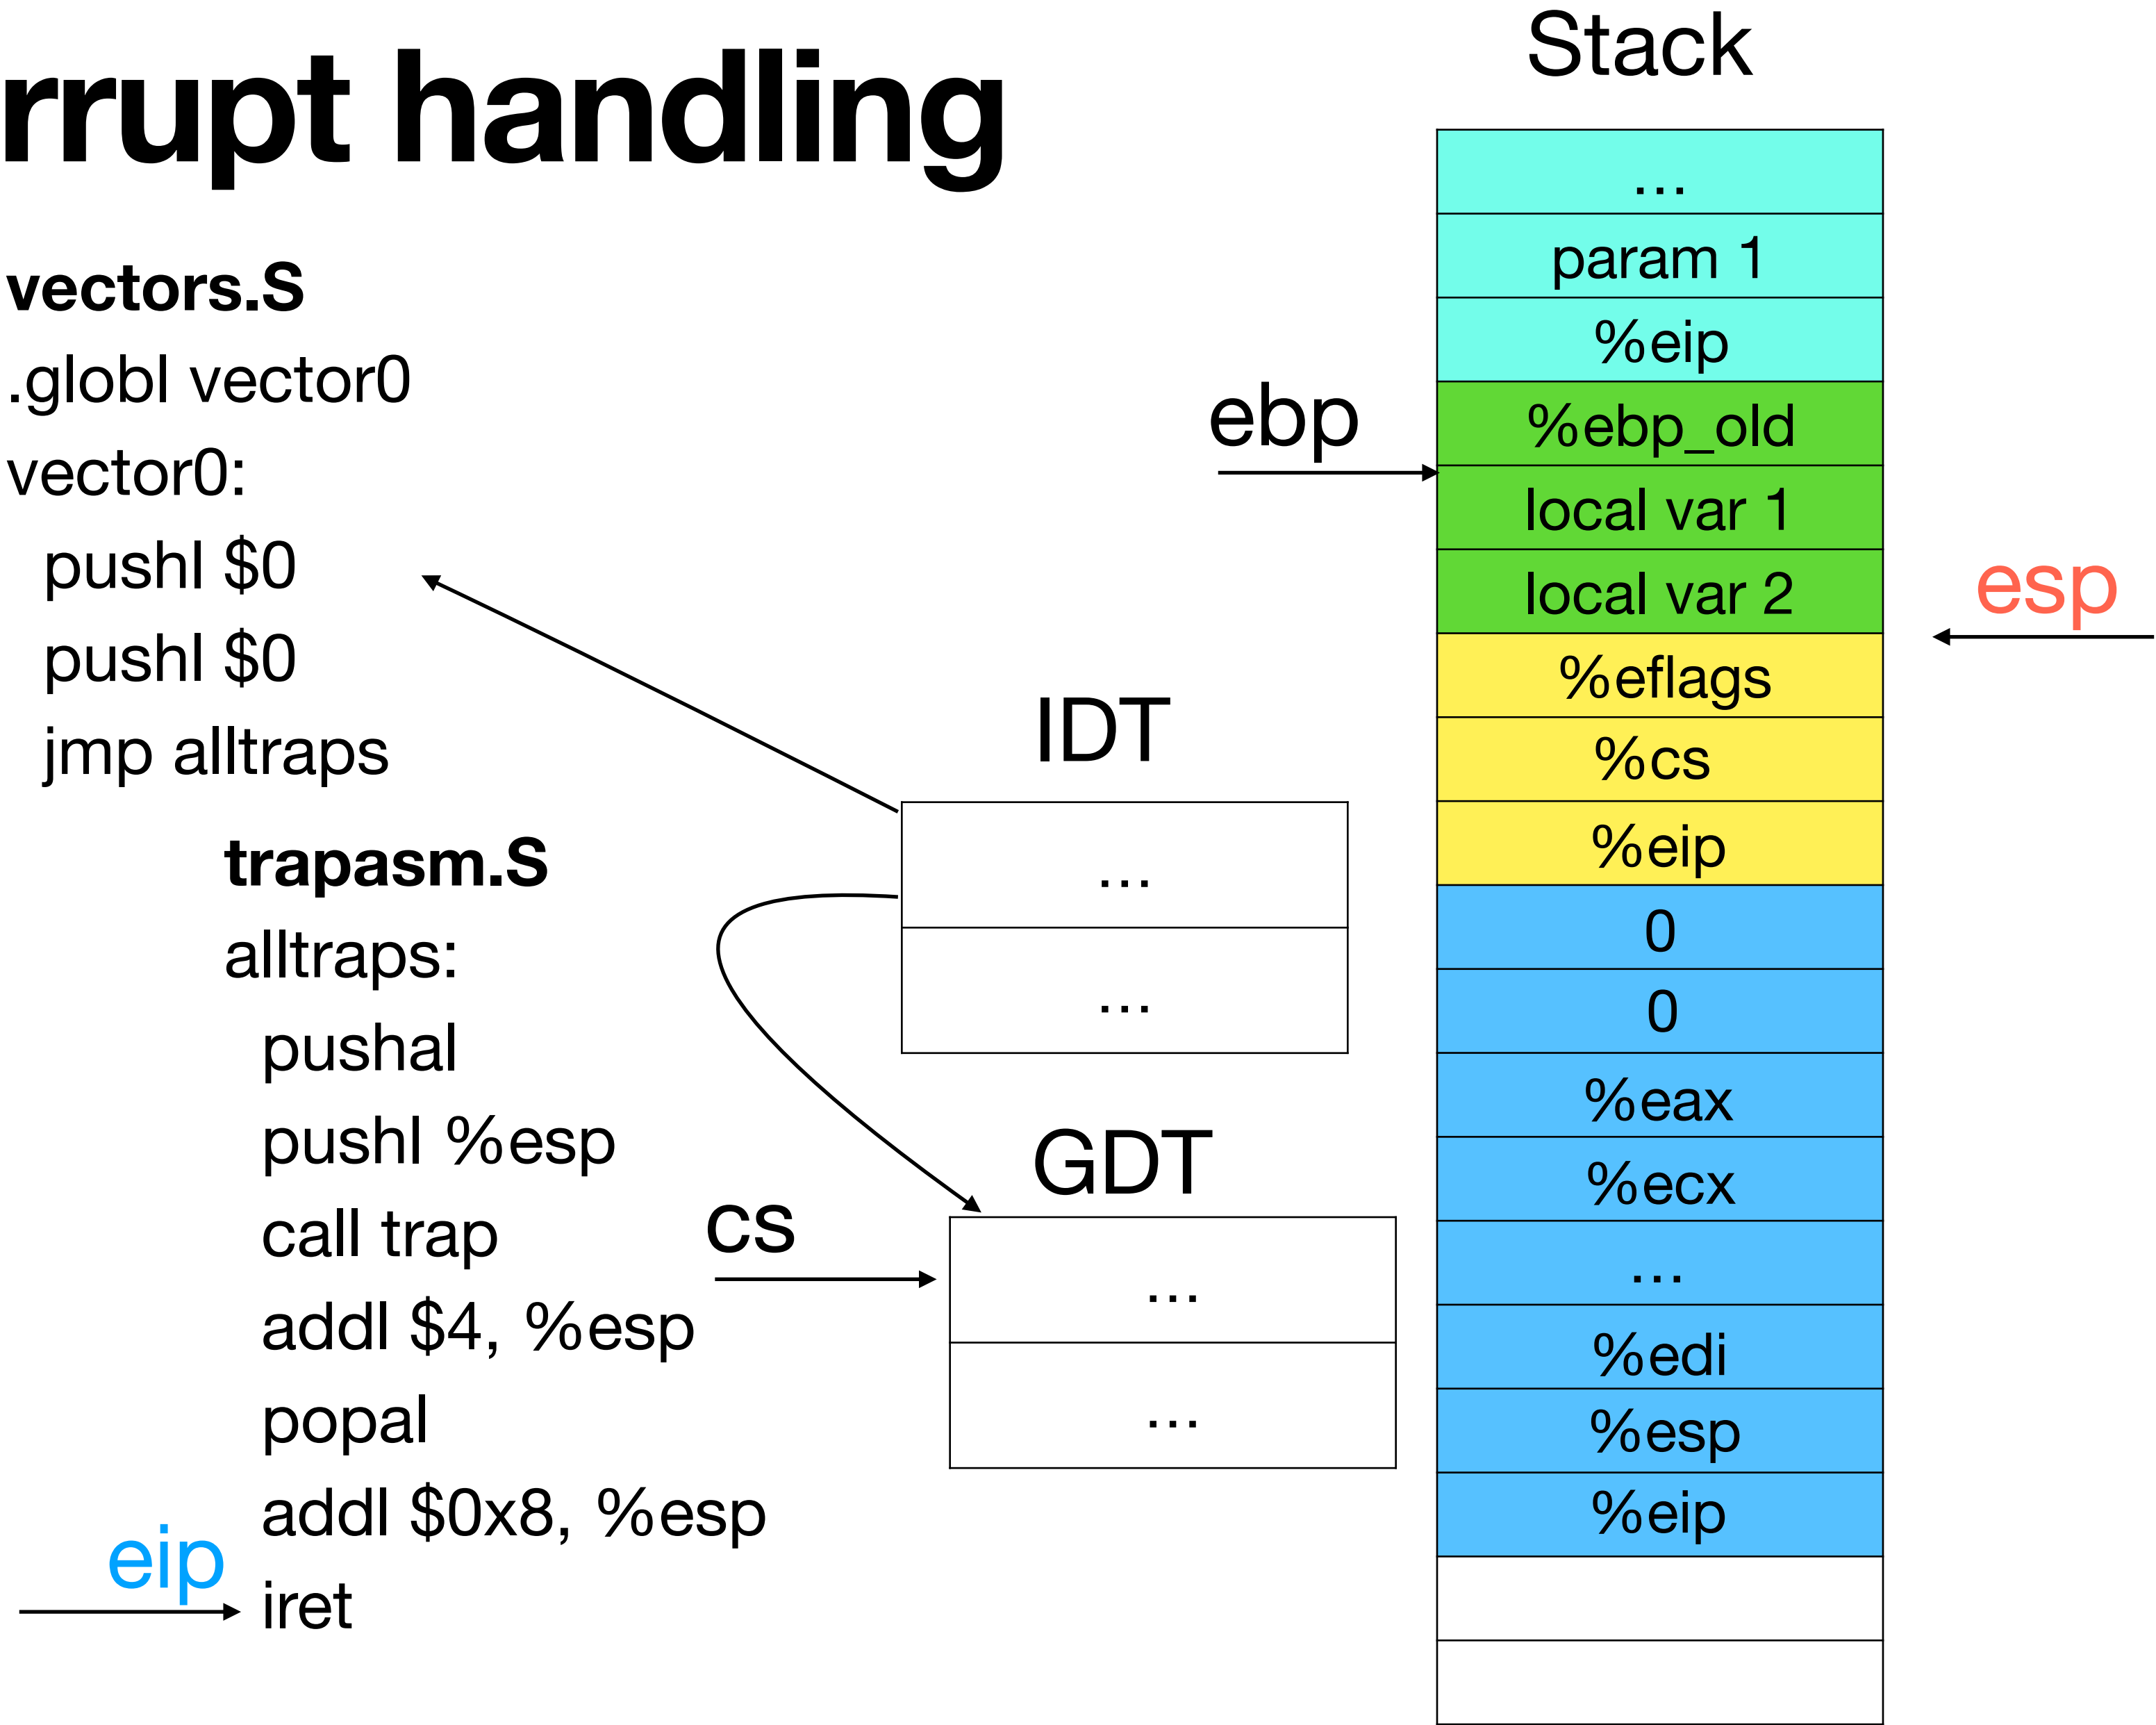
Visualizing interrupt handling

```
for(;;)
;

trap.c
void
trap(struct trapframe *tf)
{
    switch(tf->trapno){
    case T_IRQ0 + IRQ_TIMER:
        ticks++;
        cprintf("Tick! %d\n\0", ticks);
        lapiceoi();
    ..
    return
```

```
vectors.S
.globl vector0
vector0:
    pushl $0
    pushl $0
    jmp alltraps

trapasm.S
alltraps:
    pushal
    pushl %esp
    call trap
    addl $4, %esp
    popal
    addl $0x8, %esp
    iret
```



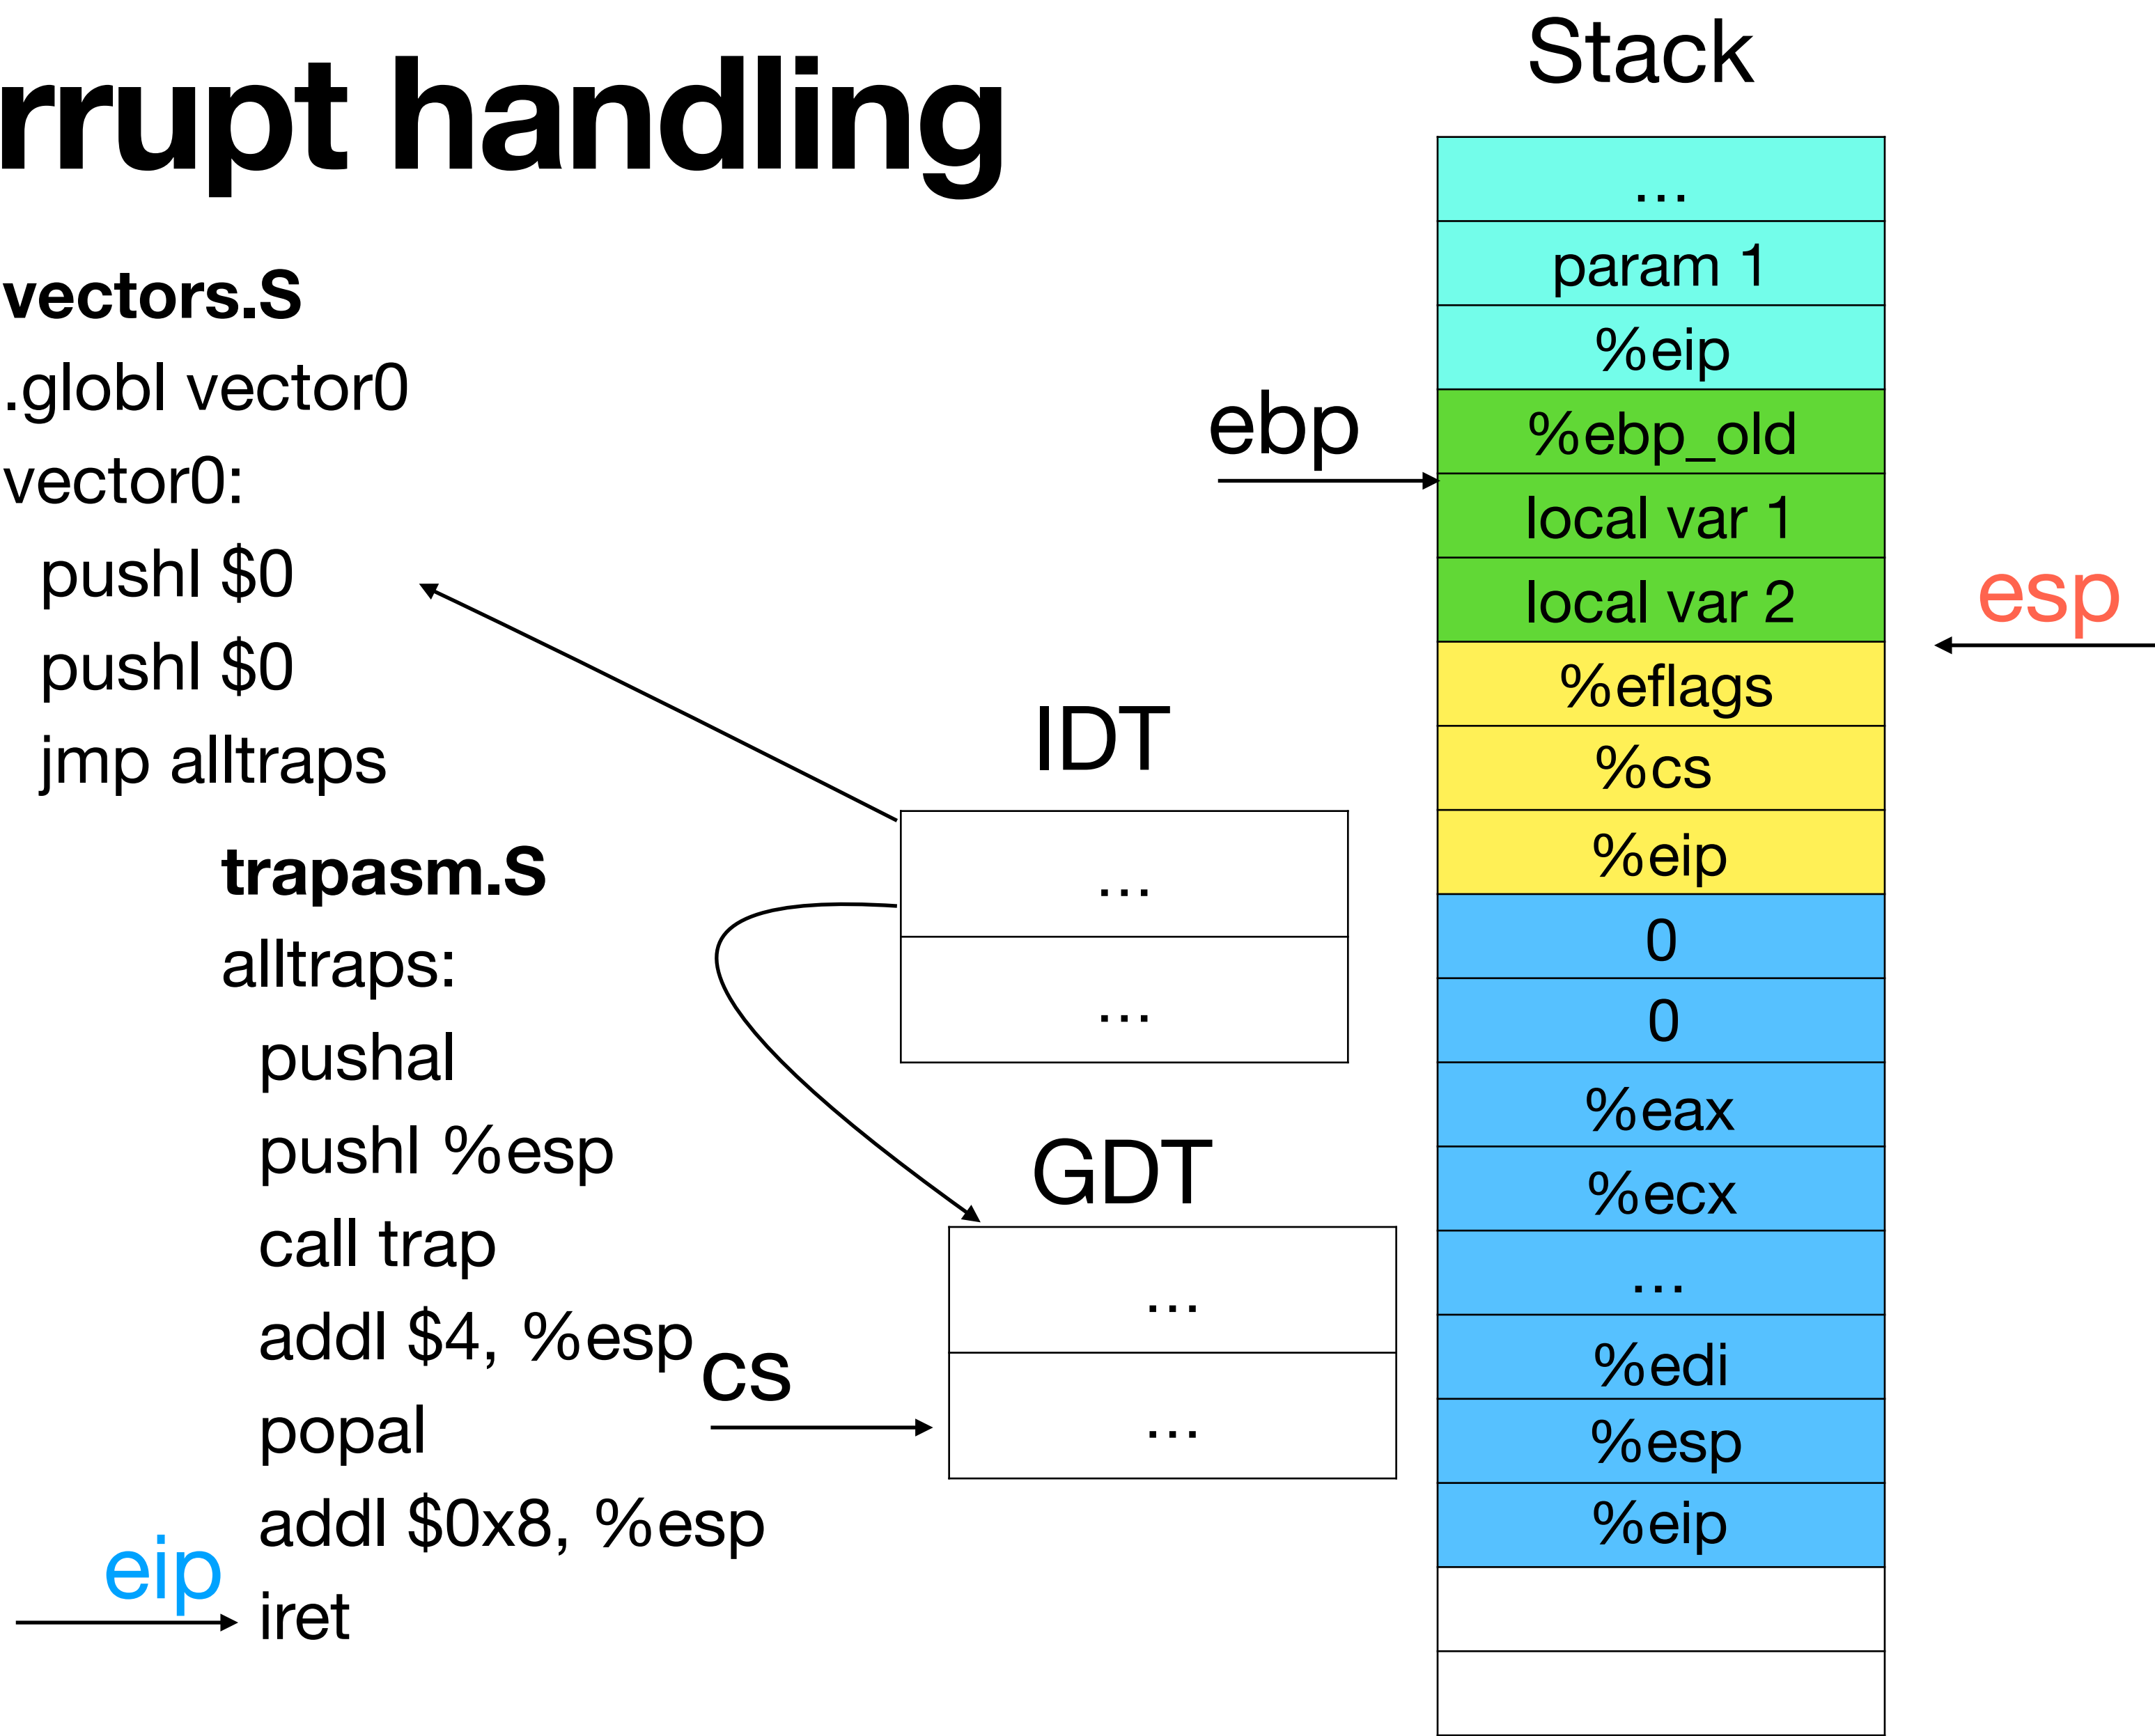
Visualizing interrupt handling

```
for(;;)
;

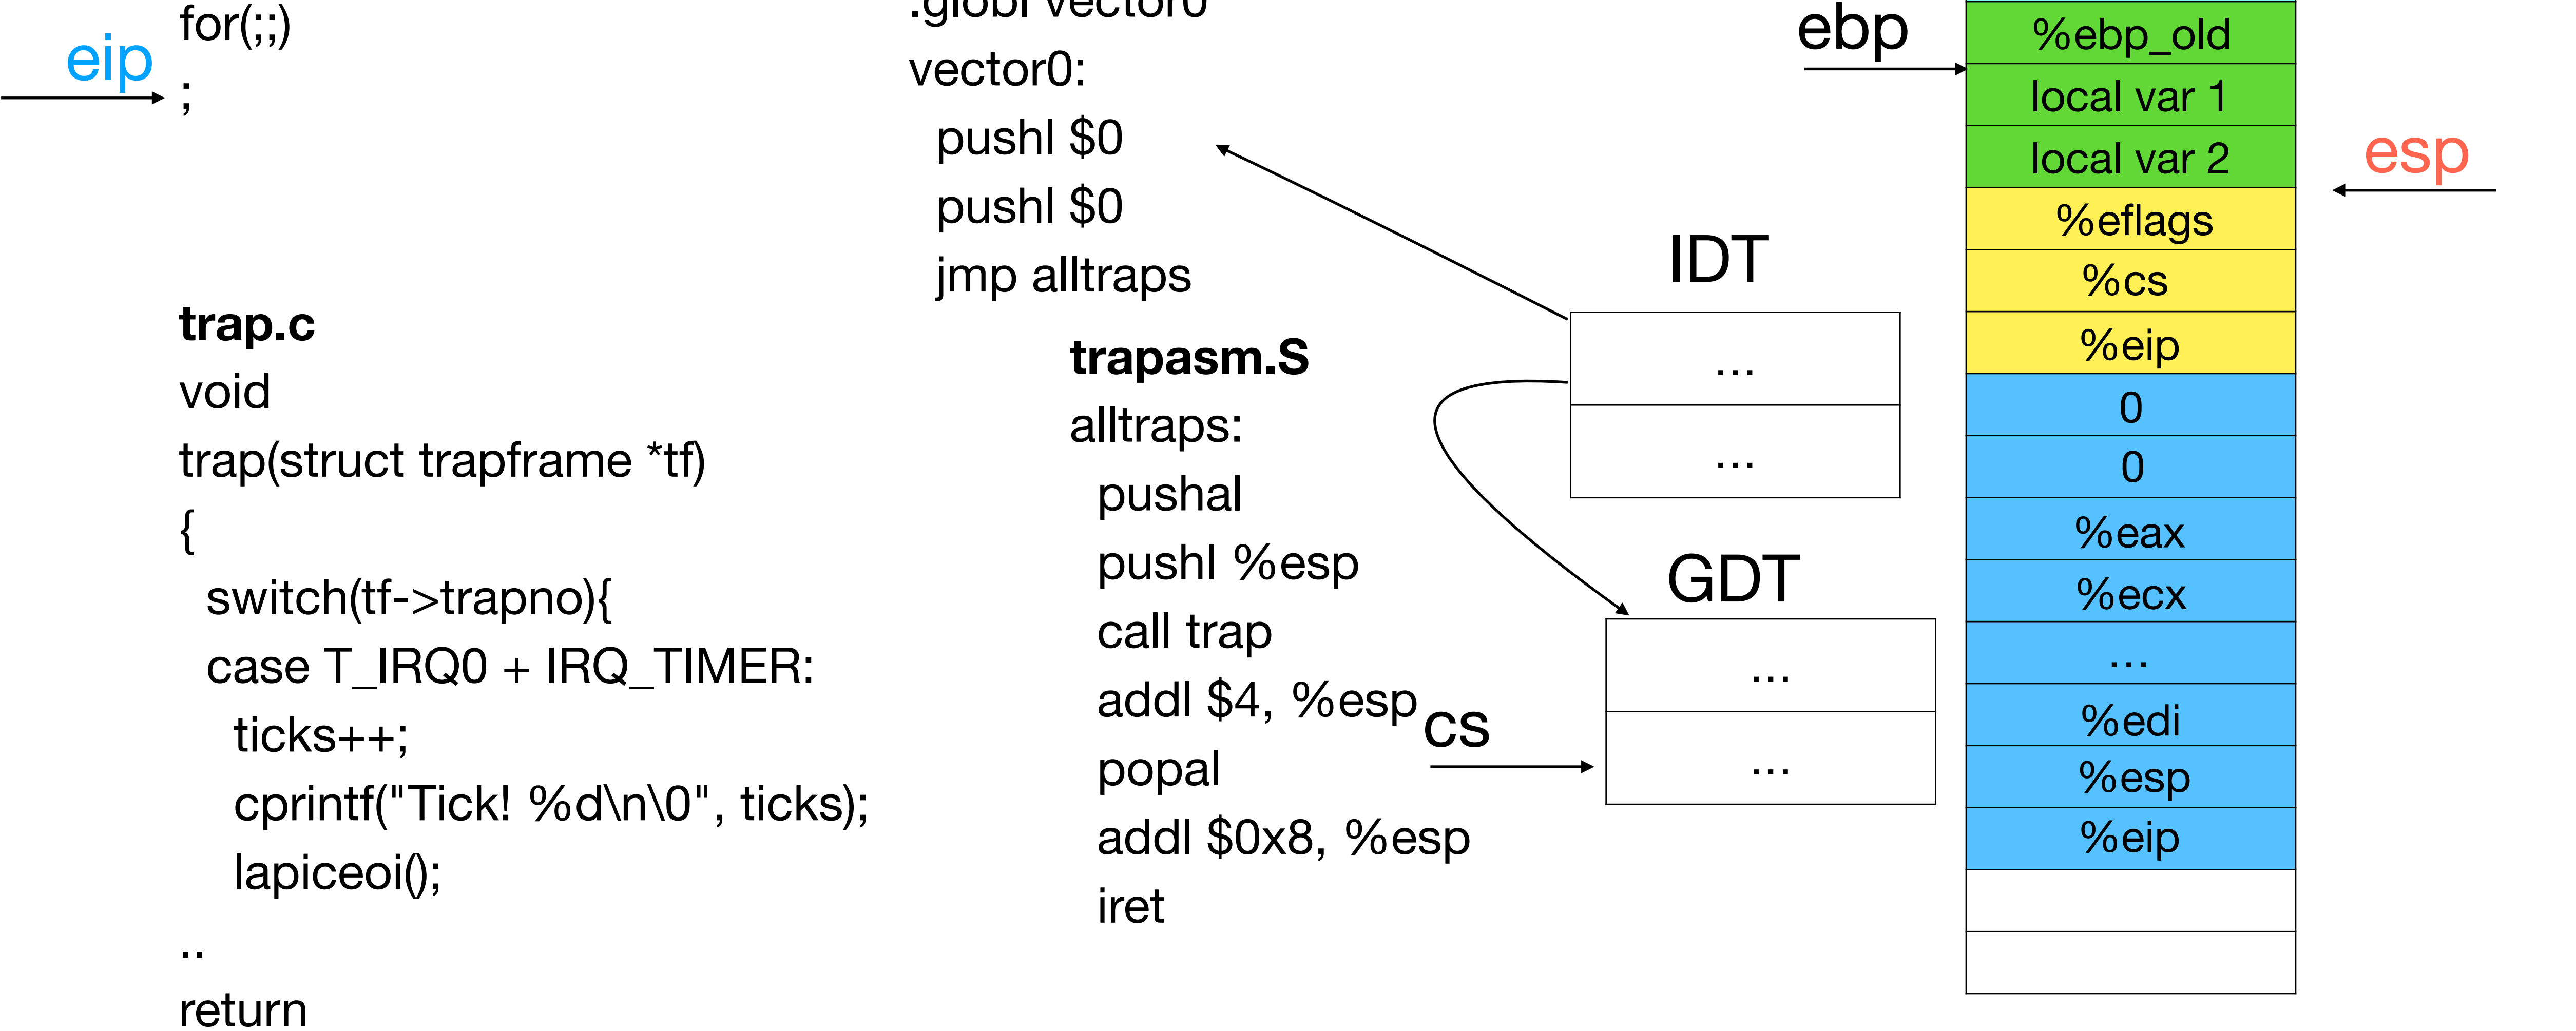
trap.c
void
trap(struct trapframe *tf)
{
    switch(tf->trapno){
    case T_IRQ0 + IRQ_TIMER:
        ticks++;
        cprintf("Tick! %d\n\0", ticks);
        lapiceoi();
    ..
    return
```

```
vectors.S
.globl vector0
vector0:
    pushl $0
    pushl $0
    jmp alltraps

trapasm.S
alltraps:
    pushal
    pushl %esp
    call trap
    addl $4, %esp
    popal
    addl $0x8, %esp
    iret
```



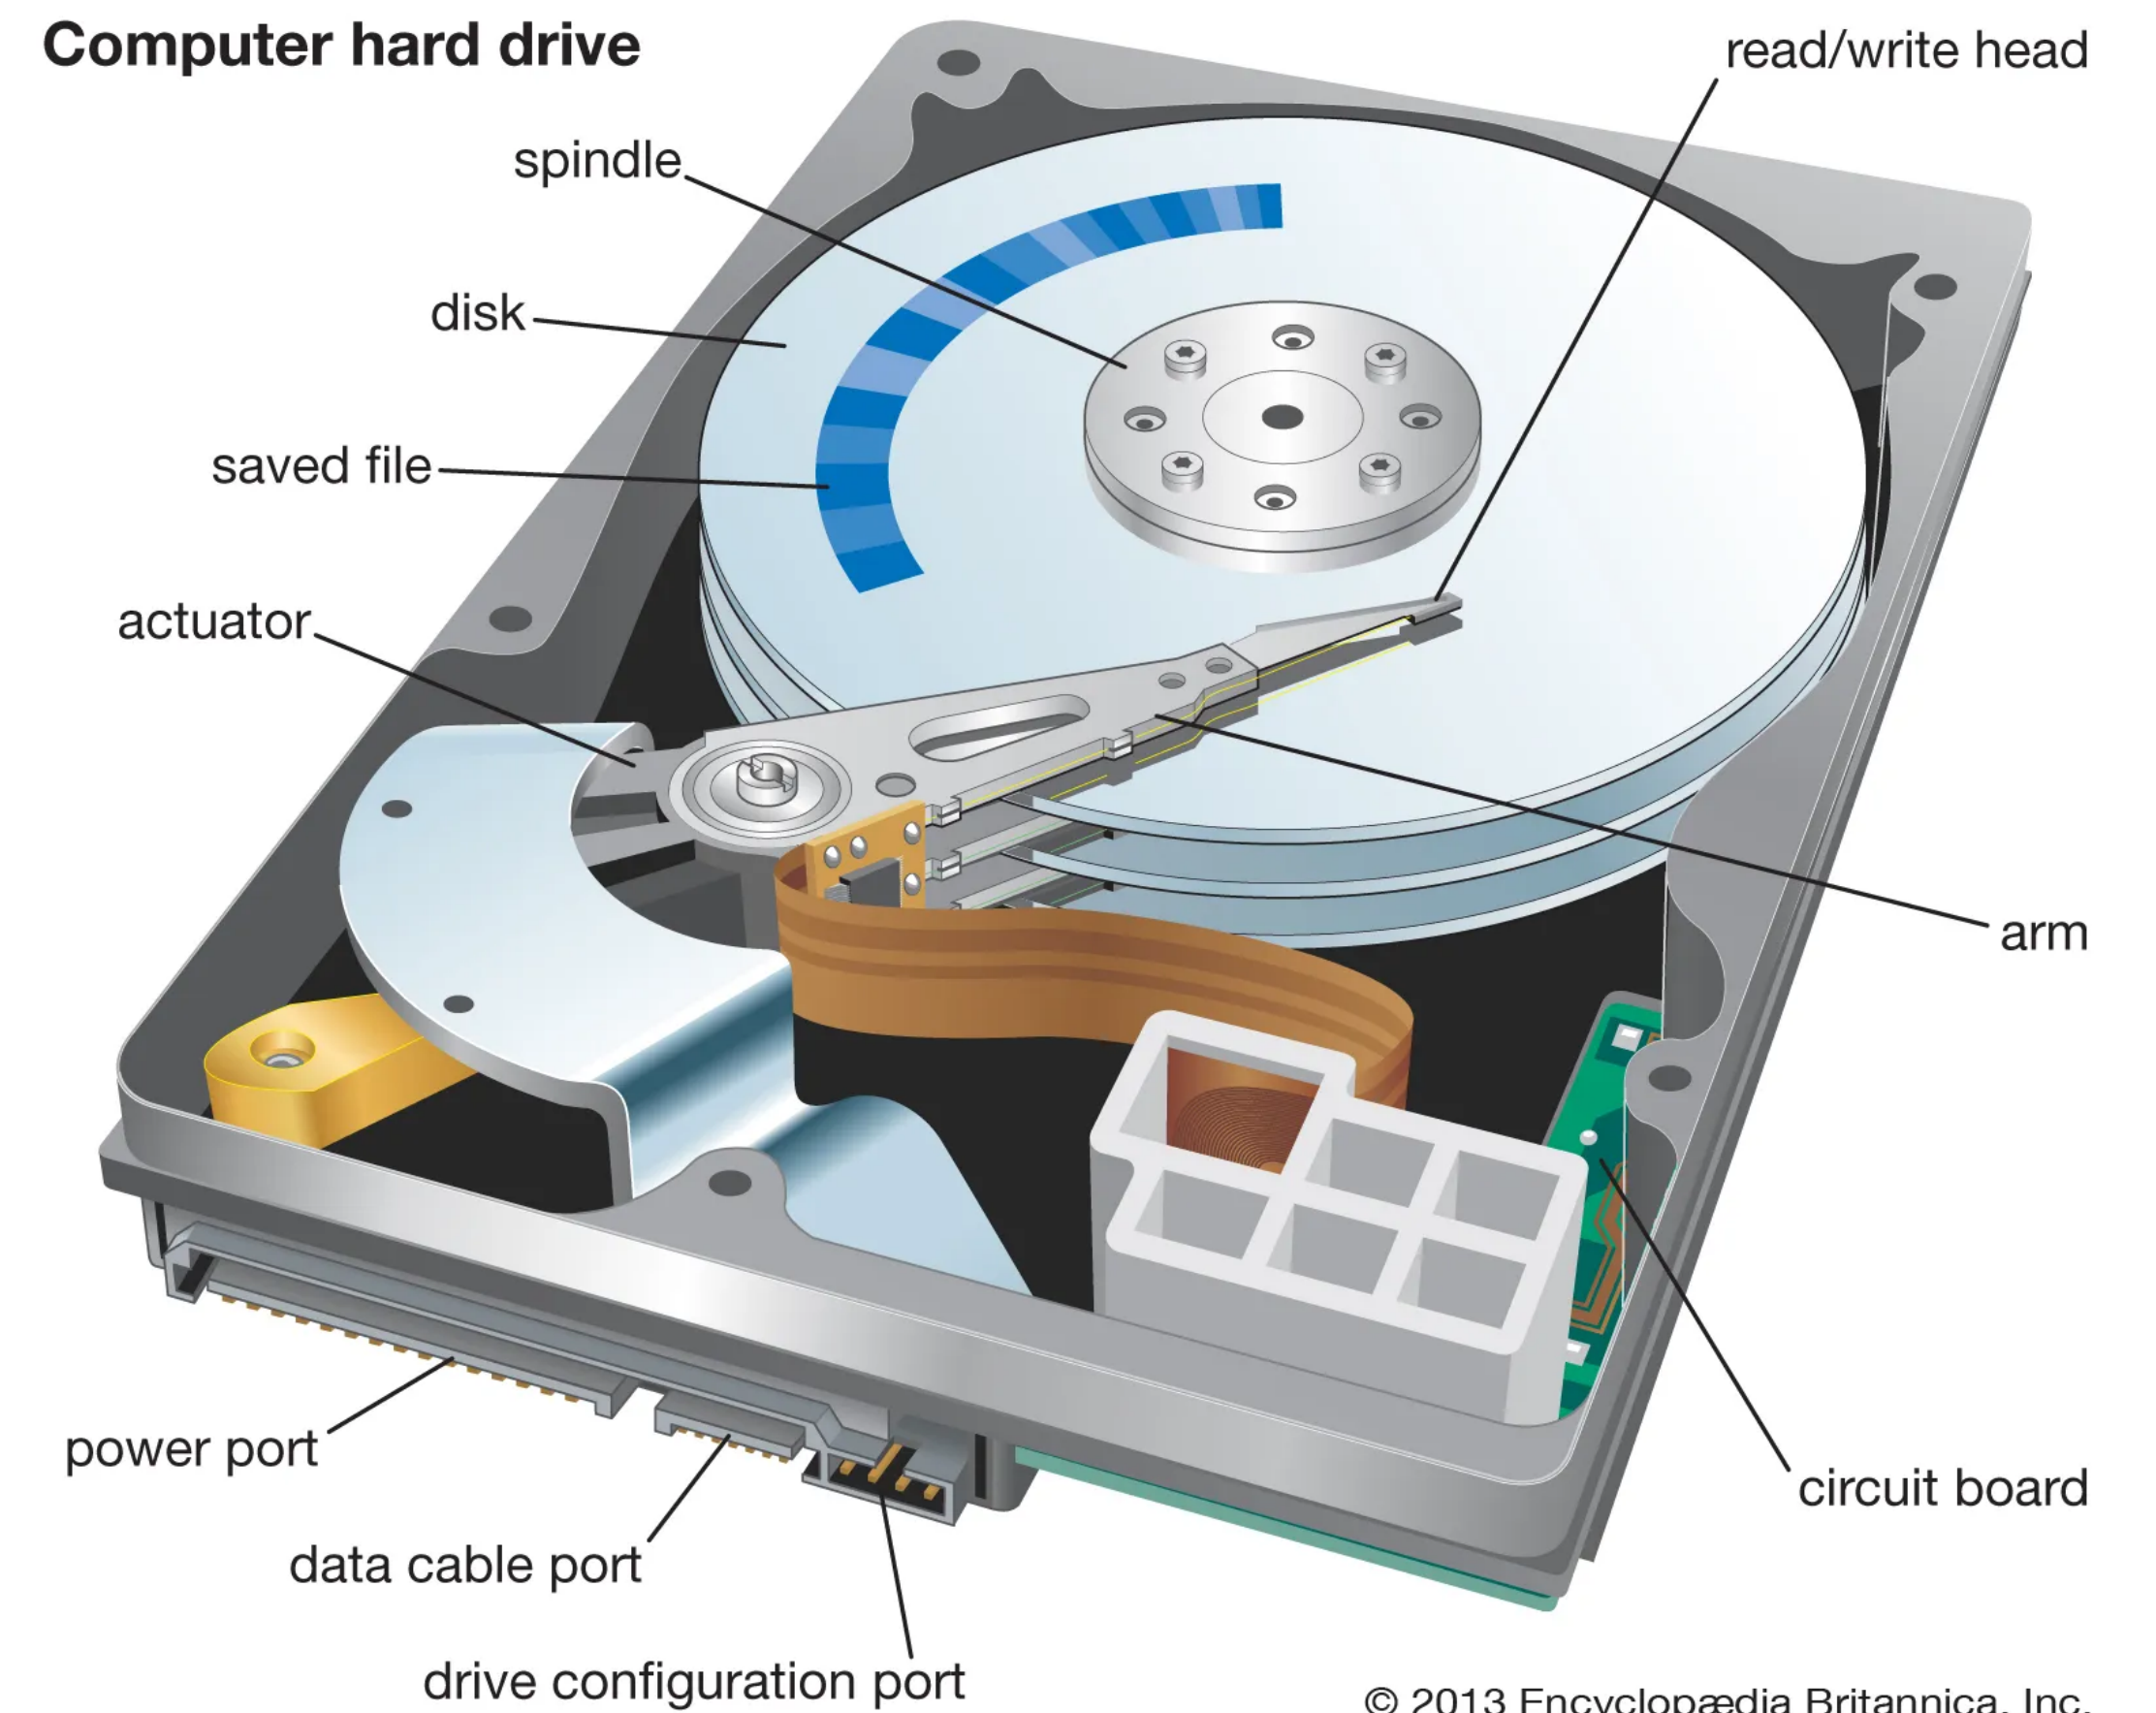
Visualizing interrupt handling



Hard disk drive

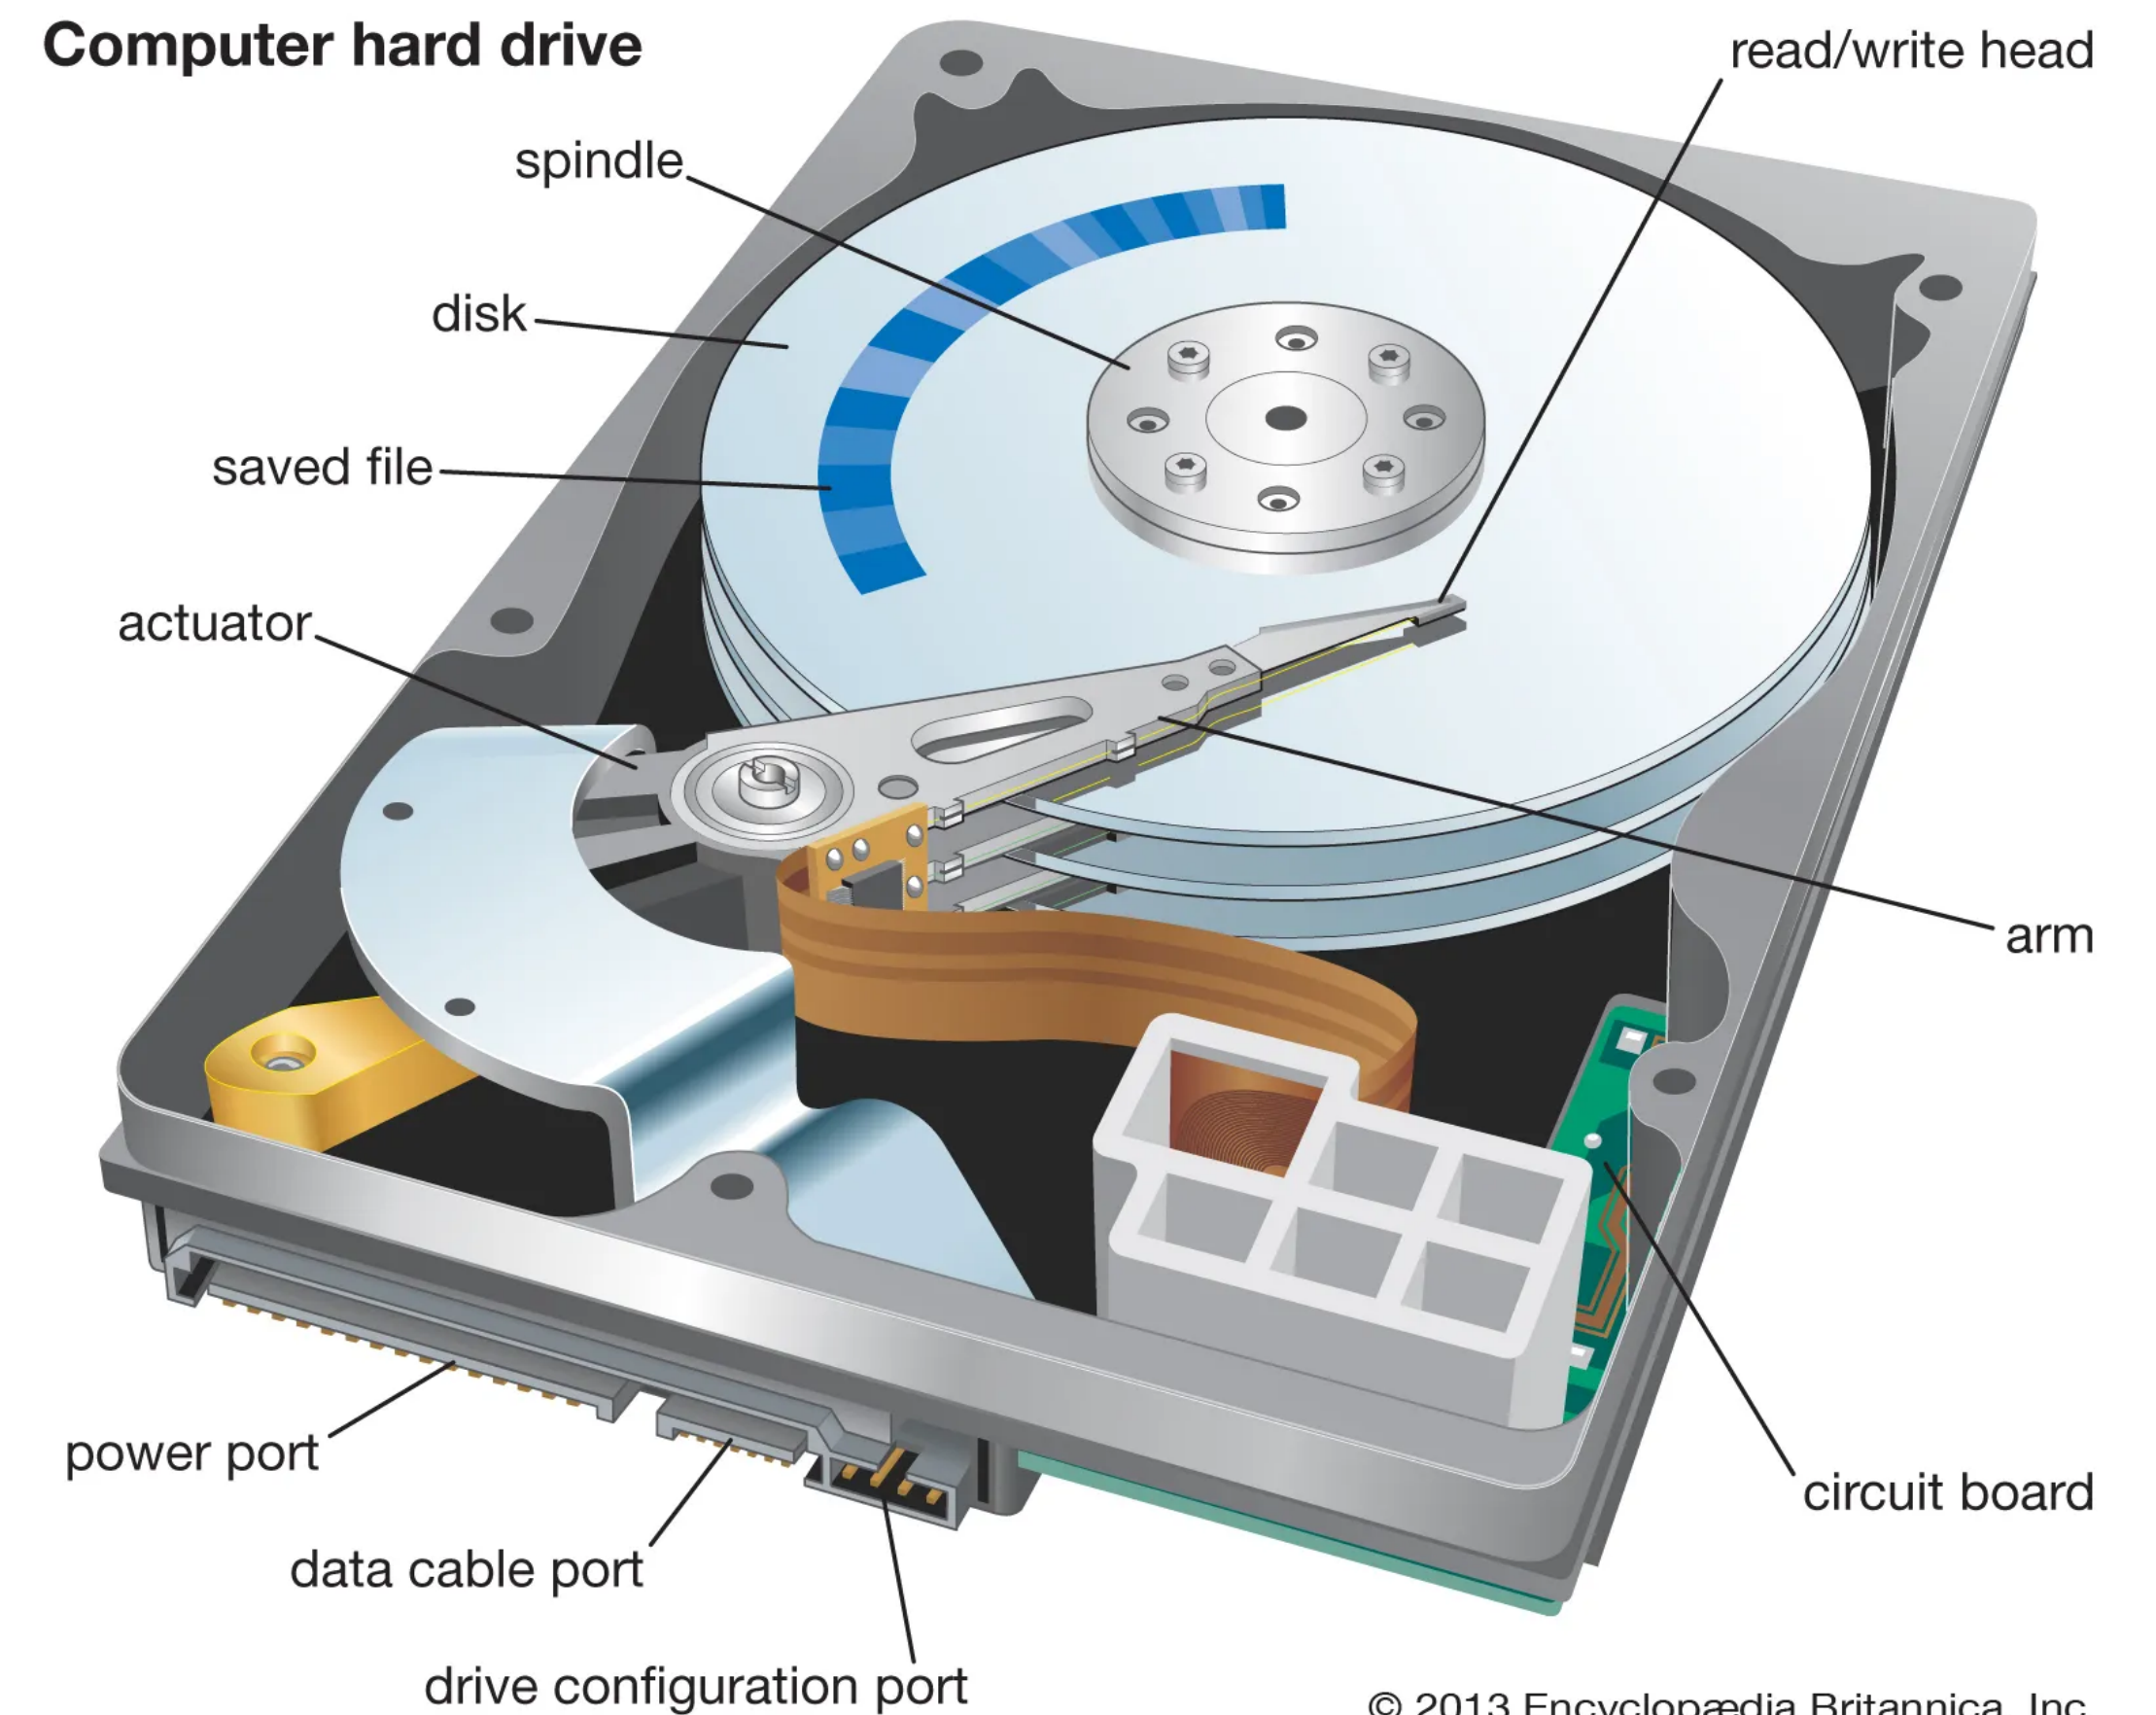
Ch. 37 OSTEP book

Disk geometry



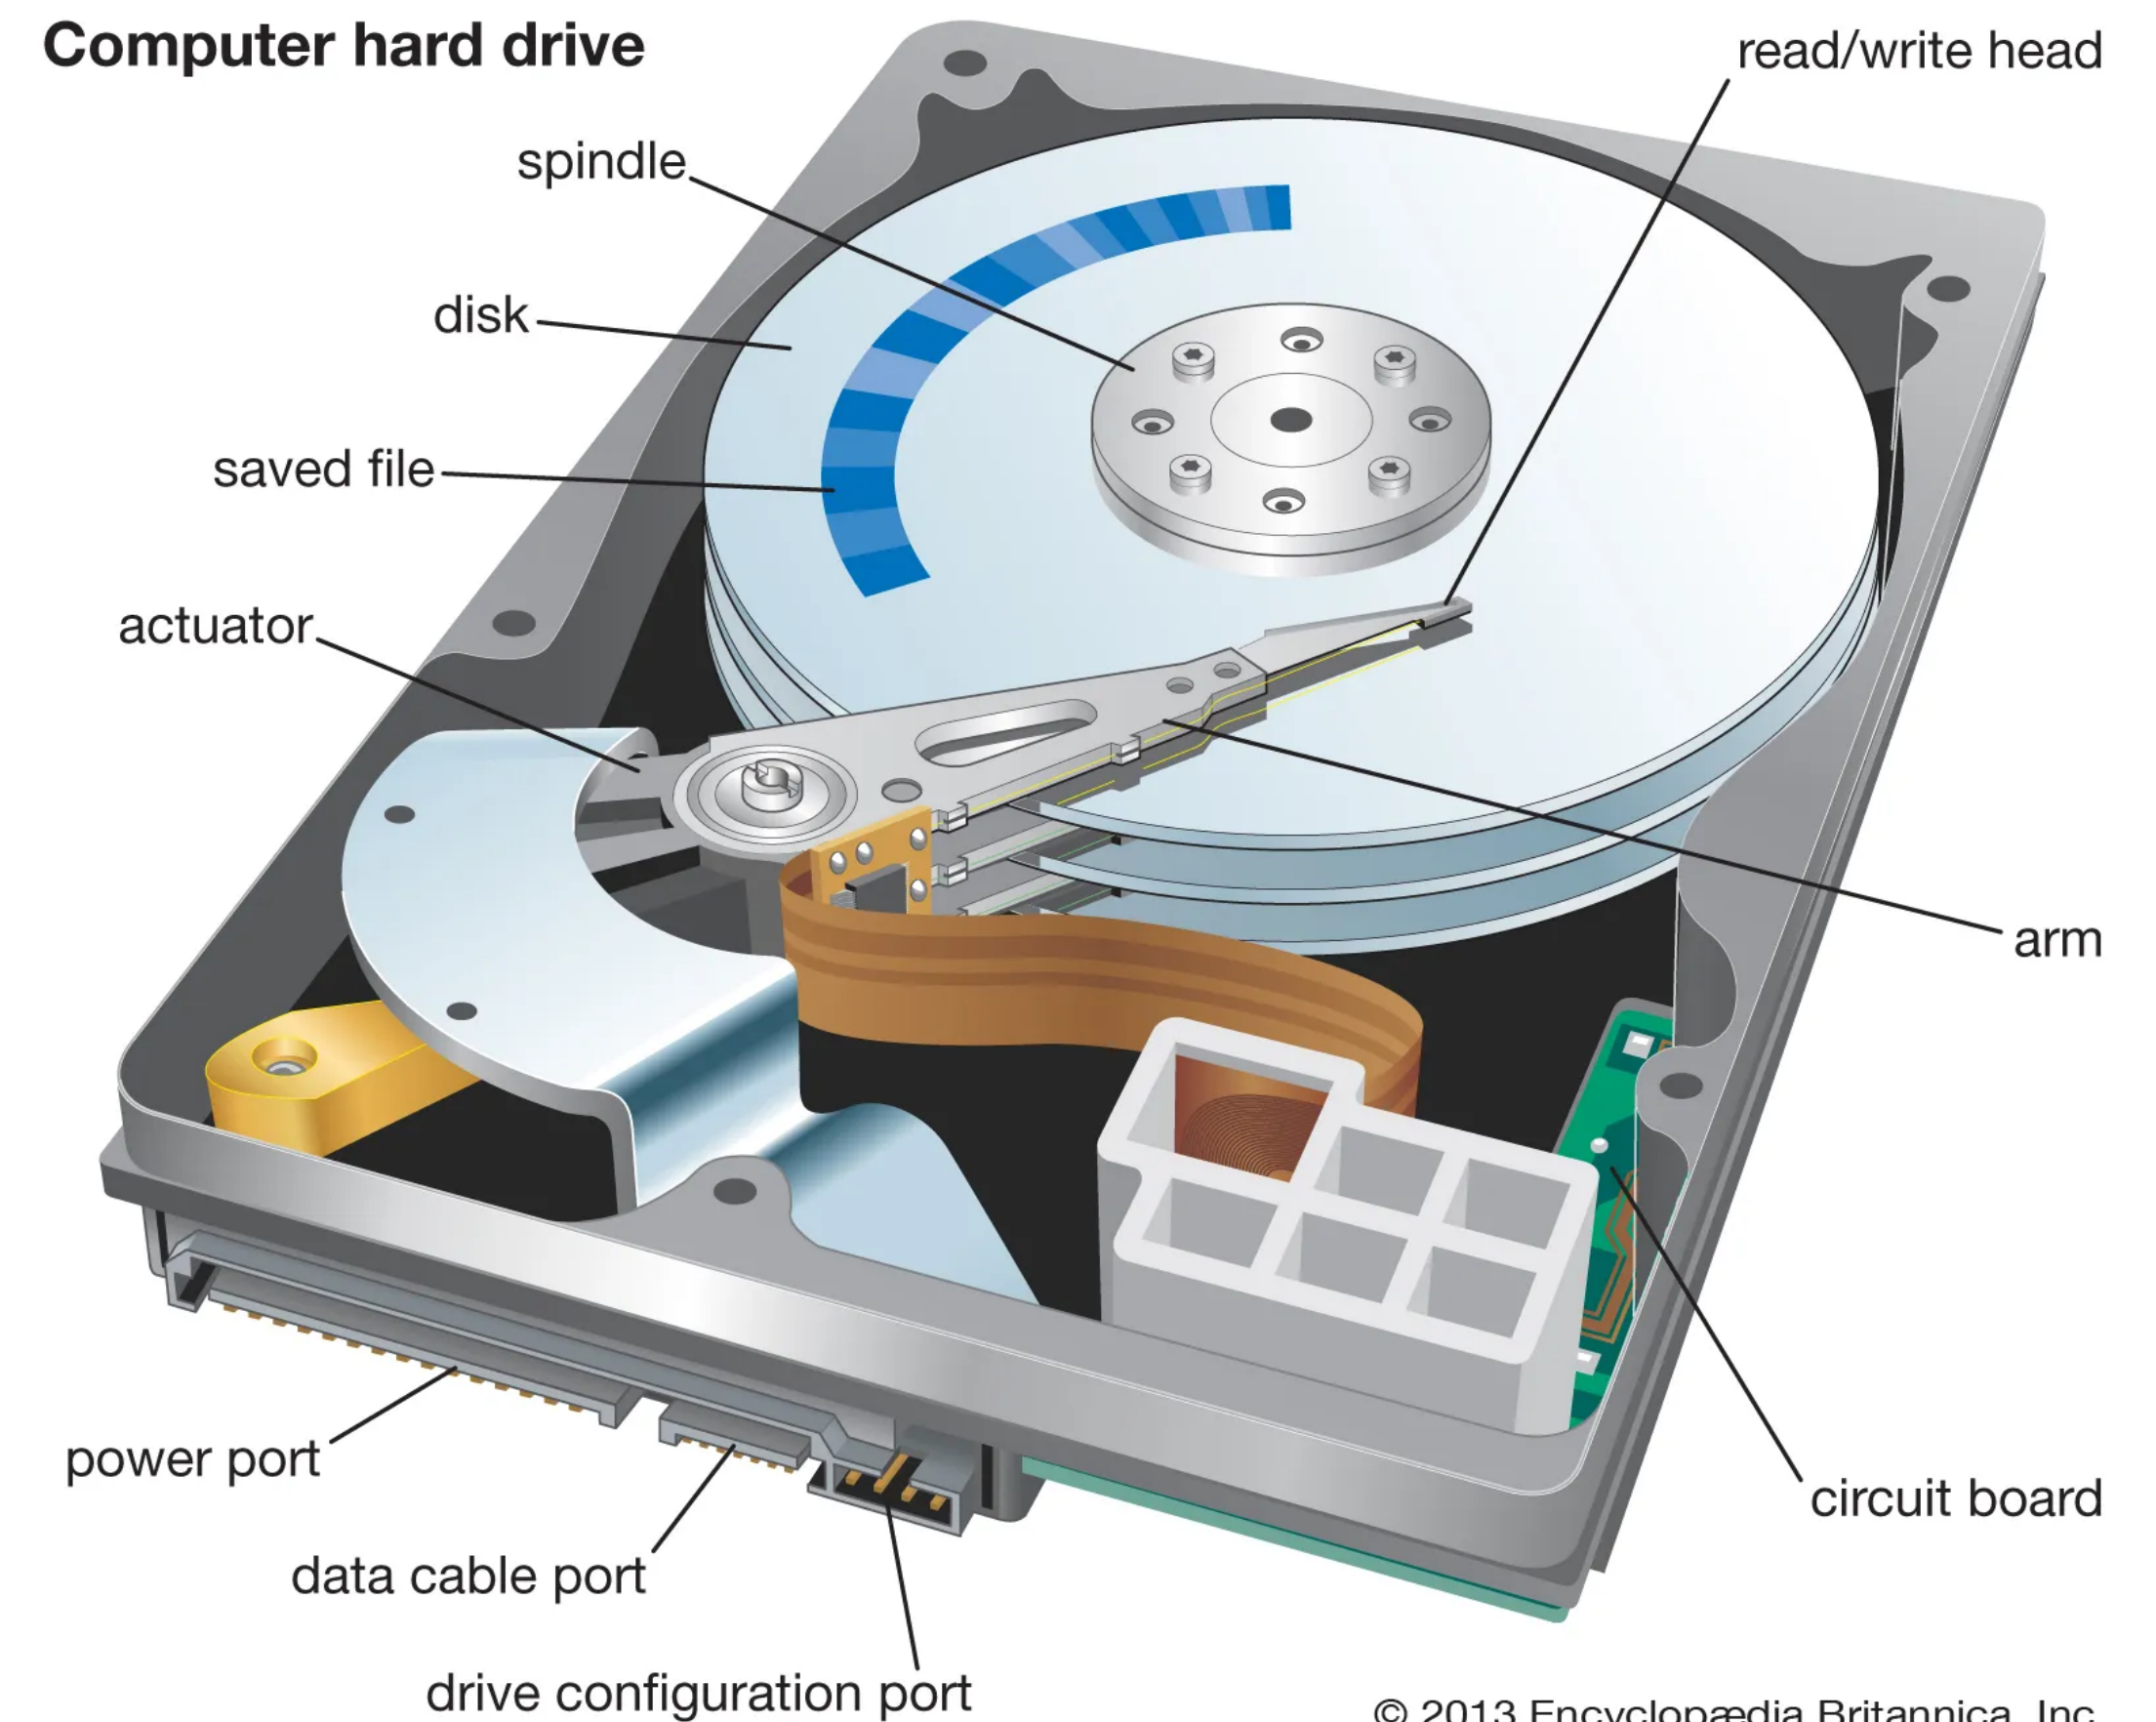
Disk geometry

- Many platters spinning on a spindle (~10,000 RPM)



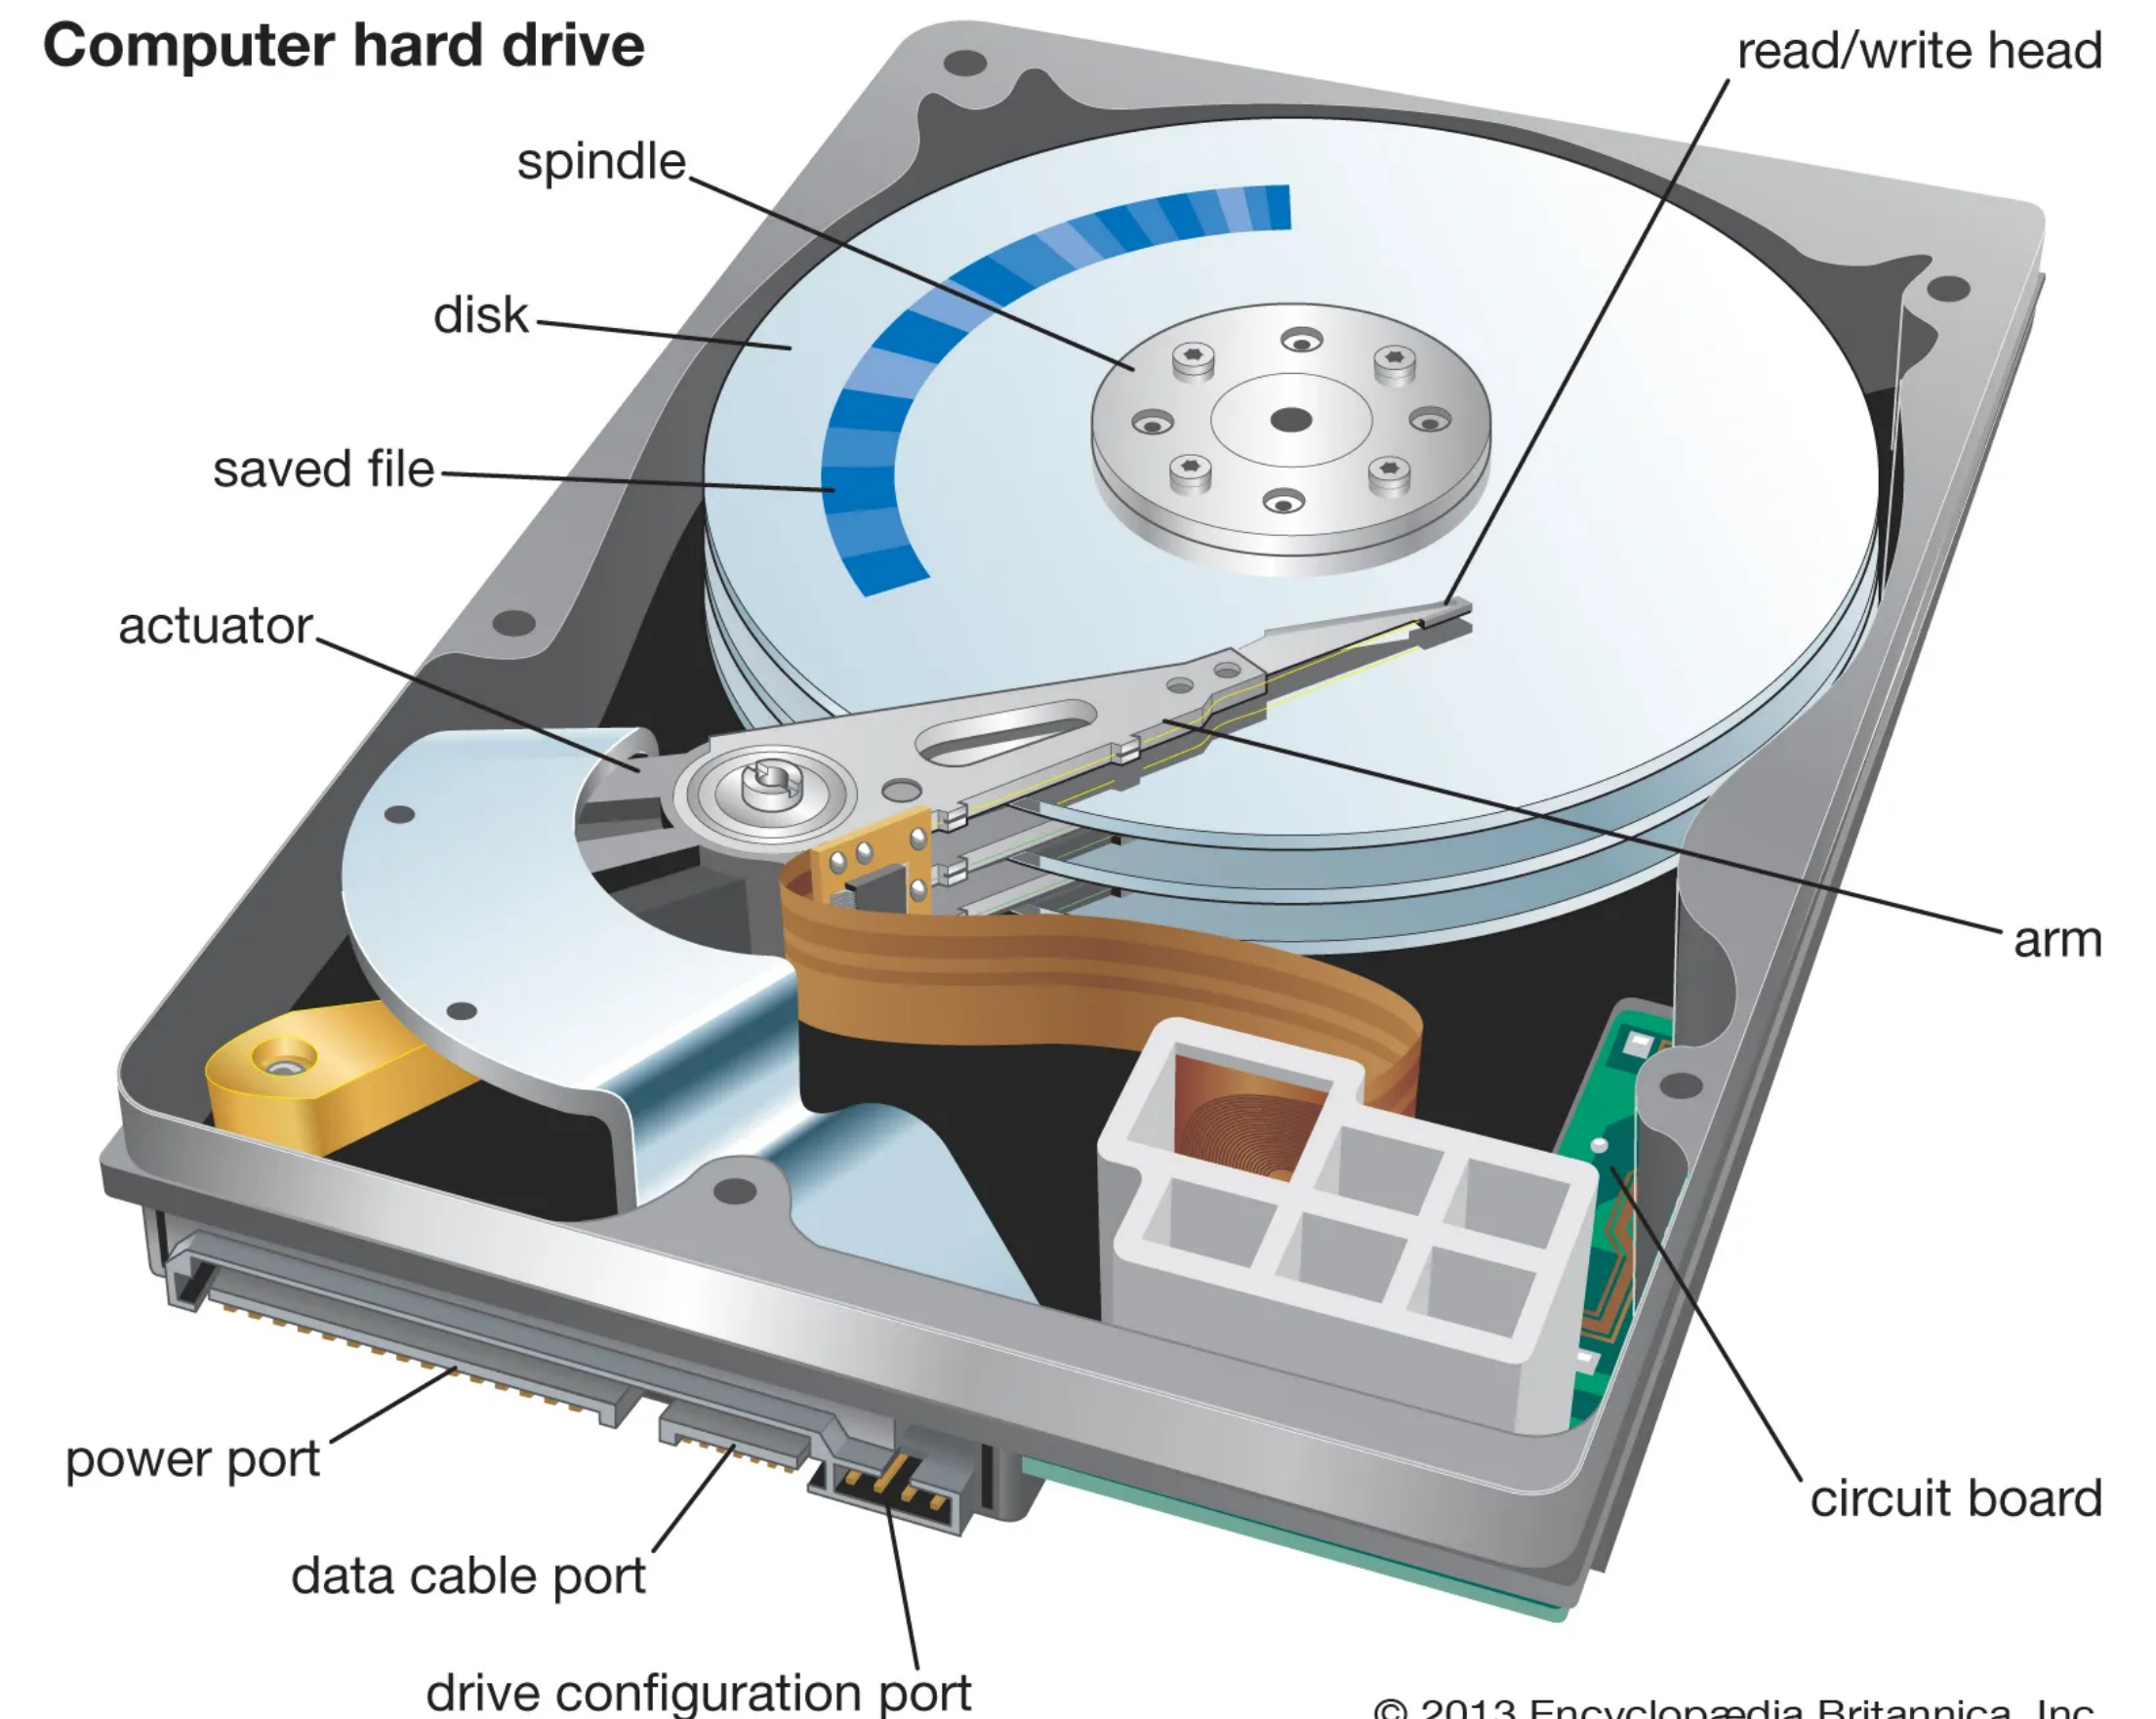
Disk geometry

- Many platters spinning on a spindle (~10,000 RPM)
- Each platter has two disk heads, one for each surface



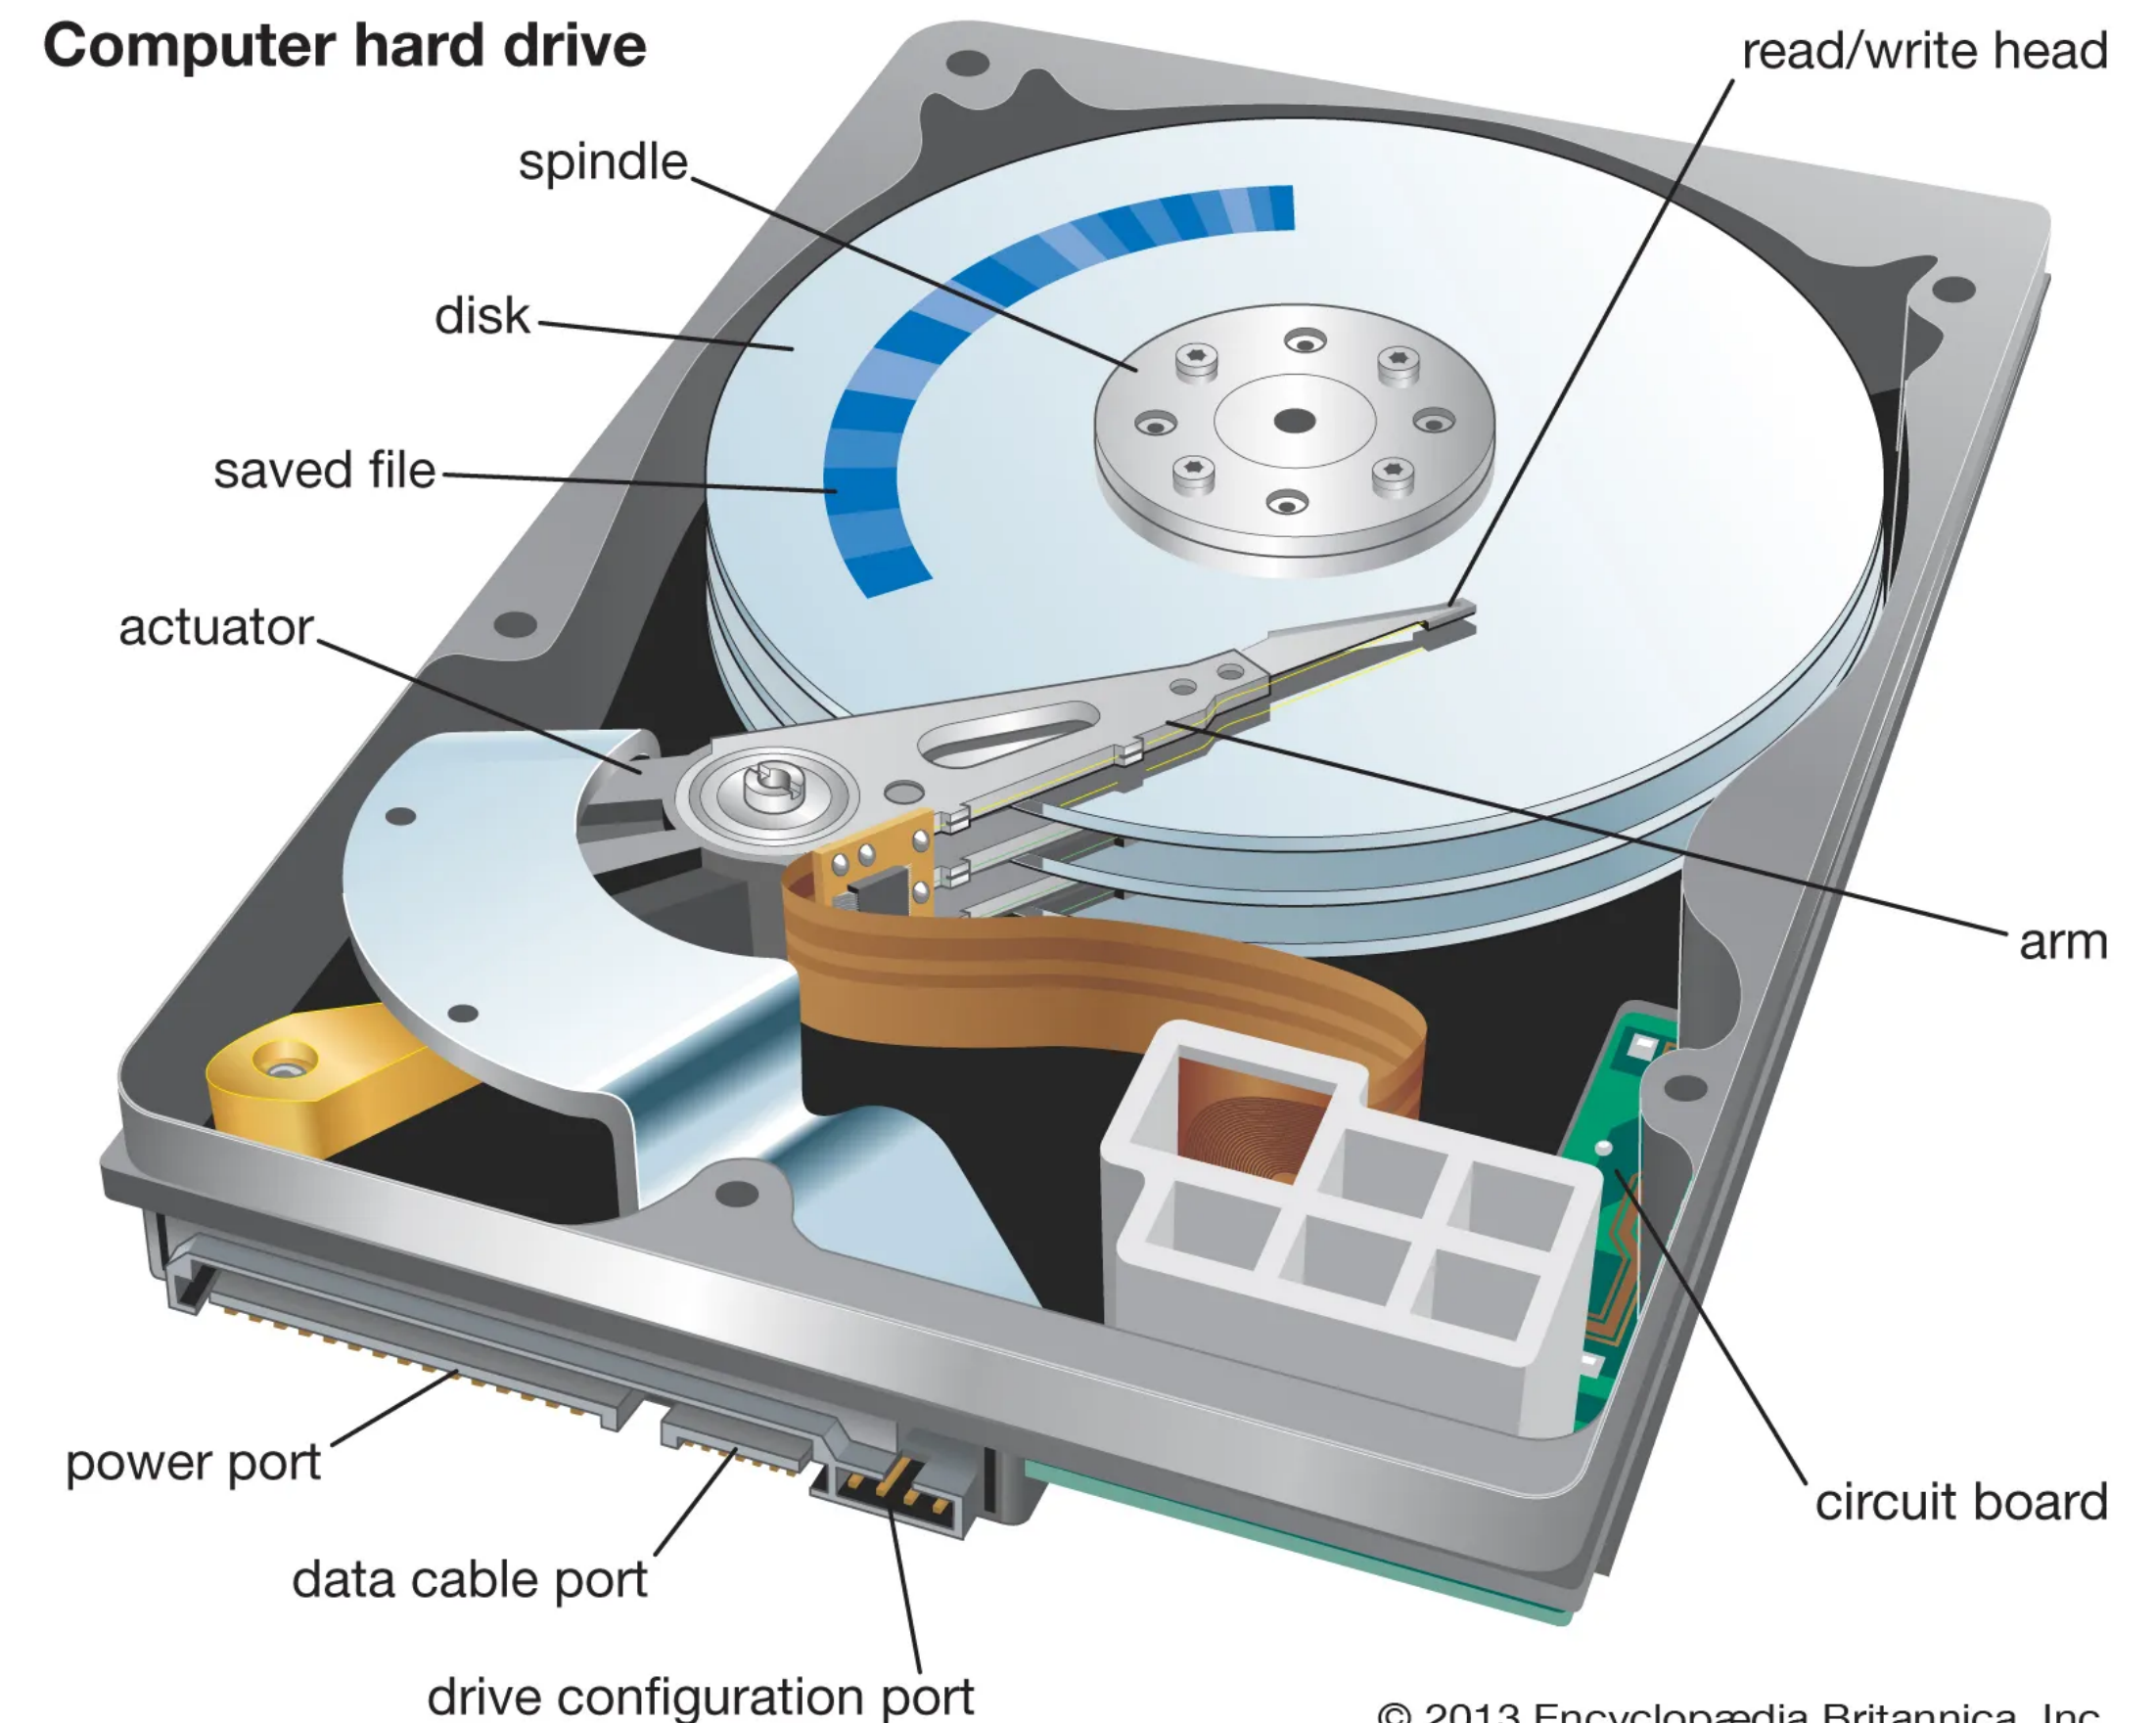
Disk geometry

- Many platters spinning on a spindle (~10,000 RPM)
- Each platter has two disk heads, one for each surface
- Disk heads are controlled by actuator



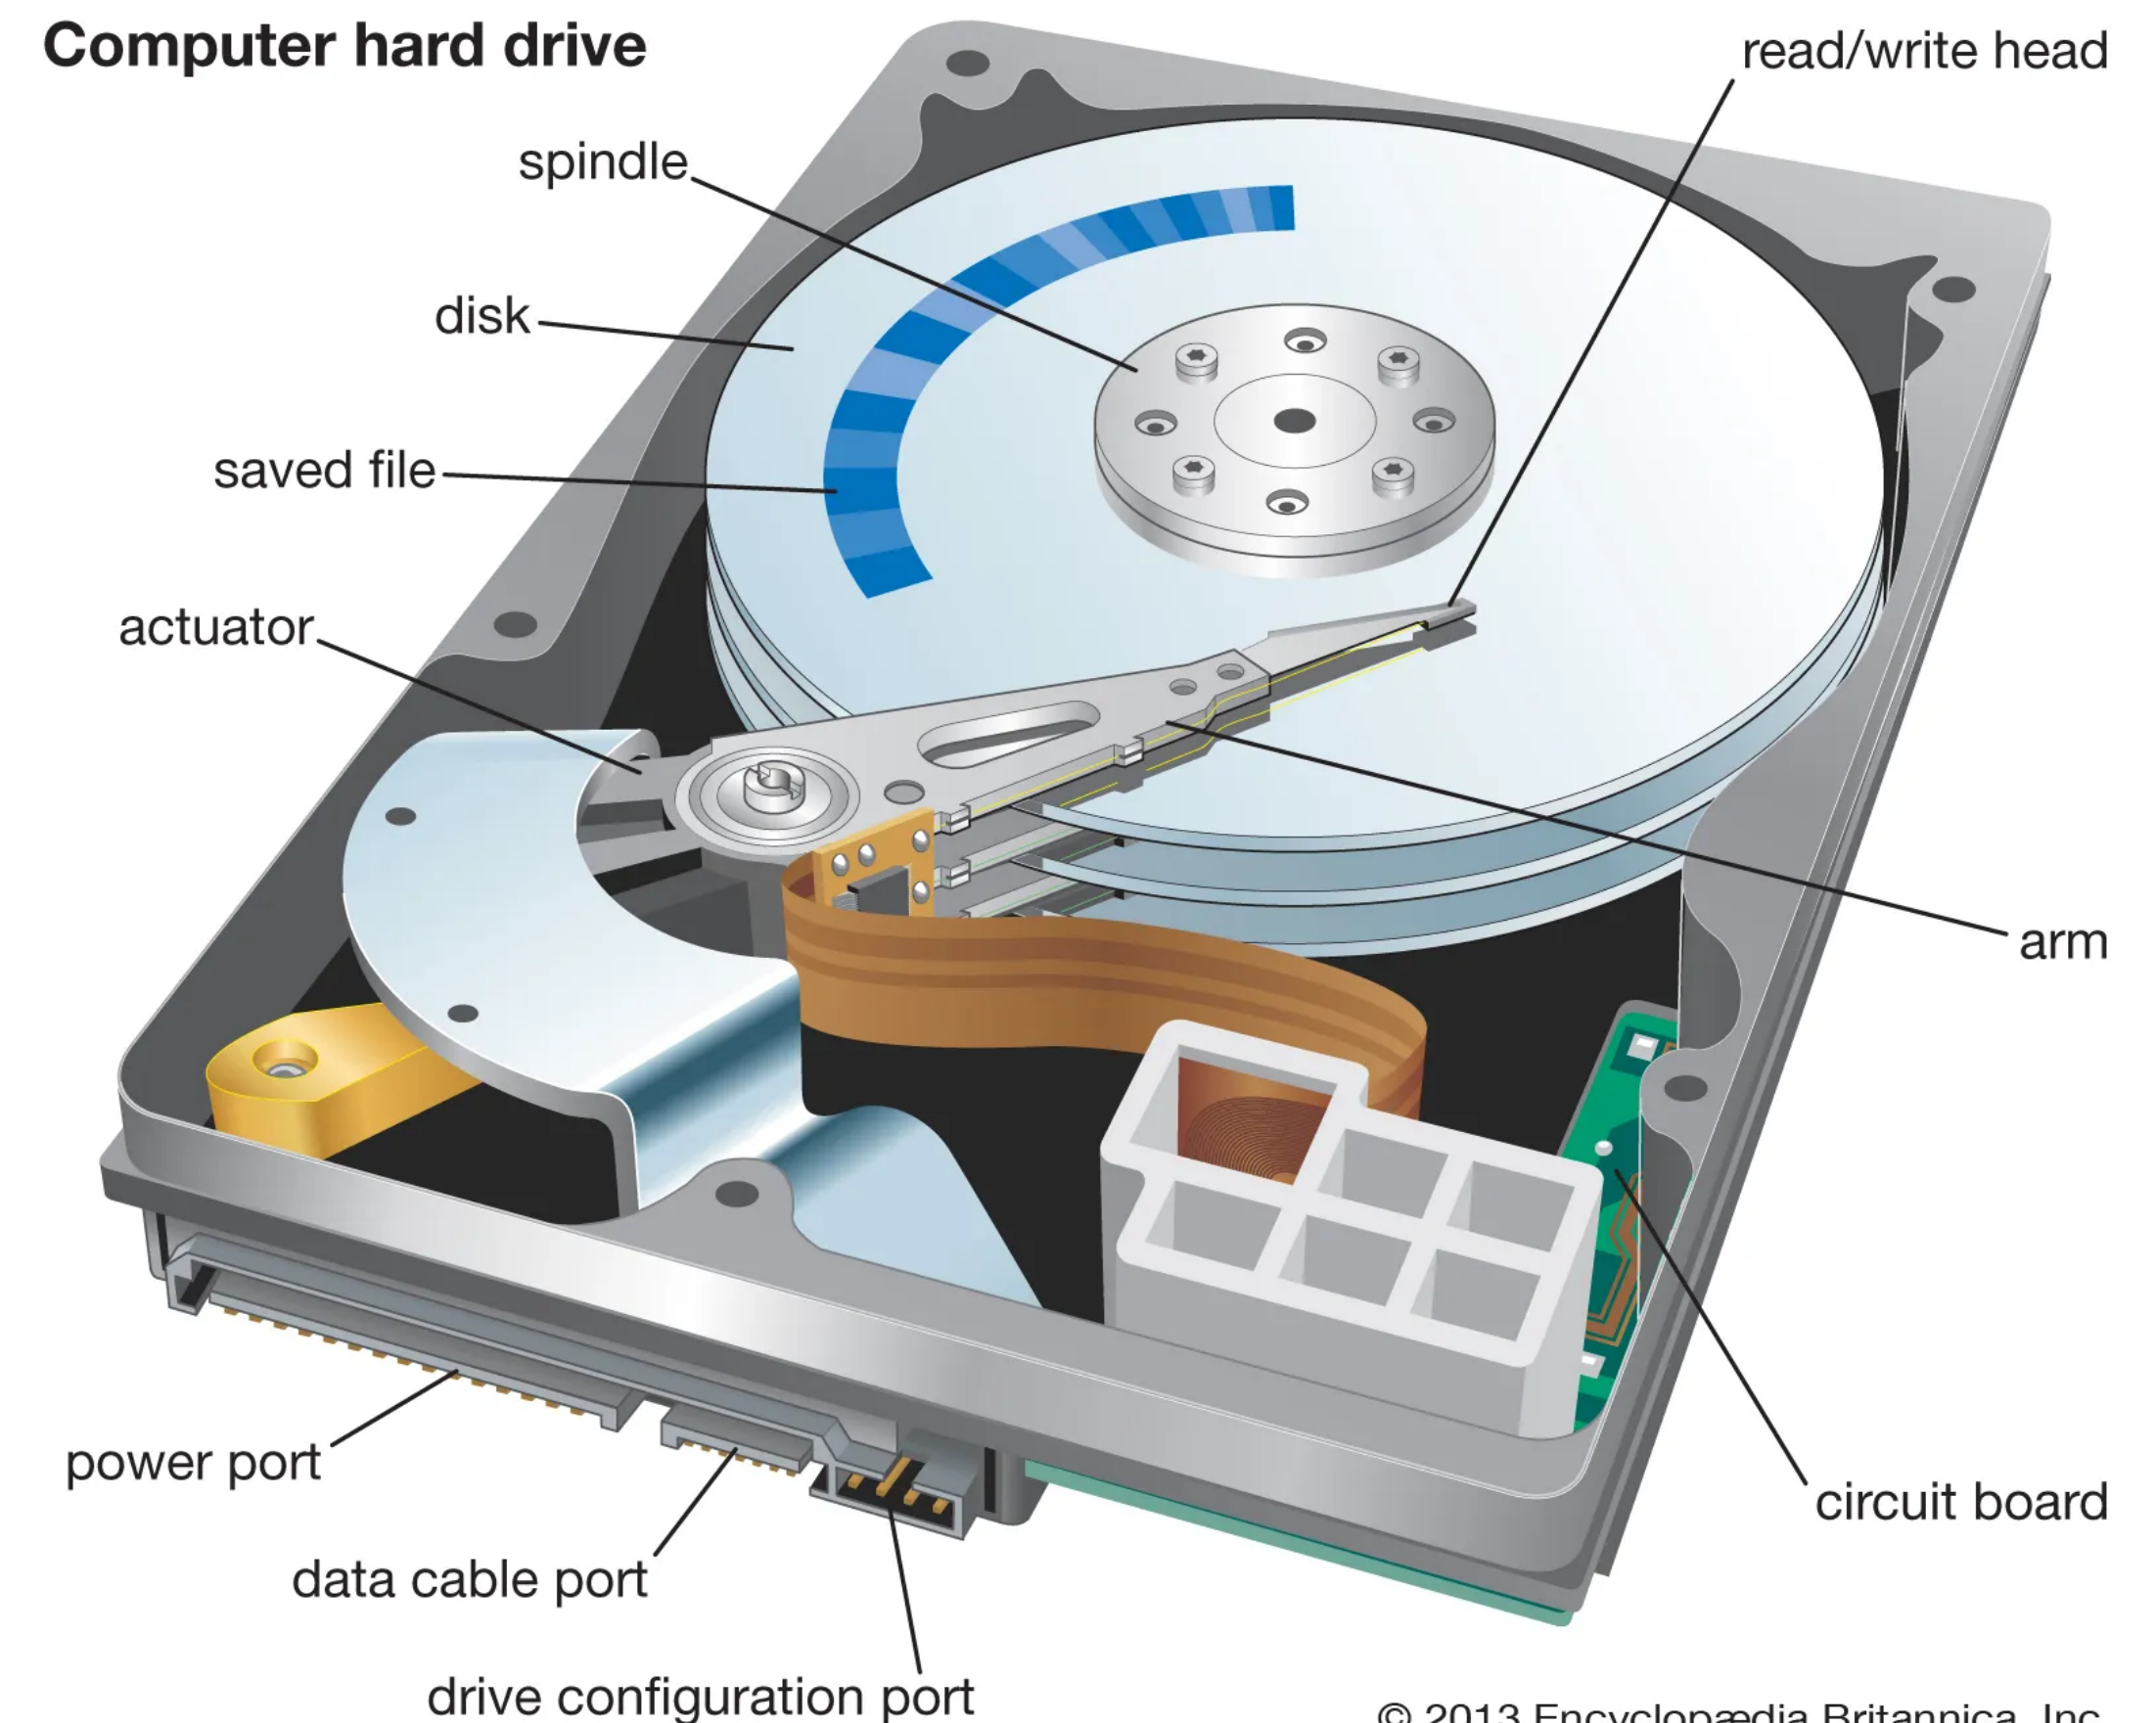
Disk geometry

- Many platters spinning on a spindle (~10,000 RPM)
- Each platter has two disk heads, one for each surface
- Disk heads are controlled by actuator
- One circle is called a track. Data is stored in sectors

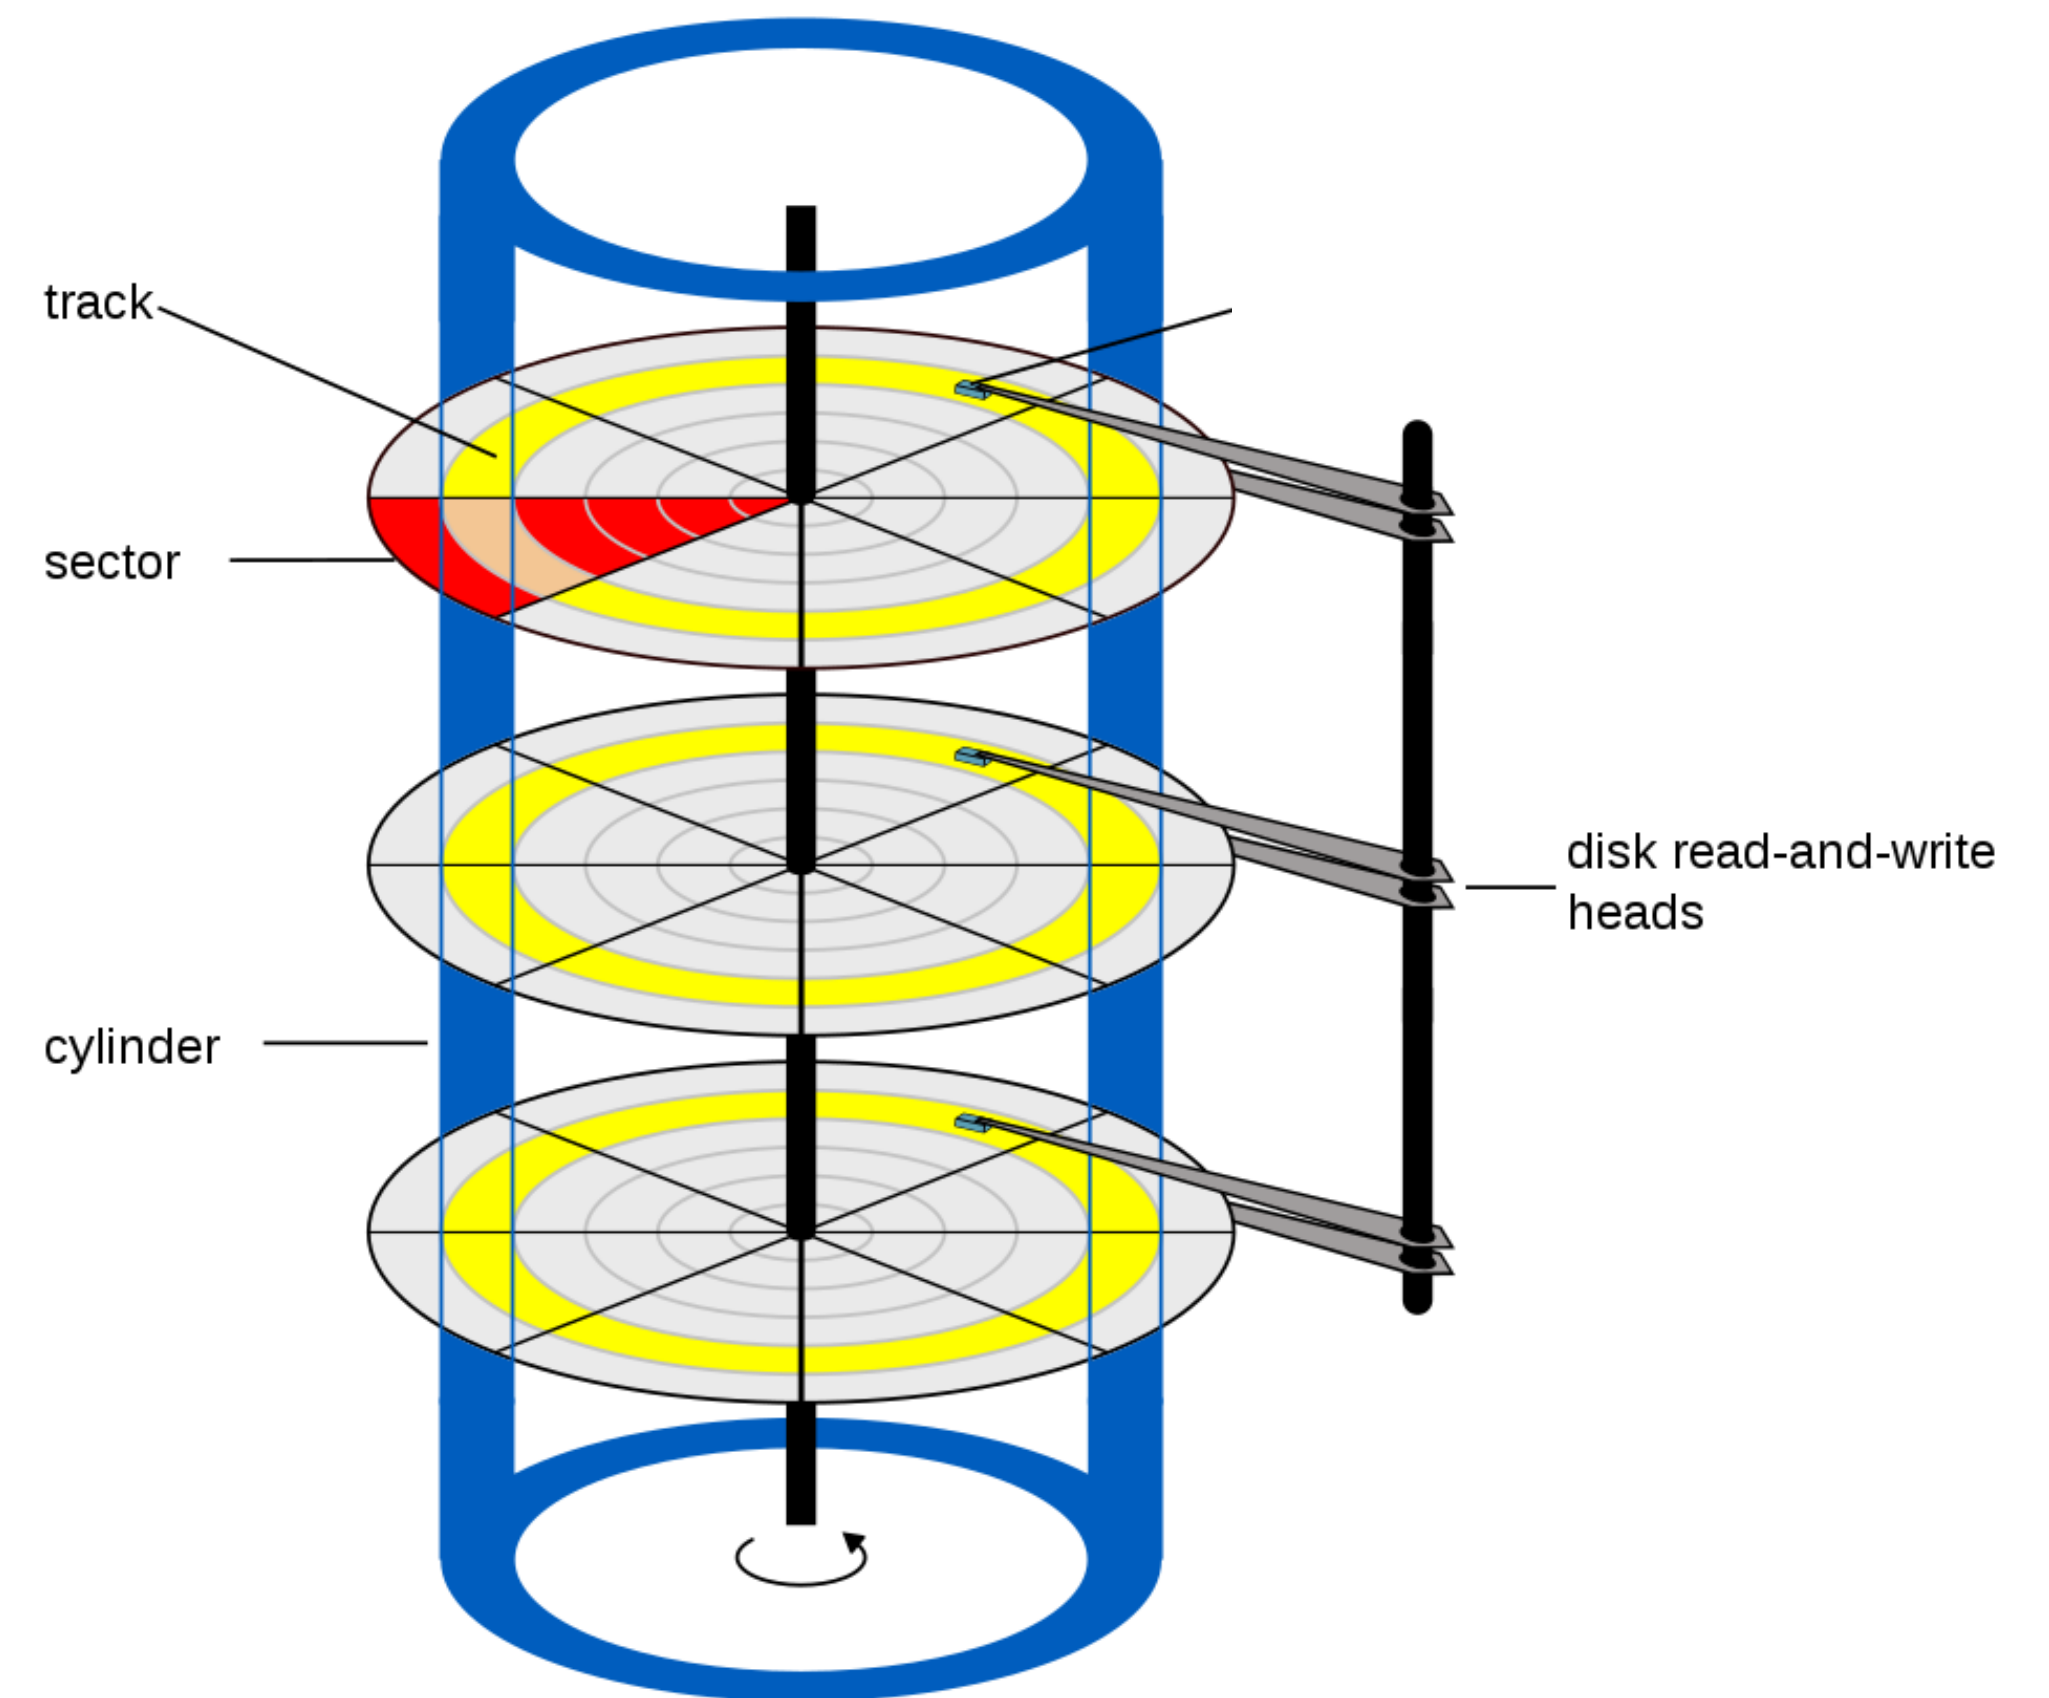


Disk geometry

- Many platters spinning on a spindle (~10,000 RPM)
- Each platter has two disk heads, one for each surface
- Disk heads are controlled by actuator
- One circle is called a track. Data is stored in sectors
- When the head is above a sector, it can read/write data

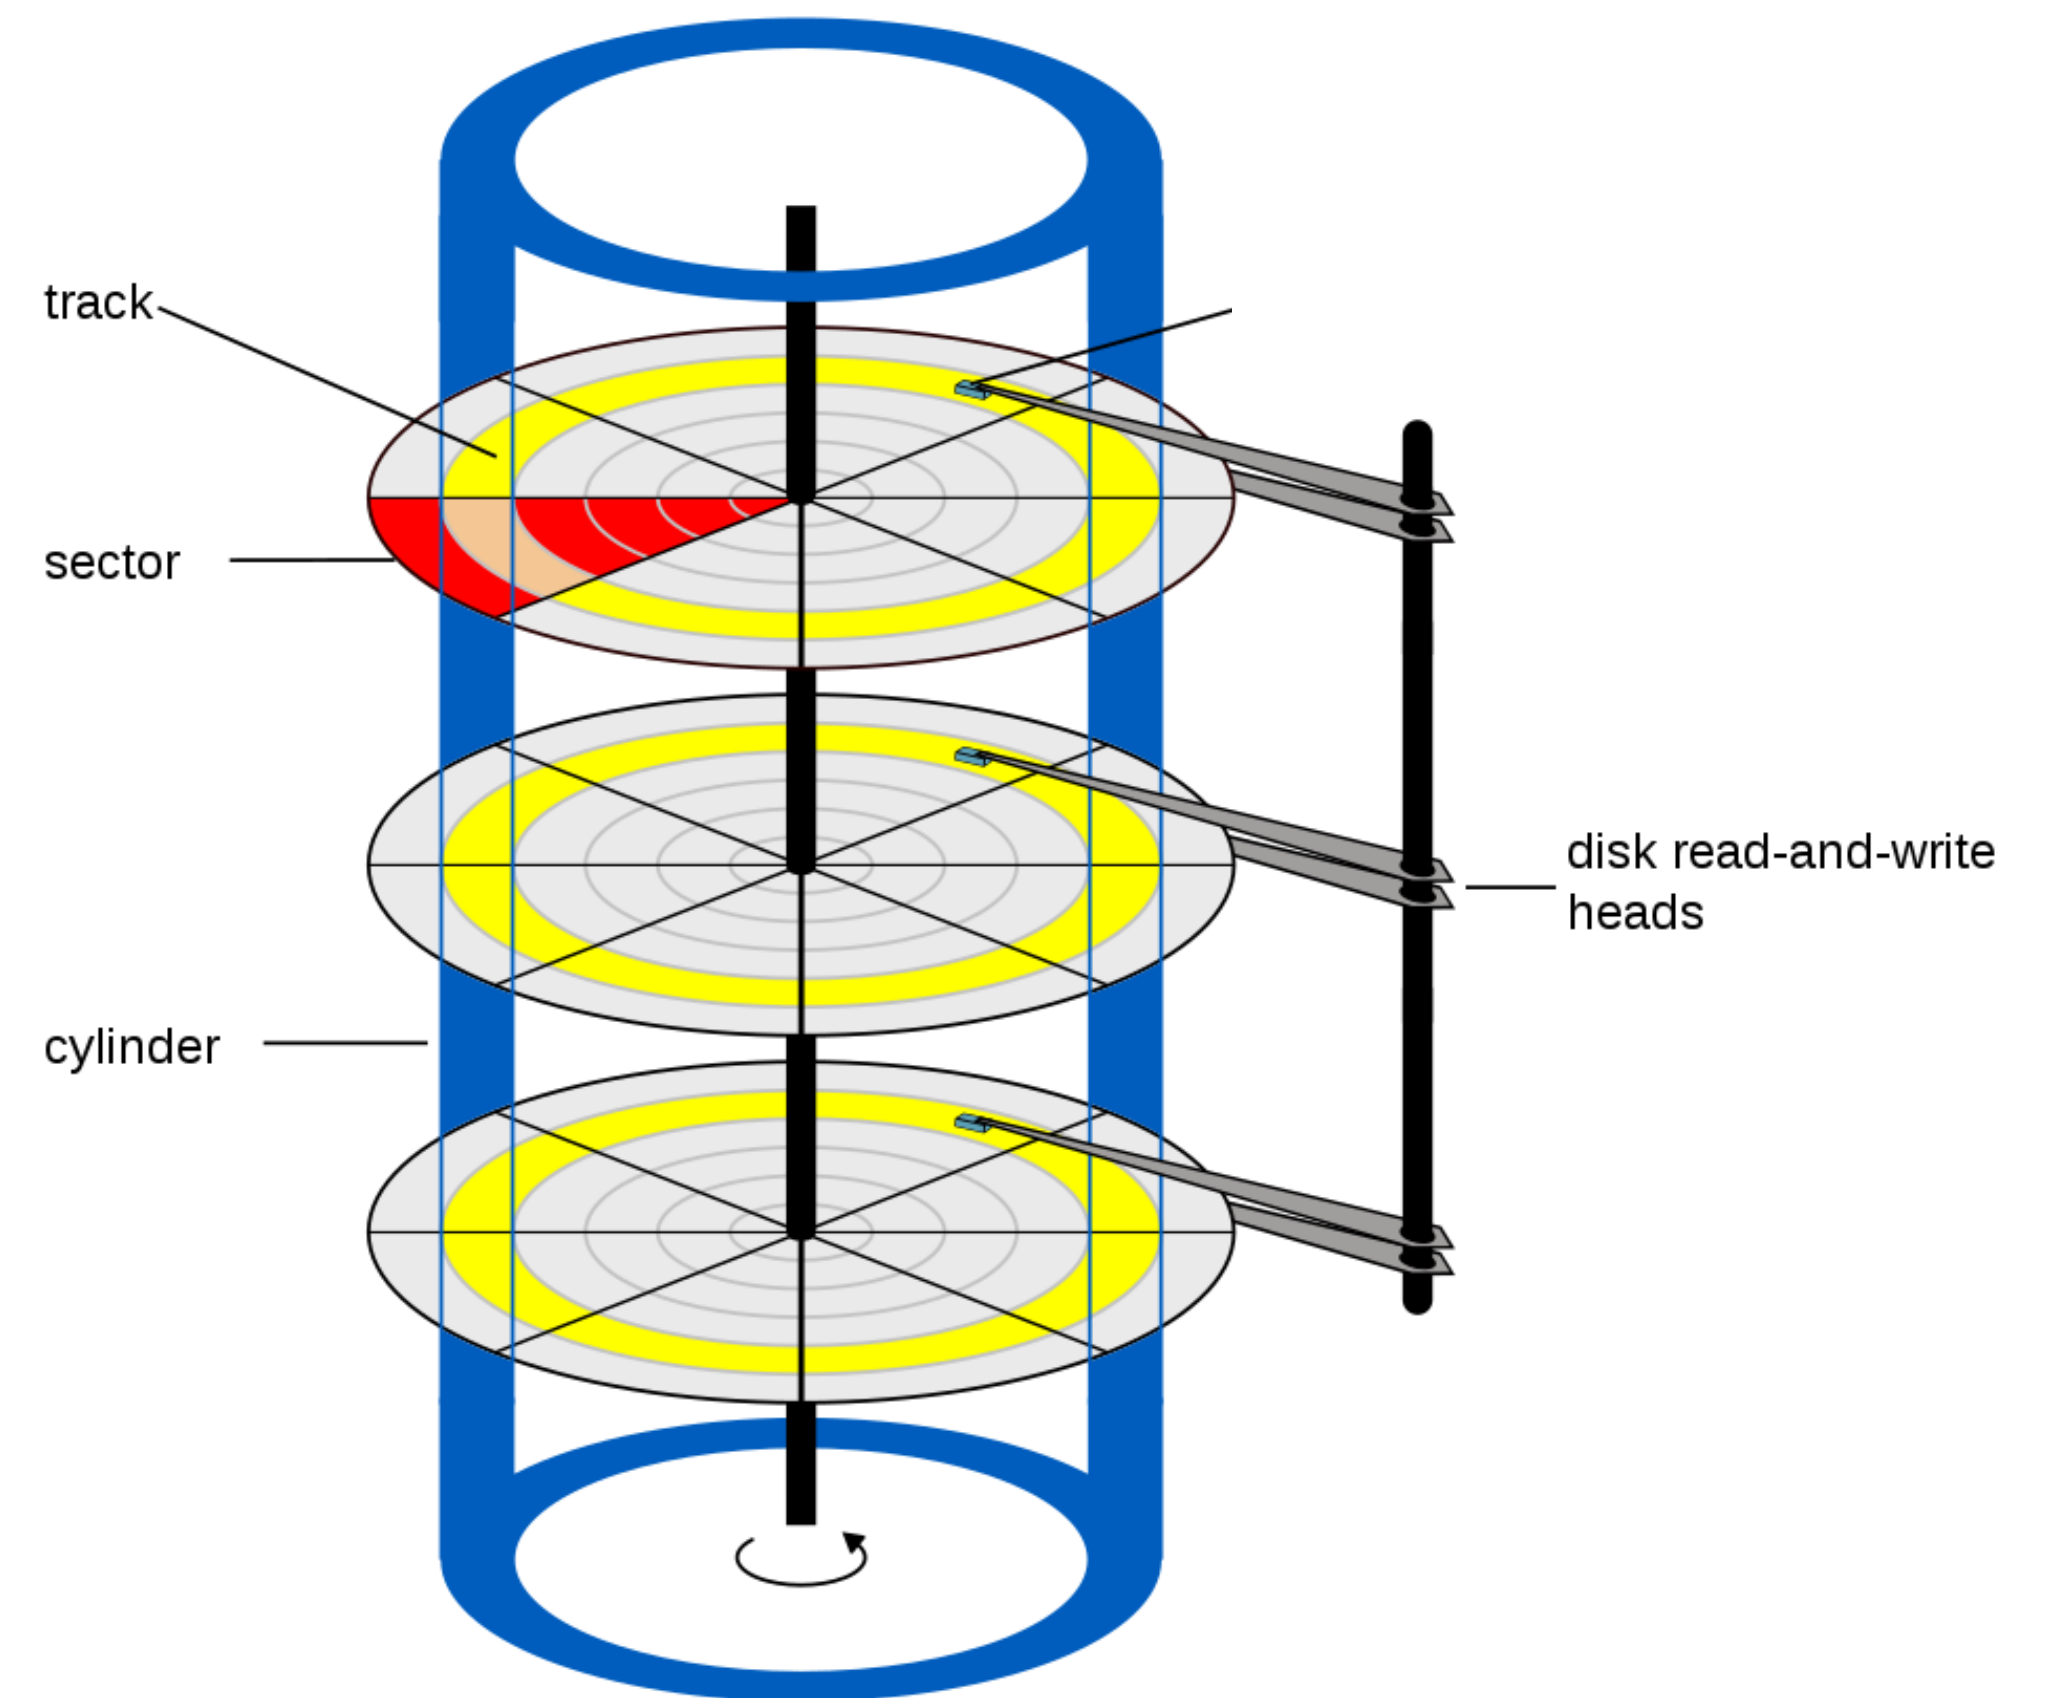


CPU-disk interface: Cylinder-head-sector (CHS) addressing



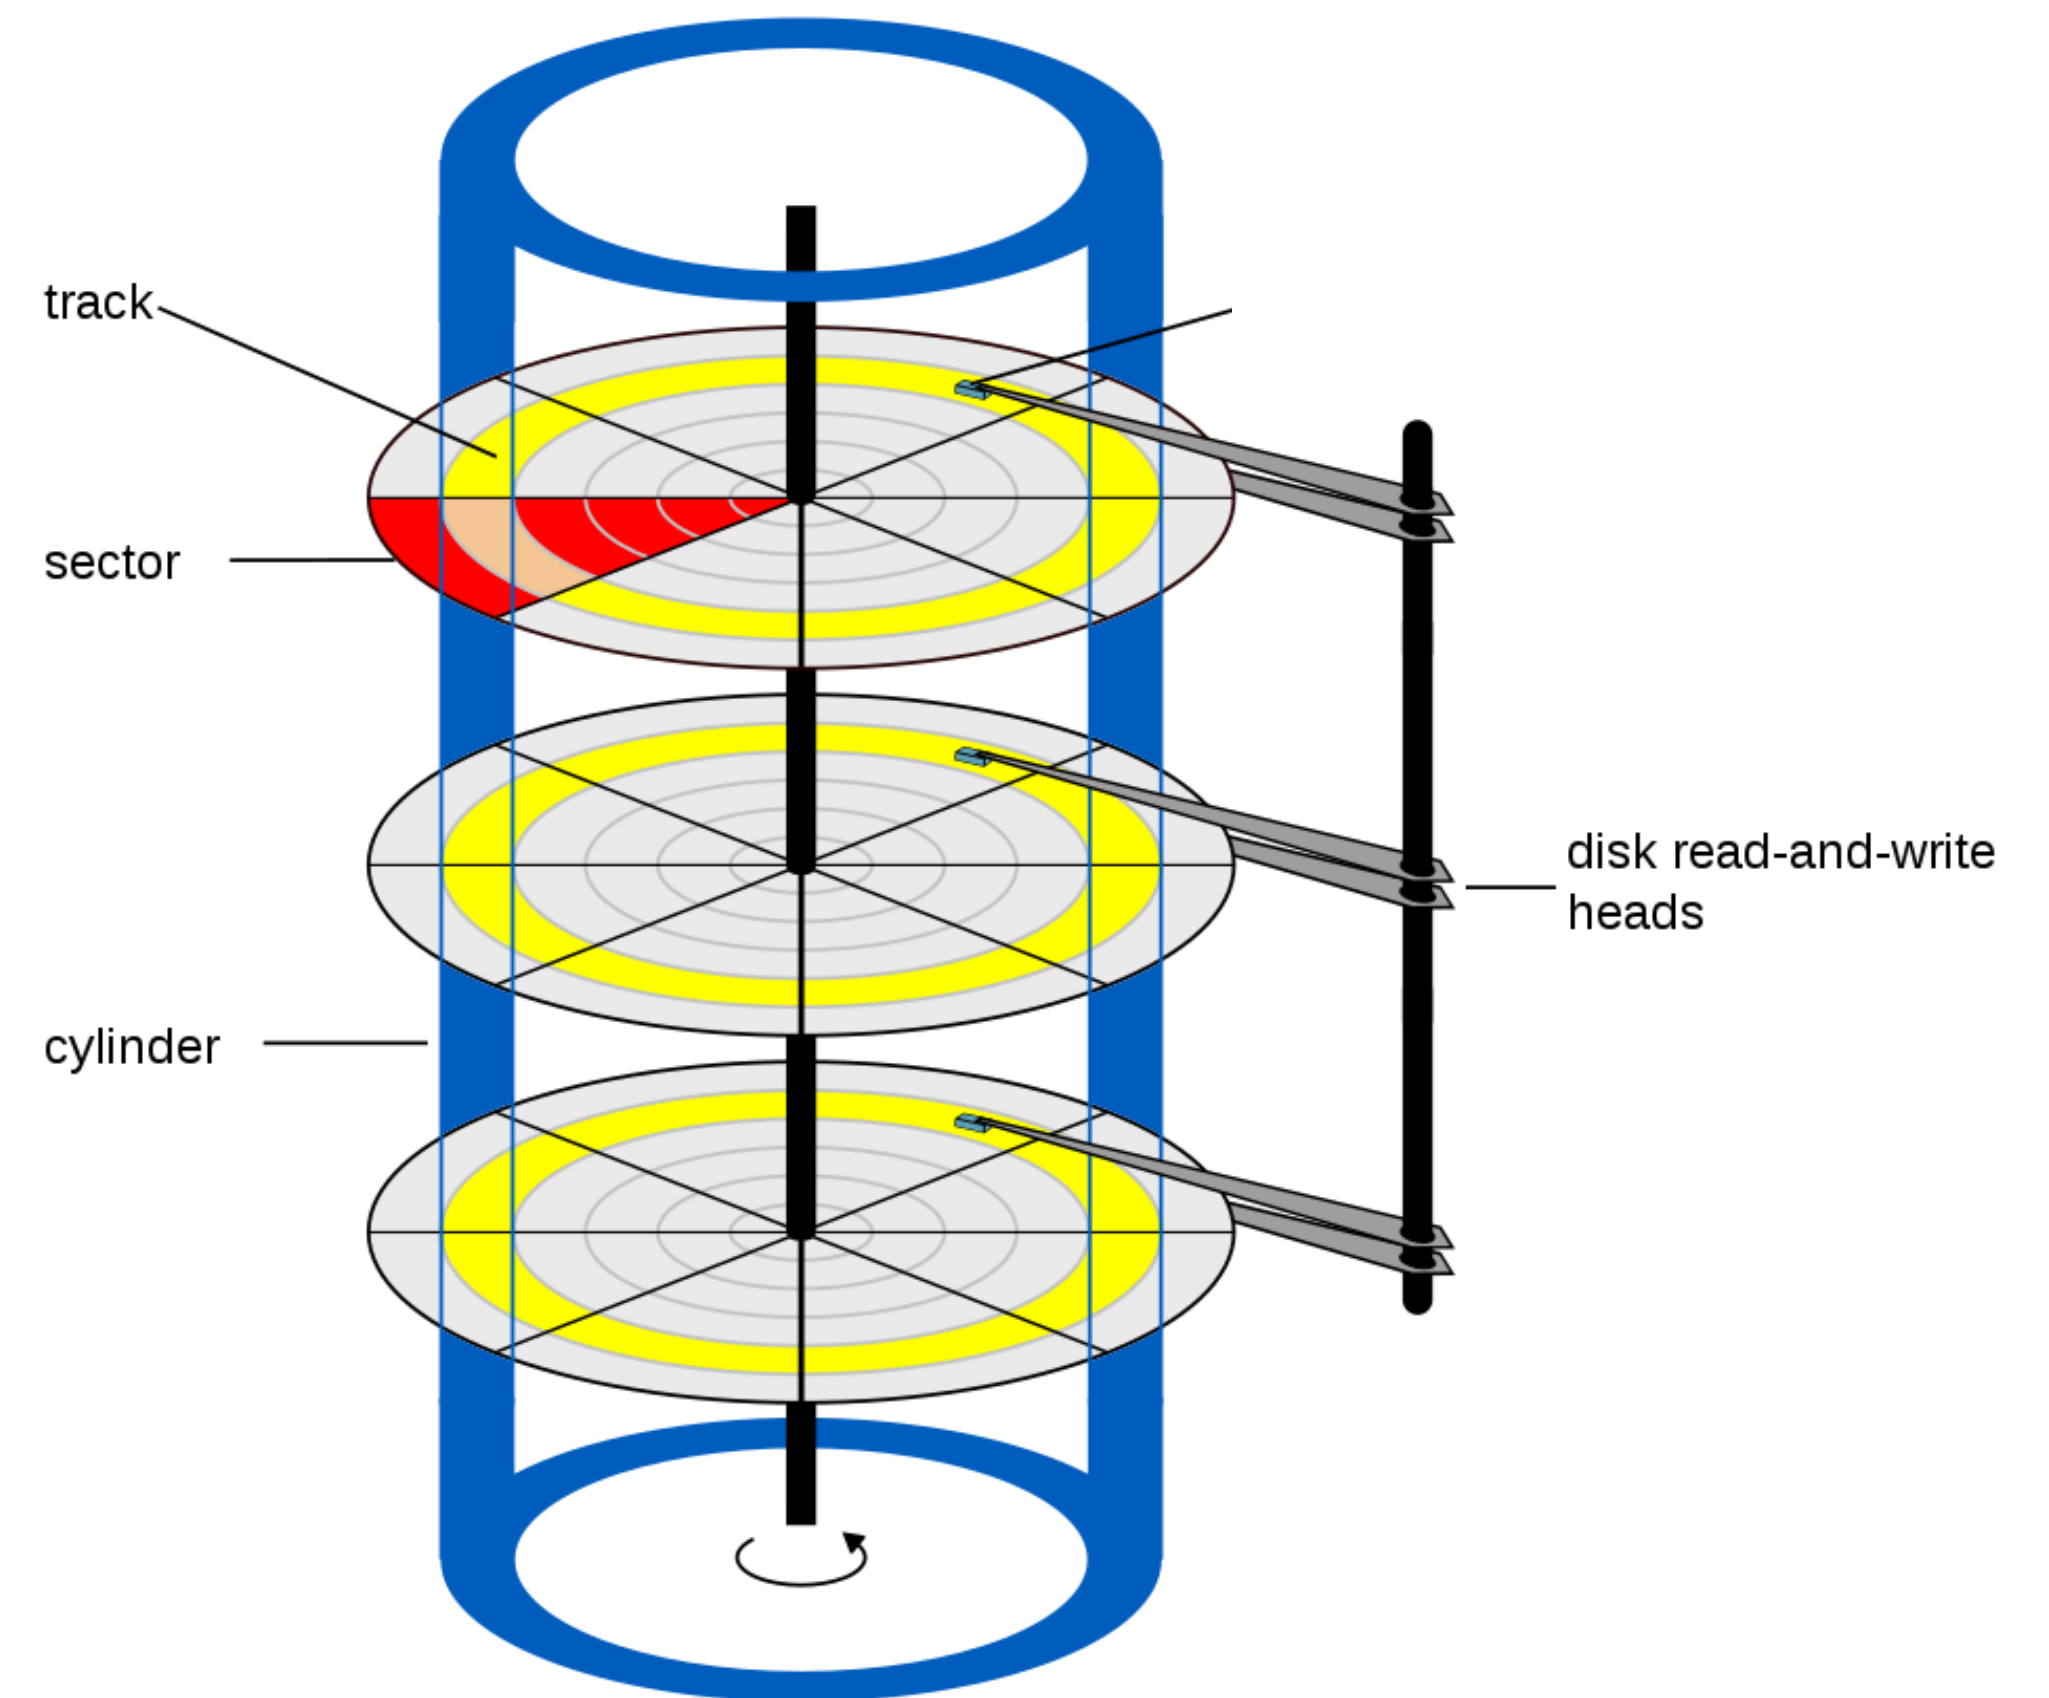
CPU-disk interface: Cylinder-head-sector (CHS) addressing

- C: cylinder number. 1024 cylinders.



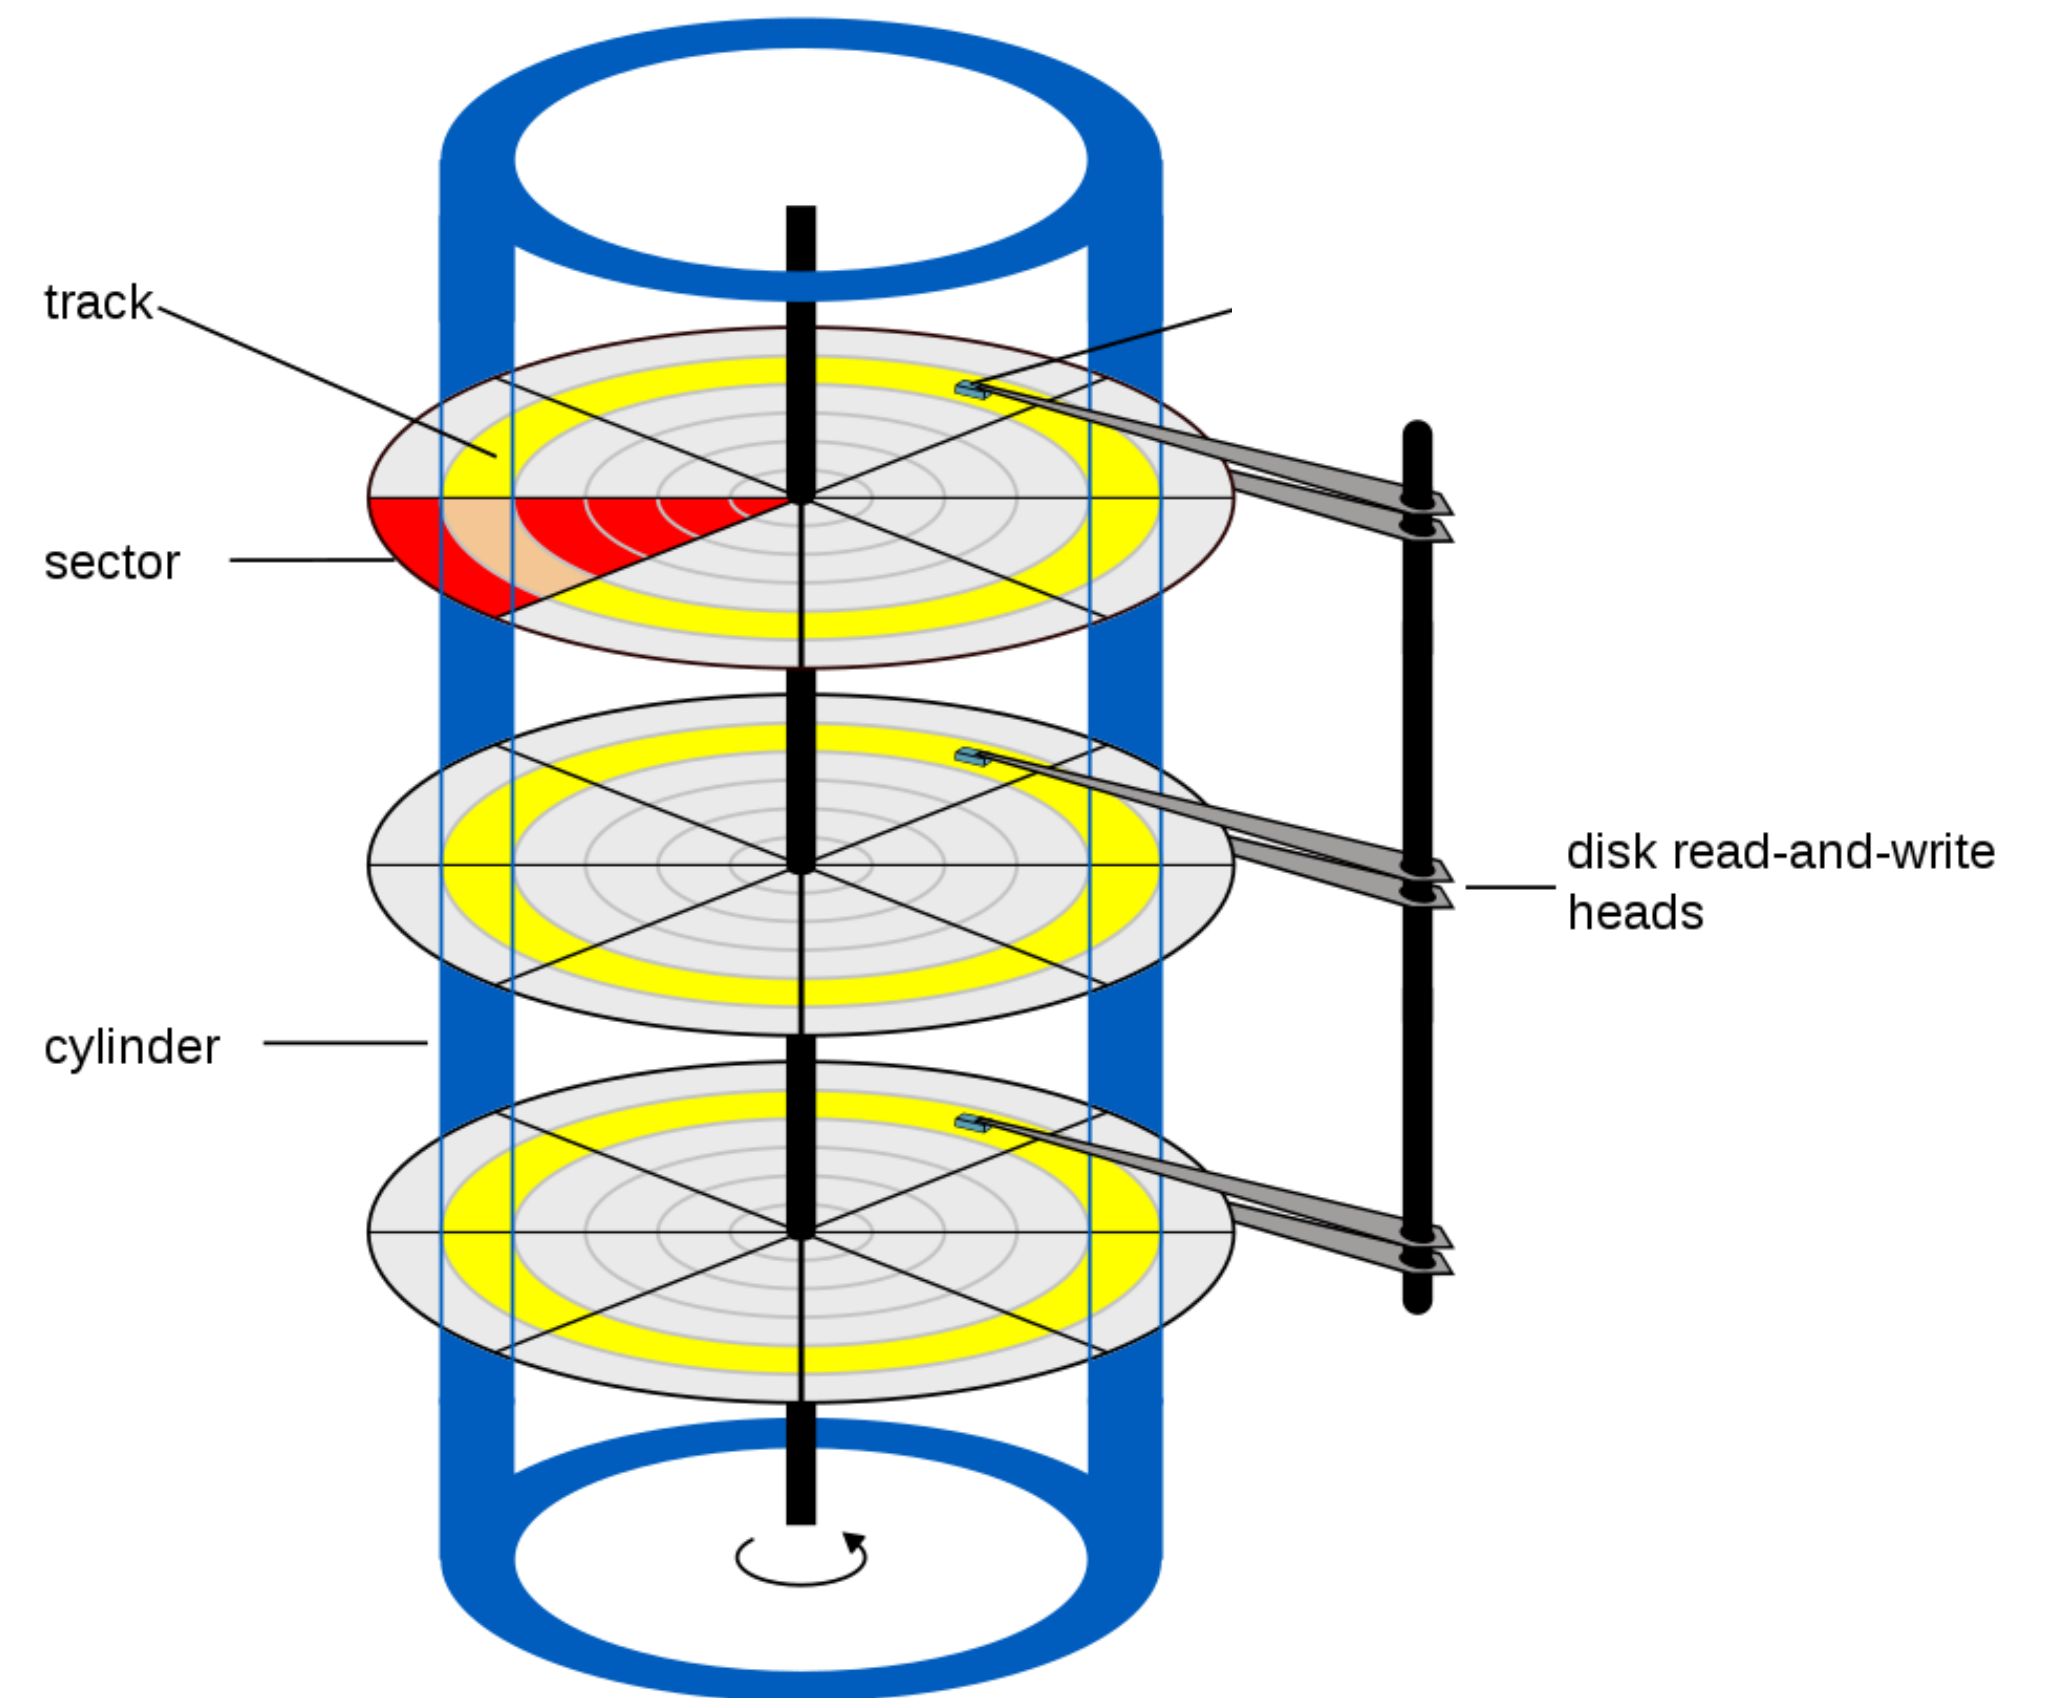
CPU-disk interface: Cylinder-head-sector (CHS) addressing

- C: cylinder number. 1024 cylinders.
- H: head number. 255 heads.



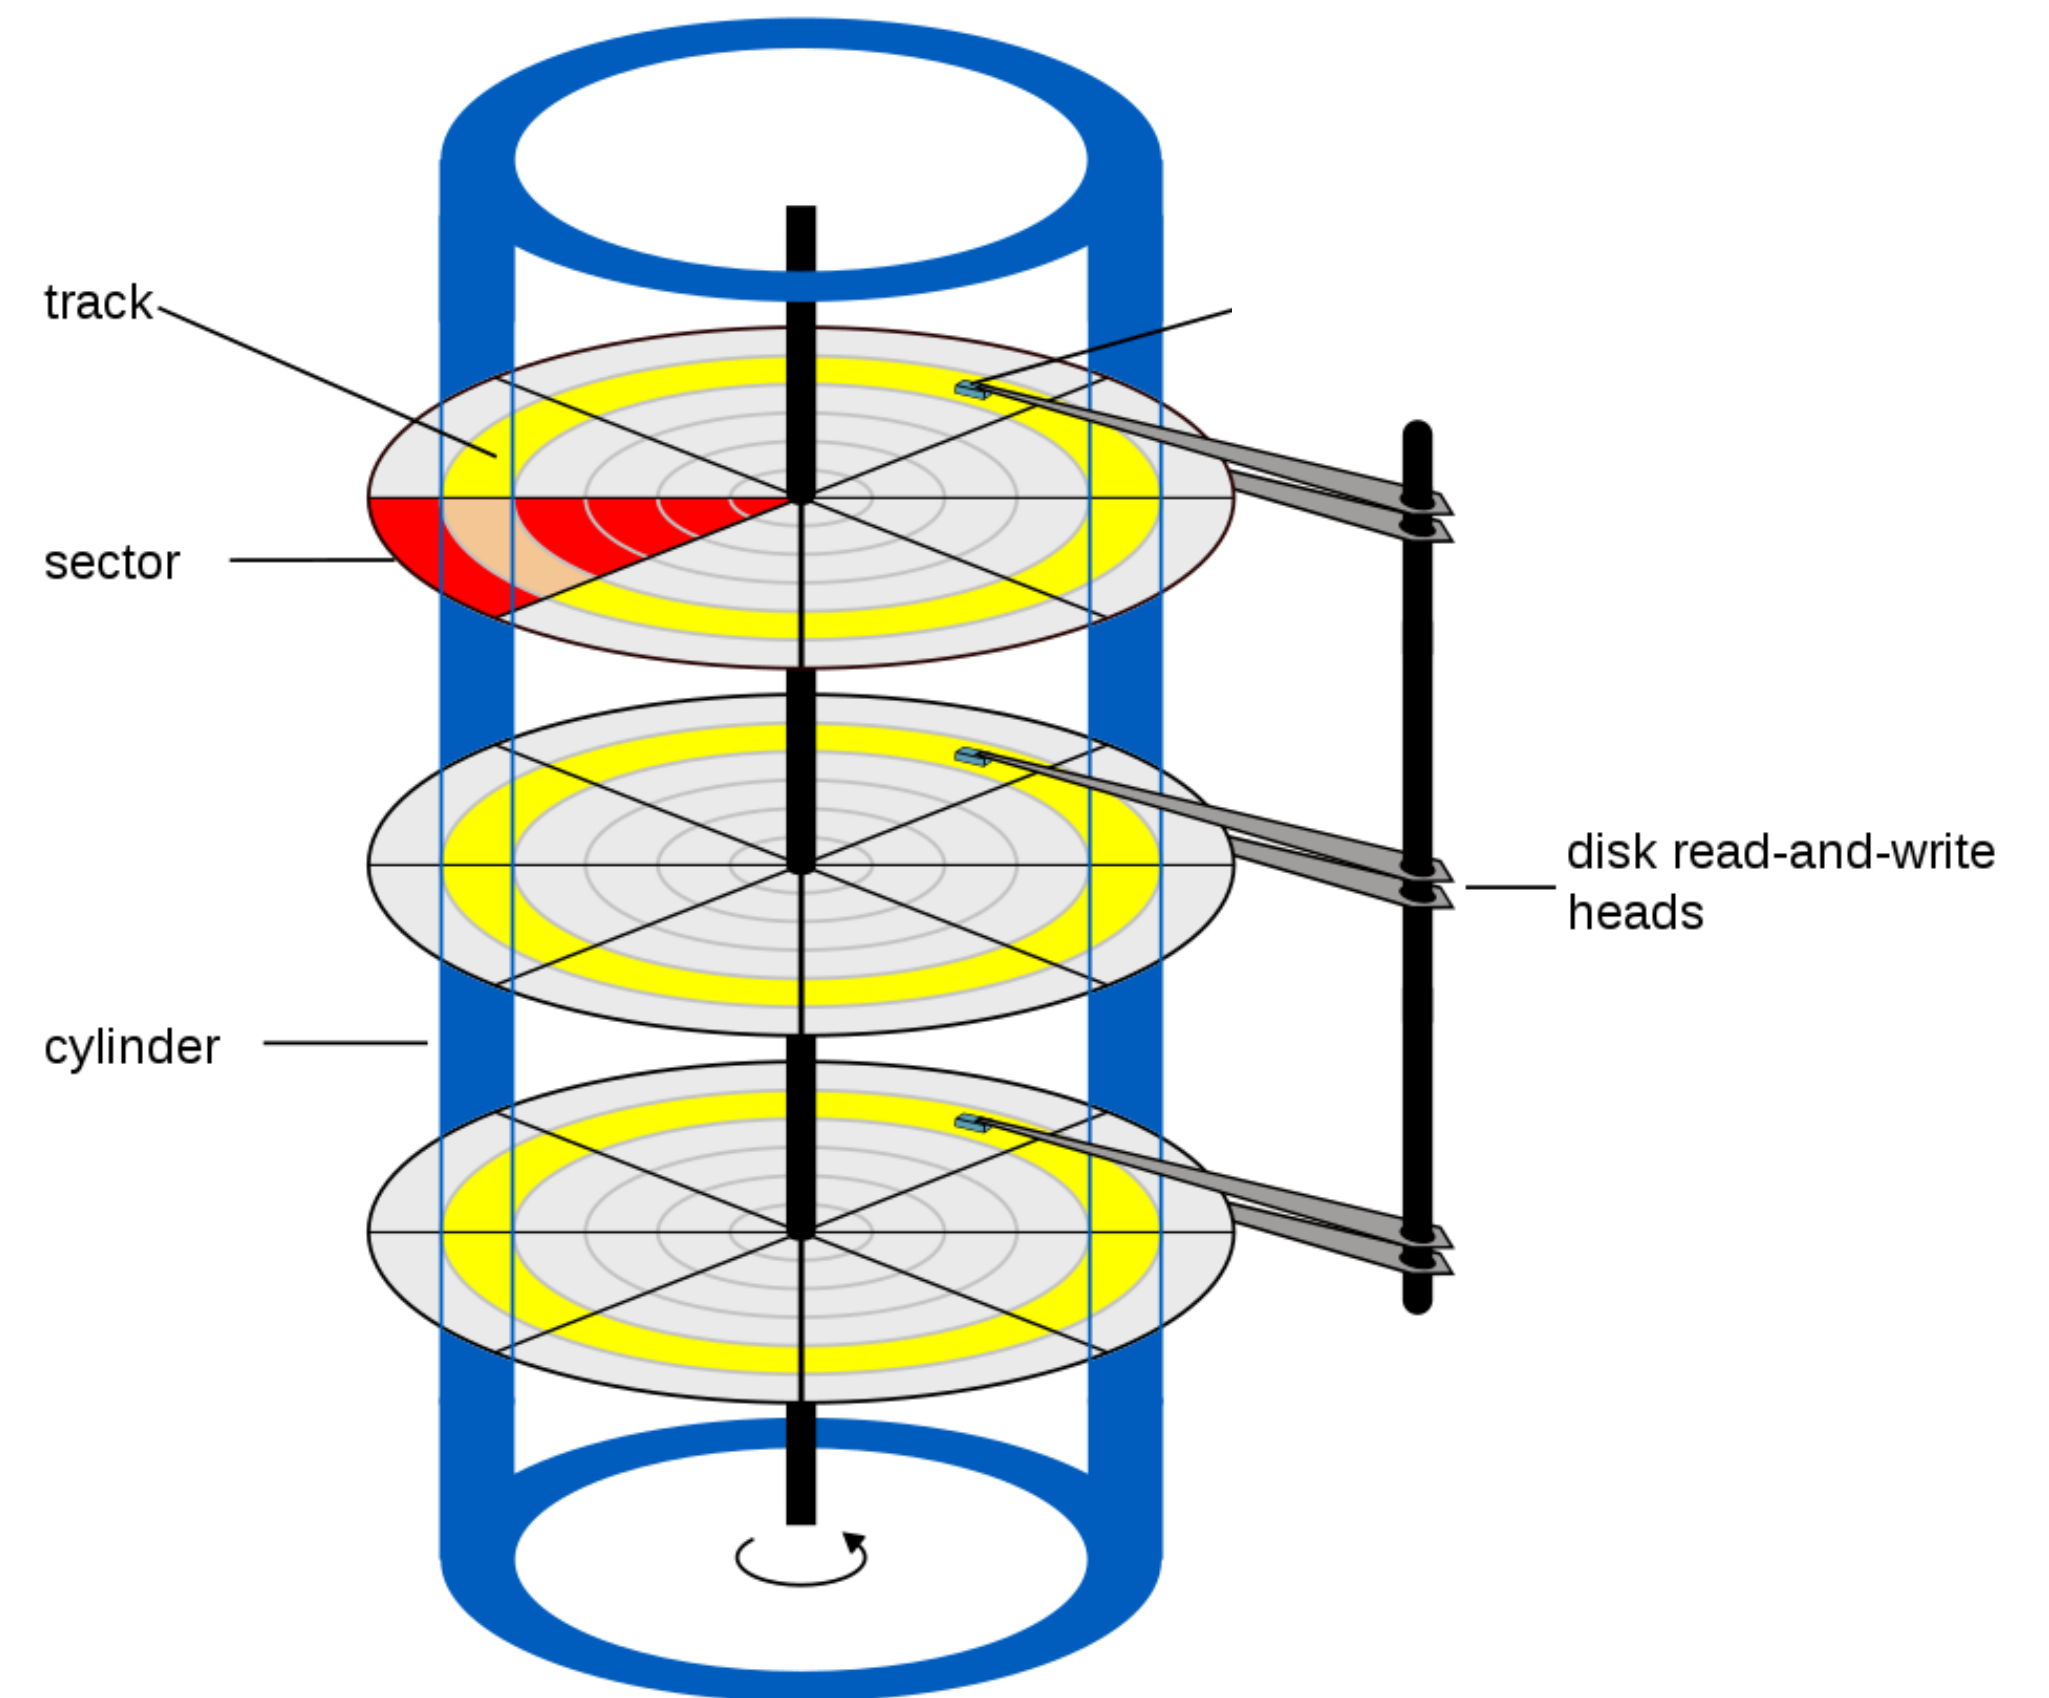
CPU-disk interface: Cylinder-head-sector (CHS) addressing

- C: cylinder number. 1024 cylinders.
- H: head number. 255 heads.
- S: sector number. 63 sectors per track.



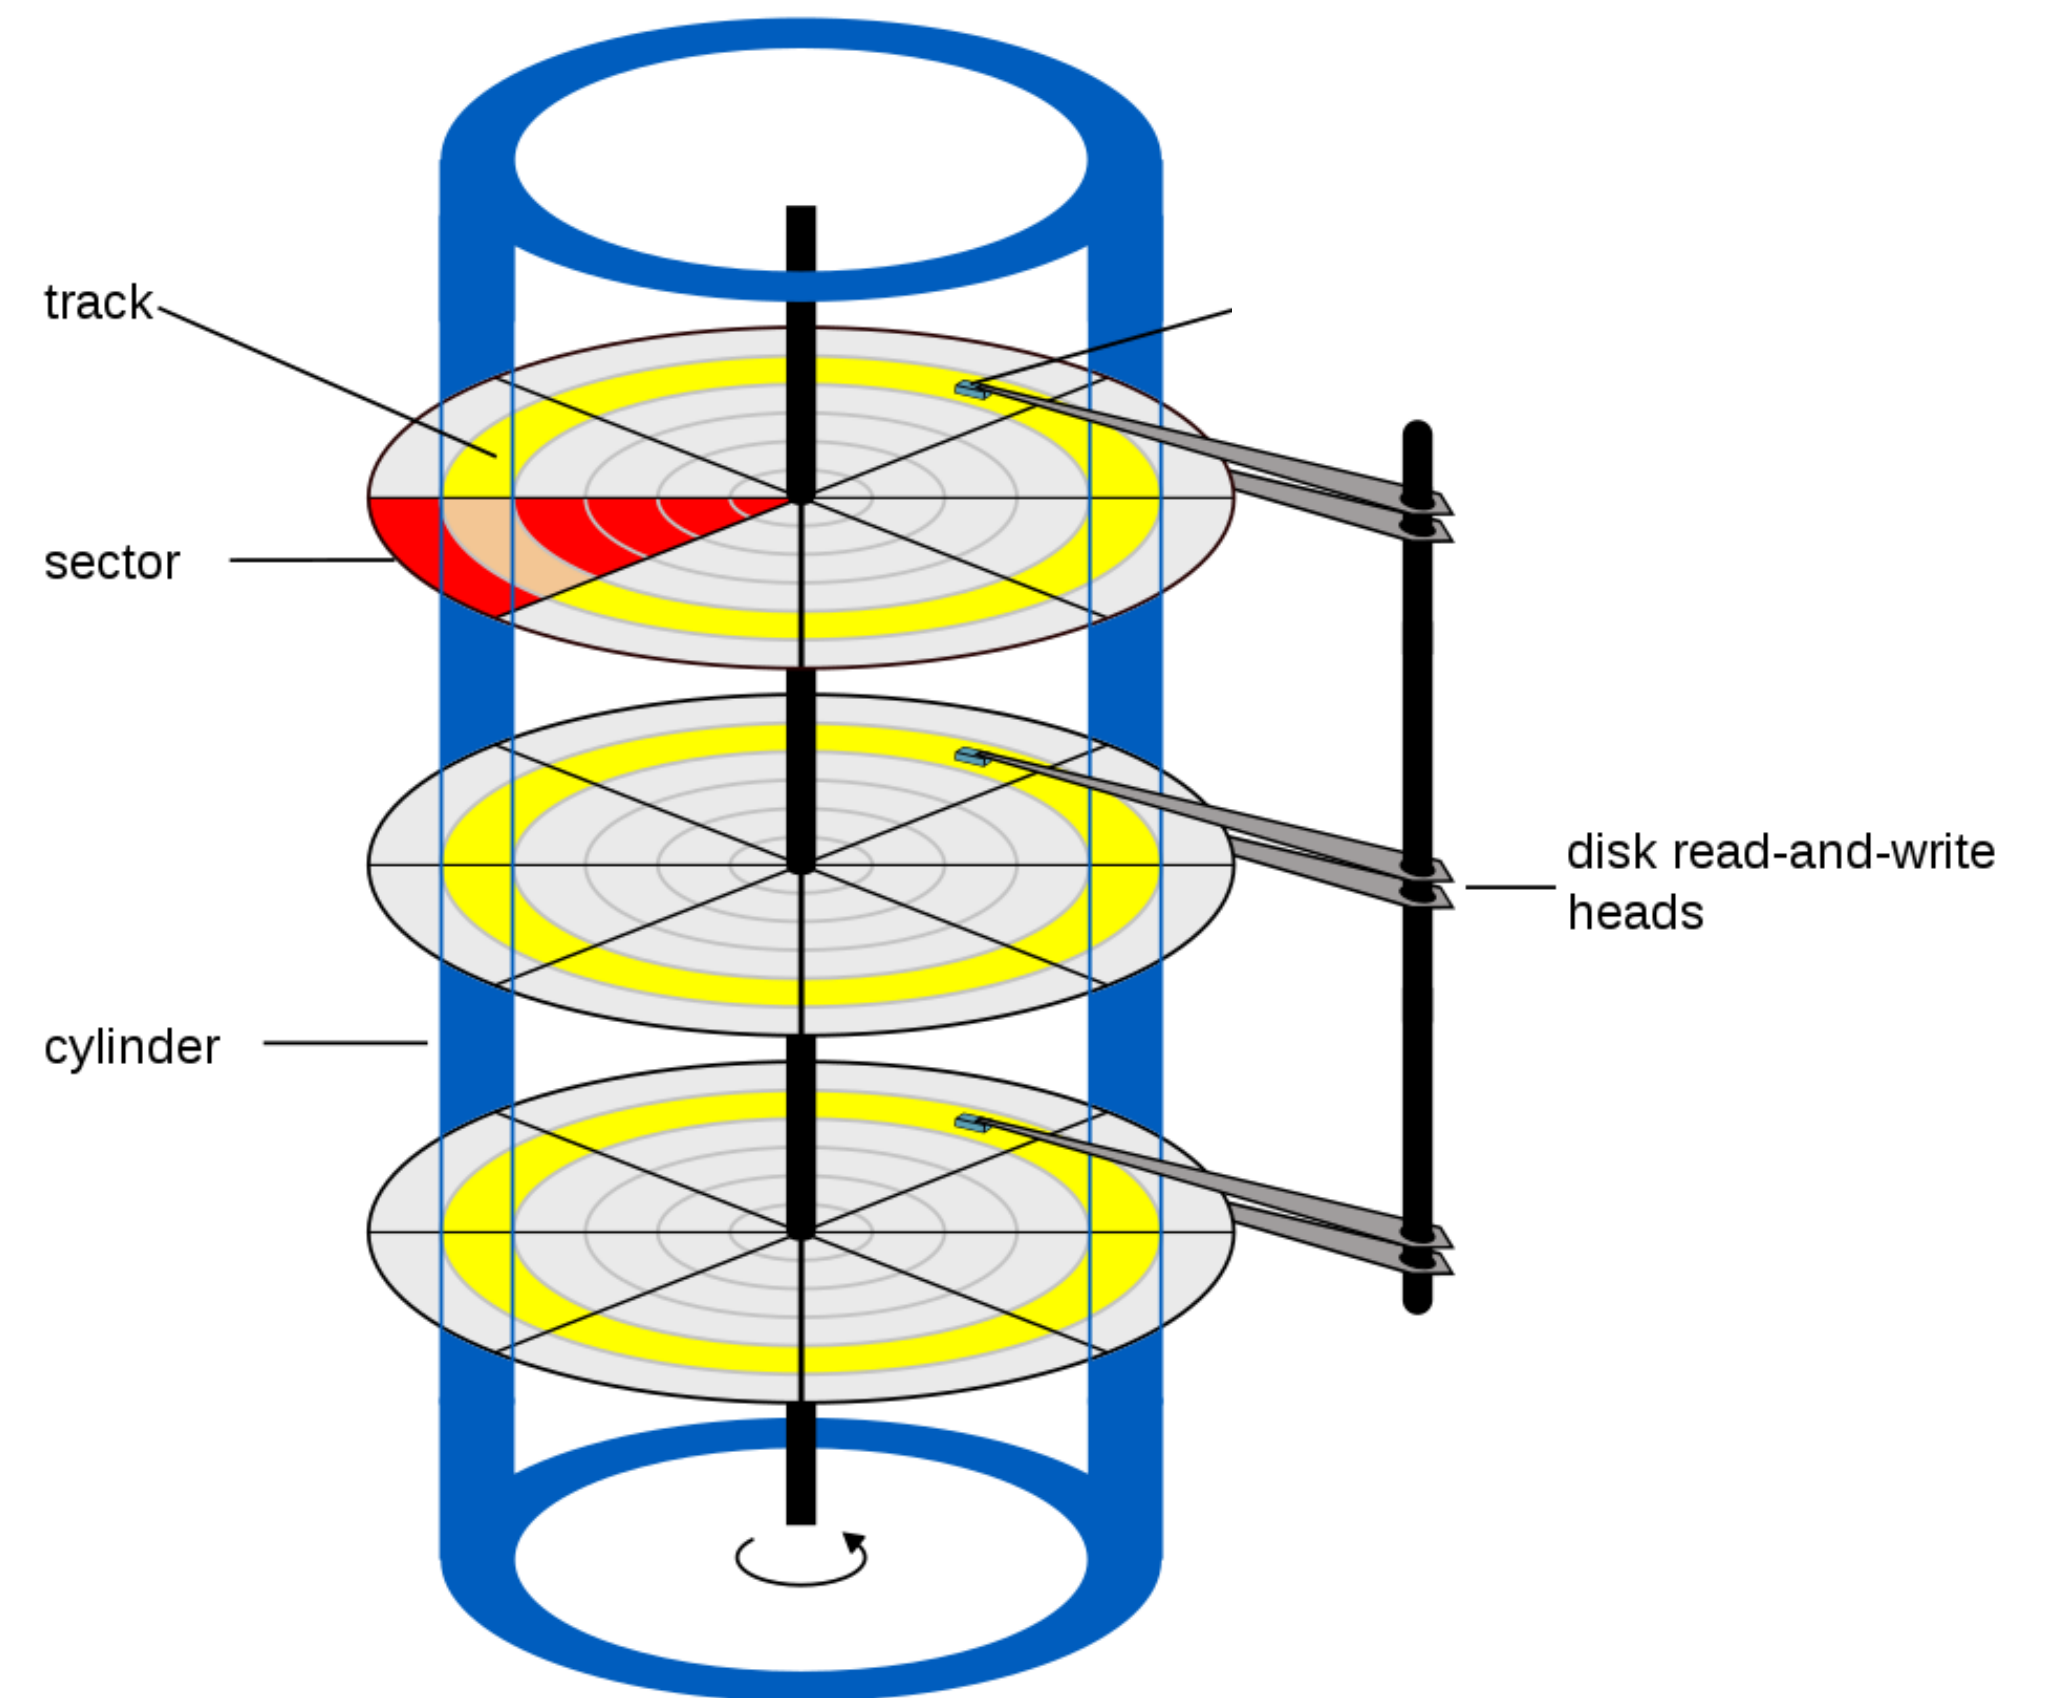
CPU-disk interface: Cylinder-head-sector (CHS) addressing

- C: cylinder number. 1024 cylinders.
- H: head number. 255 heads.
- S: sector number. 63 sectors per track.
- 512 bytes in each sector



CPU-disk interface: Cylinder-head-sector (CHS) addressing

- C: cylinder number. 1024 cylinders.
- H: head number. 255 heads.
- S: sector number. 63 sectors per track.
- 512 bytes in each sector
- Example: read 40th cylinder's 26th sector using 7th head.



Example of reads

	Cheetah 15K.5
Capacity	300 GB
RPM	15,000
Average Seek	4 ms
Max Transfer	125 MB/s
Platters	4
Cache	16 MB

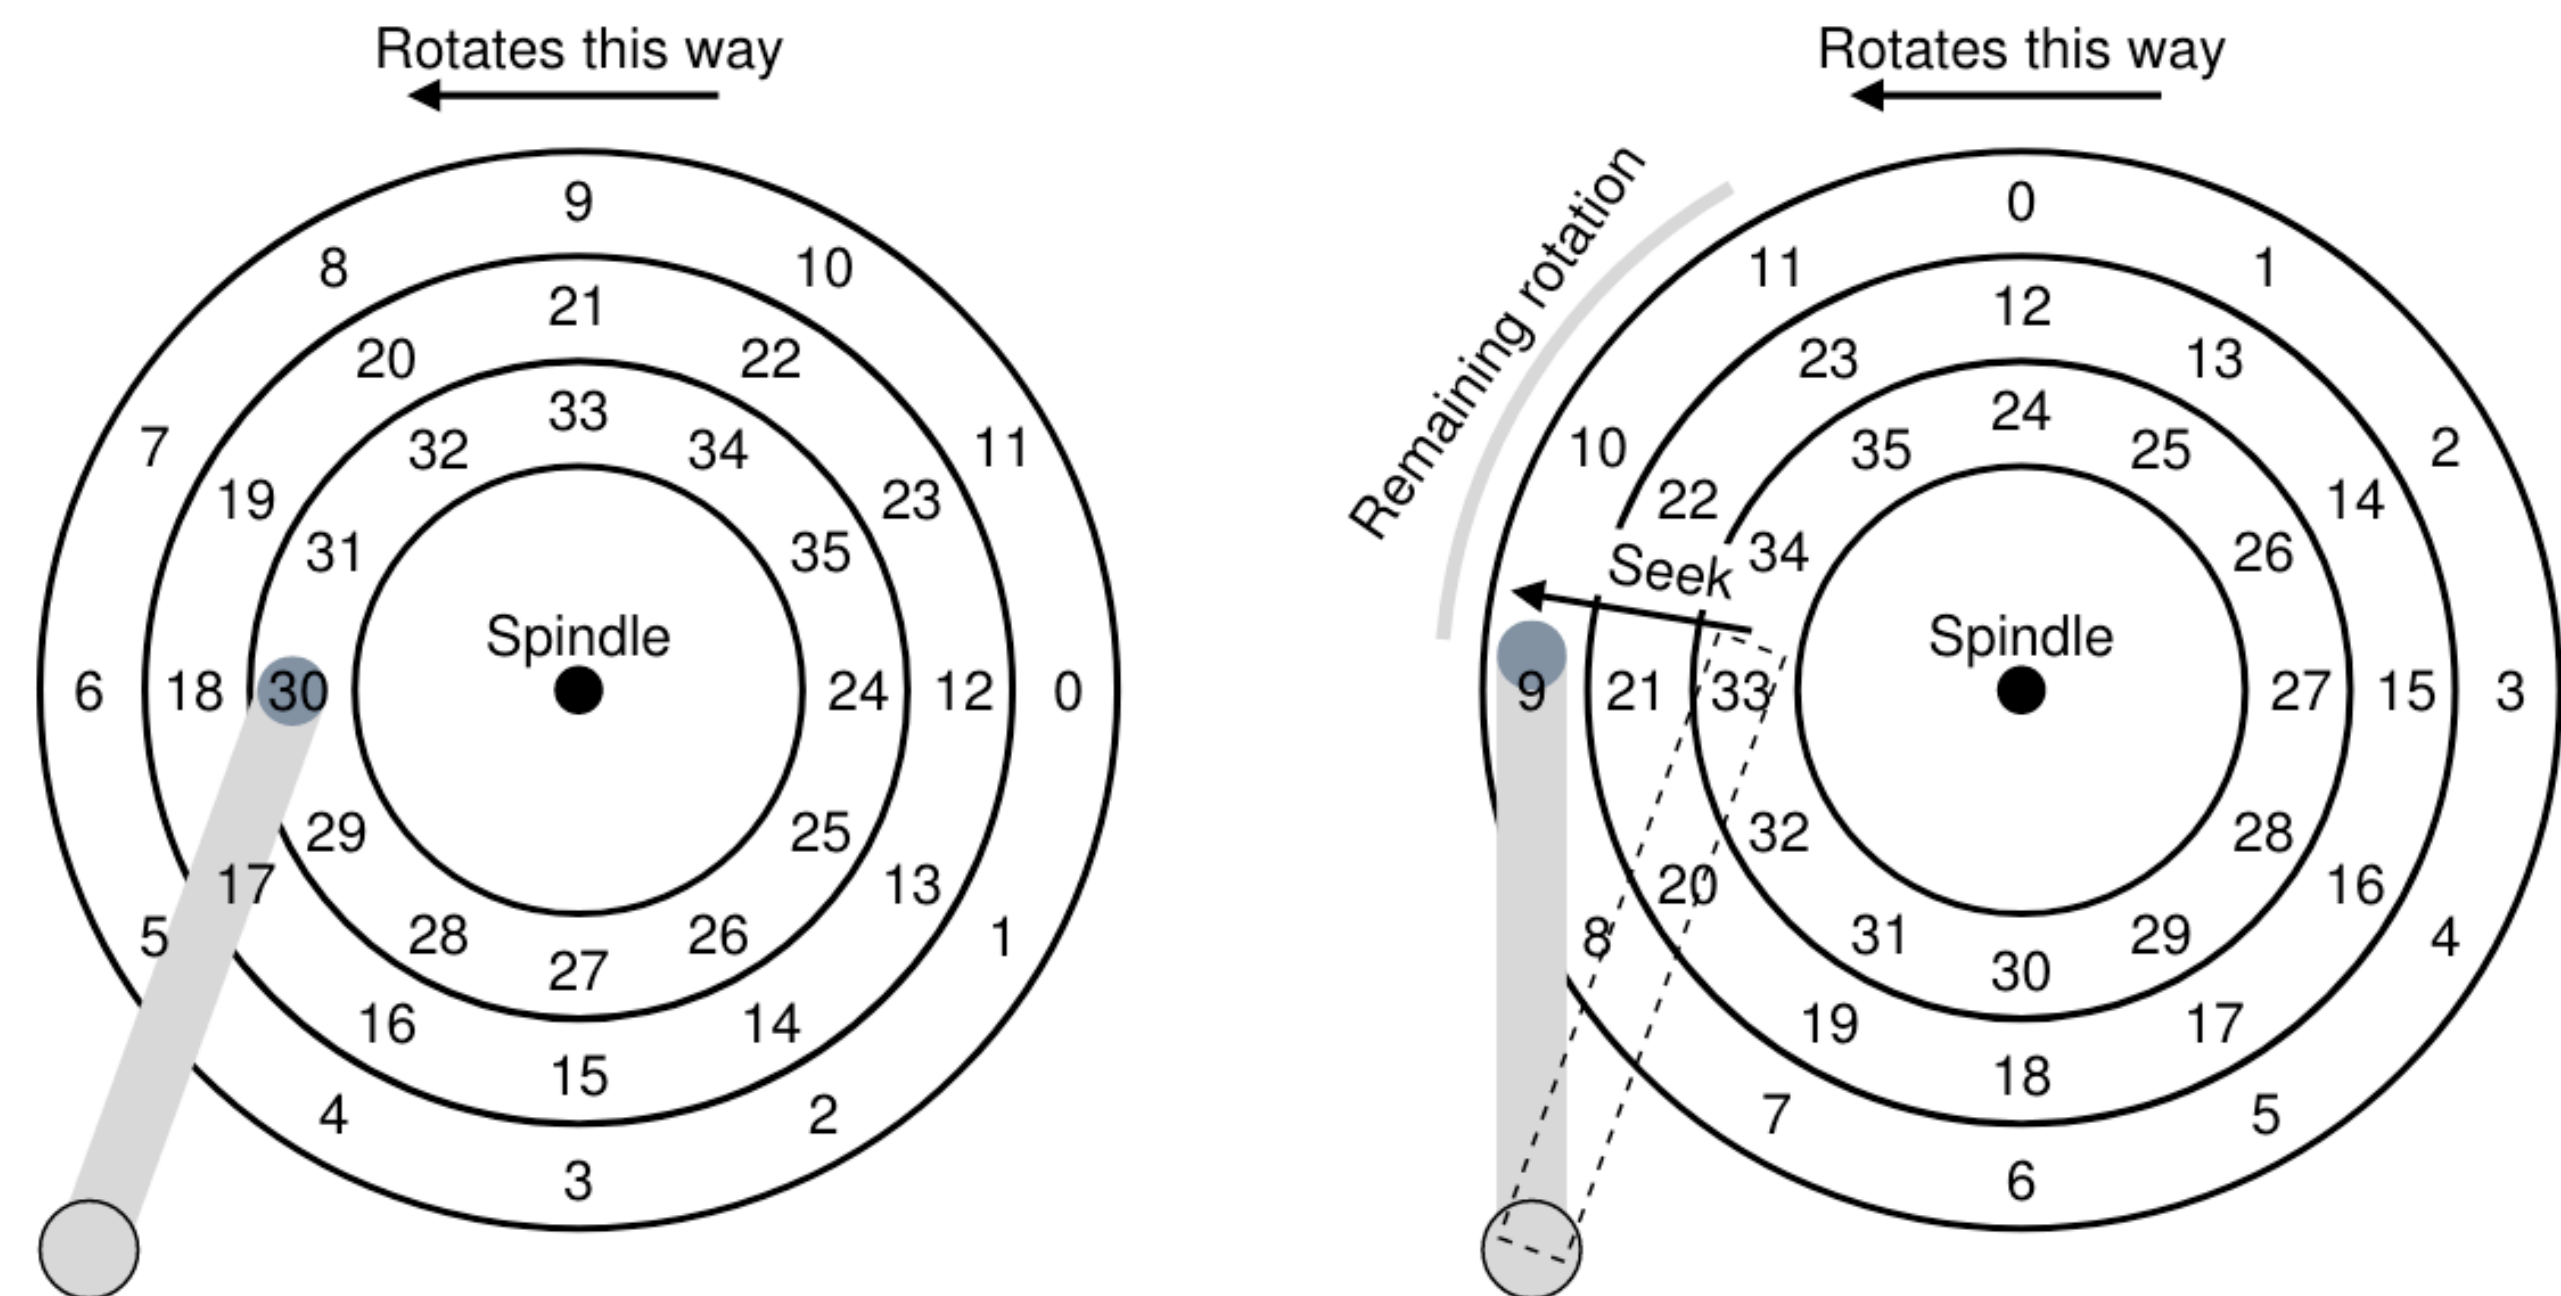


Figure 37.3: Three Tracks Plus A Head (Right: With Seek)

Example of reads

	Cheetah 15K.5
Capacity	300 GB
RPM	15,000
Average Seek	4 ms
Max Transfer	125 MB/s
Platters	4
Cache	16 MB

- Seek delay (4ms)

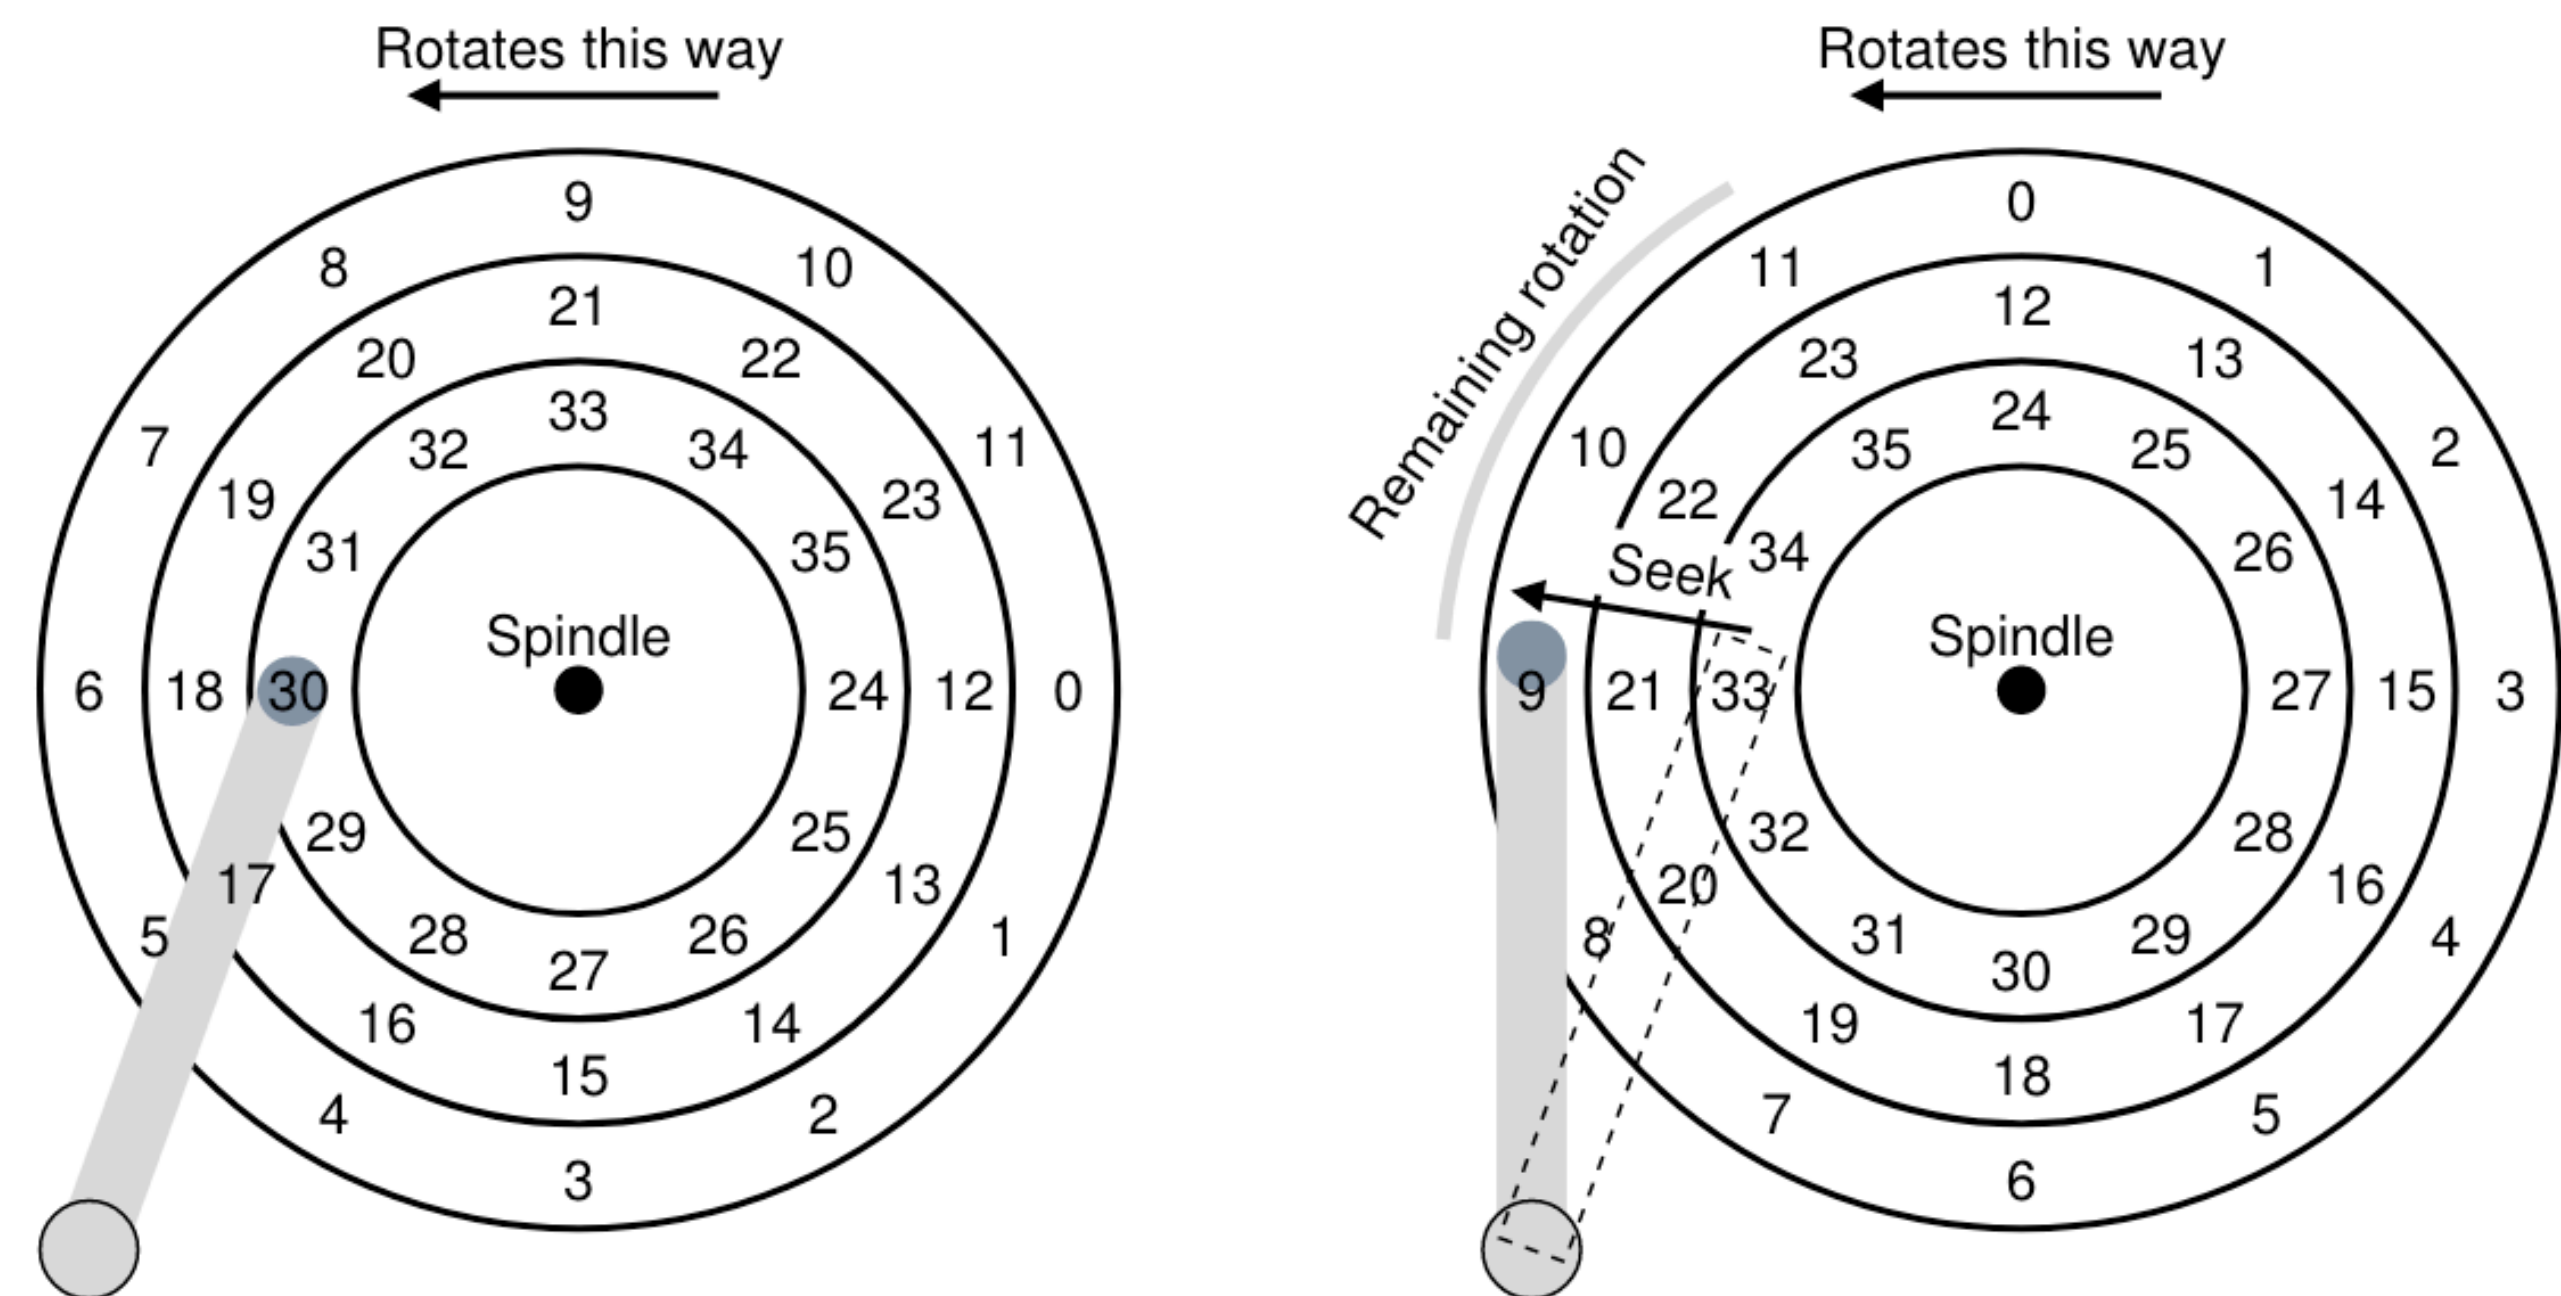


Figure 37.3: Three Tracks Plus A Head (Right: With Seek)

Example of reads

	Cheetah 15K.5
Capacity	300 GB
RPM	15,000
Average Seek	4 ms
Max Transfer	125 MB/s
Platters	4
Cache	16 MB

- Seek delay (4ms)
- Rotation delay: $(60 \times 1000 / 15,000) / 2 = 2\text{ms}$

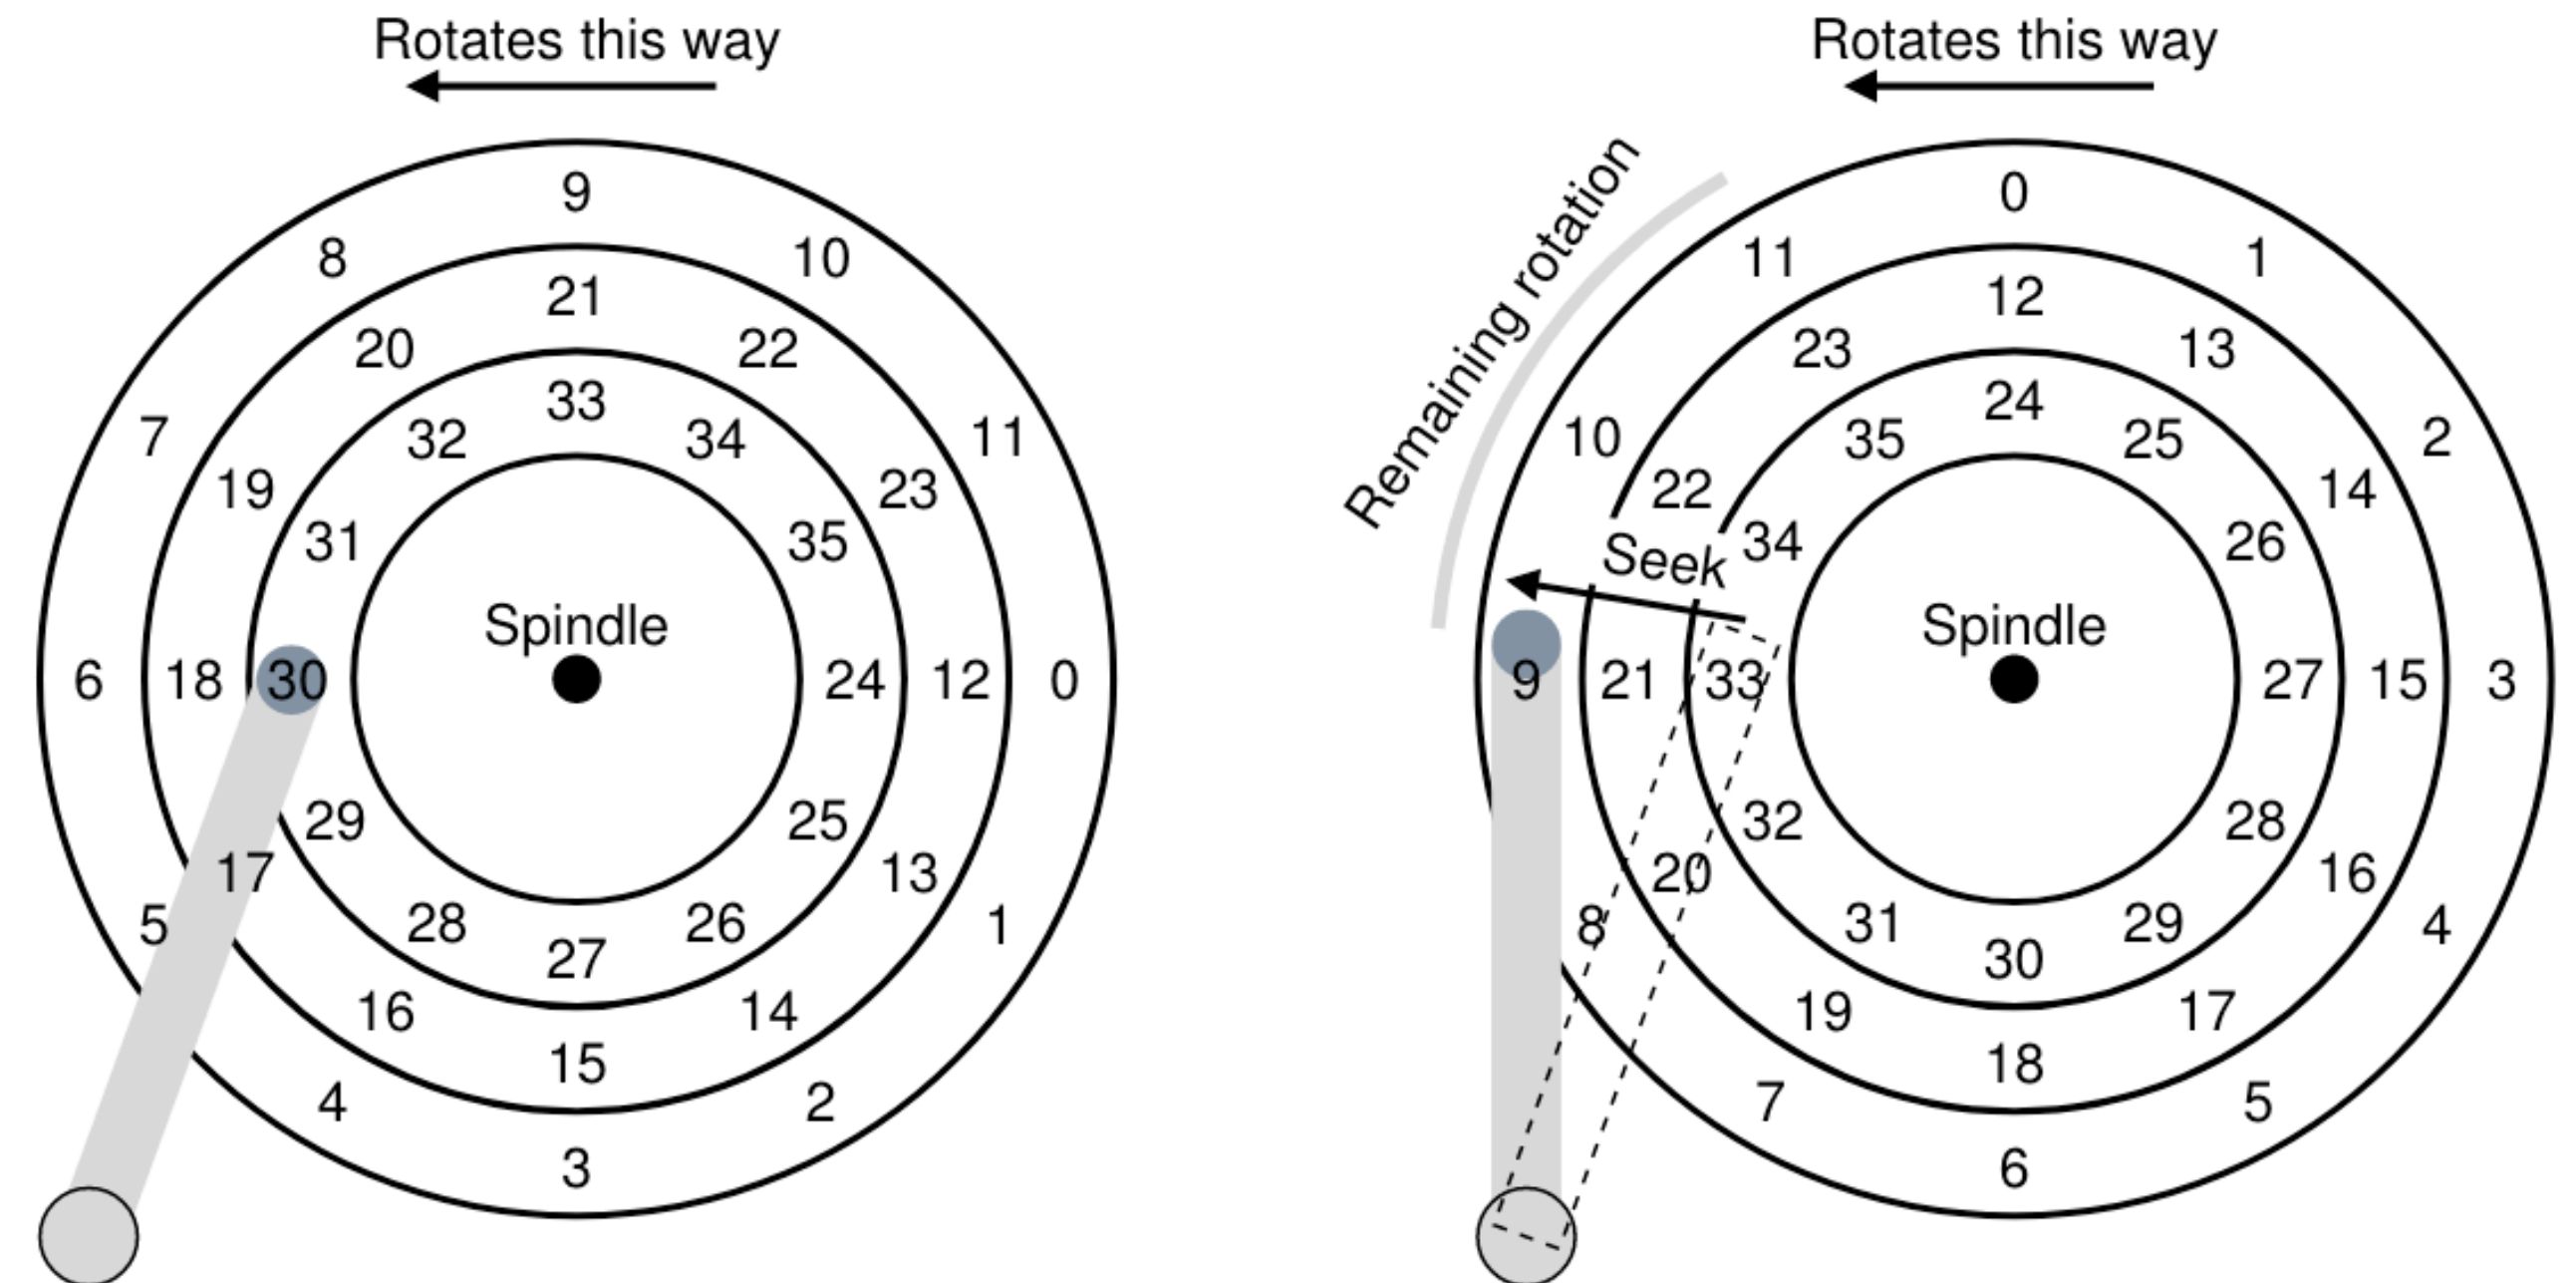


Figure 37.3: Three Tracks Plus A Head (Right: With Seek)

Example of reads

	Cheetah 15K.5
Capacity	300 GB
RPM	15,000
Average Seek	4 ms
Max Transfer	125 MB/s
Platters	4
Cache	16 MB

- Seek delay (4ms)
- Rotation delay: $(60 \times 1000 / 15,000) / 2 = 2\text{ms}$
- Transfer delay

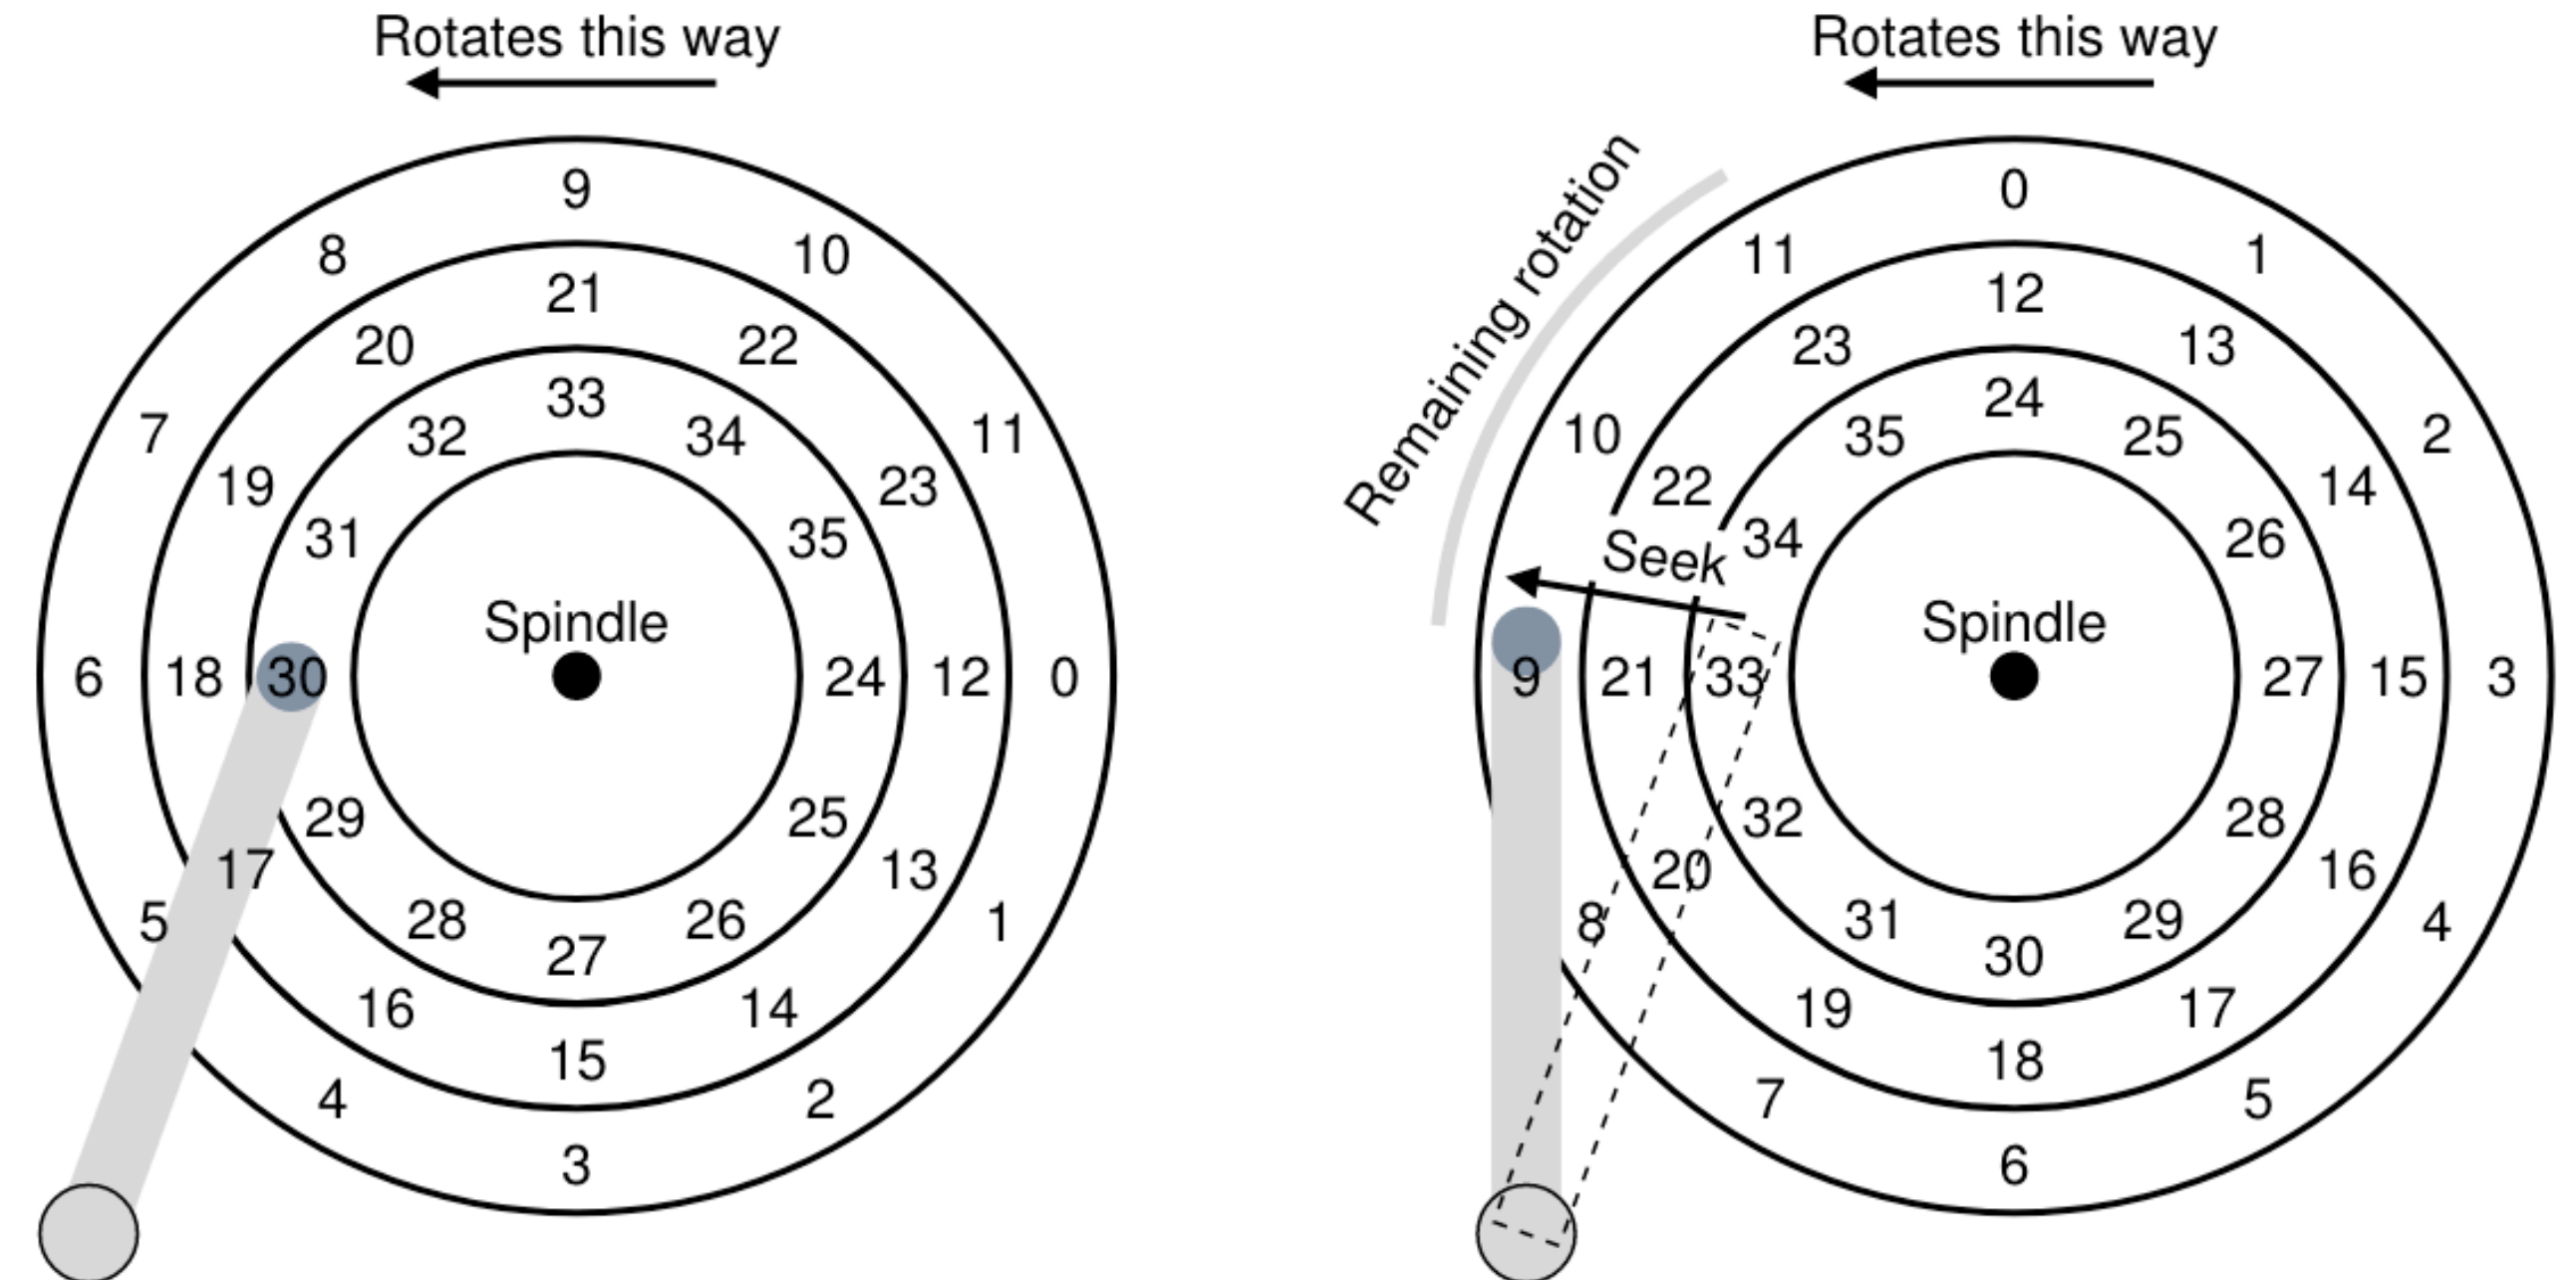


Figure 37.3: Three Tracks Plus A Head (Right: With Seek)

Example of reads

	Cheetah 15K.5
Capacity	300 GB
RPM	15,000
Average Seek	4 ms
Max Transfer	125 MB/s
Platters	4
Cache	16 MB

- Seek delay (4ms)
- Rotation delay: $(60 \times 1000 / 15,000) / 2 = 2\text{ms}$
- Transfer delay
 - 125MBps = 125 bytes per us.

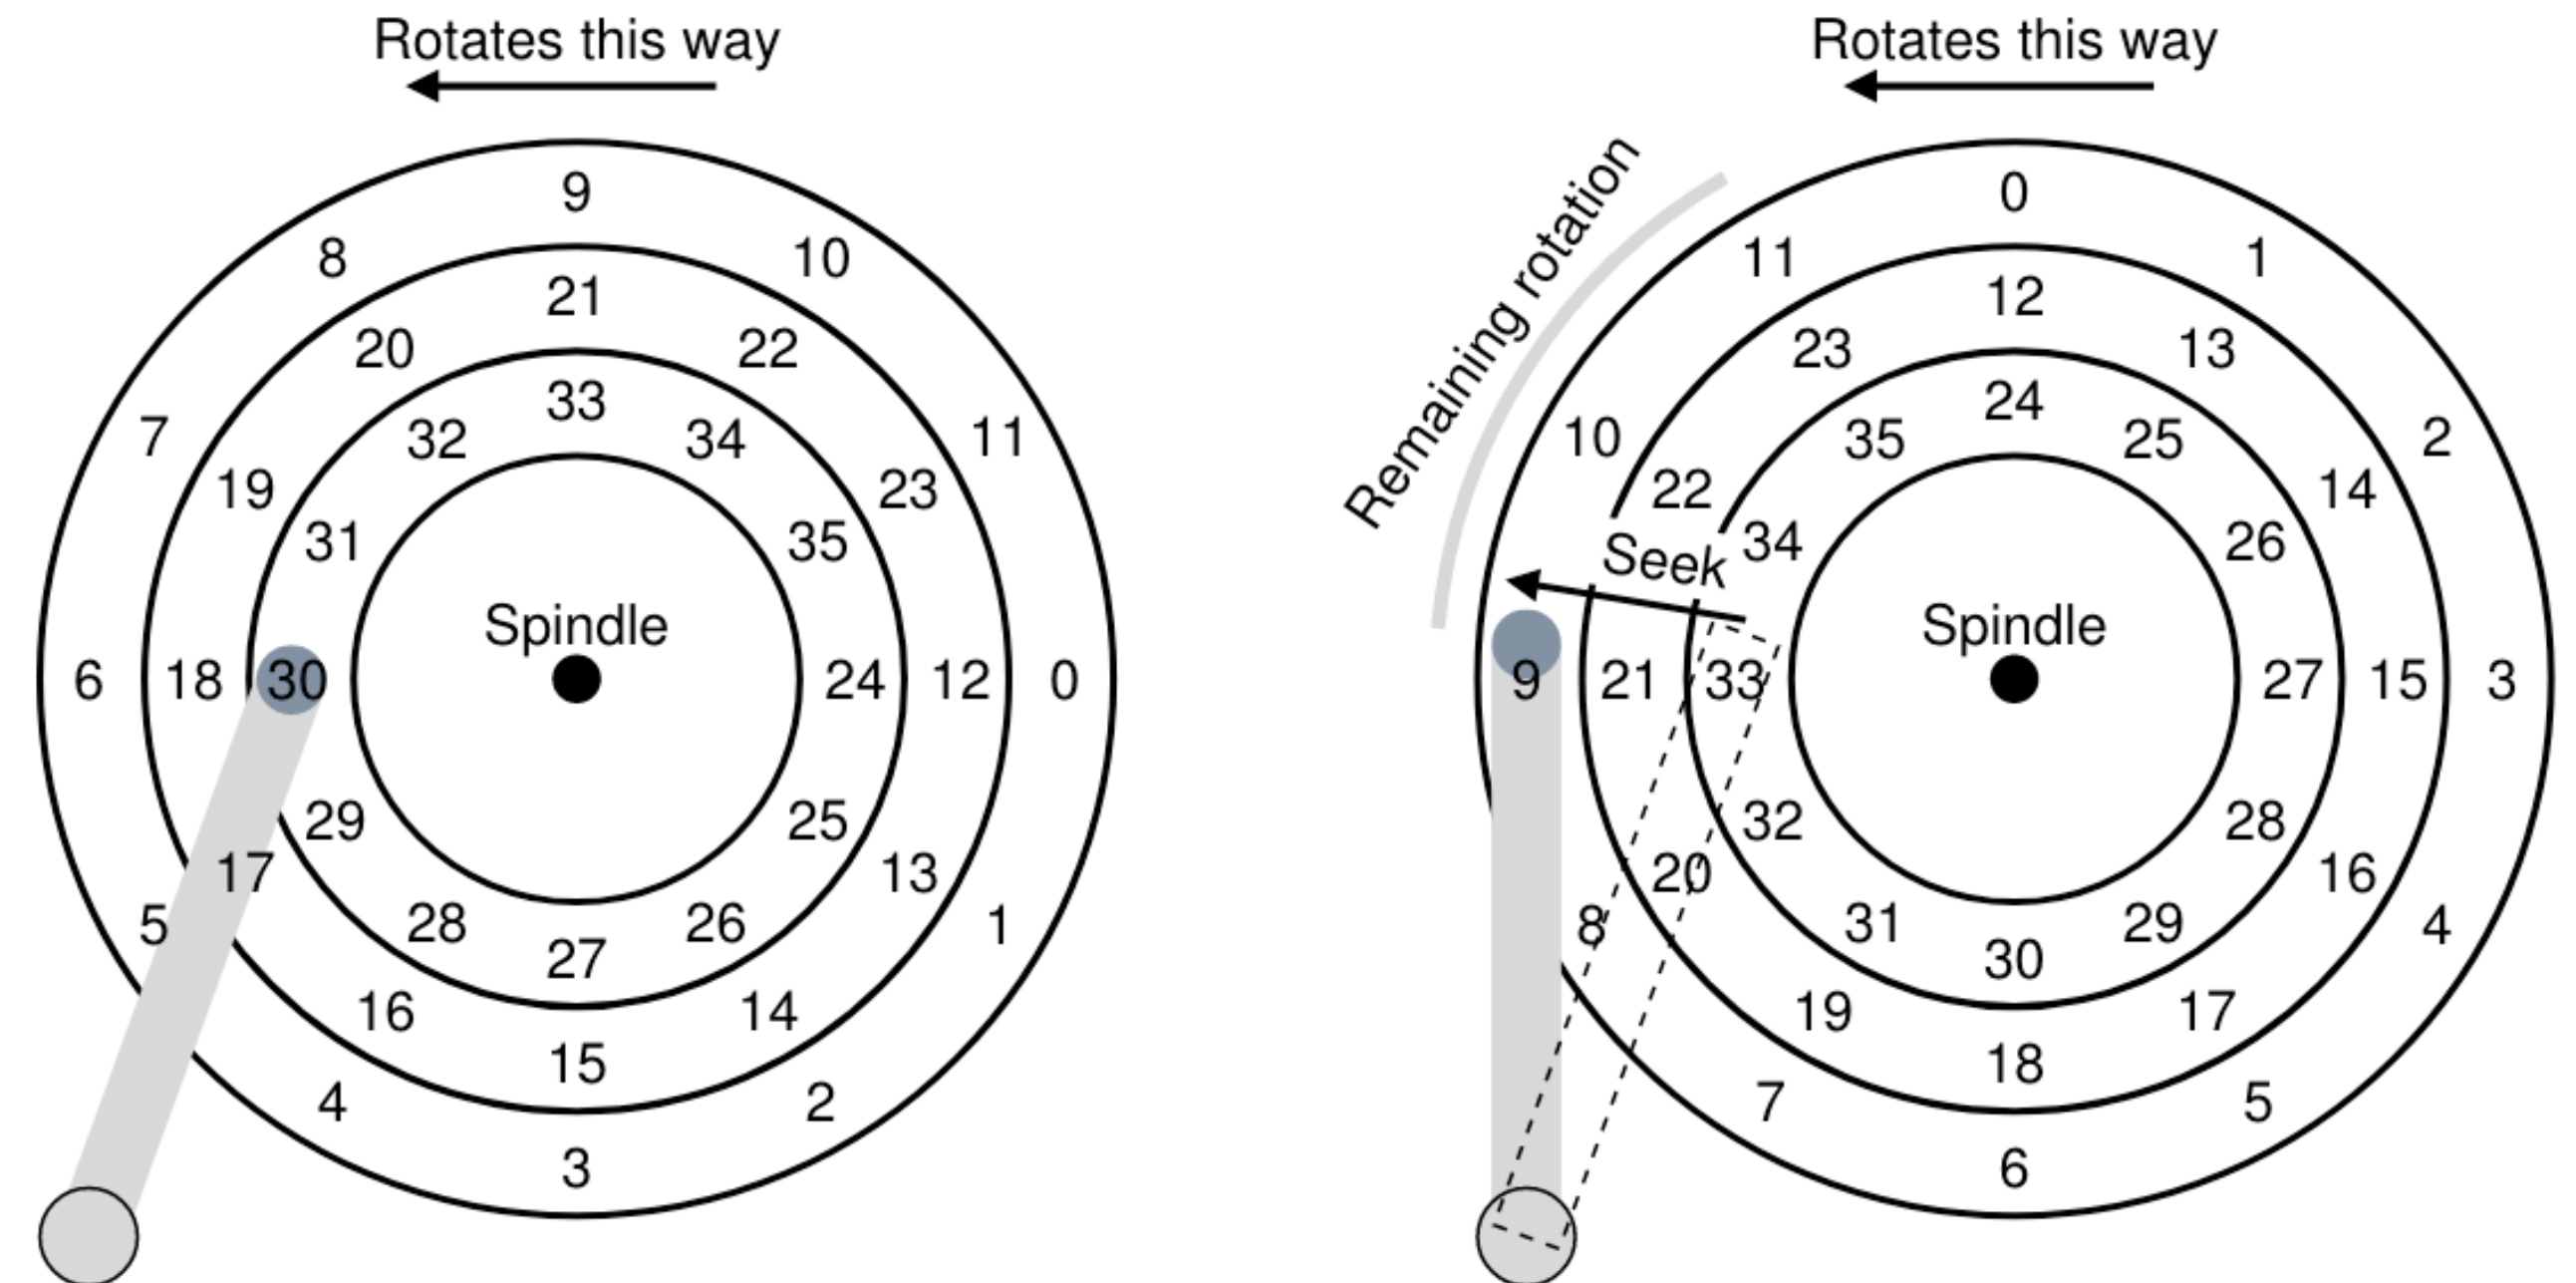


Figure 37.3: Three Tracks Plus A Head (Right: With Seek)

Example of reads

	Cheetah 15K.5
Capacity	300 GB
RPM	15,000
Average Seek	4 ms
Max Transfer	125 MB/s
Platters	4
Cache	16 MB

- Seek delay (4ms)
- Rotation delay: $(60 \times 1000 / 15,000) / 2 = 2\text{ms}$
- Transfer delay
 - 125MBps = 125 bytes per us.
 - Time take to read 4KB: $4096 / 125 \sim 30\text{us}$

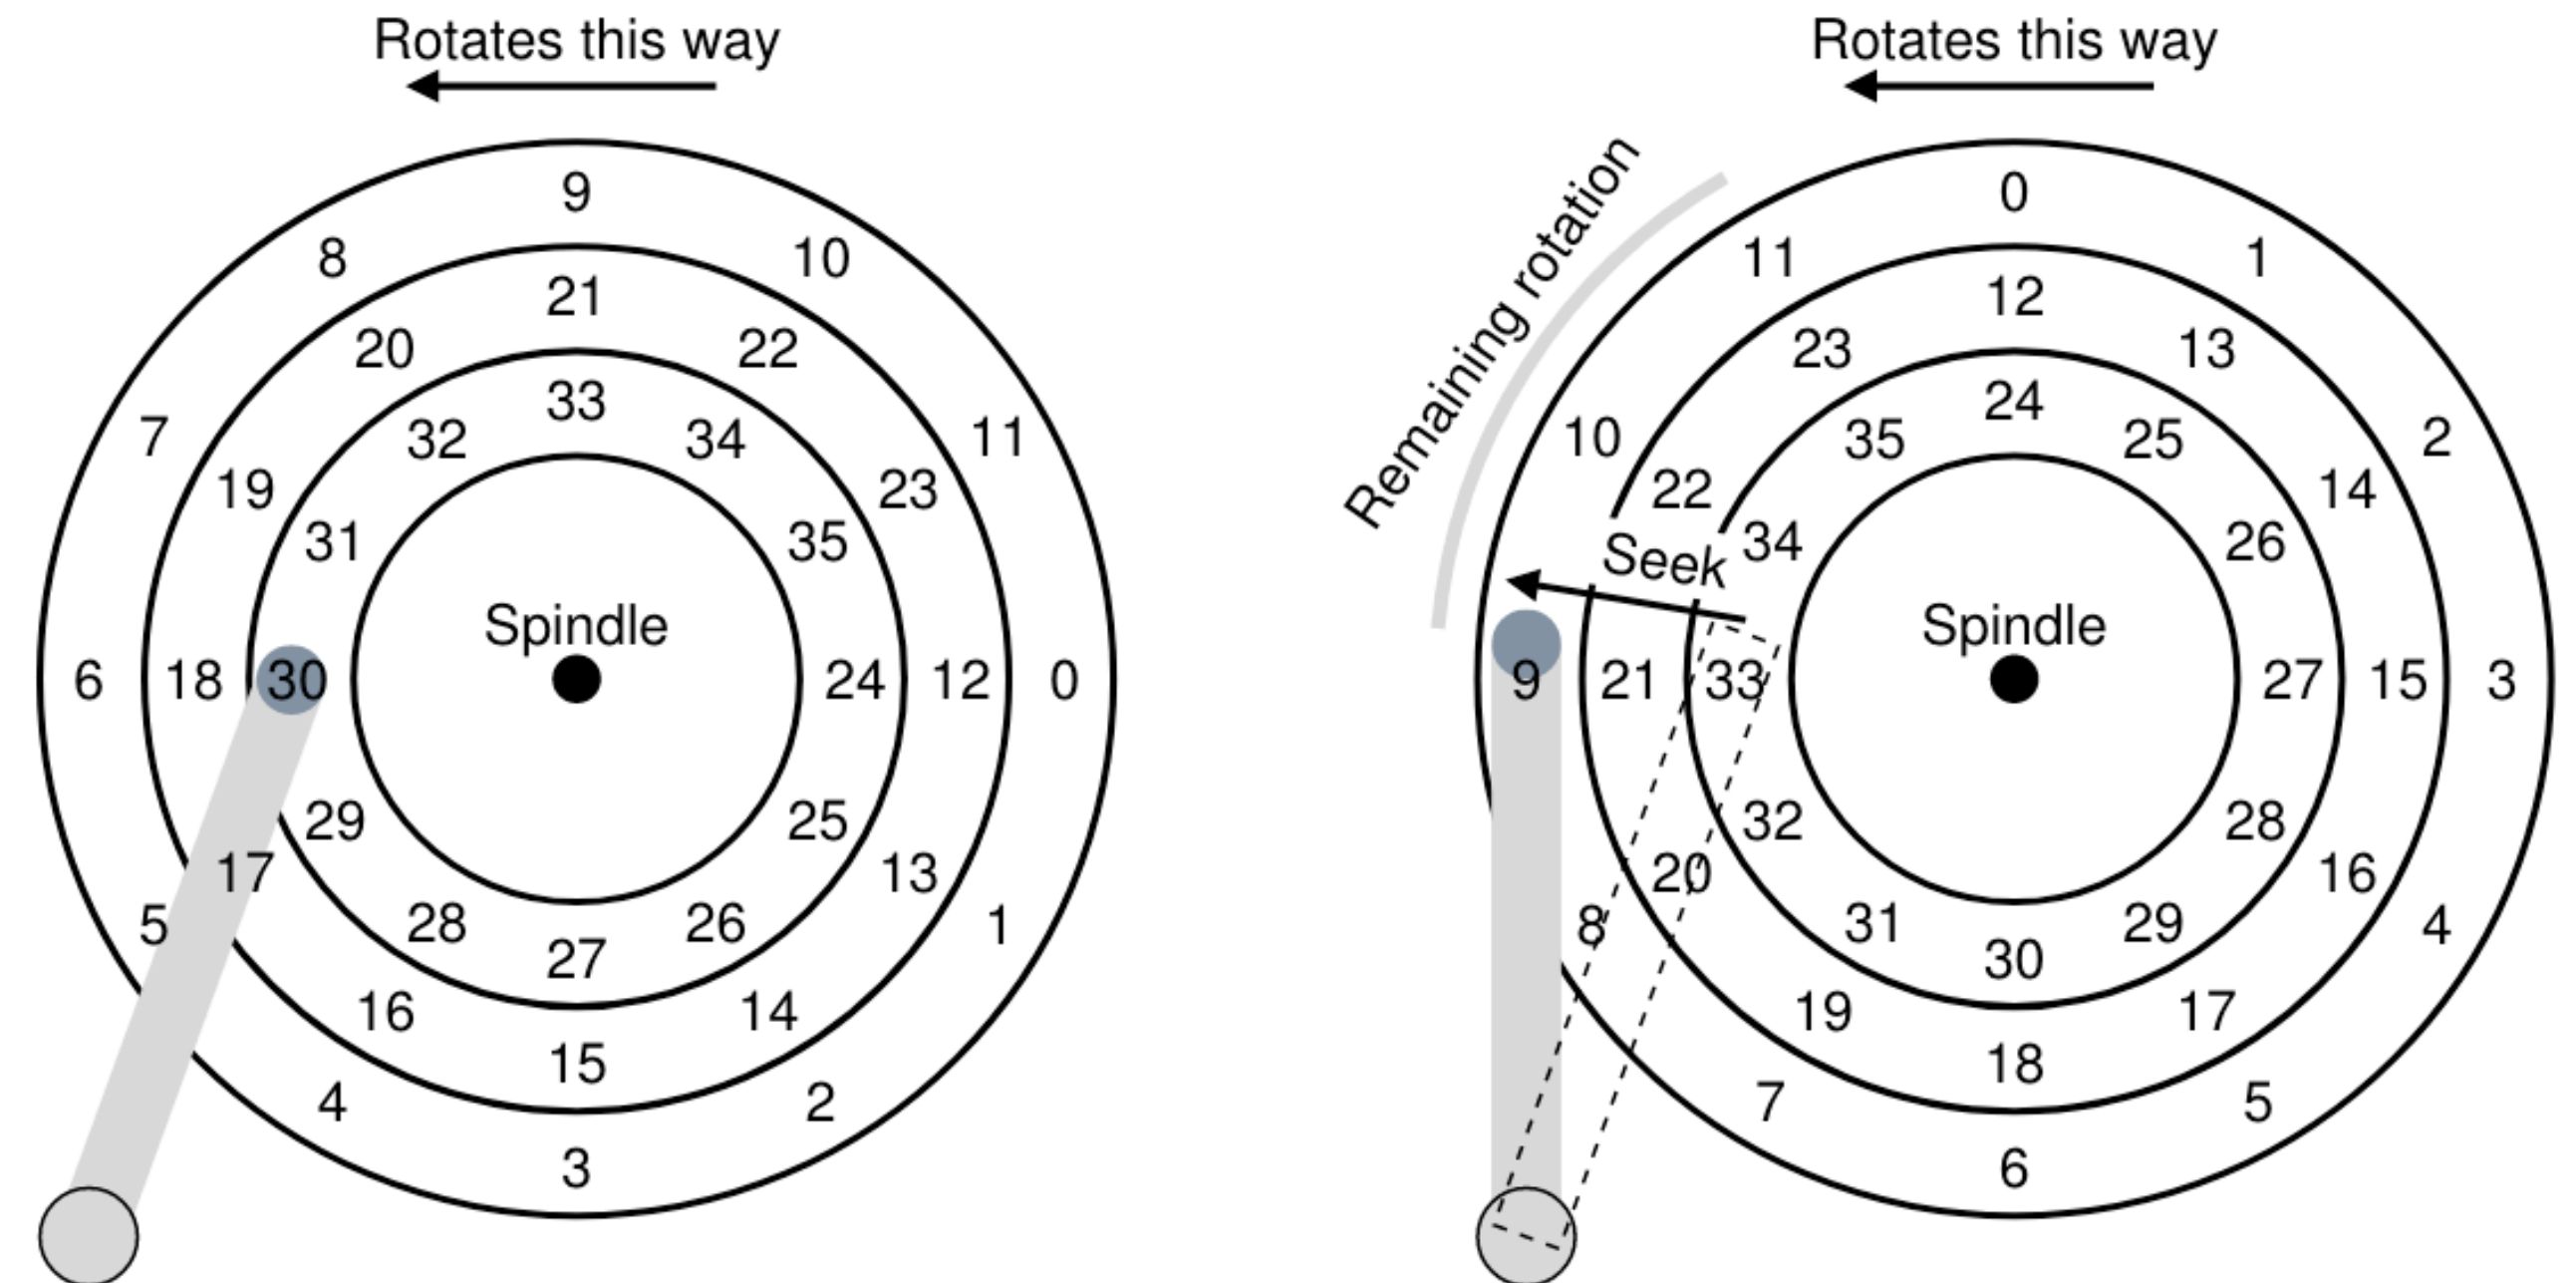


Figure 37.3: Three Tracks Plus A Head (Right: With Seek)

Example of reads

	Cheetah 15K.5
Capacity	300 GB
RPM	15,000
Average Seek	4 ms
Max Transfer	125 MB/s
Platters	4
Cache	16 MB

- Seek delay (4ms)
- Rotation delay: $(60 \times 1000 / 15,000) / 2 = 2\text{ms}$
- Transfer delay
 - 125MBps = 125 bytes per us.
 - Time take to read 4KB: $4096 / 125 \sim 30\text{us}$
- 4KB random read: 4ms (seek) + 2ms (rotation) + 30us (transfer) $\sim 6\text{ms}$

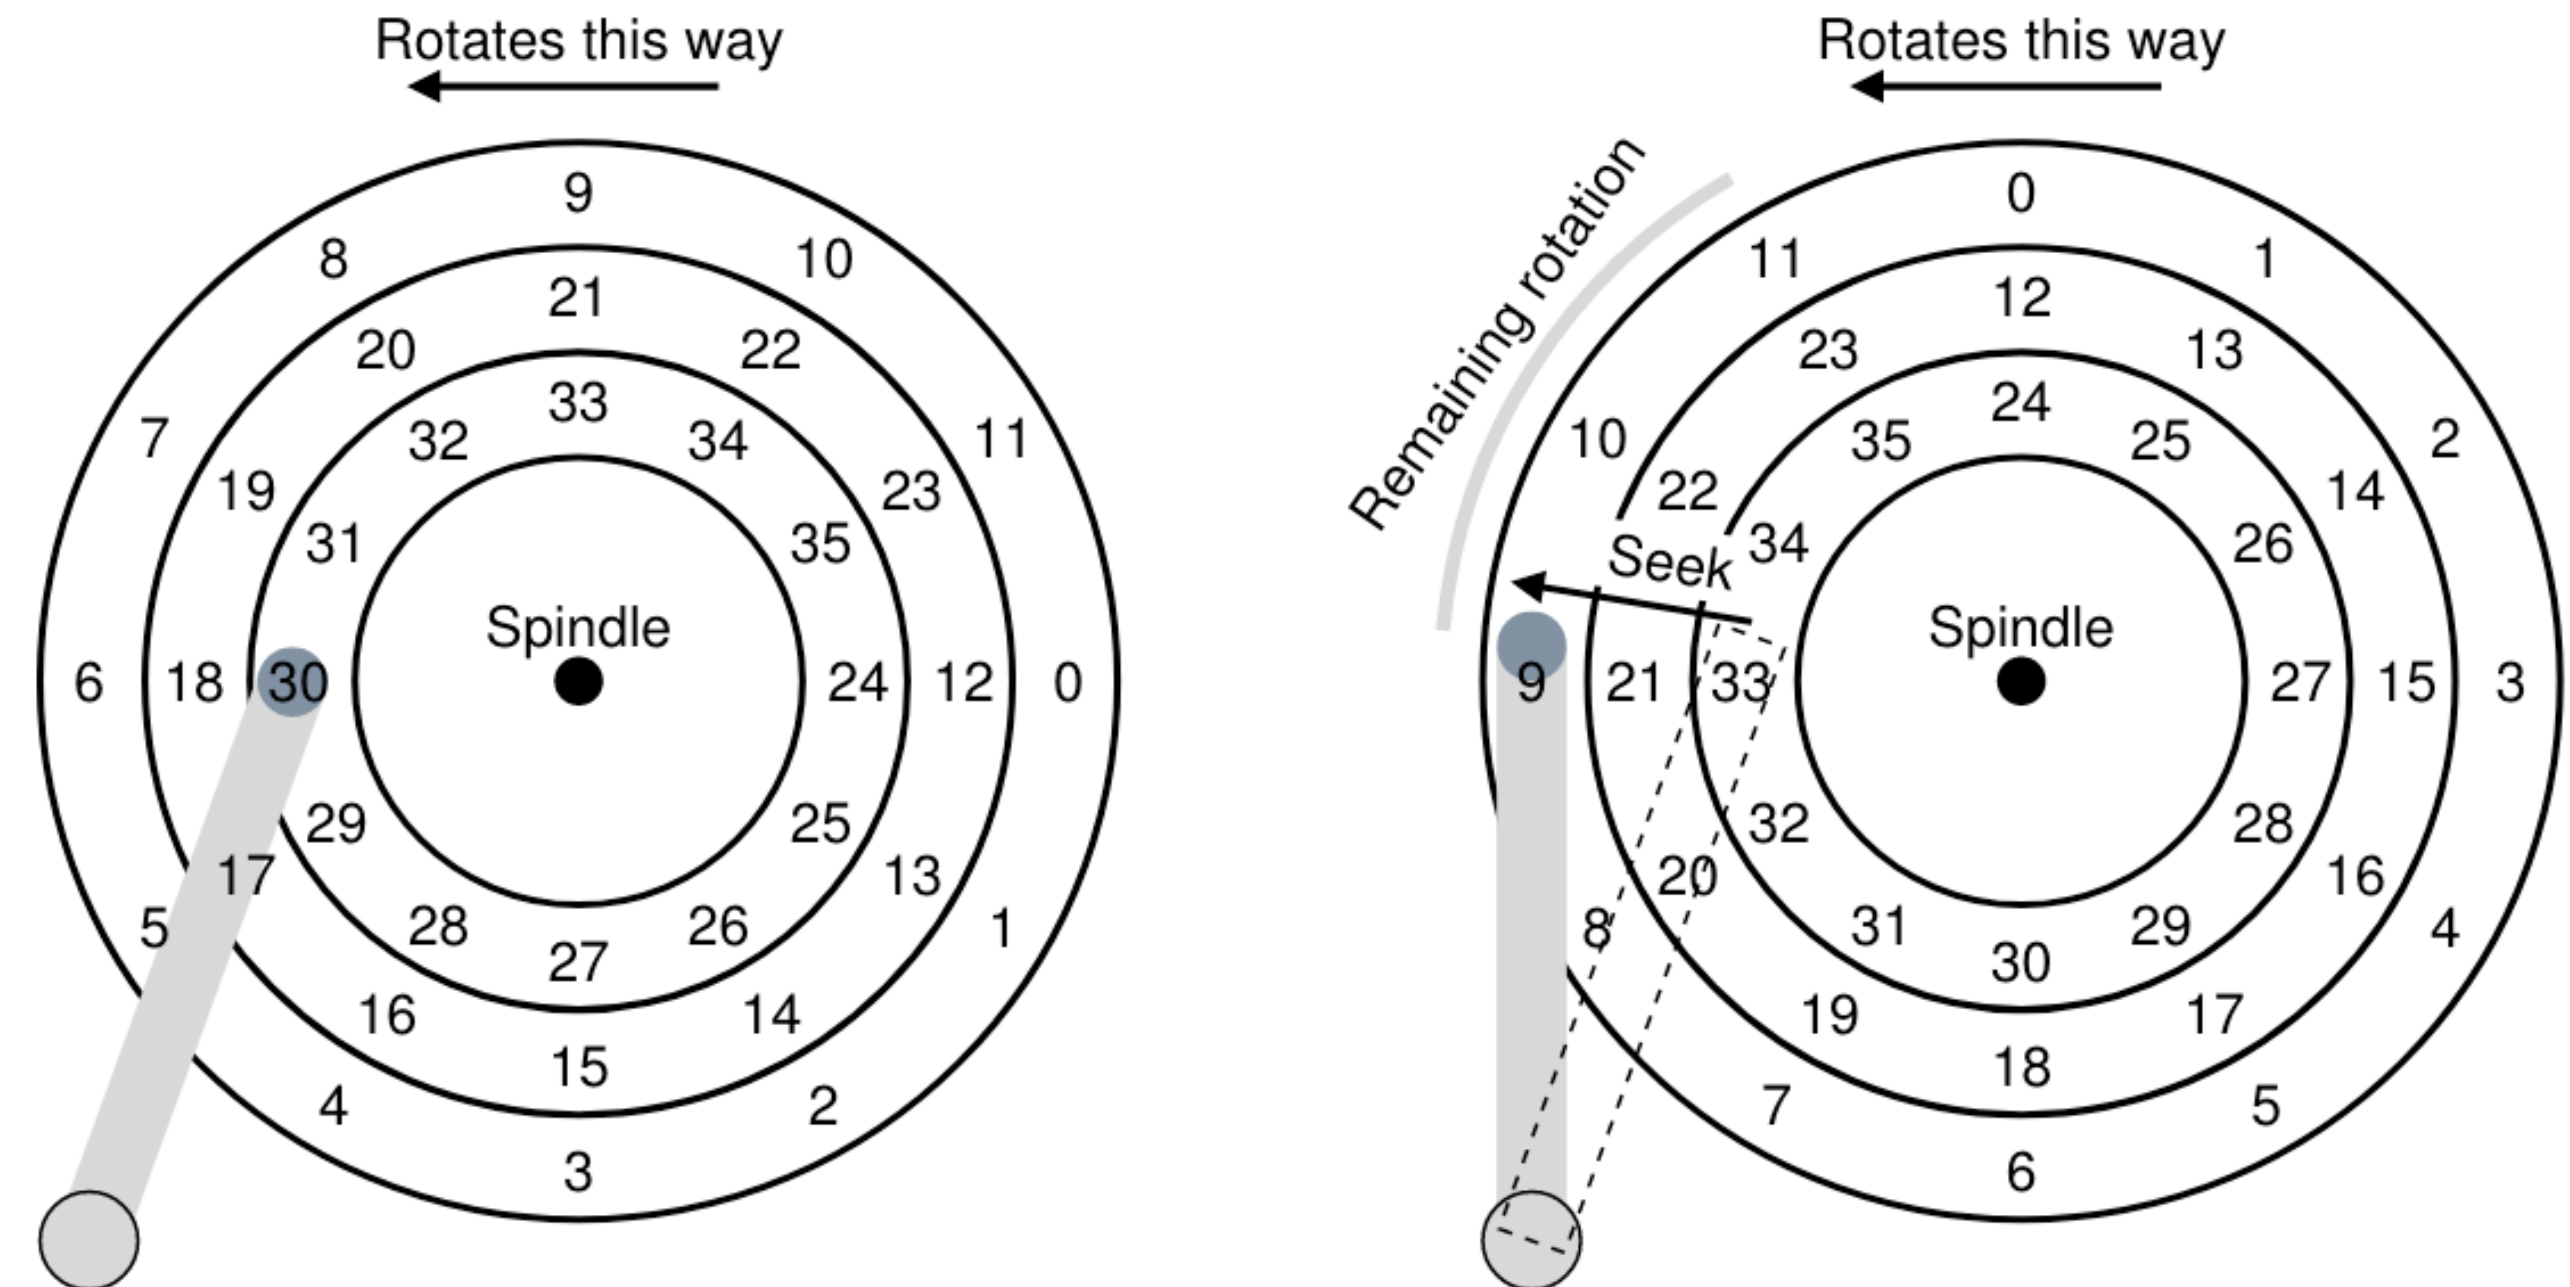


Figure 37.3: Three Tracks Plus A Head (Right: With Seek)

Random rws are ~100x slower than sequential rws!

- Random read: 4ms (seek time) + 2ms (rotation time) + 30us (transfer time) ~ 6ms
- Sequential read: 30us (transfer time)

Disk scheduling problem

- `python3 disk.py -a 10,15,32,11,33,16 -G`
- Given a sequence of requests, reorder requests to service them quicker

